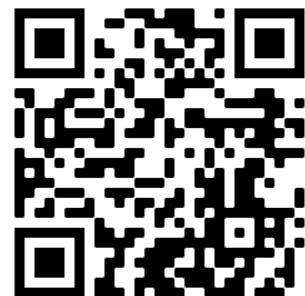


Artificial Intelligence and Intelligent Agents; Problem Solving

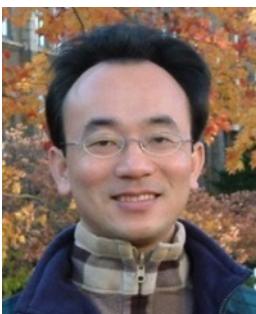
1141AI02

MBA, IM, NTPU (M5276) (Fall 2025)

Tue 2, 3, 4 (9:10-12:00) (B3F17)



[https://meet.google.com/
paj-zhhj-mya](https://meet.google.com/paj-zhhj-mya)



Min-Yuh Day, Ph.D,
Professor and Director

Institute of Information Management, National Taipei University

<https://web.ntpu.edu.tw/~myday>



Syllabus

Week Date Subject/Topics

1 2025/09/09 Introduction to Artificial Intelligence

2 2025/09/16 Artificial Intelligence and Intelligent Agents;
Problem Solving

3 2025/09/23 Knowledge, Reasoning and Knowledge Representation;
Uncertain Knowledge and Reasoning

4 2025/09/30 Case Study on Artificial Intelligence I

5 2025/10/07 Machine Learning: Supervised and Unsupervised Learning;
The Theory of Learning and Ensemble Learning

Syllabus

Week Date Subject/Topics

6 2025/10/14 NVIDIA Fundamentals of Deep Learning I:
Deep Learning; Neural Networks

7 2025/10/21 NVIDIA Fundamentals of Deep Learning II:
Convolutional Neural Networks;
Data Augmentation and Deployment

8 2025/10/28 Self-Learning

9 2025/11/04 Midterm Project Report

10 2025/11/11 NVIDIA Fundamentals of Deep Learning III:
Pre-trained Models; Natural Language Processing

Syllabus

Week Date Subject/Topics

11 2025/11/18 Case Study on Artificial Intelligence II

12 2025/11/25 Computer Vision and Robotics

13 2025/12/02 Generative AI, Agentic AI, and Physical AI

14 2025/12/09 Philosophy and Ethics of AI and the Future of AI

15 2025/12/16 Final Project Report I

16 2025/12/23 Final Project Report II

Artificial Intelligence

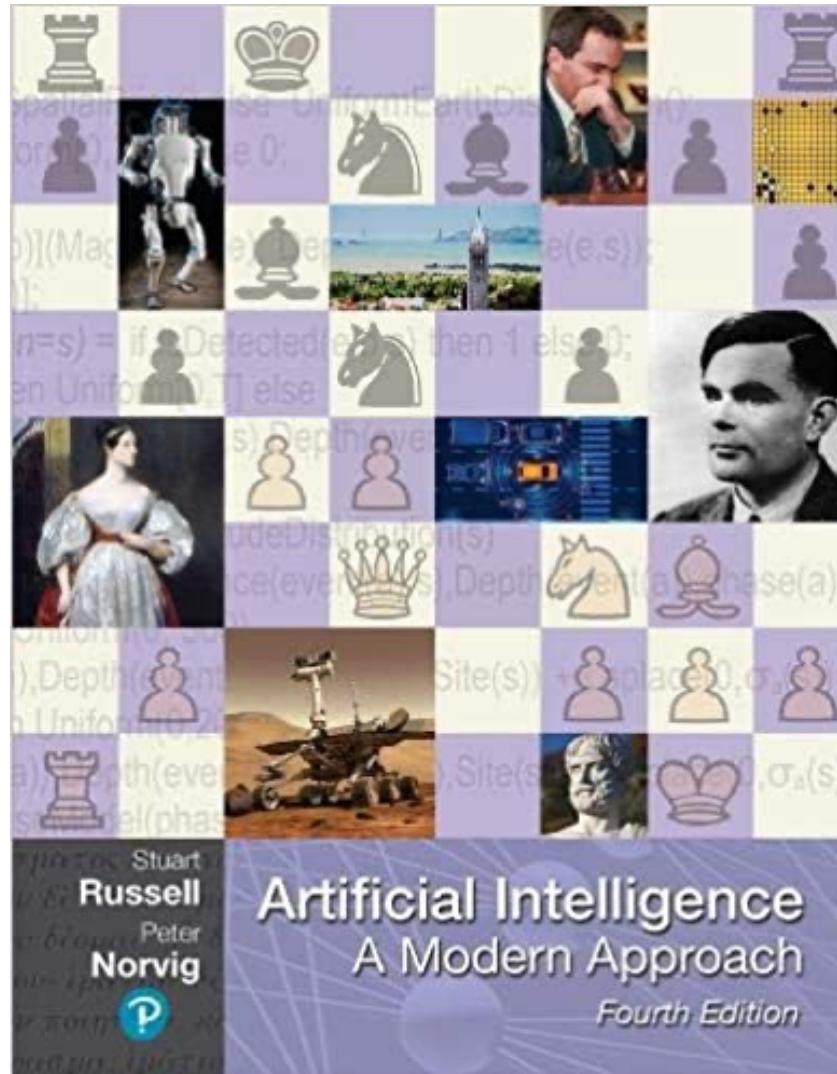
Intelligent Agents

Problem Solving

Outline

- Artificial Intelligence
- Intelligent Agents
- Problem Solving

Stuart Russell and Peter Norvig (2020),
Artificial Intelligence: A Modern Approach,
4th Edition, Pearson



Source: Stuart Russell and Peter Norvig (2020), Artificial Intelligence: A Modern Approach, 4th Edition, Pearson

<https://www.amazon.com/Artificial-Intelligence-A-Modern-Approach/dp/0134610997/>

Artificial Intelligence: A Modern Approach

- 1. Artificial Intelligence**
- 2. Problem Solving**
- 3. Knowledge and Reasoning**
- 4. Uncertain Knowledge and Reasoning**
- 5. Machine Learning**
- 6. Communicating, Perceiving, and Acting**
- 7. Philosophy and Ethics of AI**

Artificial Intelligence: Intelligent Agents

Artificial Intelligence:

2. Problem Solving

- Solving Problems by Searching**
- Search in Complex Environments**
- Adversarial Search and Games**
- Constraint Satisfaction Problems**

Artificial Intelligence:

3. Knowledge and Reasoning

- **Logical Agents**
- **First-Order Logic**
- **Inference in First-Order Logic**
- **Knowledge Representation**
- **Automated Planning**

Artificial Intelligence:

4. Uncertain Knowledge and Reasoning

- Quantifying Uncertainty
- Probabilistic Reasoning
- Probabilistic Reasoning over Time
- Probabilistic Programming
- Making Simple Decisions
- Making Complex Decisions
- Multiagent Decision Making

Artificial Intelligence:

5. Machine Learning

- **Learning from Examples**
- **Learning Probabilistic Models**
- **Deep Learning**
- **Reinforcement Learning**

Artificial Intelligence:

6. Communicating, Perceiving, and Acting

- **Natural Language Processing**
- **Deep Learning for Natural Language Processing**
- **Computer Vision**
- **Robotics**

Artificial Intelligence: Philosophy and Ethics of AI The Future of AI

NVIDIA Developer Program

<https://developer.nvidia.com/join-nvidia-developer-program>

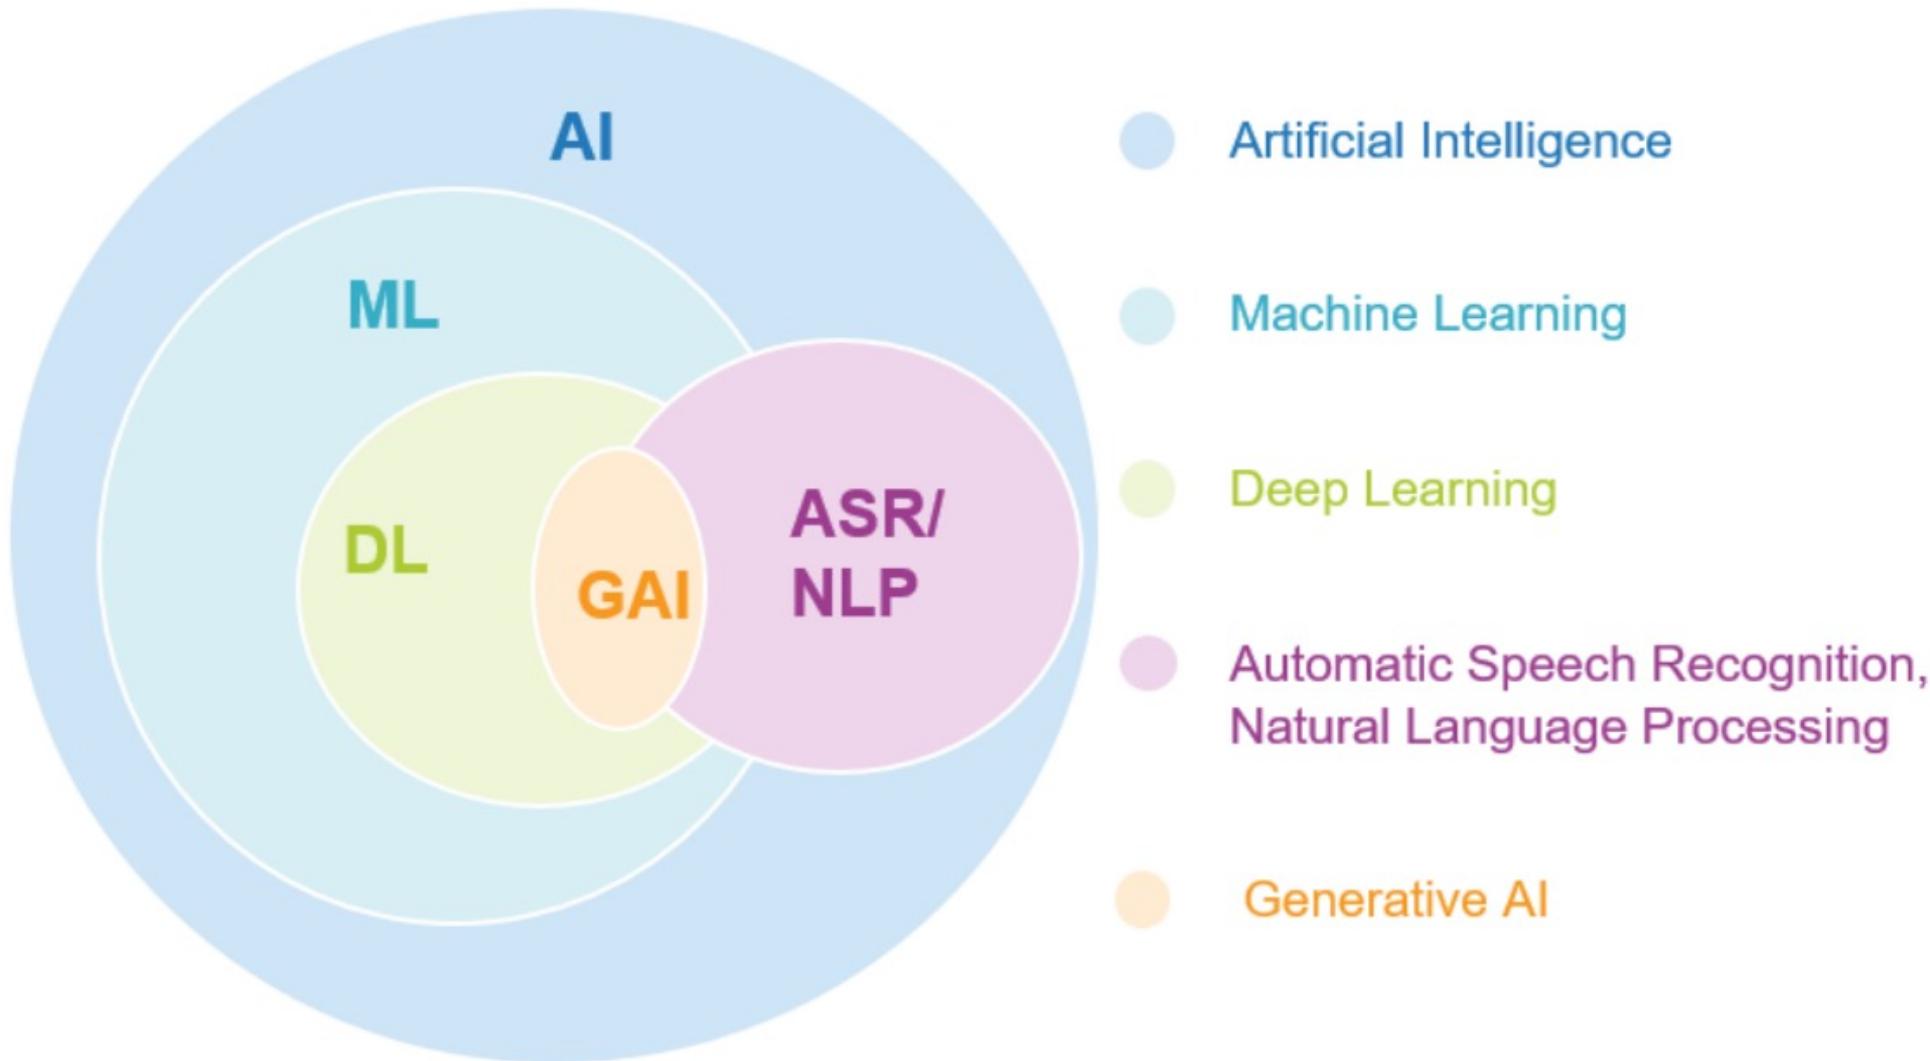
NVIDIA

Deep Learning Institute (DLI)

<https://learn.nvidia.com/>

Artificial Intelligence (AI)

AI, ML, DL, Generative AI



Generative AI, Agentic AI, Physical AI



Physical AI

Self-driving cars
General robotics



Agentic AI

Coding assistants
Customer service
Patient care



Generative AI

Digital marketing
Content creation



Perception AI

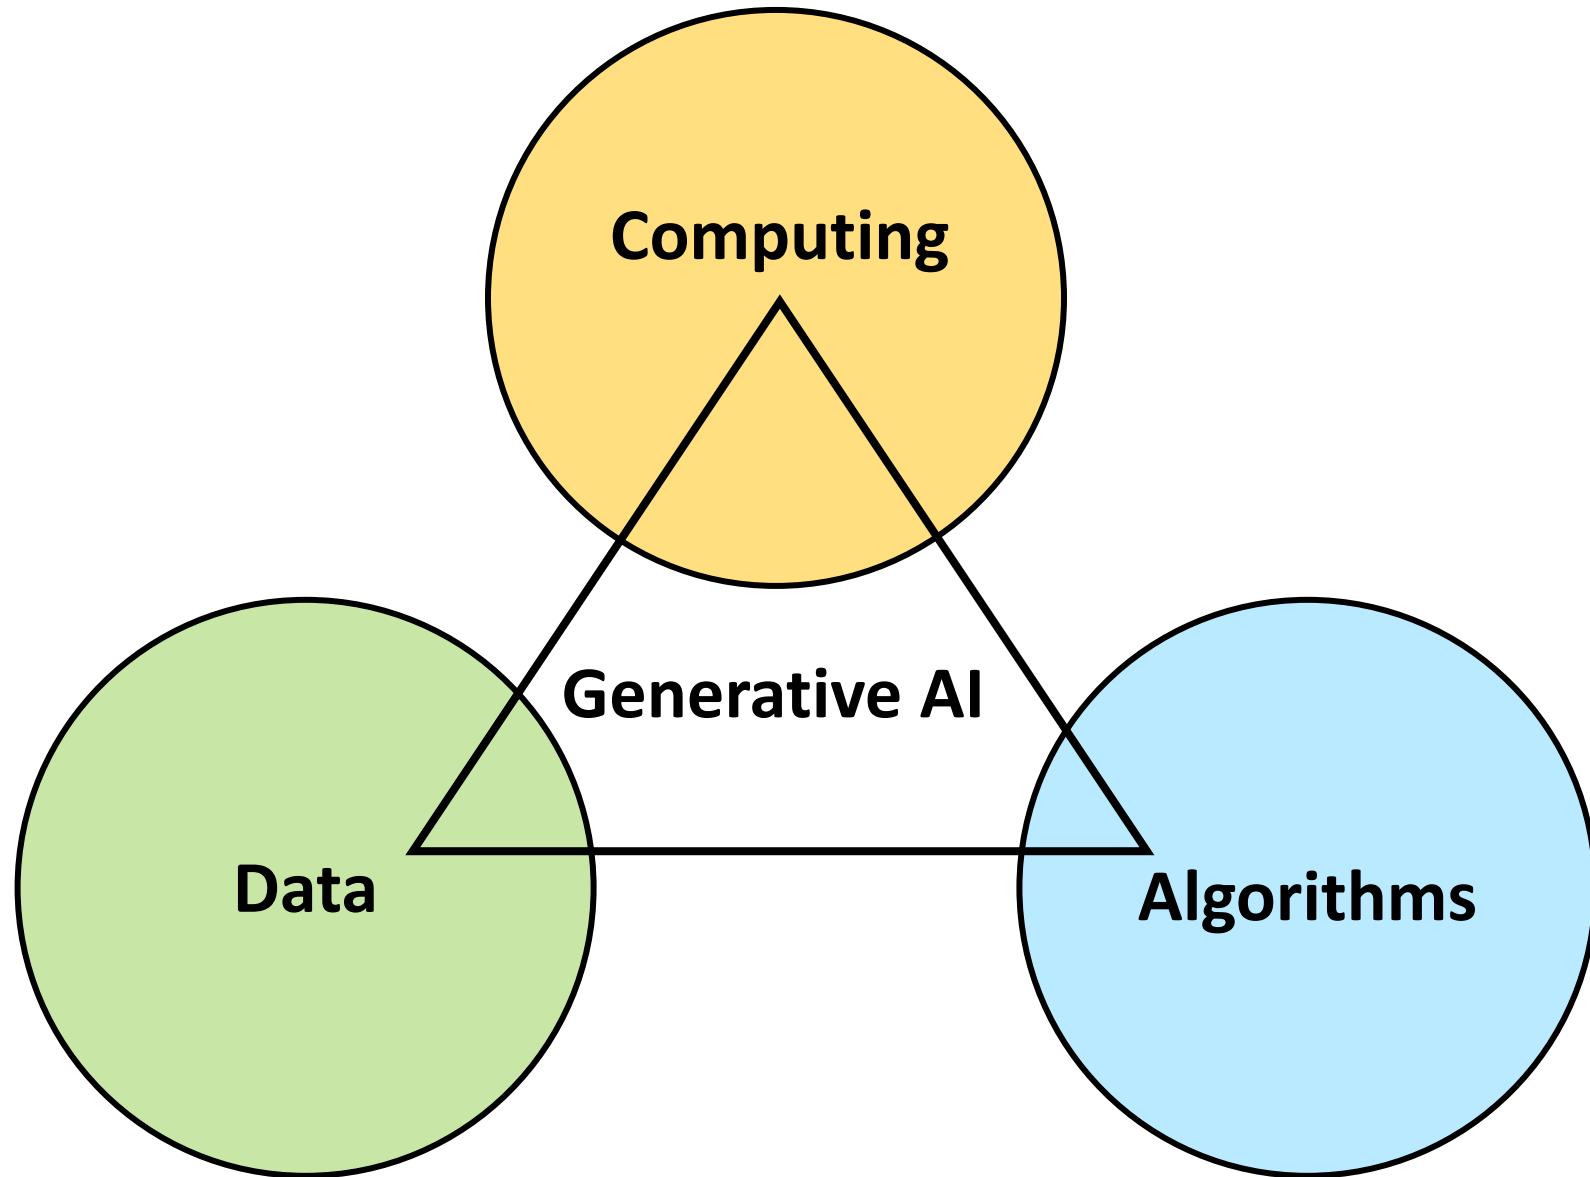
Speech recognition
Deep recommender systems
Medical imaging

2012 AlexNet

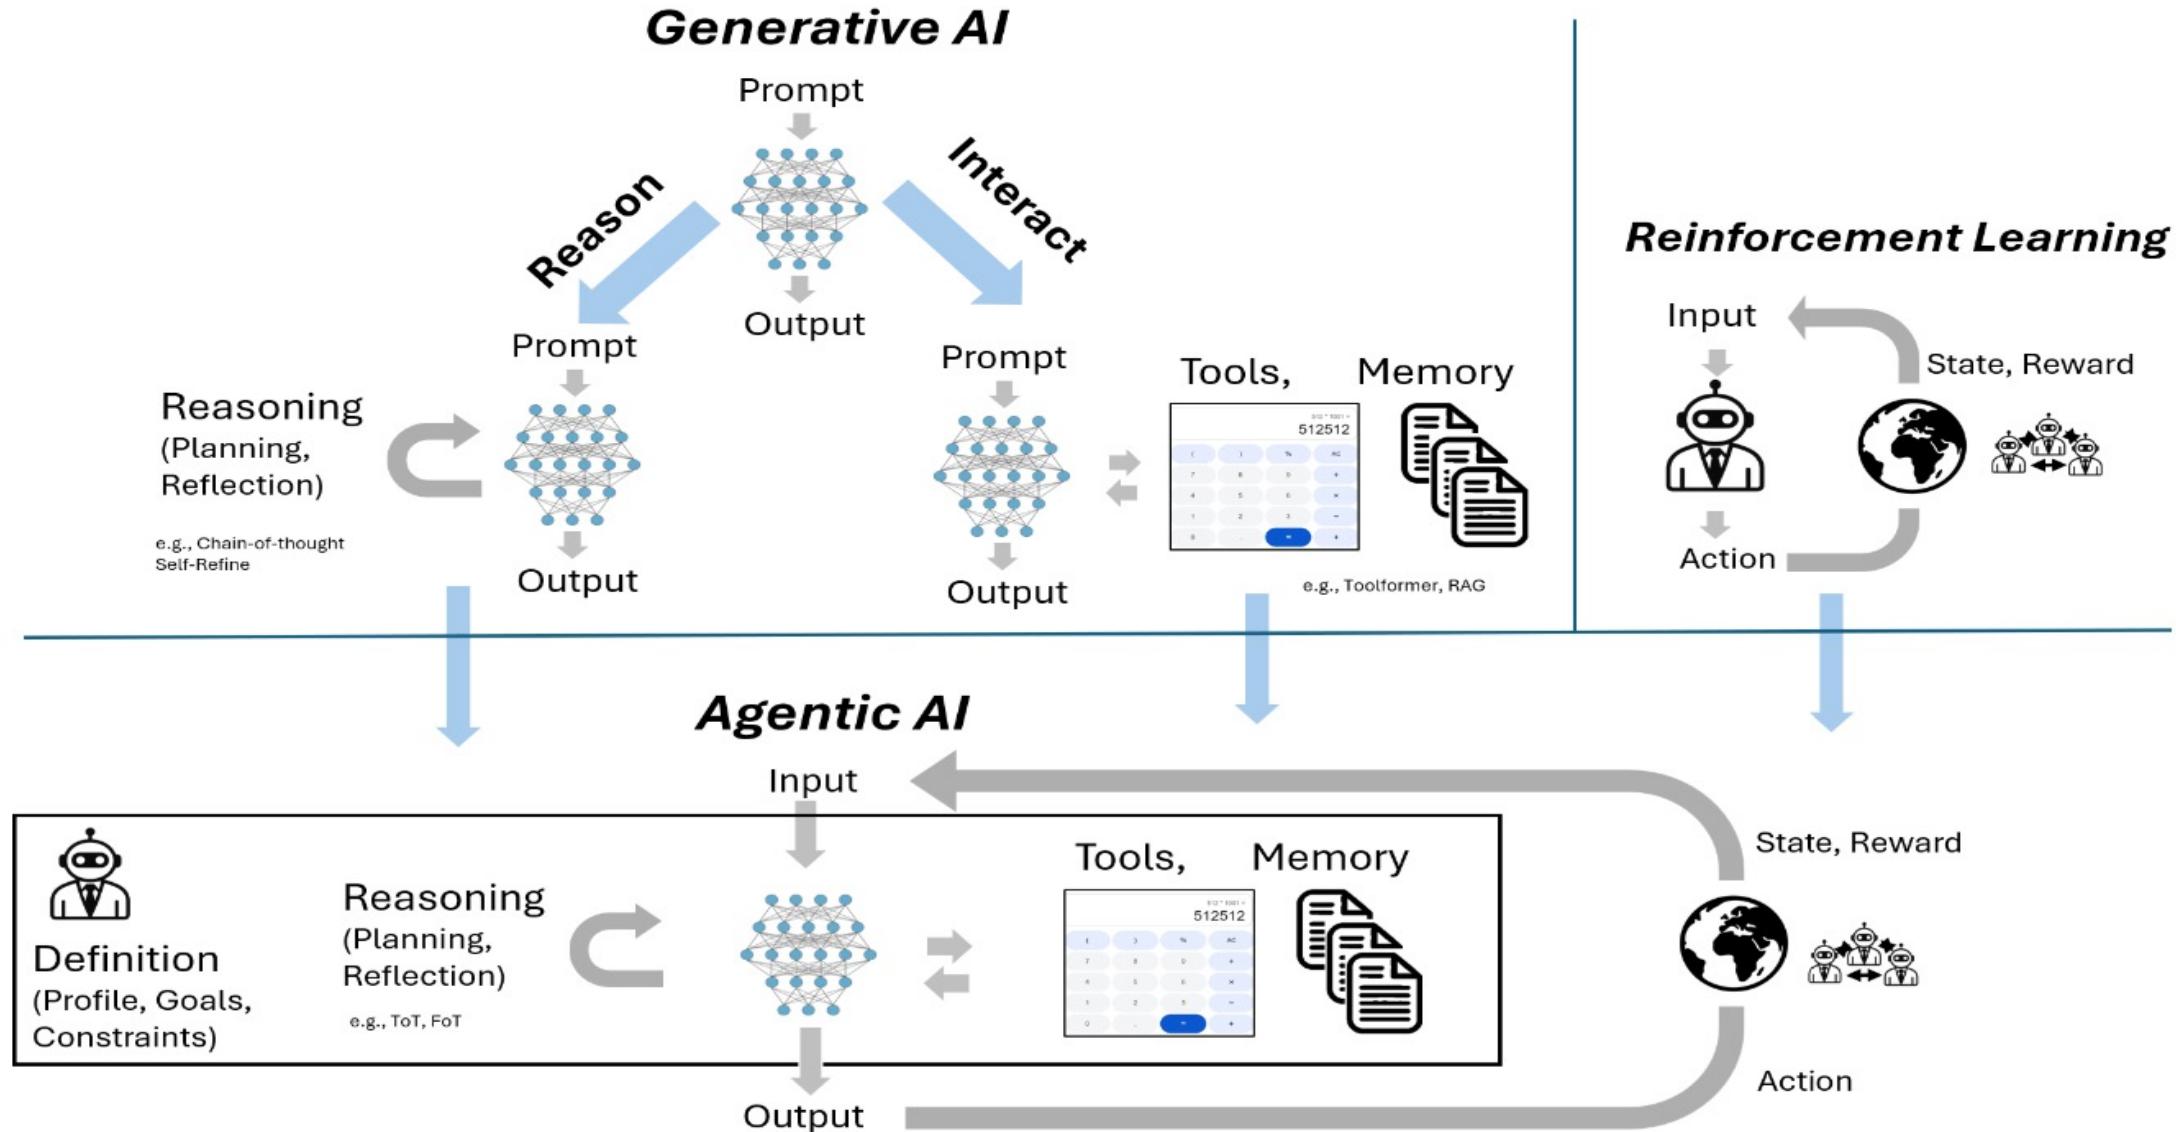


Deep learning breakthrough

Generative AI



From Generative AI to Agentic AI

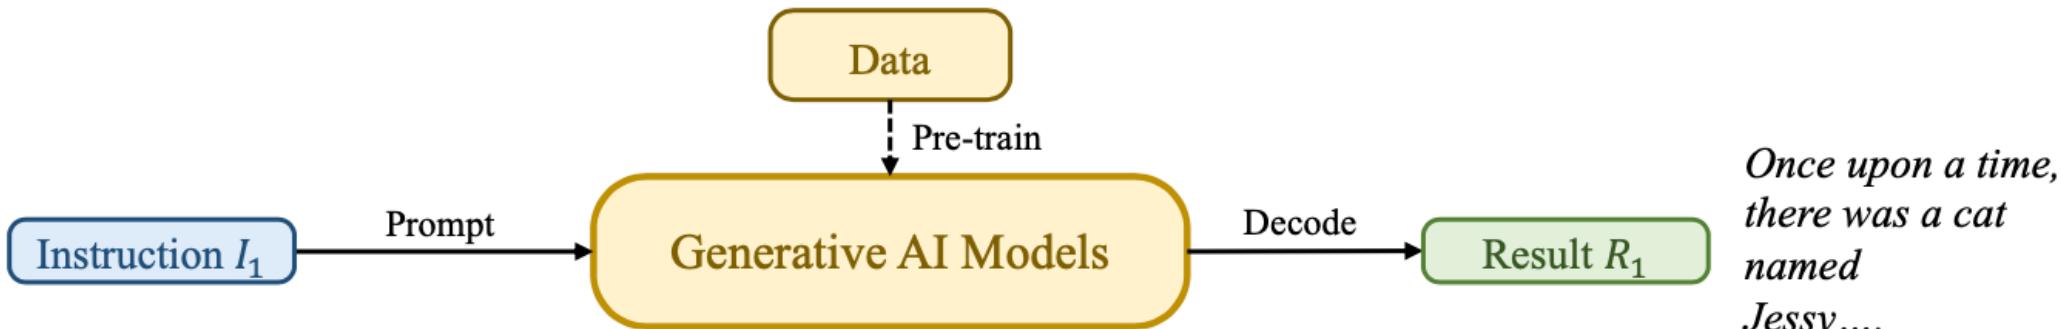


Generative AI (Gen AI)

AI Generated Content (AIGC)

Unimodal

Please write a story about a cat.



Multimodal



Describe this picture.

Instruction I_2

Prompt

Draw a picture of a cat.

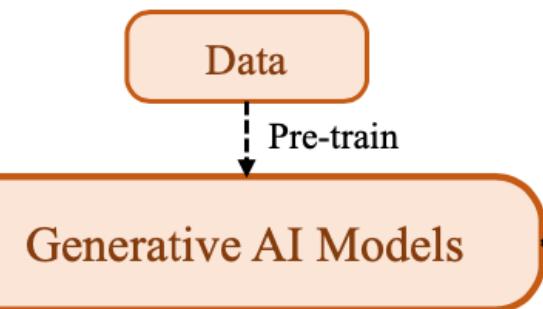
Instruction I_3

Prompt

Write a song about a cat.

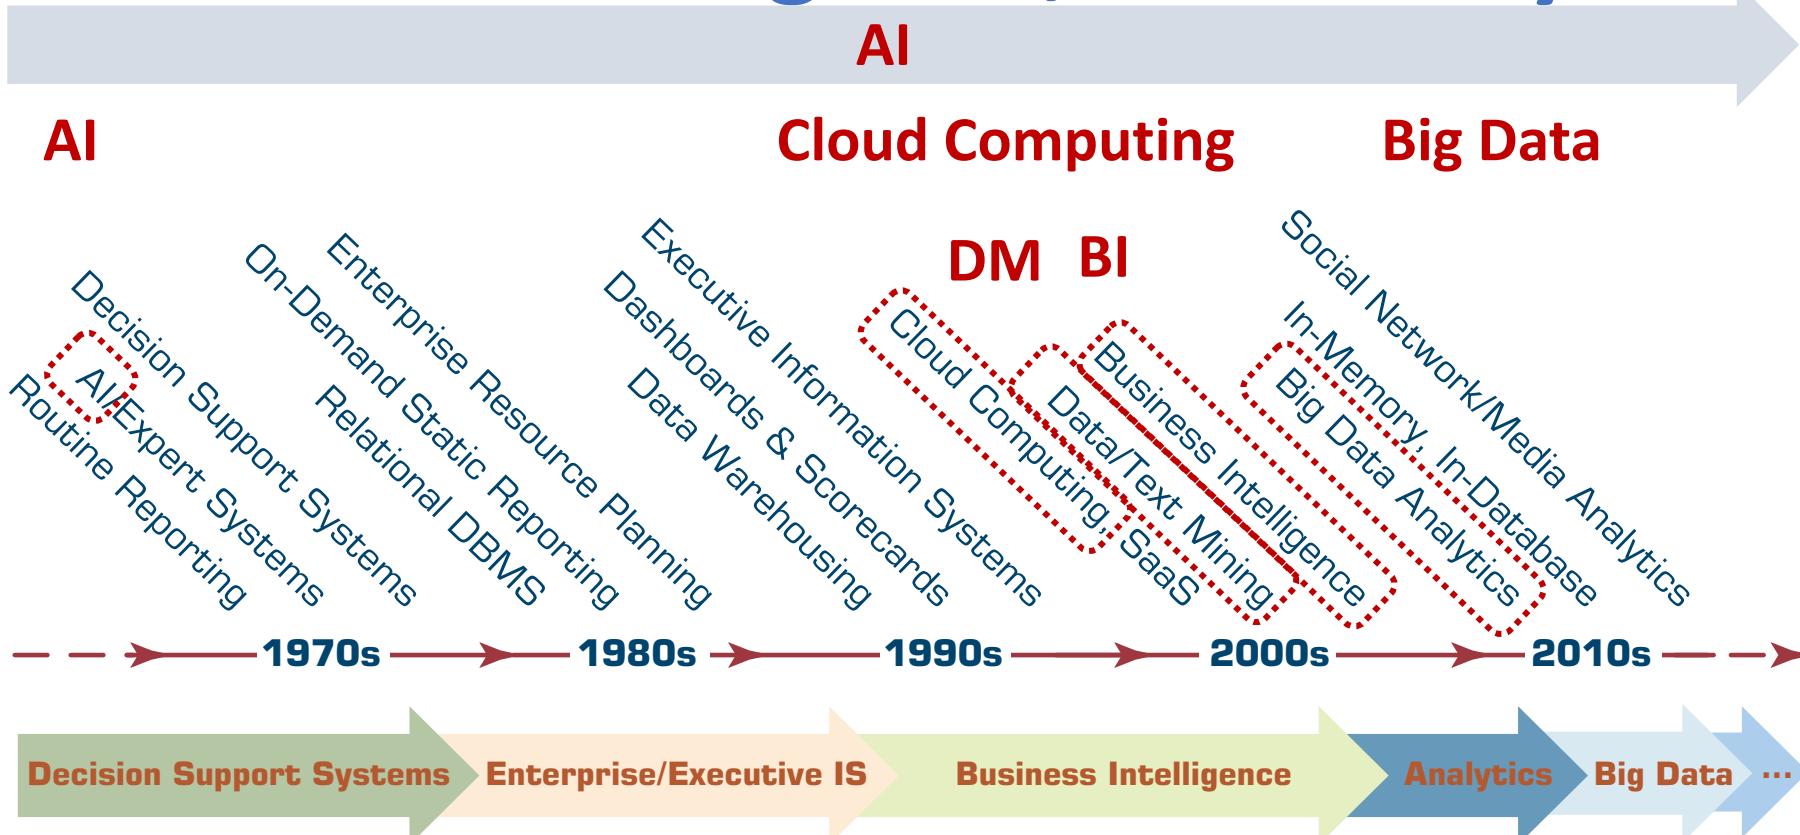
Instruction I_4

Prompt

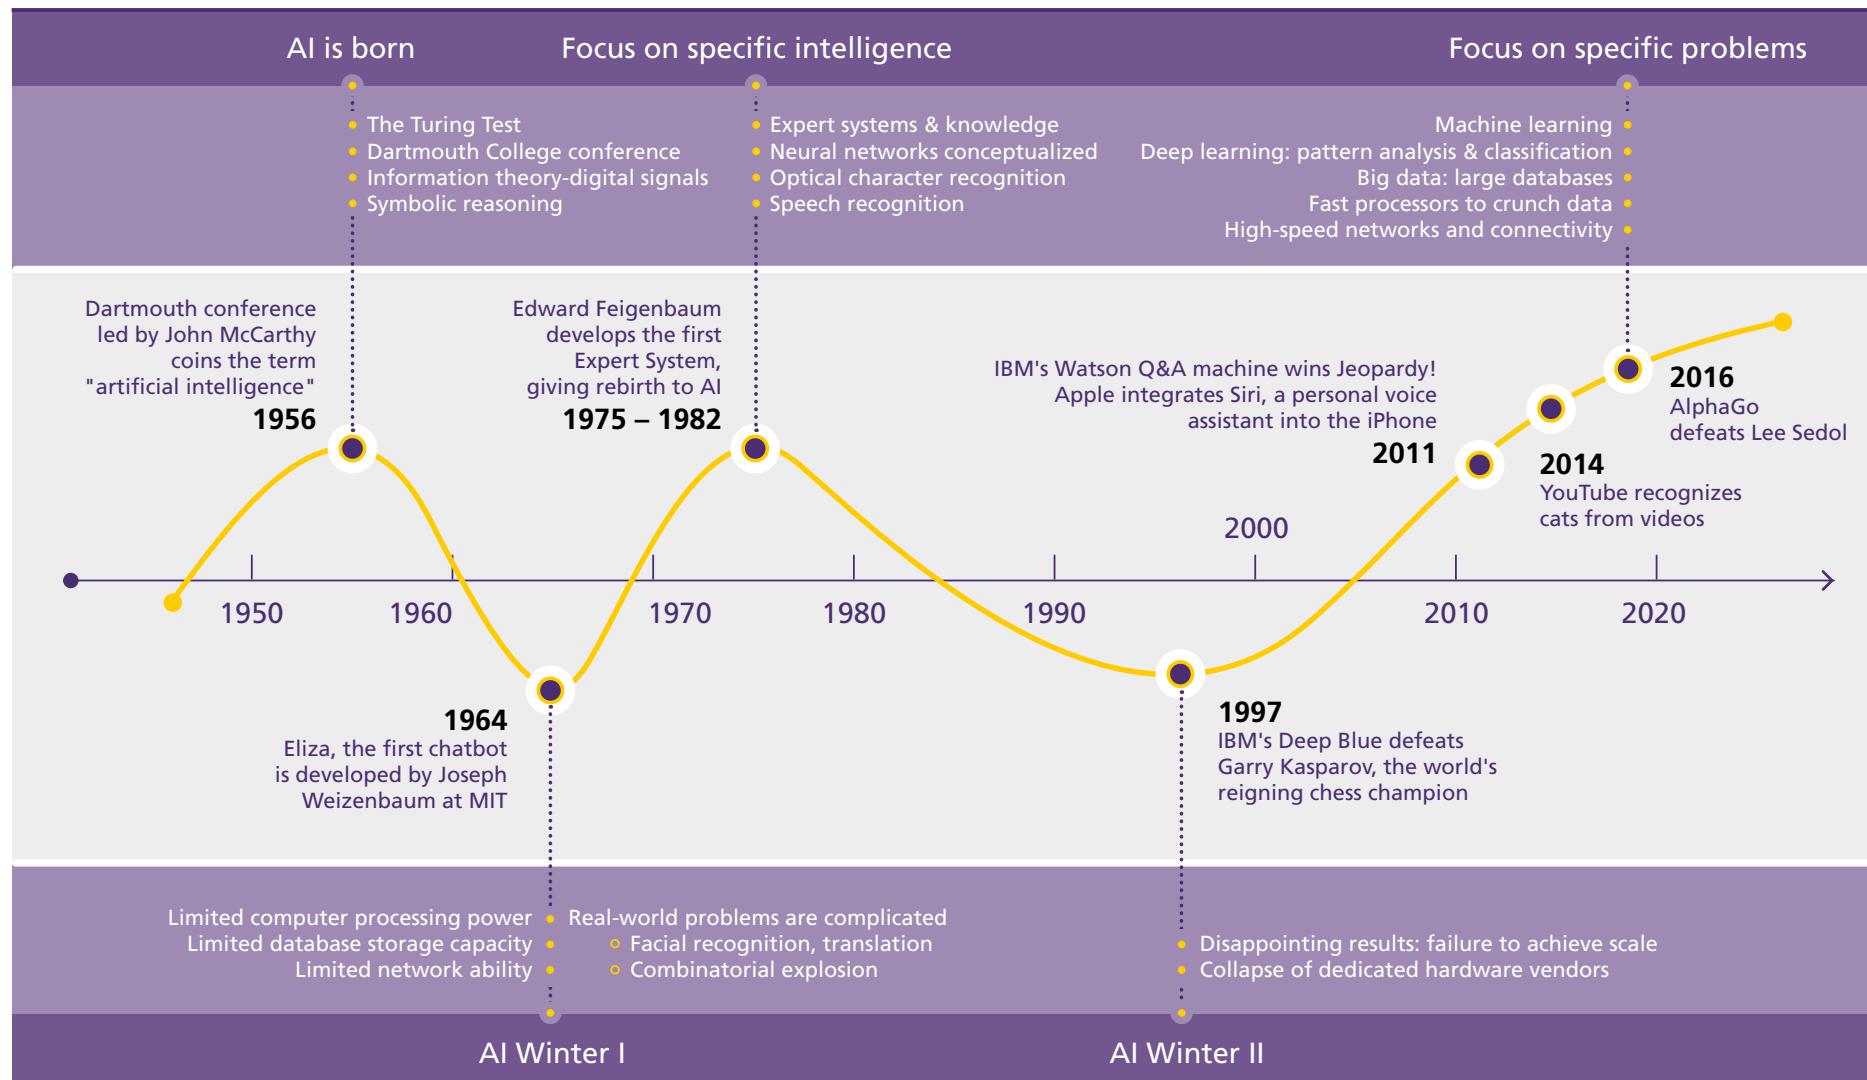


AI, Big Data, Cloud Computing

Evolution of Decision Support, Business Intelligence, and Analytics



The Rise of AI



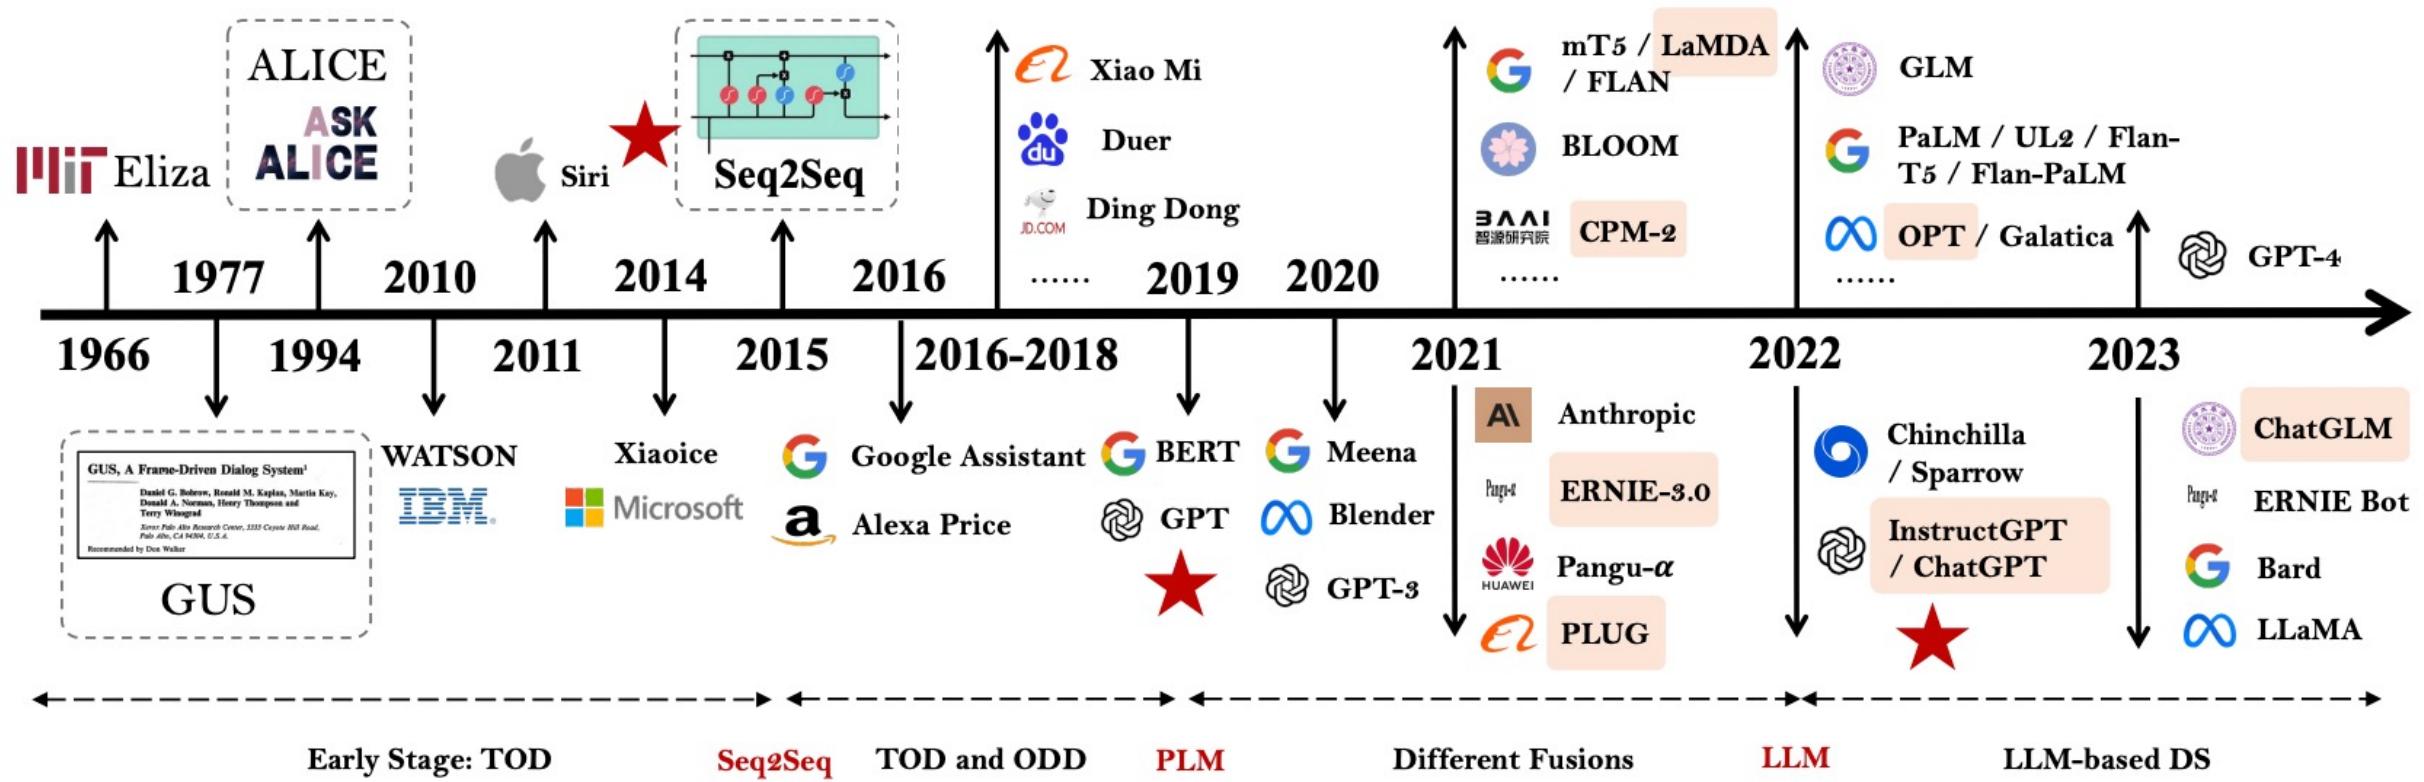
The Development of LM-based Dialogue Systems

1) Early Stage (1966 - 2015)

2) The Independent Development of TOD and ODD (2015 - 2019)

3) Fusions of Dialogue Systems (2019 - 2022)

4) LLM-based DS (2022 - Now)



Task-oriented DS (TOD), Open-domain DS (ODD)

Definition of Artificial Intelligence (A.I.)

Artificial Intelligence

“... the science and
engineering
of
making
intelligent machines”
(John McCarthy, 1955)

Artificial Intelligence

“... technology that
thinks and acts
like humans”

Artificial Intelligence

“... intelligence
exhibited by machines
or software”

4 Approaches of AI

Thinking Humanly	Thinking Rationally
Acting Humanly	Acting Rationally

4 Approaches of AI

<p>2.</p> <p>Thinking Humanly: The Cognitive Modeling Approach</p>	<p>3.</p> <p>Thinking Rationally: The “Laws of Thought” Approach</p>
<p>1.</p> <p>Acting Humanly: The Turing Test Approach <small>(1950)</small></p>	<p>4.</p> <p>Acting Rationally: The Rational Agent Approach</p>

AI Acting Humanly: The Turing Test Approach

(Alan Turing, 1950)

- Knowledge Representation
- Automated Reasoning
- Machine Learning (ML)
 - Deep Learning (DL)
- Computer Vision (Image, Video)
- Natural Language Processing (NLP)
- Robotics

AI, ML, DL

Artificial Intelligence (AI)

Machine Learning (ML)

Supervised
Learning

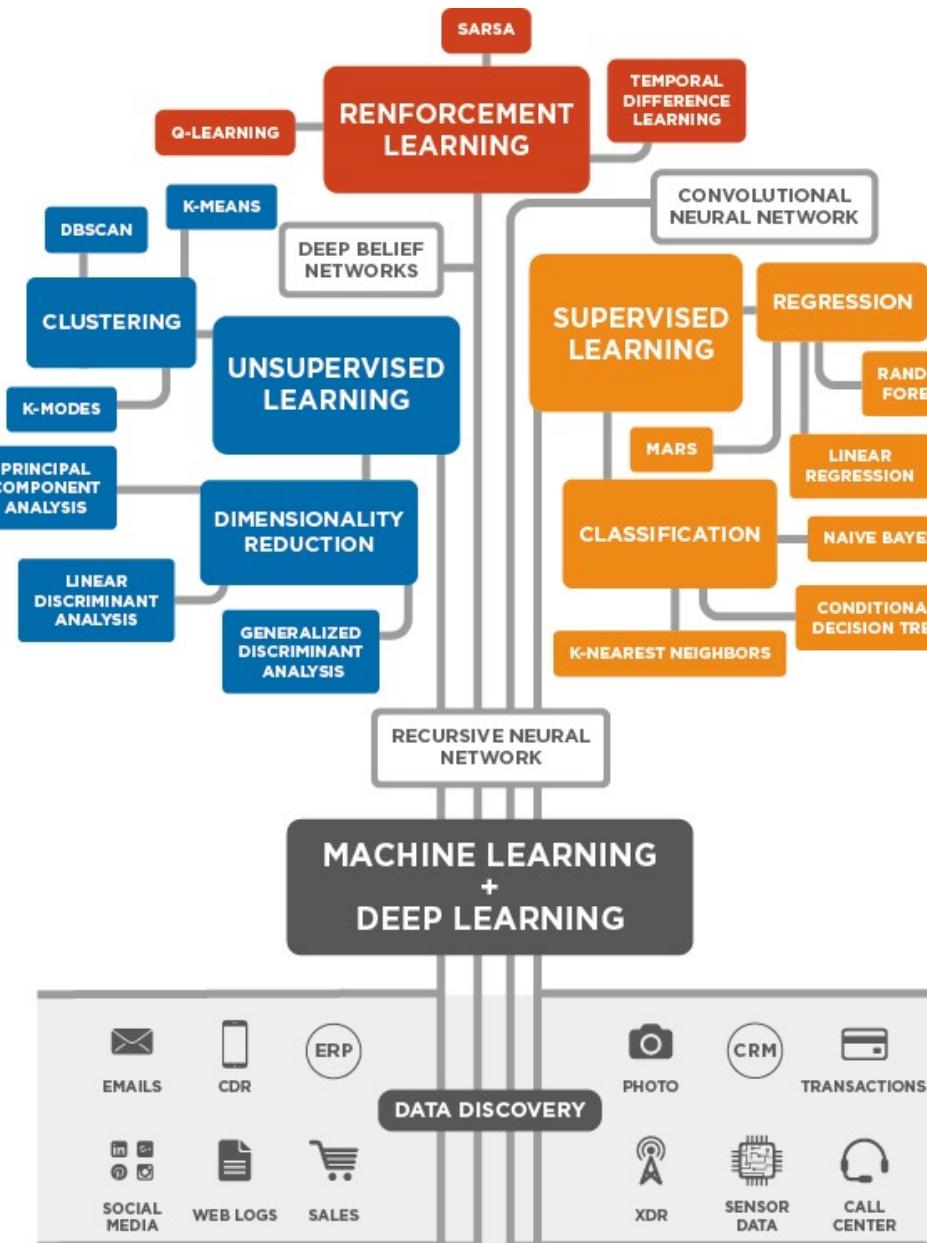
Unsupervised
Learning

Deep Learning (DL)
CNN
RNN LSTM GRU
GAN

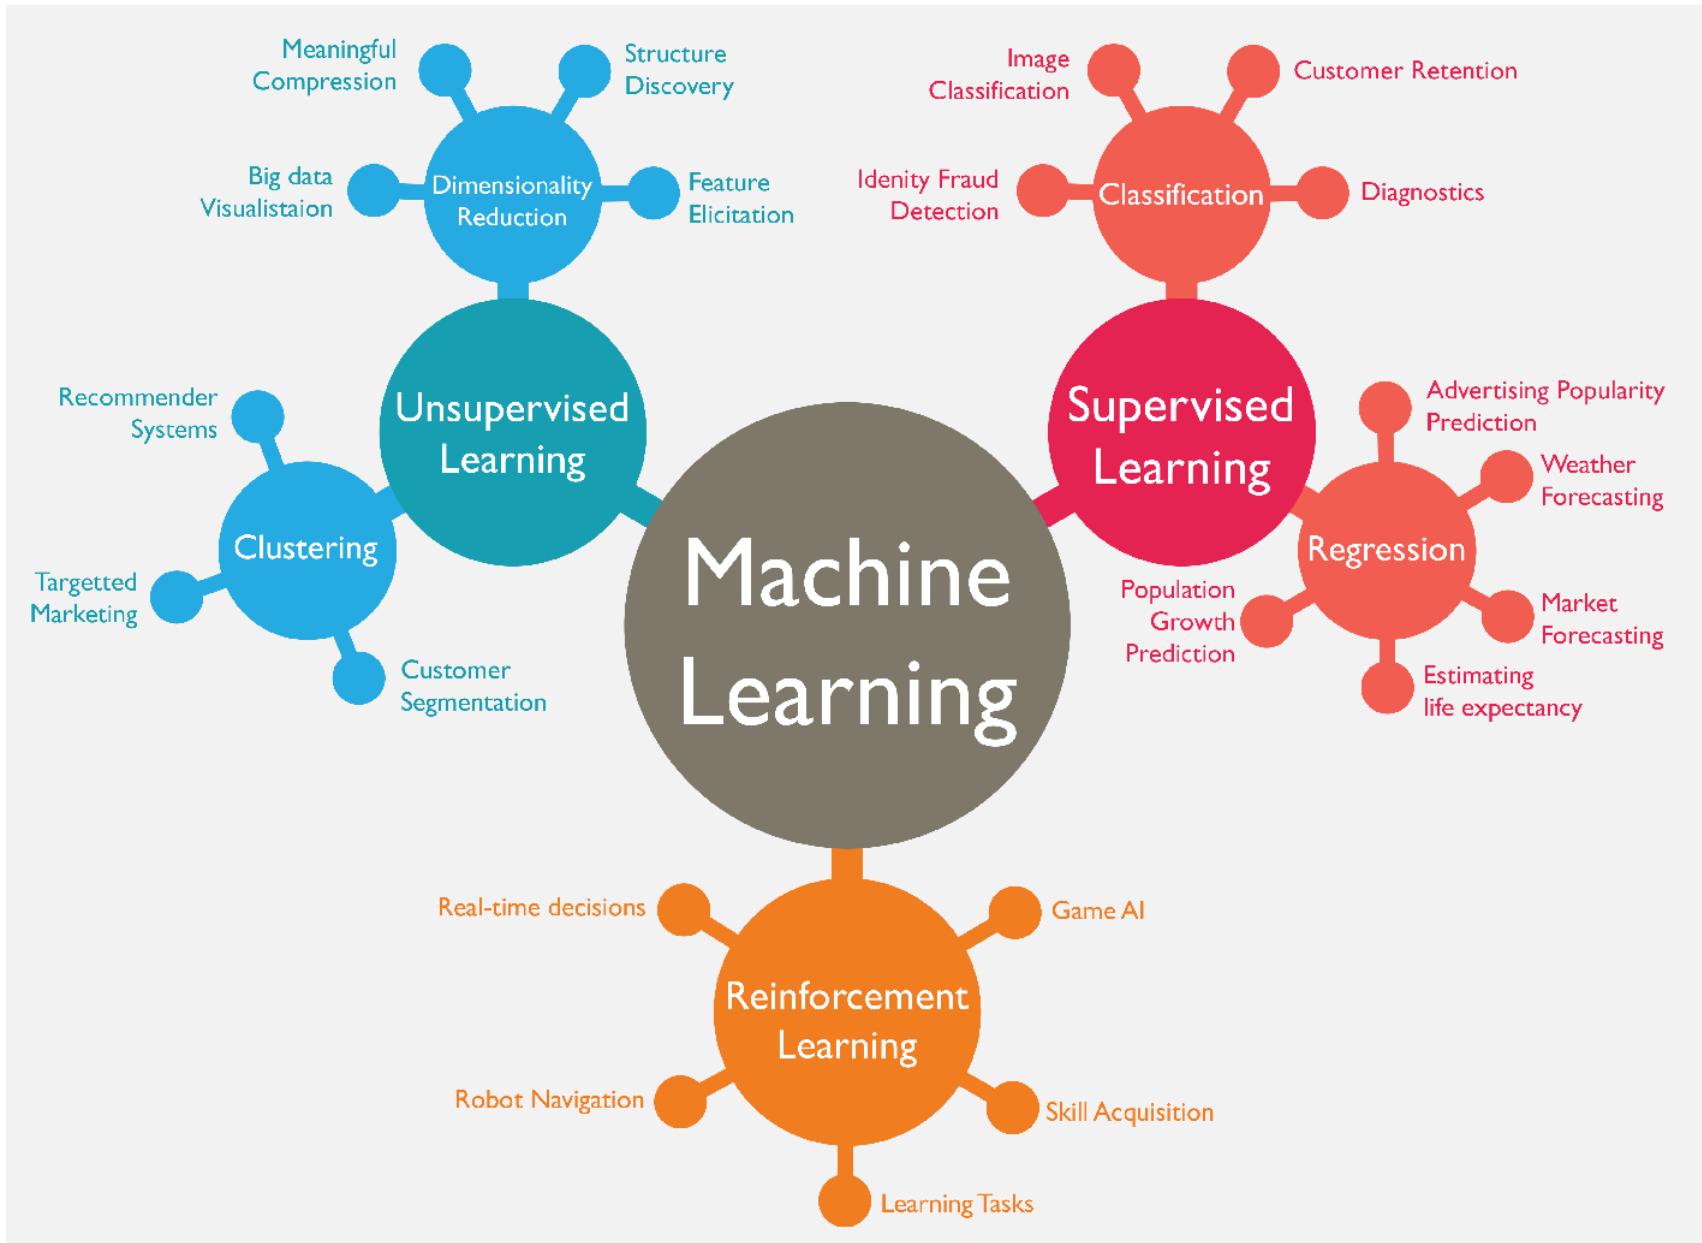
Semi-supervised
Learning

Reinforcement
Learning

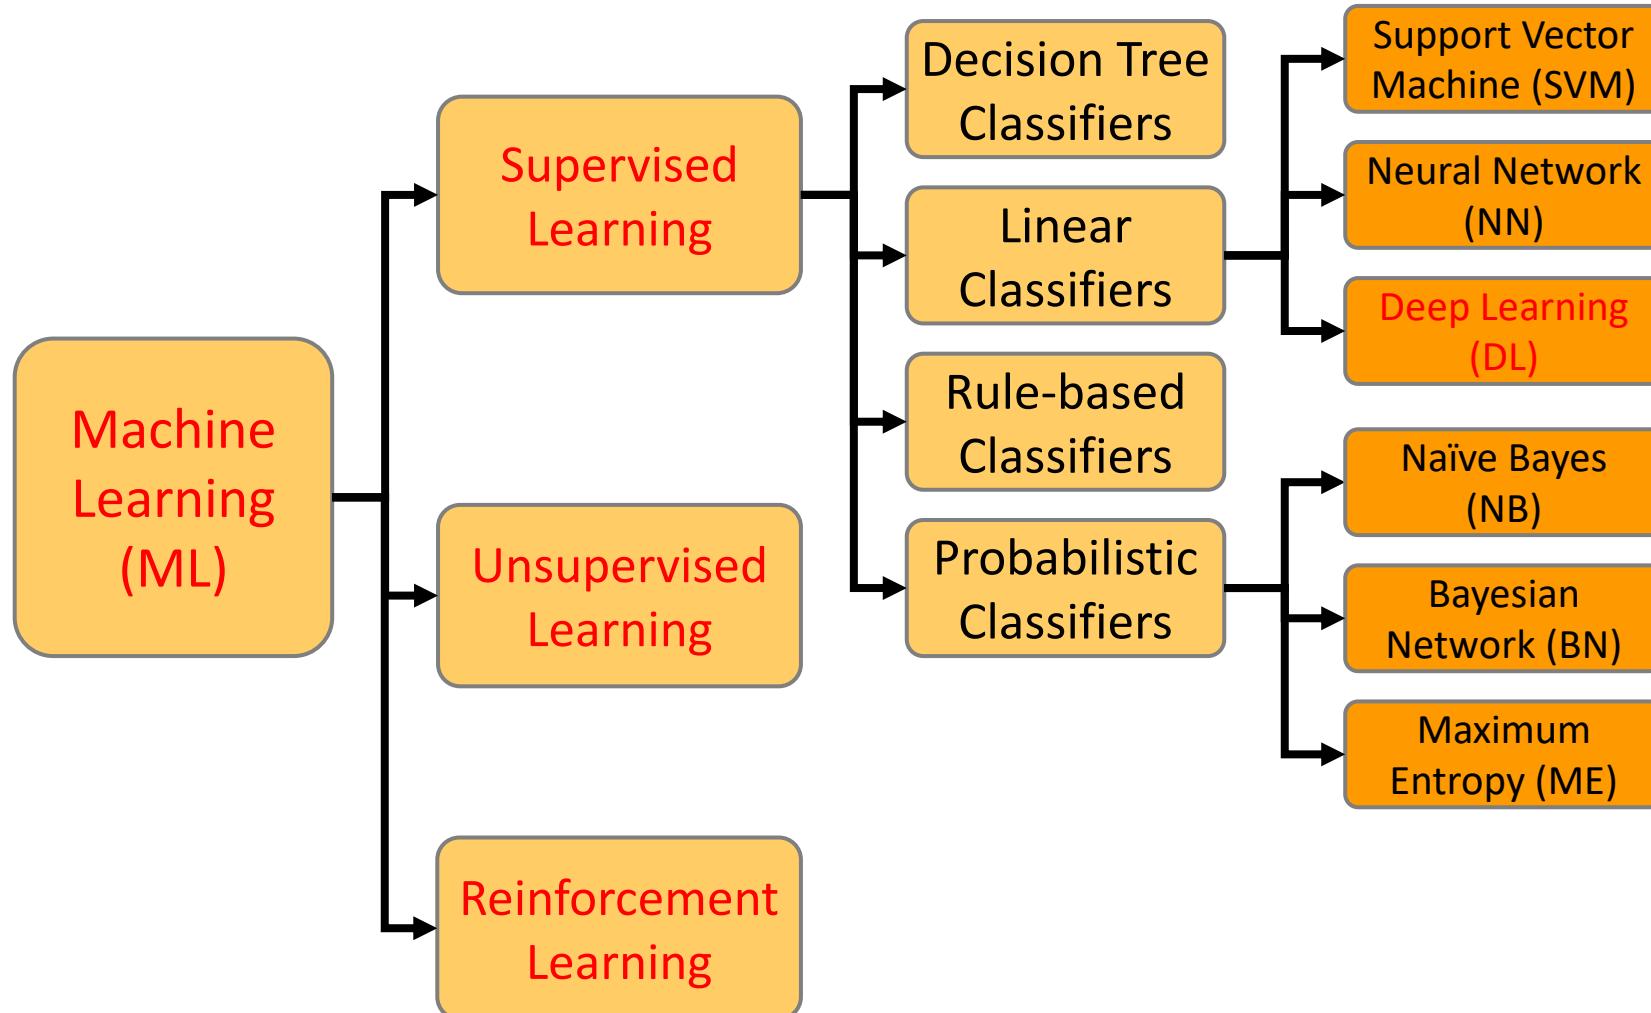
3 Machine Learning Algorithms



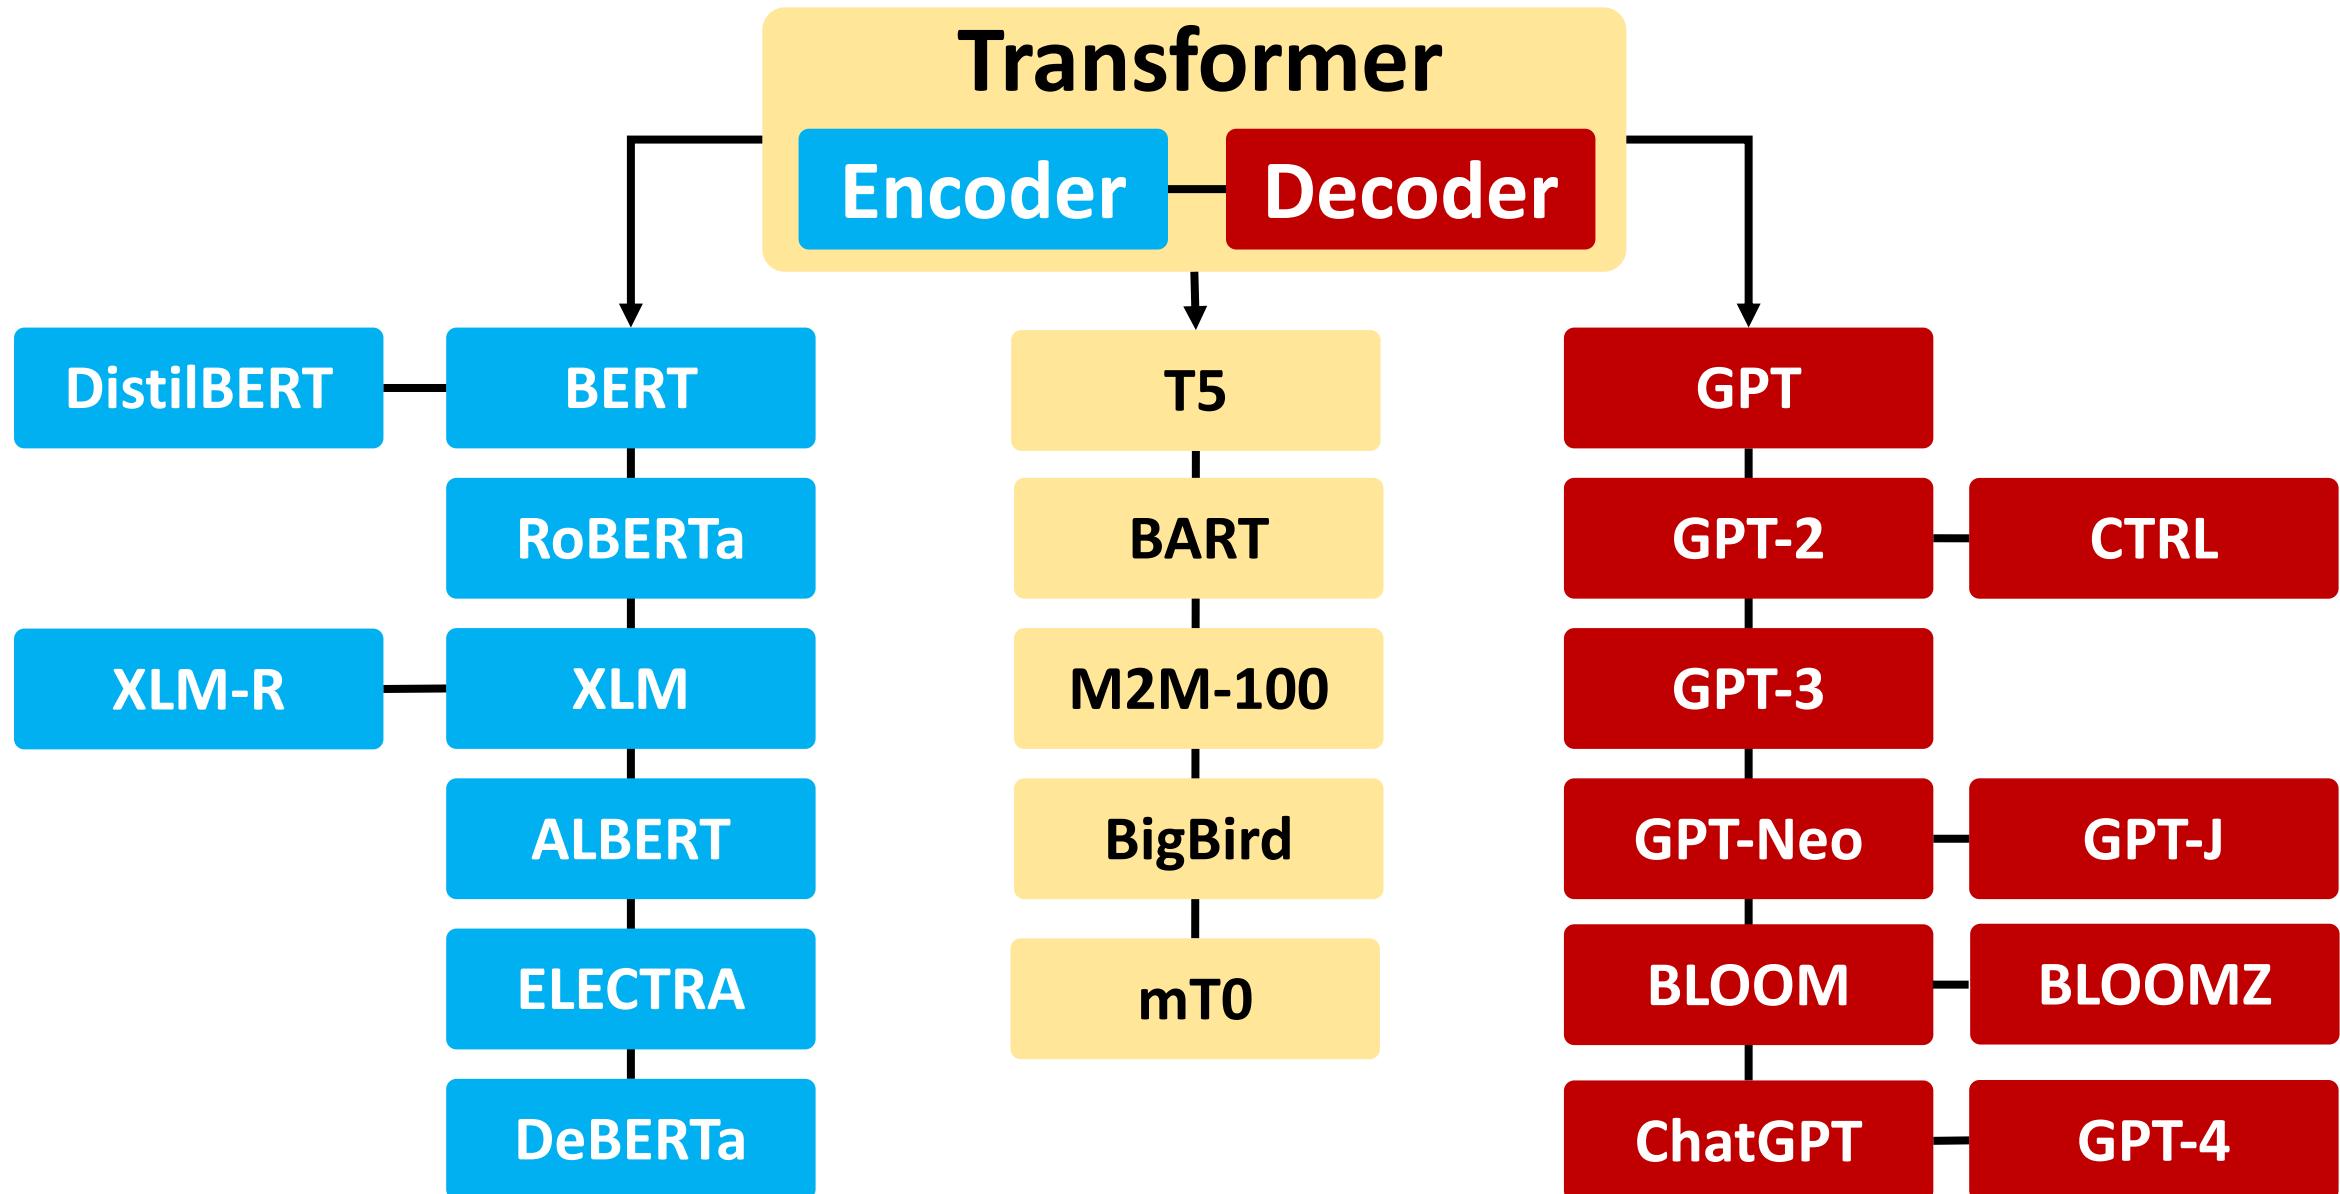
Machine Learning (ML)



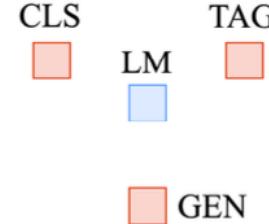
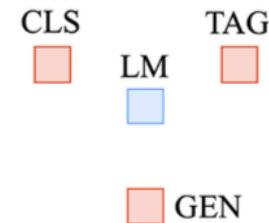
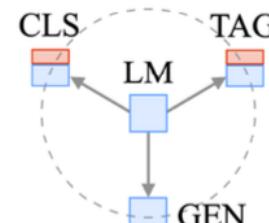
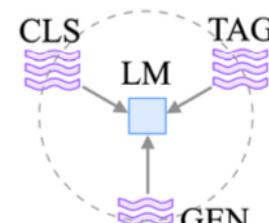
Machine Learning (ML) / Deep Learning (DL)



Transformer Models



Four Paradigms in NLP (LM)

Paradigm	Engineering	Task Relation
a. Fully Supervised Learning (Non-Neural Network)	Feature (e.g. word identity, part-of-speech, sentence length)	
b. Fully Supervised Learning (Neural Network)	Architecture (e.g. convolutional, recurrent, self-attentional)	
c. Pre-train, Fine-tune	Objective (e.g. masked language modeling, next sentence prediction)	
Transfer Learning: Pre-training, Fine-Tuning (FT)		
d. Pre-train, Prompt, Predict	Prompt (e.g. cloze, prefix)	

Generative AI

Text, Image, Video, Audio

Applications

Comparison of Generative AI and Traditional AI

Feature	Generative AI	Traditional AI
Output type	New content	Classification/Prediction
Creativity	High	Low
Interactivity	Usually more natural	Limited

Generative AI

- **Generative AI: The Art of Creation**
- **Definition: AI systems capable of creating new content**
- **Characteristics: Creativity, interactivity**

LMArena Leaderboard

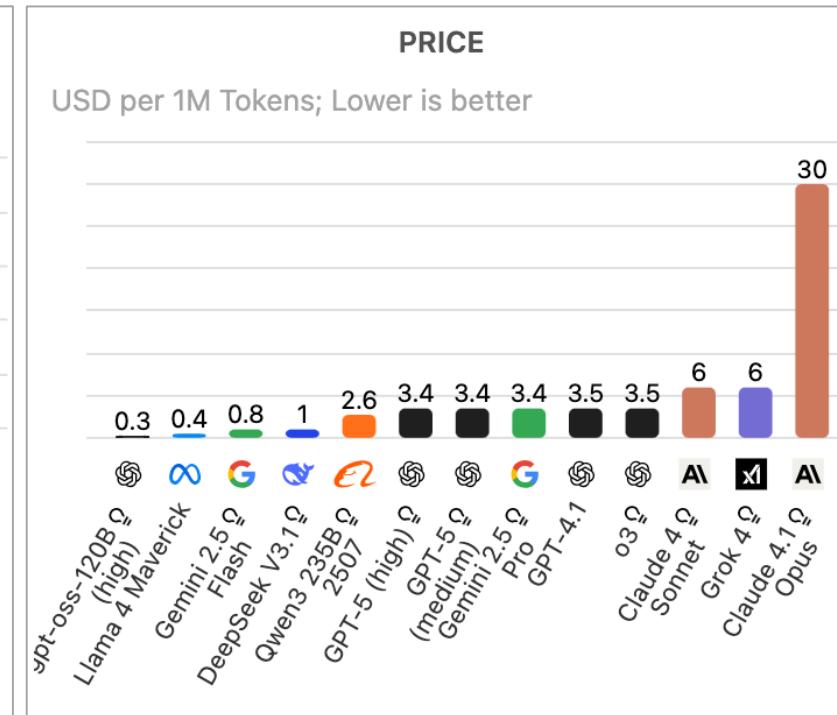
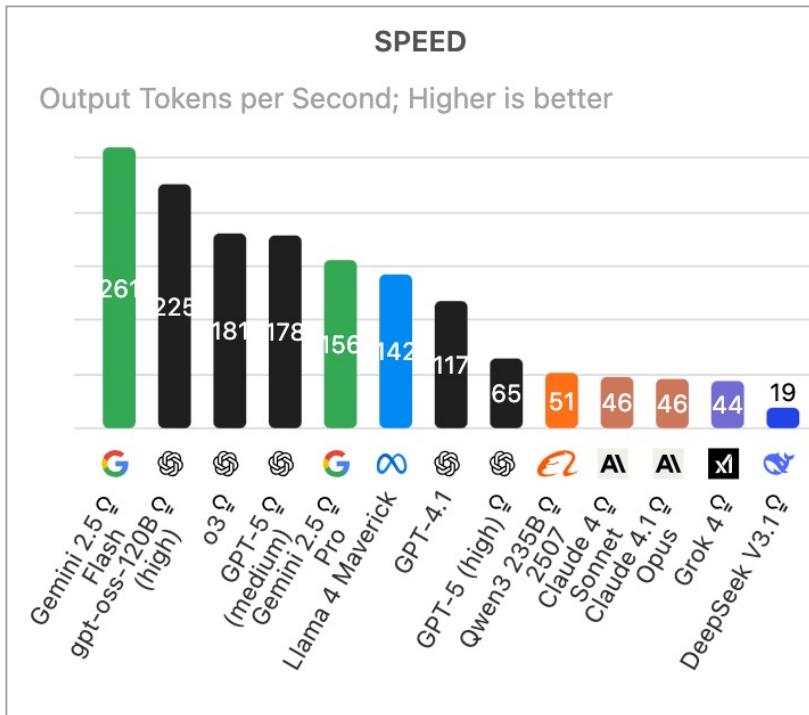
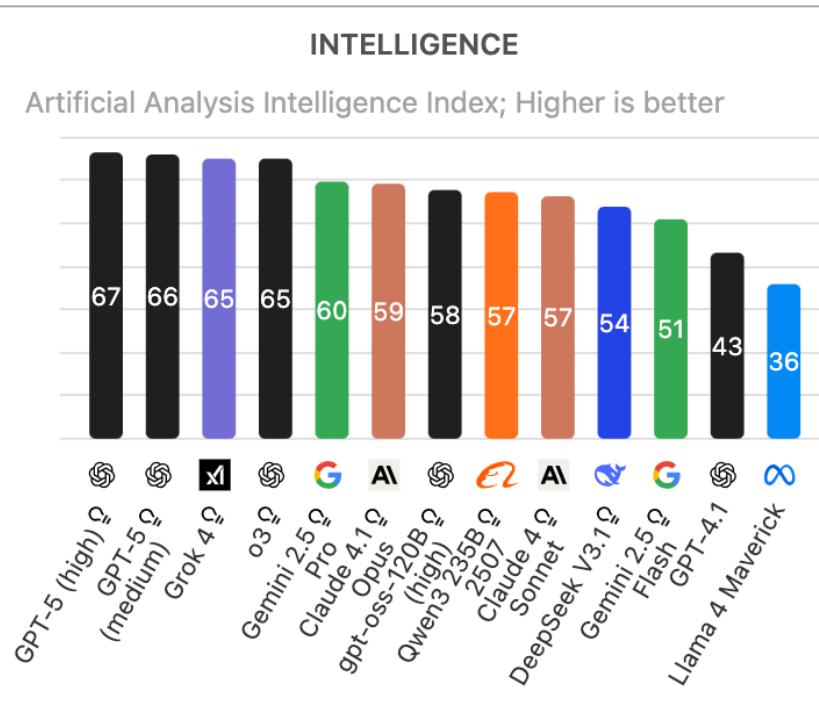
Rank (UB) ↑	Model ↓	Score ↑	95% CI (±) ↑↓	Votes ↑↓	Organization ↑↓	License ↑↓
1	 gemini-2.5-pro	1455	±5	41,731	Google	Proprietary
1	 claude-opus-4-1-20250805-thinking-16k	1451	±6	11,750	Anthropic	Proprietary
2	 o3-2025-04-16	1444	±4	43,898	OpenAI	Proprietary
2	 gpt-5-high	1442	±6	15,076	OpenAI	Proprietary
2	 chatgpt-4o-latest-20250326	1441	±4	36,426	OpenAI	Proprietary
3	 gpt-4.5-preview-2025-02-27	1439	±6	15,271	OpenAI	Proprietary
3	 claude-opus-4-1-20250805	1438	±6	18,341	Anthropic	Proprietary
5	 gpt-5-chat	1430	±6	11,808	OpenAI	Proprietary
6	 qwen3-max-preview	1428	±7	8,781	Alibaba	Proprietary
8	 grok-4-0709	1422	±5	21,446	xAI	Proprietary

LMArena Leaderboard

Q Model ▾	239 / 239	Overall ↑↓	Hard Prompts ↑↓	Coding ↑↓	Math ↑↓	Creative Writing ↑↓	Instruction Following	Longer Query ↑↓	Multi-Turn ↑↓
AI claude-opus-4-1...	1	1	1	1	1	1	1	1	1
Google gemini-2.5-pro	1	2	3	1	1	1	1	1	1
Coqui chatgpt-4o-lates...	2	4	3	13	2	5	4	1	1
Coqui gpt-5-high	2	2	3	1	7	5	11	6	
Coqui o3-2025-04-16	2	4	3	1	8	6	13	7	
AI claude-opus-4-1...	3	2	1	1	1	1	1	1	1
Coqui gpt-4.5-preview-...	3	5	4	8	1	4	3	1	
Coqui gpt-5-chat	5	3	3	8	3	5	3	1	
Huggingface qwen3-max-preview	6	4	2	1	7	4	4	3	
AI claude-opus-4-20...	8	4	3	6	2	2	2	7	
Coqui deepseek-r1-0528	8	8	4	10	8	15	13	14	
Coqui deepseek-v3.1	8	6	4	1	7	6	5	9	
Coqui deepseek-v3.1-th...	8	4	3	1	2	4	1	7	
XL grok-4-0709	8	10	12	1	4	6	8	7	
Coqui kimi-k2-0711-pre...	8	10	7	13	16	24	22	7	
Coqui kimi-k2-0905-pre...	8	5	3	-	6	16	12	7	
Huggingface qwen3-235b-a22b-...	8	4	3	2	9	6	4	7	
Zglm glm-4.5	10	7	4	7	14	7	8	10	

Artificial Analysis Intelligence Index

Intelligence, Speed, Price



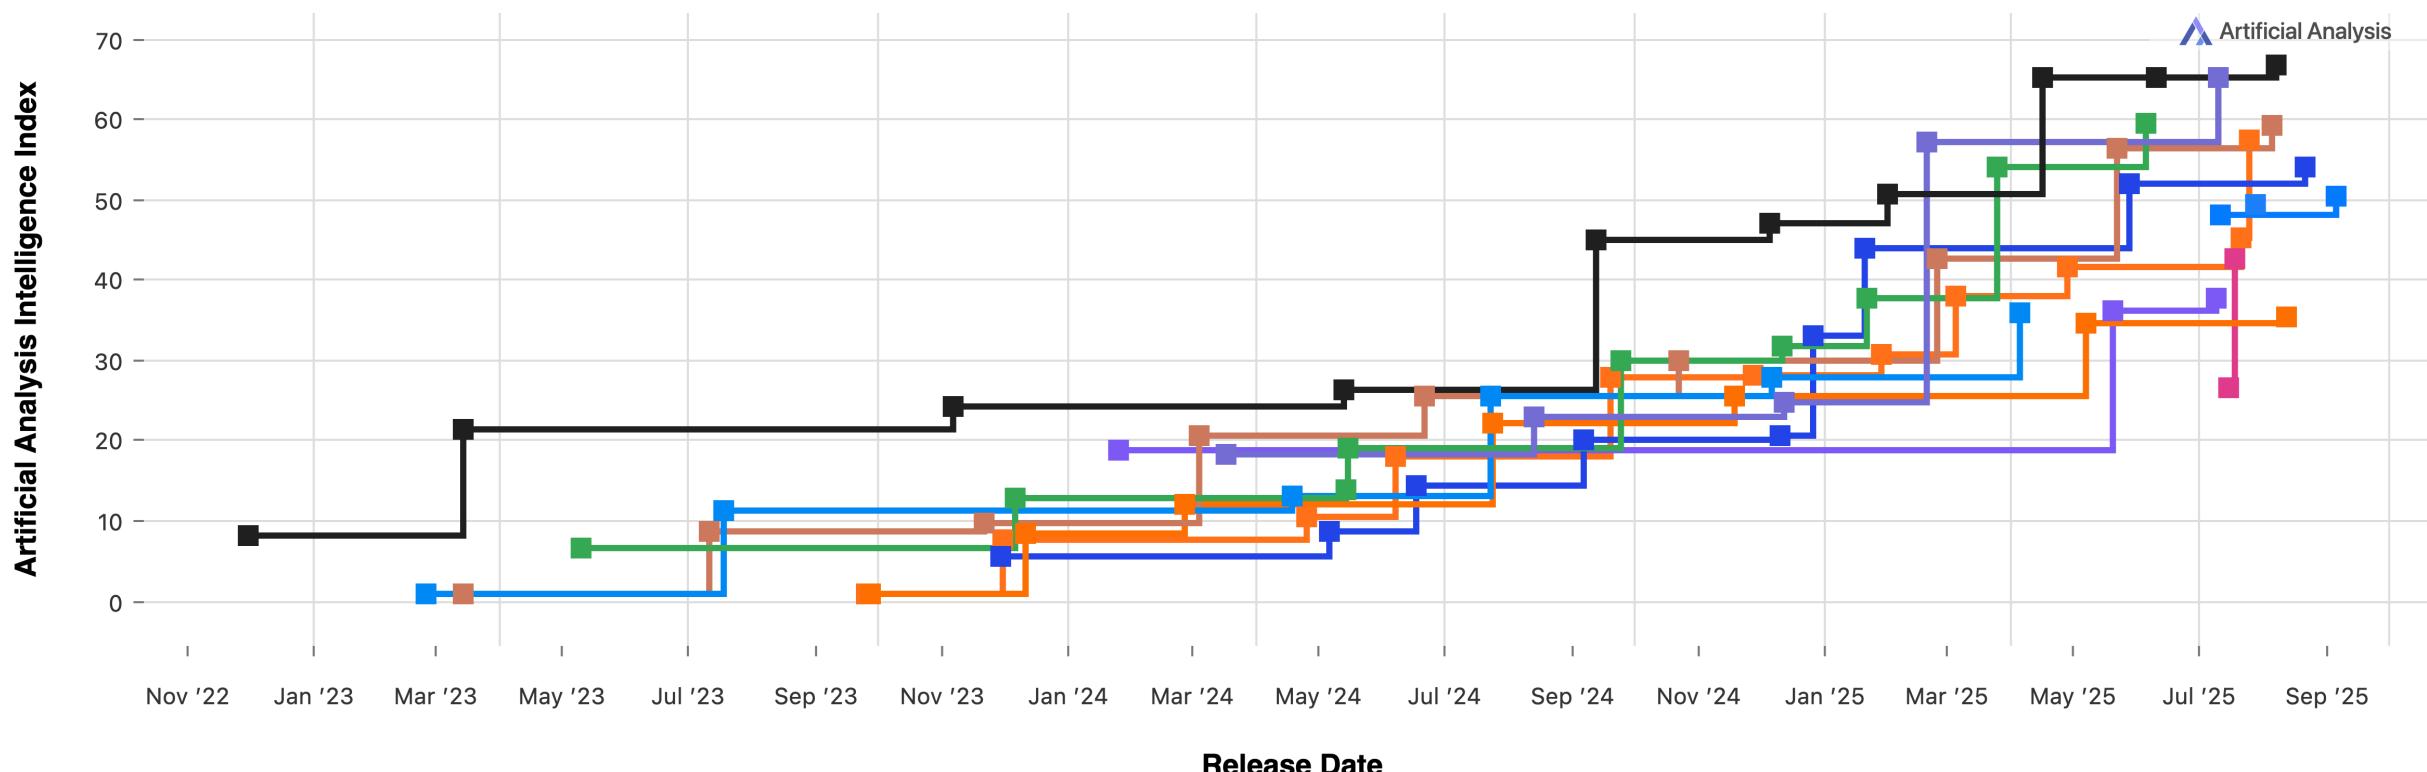
Artificial Analysis Intelligence Index

2022-2025

Frontier Language Model Intelligence, Over Time

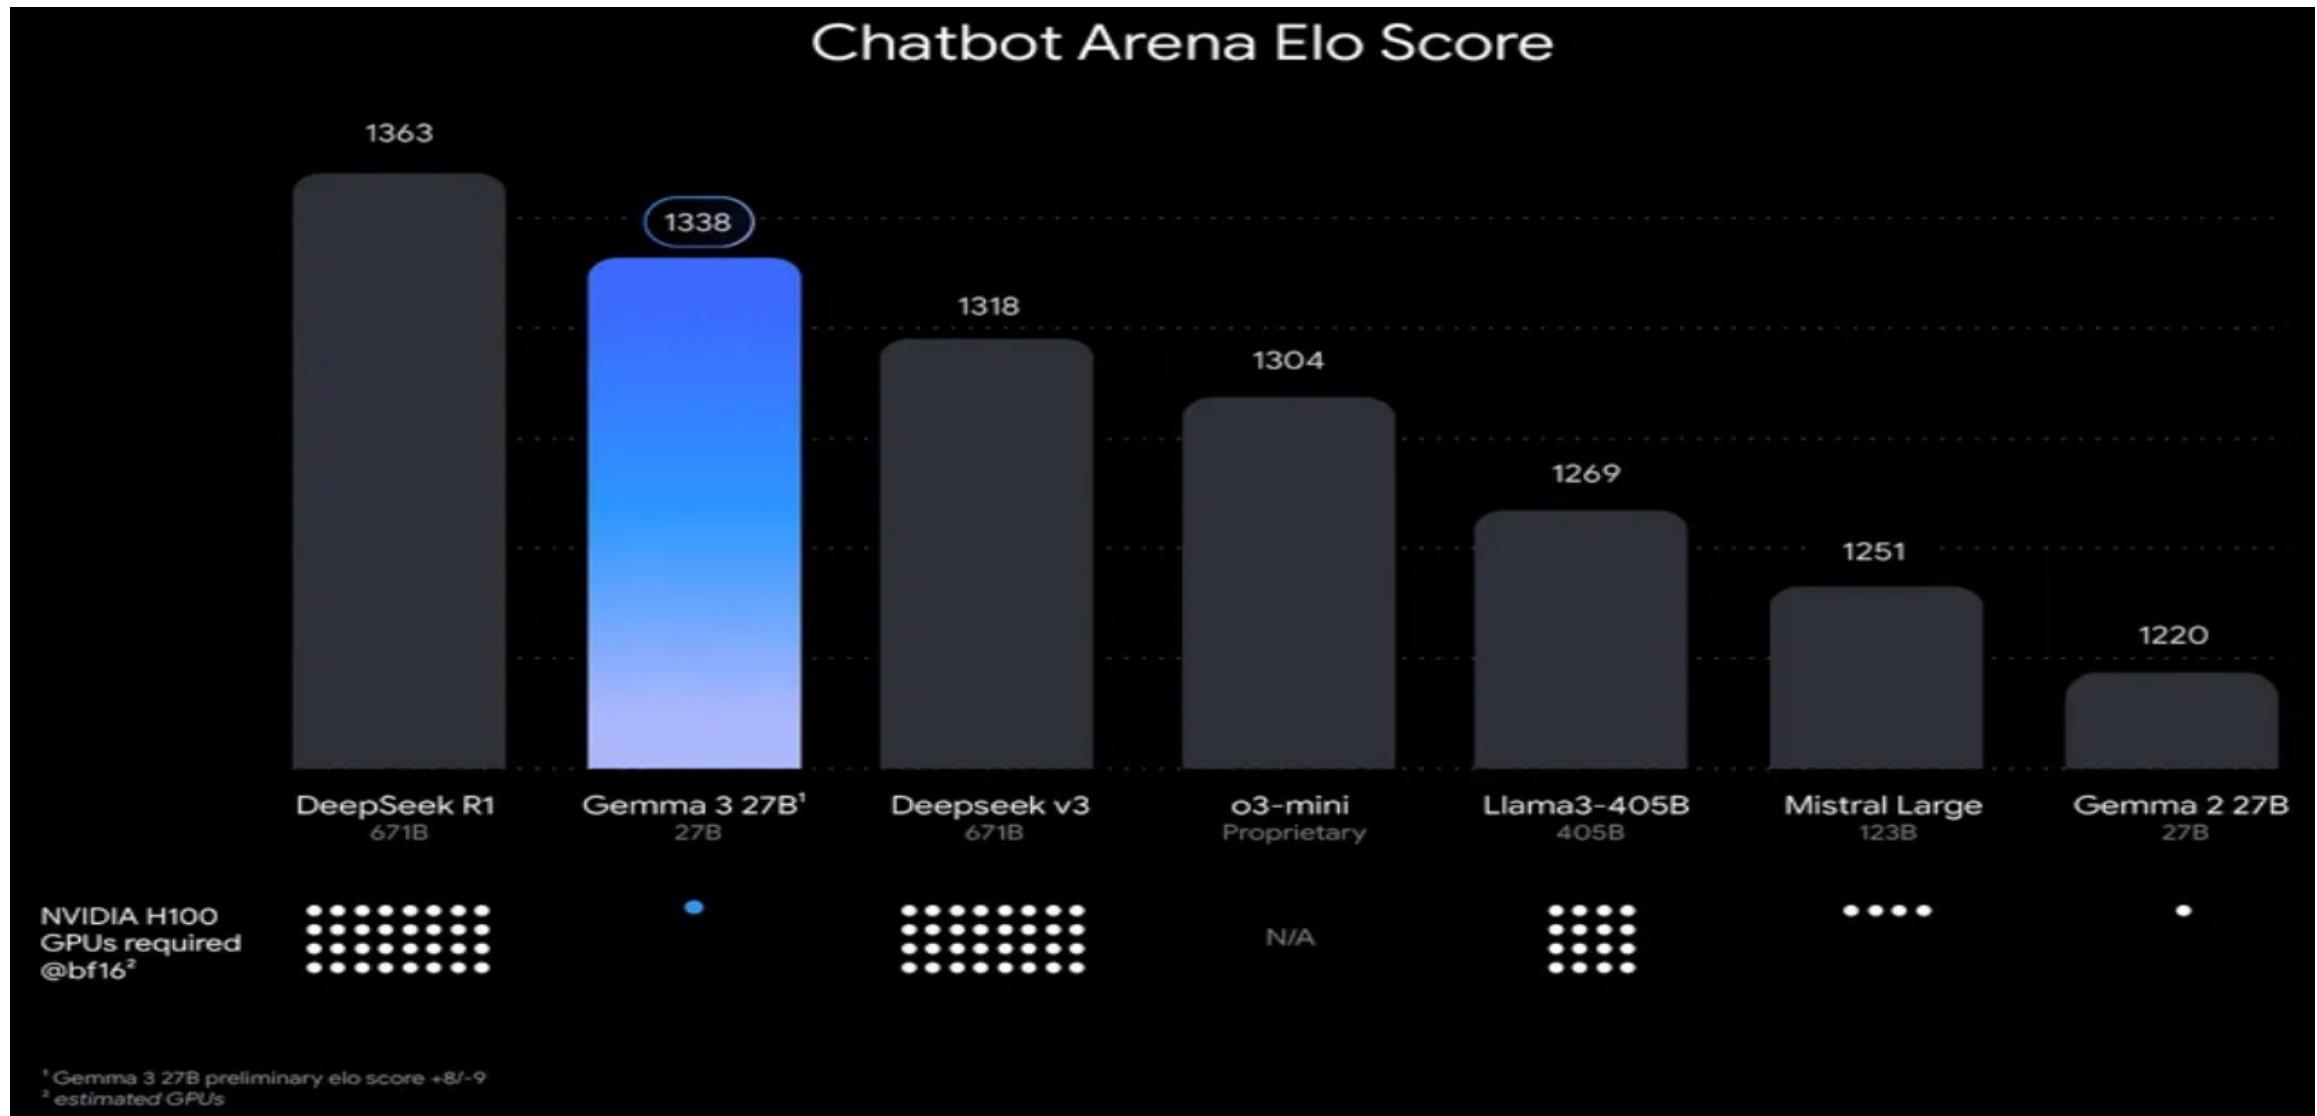
Artificial Analysis Intelligence Index v3.0 incorporates 10 evaluations: MMLU-Pro, GPQA Diamond, Humanity's Last Exam, LiveCodeBench, SciCode, AIME 2025, IFFBench, AA-LCR, Terminal-Bench Hard, τ^2 -Bench Telecom

Alibaba Anthropic DeepSeek Google LG AI Research Meta Mistral Moonshot AI OpenAI Upstage xAI Z AI



Google Gemma 3 27B

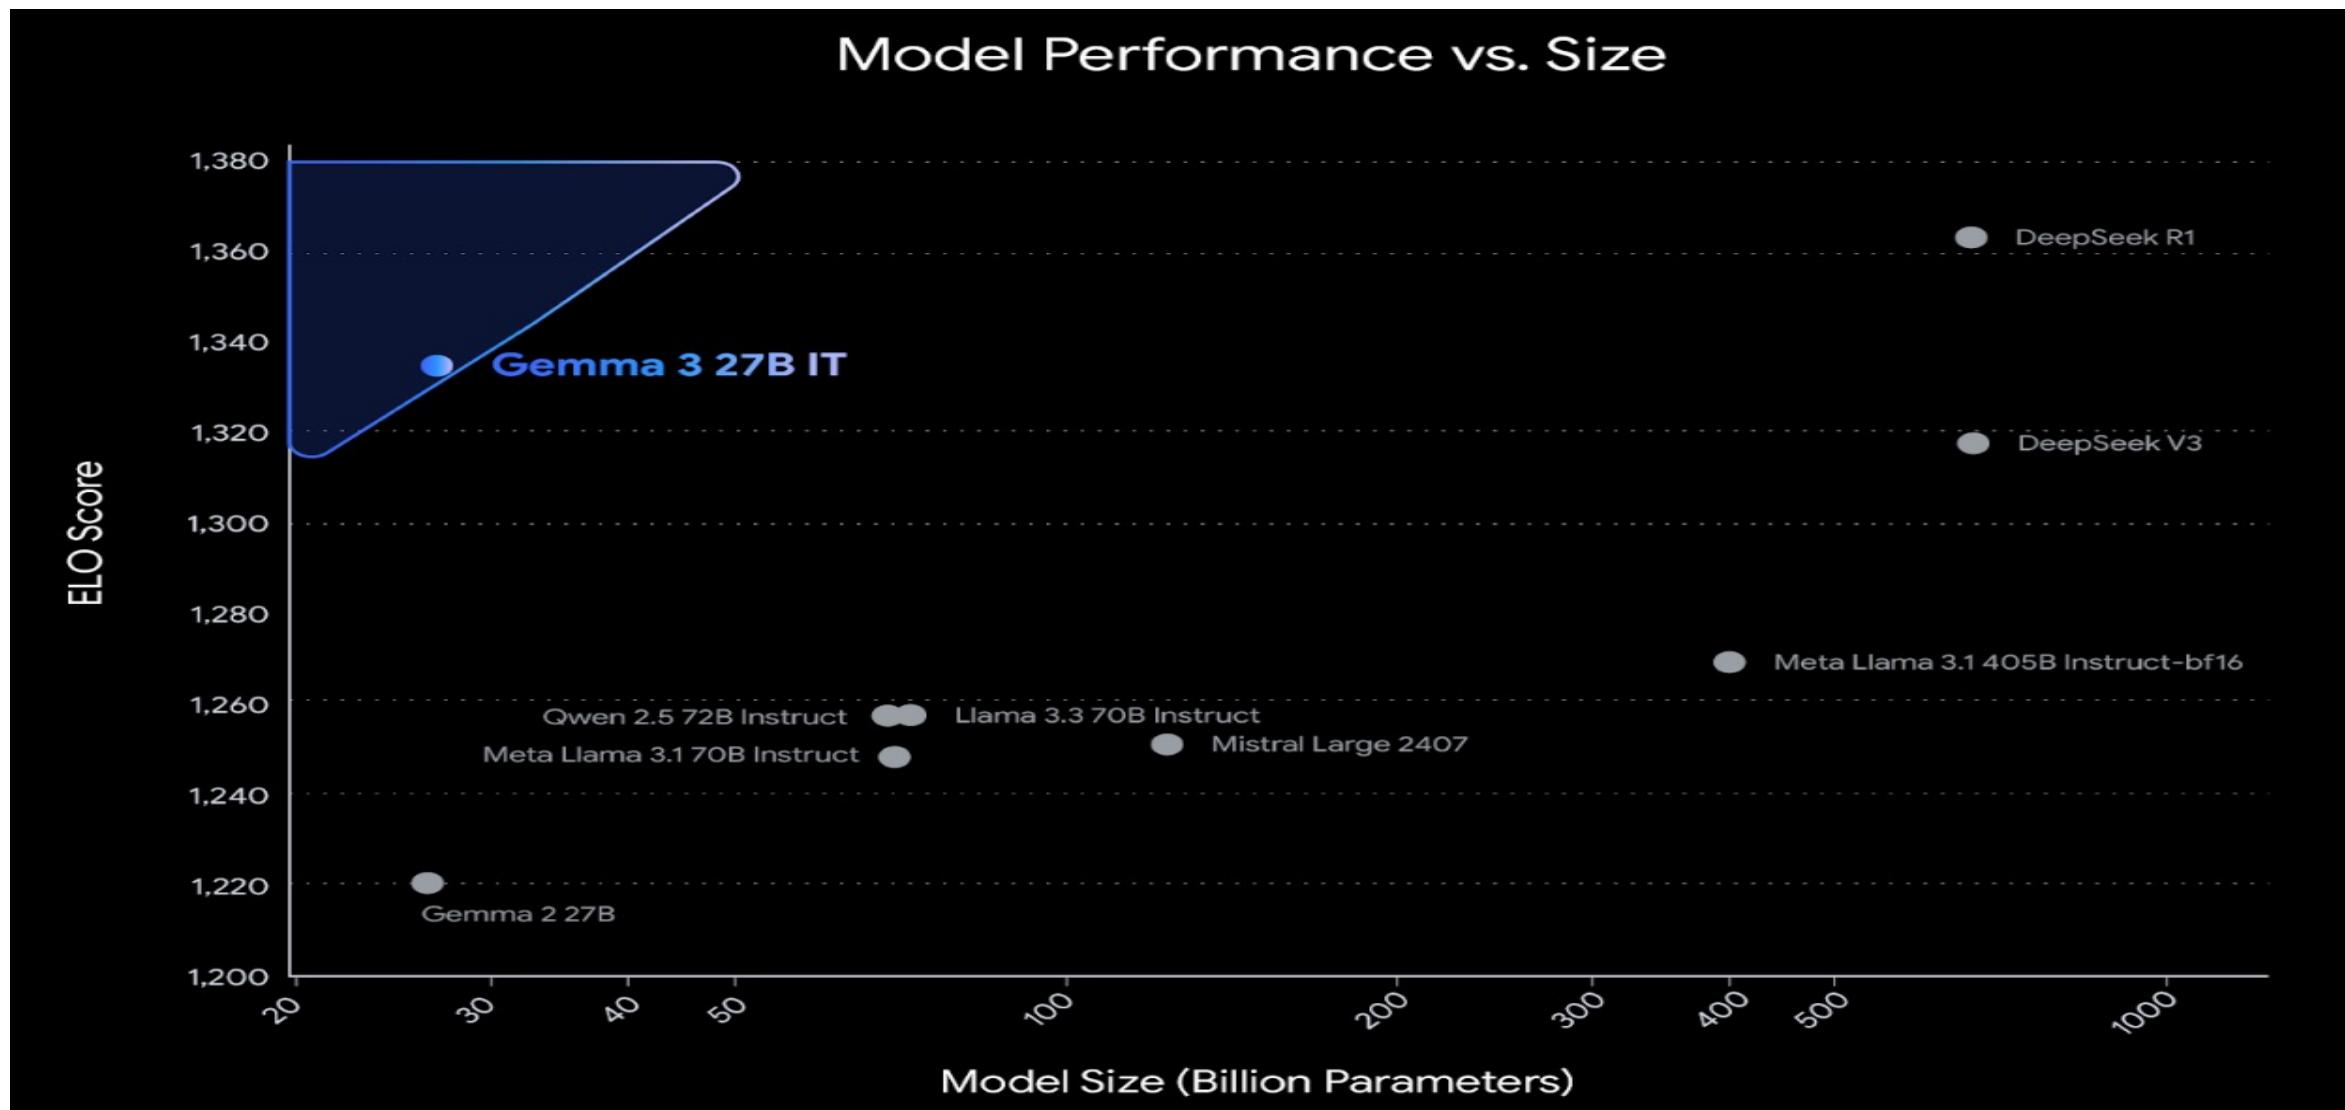
The most capable model you can run on a single GPU or TPU



Google Gemma 3 Multimodality (vision-language input and text outputs)

MODEL	SIZE (in billion parameter)	CONTEXT LENGTH	LANGUAGES	INPUT MODALITIES
Gemma 3 1B (IT)	1B	32k	English	Input: Text Output: Text
Gemma 3 4B (IT)	4B	128k	+140 Languages	Input: Text, Image Output: Text
Gemma 3 12B (IT)	12B	128k	+140 Languages	Input: Text, Image Output: Text
Gemma 3 27B (IT)	27B	128k	+140 Languages	Input: Text, Image Output: Text
Shield Gemma 2	4B	8k	+140 Languages	Input: Text, Image Output: Text

Google Gemma 3: Pre-training and Post-training (distillation, reinforcement learning, and model merging)



Google AI Studio (Gemma 3 27B)

Google AI Studio Generative AI: A Comprehensive Over... Enable Autosave Compare Open in Kaggle Open in Vertex AI

Stream Realtime Tune a Model Developer documentation Changelog Gemma models may make mistakes, so double-check its response. Gemma is provided under and subject to the Gemma Terms of Use.

Generative AI: A Deep Dive

Generative AI is one of the most exciting and rapidly evolving fields in artificial intelligence right now. Here's a comprehensive overview, covering what it is, how it works, its types, applications, limitations, and future trends:

1. What is Generative AI?

Generative AI refers to a category of machine learning models that can **create new content** – text, images, audio, video, code, and more – that resembles the data they were trained on. Unlike traditional AI which *analyzes* or *acts* on existing data, generative AI *produces* new data. Think of it as AI that can "imagine" and then manifest those imaginations.

Key Characteristics:

- **Learns Patterns:** Generative AI models learn the underlying patterns and structures within a dataset.
- **Generates Novel Content:** They use this learned knowledge to create new, original content that wasn't explicitly present in the training data.
- **Probabilistic:** Generation isn't deterministic. The models predict probabilities, leading to variations in output even with the same input.
- **Requires Large Datasets:** Effective generative AI typically needs massive amounts of data for training.

Run settings Model Gemma 3 27B Token count 1,760 / 131,072 Temperature 1 Advanced settings Safety settings Edit safety settings Add stop sequence Add stop... Output length 8192

Grok 3 Deep Search

The screenshot shows the Grok 3 Deep Search interface. At the top, there is a navigation bar with icons for refresh, upload, and user profile, and a search bar labeled "Generative AI". Below the search bar is a sidebar titled "Completed" with a timestamp of "49s · 71 Sources". The sidebar lists four tasks: "Thinking", "Exploring user intent", "Researching recent developments", and "Investigating healthcare applications", each preceded by a checkmark icon.

Exploring user intent

- The request is "Generative AI," which feels a bit unclear. It could be asking for a definition, recent updates, or use cases.

Searching for "definition of generative AI"

10 results found

- Generative artificial intelligence - Wikipedia en.wikipedia.org
- Generative AI – What is it and How Does it Work? | NVIDIA nvidia.com
- What is generative AI? - IBM Research research.ibm.com
- What is Generative AI? - Gen AI Explained - AWS aws.amazon.com
- Explained: Generative AI | MIT News | Massachusetts I... news.mit.edu

See more (5)

Browsing results

- Currently looking for a definition from the search results, including sources like Wikipedia and IBM. Generative AI creates...

How can Grok help?

DeepSearch Think Grok 3 ↑

Perplexity.ai Deep Research

perplexity ←

New Thread ⌂ K

Home

Discover

Spaces

Library

Generative AI

Deep Research Images Sources 39 54 steps ▾

linkedin The Future of Generative AI: Key Trends and Opportunities in...

ciklum Top 10 GenAI Trends to Watch in Technology 2025 and Beyond

research.aimultiple Top 100+ Generative AI Applications with Real-Life... +36 sources

Generative AI: Current State, Applications, and Future Trends in 2025

Generative AI has rapidly evolved into a transformative technology, revolutionizing content creation, business operations, and digital interactions across industries. As of early 2025, this technology has moved beyond experimental phases into mainstream adoption, with McKinsey reporting that 65% of organizations now regularly use generative AI, demonstrating its growing significance in the business landscape ⁴.

Understanding Generative AI

Generative AI refers to a sophisticated branch of artificial intelligence that employs machine learning models to generate new content, such as text, images, and audio, based on patterns learned from large datasets.

Ask follow-up

Deep Research

Networks), which have enabled increasingly sophisticated applications ¹.

Token

Tiktokenizer

gpt-4o

System

You are a helpful assistant



User

Content



Add message

```
<|im_start|>system<|im_sep|>You are a helpful  
assistant<|im_end|><|im_start|>user<|im_sep|><|im_end|>  
<|im_start|>assistant<|im_sep|>
```

Token count

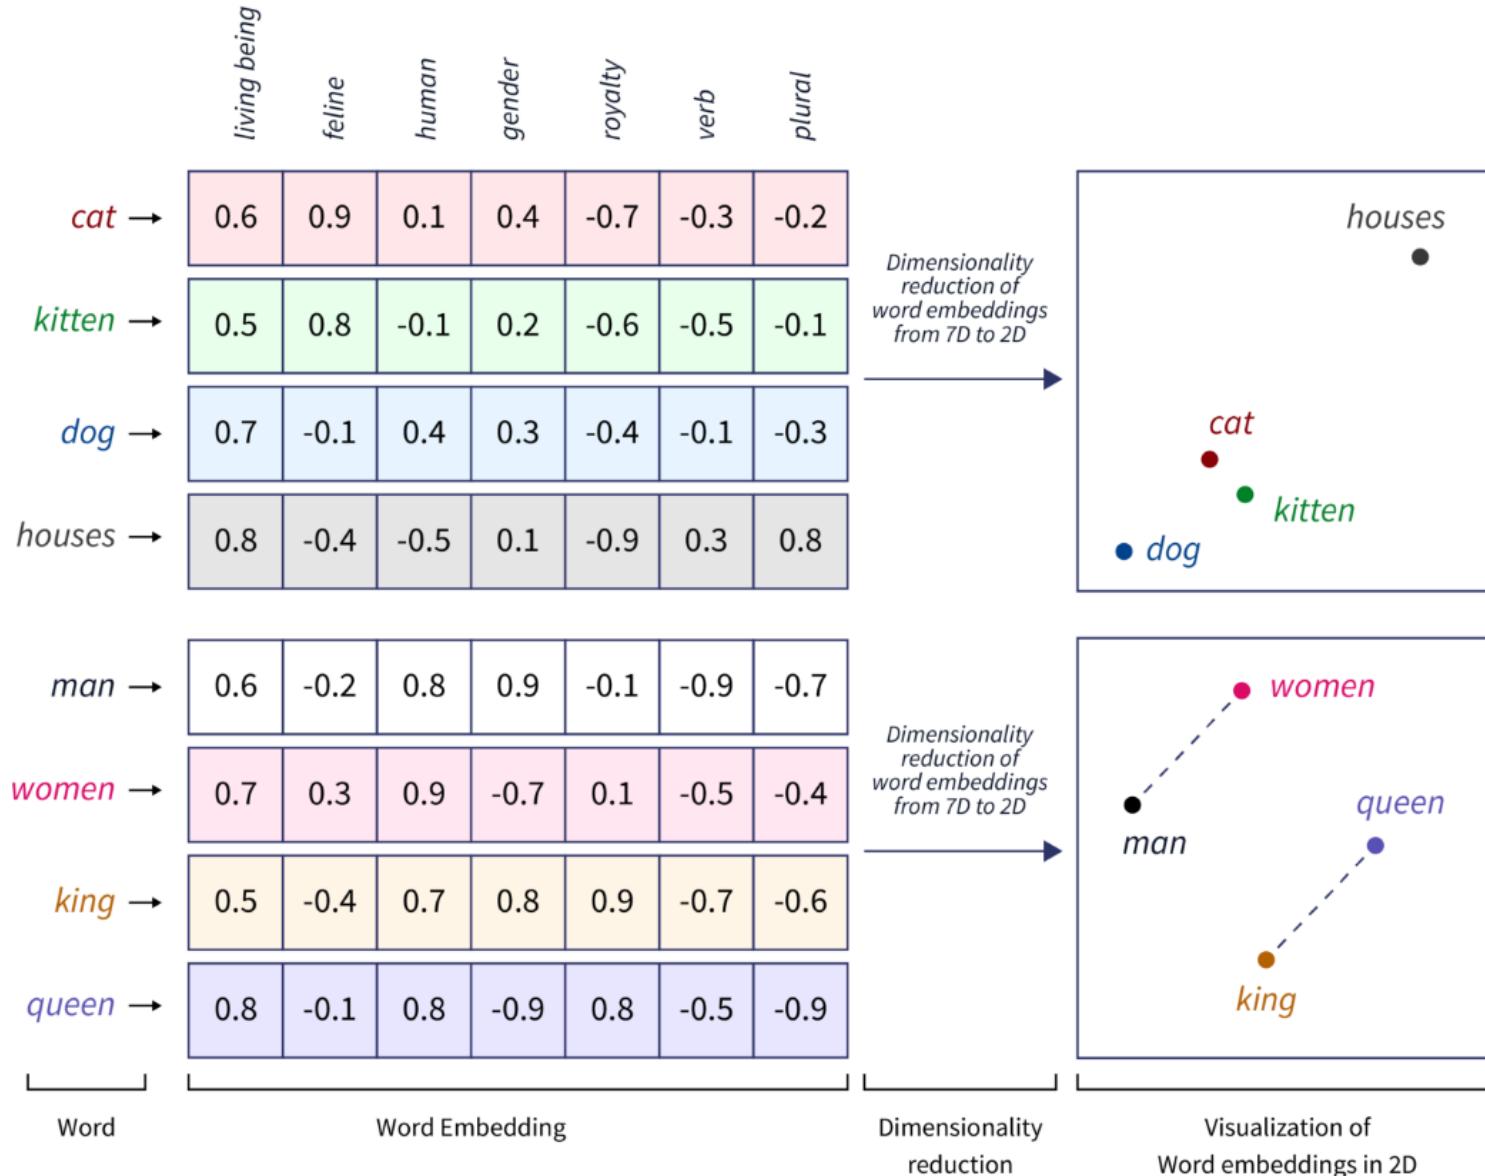
16

```
<|im_start|>system<|im_sep|>You are a helpful assistan  
t<|im_end|><|im_start|>user<|im_sep|><|im_end|><|im_st  
art|>assistant<|im_sep|>
```

```
200264, 17360, 200266, 3575, 553, 261, 10297, 29186, 2  
00265, 200264, 1428, 200266, 200265, 200264, 173781, 2  
00266
```

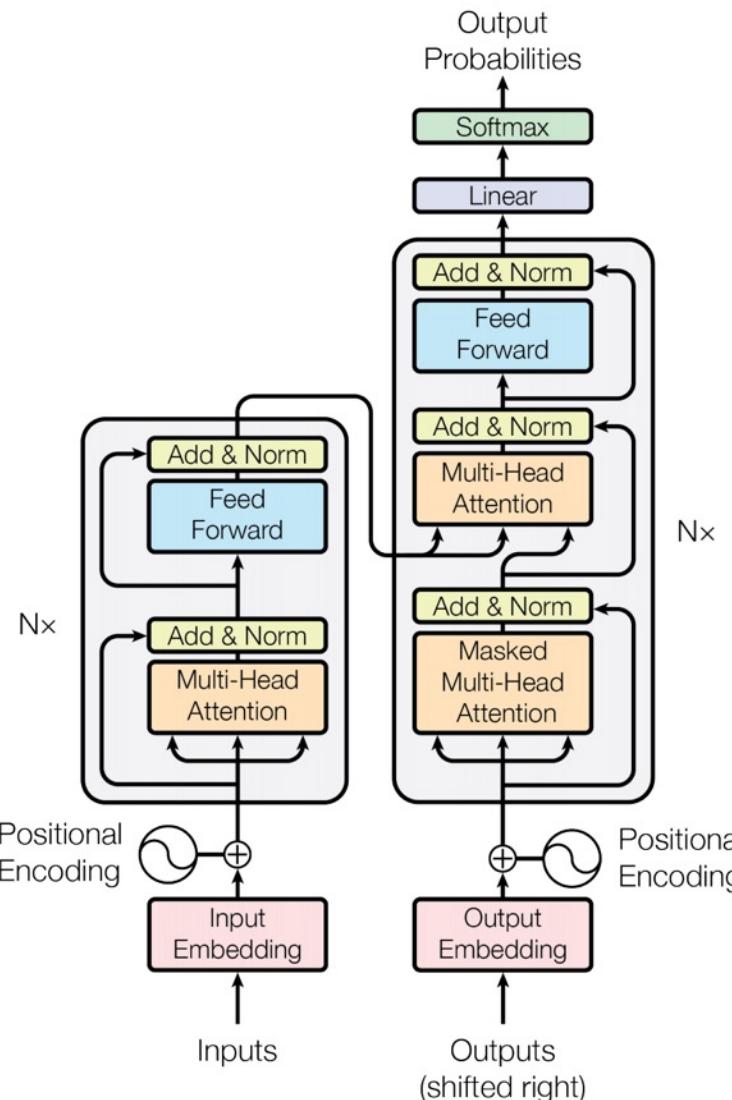
Show whitespace

Word Embeddings



Transformer (Attention is All You Need)

(Vaswani et al., 2017)

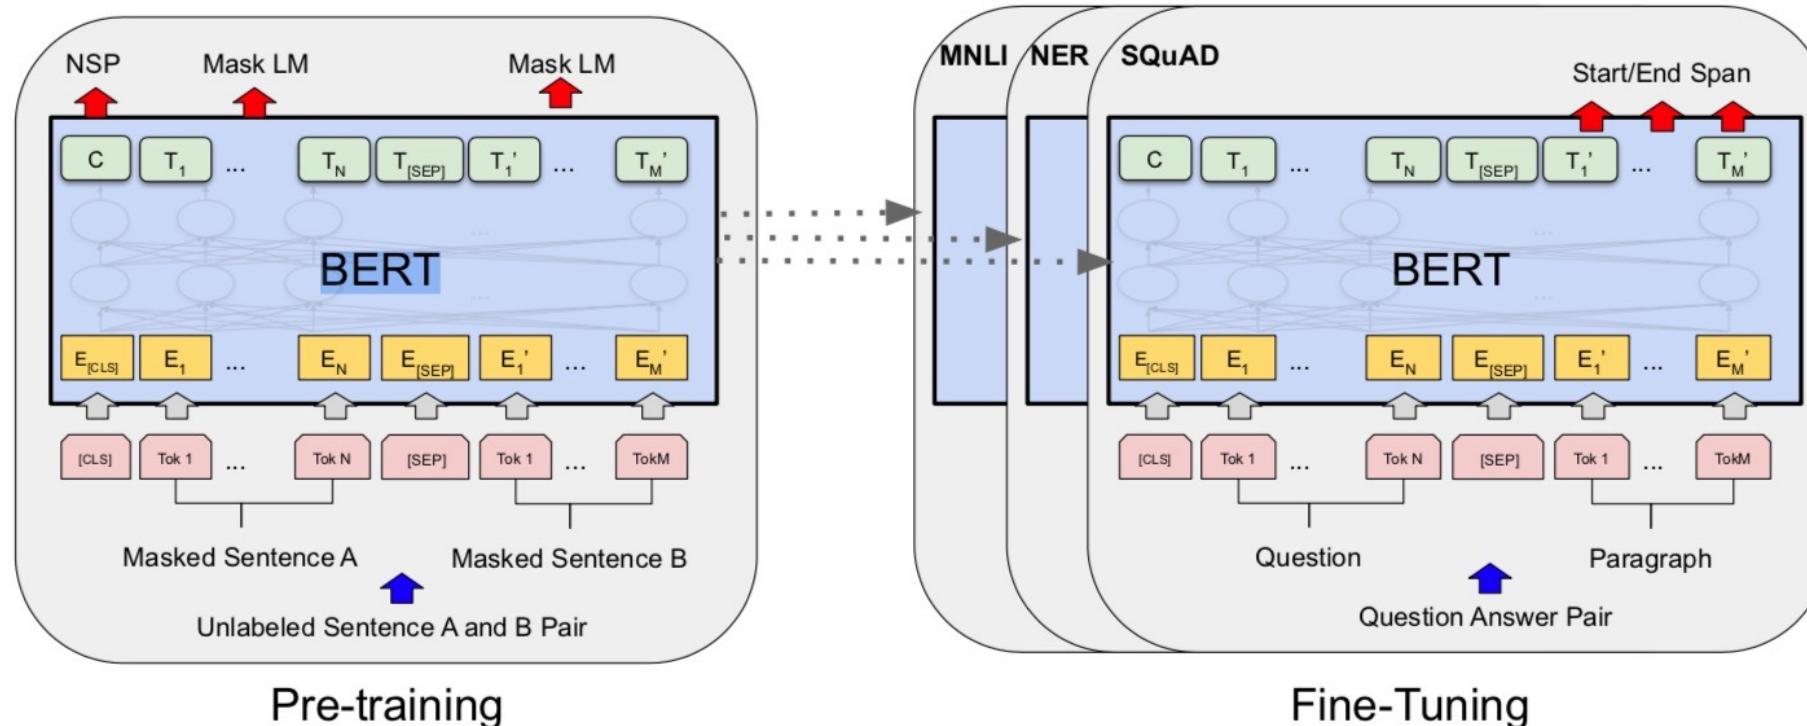


Source: Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin.
"Attention is all you need." In *Advances in neural information processing systems*, pp. 5998-6008. 2017.

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

BERT (Bidirectional Encoder Representations from Transformers)

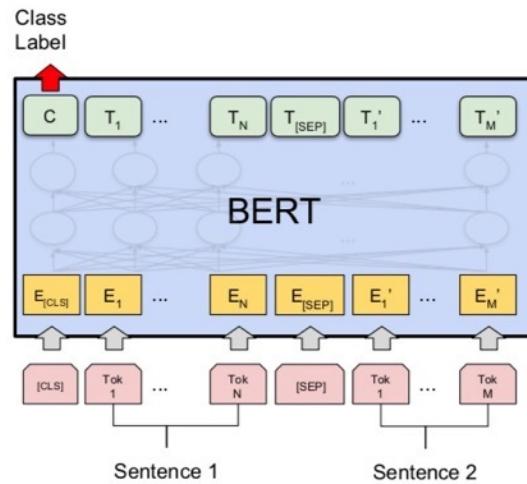
Overall pre-training and fine-tuning procedures for BERT



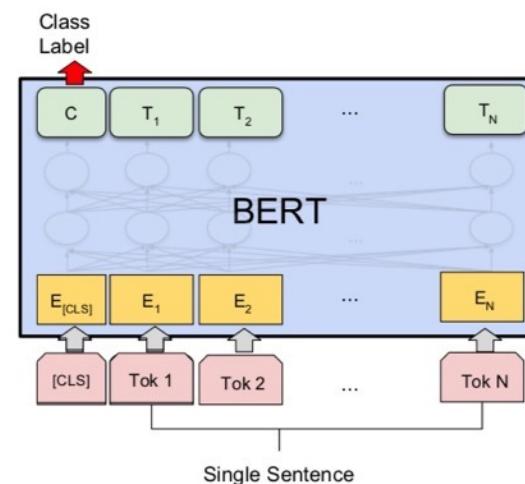
Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).

"Bert: Pre-training of deep bidirectional transformers for language understanding." arXiv preprint arXiv:1810.04805.

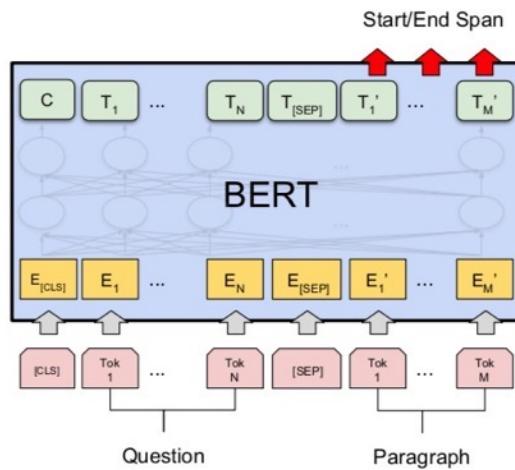
Fine-tuning BERT on Different Tasks



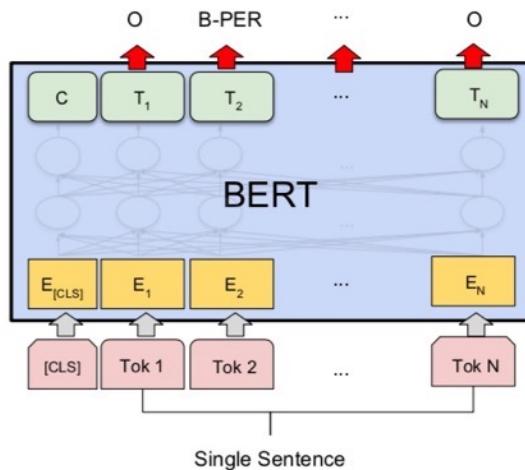
(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



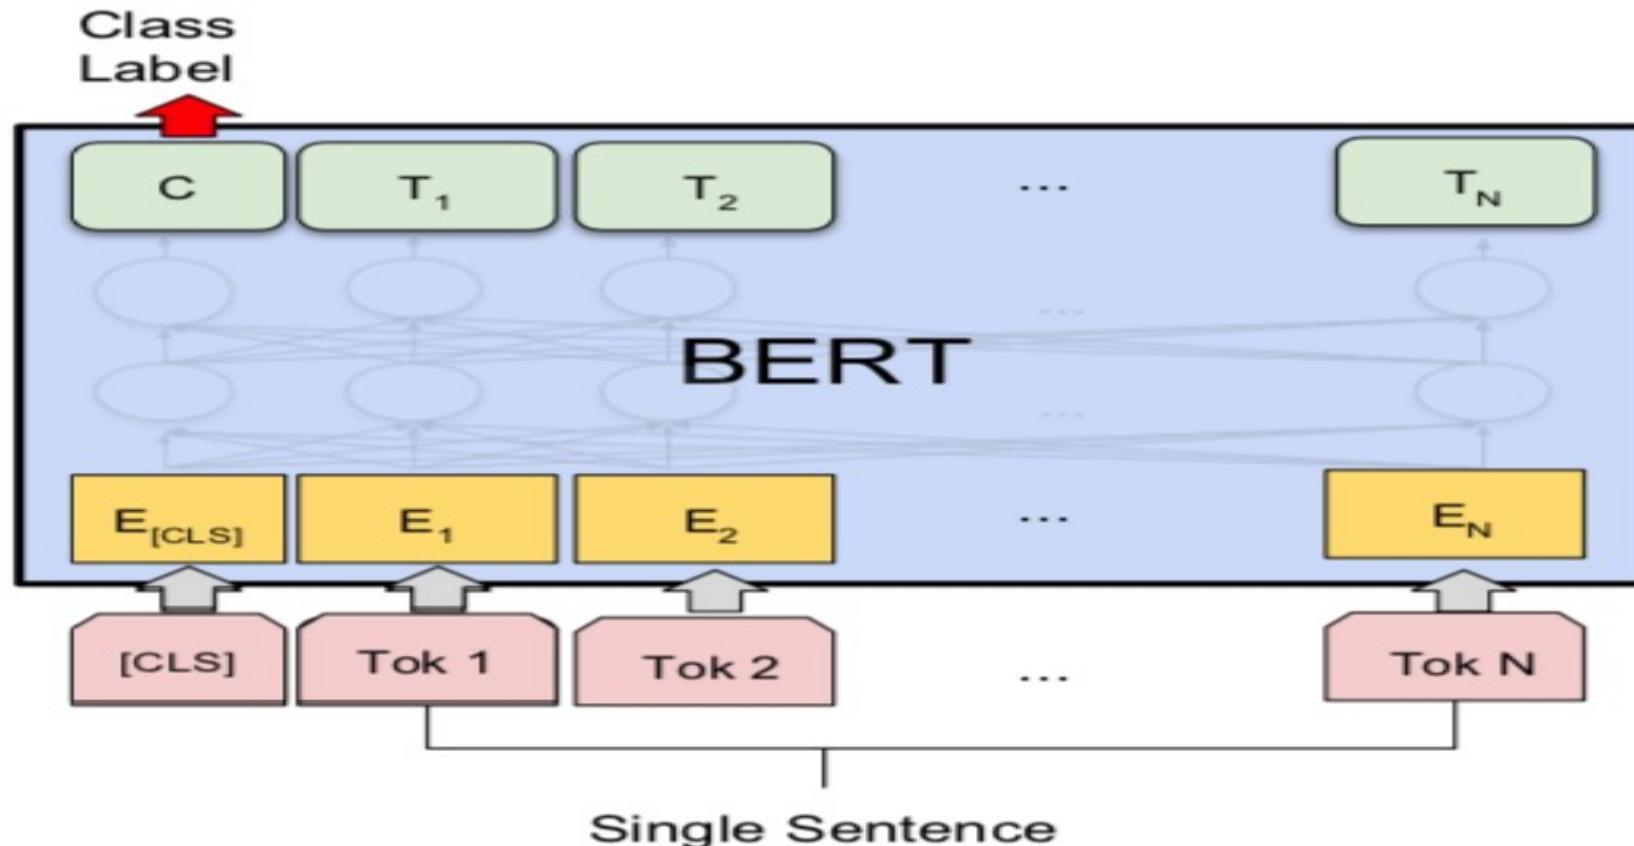
(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).
"Bert: Pre-training of deep bidirectional transformers for language understanding." arXiv preprint arXiv:1810.04805.

Sentiment Analysis: Single Sentence Classification

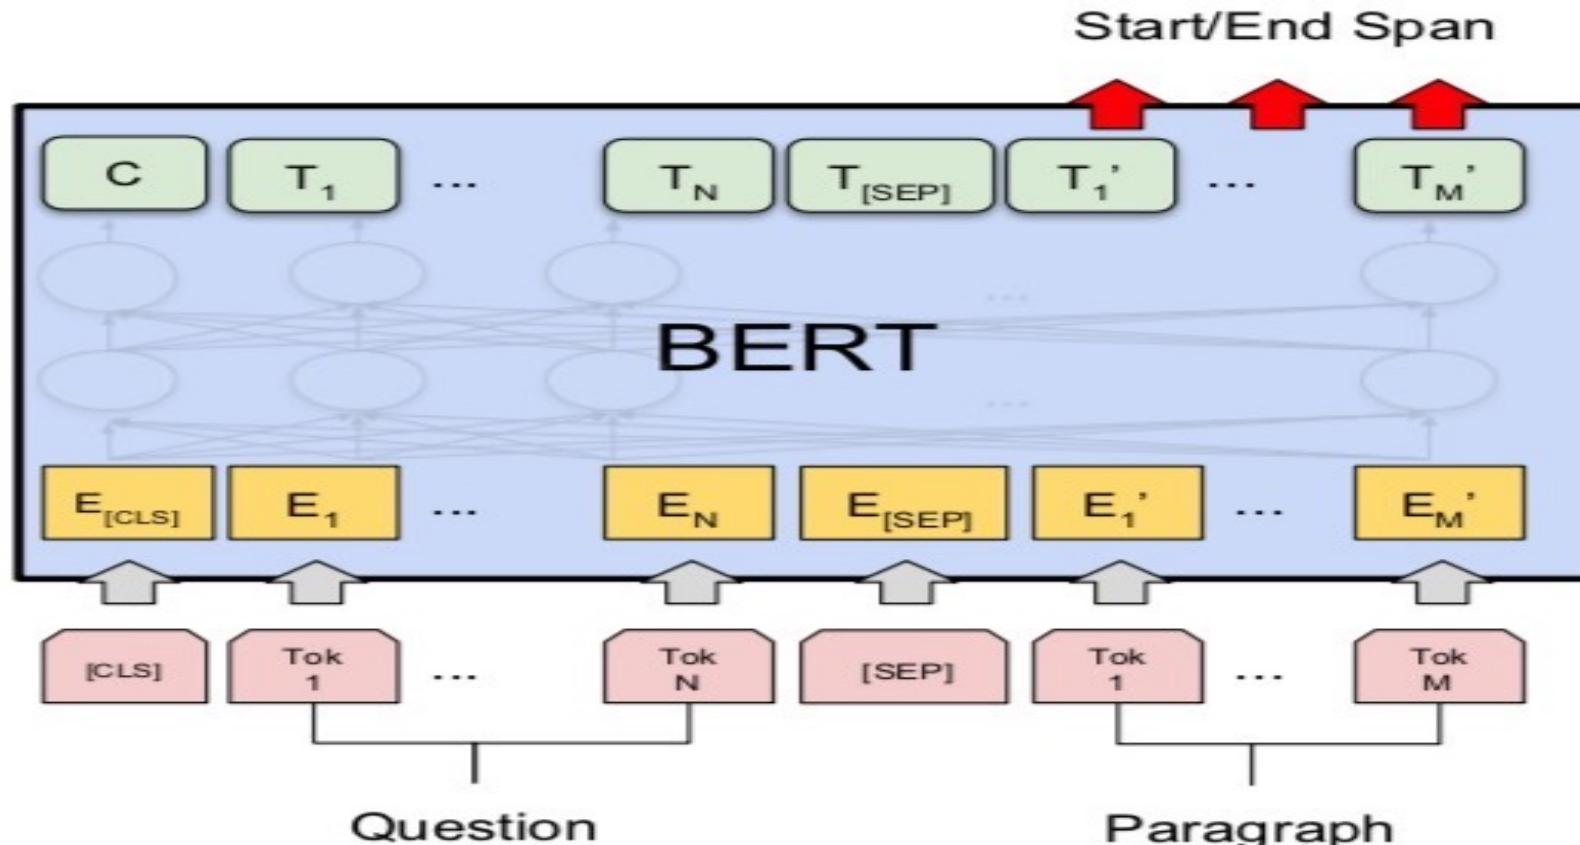


(b) Single Sentence Classification Tasks:
SST-2, CoLA

Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).

"BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." arXiv preprint arXiv:1810.04805

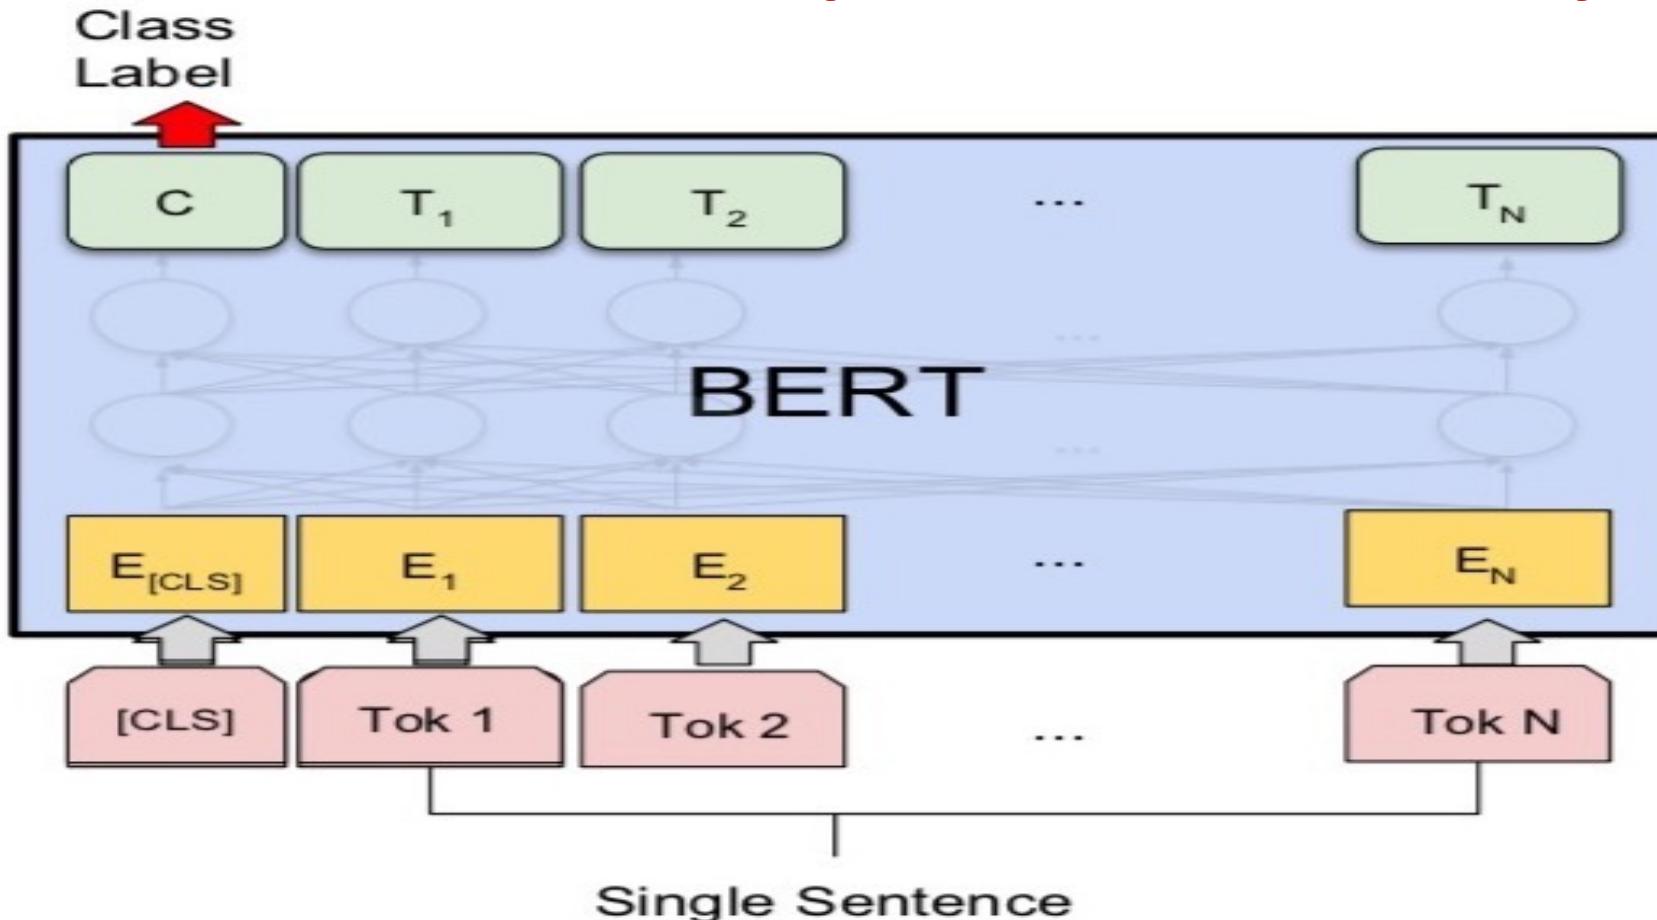
Fine-tuning BERT on Question Answering (QA)



(c) Question Answering Tasks:
SQuAD v1.1

Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).
"Bert: Pre-training of deep bidirectional transformers for language understanding." arXiv preprint arXiv:1810.04805.

Fine-tuning BERT on Dialogue Intent Detection (ID; Classification)

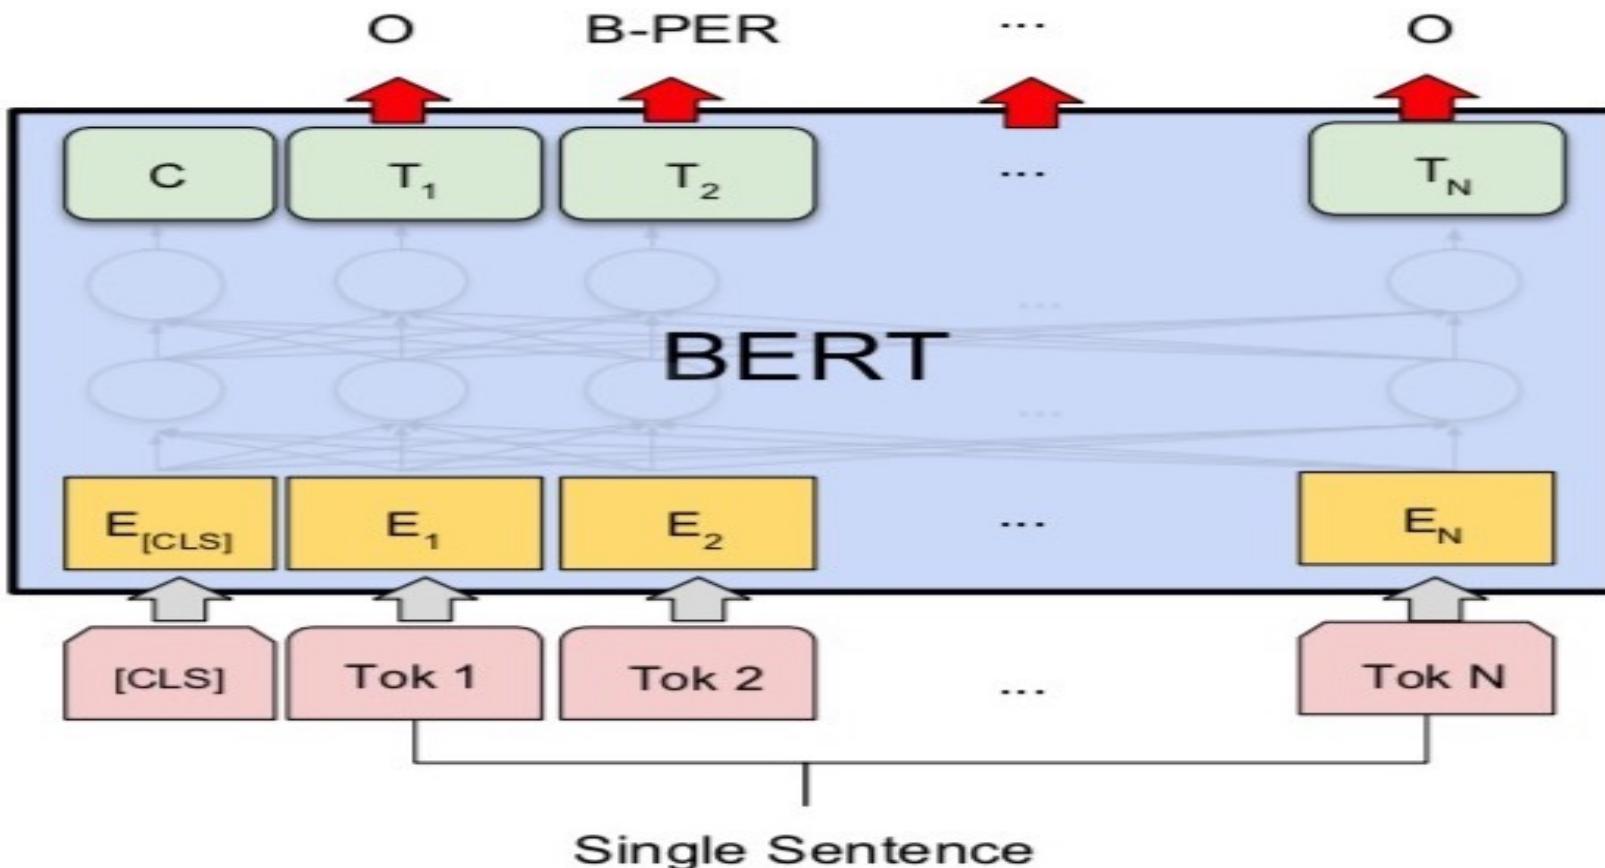


(b) Single Sentence Classification Tasks:
SST-2, CoLA

Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).

"Bert: Pre-training of deep bidirectional transformers for language understanding." arXiv preprint arXiv:1810.04805.

Fine-tuning BERT on Dialogue Slot Filling (SF)



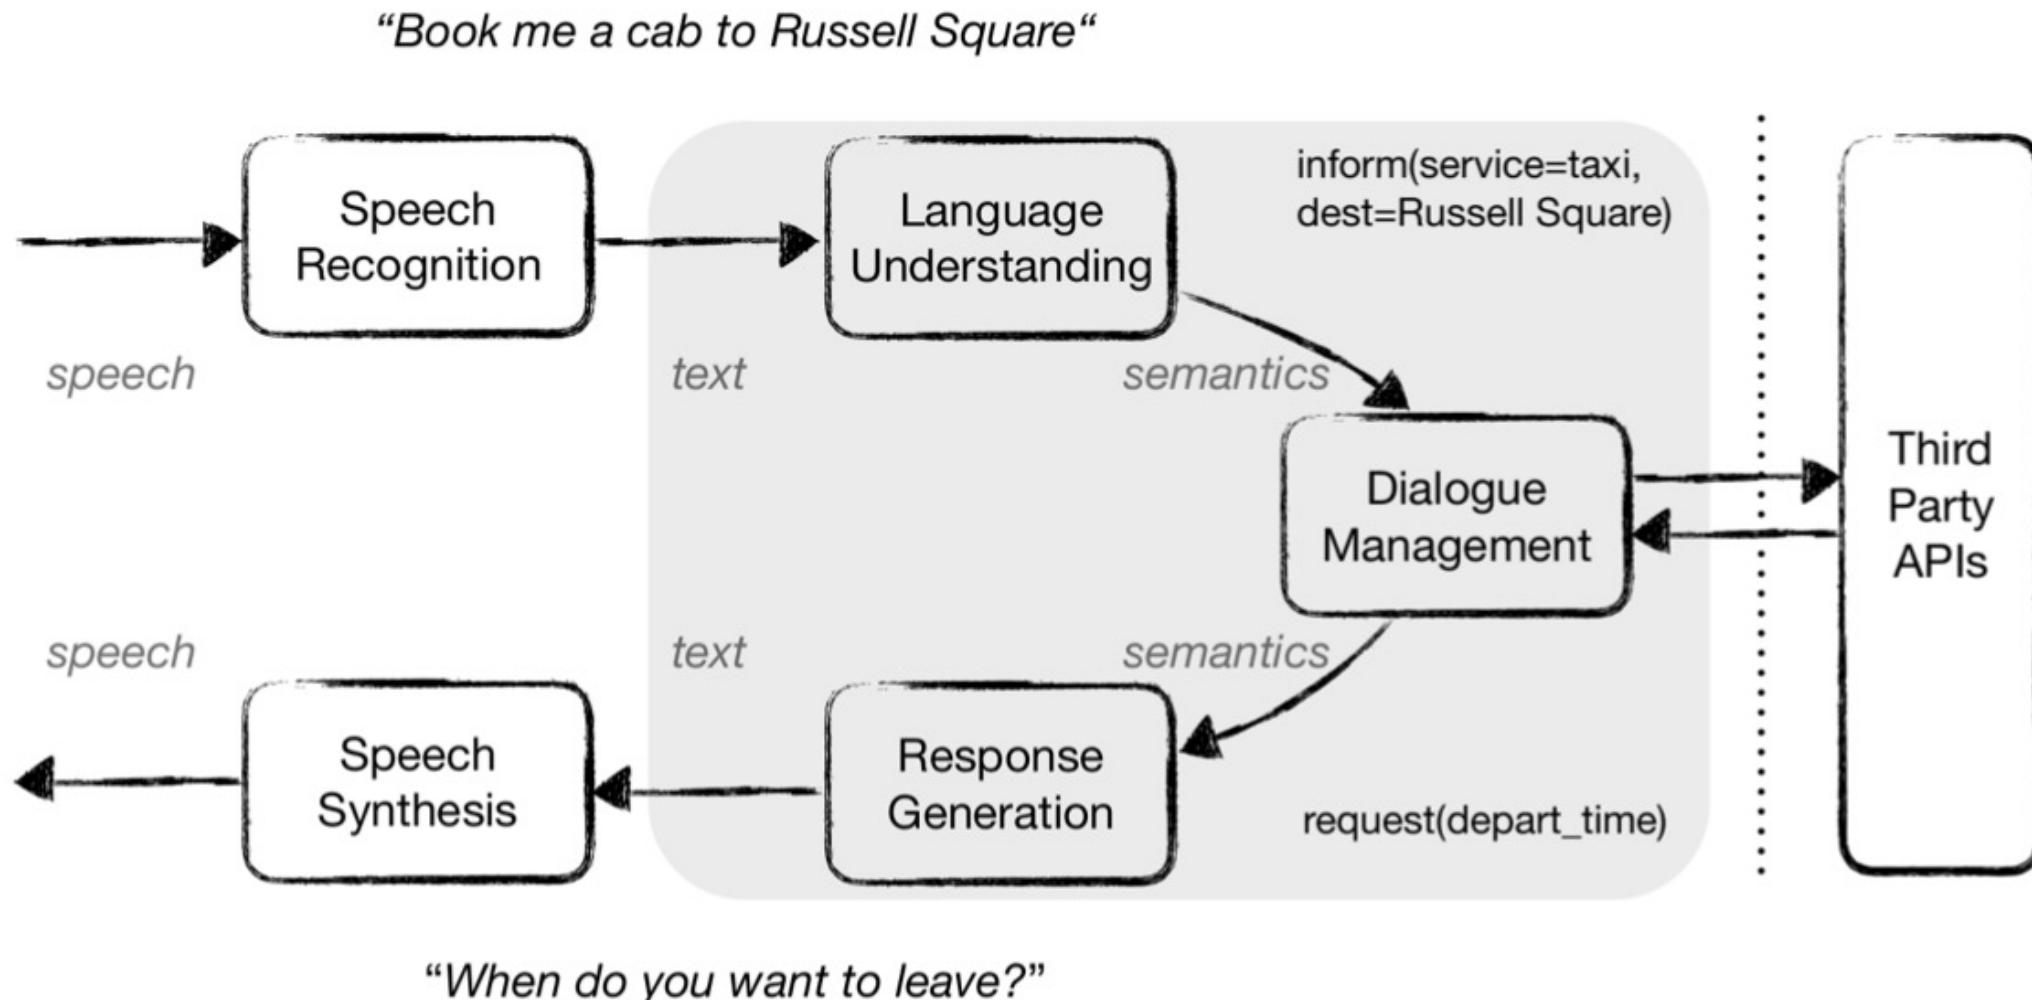
(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).

"Bert: Pre-training of deep bidirectional transformers for language understanding." arXiv preprint arXiv:1810.04805.

Task-Oriented Dialogue (ToD) System

Speech, Text, NLP

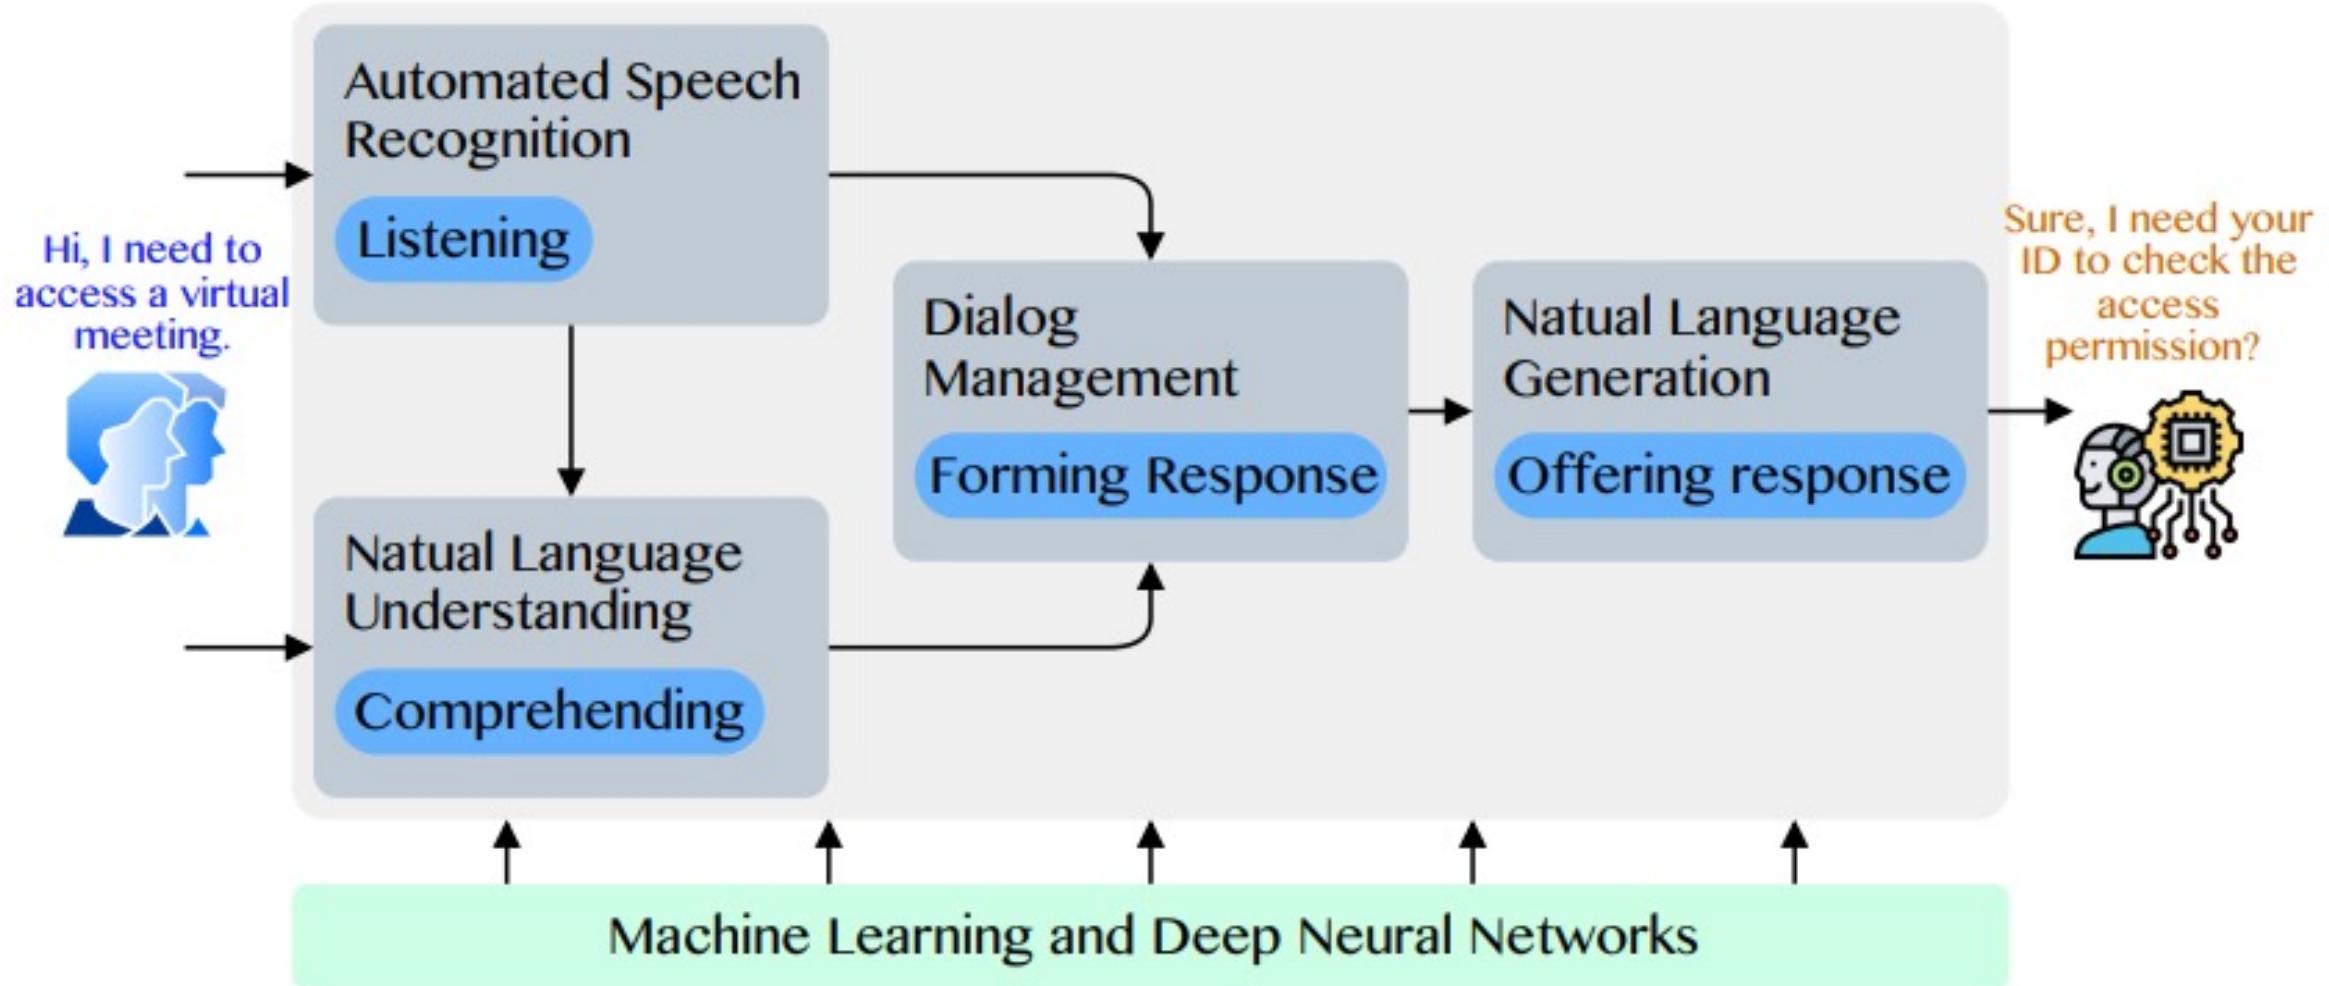


Source: Razumovskia, Evgeniia, Goran Glavas, Olga Majewska, Edoardo M. Ponti, Anna Korhonen, and Ivan Vulic.

“Crossing the conversational chasm: A primer on natural language processing for multilingual task-oriented dialogue systems.” Journal of Artificial Intelligence Research 74 (2022): 1351-1402.

Conversational AI

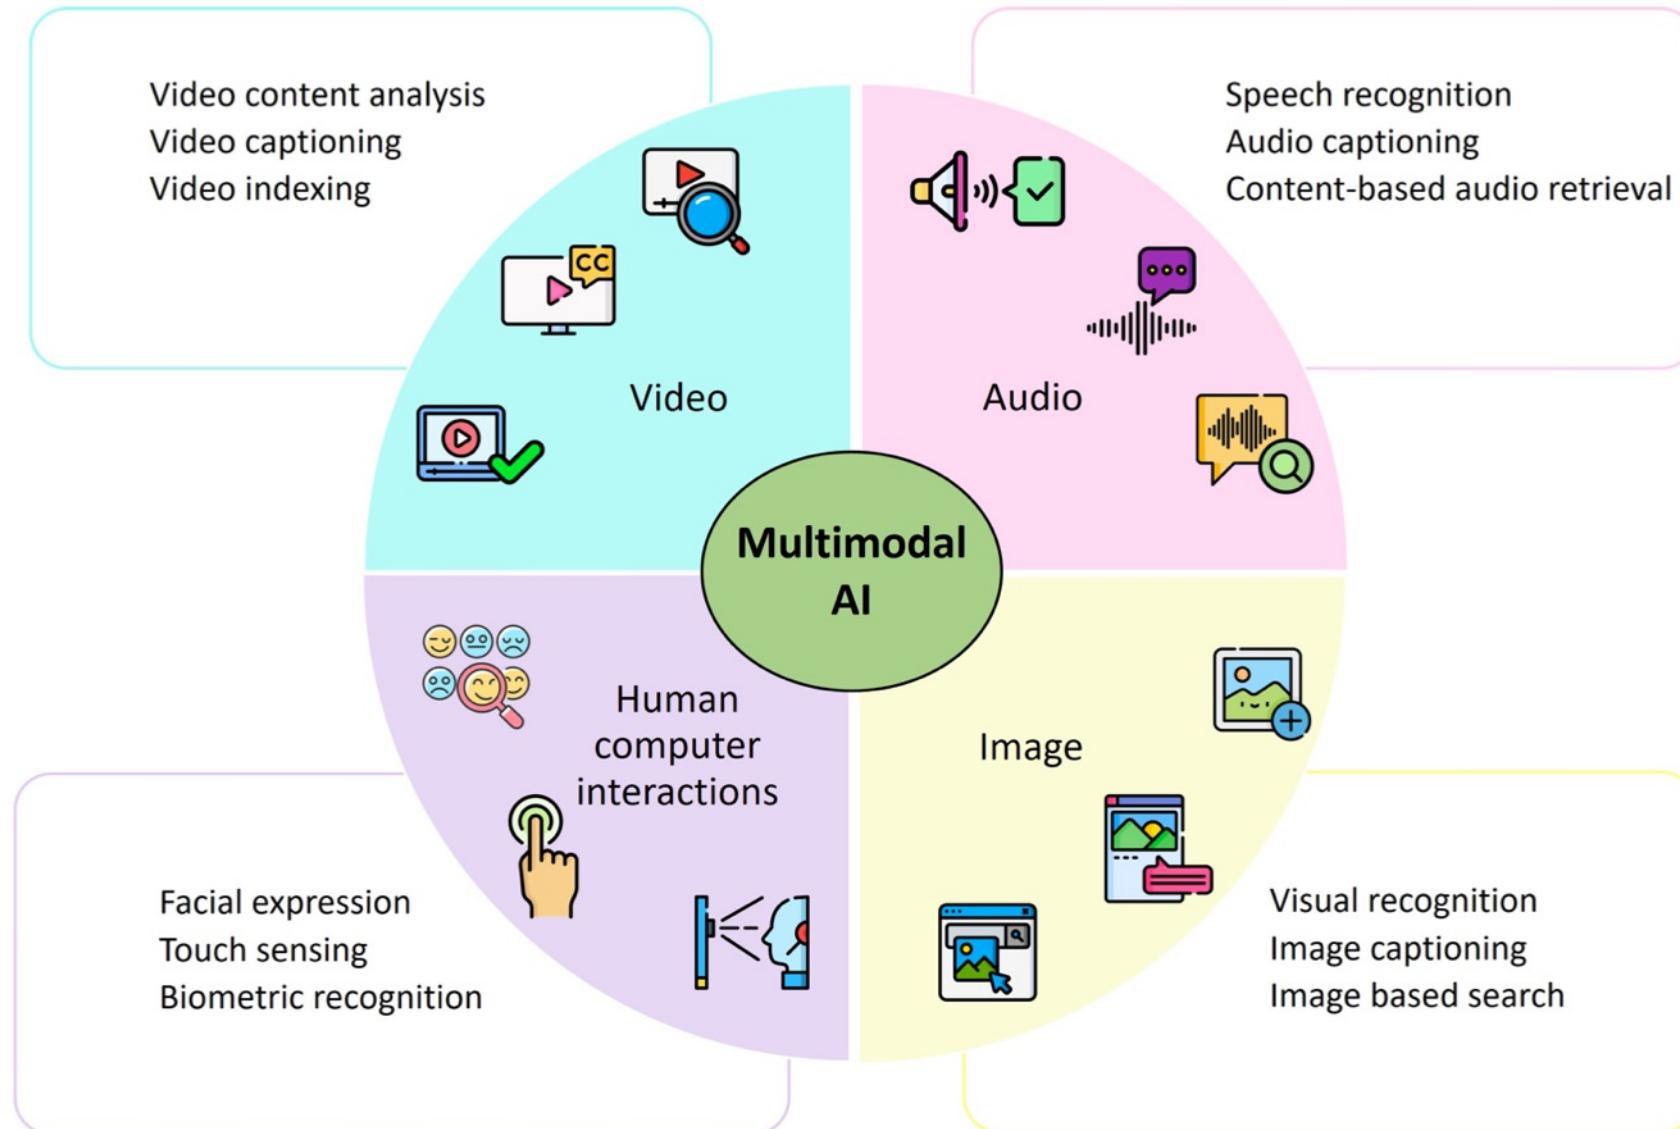
to deliver contextual and personal experience to users



Source: Huynh-The, Thien, Quoc-Viet Pham, Xuan-Qui Pham, Thanh Thi Nguyen, Zhu Han, and Dong-Seong Kim (2022).

"Artificial Intelligence for the Metaverse: A Survey." arXiv preprint arXiv:2202.10336.

Technological Integration for Multimodal AI



4 Approaches of AI

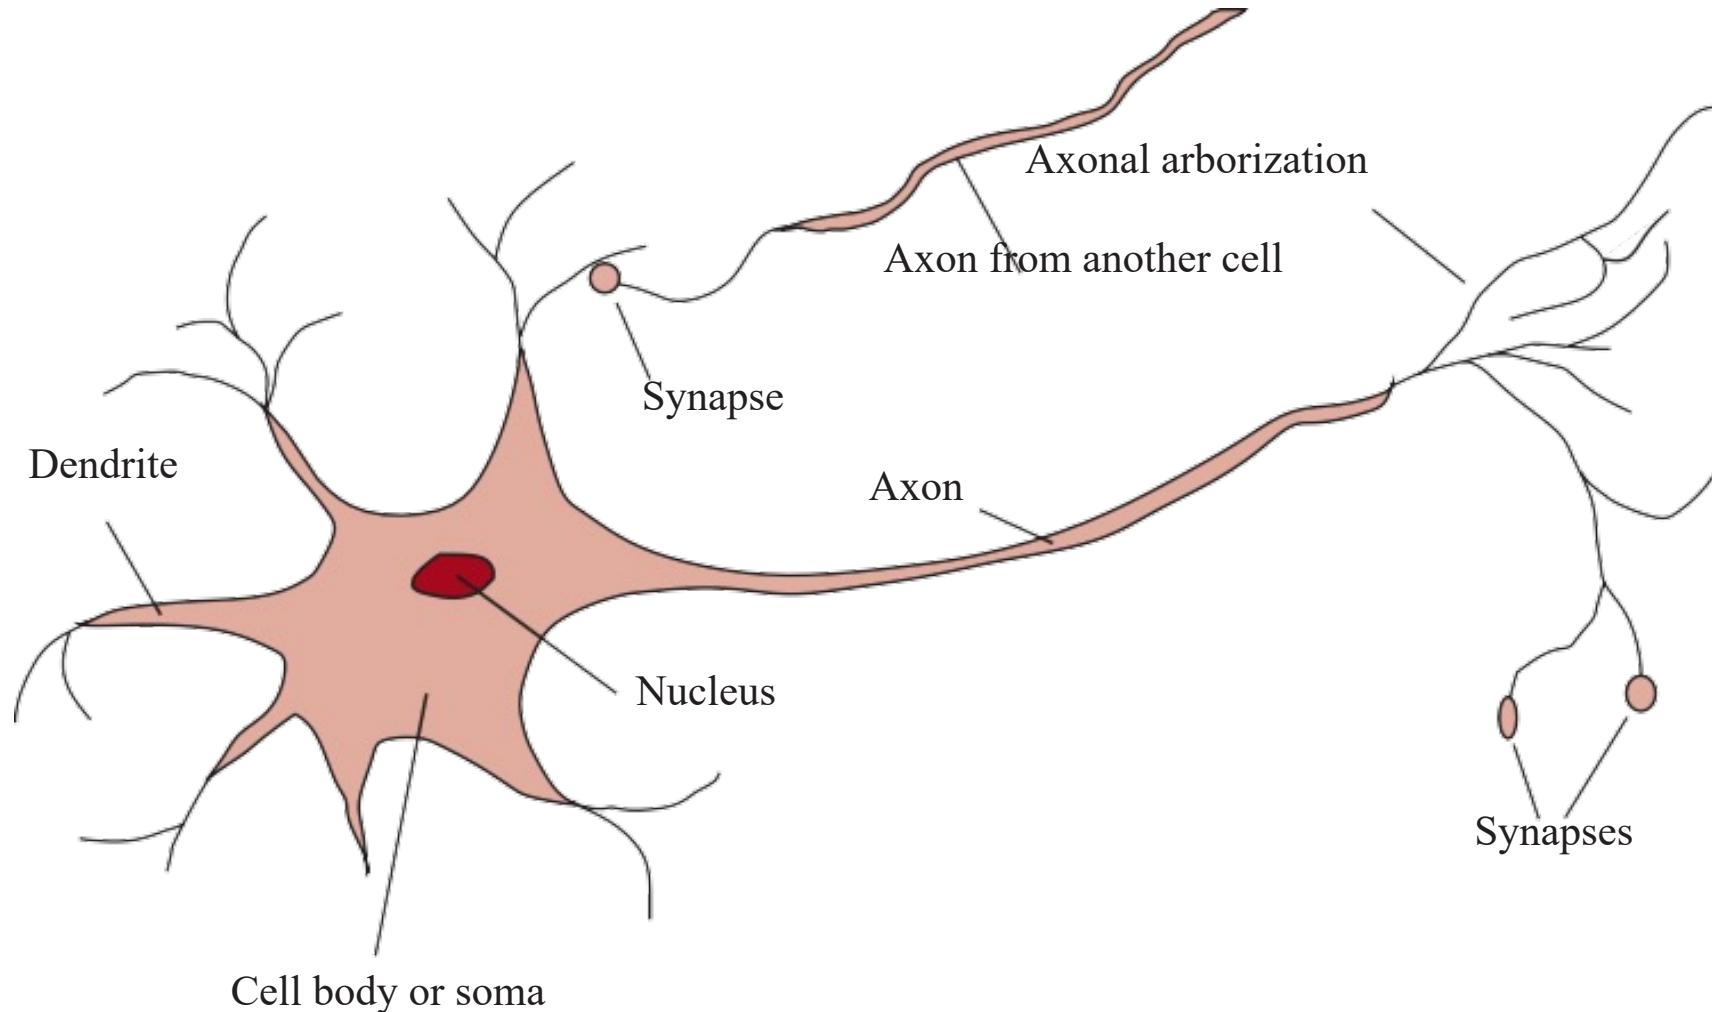
<p>2. Thinking Humanly: The Cognitive Modeling Approach</p>	<p>3. Thinking Rationally: The “Laws of Thought” Approach</p>
<p>1. Acting Humanly: The Turing Test Approach <small>(1950)</small></p>	<p>4. Acting Rationally: The Rational Agent Approach</p>

Acting Rationally: The Rational Agent Approach

- AI has focused on the study and construction of agents that **do the right thing**.
- **Standard model**

Neuroscience

The parts of a **nerve cell** or **neuron**

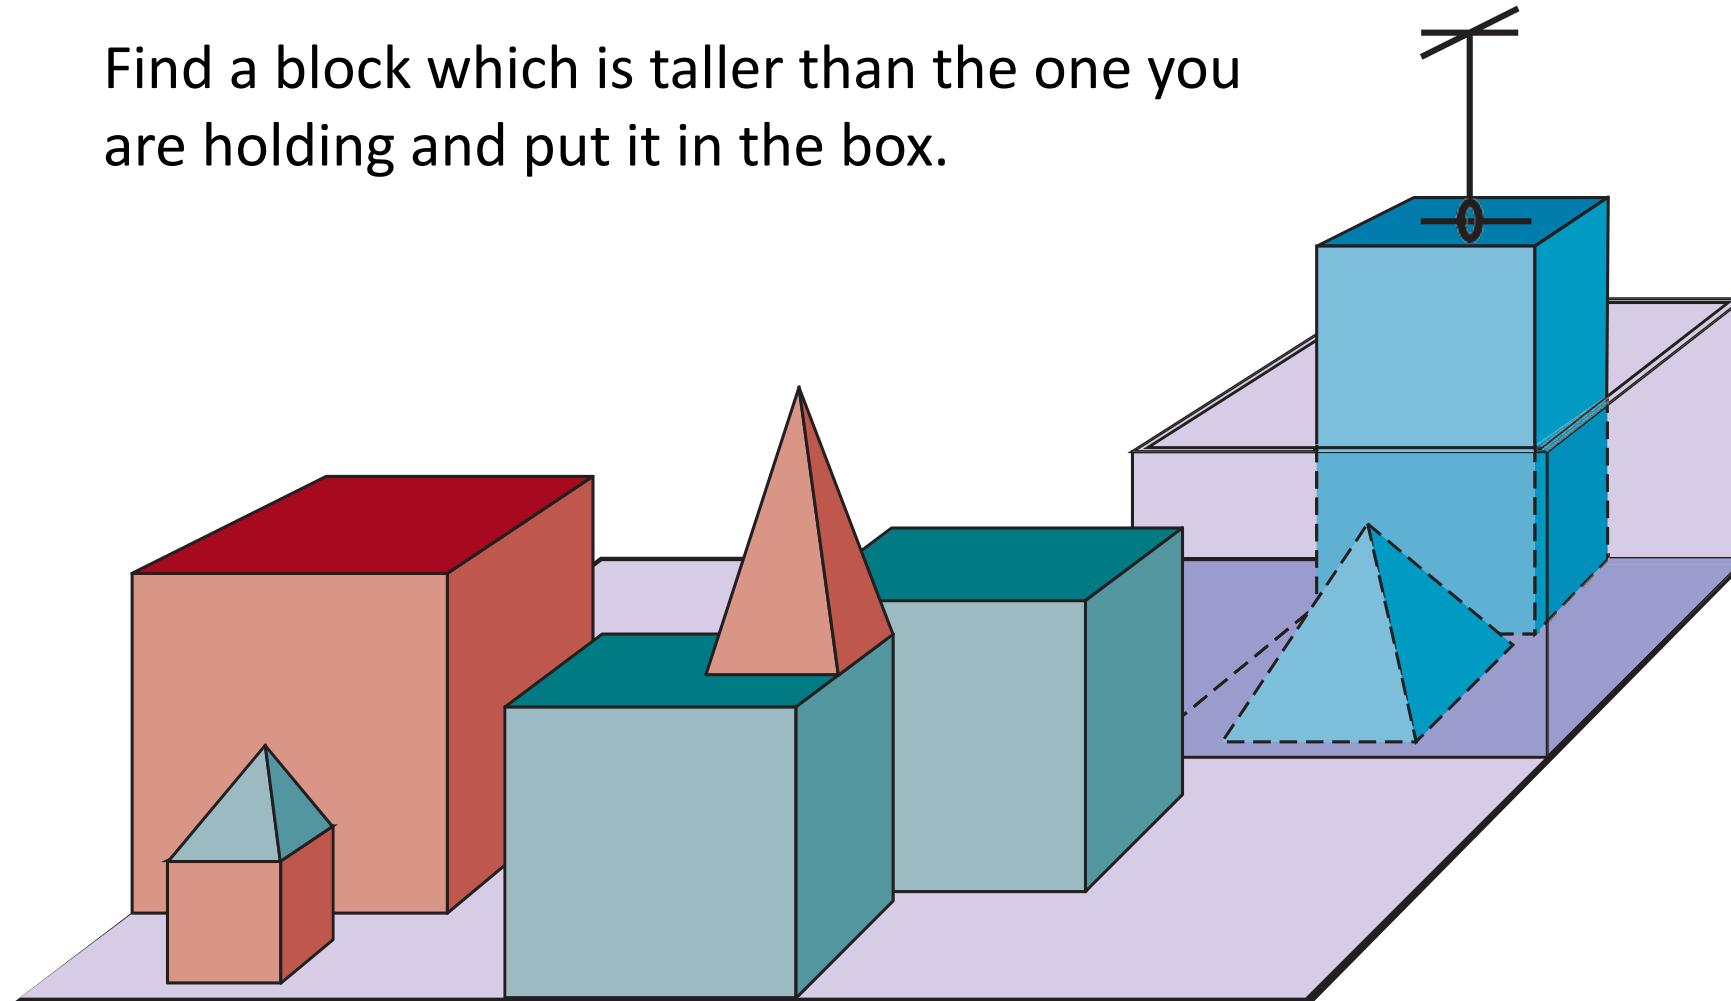


Comparison of Computer and Human Brain

	Supercomputer	Personal Computer	Human Brain
Computational units	10^6 GPUs + CPUs	8 CPU cores	10^6 columns
	10^{15} transistors	10^{10} transistors	10^{11} neurons
Storage units	10^{16} bytes RAM	10^{10} bytes RAM	10^{11} neurons
	10^{17} bytes disk	10^{12} bytes disk	10^{14} synapses
Cycle time	10^{-9} sec	10^{-9} sec	10^{-3} sec
Operations/sec	10^{18}	10^{10}	10^{17}

A scene from the blocks world

Find a block which is taller than the one you are holding and put it in the box.

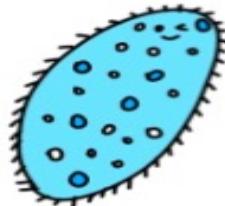


Intelligent Agents

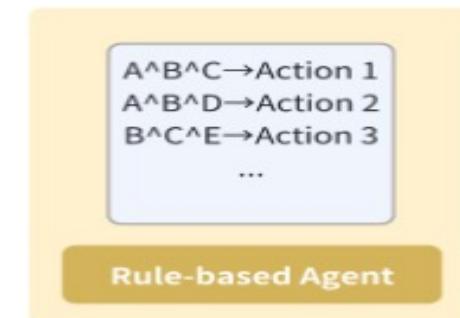
4 Approaches of AI

<p>2. Thinking Humanly: The Cognitive Modeling Approach</p>	<p>3. Thinking Rationally: The “Laws of Thought” Approach</p>
<p>1. Acting Humanly: The Turing Test Approach <small>(1950)</small></p>	<p>4. Acting Rationally: The Rational Agent Approach</p>

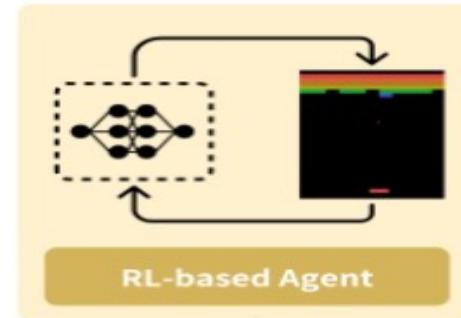
Intelligent Agents Roadmap



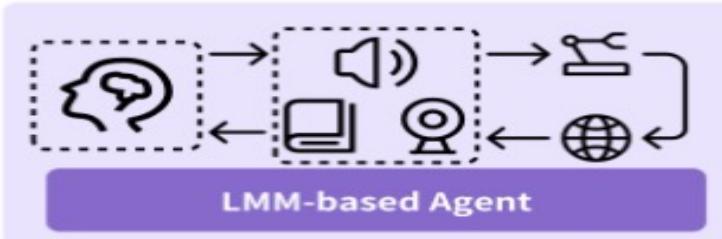
Reflex Agent ►►



►► Learning-based Agent

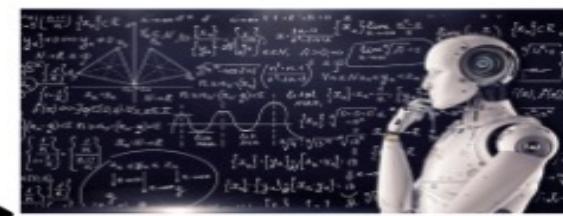


RL-based Agent



LLM-based Agent

<< Large Model-based Agent

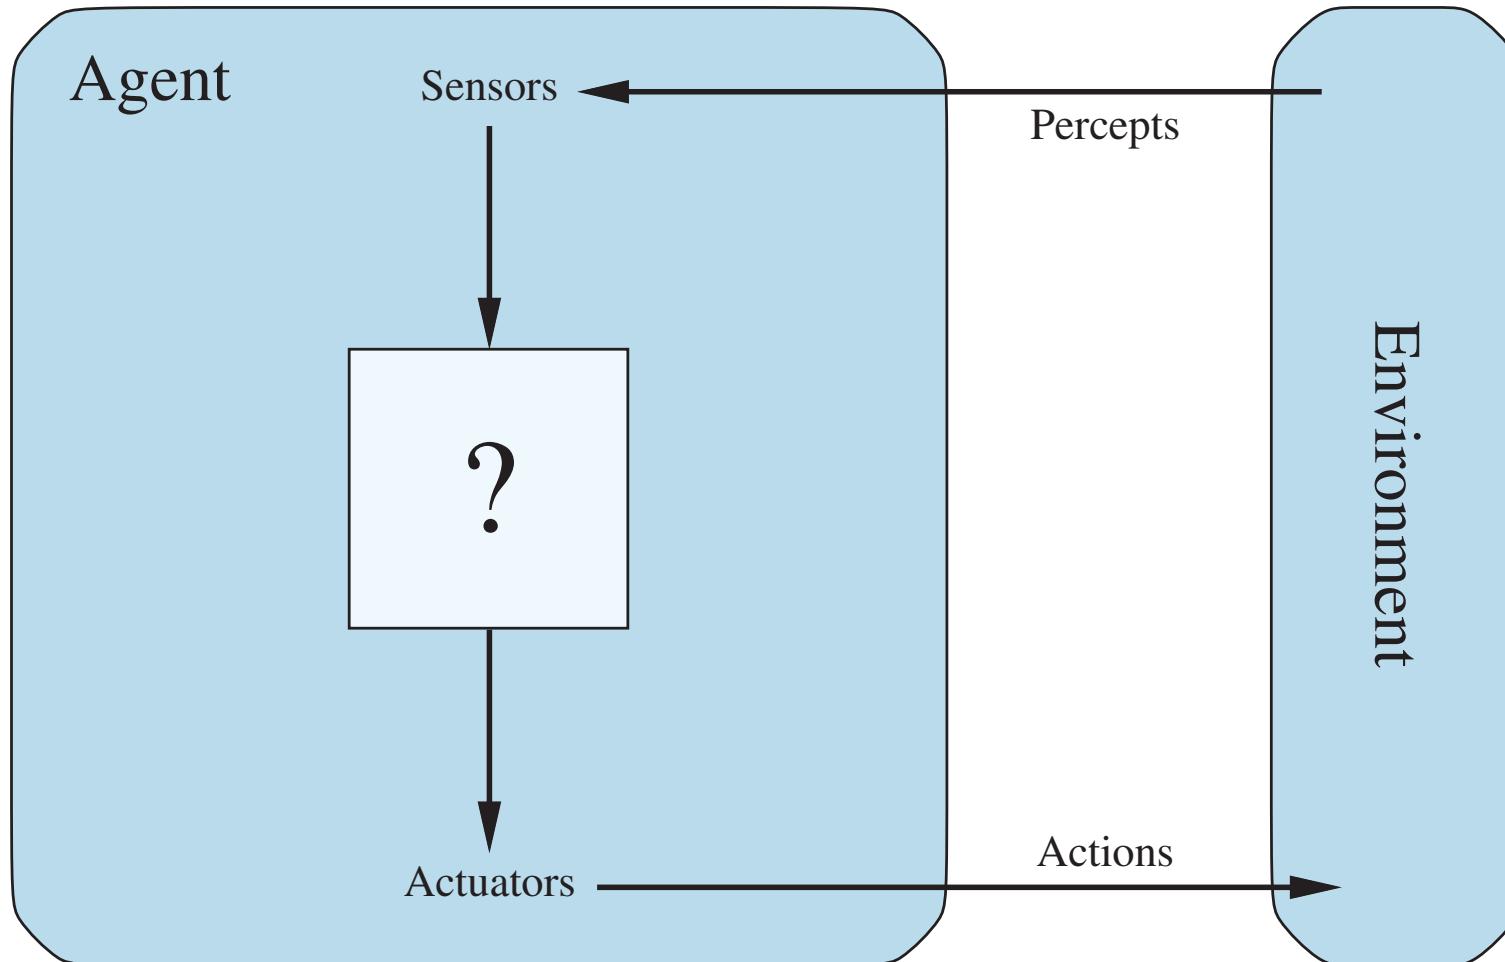


AGI Agent ►►

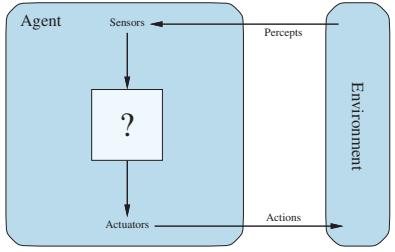
AI Agents

- Traditional AI Agents
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents
 - Learning agents
- Evolution of AI Agents
 - LLM-based Agents
 - Multi-modal agents
 - Embodied AI agents in virtual environments
 - Collaborative AI agents

Agents interact with environments through sensors and actuators



AI Agents



- **Definition:** An **AI agent** is an **entity** that **perceives** its **environment** and takes **actions** to achieve **goals**
- **Components:**
 1. **Sensors:** Perceive the environment
 2. **Actuators:** Act upon the environment
 3. **Decision-making mechanism:** Process inputs and decide on actions

Reinforcement Learning (DL)

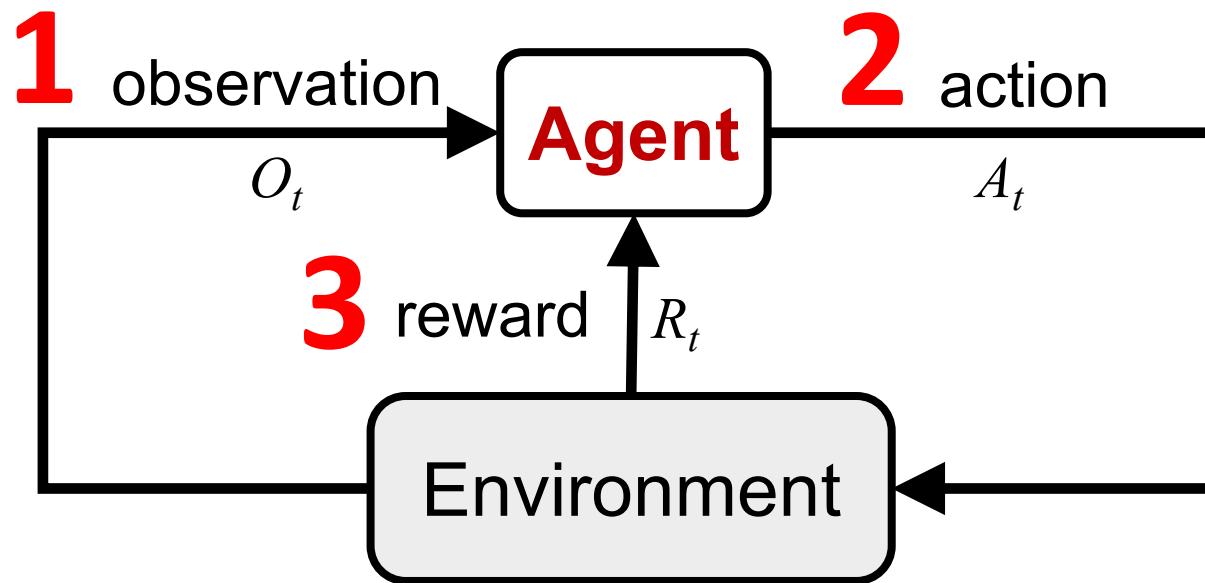
Agent

Environment

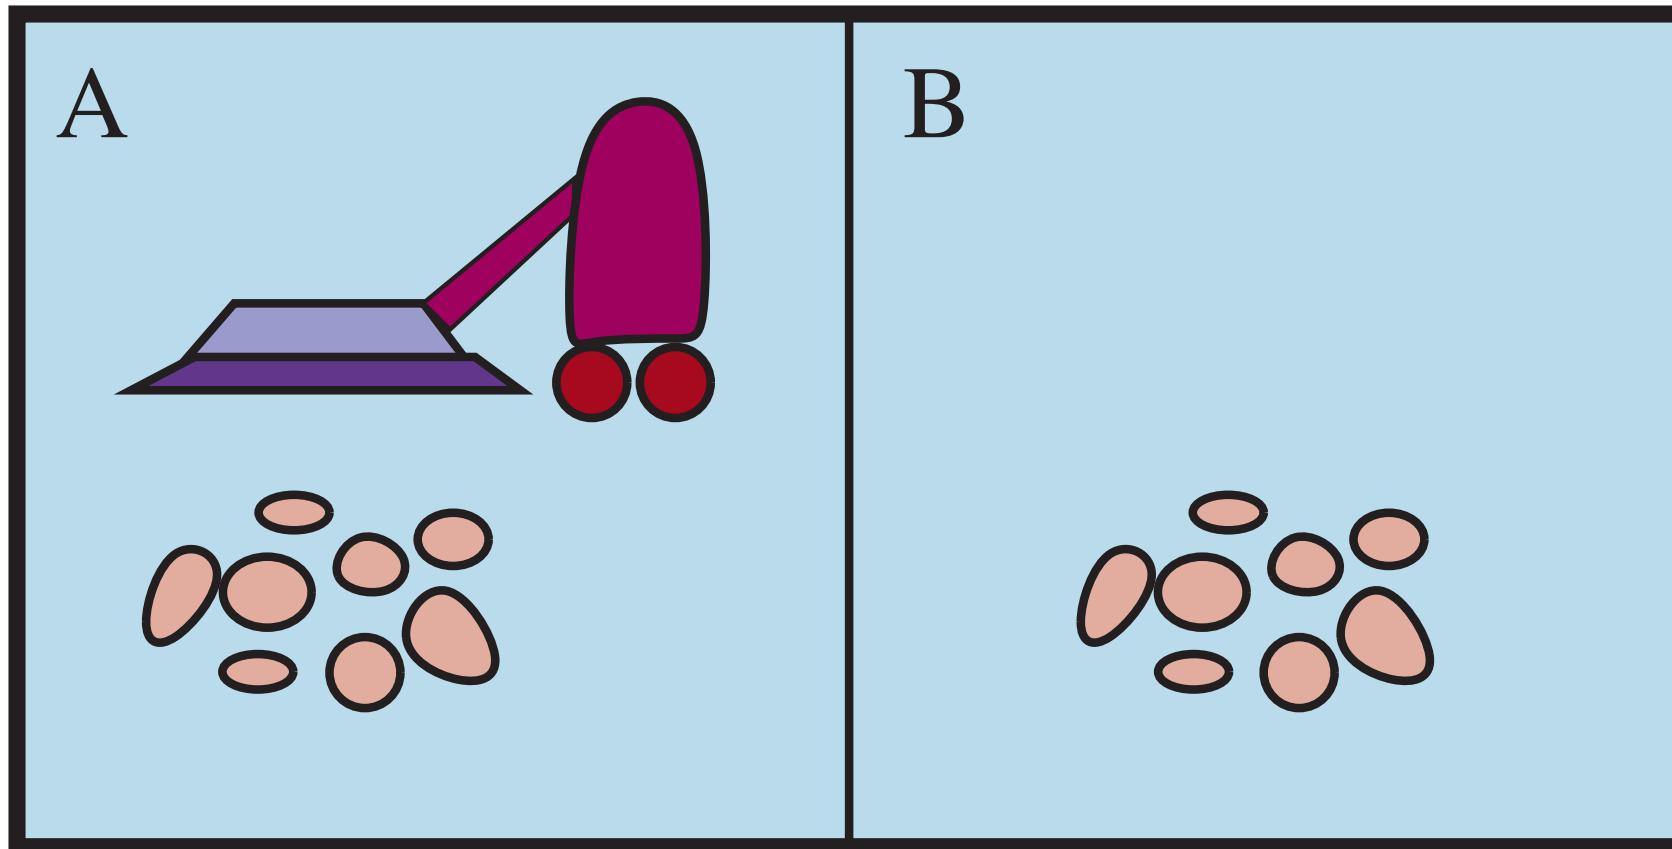
Reinforcement Learning (DL)



Reinforcement Learning (DL)



A vacuum-cleaner world with just two locations



Partial tabulation of a simple agent function for the vacuum-cleaner world

Percept sequence	Action
$[A, Clean]$	<i>Right</i>
$[A, Dirty]$	<i>Suck</i>
$[B, Clean]$	<i>Left</i>
$[B, Dirty]$	<i>Suck</i>
$[A, Clean], [A, Clean]$	<i>Right</i>
$[A, Clean], [A, Dirty]$	<i>Suck</i>
:	:
$[A, Clean], [A, Clean], [A, Clean]$	<i>Right</i>
$[A, Clean], [A, Clean], [A, Dirty]$	<i>Suck</i>
:	:

PEAS description of the task environment for an automated taxi driver

Agent Type	Performance Measure	Environment	Actuators	Sensors
Taxi driver	Safe, fast, legal, comfortable trip, maximize profits, minimize impact on other road users	Roads, other traffic, police, pedestrians, customers, weather	Steering, accelerator, brake, signal, horn, display, speech	Cameras, radar, speedometer, GPS, engine sensors, accelerometer, microphones, touchscreen

Examples of Agent Types and their PEAS descriptions

Agent Type	Performance Measure	Environment	Actuators	Sensors
Medical diagnosis system	Healthy patient, reduced costs	Patient, hospital, staff	Display of questions, tests, diagnoses, treatments	Touchscreen/voice entry of symptoms and findings
Satellite image analysis system	Correct categorization of objects, terrain	Orbiting satellite, downlink, weather	Display of scene categorization	High-resolution digital camera
Part-picking robot	Percentage of parts in correct bins	Conveyor belt with parts; bins	Jointed arm and hand	Camera, tactile and joint angle sensors
Refinery controller	Purity, yield, safety	Refinery, raw materials, operators	Valves, pumps, heaters, stirrers, displays	Temperature, pressure, flow, chemical sensors
Interactive English tutor	Student's score on test	Set of students, testing agency	Display of exercises, feedback, speech	Keyboard entry, voice

Examples of Task Environments and their Characteristics

Task Environment	Observable	Agents	Deterministic	Episodic	Static	Discrete
Crossword puzzle	Fully	Single	Deterministic	Sequential	Static	Discrete
Chess with a clock	Fully	Multi	Deterministic	Sequential	Semi	Discrete
Poker	Partially	Multi	Stochastic	Sequential	Static	Discrete
Backgammon	Fully	Multi	Stochastic	Sequential	Static	Discrete
Taxi driving	Partially	Multi	Stochastic	Sequential	Dynamic	Continuous
Medical diagnosis	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
Image analysis	Fully	Single	Deterministic	Episodic	Semi	Continuous
Part-picking robot	Partially	Single	Stochastic	Episodic	Dynamic	Continuous
Refinery controller	Partially	Single	Stochastic	Sequential	Dynamic	Continuous
English tutor	Partially	Multi	Stochastic	Sequential	Dynamic	Discrete

The TABLE-DRIVEN-AGENT program is invoked for each new percept and returns an action each time.

It retains the complete percept sequence in memory.

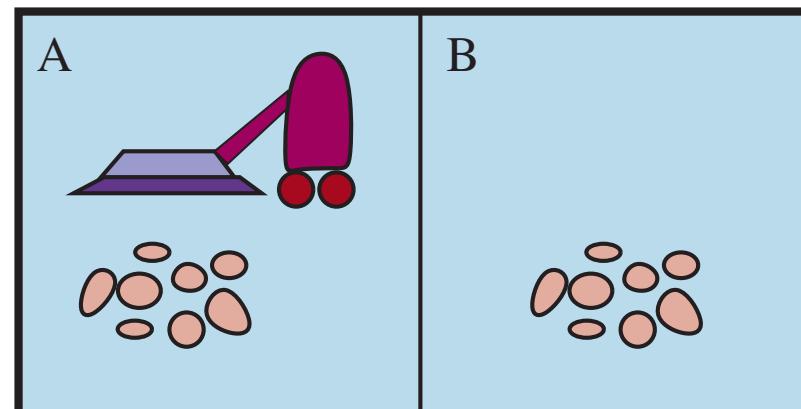
```
function TABLE-DRIVEN-AGENT(percept) returns an action
    persistent: percepts, a sequence, initially empty
                table, a table of actions, indexed by percept sequences, initially fully specified

    append percept to the end of percepts
    action  $\leftarrow$  LOOKUP(percepts, table)
    return action
```

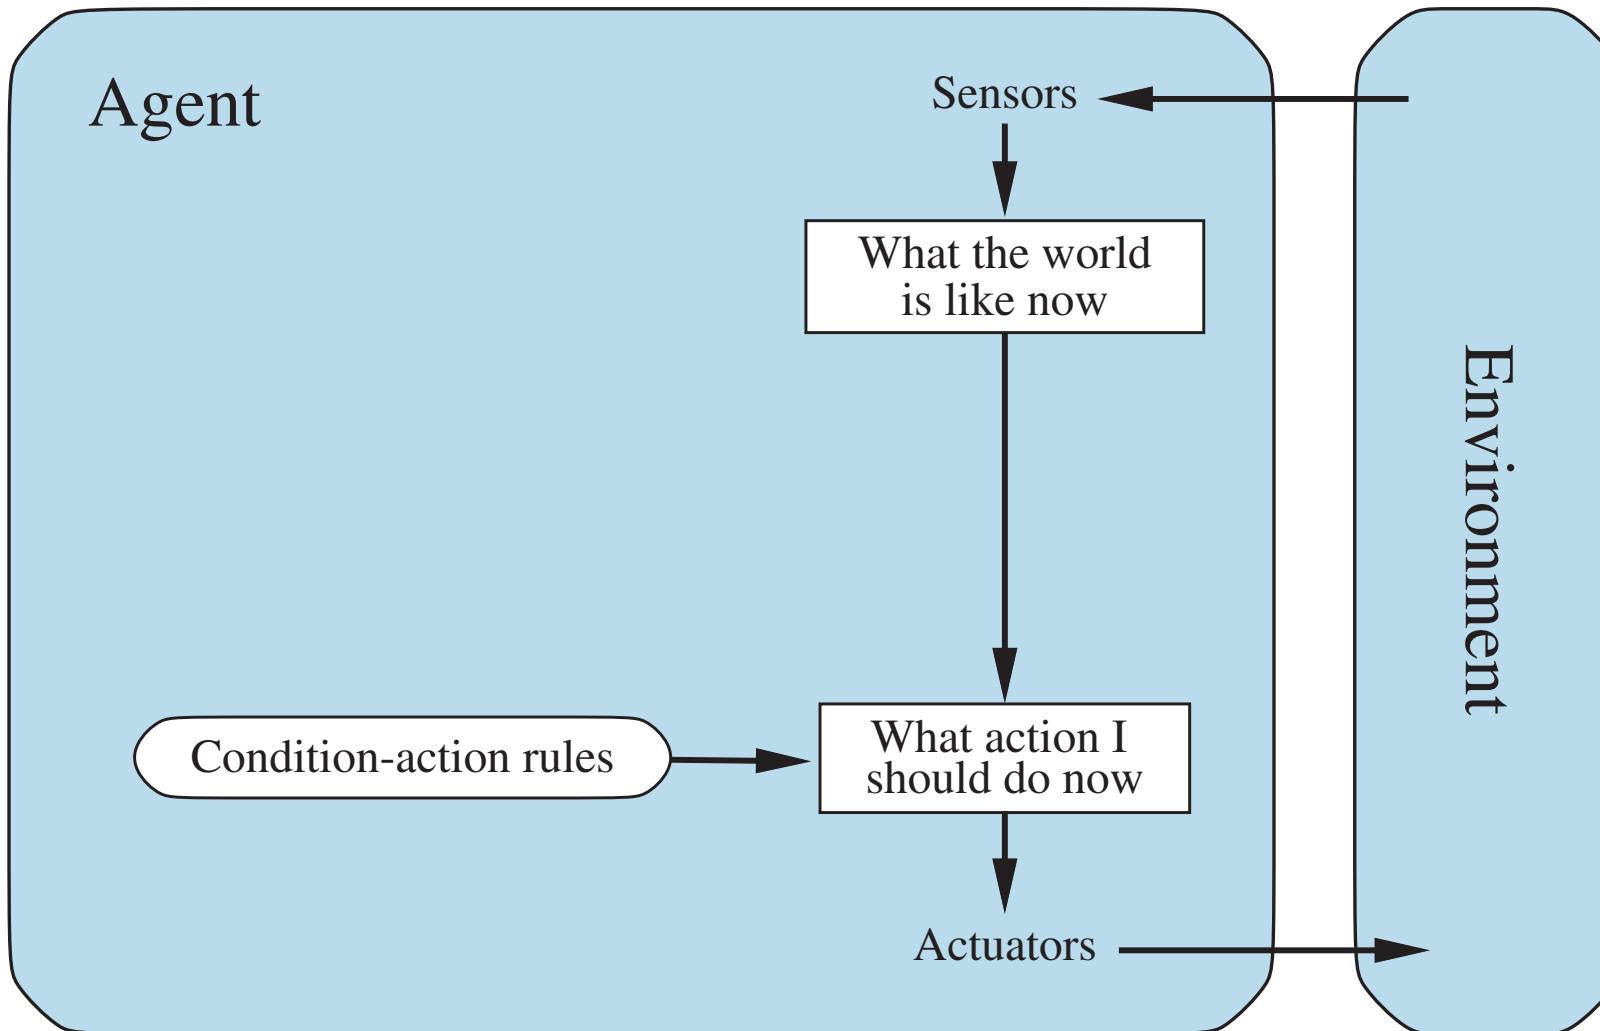
The agent program for a simple reflex agent in the two-location vacuum environment.

function REFLEX-VACUUM-AGENT(*[location, status]*) **returns** an action

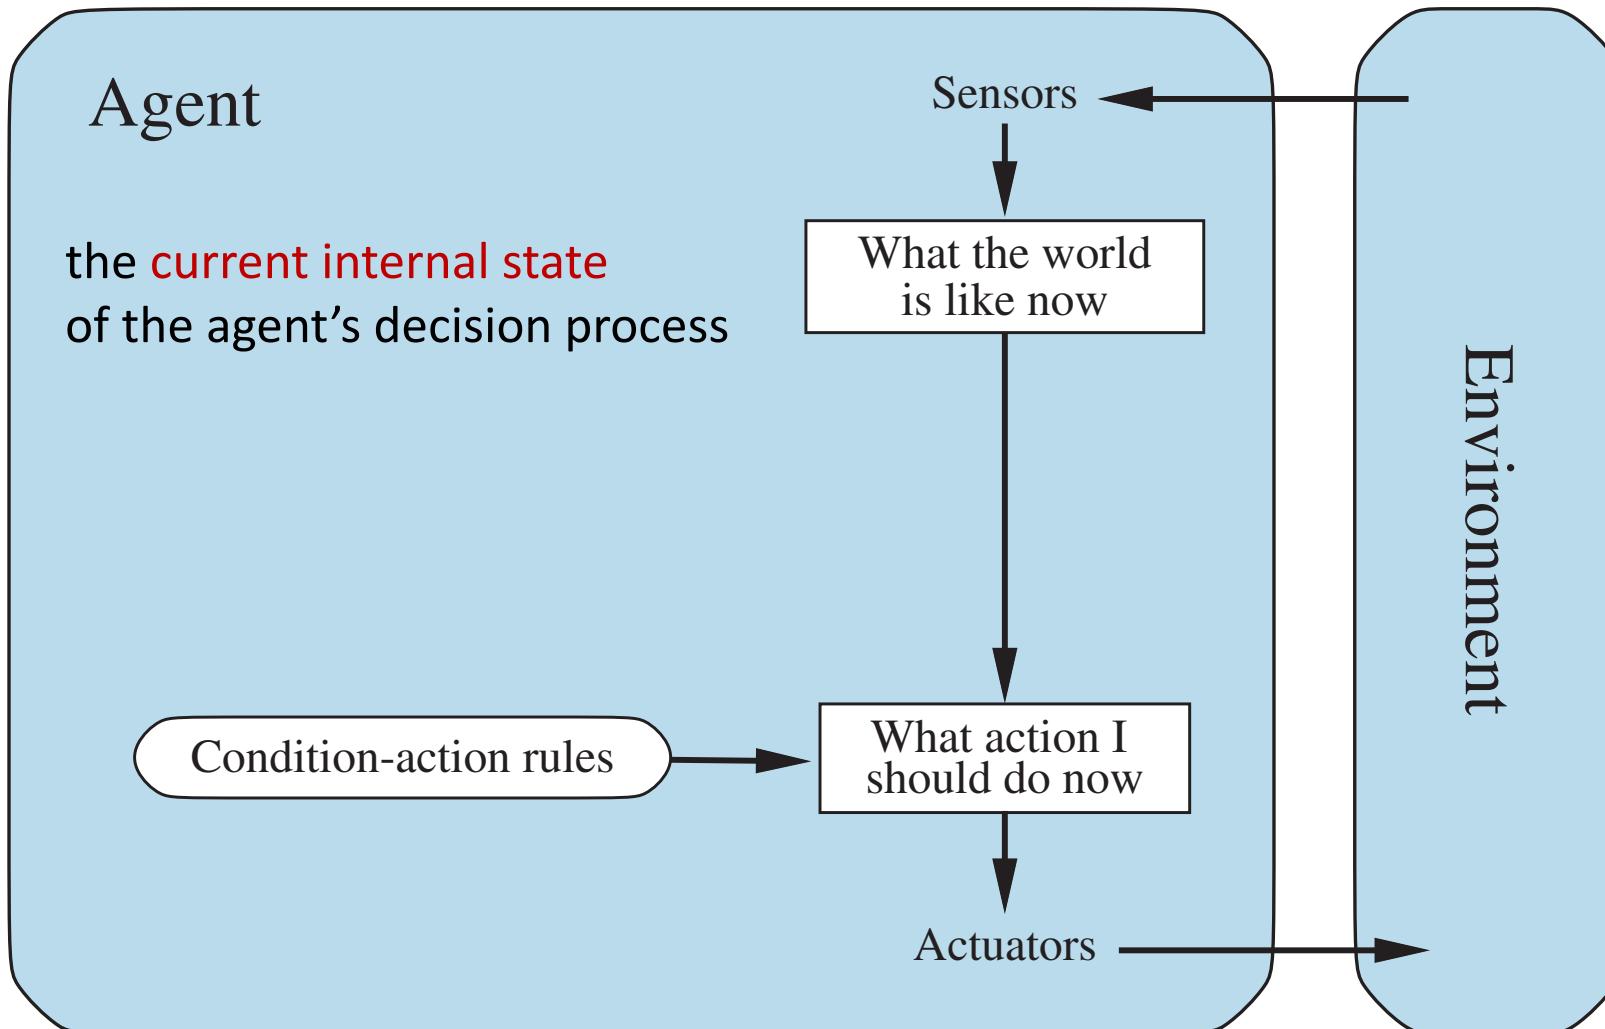
if *status* = *Dirty* **then return** *Suck*
else if *location* = *A* **then return** *Right*
else if *location* = *B* **then return** *Left*



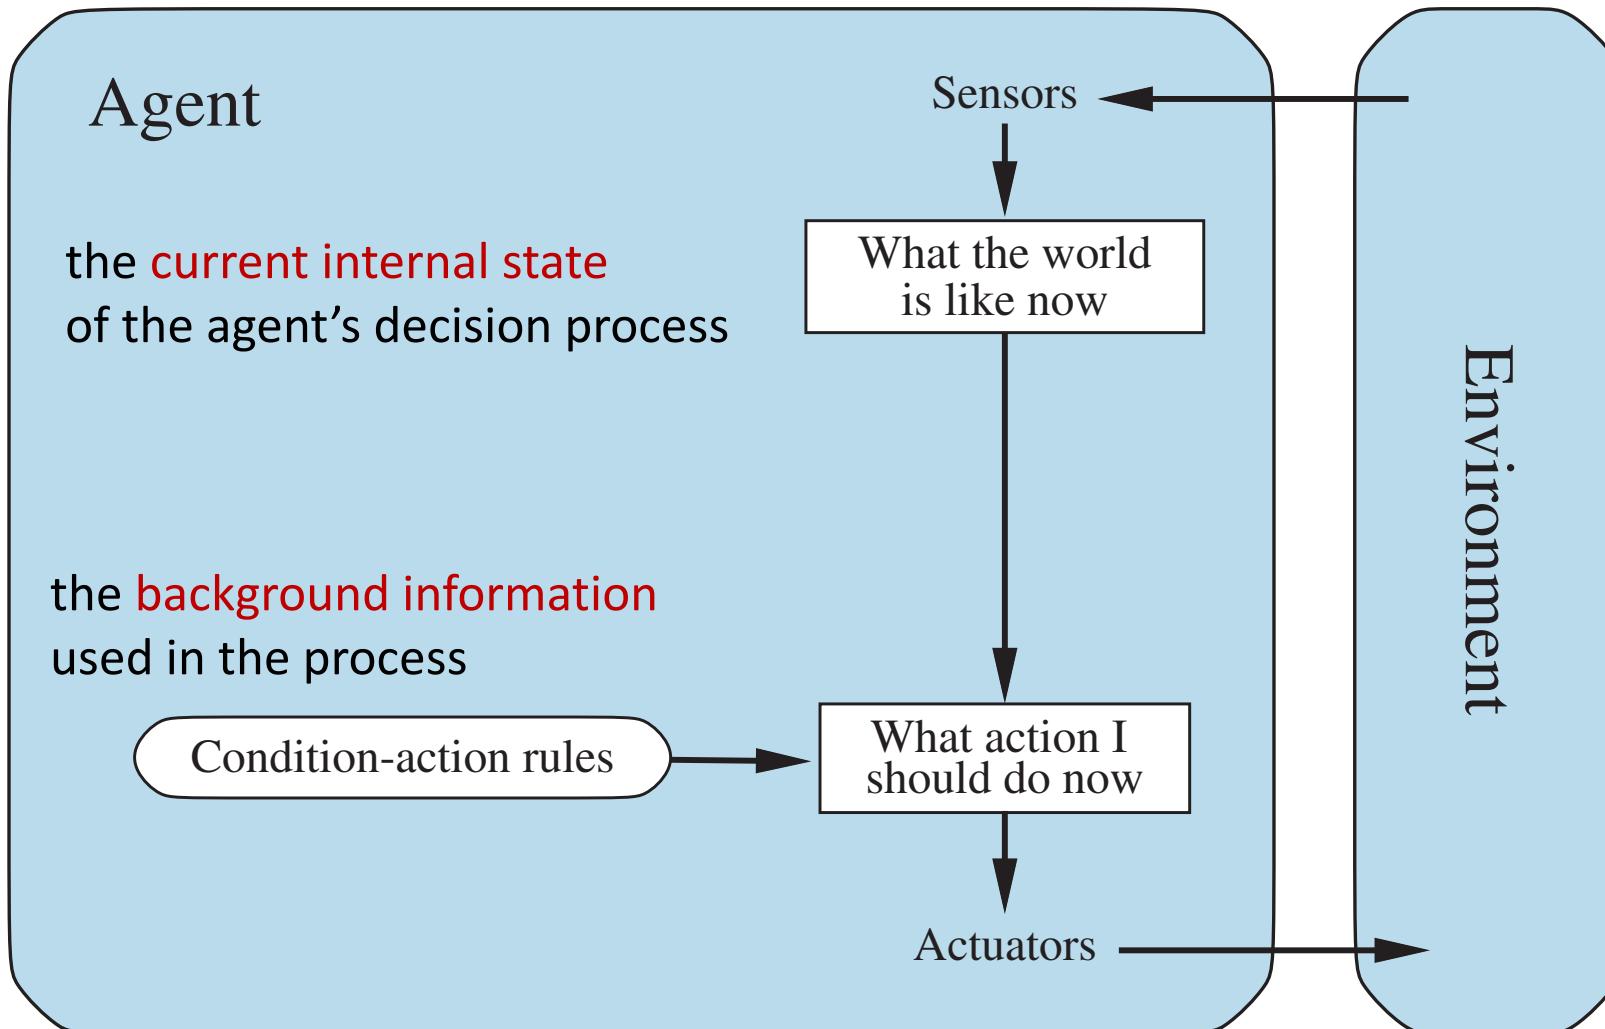
Schematic Diagram of a Simple Reflex Agent



Schematic diagram of a simple reflex agent



Schematic diagram of a simple reflex agent

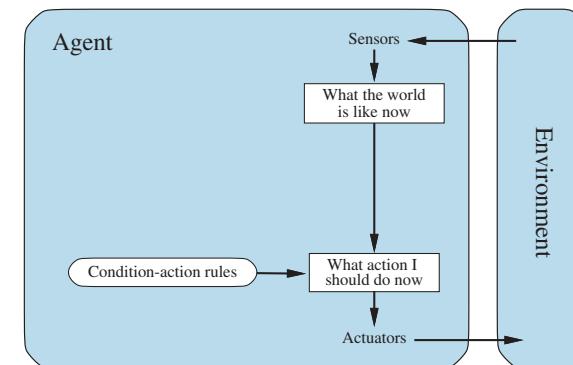


A Simple Reflex Agent

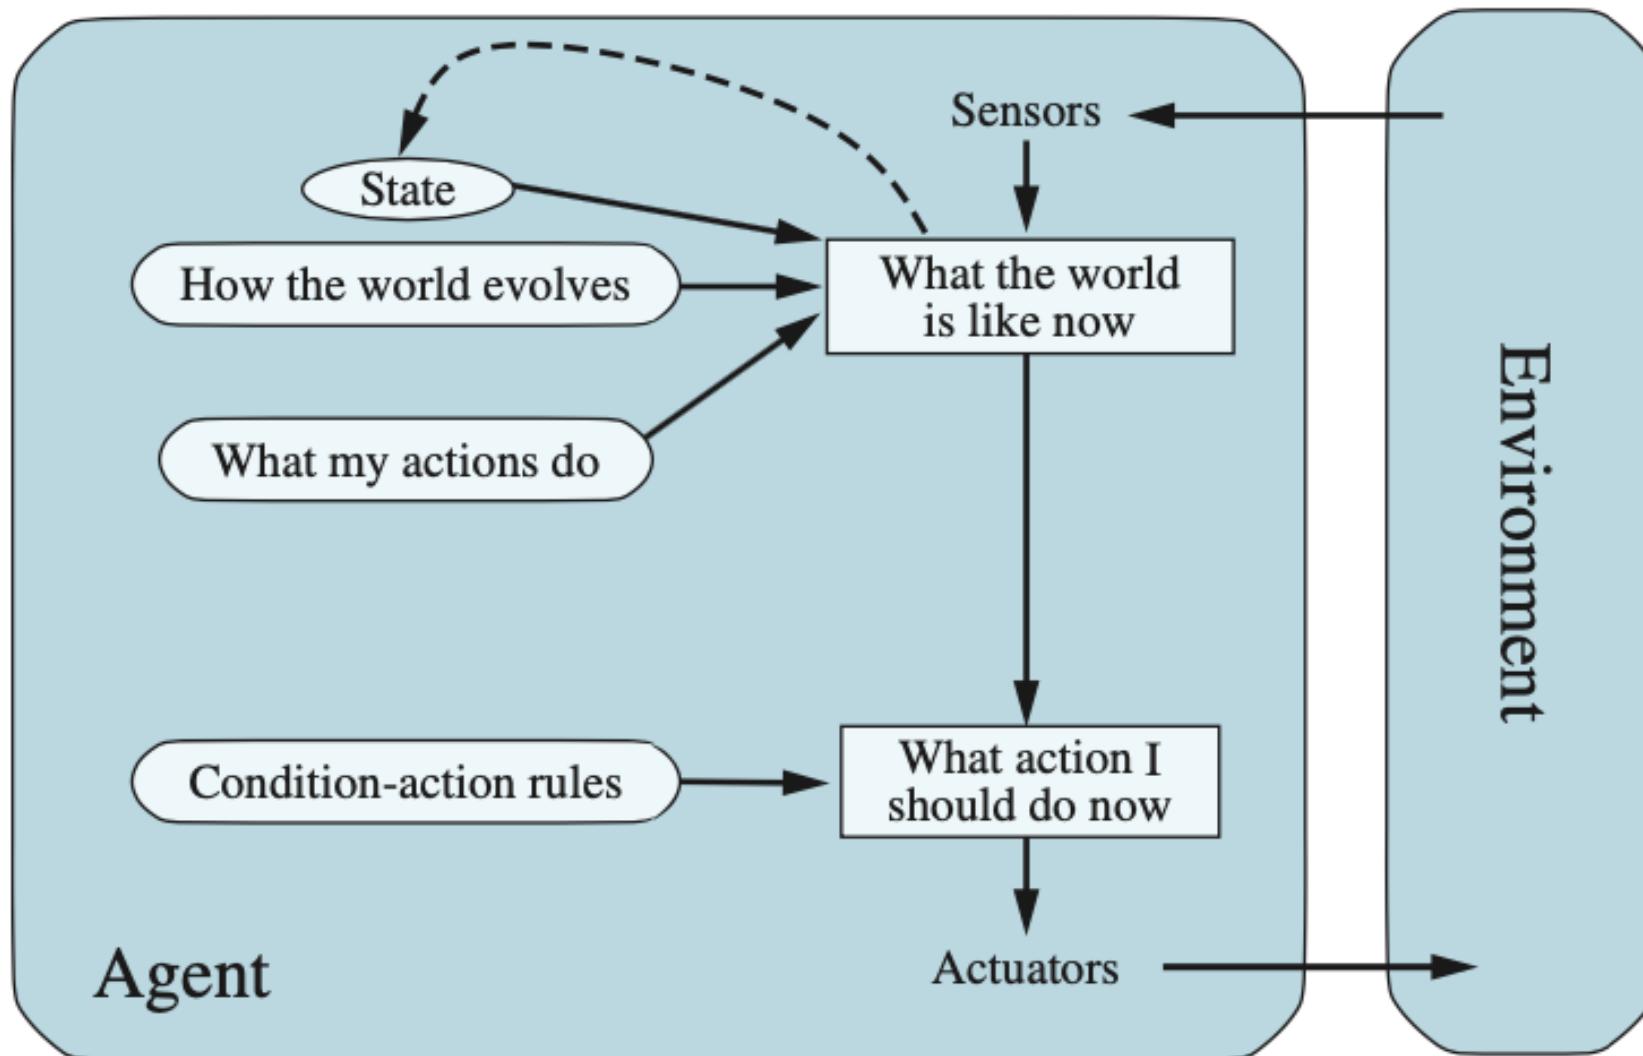
It acts according to a rule whose condition matches the current state, as defined by the percept.

function SIMPLE-REFLEX-AGENT(*percept*) **returns** an action
persistent: *rules*, a set of condition–action rules

```
state  $\leftarrow$  INTERPRET-INPUT(percept)
rule  $\leftarrow$  RULE-MATCH(state, rules)
action  $\leftarrow$  rule.ACTION
return action
```



A Model-based Reflex Agent



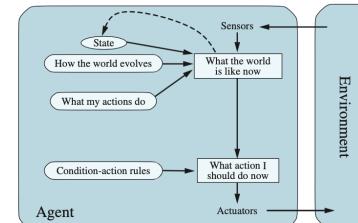
A model-based reflex agent

It keeps track of the current state of the world,
using an internal model.

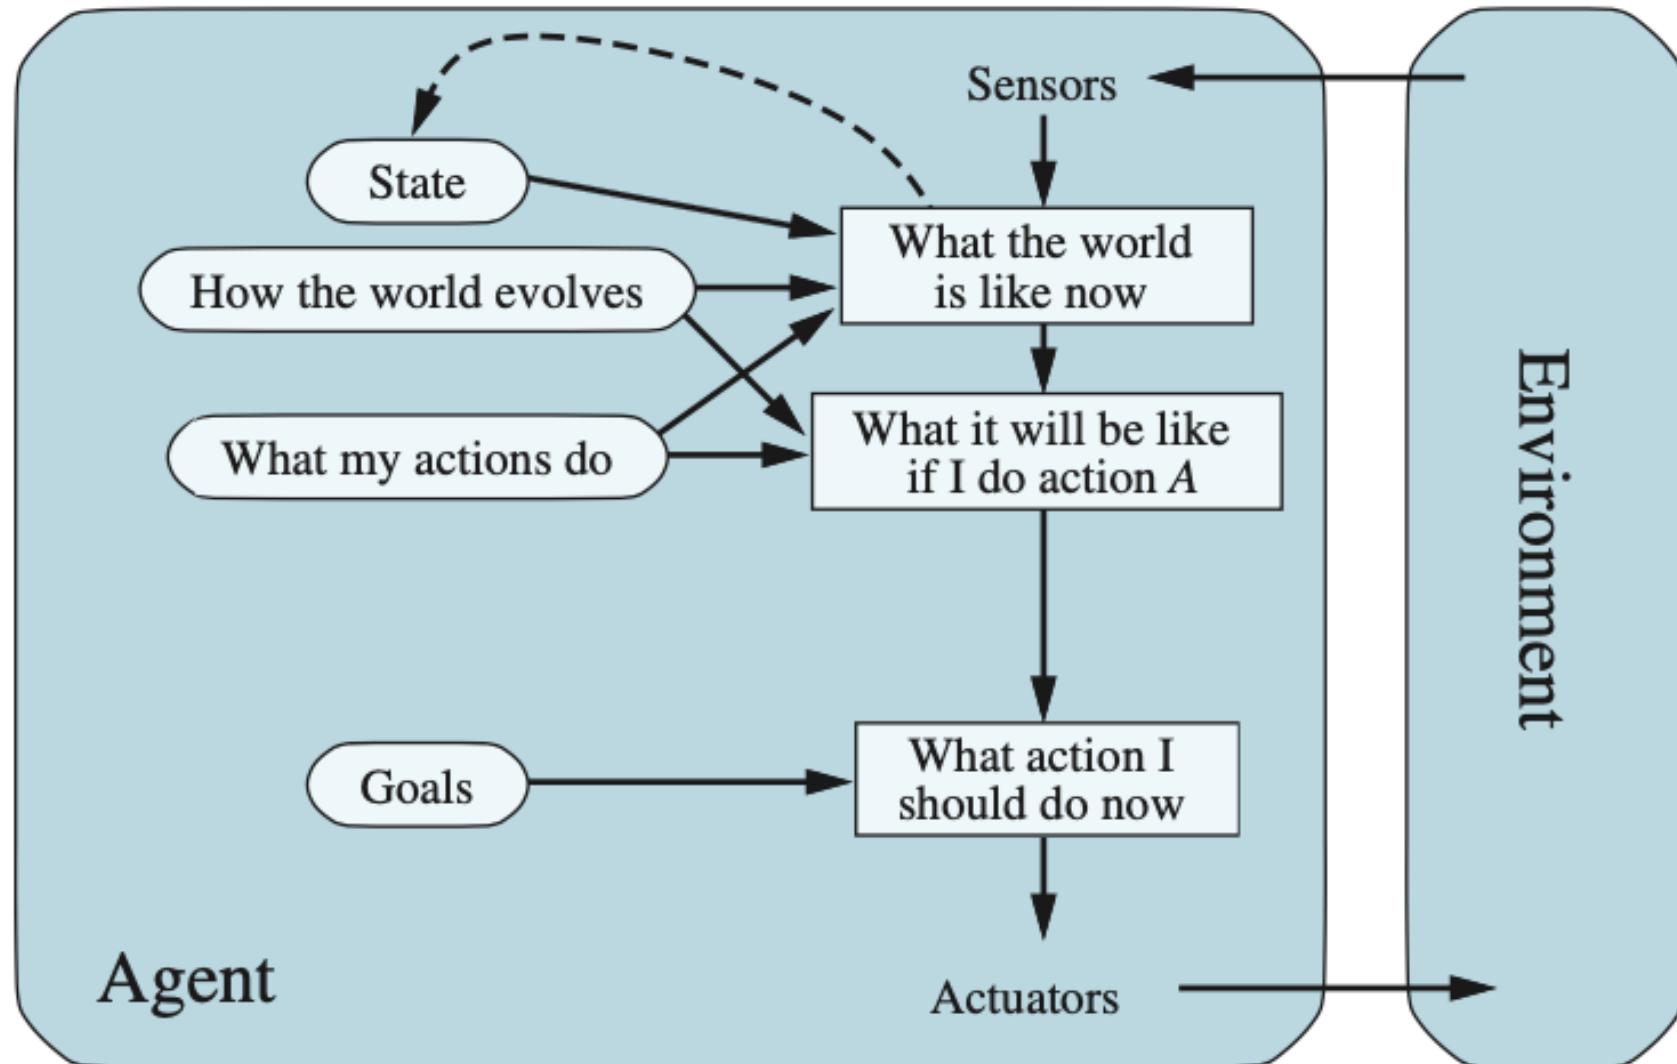
It then chooses an action in the same way as the reflex agent.

```
function MODEL-BASED-REFLEX-AGENT(percept) returns an action
    persistent: state, the agent's current conception of the world state
    transition-model, a description of how the next state depends on
        the current state and action
    sensor-model, a description of how the current world state is reflected
        in the agent's percepts
    rules, a set of condition-action rules
    action, the most recent action, initially none

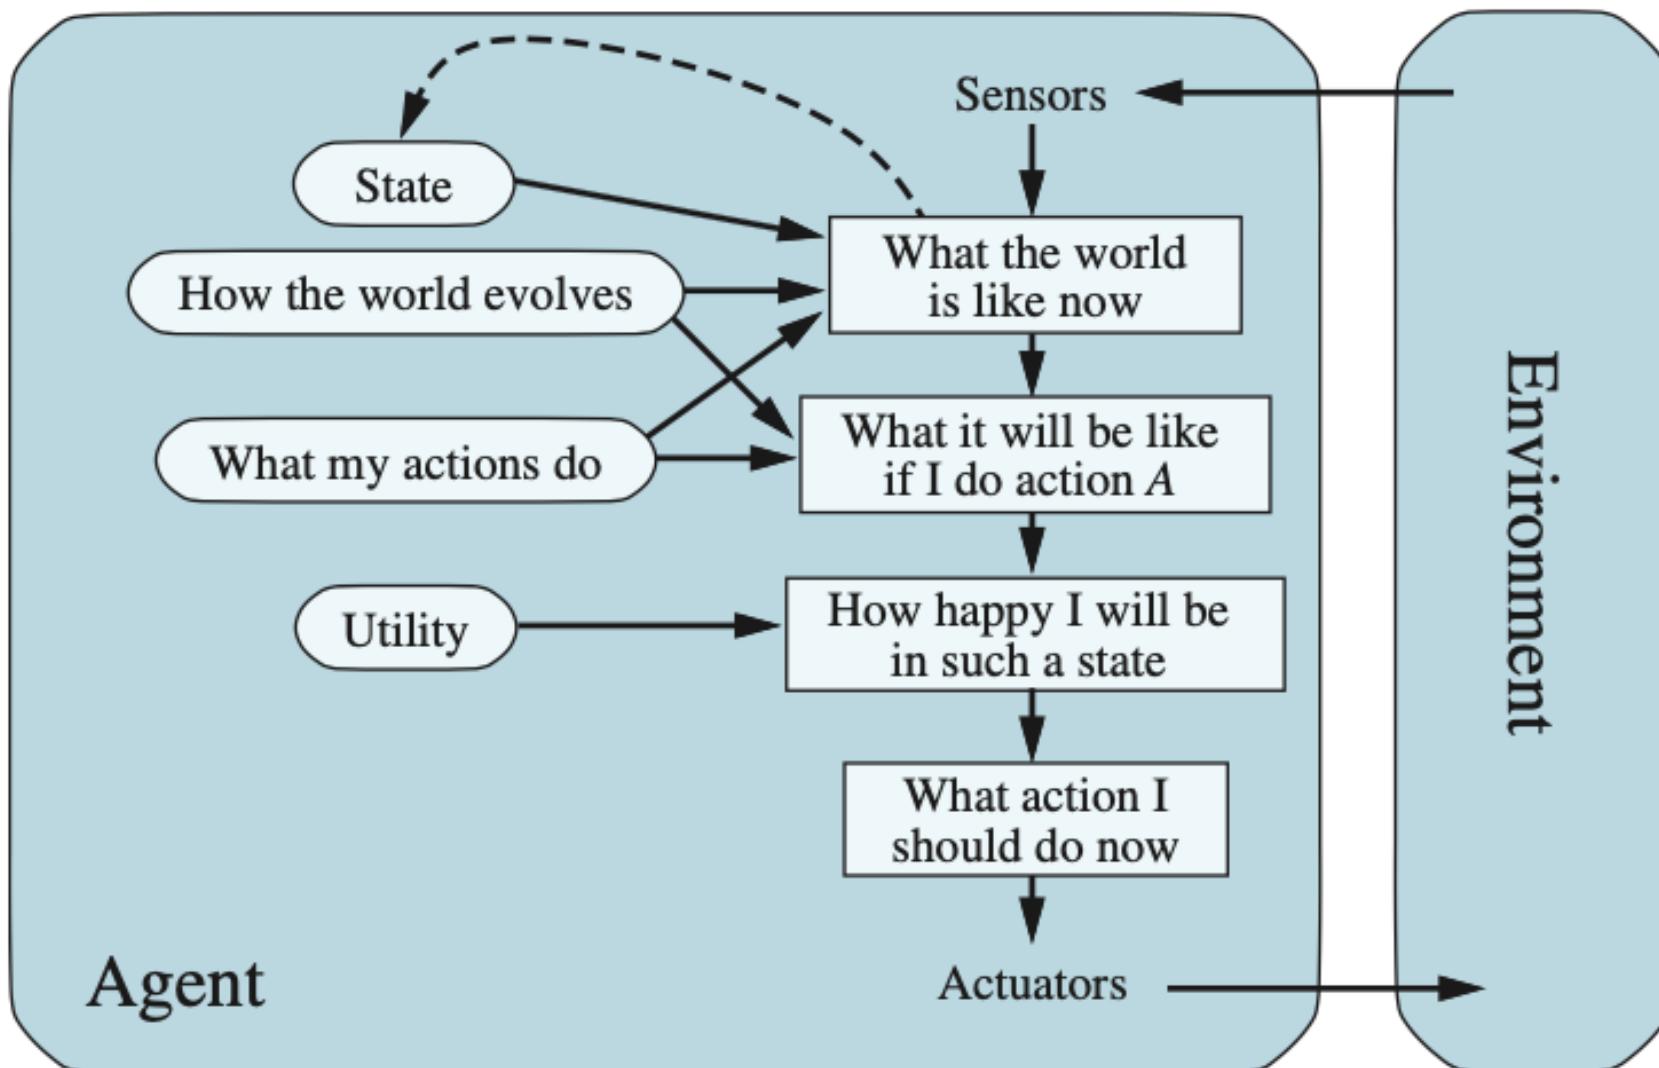
    state  $\leftarrow$  UPDATE-STATE(state, action, percept, transition-model, sensor-model)
    rule  $\leftarrow$  RULE-MATCH(state, rules)
    action  $\leftarrow$  rule.ACTION
    return action
```



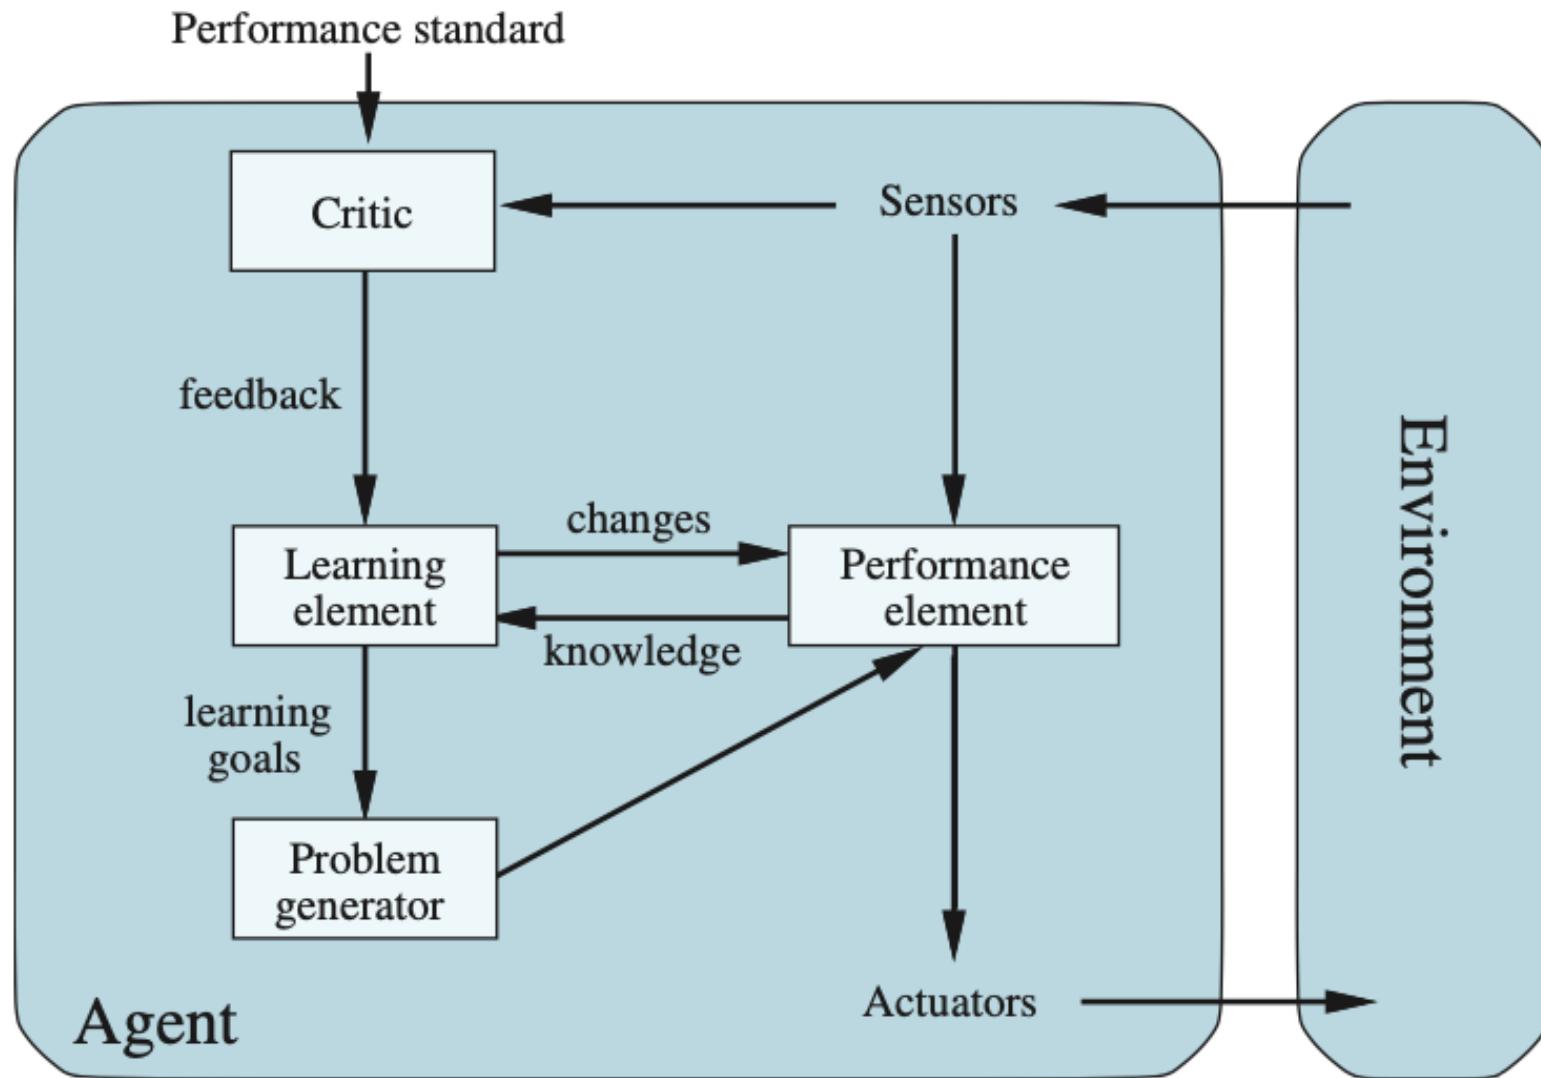
A model-based, goal-based agent



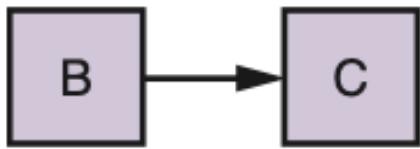
A model-based, utility-based agent



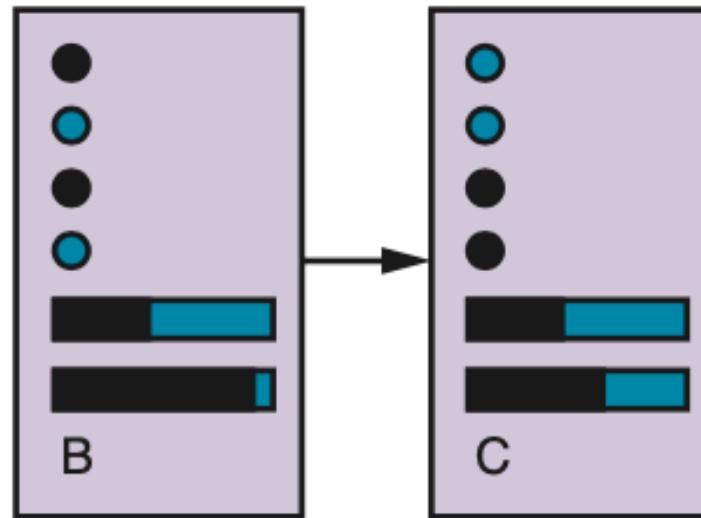
A general learning agent



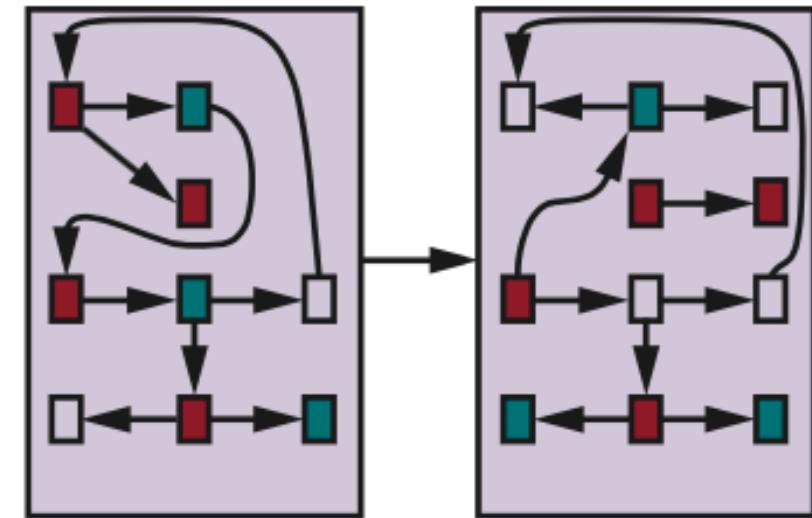
Three ways to represent states and the transitions between them



(a) Atomic

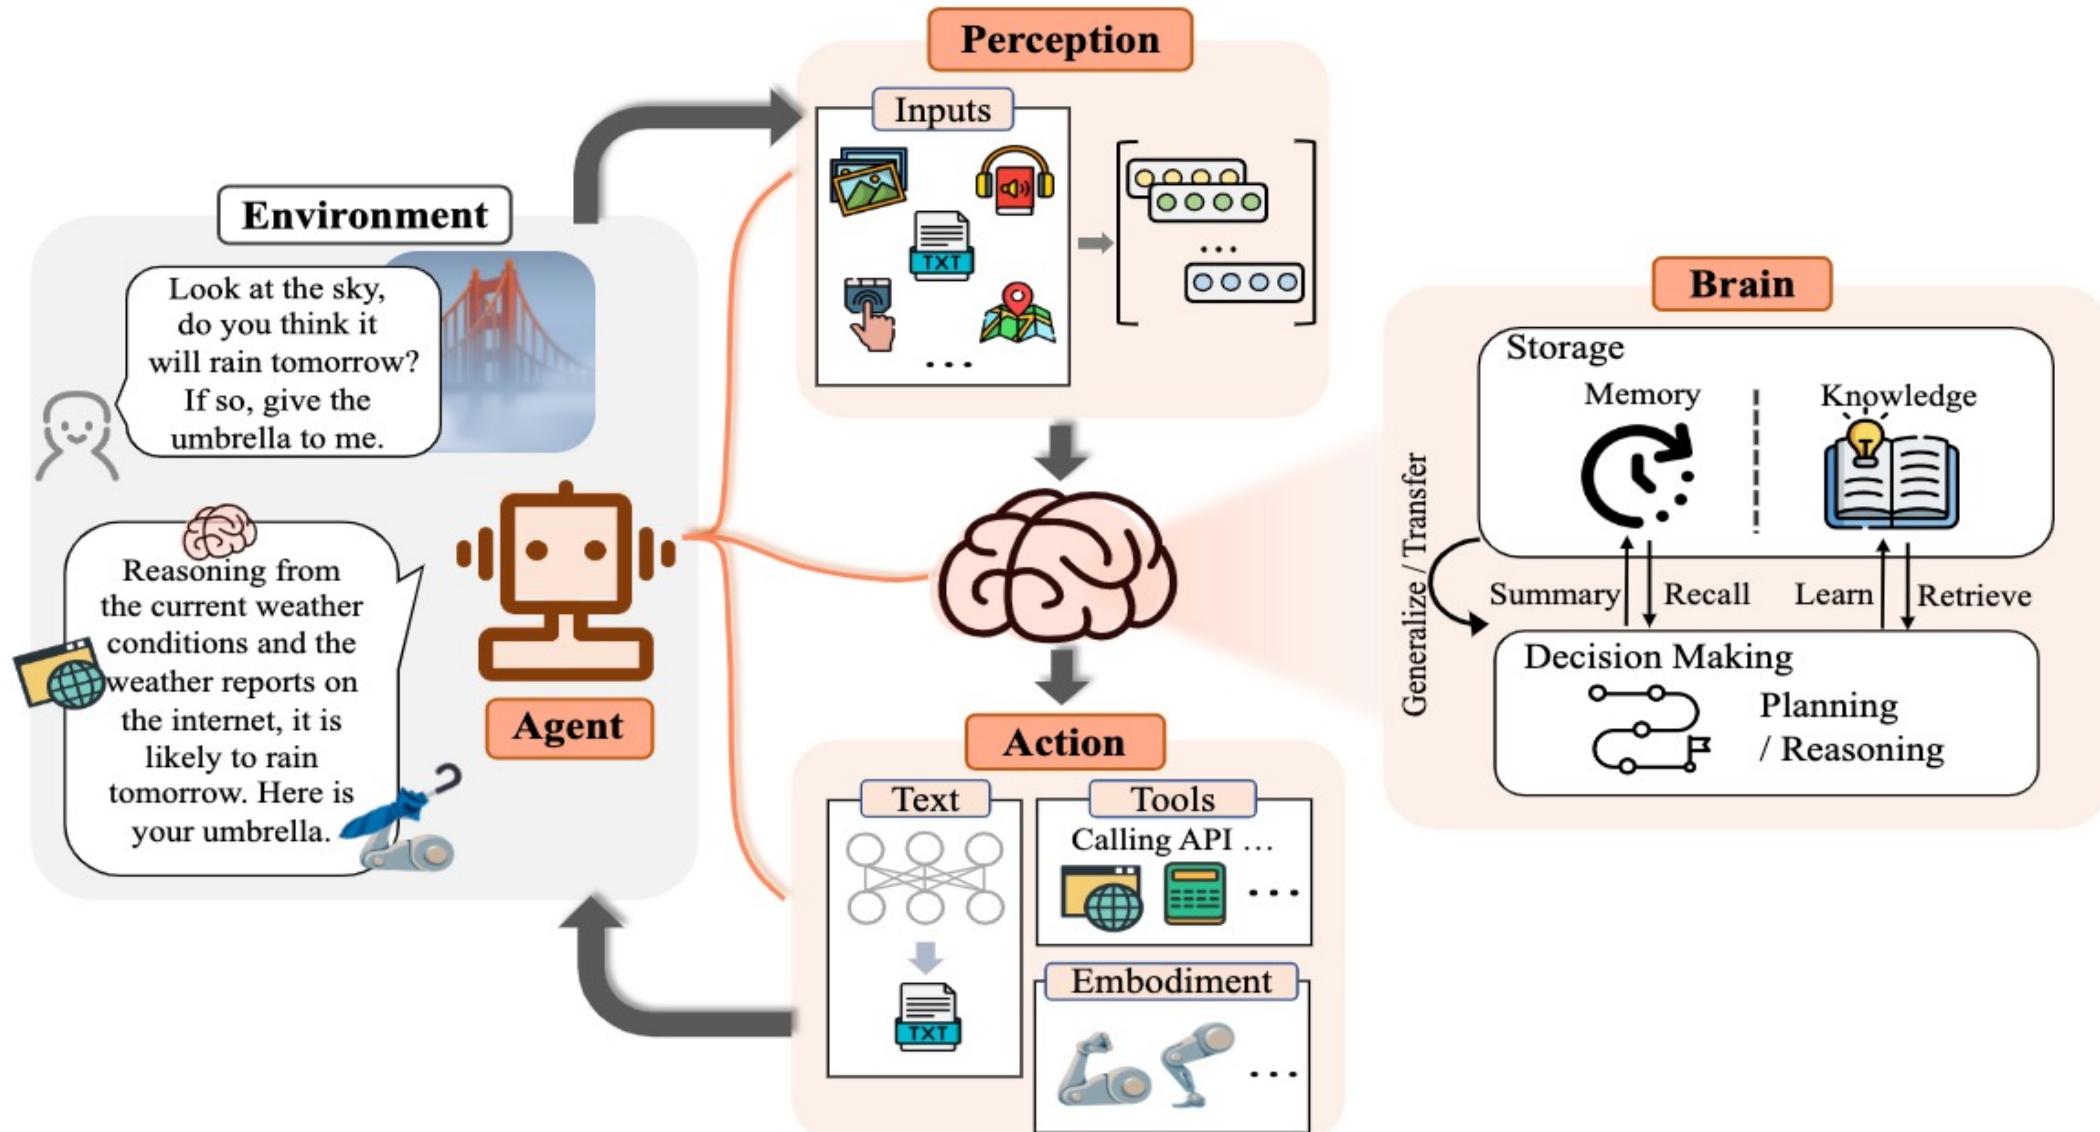


(b) Factored



(c) Structured

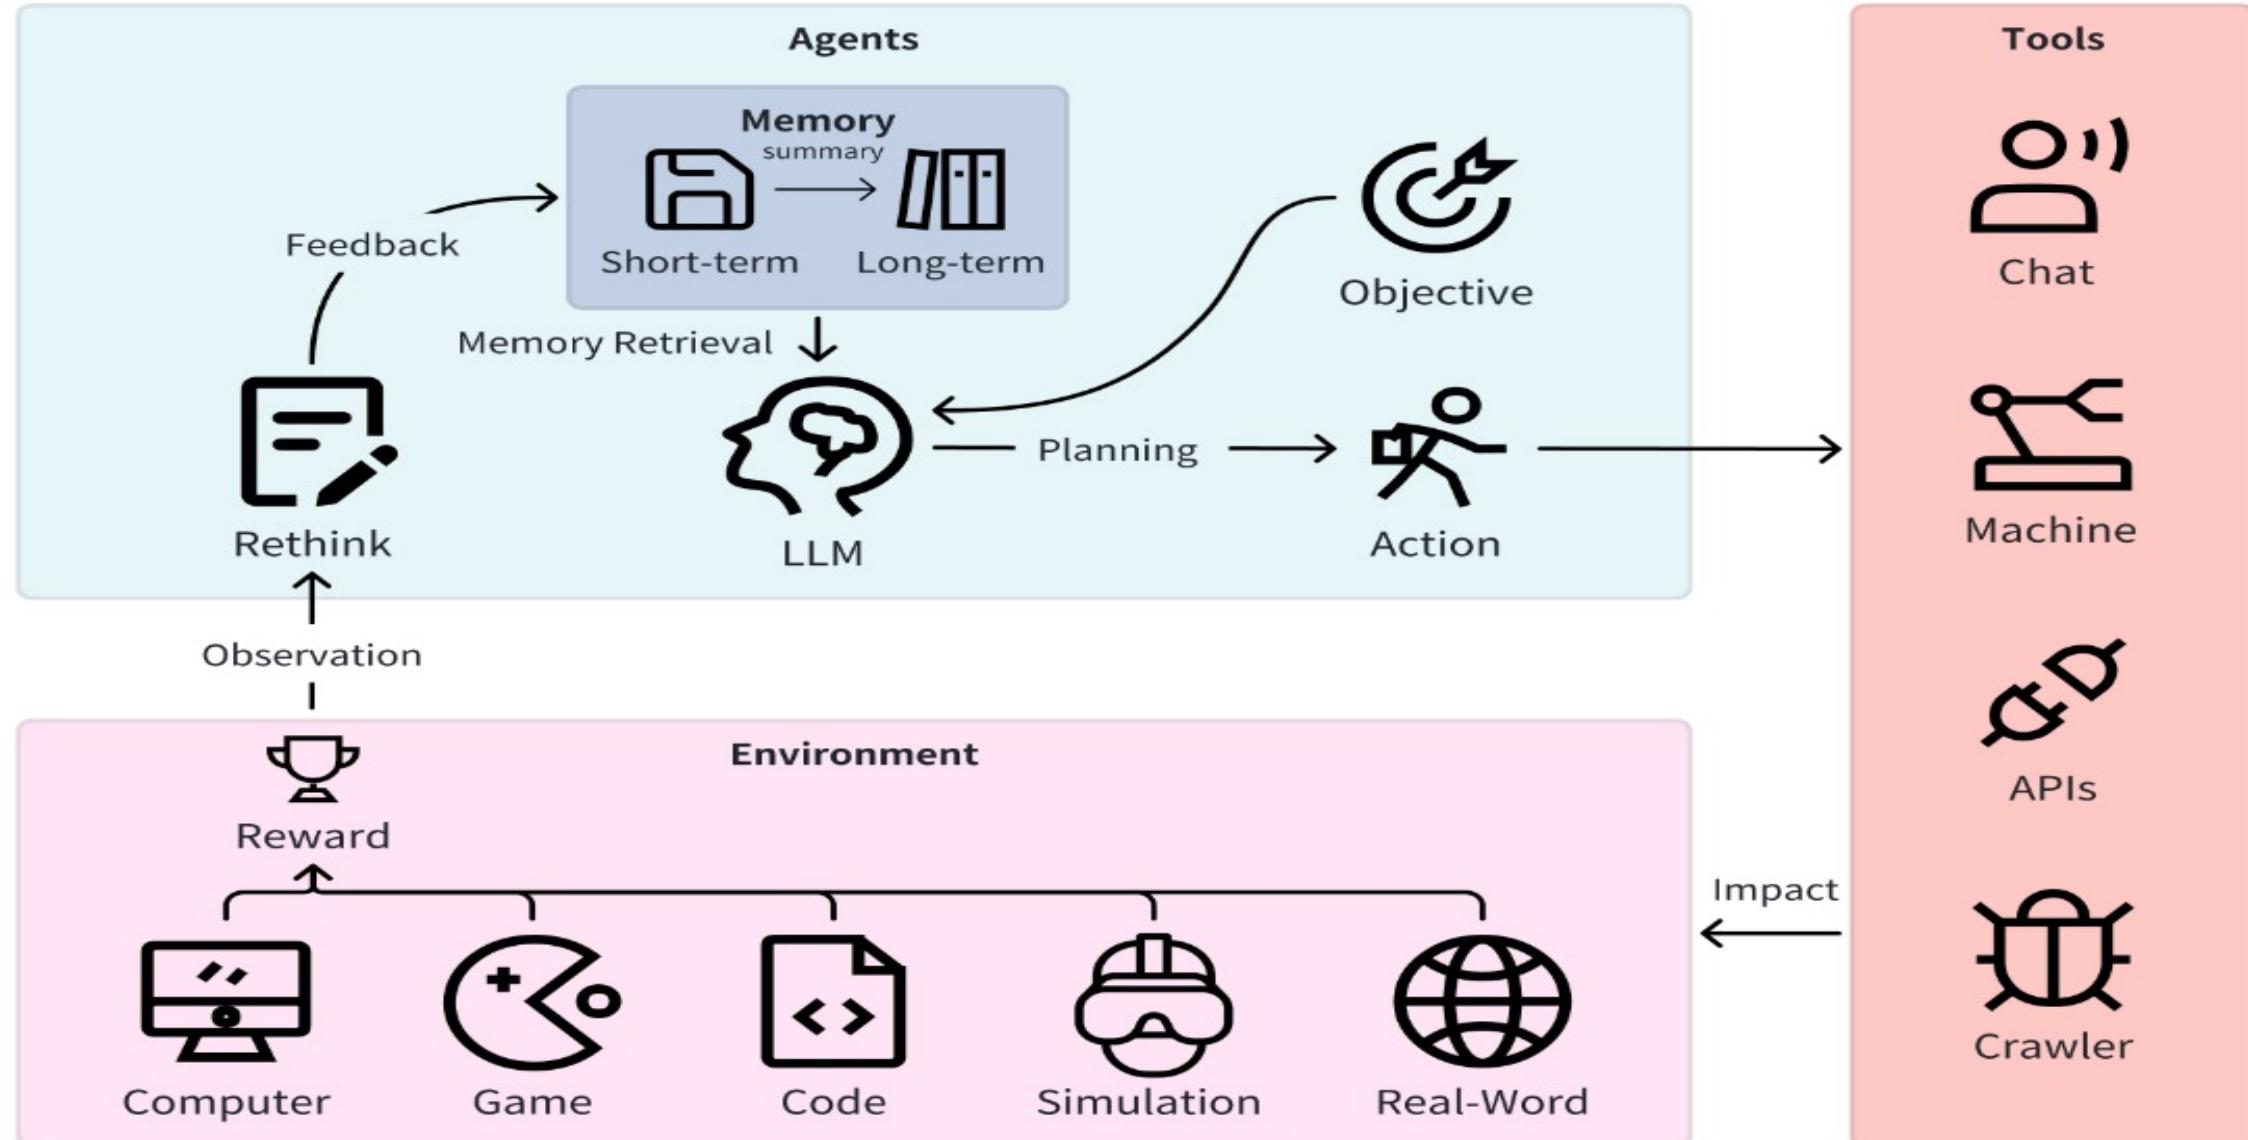
Large Language Model (LLM) based Agents



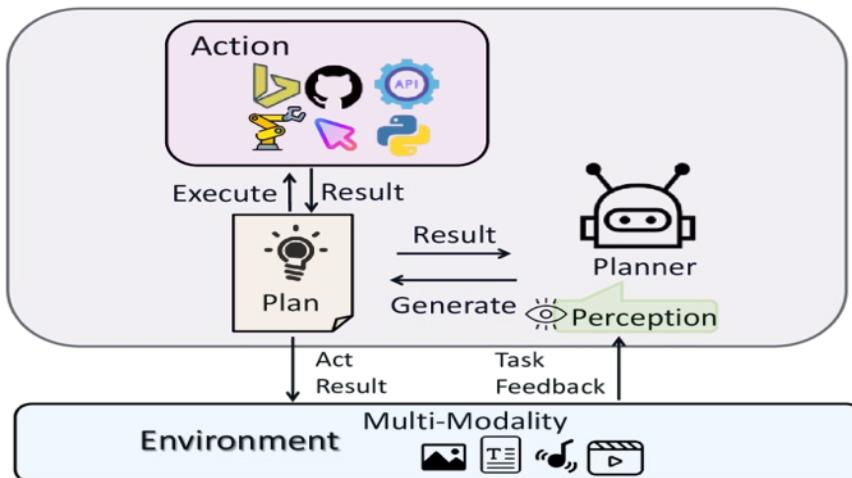
LLM-based Agents

- **Definition:** AI agents that use Large Language Models as their core decision-making mechanism
- Key Features:
 - Natural language interface
 - Vast knowledge base
 - Ability to understand context and nuance
 - Generalize to new tasks with minimal additional training

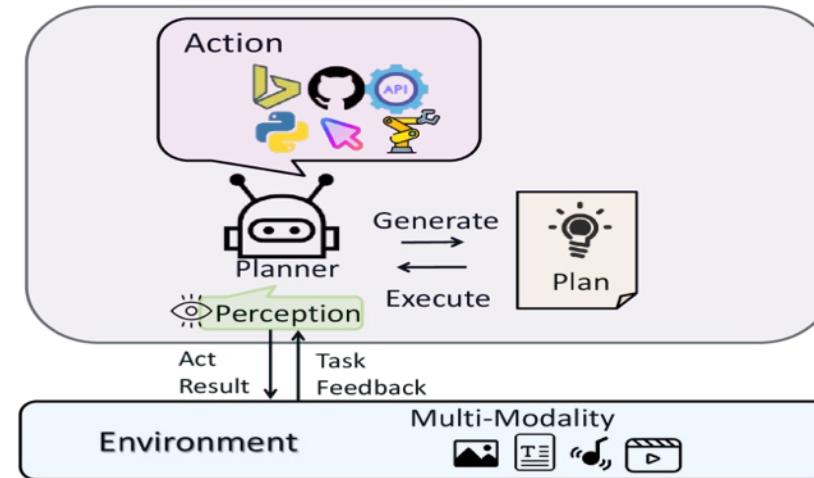
LLM-based Agents



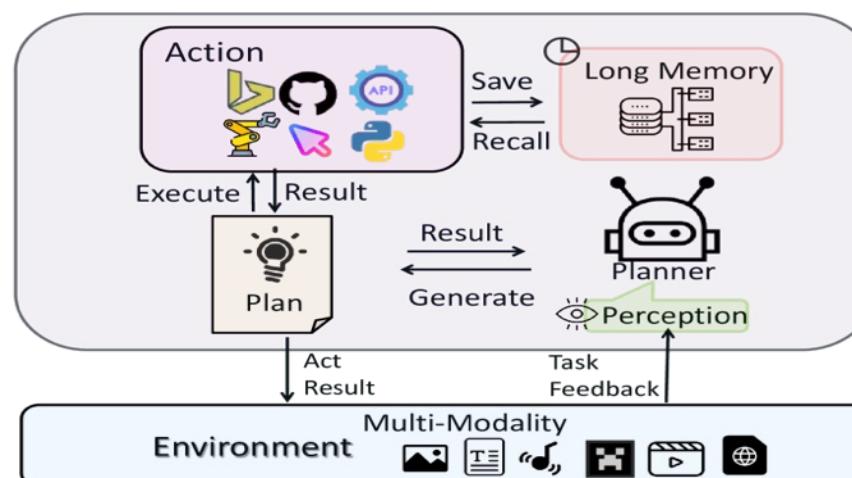
Large Multimodal Agents (LMA)



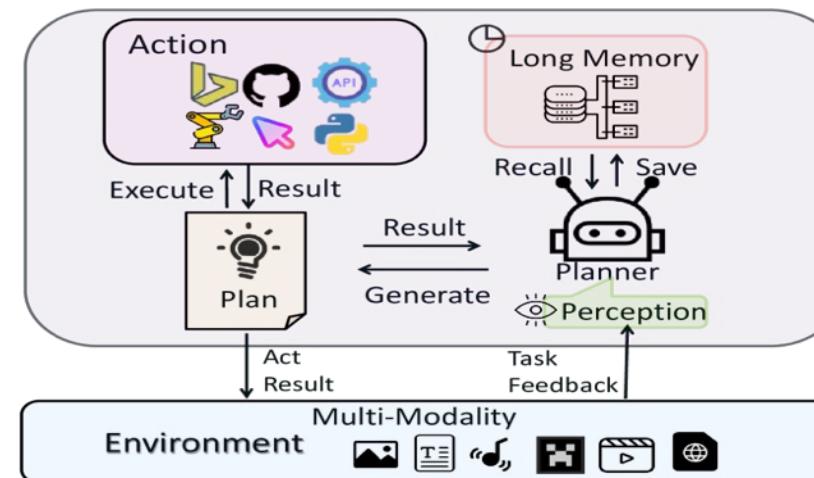
(a)



(b)

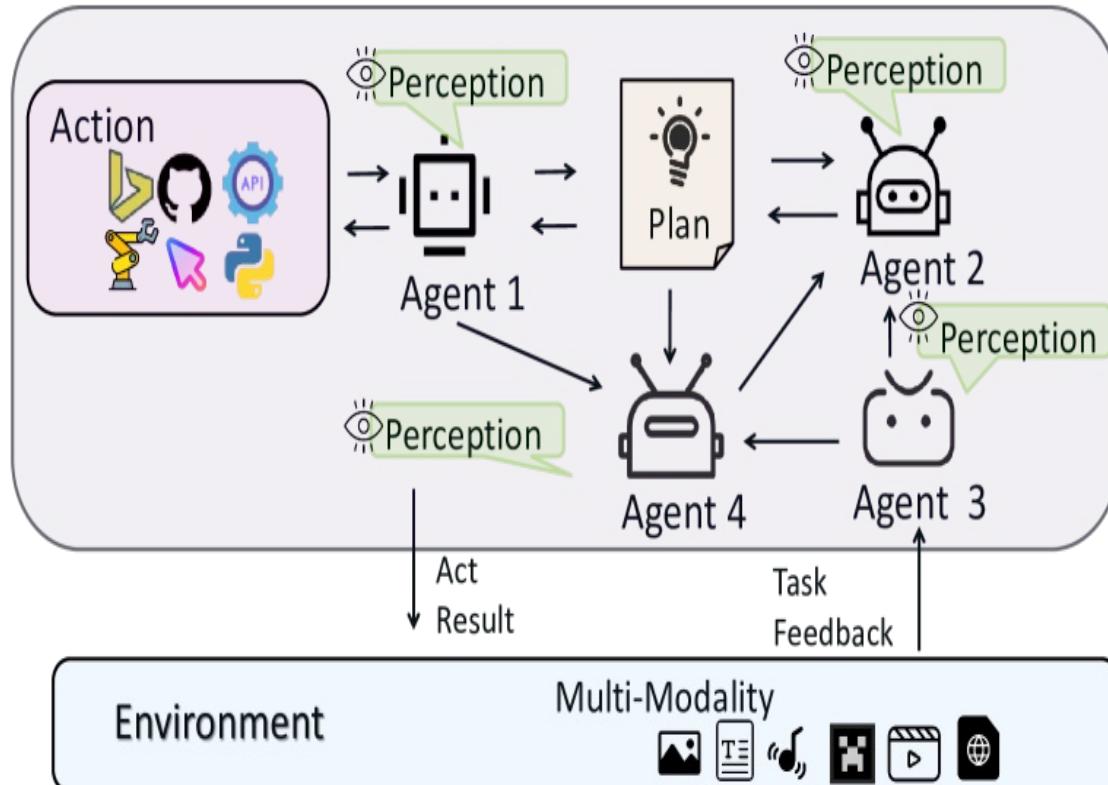


(c)

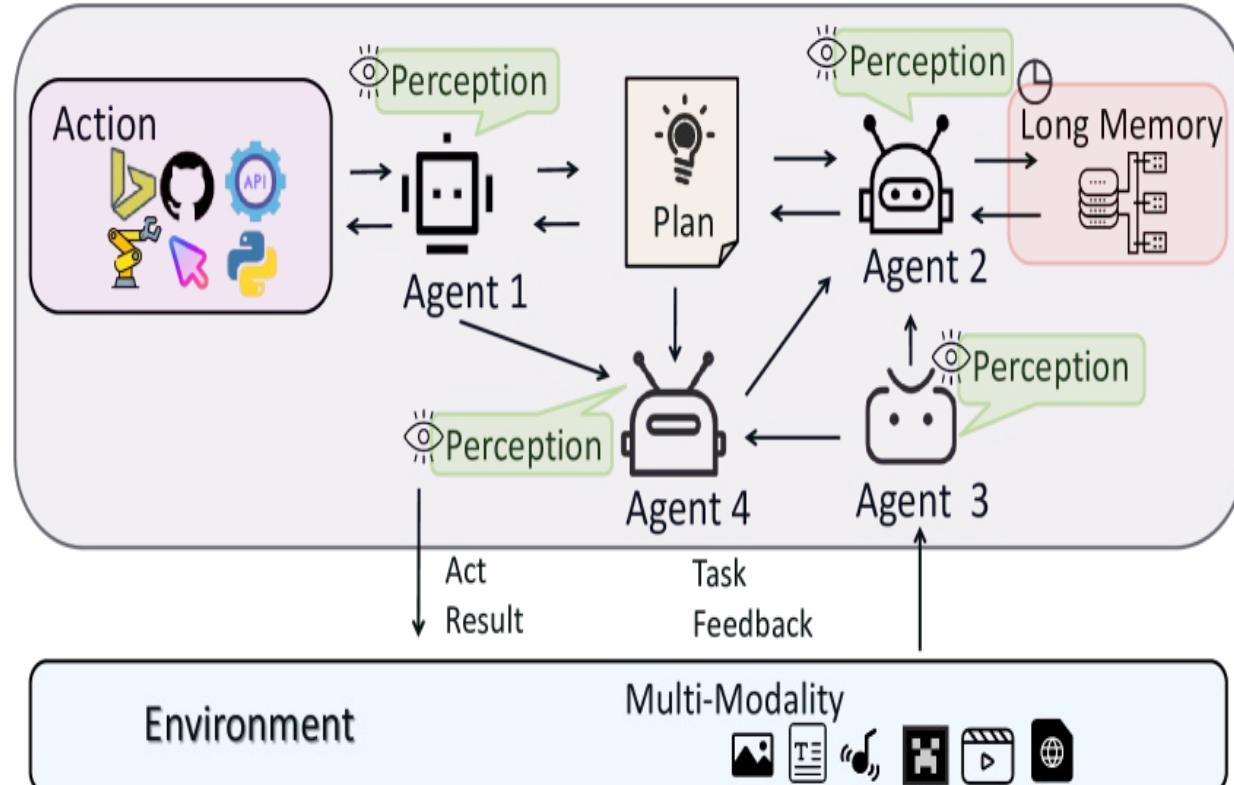


(d)

Large Multimodal Agents (LMA)



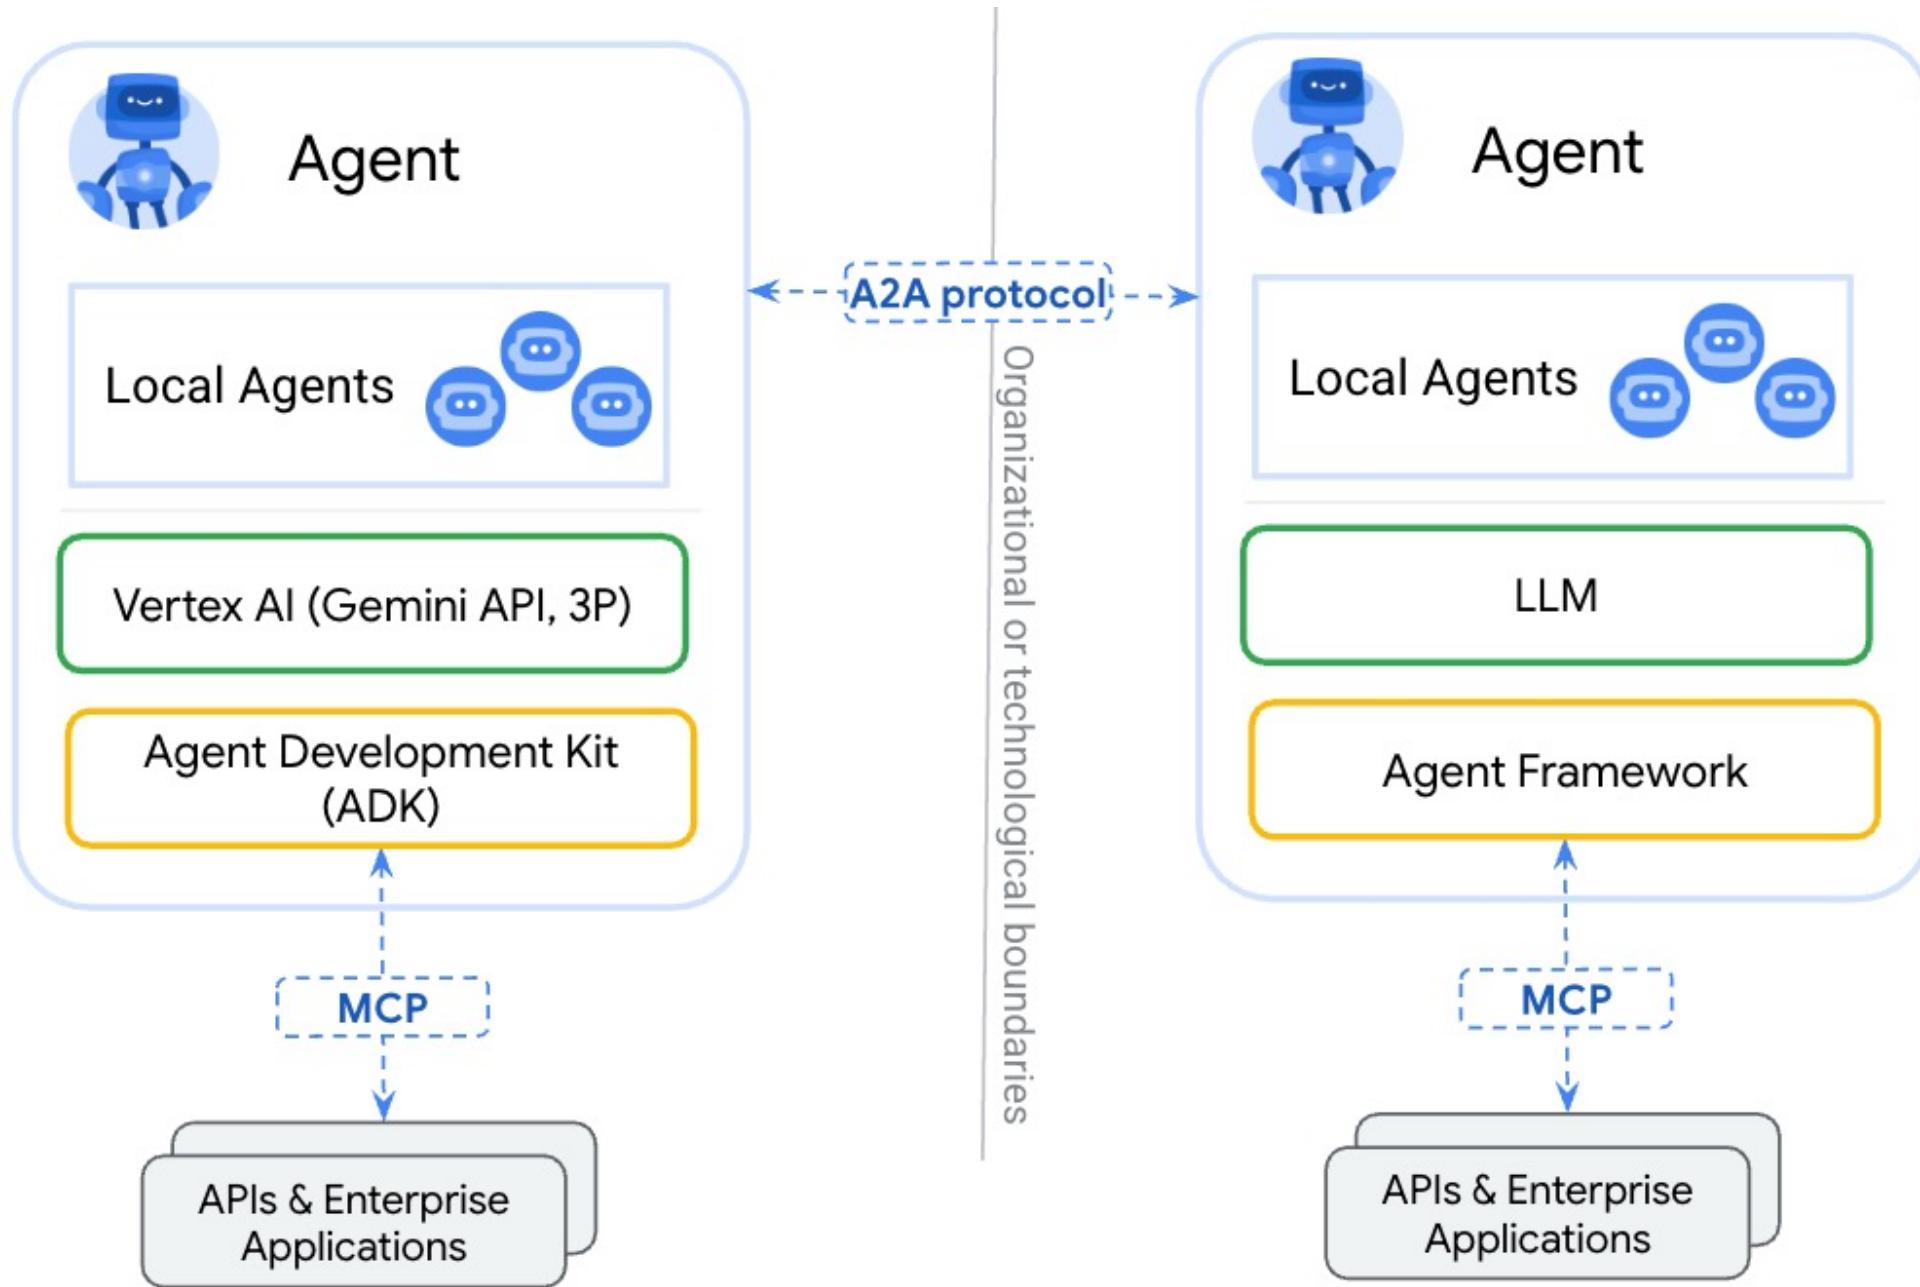
(a)



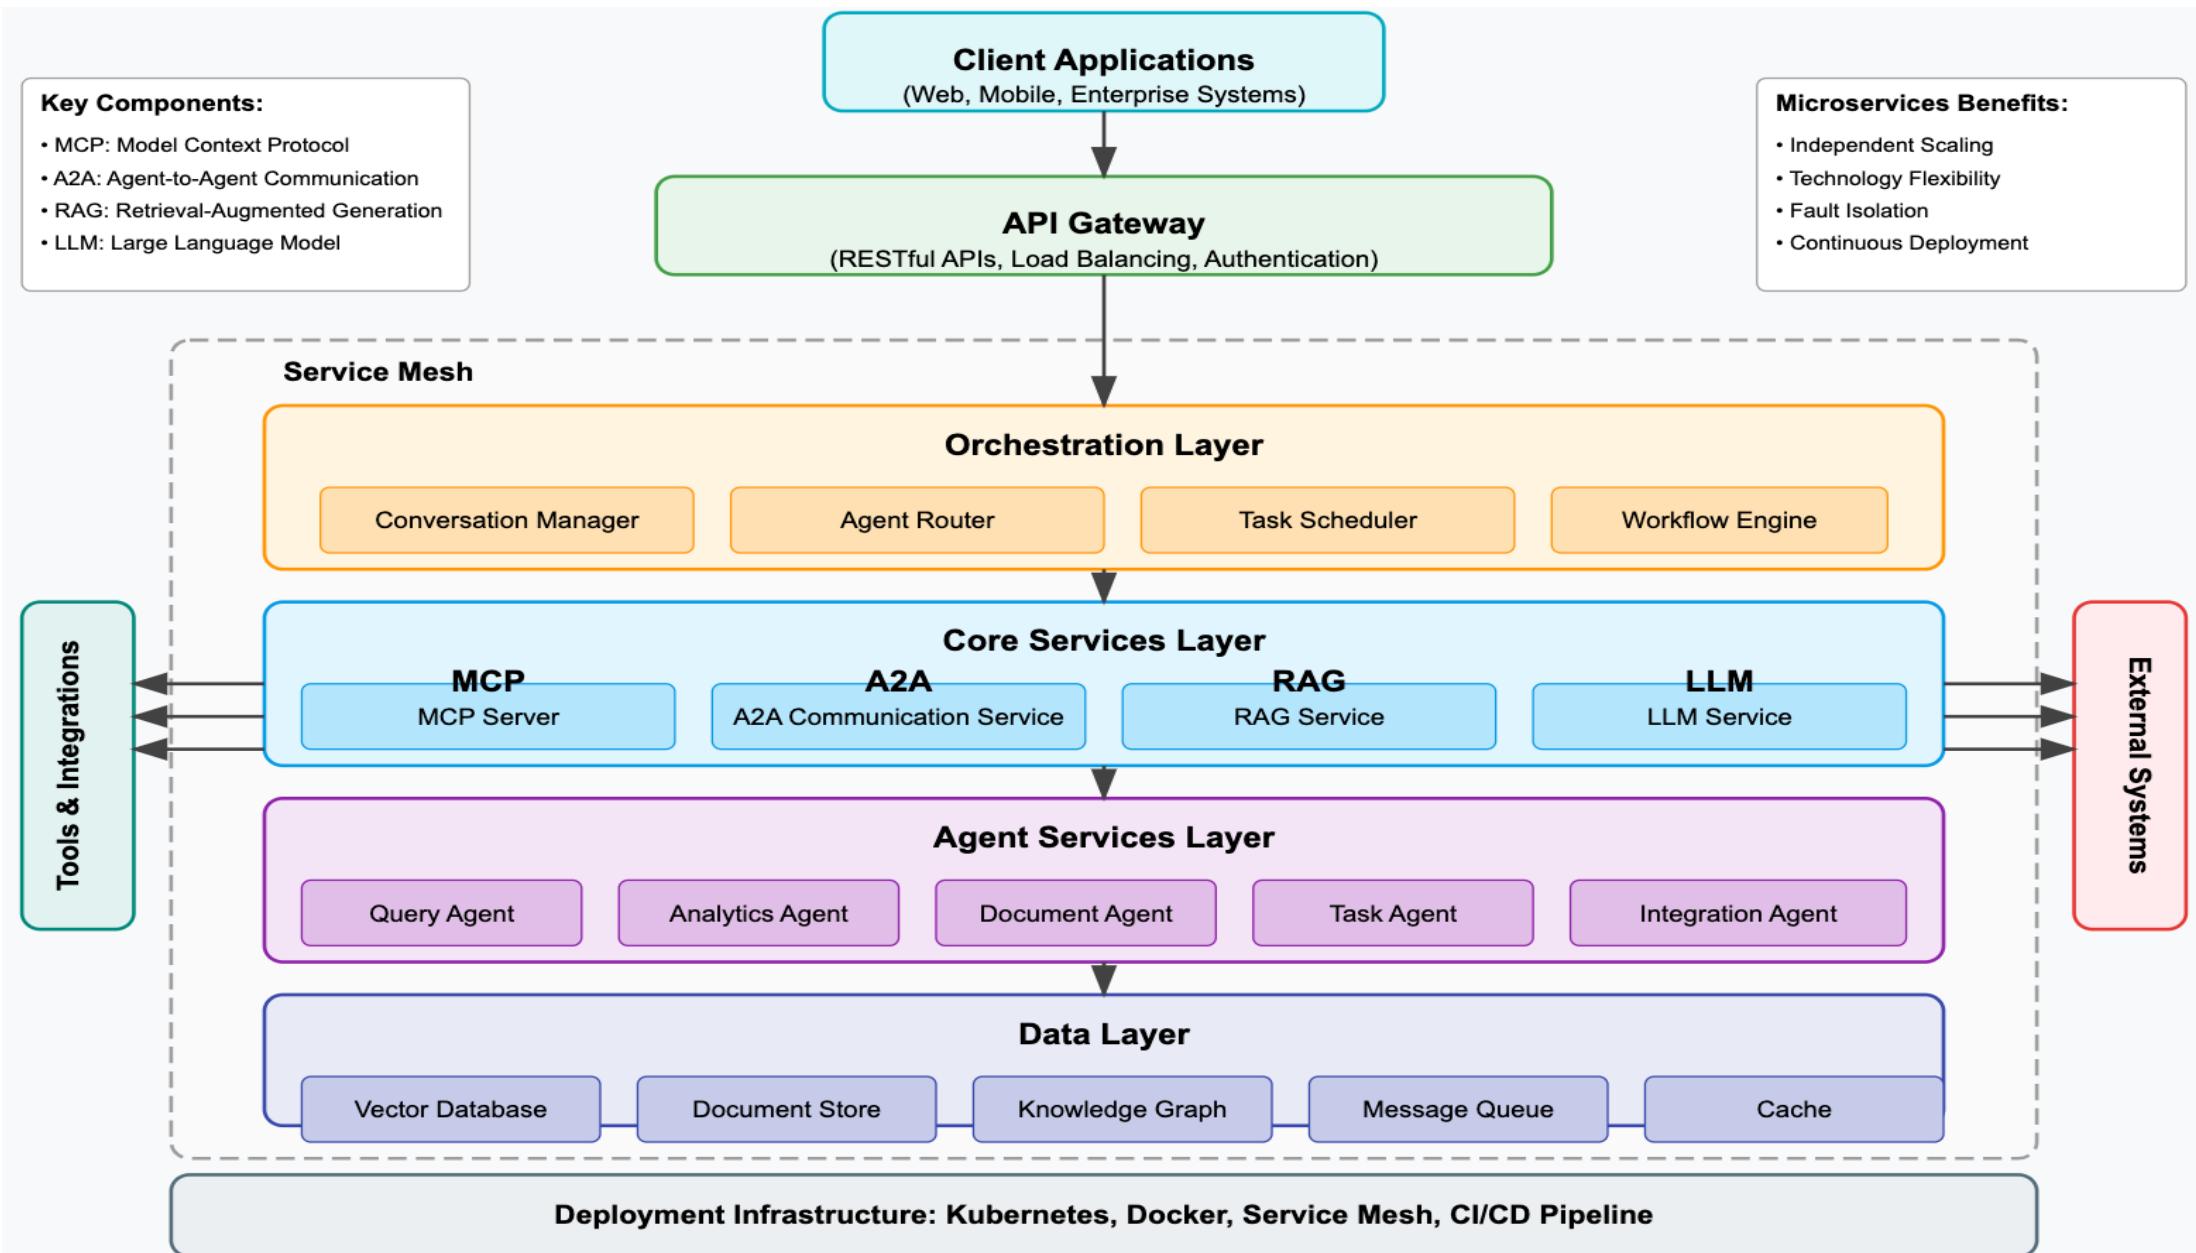
(b)

A2A (Agent2Agent Protocol) for agent-agent collaboration

MCP (Model Context Protocol) for tools and resources

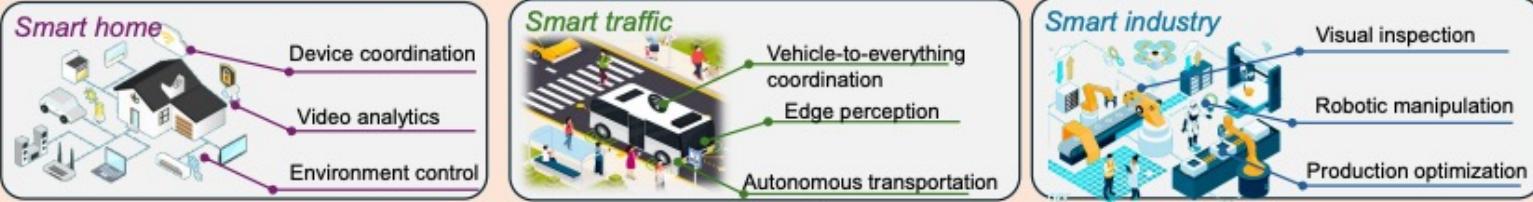


Agentic AI System with Microservices Architecture

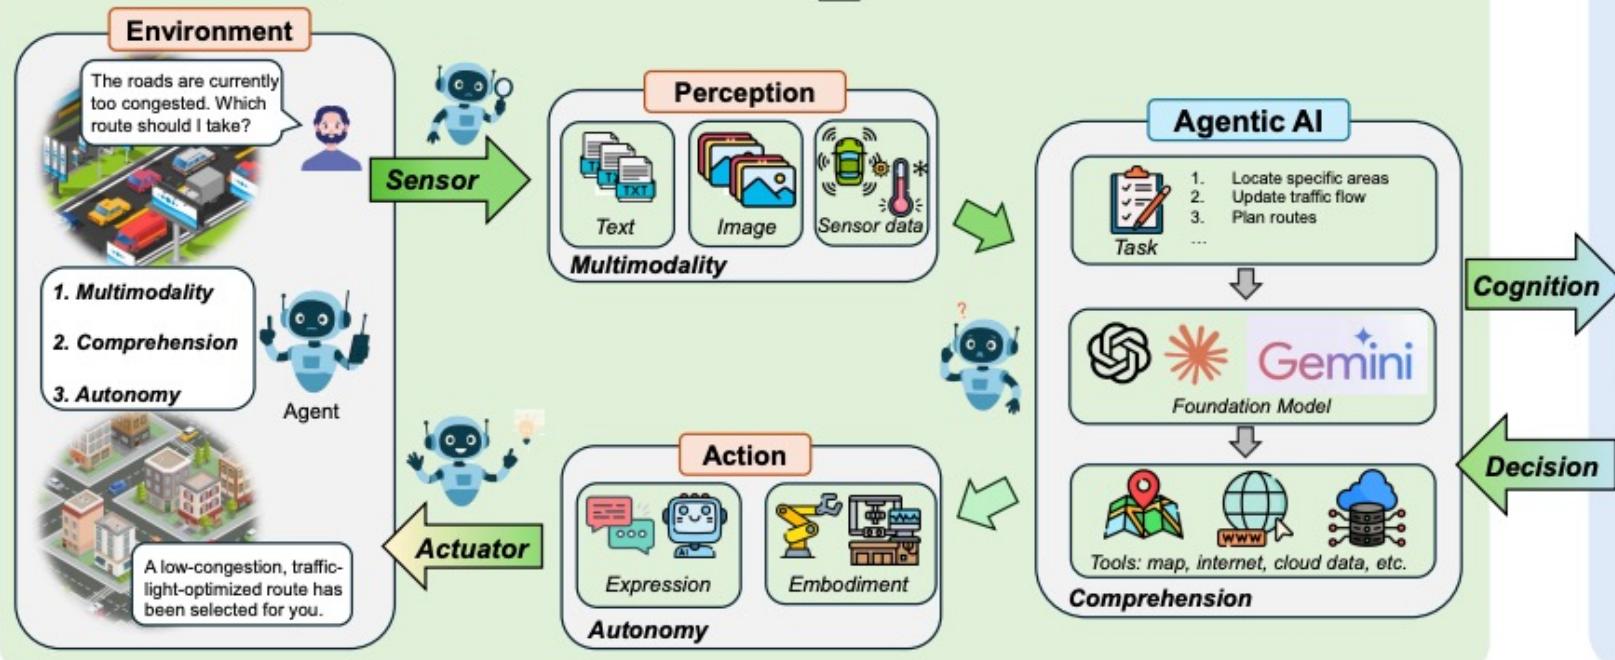


Agentic AI and World Model for Edge General Intelligence

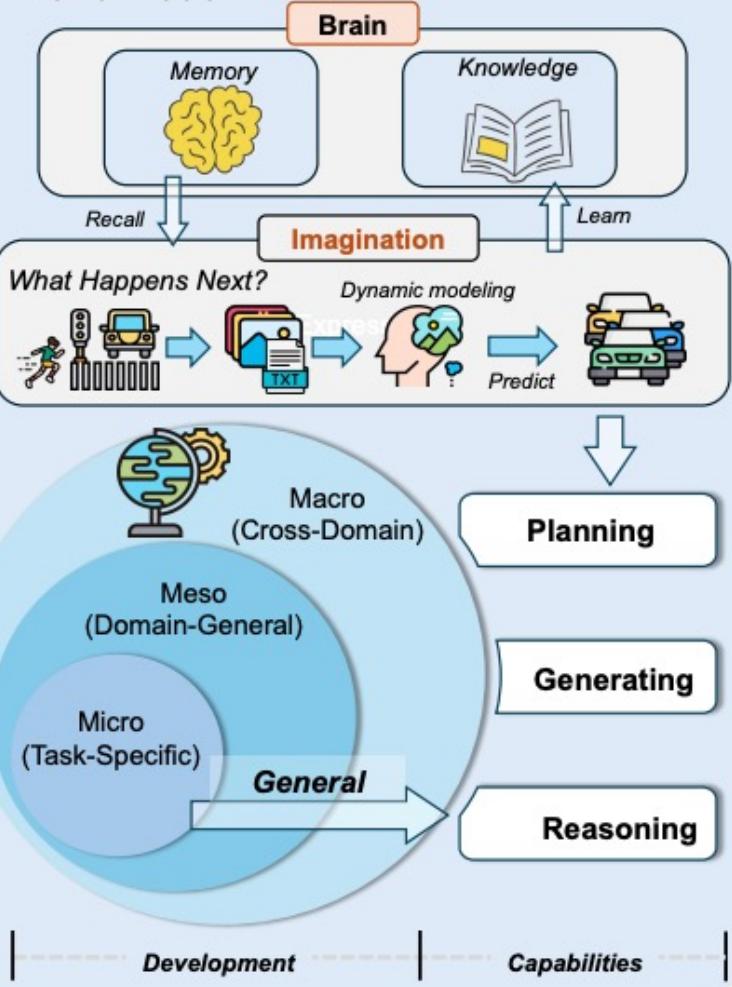
A. Edge General Intelligence



B. Workflow of agentic AI



C. World Model



Source: Changyuan Zhao, Guangyuan Liu, Ruichen Zhang, Yinqiu Liu, Jiacheng Wang, Jiawen Kang, Dusit Niyato et al (2025),

"Edge general intelligence through world models and agentic AI: Fundamentals, solutions, and challenges." arXiv preprint arXiv:2508.09561.

Artificial Intelligence Problem Solving

Artificial Intelligence: A Modern Approach

- 1. Artificial Intelligence**
- 2. Problem Solving**
- 3. Knowledge and Reasoning**
- 4. Uncertain Knowledge and Reasoning**
- 5. Machine Learning**
- 6. Communicating, Perceiving, and Acting**
- 7. Philosophy and Ethics of AI**

Artificial Intelligence:

2. Problem Solving

- Solving Problems by Searching
- Search in Complex Environments
- Adversarial Search and Games
- Constraint Satisfaction Problems

LLM-enhanced Problem Solving

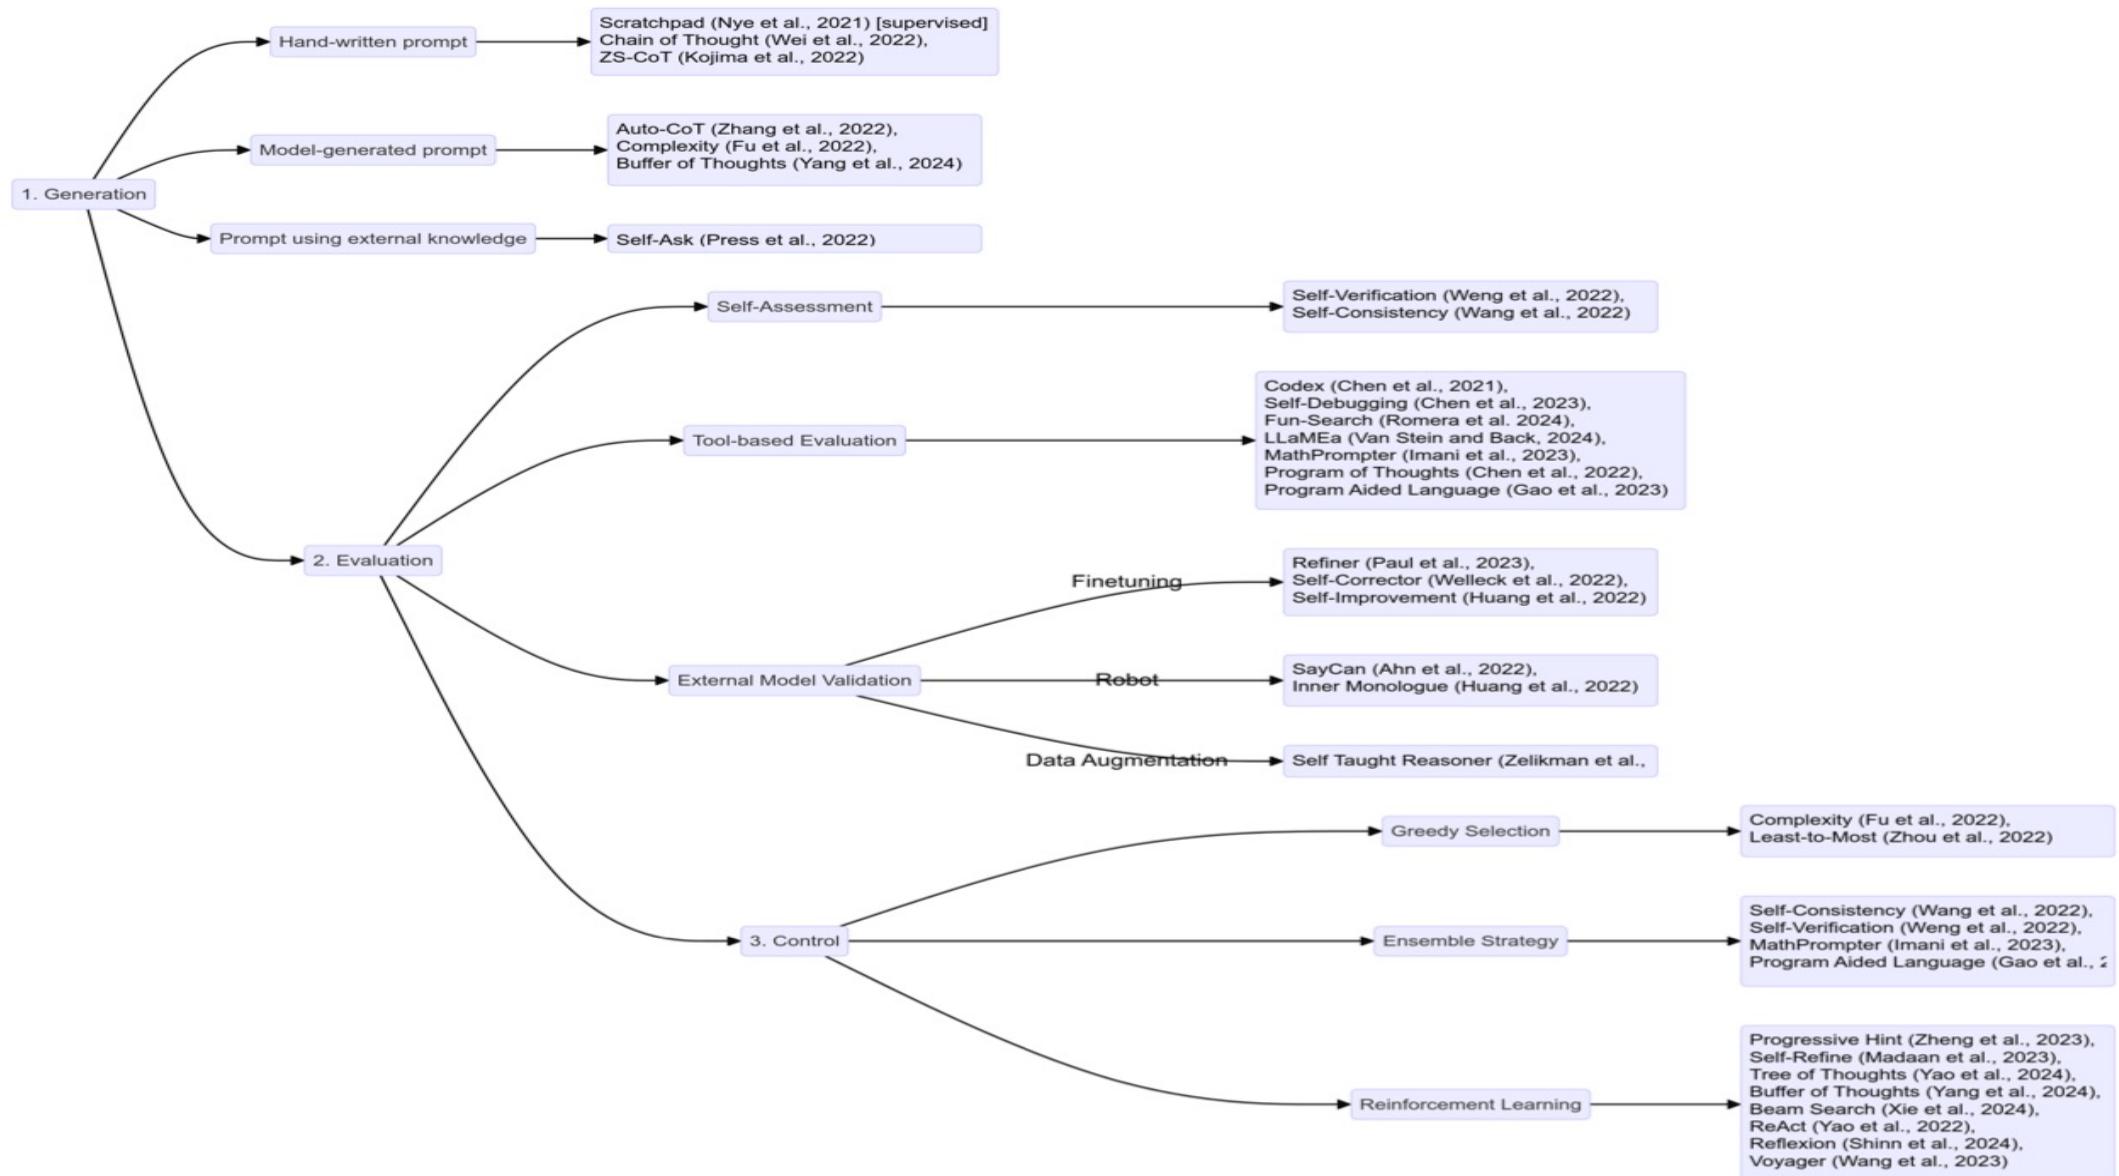
Problem Solving and LLM-enhanced Techniques

- Traditional Search Algorithms
 - Breadth-first Search (BFS)
 - Depth-first Search (DFS)
 - A* Search
- LLM-enhanced Problem Solving
 - Chain-of-thought (CoT) prompting
 - Few-shot learning for problem decomposition
 - Integration with external tools and APIs

LLM-based Agents for Complex Problem Solving

- **ReAct: Reasoning and Acting in Language Models**
- **MRKL (Modular Reasoning, Knowledge and Language) systems**
- **LLM-powered planning and decision making**
- **Chain-of-thought Prompting**
 - **Solving a complex problem
using chain-of-thought prompting**

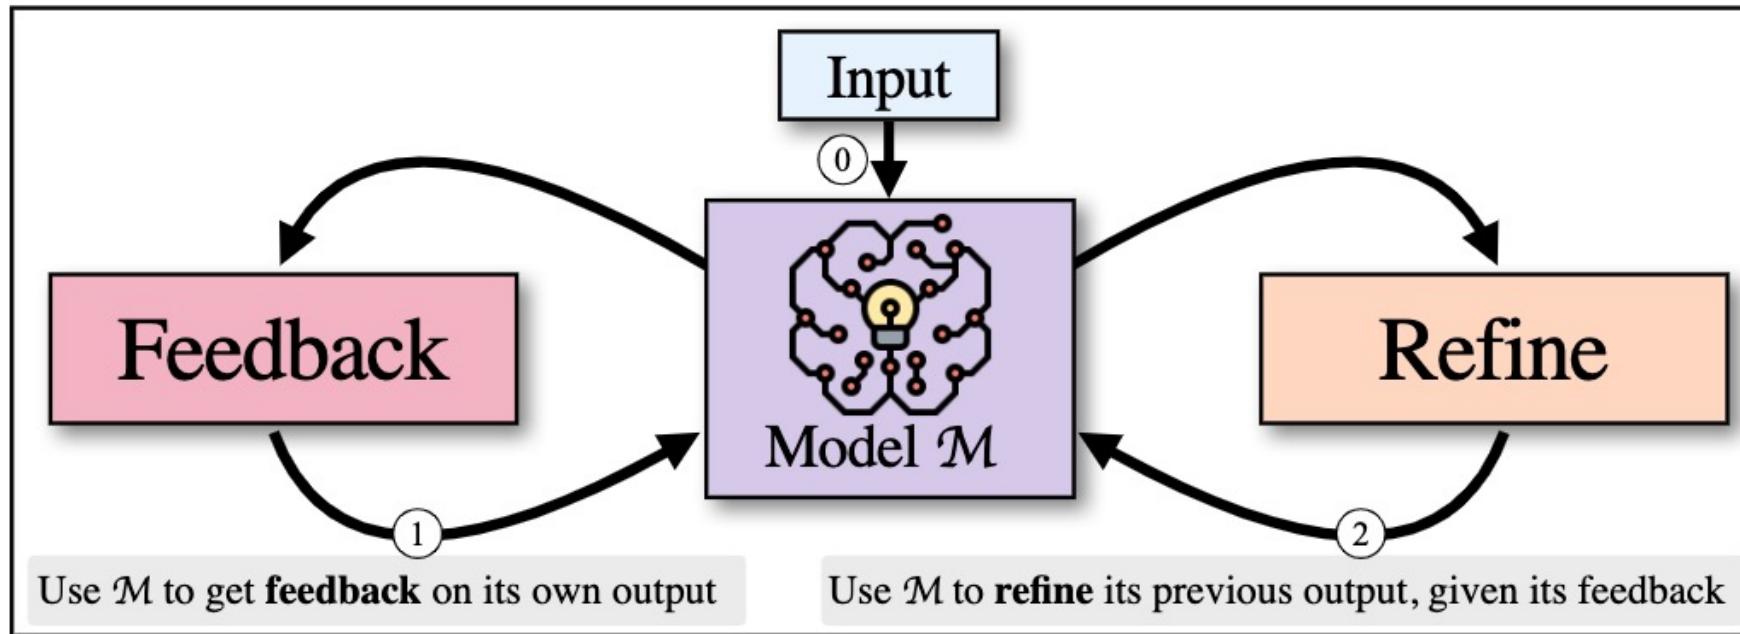
LLM-Reasoning Approaches: Prompt Generation, Evaluation, and Control



LLM-Reasoning Approaches: Prompt Generation, Evaluation, and Control

Approach	Domain	Step generation	Step evaluation	Step control
Scratchpad [Nye et al., 2021]	math word	hand-wr/supervised	-	greedy/1prompt
Chain-of-thought [Wei et al., 2022b]	math word	hand-written	-	greedy/1prompt
ZS-CoT [Kojima et al., 2022]	math word	hand-written	-	greedy/1prompt
Auto-CoT [Zhang et al., 2022]	math word	model-generated	-	clustering
Complexity [Fu et al., 2022]	math word	hand-written	self-consistency	greedy/1prompt
Self-ask [Press et al., 2022]	math word	external knowledge	LLM	multi-hop questions
Self-verification [Weng et al., 2022]	math word	hand-written	back-verify	ensemble
Self-consistency [Wang et al., 2022b]	math word	hand-written	majority	ensemble
Codex [Chen et al., 2021]	code	-	tool-based	-
Self-debugging [Chen et al., 2023]	code	hand-written	tool-based	greedy
Fun-search [Romera-Paredes et al., 2024]	code	hand-written	tool-based	evolutionary algorithm
LLaM Ea [van Stein and Bäck, 2024]	code	hand-written	tool-based	evolutionary algorithm
MathPrompter [Imani et al., 2023]	math	hand-written	tool-based	ensemble
Program-of-thoughts [Chen et al., 2022]	math word	hand-written, Codex	Python+Consist.	decouple reason/compute
Program-aided-language [Gao et al., 2023]	math word	hand-written, Codex	NLP/Python	ensemble
Refiner [Paul et al., 2023]	math word	finetune	critic model	gen/crit feedback
Self-corrector [Welleck et al., 2022]	math word	finetune	corrector model	gen/corr feedback
Self-improvement [Huang et al., 2022a]	math word	finetune	self-assessment	CoT/consistency
Say-can [Ahn et al., 2022]	robot	model-generated	external model	greedy
Inner-monologue [Huang et al., 2022b]	robot	hand-written	various	greedy
Self-taught-reasoner [Zelikman et al., 2022]	math word	finetune	augmentation	greedy/feedback
Least-to-most [Zhou et al., 2022]	math word	hand-written	self-assessment	curriculum
Progressive-hint [Zheng et al., 2023]	math word	model-generated	self-assessment	stable prompt
Self-refine [Madaan et al., 2023]	math word	model-generated	self-assessment	greedy/feedback
Tree-of-thoughts [Yao et al., 2024]	puzzles	model-generated	self-assessment	BFS/DFS
Buffer-of-thoughts [Yang et al., 2024]	math word	thought template	self-assessment	buffer manager
Beam-search [Xie et al., 2024]	math word	model-generated	self-assessment	Beam Search
ReAct [Yao et al., 2022]	action	external knowledge	self-assessment	reinforcement learning
Reflexion [Shinn et al., 2024]	decision	model-generated	ext model	reinforcement learning
Voyager [Wang et al., 2023]	Minecraft	model-generated	Minecraft	reinforcement learning

Self-Refine: Iterative refinement with self-feedback



(a) Dialogue: x, y_t

User: I am interested in playing Table tennis.

Response: I'm sure it's a great way to socialize, stay active

(b) FEEDBACK fb

Engaging: Provides no information about table tennis or how to play it.

User understanding: Lacks understanding of user's needs and state of mind.

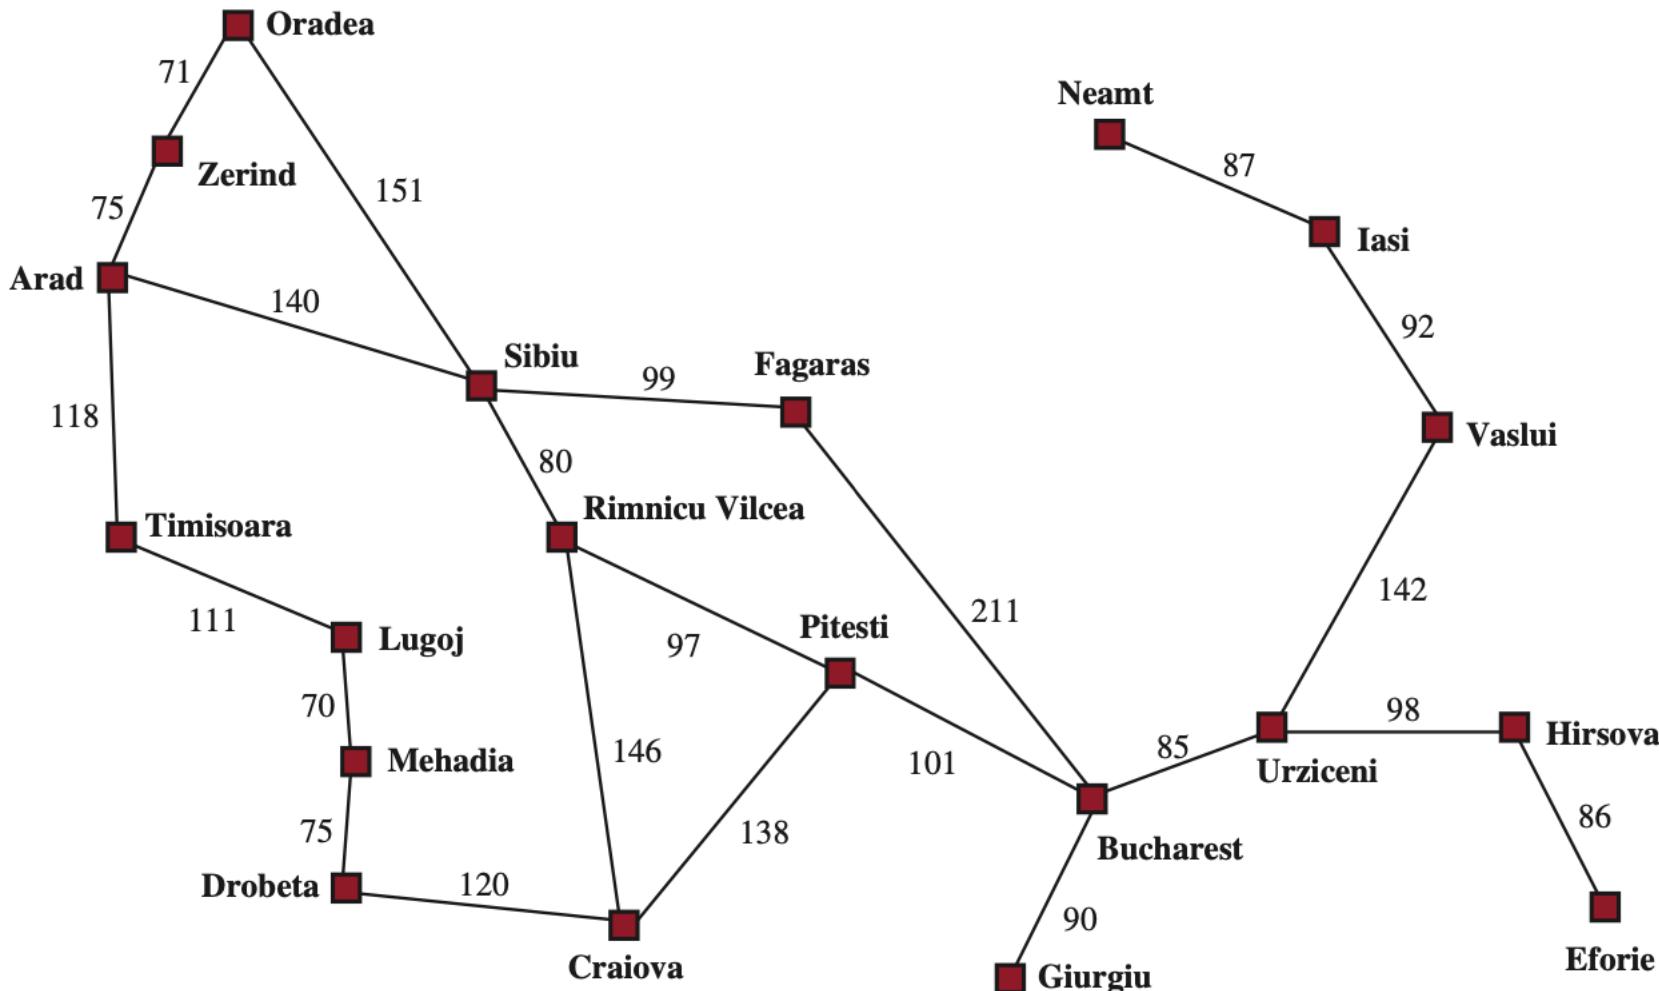
(c) REFINE y_{t+1}

Response (refined): That's great to hear (...) ! It's a fun sport requiring quick reflexes and good hand-eye coordination. Have you played before, or are you looking to learn?

Solving Problems by Searching

AI: Solving Problems by Searching

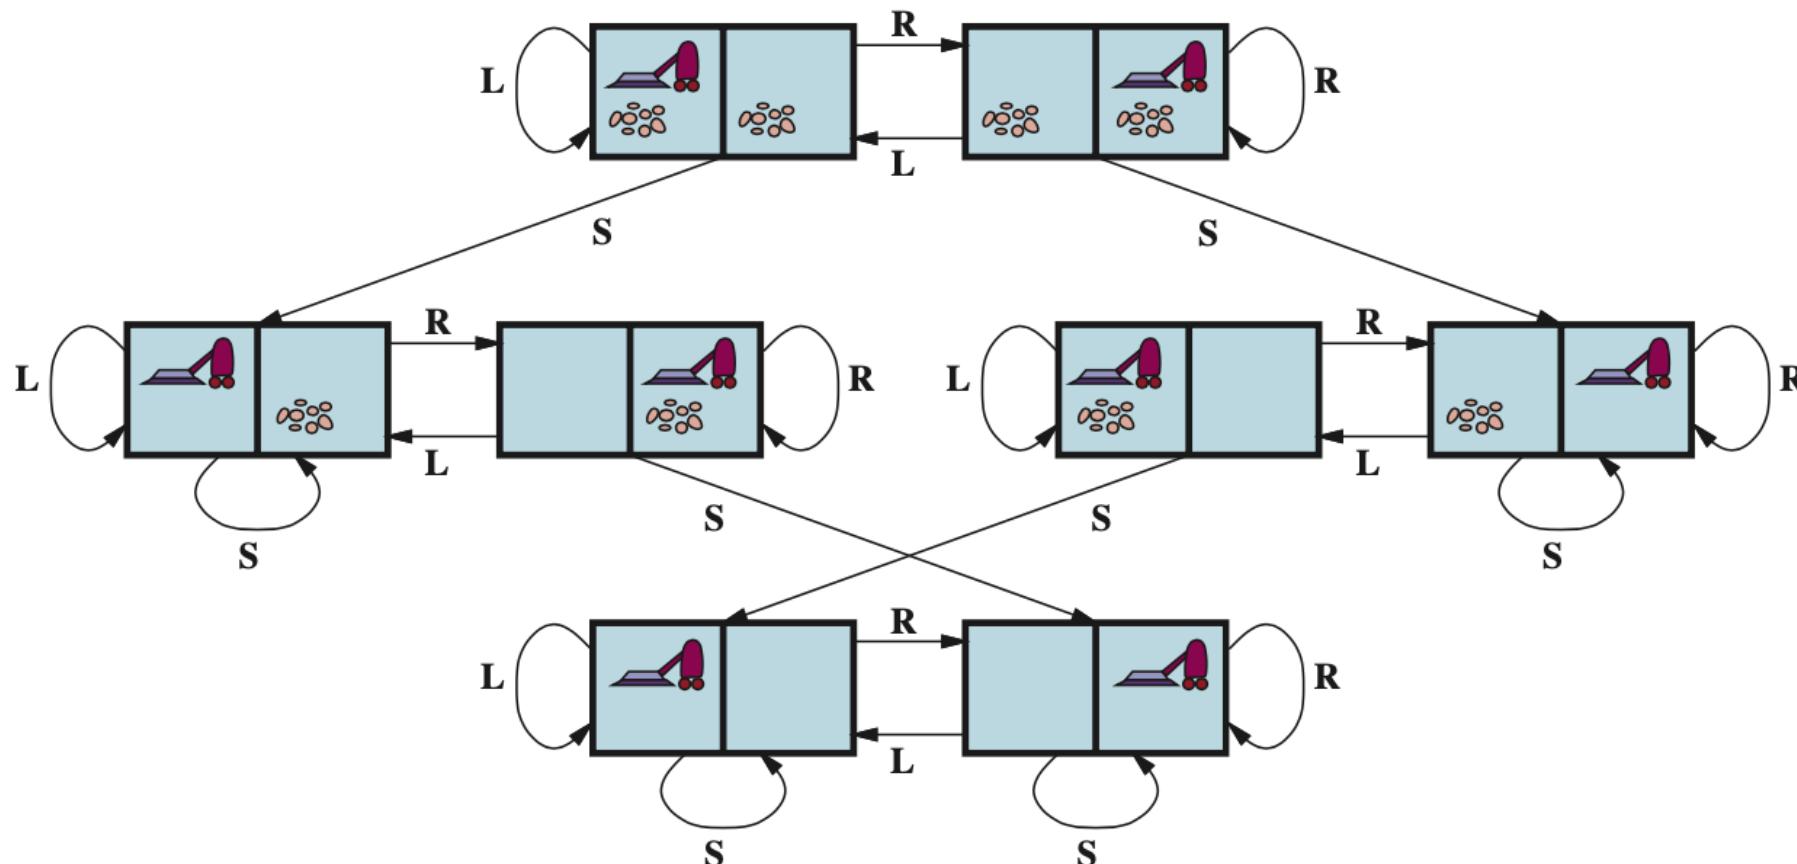
A simplified road map of part of Romania, with road distances in miles.



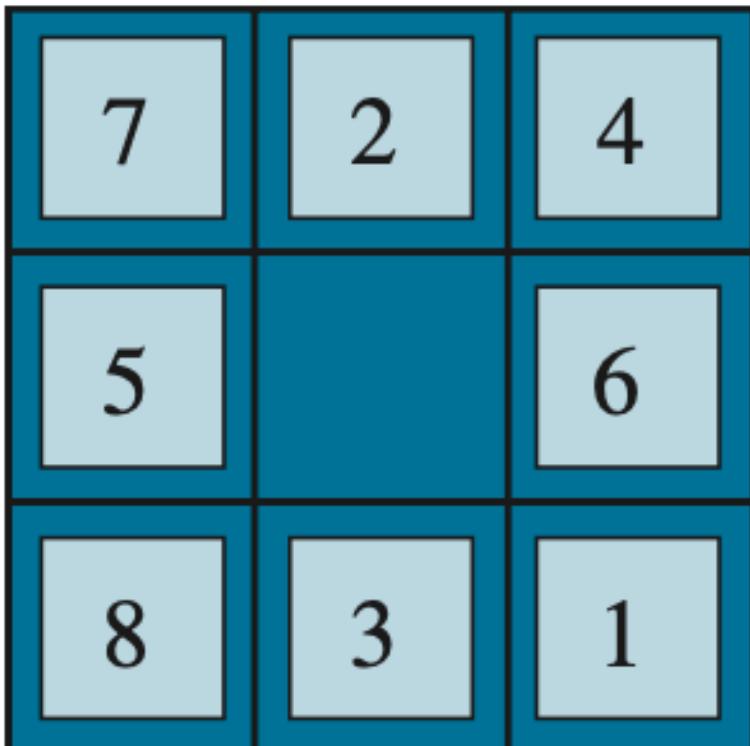
The state-space graph for the two-cell vacuum world

There are 8 states and three actions for each state:

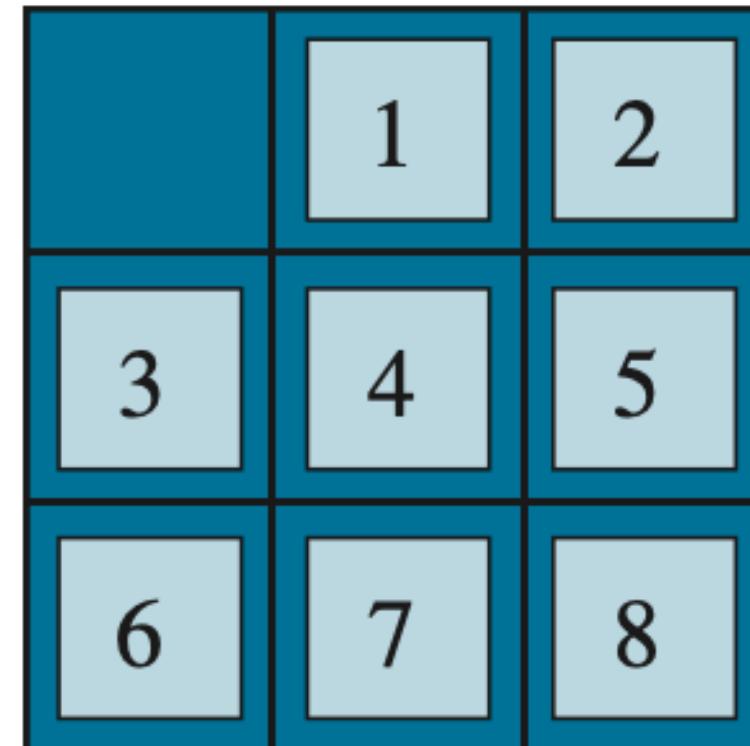
$L = Left$, $R = Right$, $S = Suck$.



A typical instance of the 8-puzzle

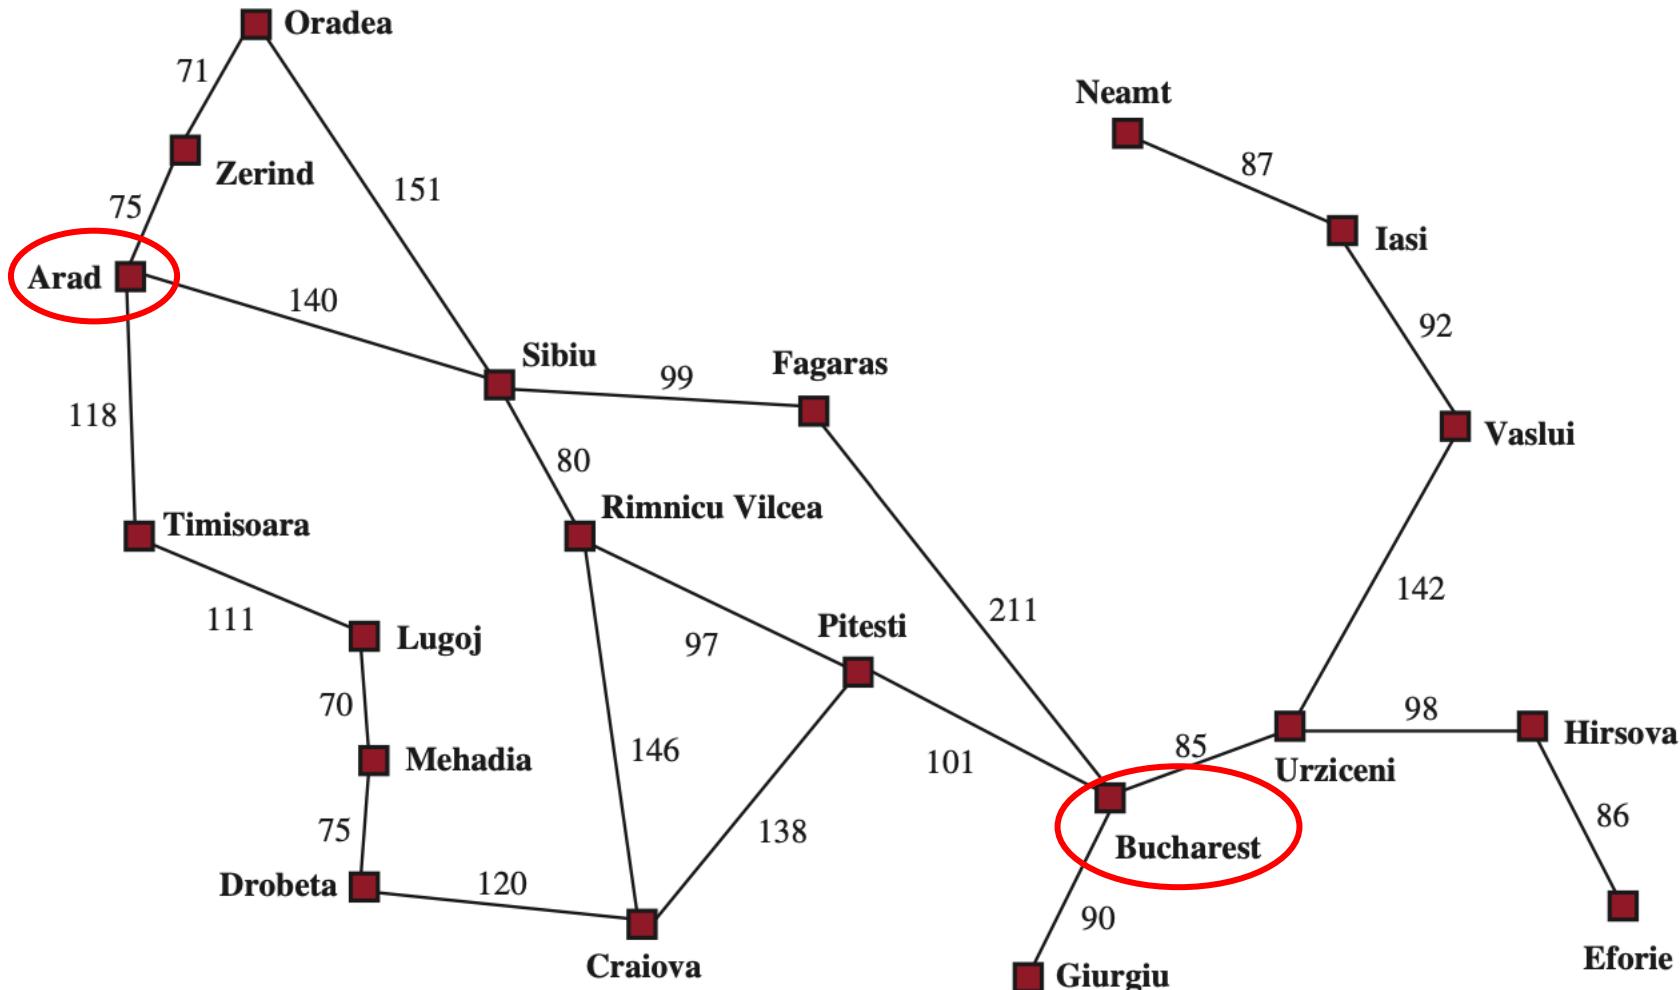


Start State

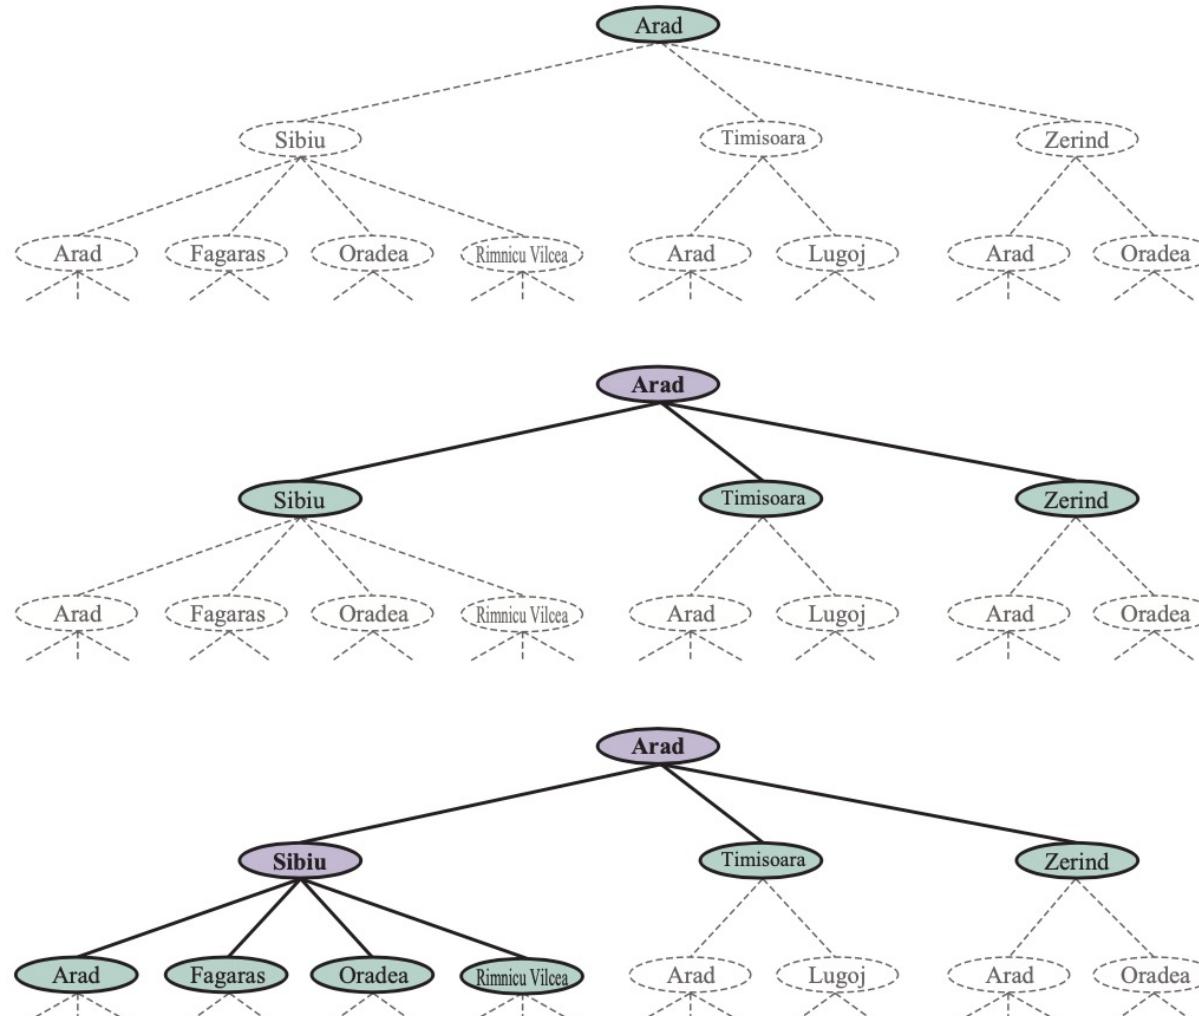


Goal State

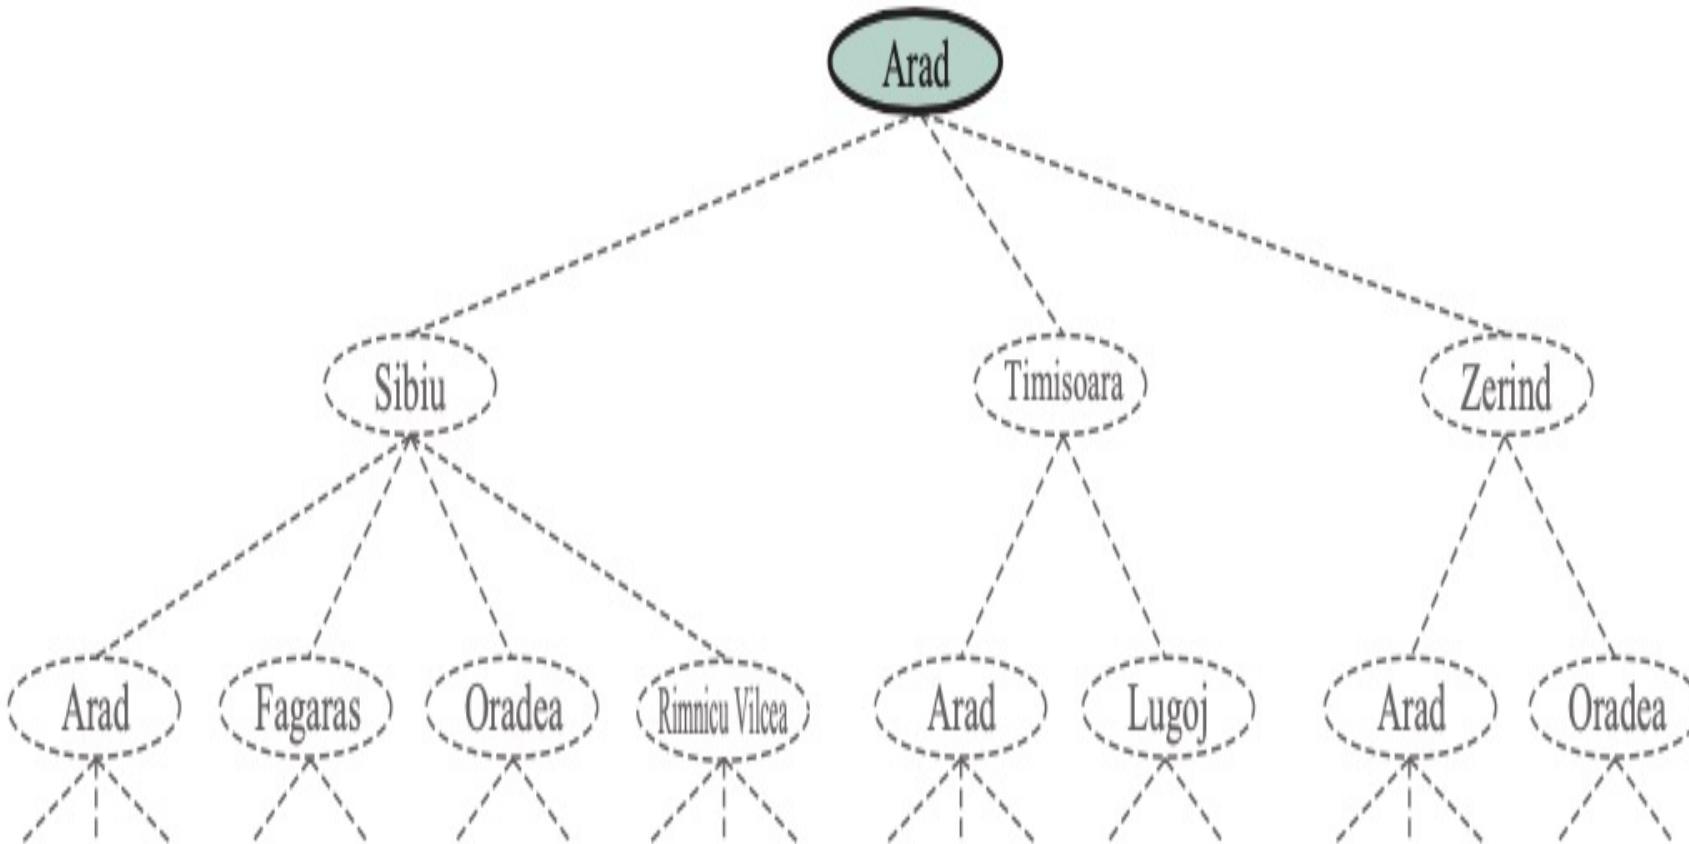
Arad to Bucharest



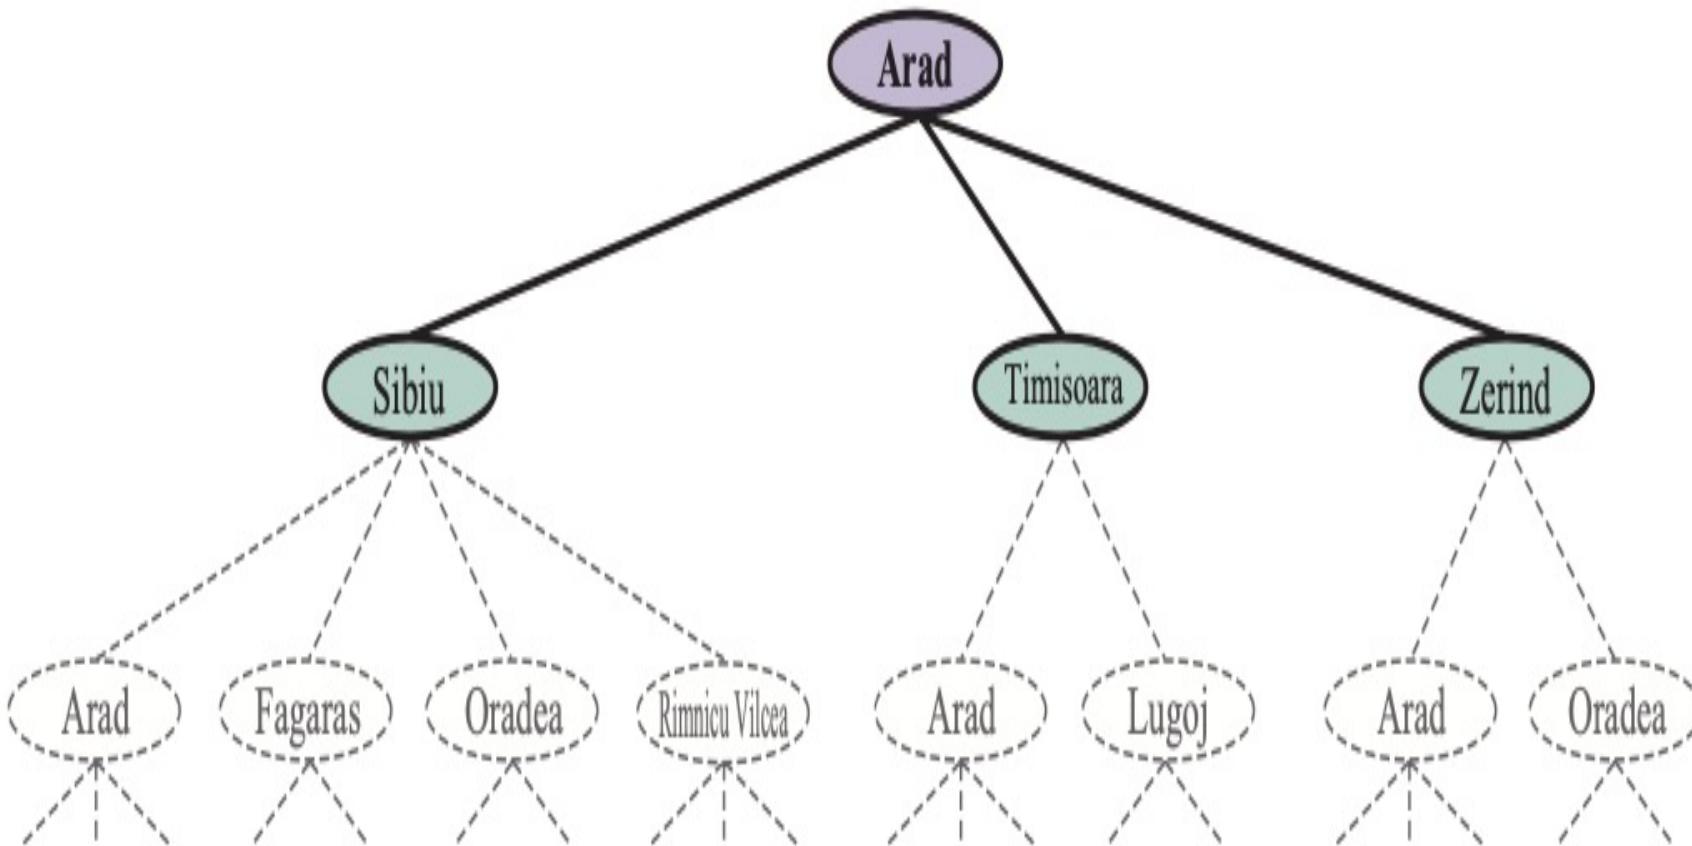
Three partial search trees for finding a route from Arad to Bucharest



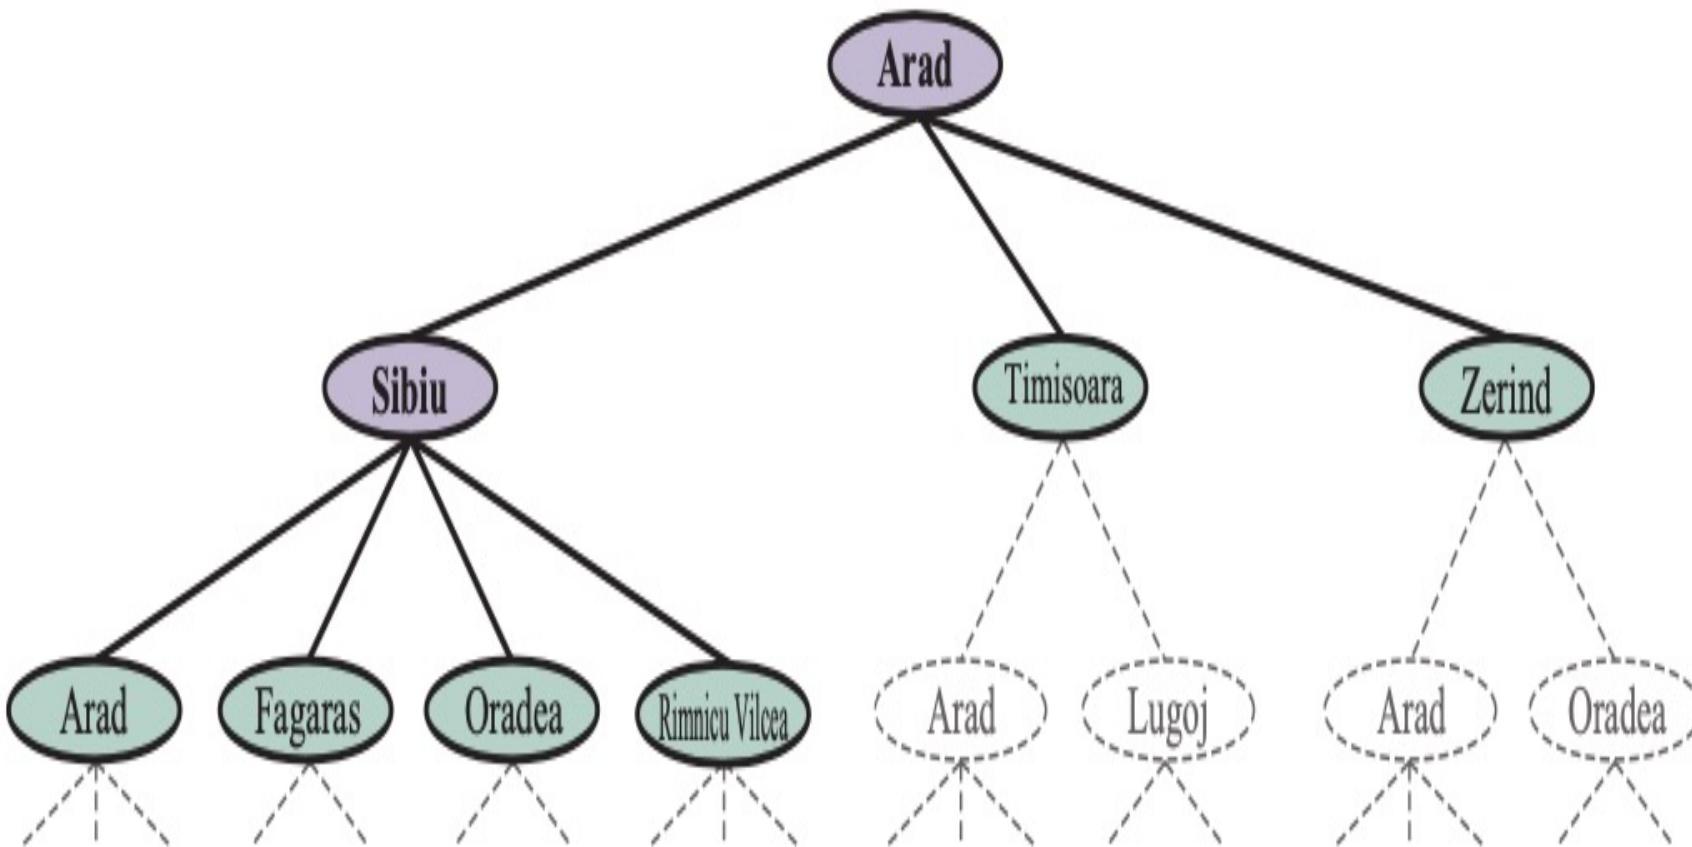
Three partial search trees for finding a route from Arad to Bucharest



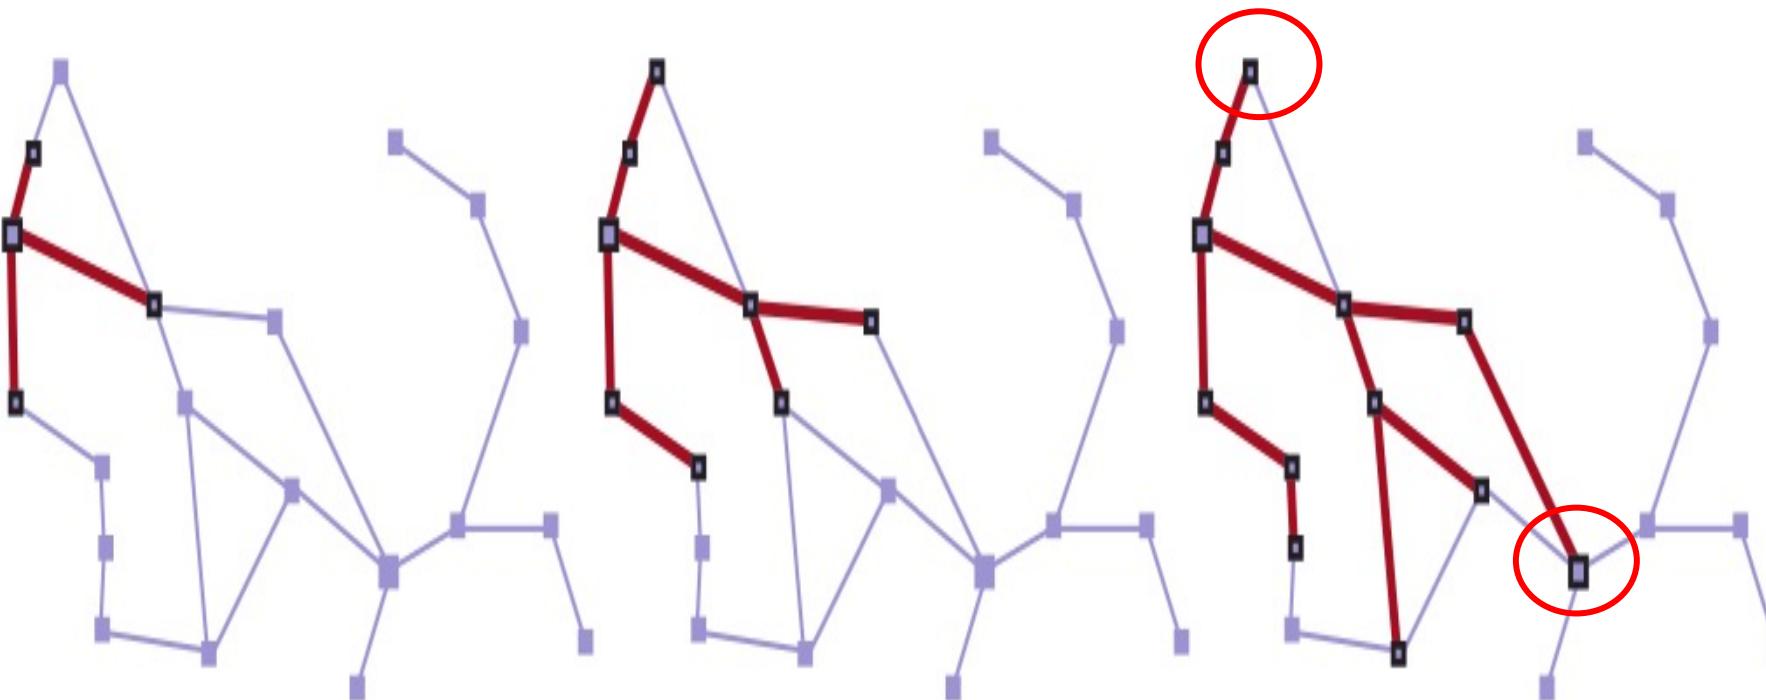
Three partial search trees for finding a route from Arad to Bucharest



Three partial search trees for finding a route from Arad to Bucharest



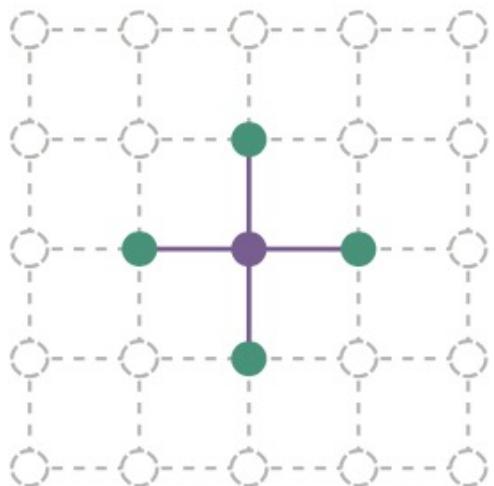
A sequence of search trees generated by a graph search on the Romania problem



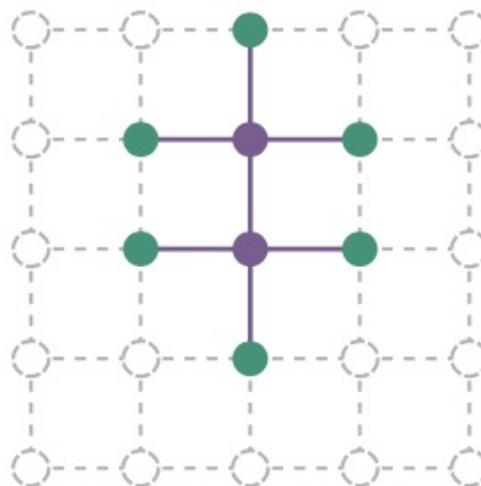
The Separation Property of Graph Search

illustrated on a rectangular-grid problem

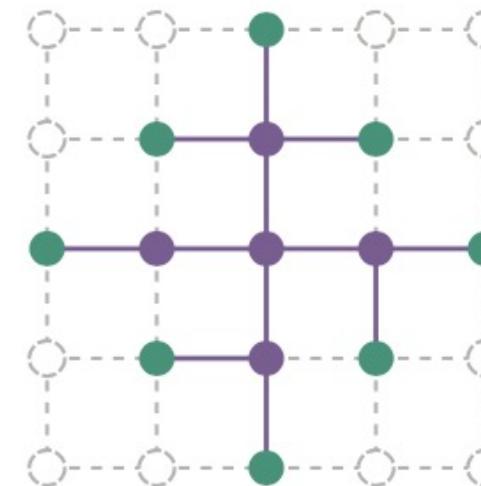
The frontier (green) separates
the interior (lavender) from
the exterior (faint dashed)



(a)



(b)



(c)

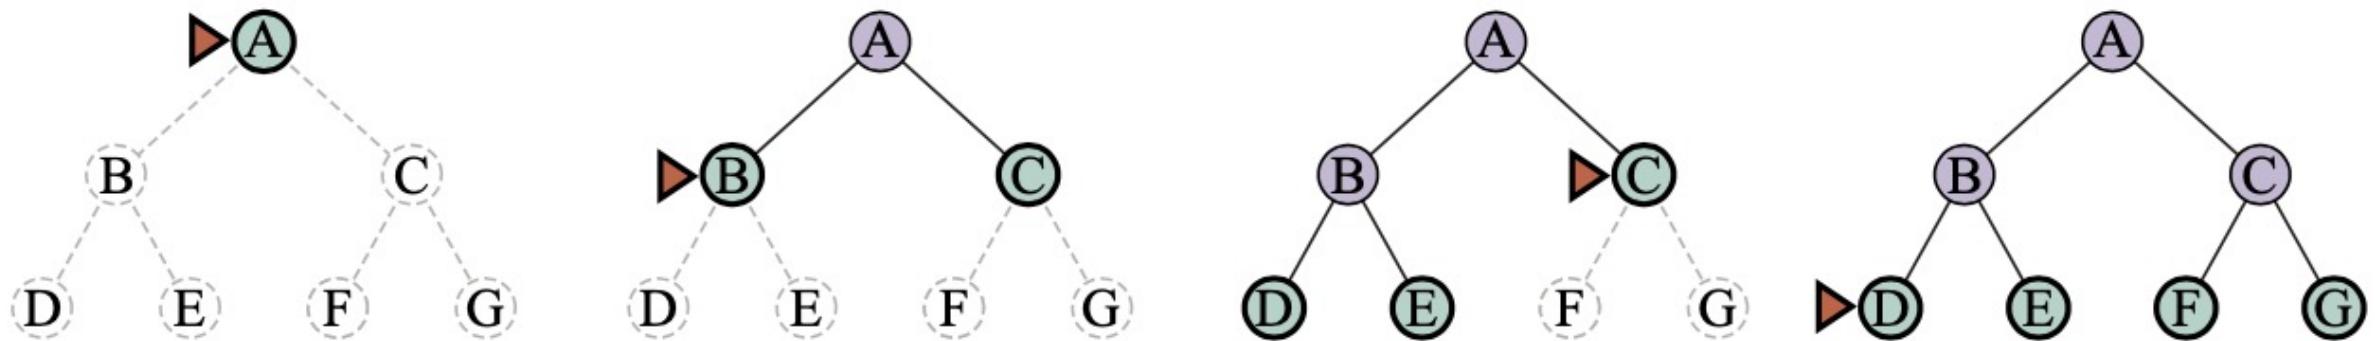
The Best-First Search (BFS) Algorithm

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
    node  $\leftarrow$  NODE(STATE=problem.INITIAL)
    frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
    reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
    while not Is-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        if problem.IS-GOAL(node.STATE) then return node
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
                reached[s]  $\leftarrow$  child
                add child to frontier
    return failure

function EXPAND(problem, node) yields nodes
    s  $\leftarrow$  node.STATE
    for each action in problem.ACTIONS(s) do
        s'  $\leftarrow$  problem.RESULT(s, action)
        cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

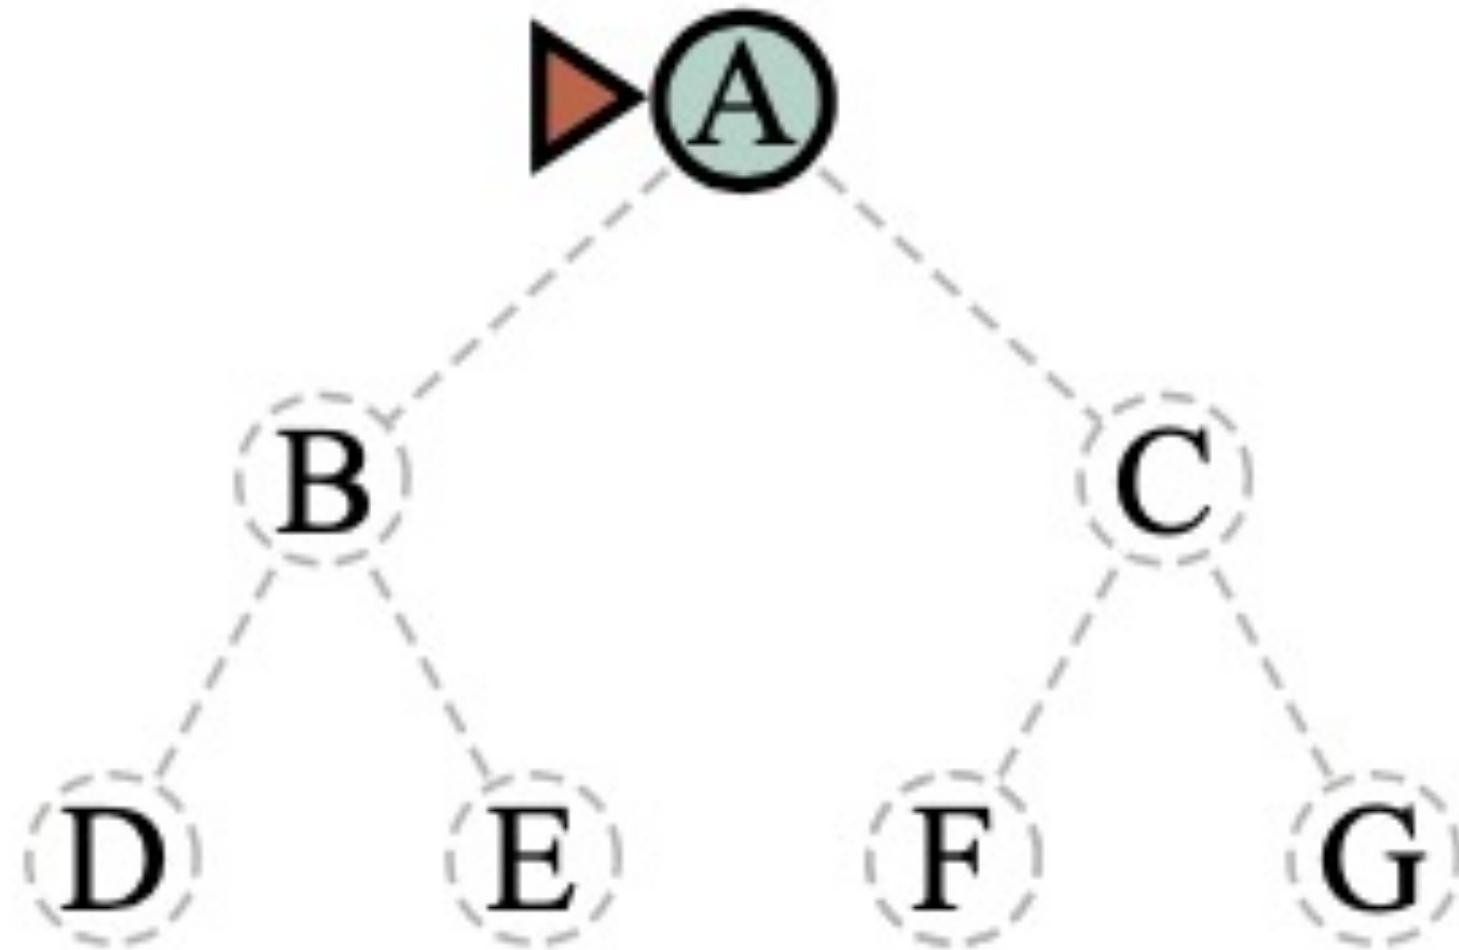
Breadth-First Search Search (BFS)

Breadth-First Search on a Simple Binary Tree



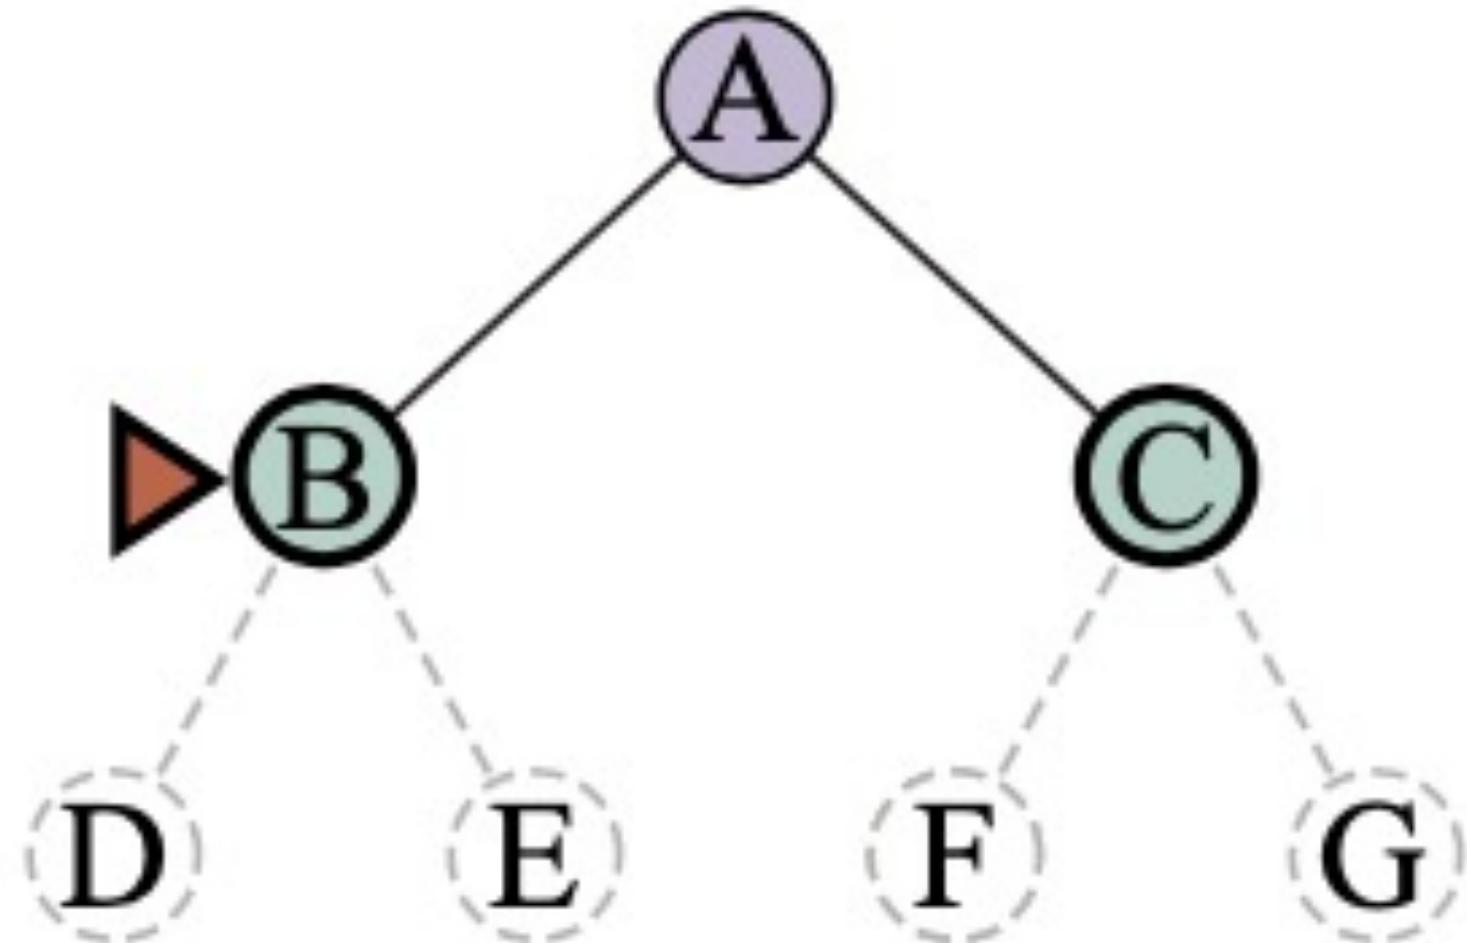
Breadth-First Search (BFS)

Breadth-First Search on a Simple Binary Tree



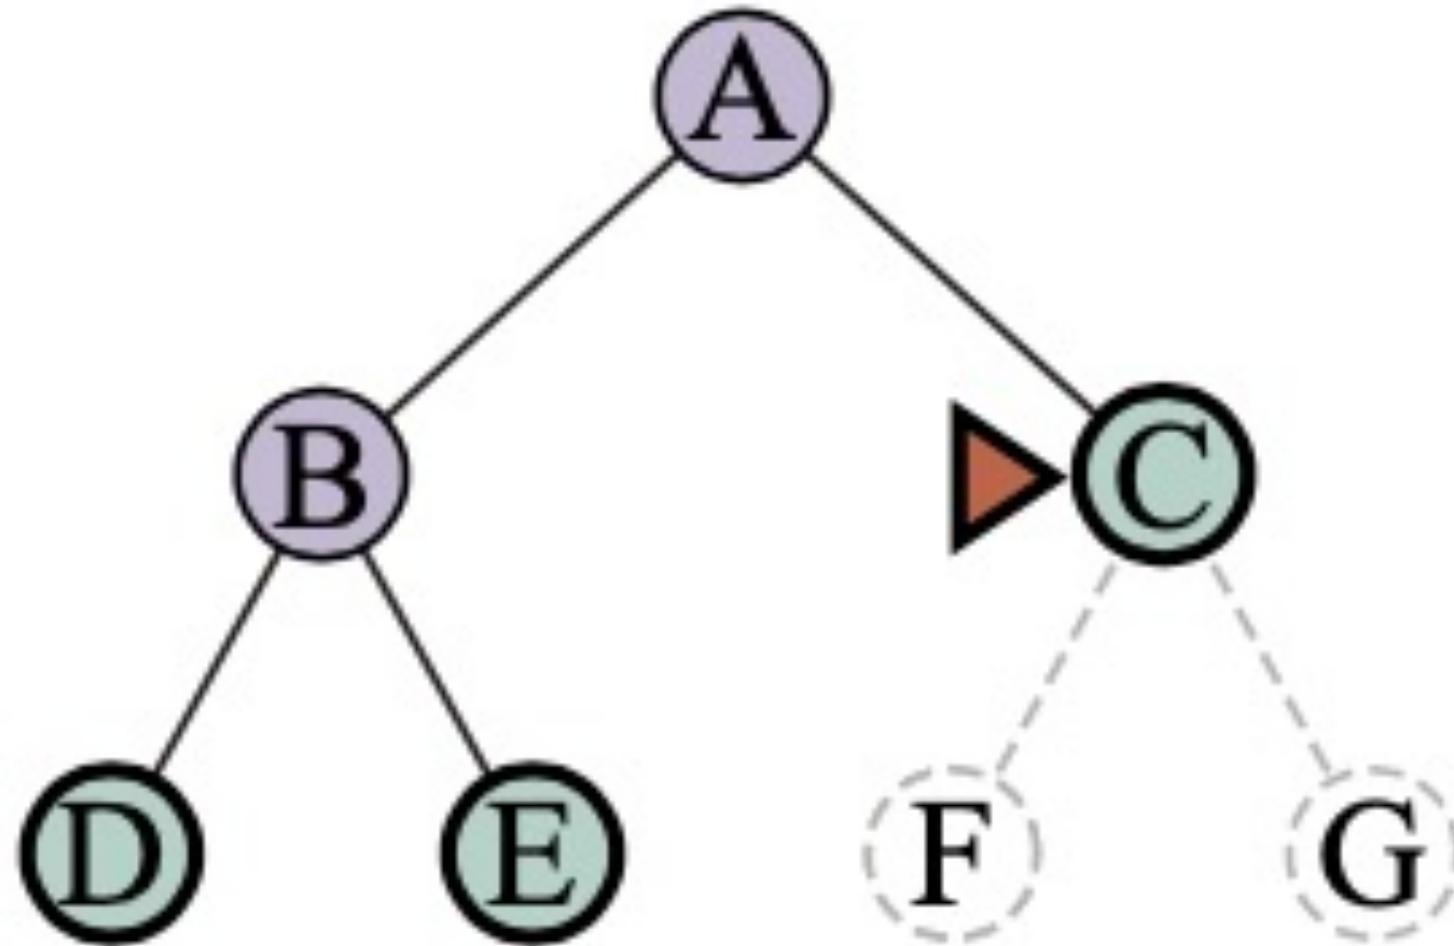
Breadth-First Search (BFS)

Breadth-First Search on a Simple Binary Tree



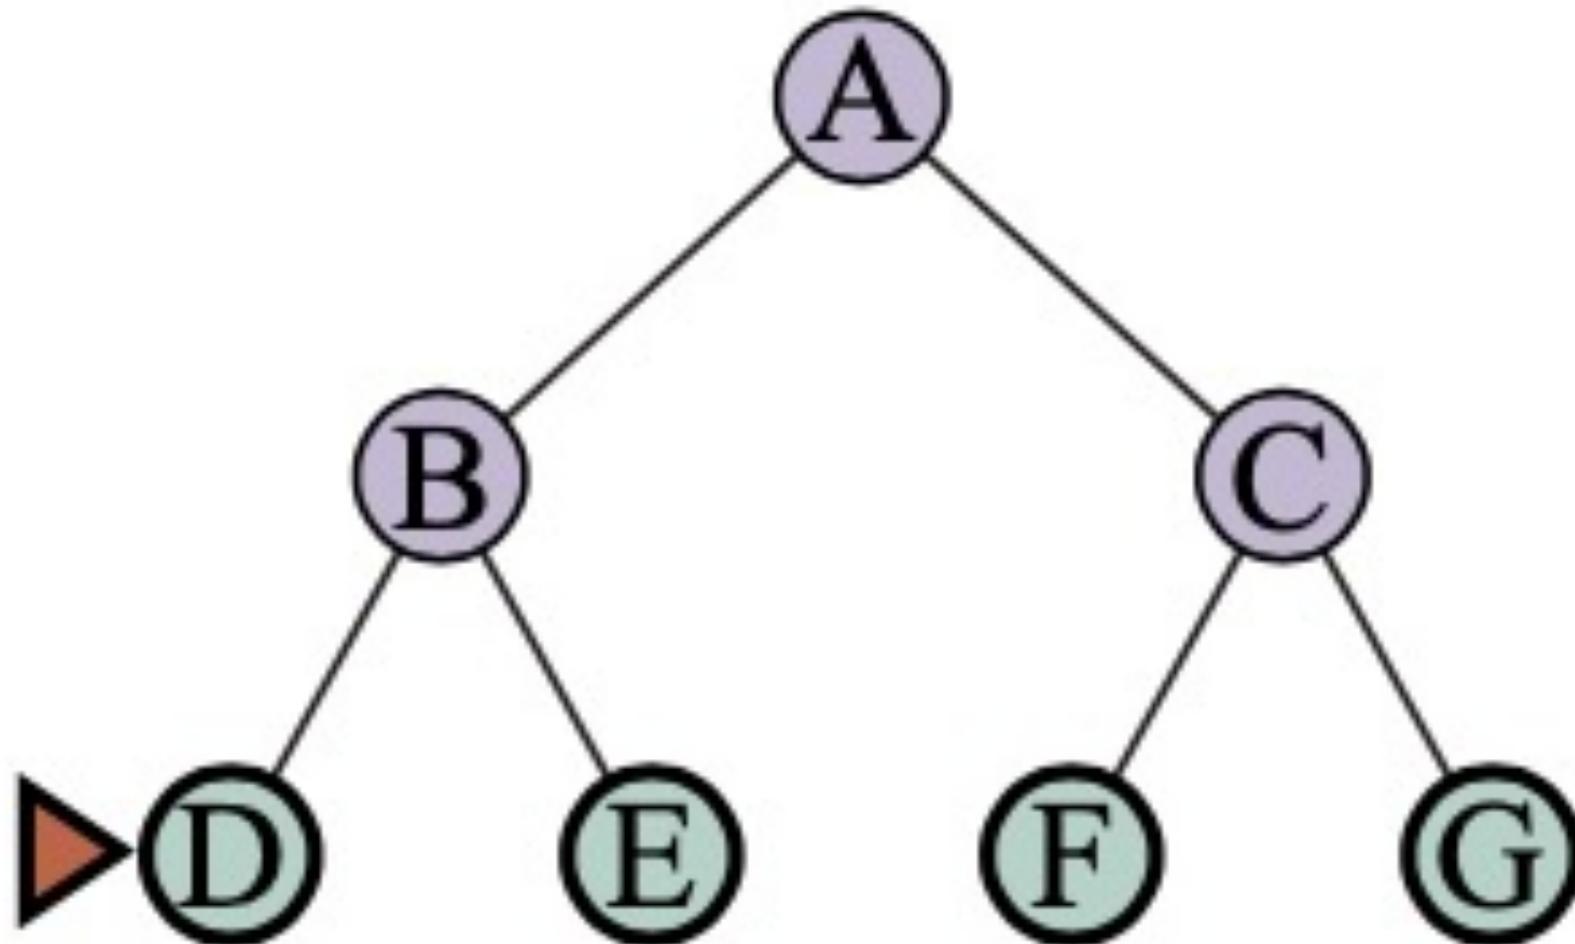
Breadth-First Search (BFS)

Breadth-First Search on a Simple Binary Tree



Breadth-First Search (BFS)

Breadth-First Search on a Simple Binary Tree

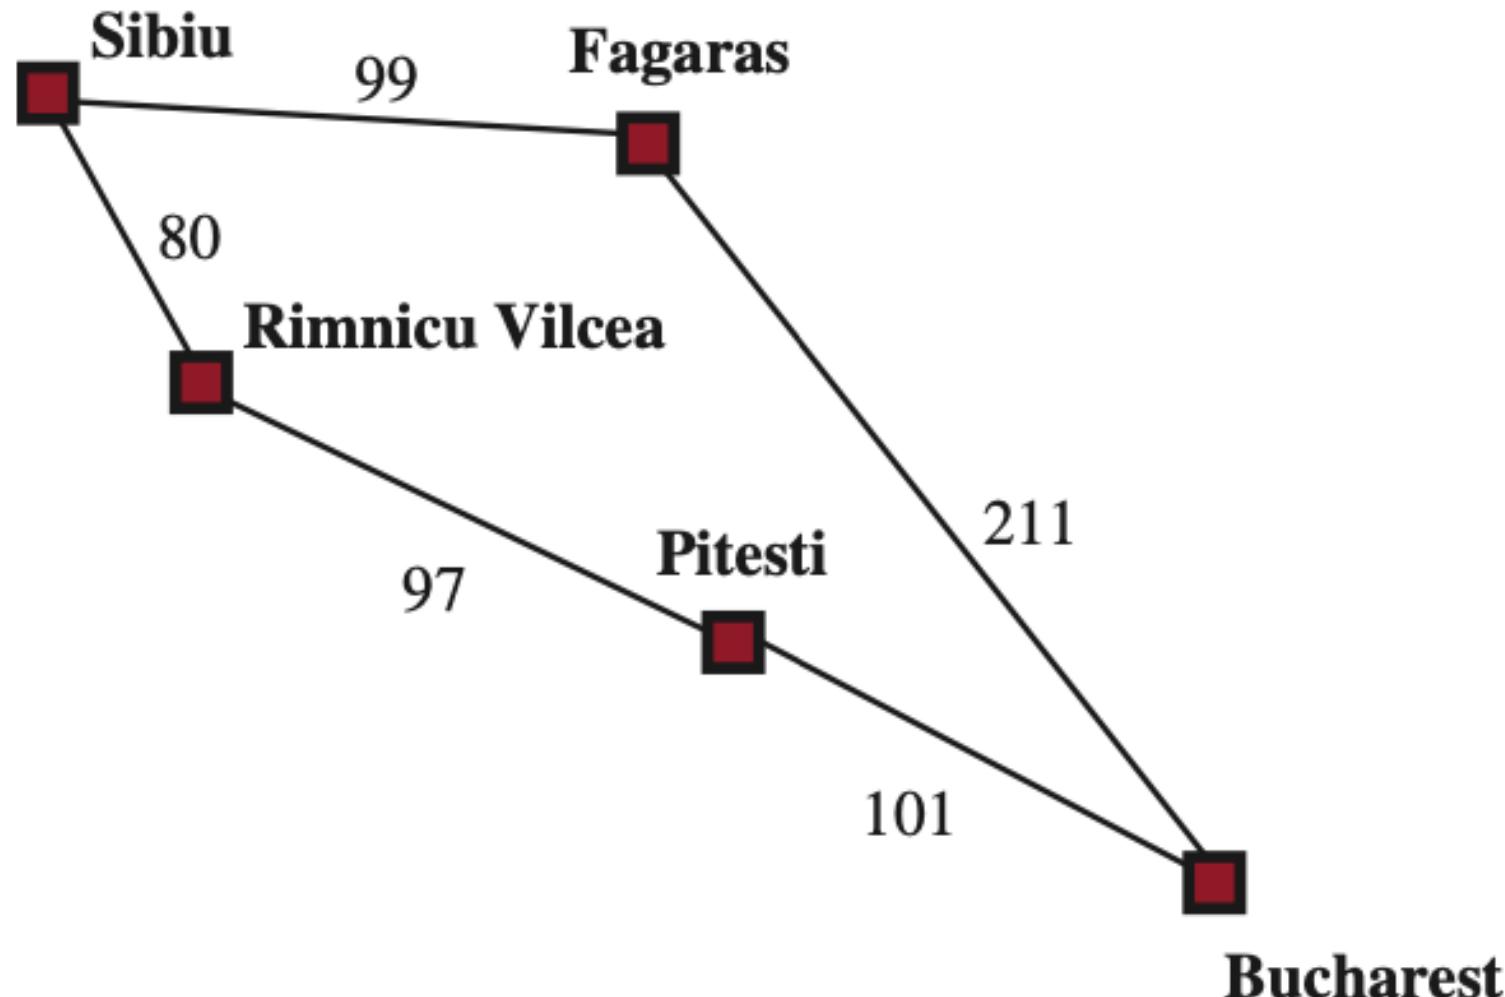


Breadth-First Search and Uniform-Cost Search Algorithms

```
function BREADTH-FIRST-SEARCH(problem) returns a solution node or failure
    node  $\leftarrow$  NODE(problem.INITIAL)
    if problem.Is-GOAL(node.STATE) then return node
    frontier  $\leftarrow$  a FIFO queue, with node as an element
    reached  $\leftarrow$  {problem.INITIAL}
    while not Is-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if problem.Is-GOAL(s) then return child
            if s is not in reached then
                add s to reached
                add child to frontier
    return failure
```

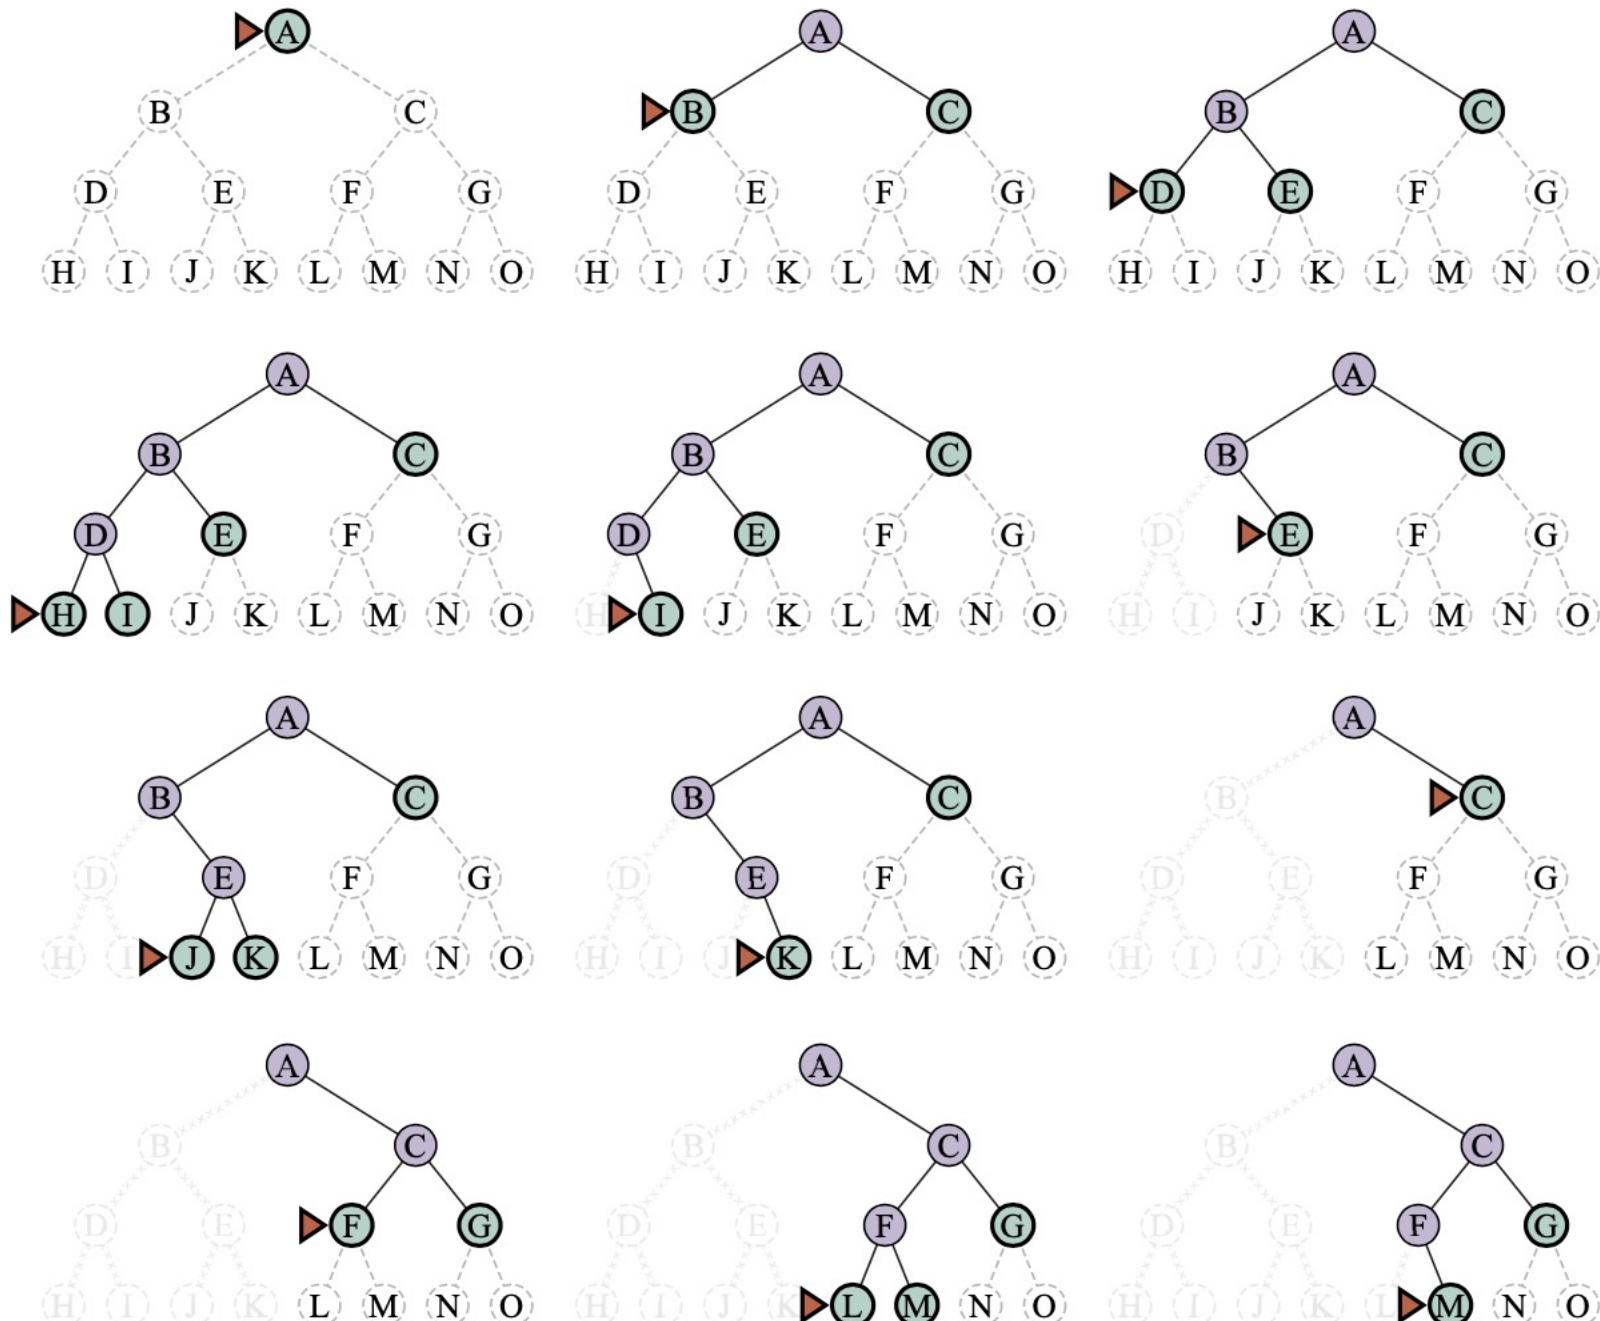
```
function UNIFORM-COST-SEARCH(problem) returns a solution node, or failure
    return BEST-FIRST-SEARCH(problem, PATH-COST)
```

Part of the Romania State Space Uniform-Cost Search



Depth-First Search (DFS)

Depth-First Search (DFS)

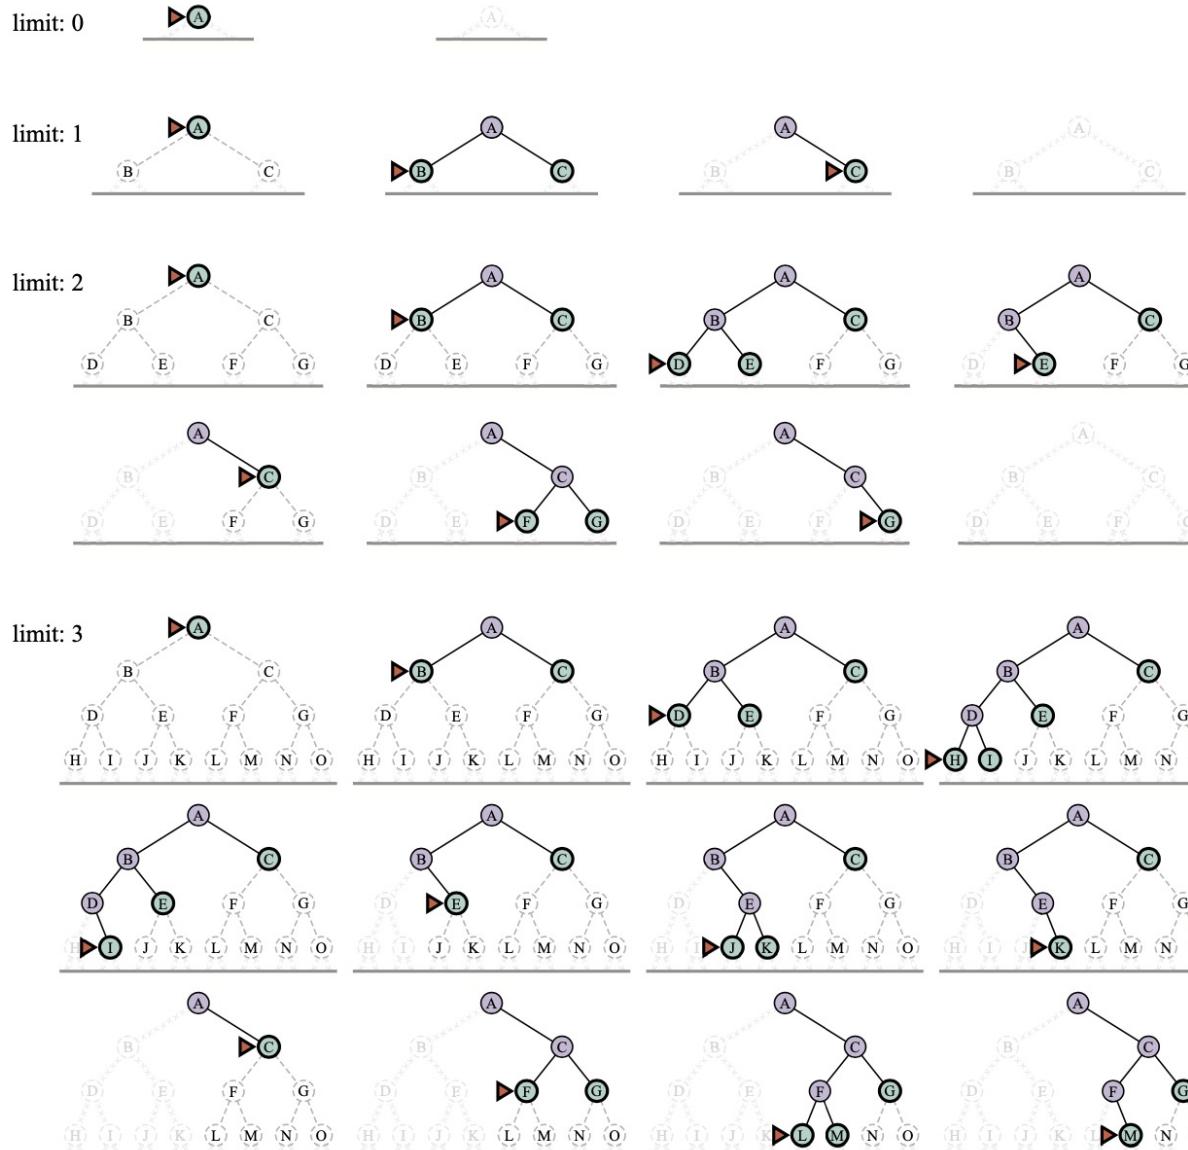


Iterative deepening and depth-limited tree-like search

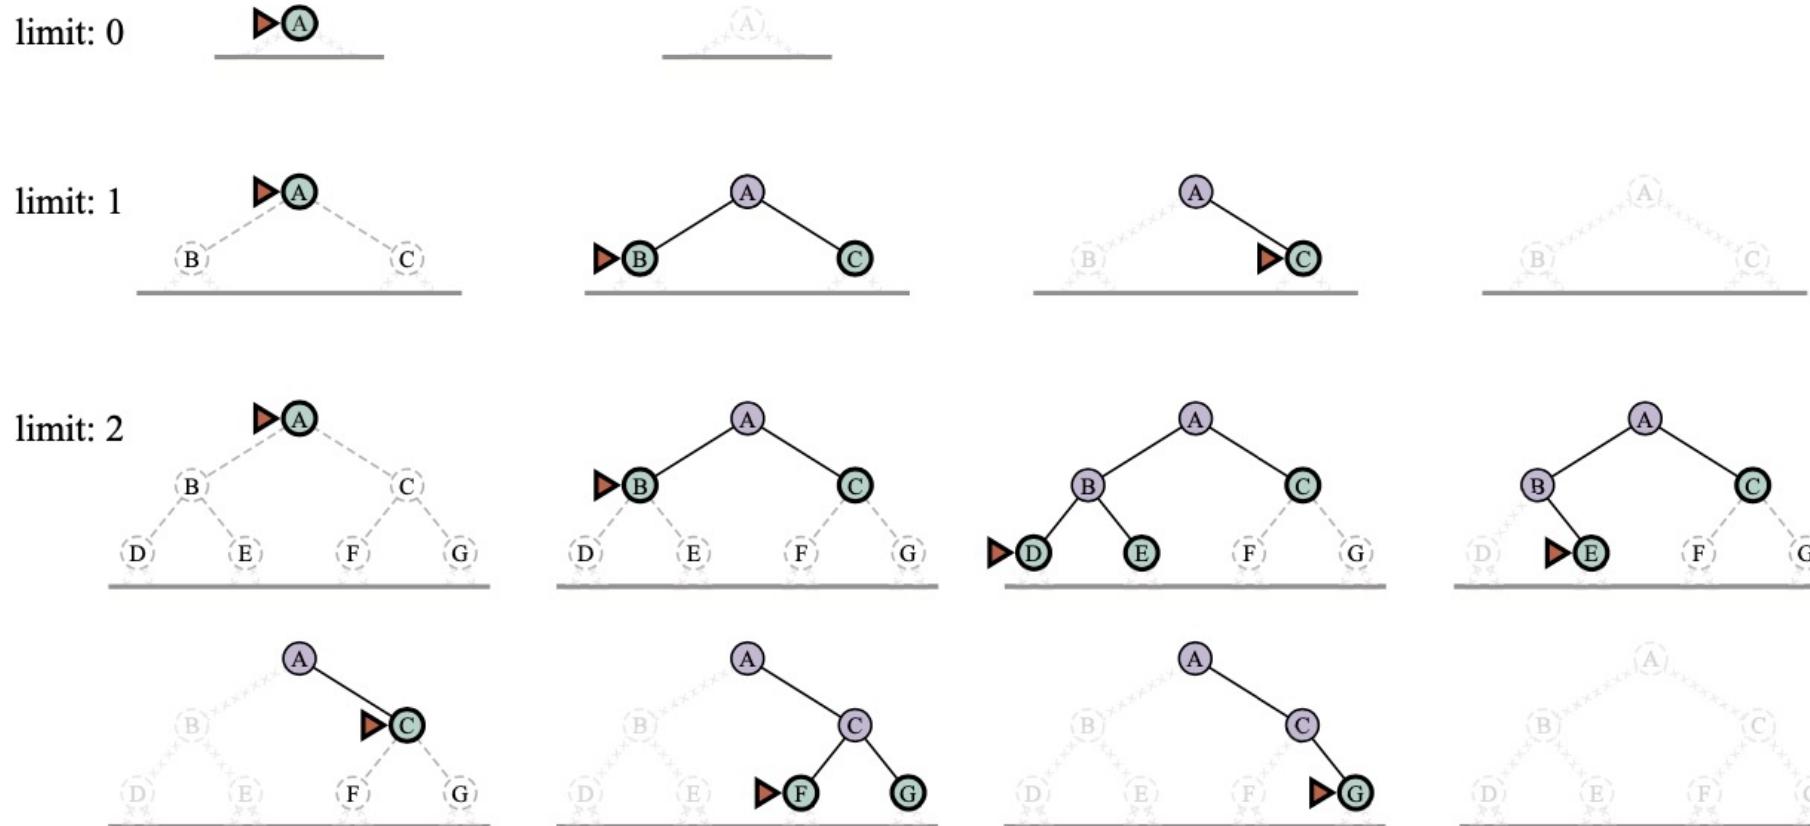
```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution node or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result

function DEPTH-LIMITED-SEARCH(problem,  $\ell$ ) returns a node or failure or cutoff
  frontier  $\leftarrow$  a LIFO queue (stack) with NODE(problem.INITIAL) as an element
  result  $\leftarrow$  failure
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    if DEPTH(node)  $>$   $\ell$  then
      result  $\leftarrow$  cutoff
    else if not IS-CYCLE(node) do
      for each child in EXPAND(problem, node) do
        add child to frontier
  return result
```

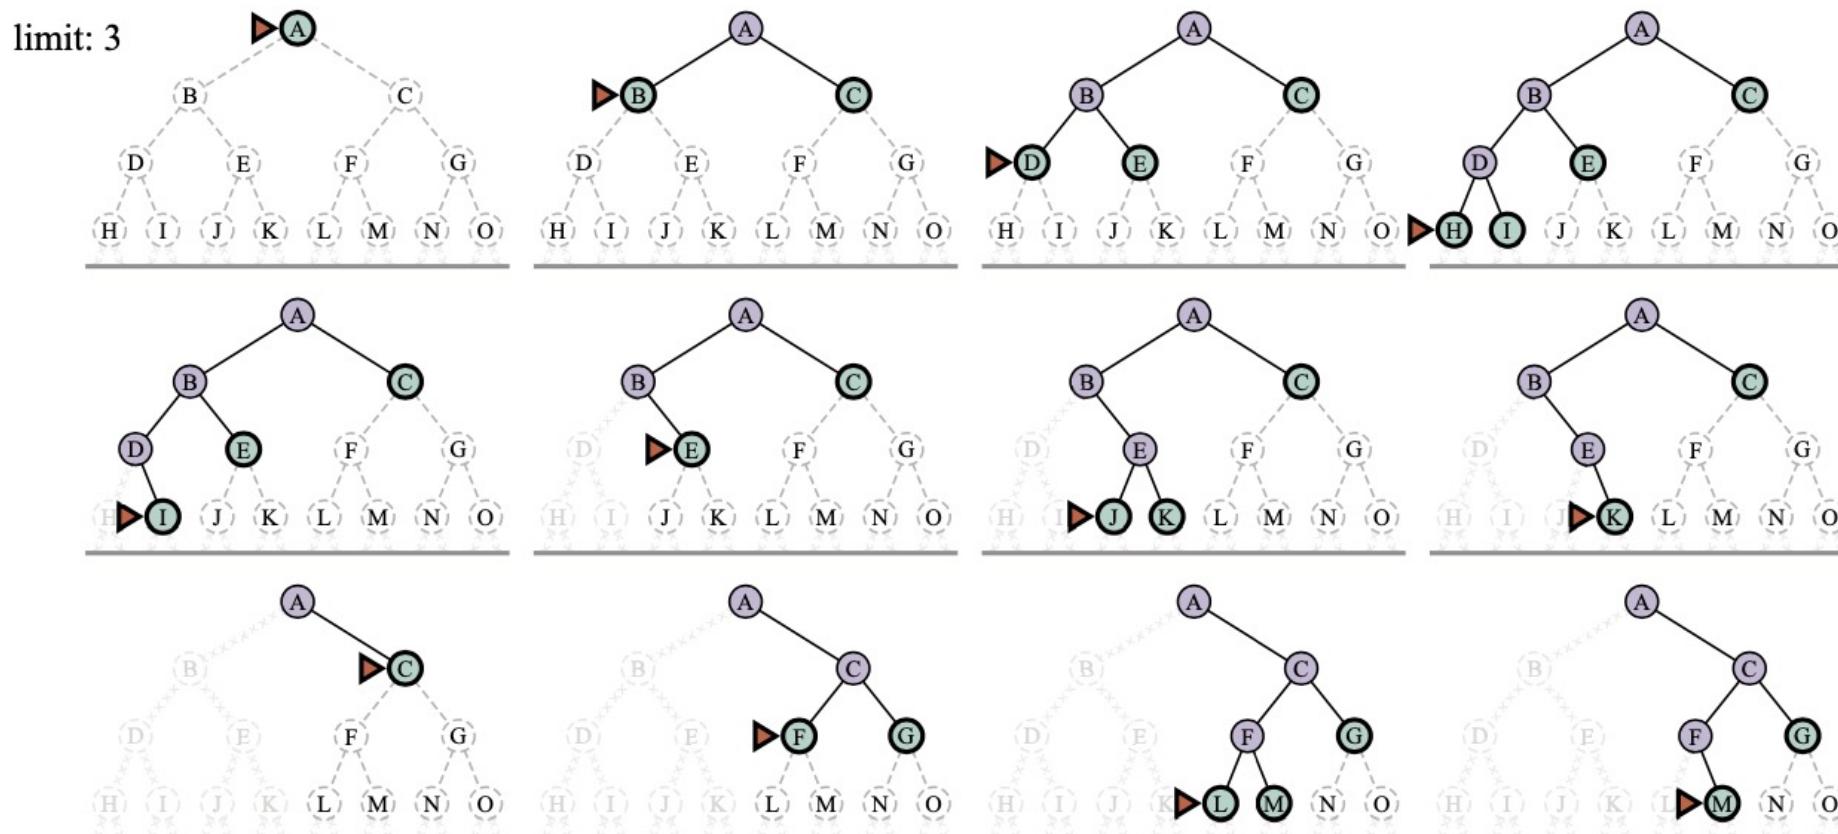
Four iterations of iterative deepening search



Four iterations of iterative deepening search



Four iterations of iterative deepening search



Bidirectional Best-First Search

keeps two frontiers and two tables of reached states

```
function BIBF-SEARCH( $problem_F, f_F, problem_B, f_B$ ) returns a solution node, or failure
   $node_F \leftarrow \text{NODE}(problem_F.\text{INITIAL})$  // Node for a start state
   $node_B \leftarrow \text{NODE}(problem_B.\text{INITIAL})$  // Node for a goal state
   $frontier_F \leftarrow$  a priority queue ordered by  $f_F$ , with  $node_F$  as an element
   $frontier_B \leftarrow$  a priority queue ordered by  $f_B$ , with  $node_B$  as an element
   $reached_F \leftarrow$  a lookup table, with one key  $node_F.\text{STATE}$  and value  $node_F$ 
   $reached_B \leftarrow$  a lookup table, with one key  $node_B.\text{STATE}$  and value  $node_B$ 
   $solution \leftarrow \text{failure}$ 
  while not TERMINATED( $solution, frontier_F, frontier_B$ ) do
    if  $f_F(\text{TOP}(frontier_F)) < f_B(\text{TOP}(frontier_B))$  then
       $solution \leftarrow \text{PROCEED}(F, problem_F, frontier_F, reached_F, reached_B, solution)$ 
    else
       $solution \leftarrow \text{PROCEED}(B, problem_B, frontier_B, reached_B, reached_F, solution)$ 
  return  $solution$ 
```

Bidirectional Best-First Search

keeps two frontiers and two tables of reached states

```
function PROCEED(dir, problem, frontier, reached, reached2, solution) returns a solution
    // Expand node on frontier; check against the other frontier in reached2.
    // The variable “dir” is the direction: either F for forward or B for backward.
    node  $\leftarrow$  POP(frontier)
    for each child in EXPAND(problem, node) do
        s  $\leftarrow$  child.STATE
        if s not in reached or PATH-COST(child) < PATH-COST(reached[s]) then
            reached[s]  $\leftarrow$  child
            add child to frontier
            if s is in reached2 then
                solution2  $\leftarrow$  JOIN-NODES(dir, child, reached2[s]))
                if PATH-COST(solution2) < PATH-COST(solution) then
                    solution  $\leftarrow$  solution2
    return solution
```

Evaluation of search algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ¹	Yes ^{1,2}	No	No	Yes ¹	Yes ^{1,4}
Optimal cost?	Yes ³	Yes	No	No	Yes ³	Yes ^{3,4}
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$

b is the branching factor; m is the maximum depth of the search tree;

d is the depth of the shallowest solution, or is m when there is no solution;

ℓ is the depth limit

Values of *hSLD* —straight-line distances to Bucharest.

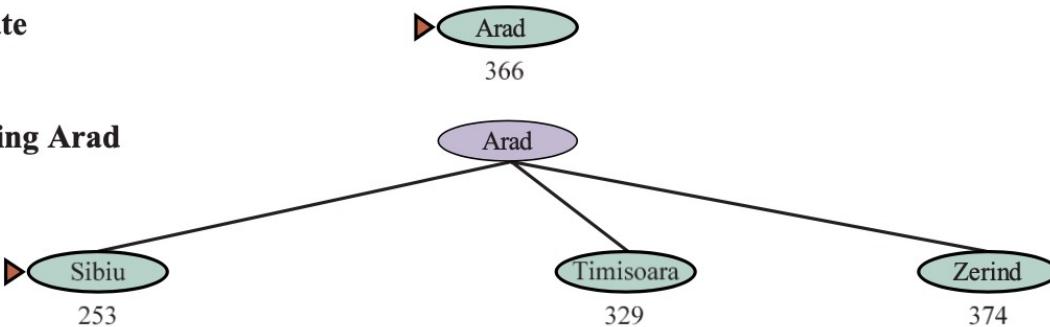
Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

A* search

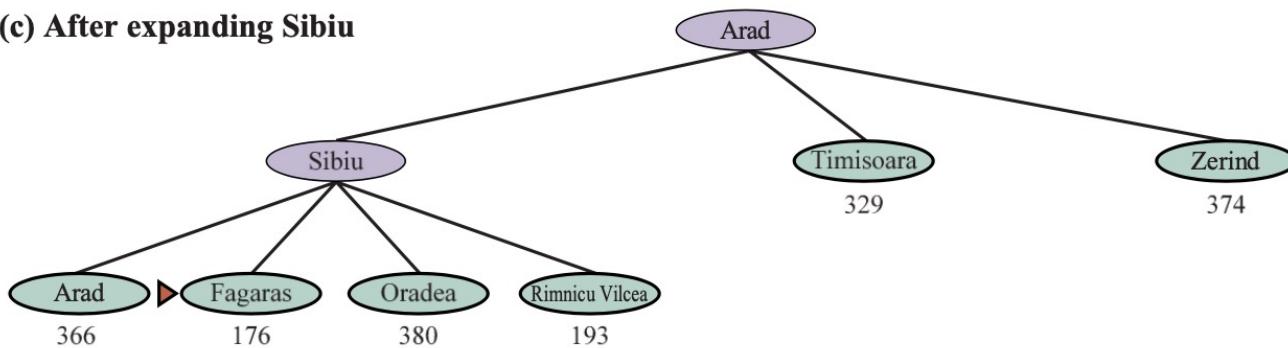
(a) The initial state



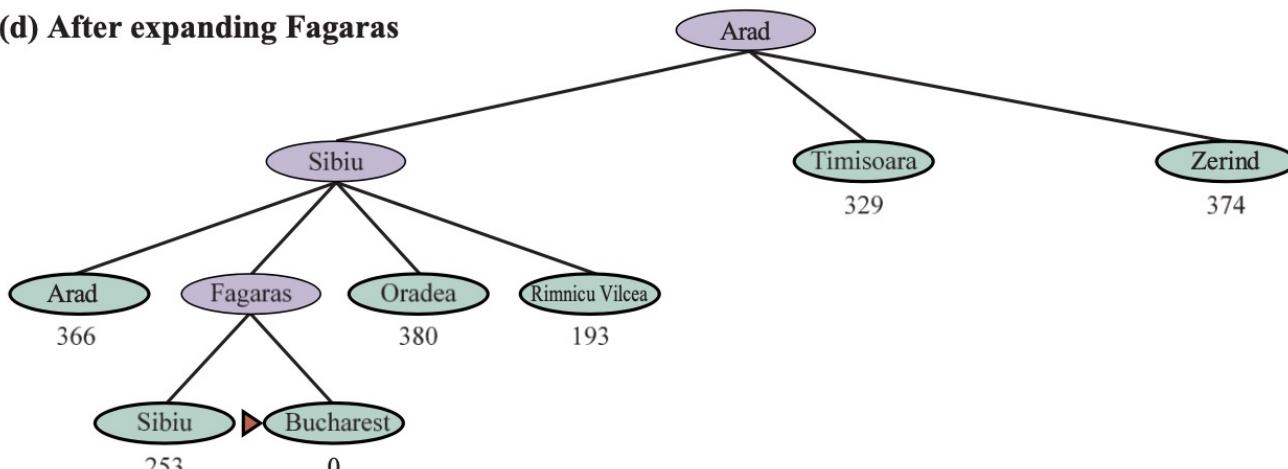
(b) After expanding Arad



(c) After expanding Sibiu



(d) After expanding Fagaras



A* search

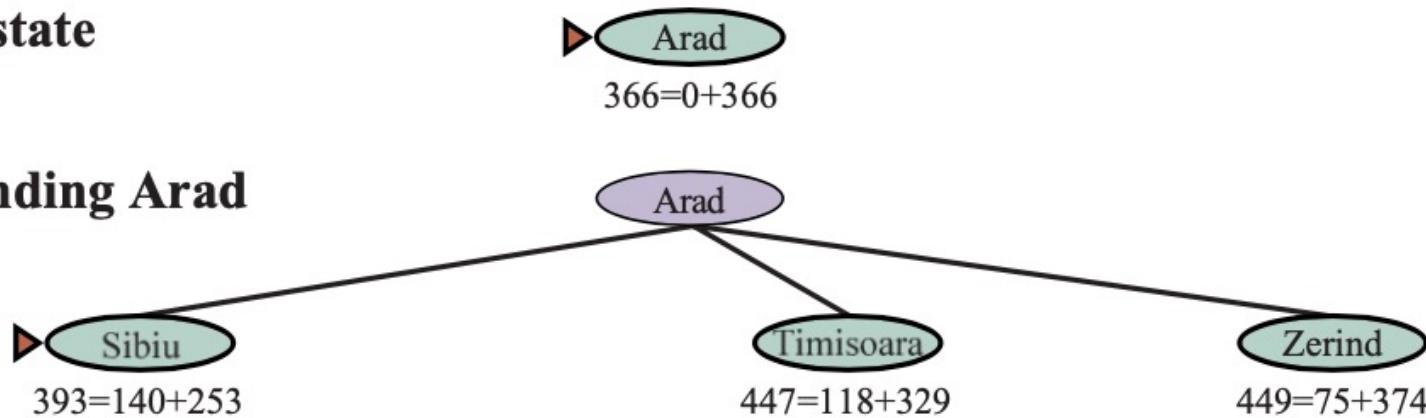
Nodes are labeled with $f = g + h$.

The h values are the Straight-Line Distances heuristic h_{SLD}

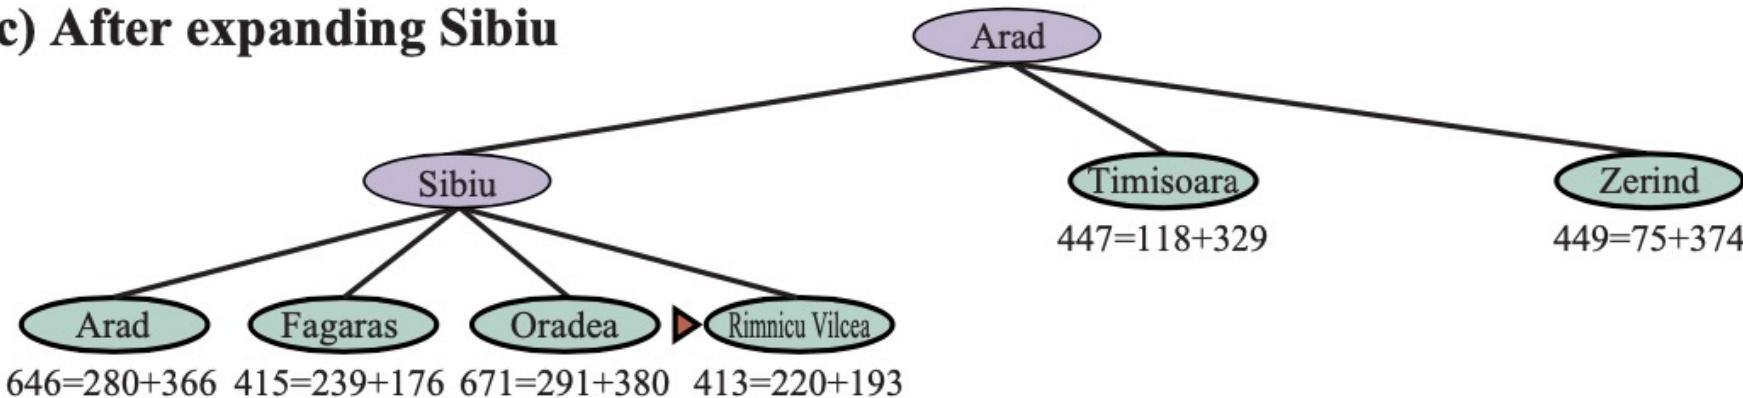
(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu

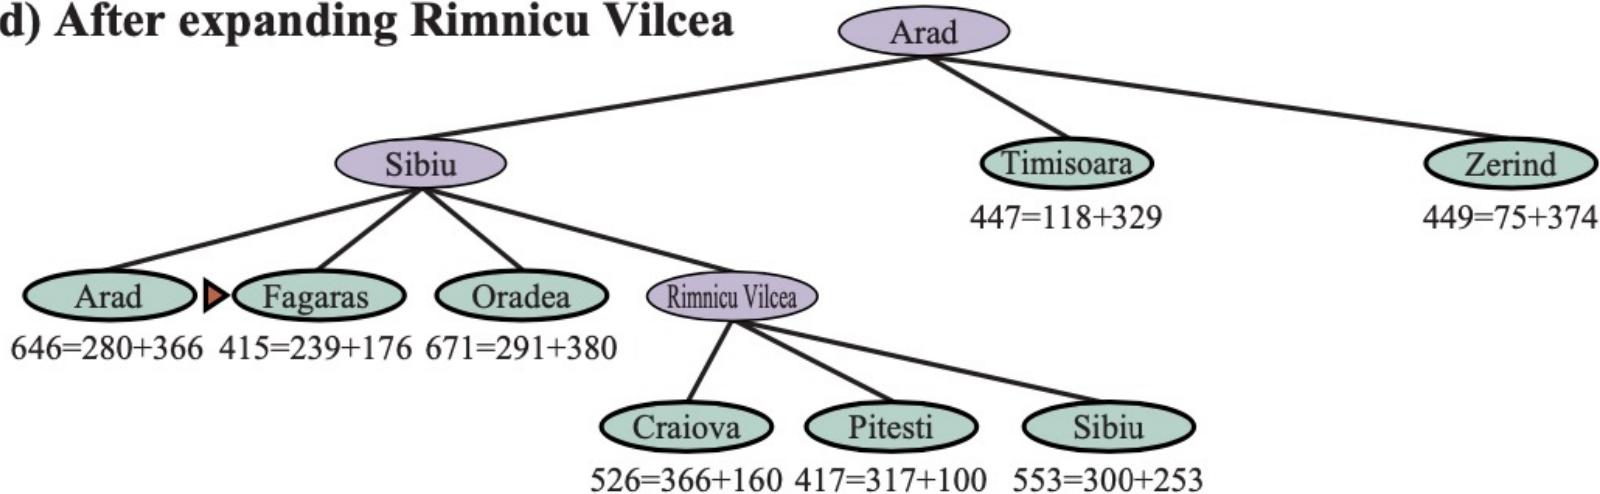


A* search

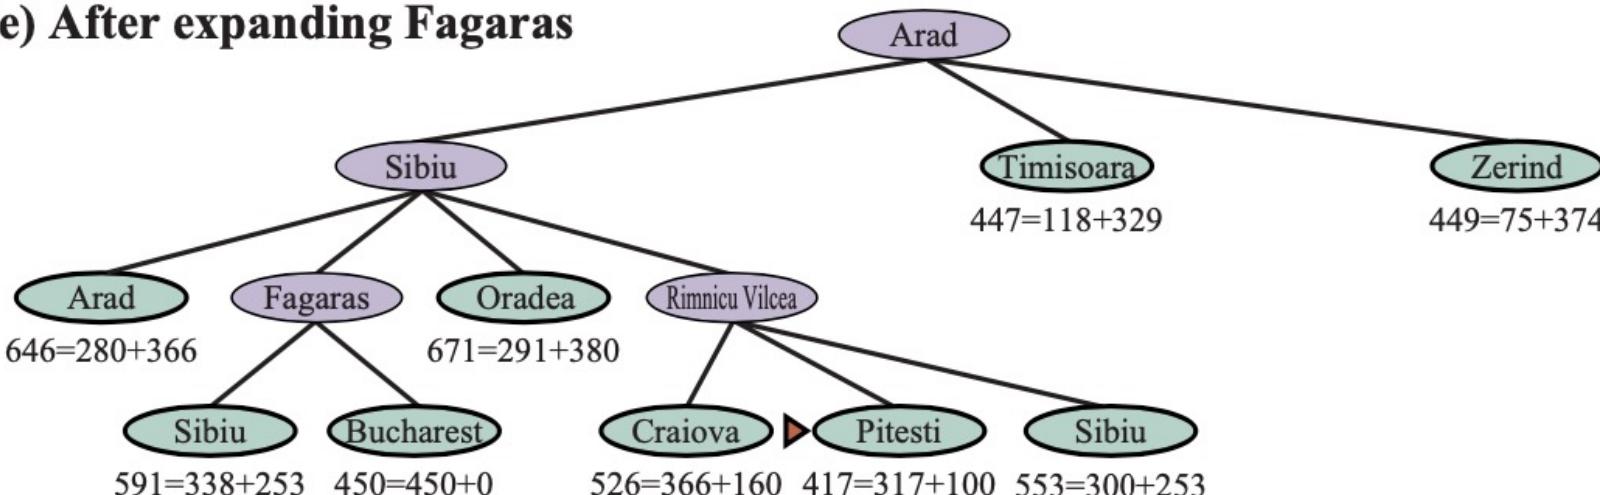
Nodes are labeled with $f = g + h$.

The h values are the Straight-Line Distances heuristic h_{SLD}

(d) After expanding Rimnicu Vilcea



(e) After expanding Fagaras

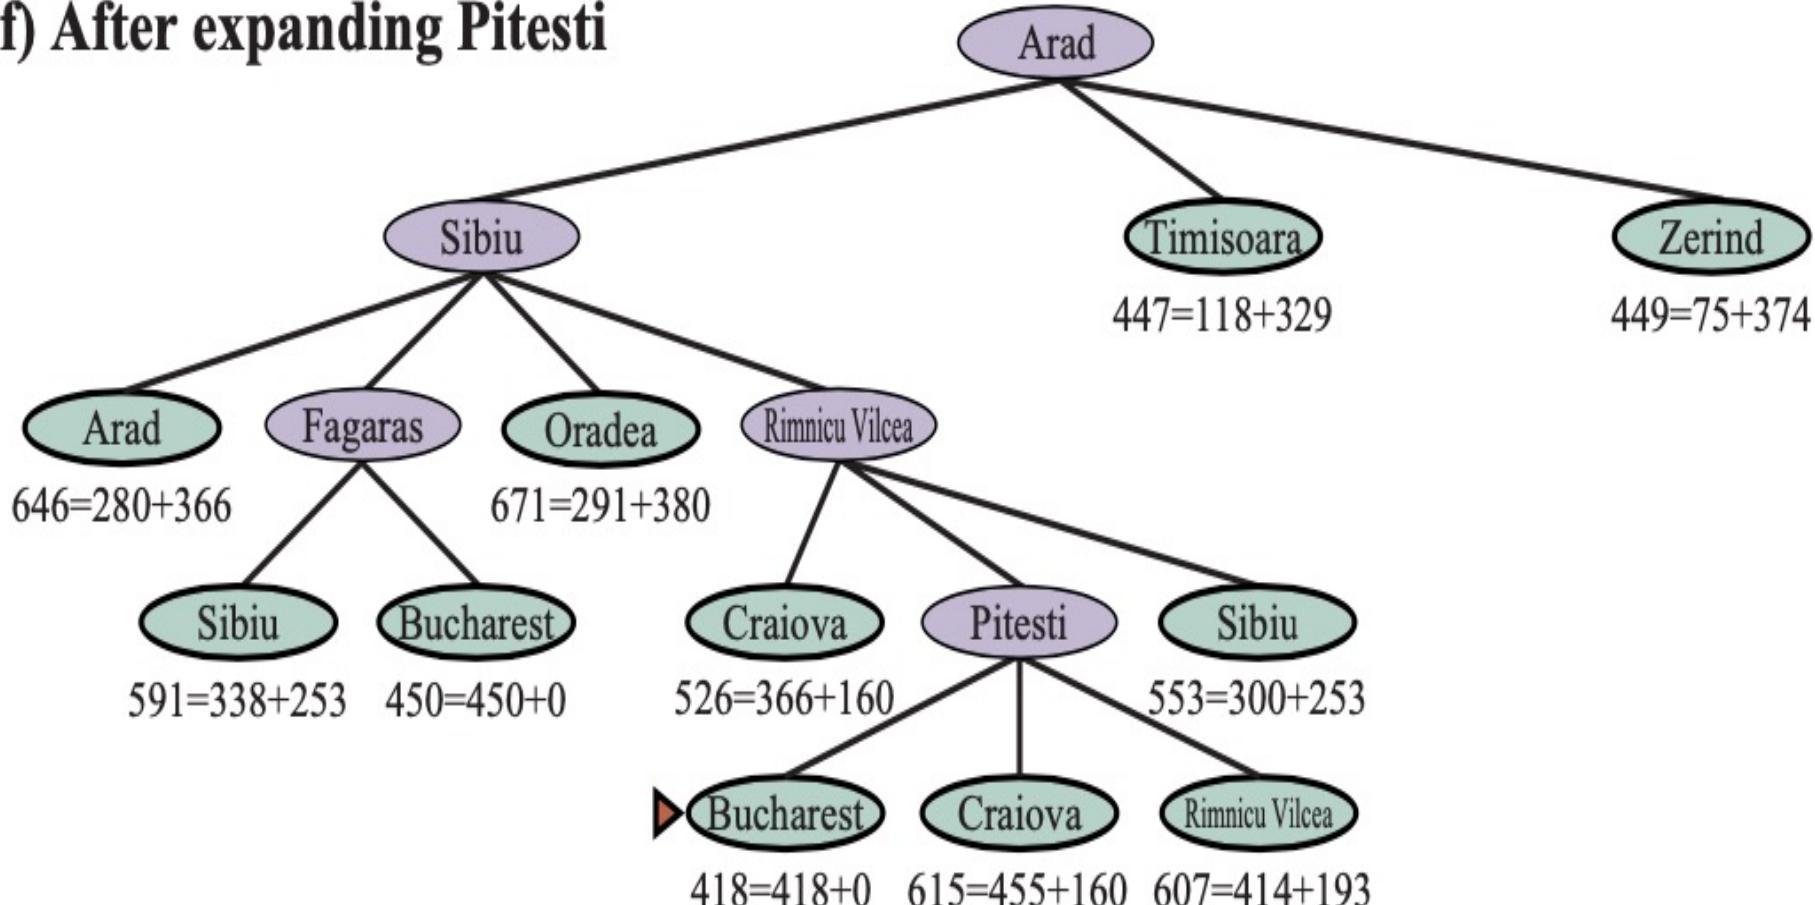


A* search

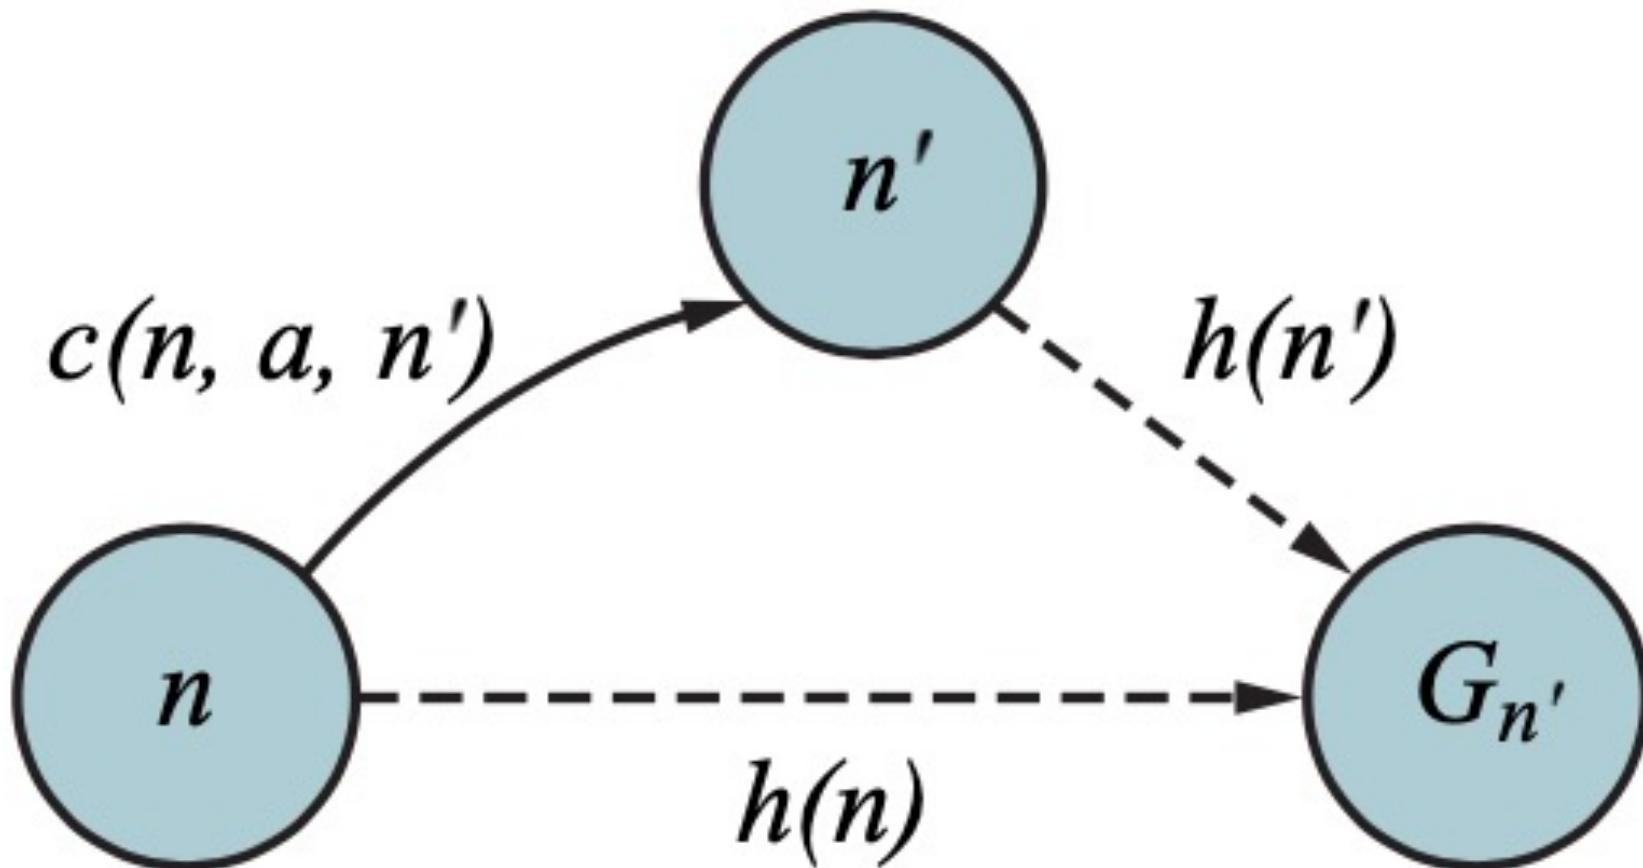
Nodes are labeled with $f = g + h$.

The h values are the Straight-Line Distances heuristic h_{SLD}

(f) After expanding Pitesti

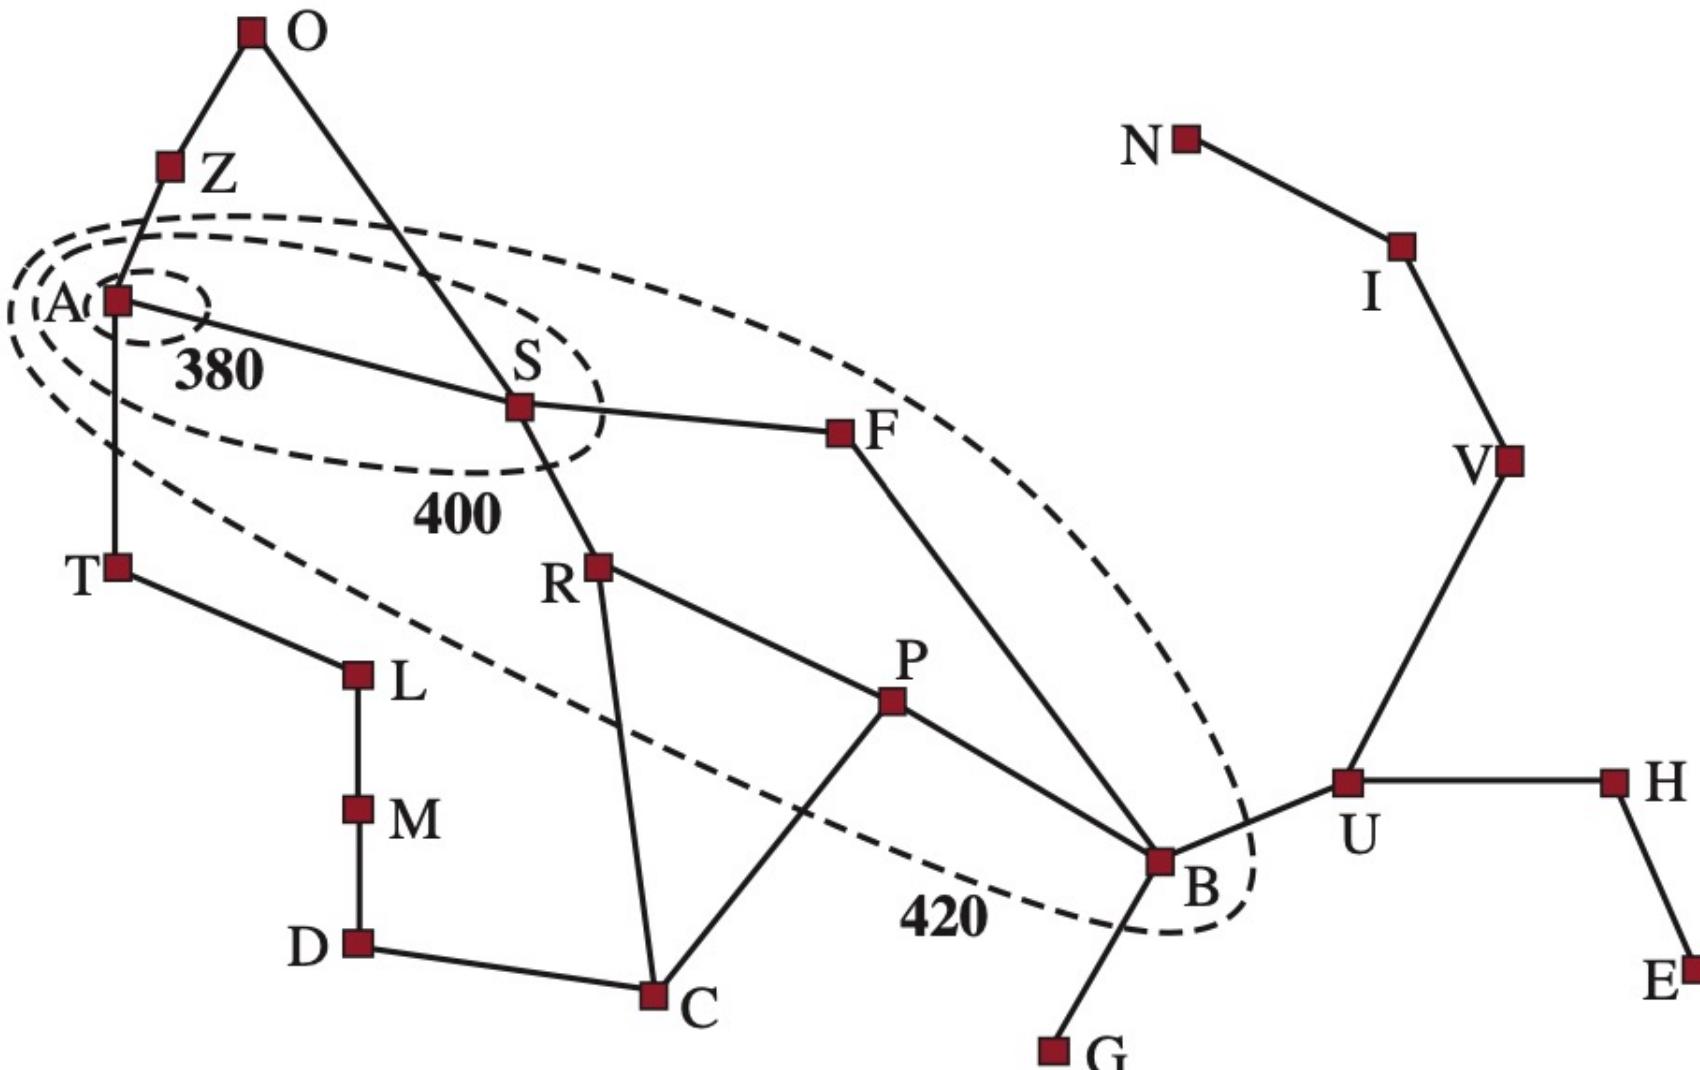


Triangle Inequality



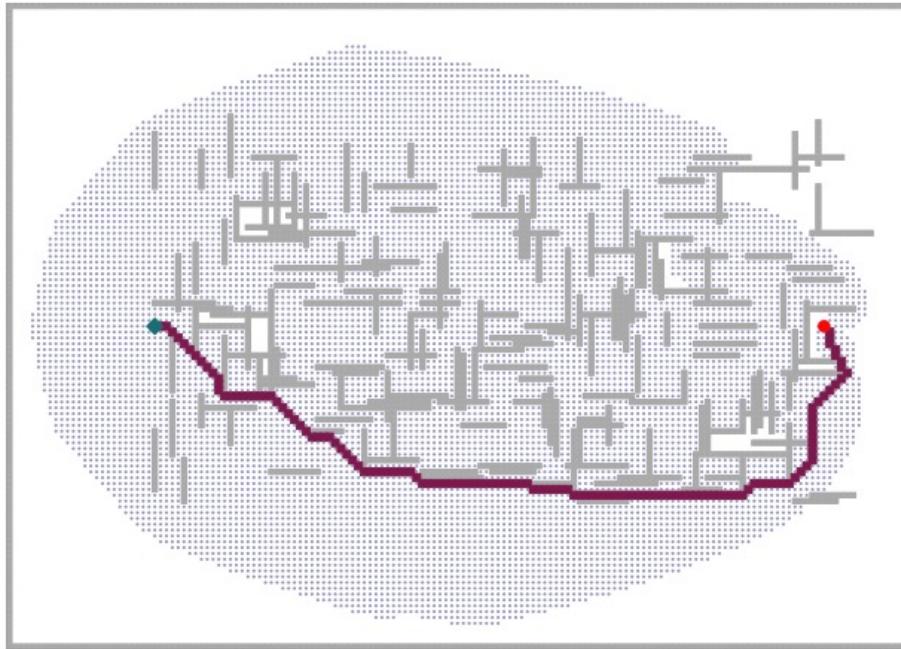
If the heuristic h is consistent, then the single number $h(n)$ will be less than the sum of the cost $c(n, a, a')$ of the action from n to n' plus the heuristic estimate $h(n')$.

Map of Romania showing contours at $f = 380$, $f = 400$, and $f = 420$, with Arad as the start state

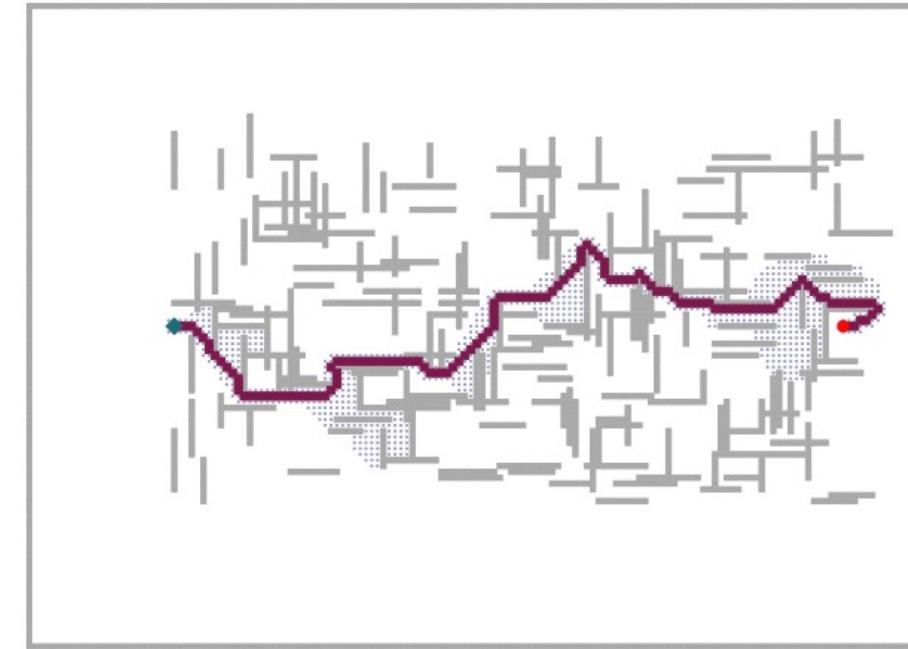


(a) A* Search

(b) Weighted A* Search



(a)



(b)

The gray bars are obstacles, the purple line is the path from the green start to red goal, and the small dots are states that were reached by each search.

On this particular problem, weighted A* explores 7 times fewer states and finds a path that is 5% more costly.

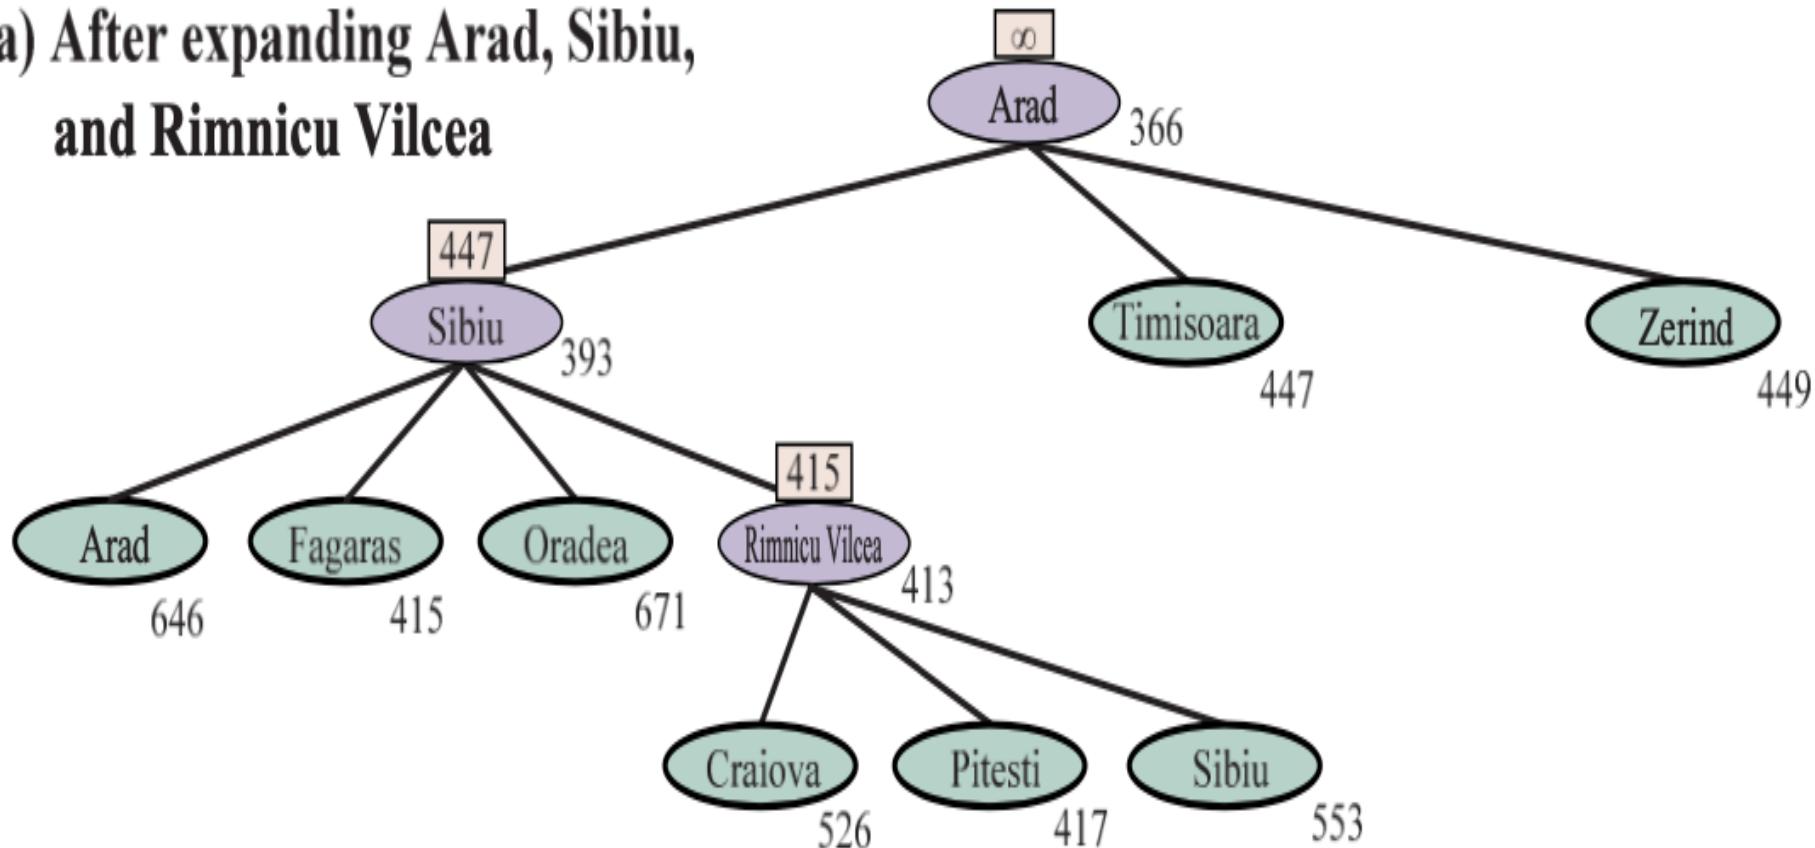
Recursive Best-First Search (RBFS) Algorithm

```
function RECURSIVE-BEST-FIRST-SEARCH(problem) returns a solution or failure
    solution, fvalue  $\leftarrow$  RBFS(problem, NODE(problem.INITIAL),  $\infty$ )
    return solution

function RBFS(problem, node, f_limit) returns a solution or failure, and a new f-cost limit
    if problem.IS-GOAL(node.STATE) then return node
    successors  $\leftarrow$  LIST(EXPAND(node))
    if successors is empty then return failure,  $\infty$ 
    for each s in successors do      // update f with value from previous search
        s.f  $\leftarrow$  max(s.PATH-COST + h(s), node.f))
    while true do
        best  $\leftarrow$  the node in successors with lowest f-value
        if best.f > f_limit then return failure, best.f
        alternative  $\leftarrow$  the second-lowest f-value among successors
        result, best.f  $\leftarrow$  RBFS(problem, best, min(f_limit, alternative))
        if result  $\neq$  failure then return result, best.f
```

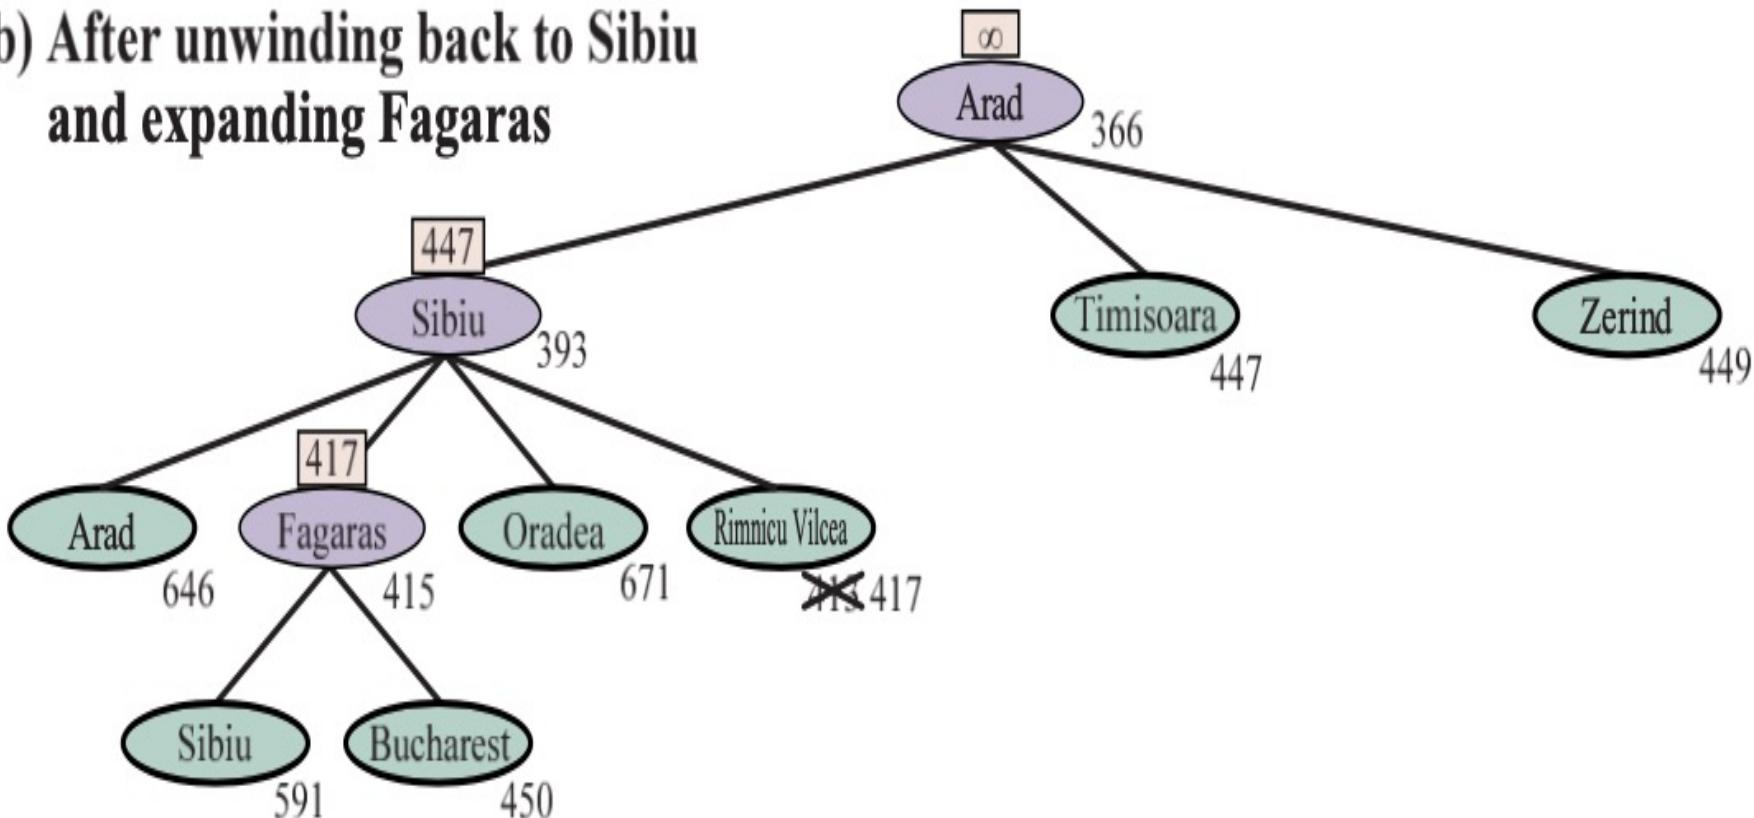
Recursive Best-First Search (RBFS)

(a) After expanding Arad, Sibiu, and Rimnicu Vilcea



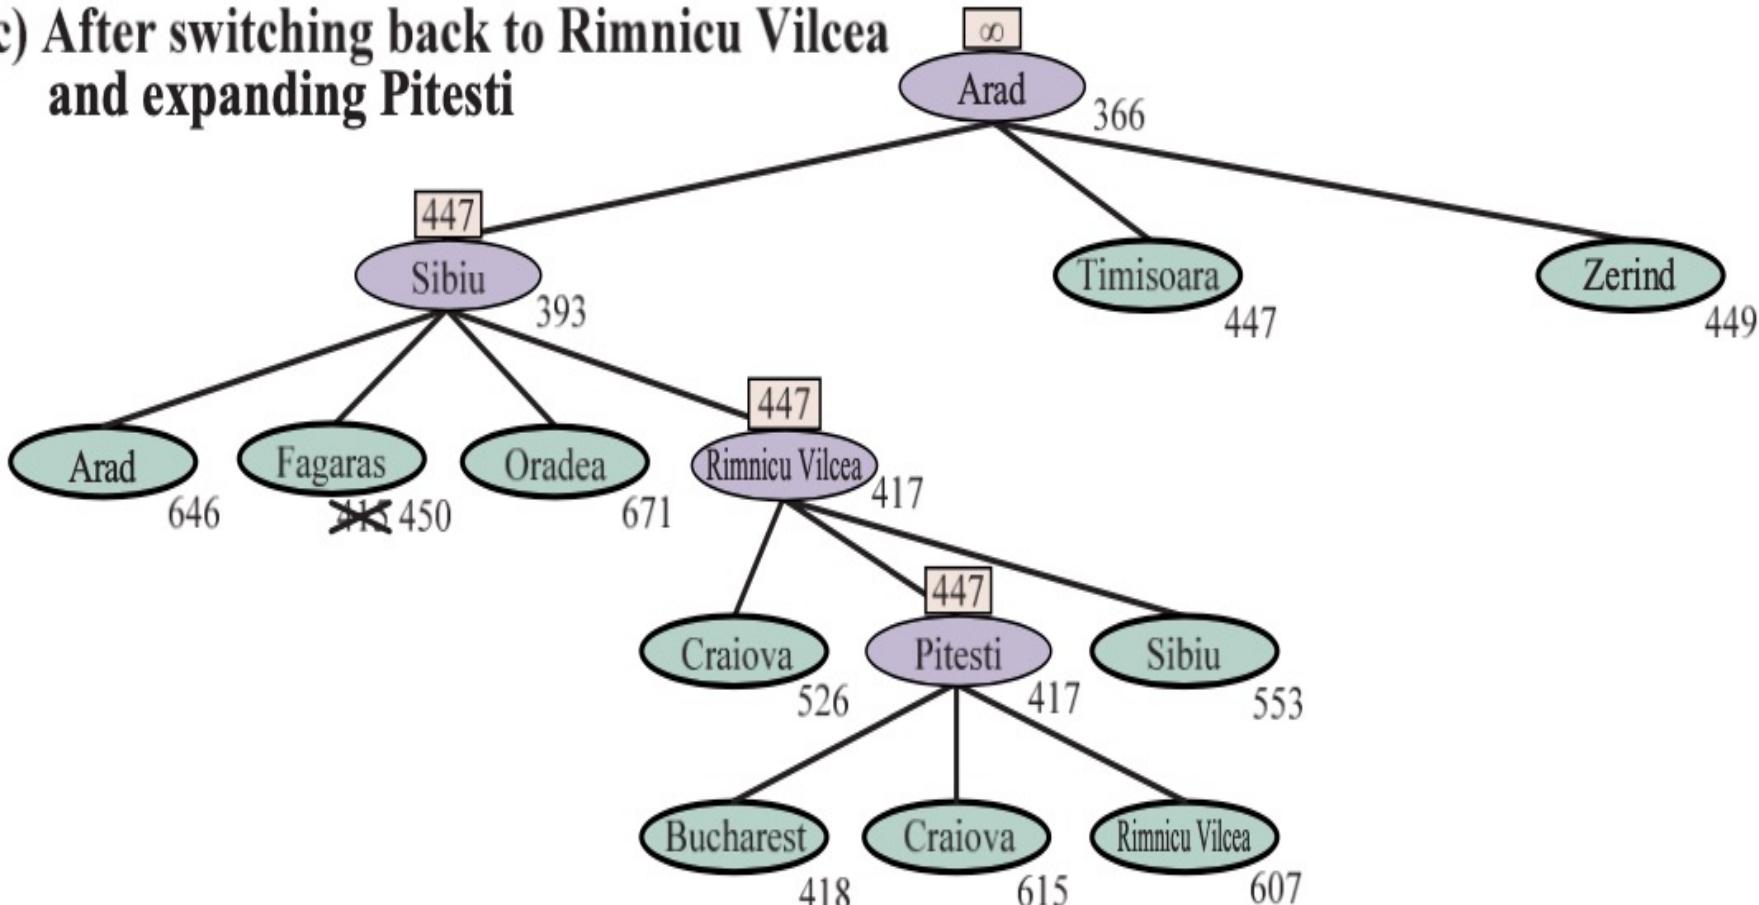
Recursive Best-First Search (RBFS)

(b) After unwinding back to Sibiu
and expanding Fagaras

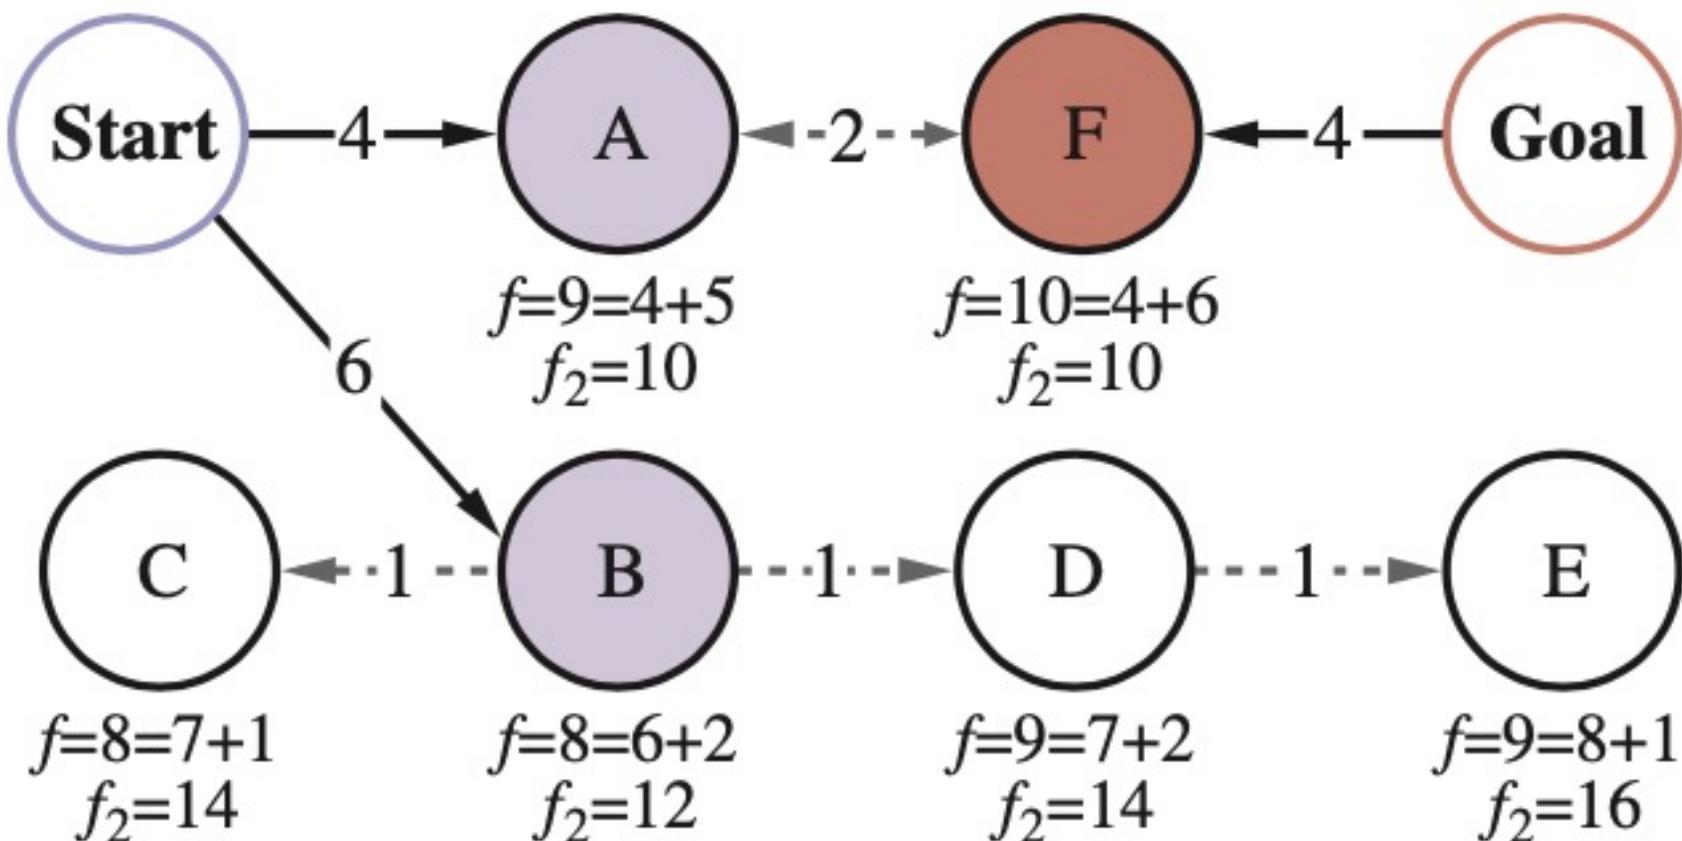


Recursive Best-First Search (RBFS)

(c) After switching back to Rimnicu Vilcea
and expanding Pitesti



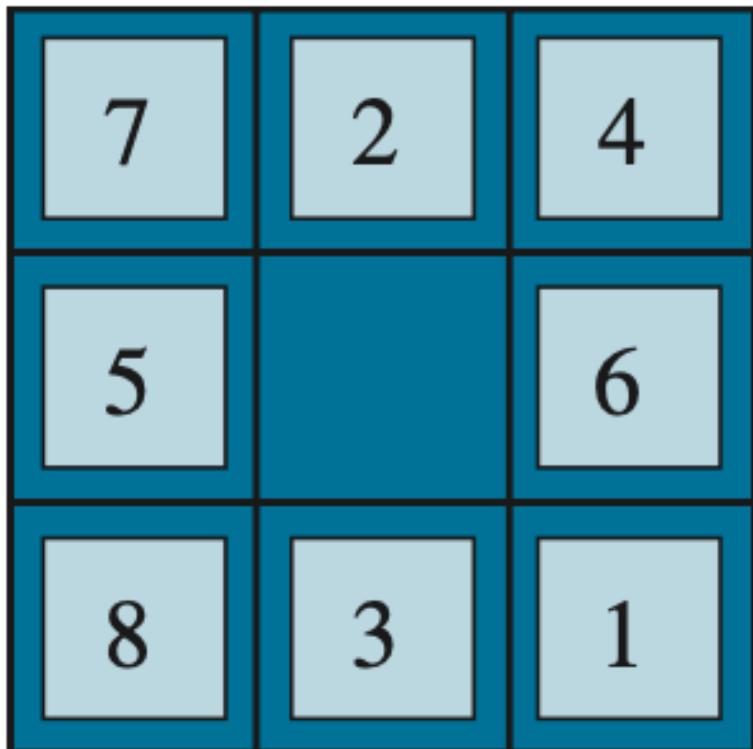
Bidirectional Search maintains two frontiers



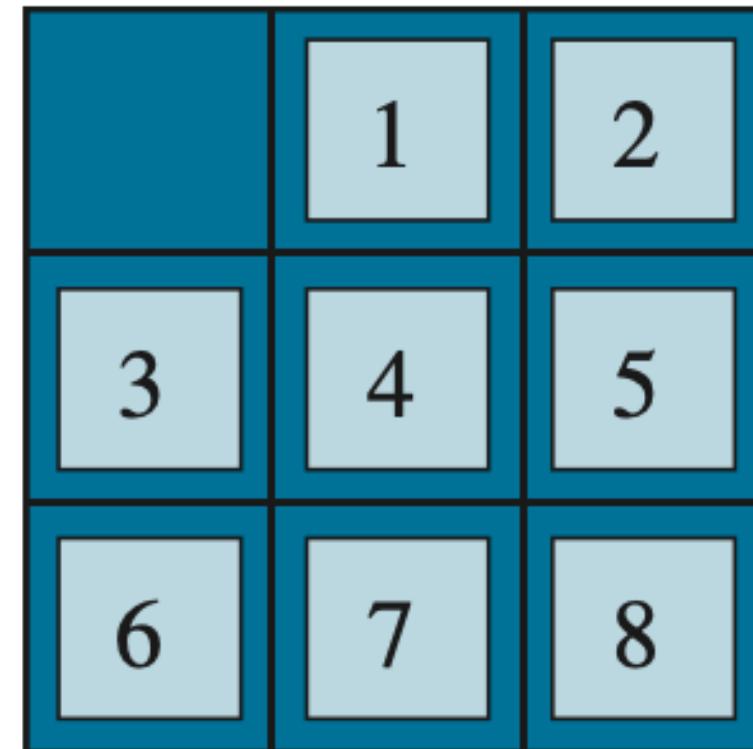
On the left, nodes A and B are successors of Start;
on the right, node F is an inverse successor of Goal

A typical instance of the 8-puzzle

The shortest solution is 26 actions long



Start State

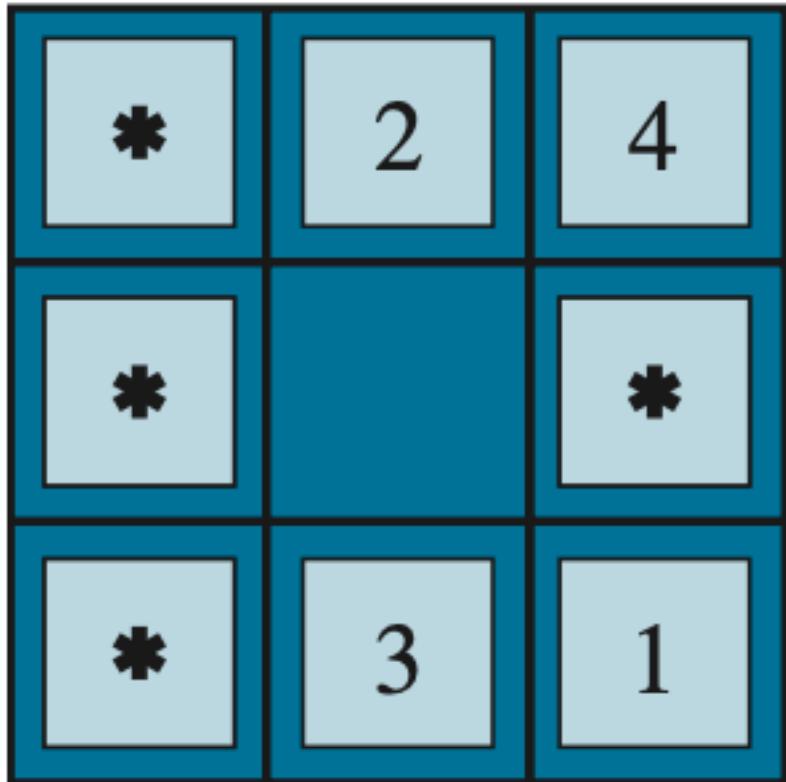


Goal State

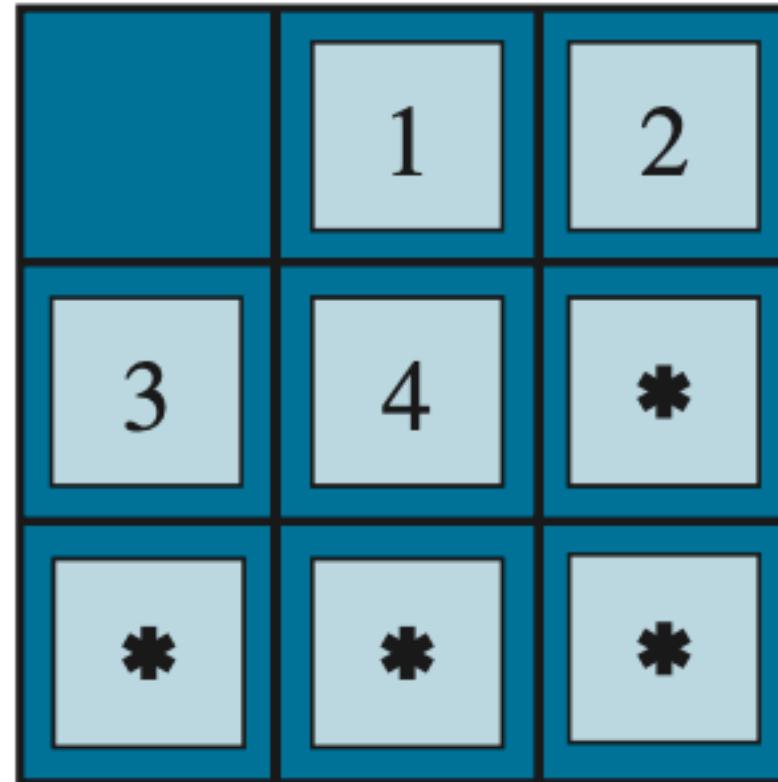
Comparison of the search costs and effective branching factors for 8-puzzle problems

d	Search Cost (nodes generated)			Effective Branching Factor		
	BFS	$A^*(h_1)$	$A^*(h_2)$	BFS	$A^*(h_1)$	$A^*(h_2)$
6	128	24	19	2.01	1.42	1.34
8	368	48	31	1.91	1.40	1.30
10	1033	116	48	1.85	1.43	1.27
12	2672	279	84	1.80	1.45	1.28
14	6783	678	174	1.77	1.47	1.31
16	17270	1683	364	1.74	1.48	1.32
18	41558	4102	751	1.72	1.49	1.34
20	91493	9905	1318	1.69	1.50	1.34
22	175921	22955	2548	1.66	1.50	1.34
24	290082	53039	5733	1.62	1.50	1.36
26	395355	110372	10080	1.58	1.50	1.35
28	463234	202565	22055	1.53	1.49	1.36

A subproblem of the 8-puzzle



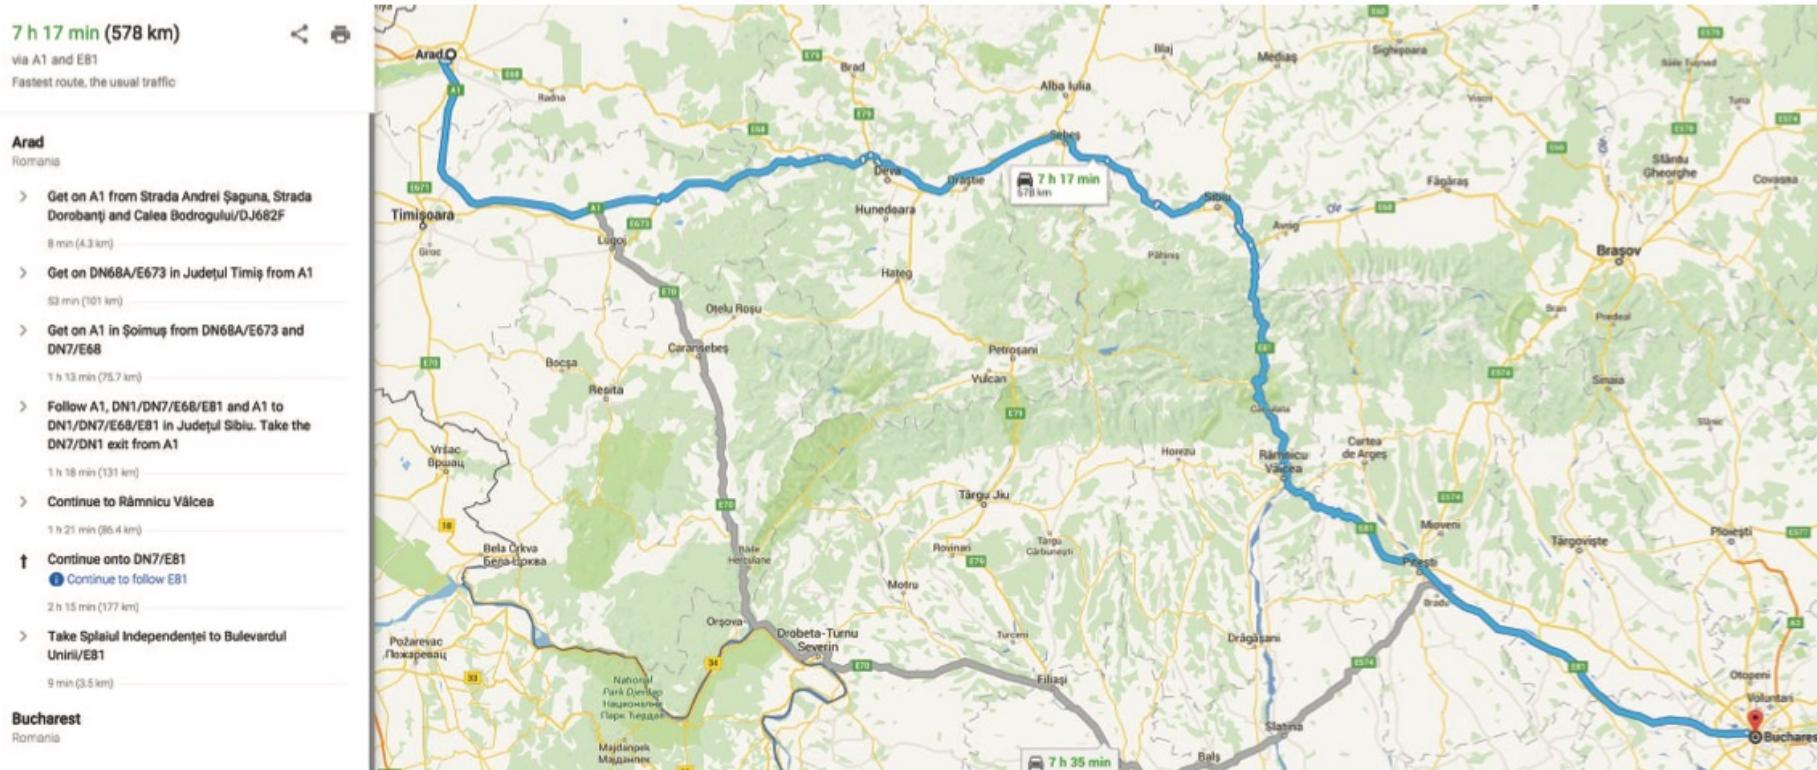
Start State



Goal State

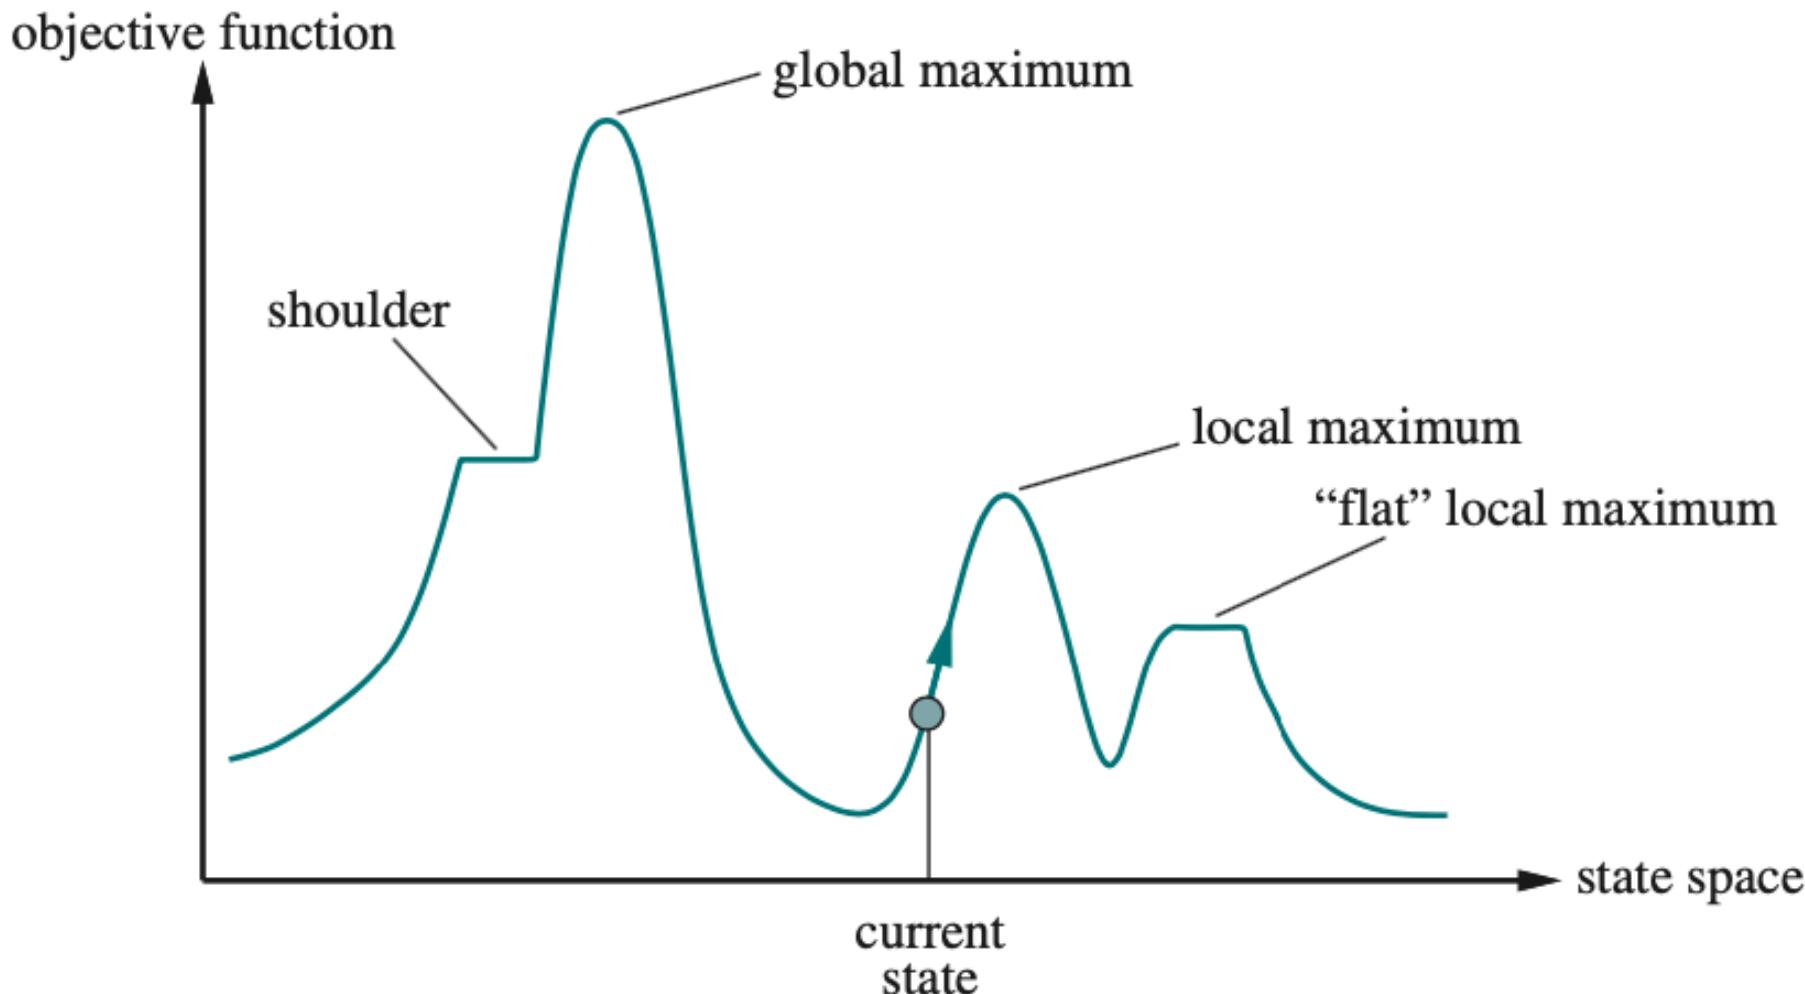
The task is to get tiles 1, 2, 3, 4, and the blank into their correct positions, without worrying about what happens to the other tiles

A Web service providing driving directions, computed by a search algorithm.



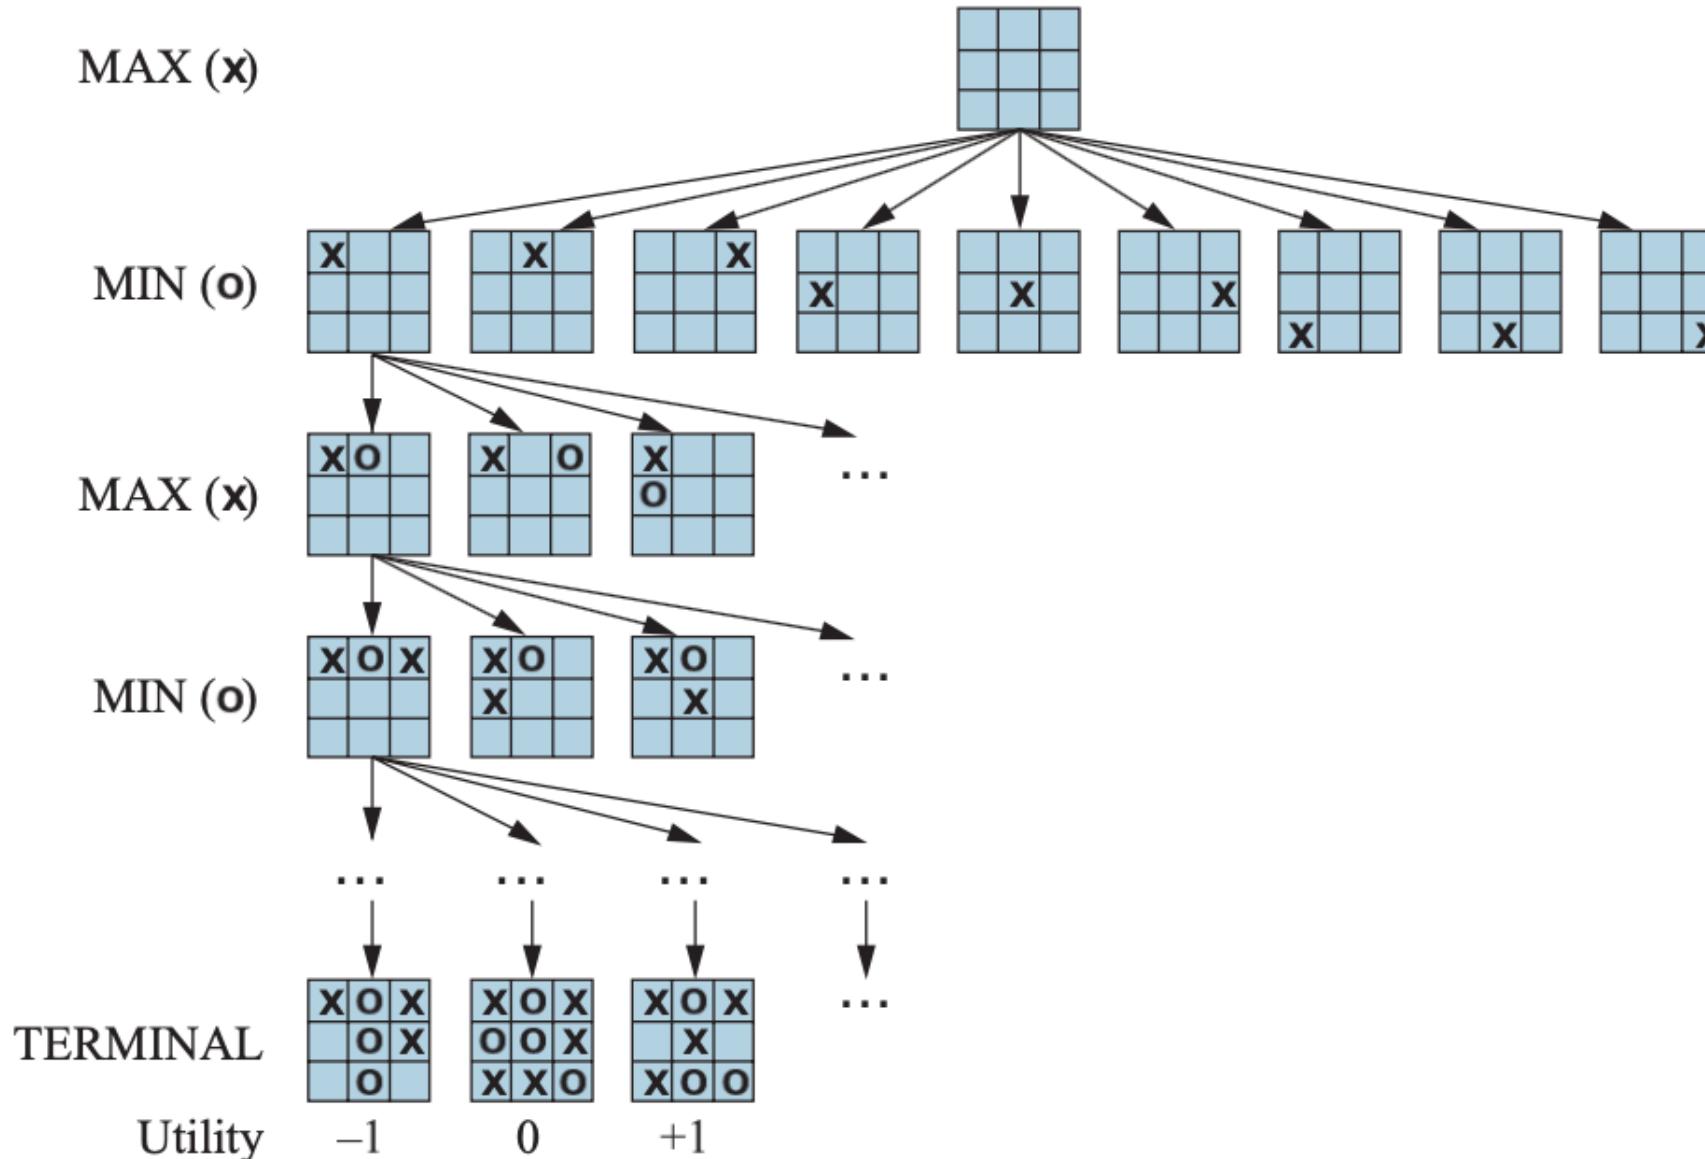
Search in Complex Environments

A one-dimensional state-space landscape



Adversarial Search and Games

Game Tree for the Game of Tic-tac-toe

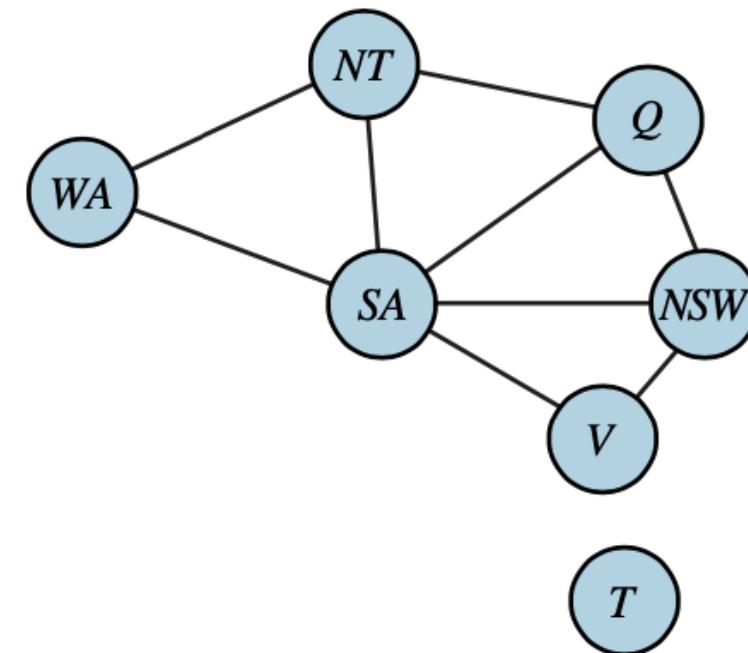


Constraint Satisfaction Problems

The Map-Coloring Problem Represented as a Constraint Graph

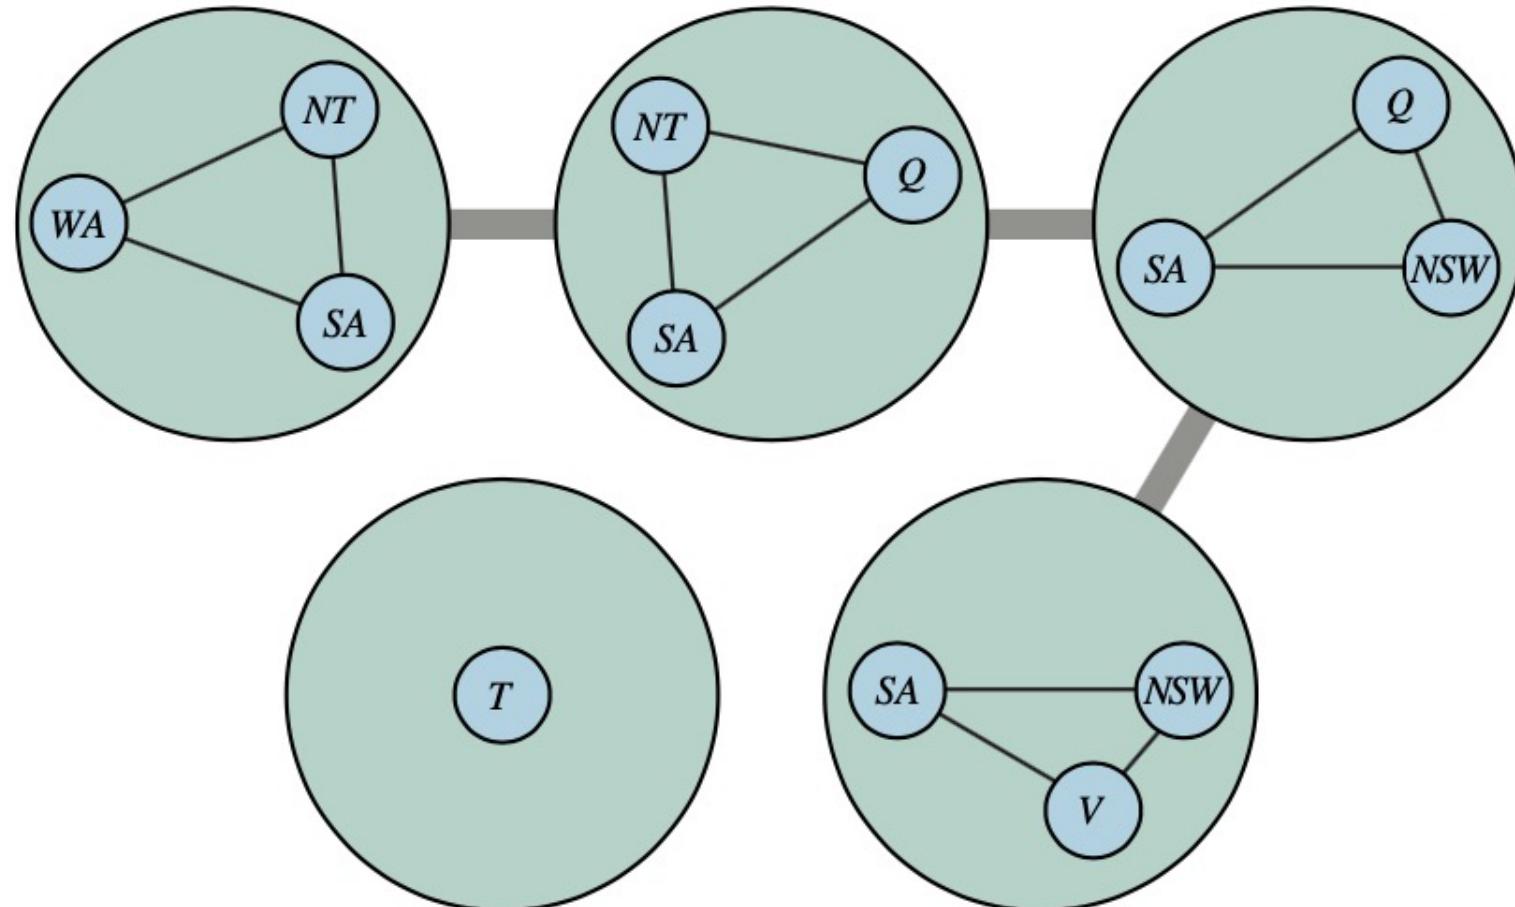


(a)



(b)

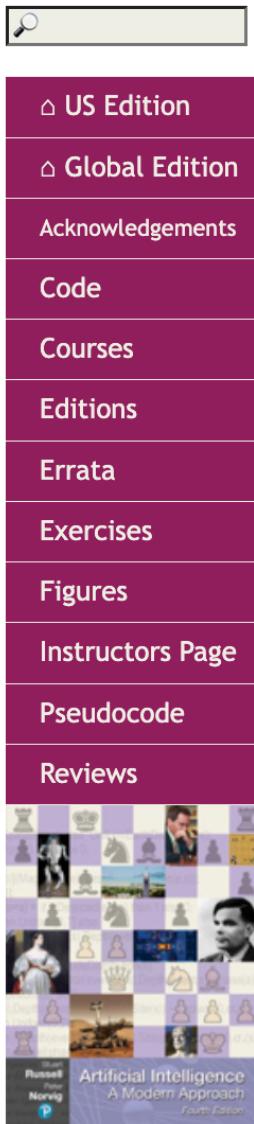
A Tree Decomposition of the Constraint Graph



Artificial Intelligence: A Modern Approach (AIMA)

- Artificial Intelligence: A Modern Approach (AIMA)
 - <http://aima.cs.berkeley.edu/>
- AIMA Python
 - <http://aima.cs.berkeley.edu/python/readme.html>
 - <https://github.com/aimacode/aima-python>
- Search
 - <http://aima.cs.berkeley.edu/python/search.html>
- Games: Adversarial Search
<http://aima.cs.berkeley.edu/python/games.html>
- CSP (Constraint Satisfaction Problems)
 - <http://aima.cs.berkeley.edu/python/csp.html>

Artificial Intelligence: A Modern Approach (AIMA)



Artificial Intelligence: A Modern Approach, 4th US ed.

by Stuart Russell and Peter Norvig

The authoritative, most-used AI textbook, adopted by over 1500 schools.

Table of Contents for the US Edition (or see the Global Edition)

Preface (pdf); Contents with subsections

I Artificial Intelligence

- 1 Introduction ... 1
- 2 Intelligent Agents ... 36

II Problem-solving

- 3 Solving Problems by Searching ... 63
- 4 Search in Complex Environments ... 110
- 5 Adversarial Search and Games ... 146
- 6 Constraint Satisfaction Problems ... 180

III Knowledge, reasoning, and planning

- 7 Logical Agents ... 208
- 8 First-Order Logic ... 251
- 9 Inference in First-Order Logic ... 280
- 10 Knowledge Representation ... 314
- 11 Automated Planning ... 344

IV Uncertain knowledge and reasoning

- 12 Quantifying Uncertainty ... 385
- 13 Probabilistic Reasoning ... 412
- 14 Probabilistic Reasoning over Time ... 461
- 15 Probabilistic Programming ... 500
- 16 Making Simple Decisions ... 528
- 17 Making Complex Decisions ... 562
- 18 Multiagent Decision Making ... 599

V Machine Learning

- 19 Learning from Examples ... 651
- 20 Learning Probabilistic Models ... 721
- 21 Deep Learning ... 750
- 22 Reinforcement Learning ... 789

VI Communicating, perceiving, and acting

- 23 Natural Language Processing ... 823
- 24 Deep Learning for Natural Language Processing ... 856
- 25 Computer Vision ... 881
- 26 Robotics ... 925

VII Conclusions

- 27 Philosophy, Ethics, and Safety of AI ... 981
- 28 The Future of AI ... 1012
- Appendix A: Mathematical Background ... 1023
- Appendix B: Notes on Languages and Algorithms ... 1030
- Bibliography ... 1033 ([pdf](#) and [LaTeX .bib file](#) and [bib data](#))
- Index ... 1069 ([pdf](#))

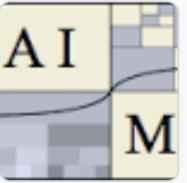
Exercises (website)

Figures (pdf)

Code (website); Pseudocode (pdf)

Covers: US, Global

AIMA Code



aimacode
Code for the book "Artificial Intelligence: A Modern Approach"
358 followers | Berkeley, CA | http://aima.cs.berkeley.edu | peter@norvig.com

[Overview](#) [Repositories 13](#) [Projects](#) [Packages](#) [People](#)

Popular repositories

aima-python Public
Python implementation of algorithms from Russell And Norvig's "Artificial Intelligence - A Modern Approach"
Jupyter Notebook 6.6k stars 3.2k forks

aima-java Public
Java implementation of algorithms from Russell And Norvig's "Artificial Intelligence - A Modern Approach"
Java 1.4k stars 767 forks

aima-pseudocode Public
Pseudocode descriptions of the algorithms from Russell And Norvig's "Artificial Intelligence - A Modern Approach"
740 stars 386 forks

aima-exercises Public
Exercises for the book Artificial Intelligence: A Modern Approach
HTML 611 stars 353 forks

aima-javascript Public
Javascript visualization of algorithms from Russell And Norvig's "Artificial Intelligence - A Modern Approach"
JavaScript 495 stars 208 forks

aima-lisp Public
Common Lisp implementation of algorithms from Russell And Norvig's "Artificial Intelligence - A Modern Approach"
Common Lisp 342 stars 95 forks

[Follow](#)

You can now follow organizations
Organization activity like new discussions, sponsorships, and repositories will appear in [your dashboard feed](#).

[OK, got it!](#)

members. You must be a member to see who's a part of this organization.

Top languages

JavaScript Python Java
Common Lisp Scala

[Report abuse](#)

<https://github.com/aimacode>

AIMA Python

Code Issues 120 Pull requests 79 Actions Projects Wiki Security Insights

master 1 branch 0 tags Go to file Add file Code

mcventur	Fixed bug in treatment of repeated nodes in frontier...	61d695b on Dec 5, 2021	1,190 commits
aima-data @ f6cbea6	updating submodule (#994)	4 years ago	
gui	fixed tests (#1191)	2 years ago	
images	add perception and tests (#1091)	3 years ago	
js	Added TicTacToe to notebook (#213)	7 years ago	
notebooks	Image Rendering problem resolved (#1178)	3 years ago	
tests	fixed tests (#1191)	2 years ago	
.coveragerc	Added coverage report generation to Travis (#1058)	3 years ago	
.flake8	Fix flake8 warnings (#508)	5 years ago	
.gitignore	Reworked PriorityQueue and Added Tests (#1025)	4 years ago	
.gitmodules	Updating Submodule (#647)	5 years ago	
.travis.yml	fixed svm for not posdef kernel matrix, updated .travis.yml wi...	2 years ago	

About

Python implementation of algorithms from Russell And Norvig's "Artificial Intelligence - A Modern Approach"

Readme MIT license 6.6k stars 337 watching 3.2k forks

Releases

No releases published

Packages

No packages published

Papers with Code

State-of-the-Art (SOTA)



Search for papers, code and tasks



Browse State-of-the-Art

Follow

Discuss

Trends

About

Log In/Register

Browse State-of-the-Art

1509 leaderboards • 1327 tasks • 1347 datasets • 17810 papers with code

Follow on Twitter for updates

Computer Vision



Semantic
Segmentation

33 leaderboards

667 papers with code



Image
Classification

52 leaderboards

564 papers with code



Object
Detection

54 leaderboards

467 papers with code



Image
Generation

51 leaderboards

231 papers with code



Pose
Estimation

40 leaderboards

231 papers with code

▶ See all 707 tasks

Natural Language Processing



Machine
Translation



Language
Modelling



Question
Answering



Sentiment
Analysis



Text
Generation

Python in Google Colab (Python101)

<https://colab.research.google.com/drive/1FEG6DnGvwfUbeo4zJ1zTunjMqf2RkCrT>

The screenshot shows a Google Colab notebook titled "python101.ipynb". The interface includes a toolbar with file operations, a sidebar with a "CODE" tab selected, and a main workspace with four code cells and their corresponding outputs.

Cell 1:

```
# Future Value
pv = 100
r = 0.1
n = 7
fv = pv * ((1 + (r)) ** n)
print(round(fv, 2))
```

Output: 194.87

Cell 11:

```
amount = 100
interest = 10 #10% = 0.01 * 10
years = 7
future_value = amount * ((1 + (0.01 * interest)) ** years)
print(round(future_value, 2))
```

Output: 194.87

Cell 12:

```
# Python Function def
def getfv(pv, r, n):
    fv = pv * ((1 + (r)) ** n)
    return fv
fv = getfv(100, 0.1, 7.)
print(round(fv, 2))
```

Output: 194.87

Cell 13:

```
# Python if else
score = 80
if score >=60 :
    print("Pass")
else:
    print("Fail")
```

Output: Pass

<https://tinyurl.com/aintpuppython101>

Summary

- Artificial Intelligence
- Intelligent Agents
- Problem Solving

References

- Stuart Russell and Peter Norvig (2020), Artificial Intelligence: A Modern Approach, 4th Edition, Pearson.
- Thomas R. Caldwell (2025), The Agentic AI Bible: The Complete and Up-to-Date Guide to Design, Build, and Scale Goal-Driven, LLM-Powered Agents that Think, Execute and Evolve, Independently published
- Numa Dhamani and Maggie Engler (2024), Introduction to Generative AI, Manning
- Denis Rothman (2024), Transformers for Natural Language Processing and Computer Vision - Third Edition: Explore Generative AI and Large Language Models with Hugging Face, ChatGPT, GPT-4V, and DALL-E 3, 3rd ed. Edition, Packt Publishing
- Ben Auffarth (2023), Generative AI with LangChain: Build large language model (LLM) apps with Python, ChatGPT and other LLMs, Packt Publishing.
- Aurélien Géron (2022), Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems, 3rd Edition, O'Reilly Media.
- Steven D'Ascoli (2022), Artificial Intelligence and Deep Learning with Python: Every Line of Code Explained For Readers New to AI and New to Python, Independently published.
- Nithin Buduma, Nikhil Buduma, Joe Papa (2022), Fundamentals of Deep Learning: Designing Next-Generation Machine Intelligence Algorithms, 2nd Edition, O'Reilly Media.
- Khaled Bayoudh, Raja Knani, Fayçal Hamdaoui, and Abdellatif Mtibaa (2022). "A survey on deep multimodal learning for computer vision: advances, trends, applications, and datasets." *The Visual Computer* 38, no. 8: 2939-2970.
- Abdelrahman Mohamed, Hung-yi Lee, Lasse Borgholt, Jakob D. Havtorn, Joakim Edin, Christian Igel, Katrin Kirchhoff et al. (2022). "Self-Supervised Speech Representation Learning: A Review." *arXiv preprint arXiv:2205.10643*.
- Longbing Cao (2022). "Decentralized ai: Edge intelligence and smart blockchain, metaverse, web3, and desc." *IEEE Intelligent Systems* 37, no. 3: 6-19.
- Thien Huynh-The, Quoc-Viet Pham, Xuan-Qui Pham, Thanh Thi Nguyen, Zhu Han, and Dong-Seong Kim (2022). "Artificial Intelligence for the Metaverse: A Survey." *arXiv preprint arXiv:2202.10336*.
- Thippa Reddy Gadekallu, Thien Huynh-The, Weizheng Wang, Gokul Yenduri, Pasika Ranaweera, Quoc-Viet Pham, Daniel Benevides da Costa, and Madhusanka Liyanage (2022). "Blockchain for the Metaverse: A Review." *arXiv preprint arXiv:2203.09738*.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry et al. (2021) "Learning transferable visual models from natural language supervision." In International Conference on Machine Learning, pp. 8748-8763. PMLR.
- Aske Plaat, Annie Wong, Suzan Verberne, Joost Broekens, Niki van Stein, and Thomas Back. (2024) "Reasoning with Large Language Models, a Survey." *arXiv preprint arXiv:2407.11511*.
- Madaan, Aman, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon et al. (2024) "Self-refine: Iterative refinement with self-feedback." *Advances in Neural Information Processing Systems* 36.