

Artificial Intelligence



Machine Learning: Supervised and Unsupervised Learning; The Theory of Learning and Ensemble Learning

1141AI04

MBA, IM, NTPU (M5276) (Fall 2025)

Tue 2, 3, 4 (9:10-12:00) (B3F17)



Min-Yuh Day, Ph.D,
Professor and Director

Institute of Information Management, National Taipei University

<https://web.ntpu.edu.tw/~myday>

2025-10-07



<https://meet.google.com/paj-zhhj-mya>



Syllabus

Week Date Subject/Topics

1 2025/09/09 Introduction to Artificial Intelligence

2 2025/09/16 Artificial Intelligence and Intelligent Agents;
Problem Solving

3 2025/09/23 Knowledge, Reasoning and Knowledge Representation;
Uncertain Knowledge and Reasoning

4 2025/09/30 Case Study on Artificial Intelligence I

5 2025/10/07 Machine Learning: Supervised and Unsupervised Learning;
The Theory of Learning and Ensemble Learning

Syllabus

Week Date Subject/Topics

6 2025/10/14 NVIDIA Fundamentals of Deep Learning I:
Deep Learning; Neural Networks

7 2025/10/21 NVIDIA Fundamentals of Deep Learning II:
Convolutional Neural Networks;
Data Augmentation and Deployment

8 2025/10/28 Self-Learning

9 2025/11/04 Midterm Project Report

10 2025/11/11 NVIDIA Fundamentals of Deep Learning III:
Pre-trained Models; Natural Language Processing

Syllabus

Week Date Subject/Topics

11 2025/11/18 Case Study on Artificial Intelligence II

12 2025/11/25 Computer Vision and Robotics

13 2025/12/02 Generative AI, Agentic AI, and Physical AI

14 2025/12/09 Philosophy and Ethics of AI and the Future of AI

15 2025/12/16 Final Project Report I

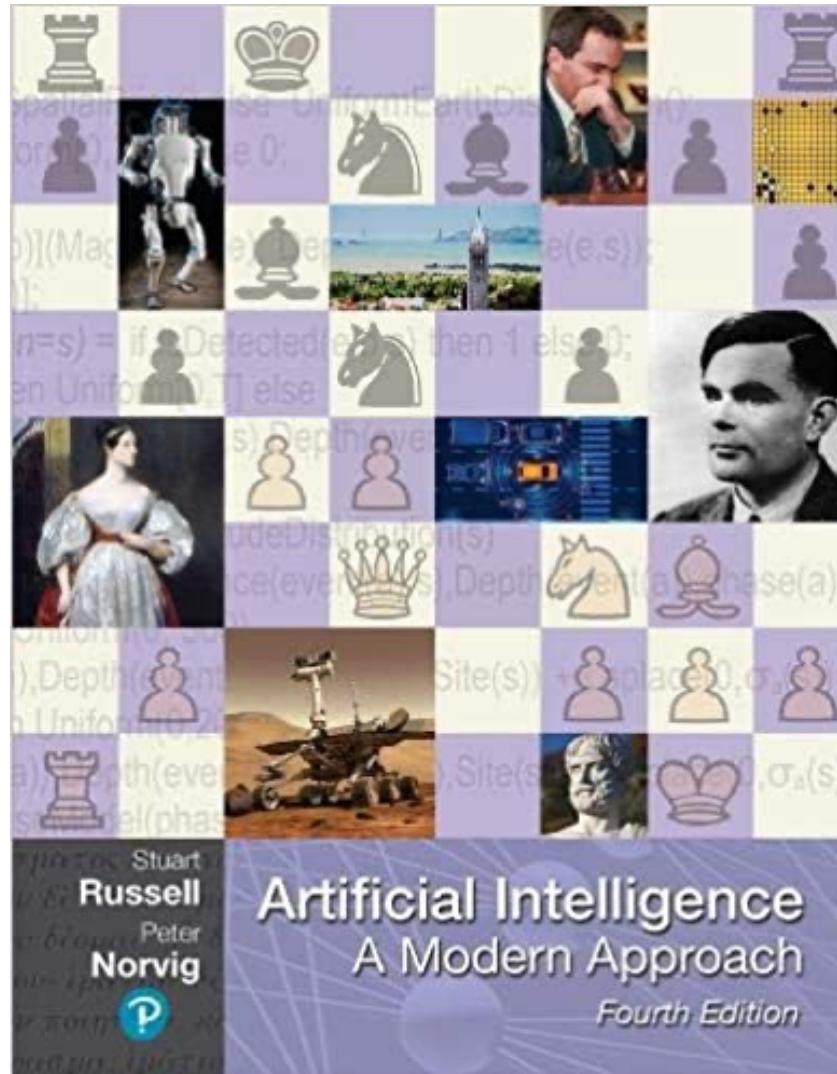
16 2025/12/23 Final Project Report II

Machine Learning: Supervised and Unsupervised Learning; The Theory of Learning and Ensemble Learning

Outline

- Machine Learning
 - Supervised Learning
 - Unsupervised Learning
- The Theory of Learning
 - Computational Learning Theory
 - Probably Approximately Correct (PAC) Learning
- Ensemble Learning
 - Bagging: Random Forests (RF)
 - Boosting: Gradient Boosting, XGBoost, LightGBM, CatBoost
 - Stacking
 - Online learning
- Meta Learning: Learning to Learn

Stuart Russell and Peter Norvig (2020),
Artificial Intelligence: A Modern Approach,
4th Edition, Pearson



Source: Stuart Russell and Peter Norvig (2020), Artificial Intelligence: A Modern Approach, 4th Edition, Pearson

<https://www.amazon.com/Artificial-Intelligence-A-Modern-Approach/dp/0134610997/>

Artificial Intelligence: A Modern Approach

- 1. Artificial Intelligence**
- 2. Problem Solving**
- 3. Knowledge and Reasoning**
- 4. Uncertain Knowledge and Reasoning**
- 5. Machine Learning**
- 6. Communicating, Perceiving, and Acting**
- 7. Philosophy and Ethics of AI**

Artificial Intelligence: Machine Learning

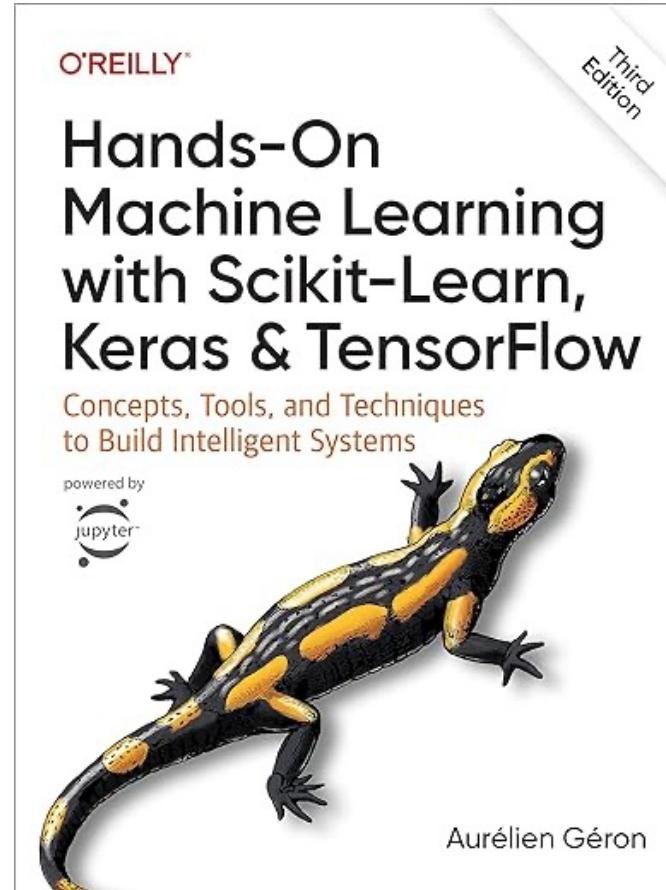
Artificial Intelligence:

5. Machine Learning

- **Learning from Examples**
- **Learning Probabilistic Models**
- **Deep Learning**
- **Reinforcement Learning**

Aurélien Géron (2022),

Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems, 3rd Edition, O'Reilly Media



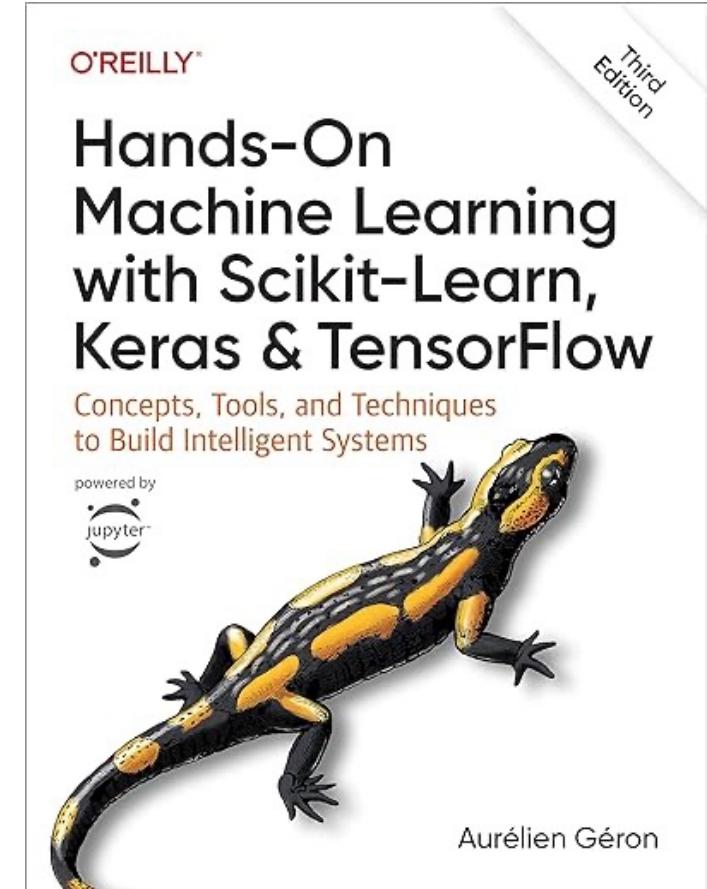
<https://github.com/ageron/handson-ml3>

Source: <https://www.amazon.com/Hands-Machine-Learning-Skikit-Learn-TensorFlow/dp/1098125975>

Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow

Notebooks

1. [The Machine Learning landscape](#)
2. [End-to-end Machine Learning project](#)
3. [Classification](#)
4. [Training Models](#)
5. [Support Vector Machines](#)
6. [Decision Trees](#)
7. [Ensemble Learning and Random Forests](#)
8. [Dimensionality Reduction](#)
9. [Unsupervised Learning Techniques](#)
10. [Artificial Neural Nets with Keras](#)
11. [Training Deep Neural Networks](#)
12. [Custom Models and Training with TensorFlow](#)
13. [Loading and Preprocessing Data](#)
14. [Deep Computer Vision Using Convolutional Neural Networks](#)
15. [Processing Sequences Using RNNs and CNNs](#)
16. [Natural Language Processing with RNNs and Attention](#)
17. [Autoencoders, GANs, and Diffusion Models](#)
18. [Reinforcement Learning](#)
19. [Training and Deploying TensorFlow Models at Scale](#)



Reinforcement Learning (DL)

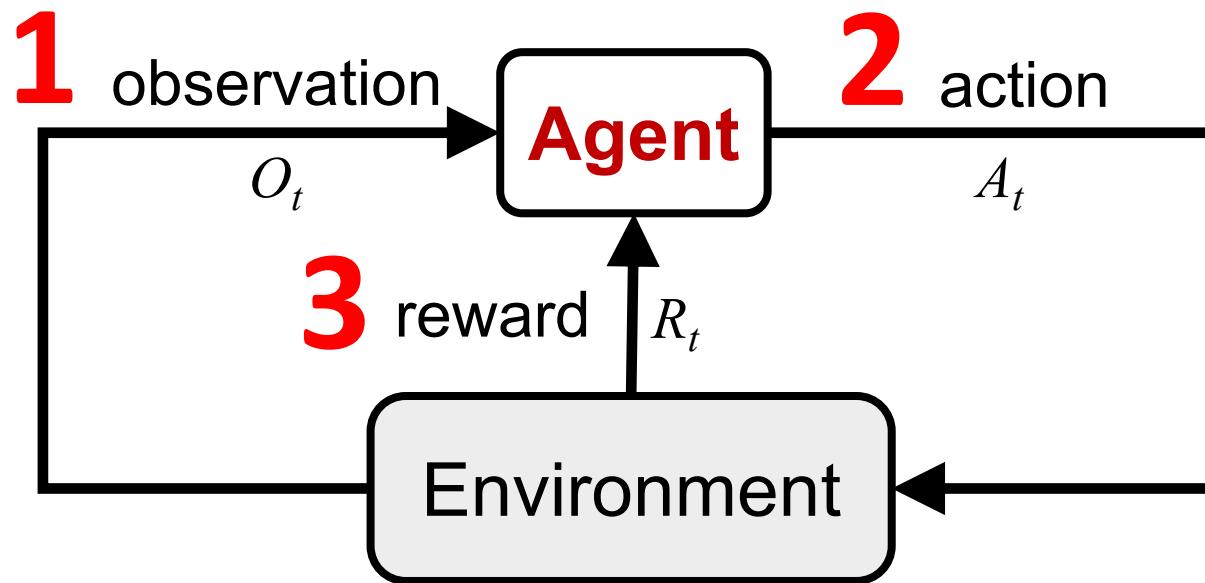
Agent

Environment

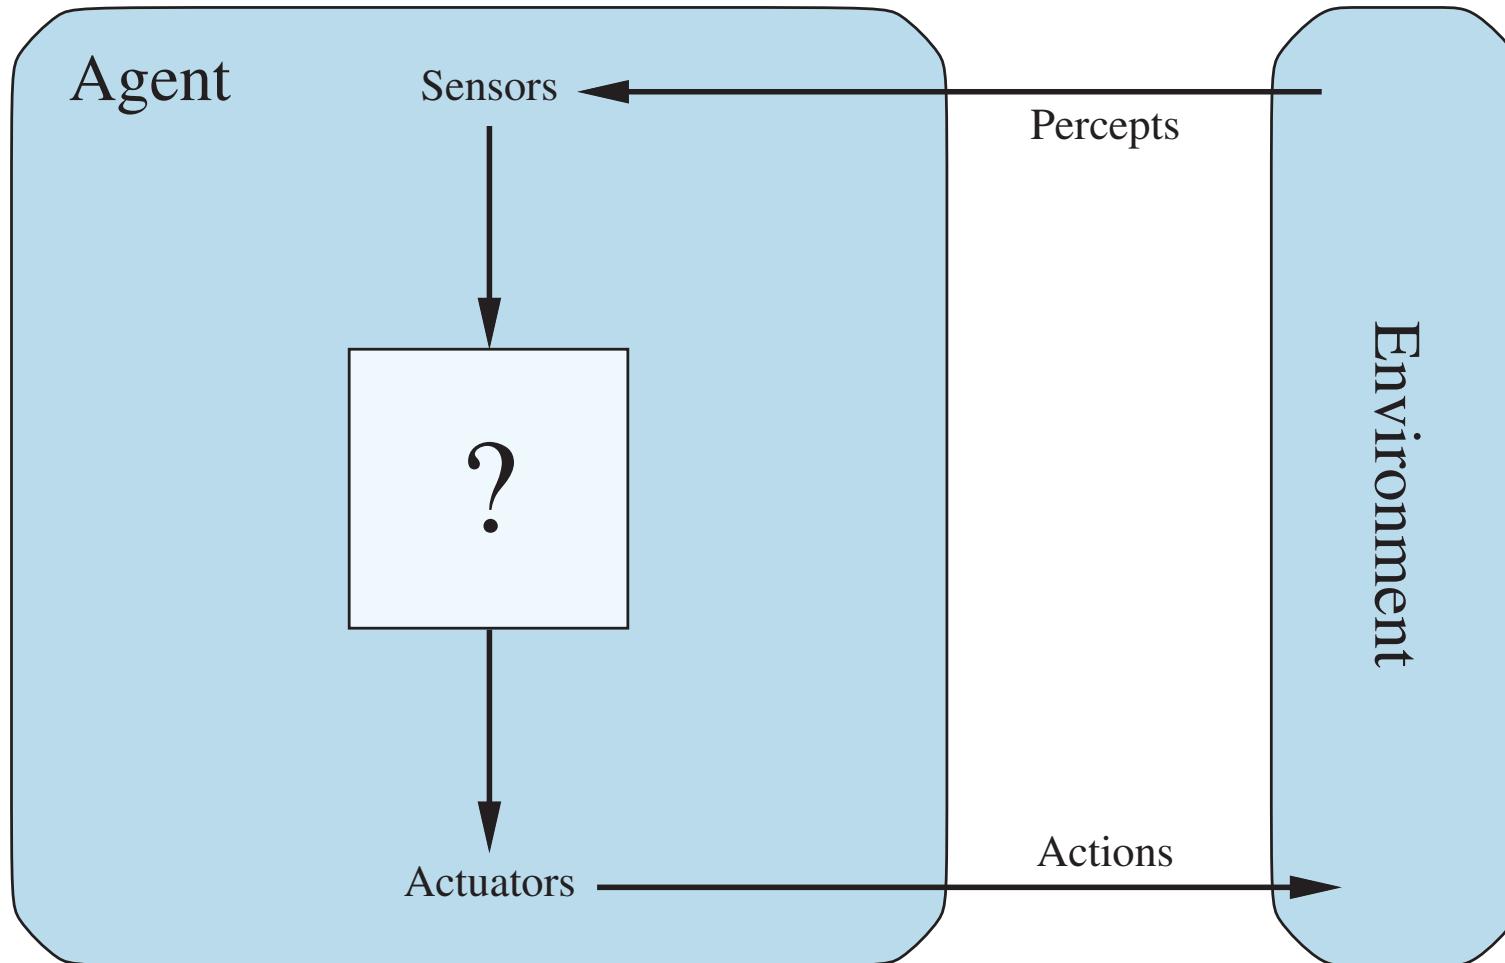
Reinforcement Learning (DL)



Reinforcement Learning (DL)



Agents interact with environments through sensors and actuators

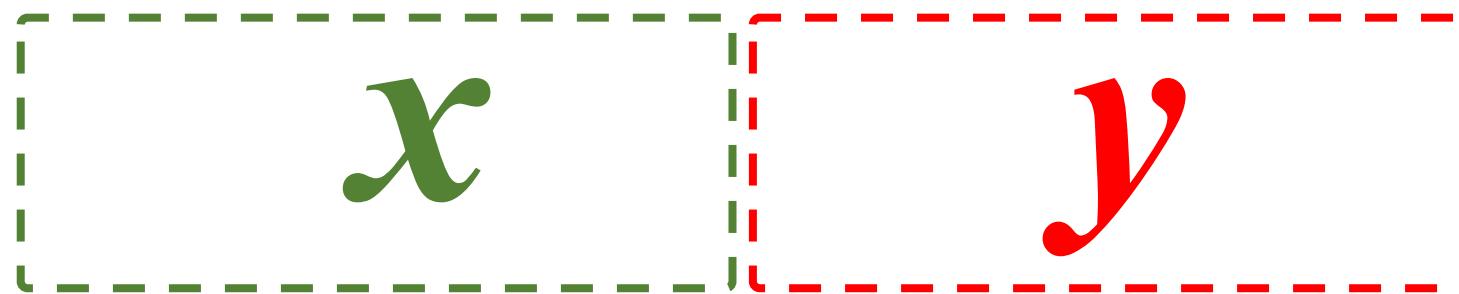


Machine Learning

Supervised Learning (Classification)

Learning from Examples

$$y = f(x)$$

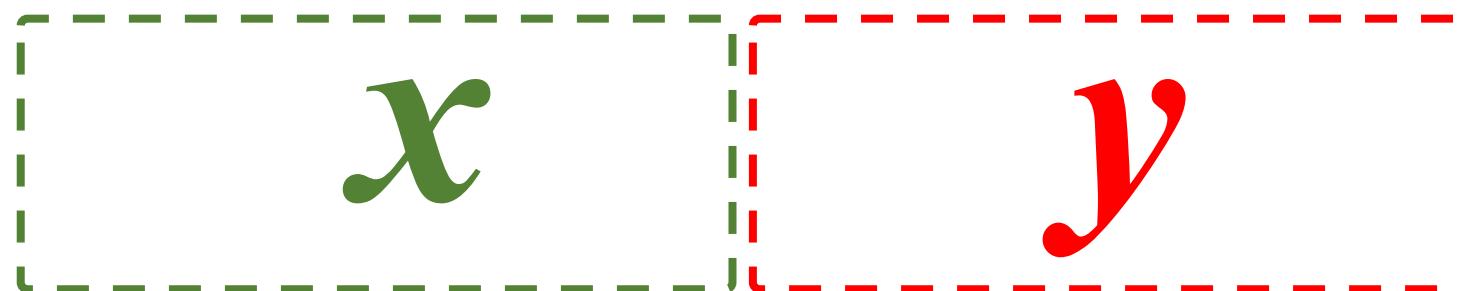


Machine Learning

Supervised Learning (Classification)

Learning from Examples

$$y = f(x)$$



input

Output

label

Iris flower data set

setosa



versicolor



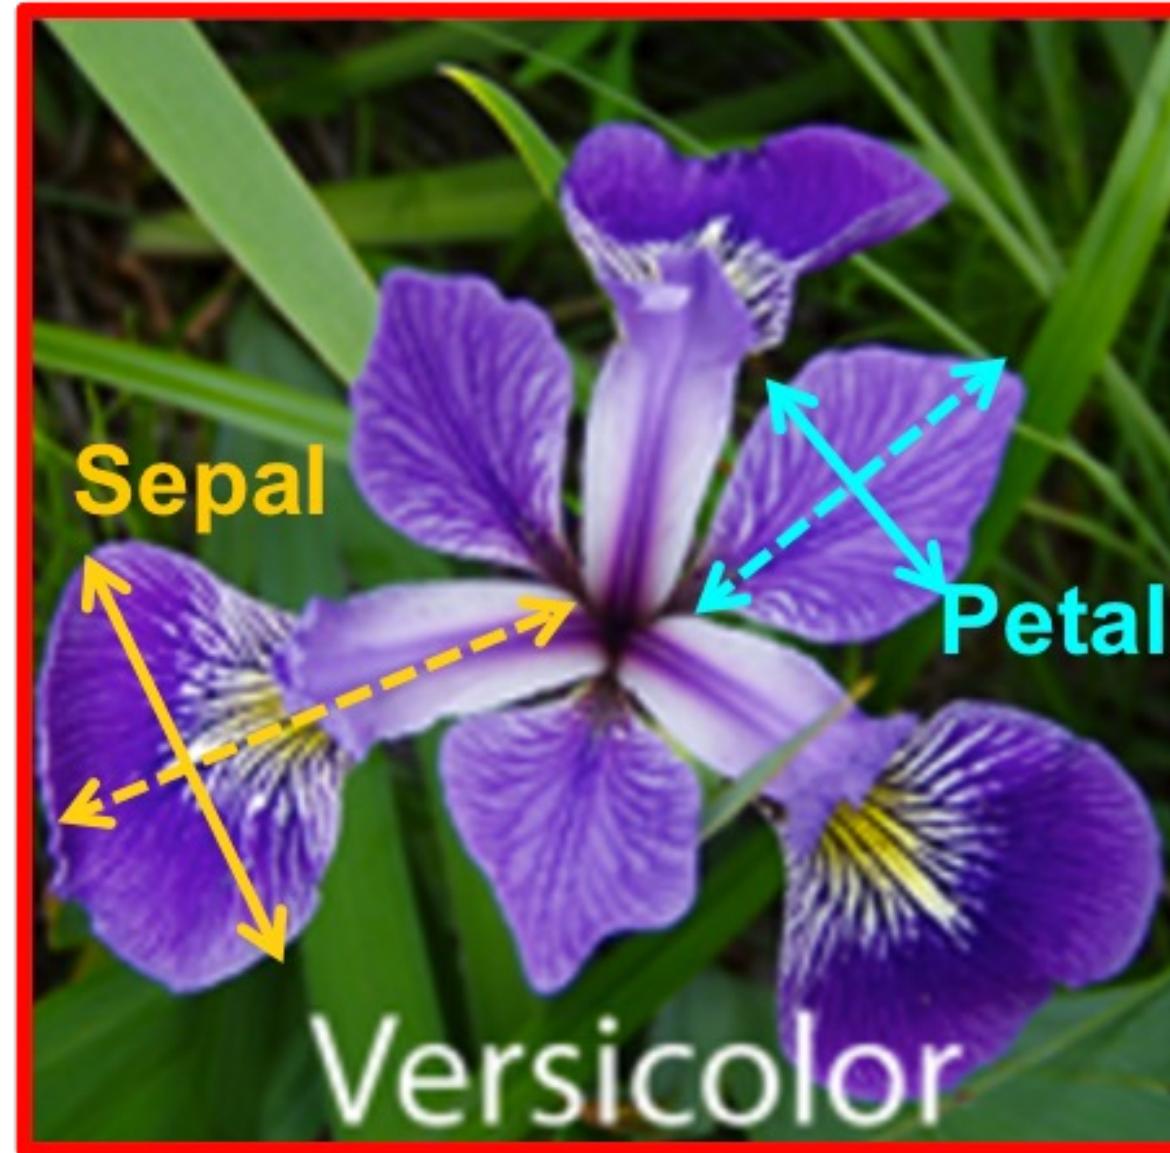
virginica



Source: https://en.wikipedia.org/wiki/Iris_flower_data_set

Source: <http://suruchifialoke.com/2016-10-13-machine-learning-tutorial-iris-classification/>

Iris Classification



iris.data

<https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data>

5.1,3.5,1.4,0.2,Iris-setosa
4.9,3.0,1.4,0.2,Iris-setosa
4.7,3.2,1.3,0.2,Iris-setosa
4.6,3.1,1.5,0.2,Iris-setosa
5.0,3.6,1.4,0.2,Iris-setosa
5.4,3.9,1.7,0.4,Iris-setosa
4.6,3.4,1.4,0.3,Iris-setosa
5.0,3.4,1.5,0.2,Iris-setosa
4.4,2.9,1.4,0.2,Iris-setosa
4.9,3.1,1.5,0.1,Iris-setosa
5.4,3.7,1.5,0.2,Iris-setosa
4.8,3.4,1.6,0.2,Iris-setosa
4.8,3.0,1.4,0.1,Iris-setosa
4.3,3.0,1.1,0.1,Iris-setosa
5.8,4.0,1.2,0.2,Iris-setosa
5.7,4.4,1.5,0.4,Iris-setosa
5.4,3.9,1.3,0.4,Iris-setosa
5.1,3.5,1.4,0.3,Iris-setosa
5.7,3.8,1.7,0.3,Iris-setosa
5.1,3.8,1.5,0.3,Iris-setosa
5.4,3.4,1.7,0.2,Iris-setosa
5.1,3.7,1.5,0.4,Iris-setosa
4.6,3.6,1.0,0.2,Iris-setosa
5.1,3.3,1.7,0.5,Iris-setosa
4.8,3.4,1.9,0.2,Iris-setosa
5.0,3.0,1.6,0.2,Iris-setosa
5.0,3.4,1.6,0.4,Iris-setosa

setosa



virginica



versicolor



Machine Learning

Supervised Learning (Classification)

Learning from Examples

$$\textcolor{red}{y} = f(\textcolor{green}{x})$$

5.1, 3.5, 1.4, 0.2, Iris-setosa
4.9, 3.0, 1.4, 0.2, Iris-setosa
4.7, 3.2, 1.3, 0.2, Iris-setosa
7.0, 3.2, 4.7, 1.4, Iris-versicolor
6.4, 3.2, 4.5, 1.5, Iris-versicolor
6.9, 3.1, 4.9, 1.5, Iris-versicolor
6.3, 3.3, 6.0, 2.5, Iris-virginica
5.8, 2.7, 5.1, 1.9, Iris-virginica
7.1, 3.0, 5.9, 2.1, Iris-virginica

Machine Learning

Supervised Learning (Classification)

Learning from Examples

$$\textcolor{red}{y} = f(\textcolor{green}{x})$$

<i>Example</i>
5.1, 3.5, 1.4, 0.2, Iris-setosa
4.9, 3.0, 1.4, 0.2, Iris-setosa
4.7, 3.2, 1.3, 0.2, Iris-setosa
7.0, 3.2, 4.7, 1.4, Iris-versicolor
6.4, 3.2, 4.5, 1.5, Iris-versicolor
6.9, 3.1, 4.9, 1.5, Iris-versicolor
6.3, 3.3, 6.0, 2.5, Iris-virginica
5.8, 2.7, 5.1, 1.9, Iris-virginica
7.1, 3.0, 5.9, 2.1, Iris-virginica

Machine Learning

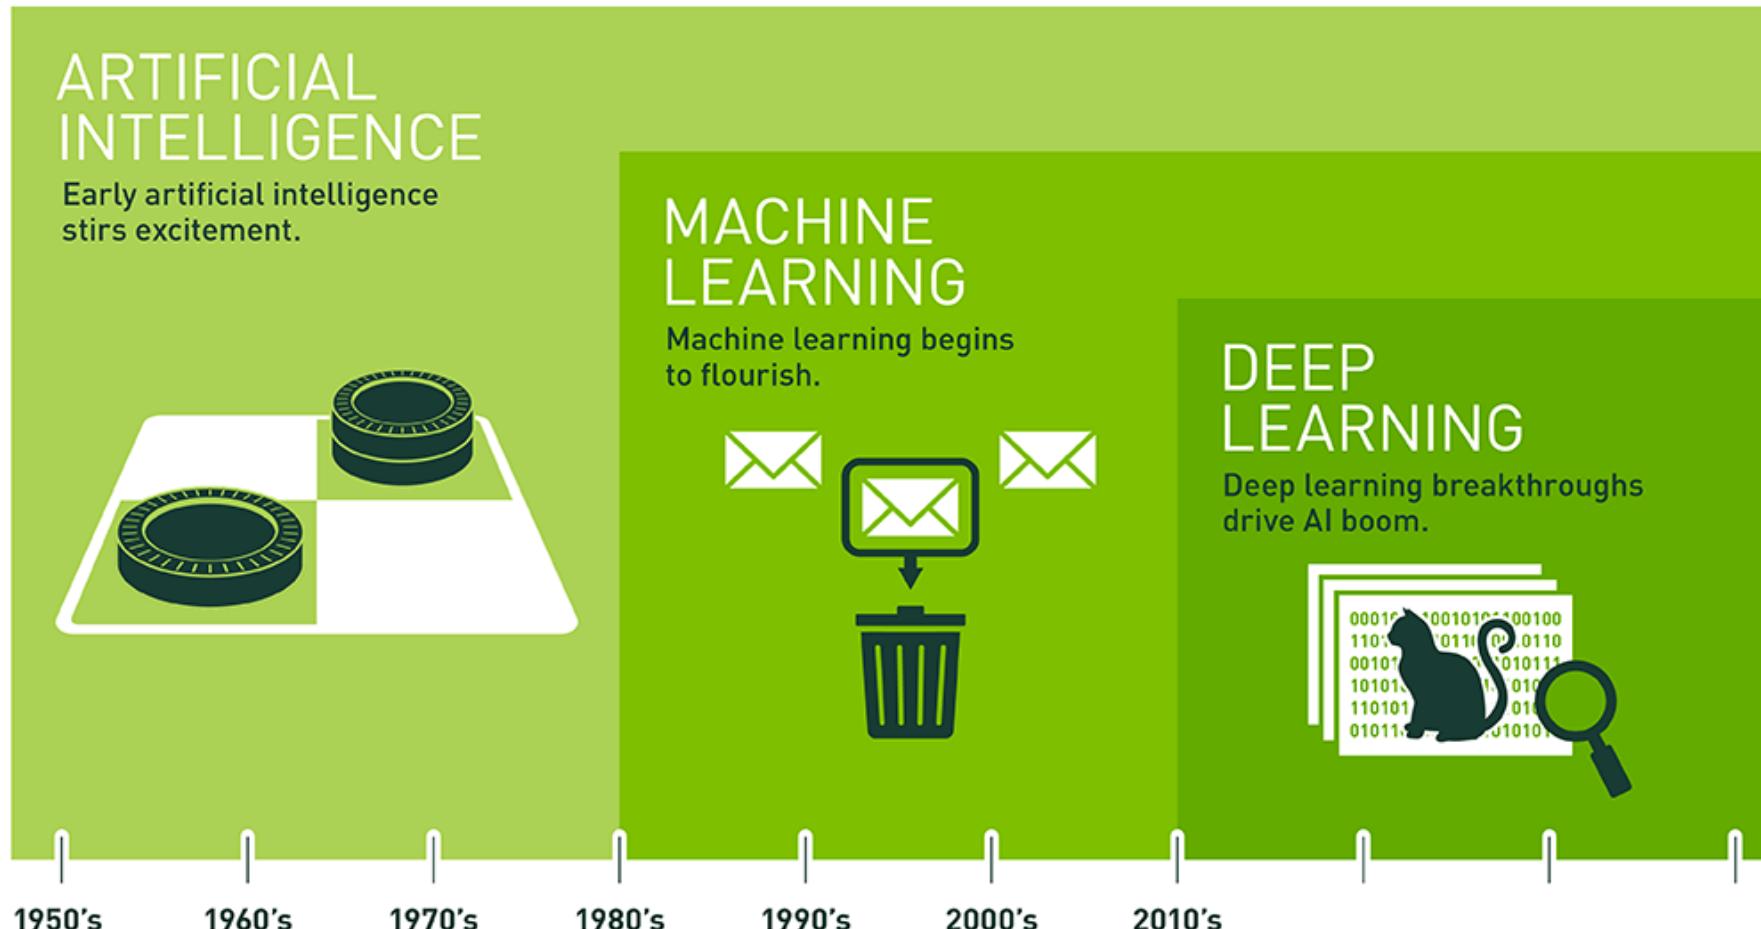
Supervised Learning (Classification)

Learning from Examples

<i>Example</i>	$y = f(x)$
x	y
5.1, 3.5, 1.4, 0.2 Iris-setosa	
4.9, 3.0, 1.4, 0.2 Iris-setosa	
4.7, 3.2, 1.3, 0.2 Iris-setosa	
7.0, 3.2, 4.7, 1.4 Iris-versicolor	
6.4, 3.2, 4.5, 1.5 Iris-versicolor	
6.9, 3.1, 4.9, 1.5 Iris-versicolor	
6.3, 3.3, 6.0, 2.5 Iris-virginica	
5.8, 2.7, 5.1, 1.9 Iris-virginica	
7.1, 3.0, 5.9, 2.1 Iris-virginica	

Artificial Intelligence

Machine Learning & Deep Learning



Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

AI, ML, DL

Artificial Intelligence (AI)

Machine Learning (ML)

Supervised
Learning

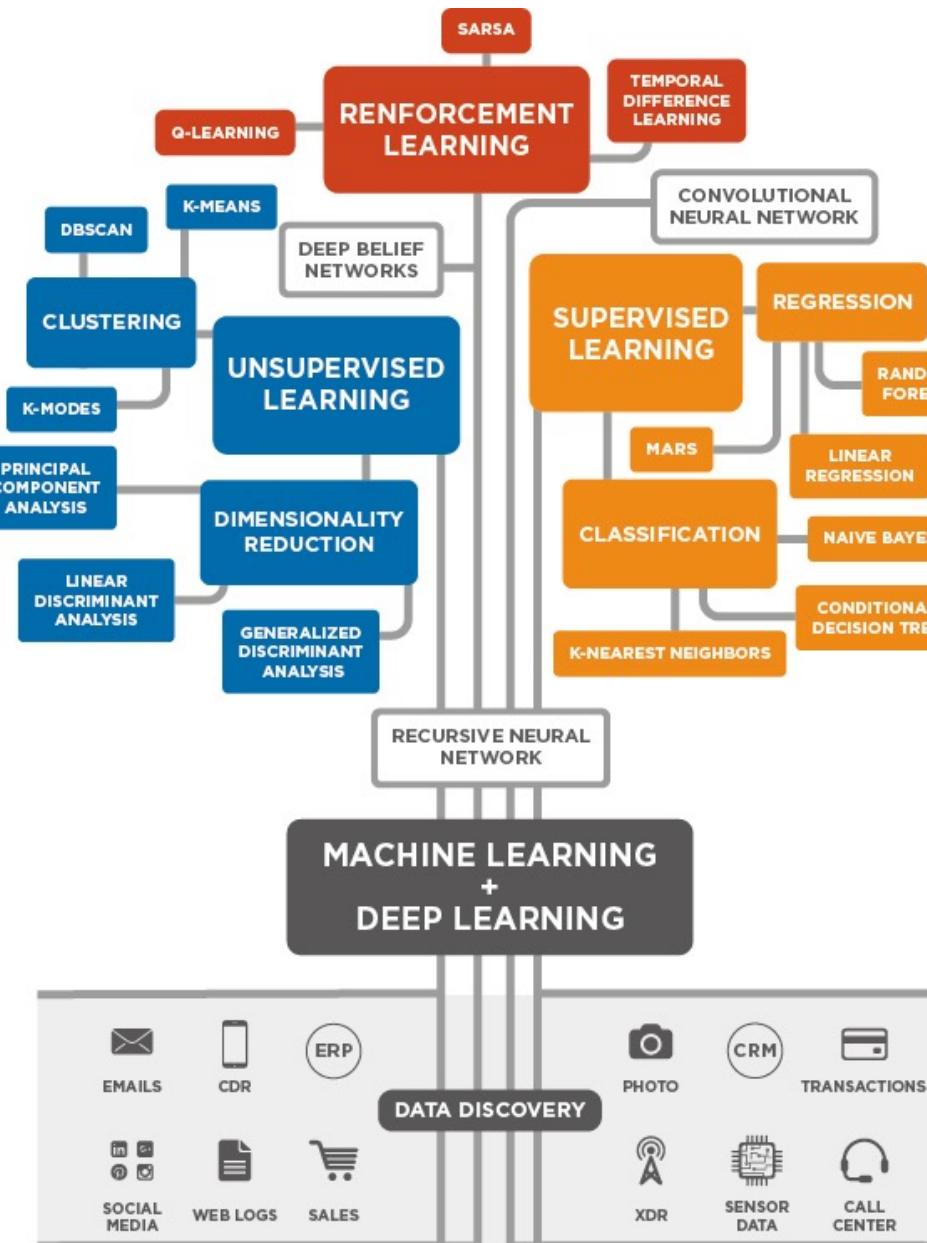
Unsupervised
Learning

Deep Learning (DL)
CNN
RNN LSTM GRU
GAN

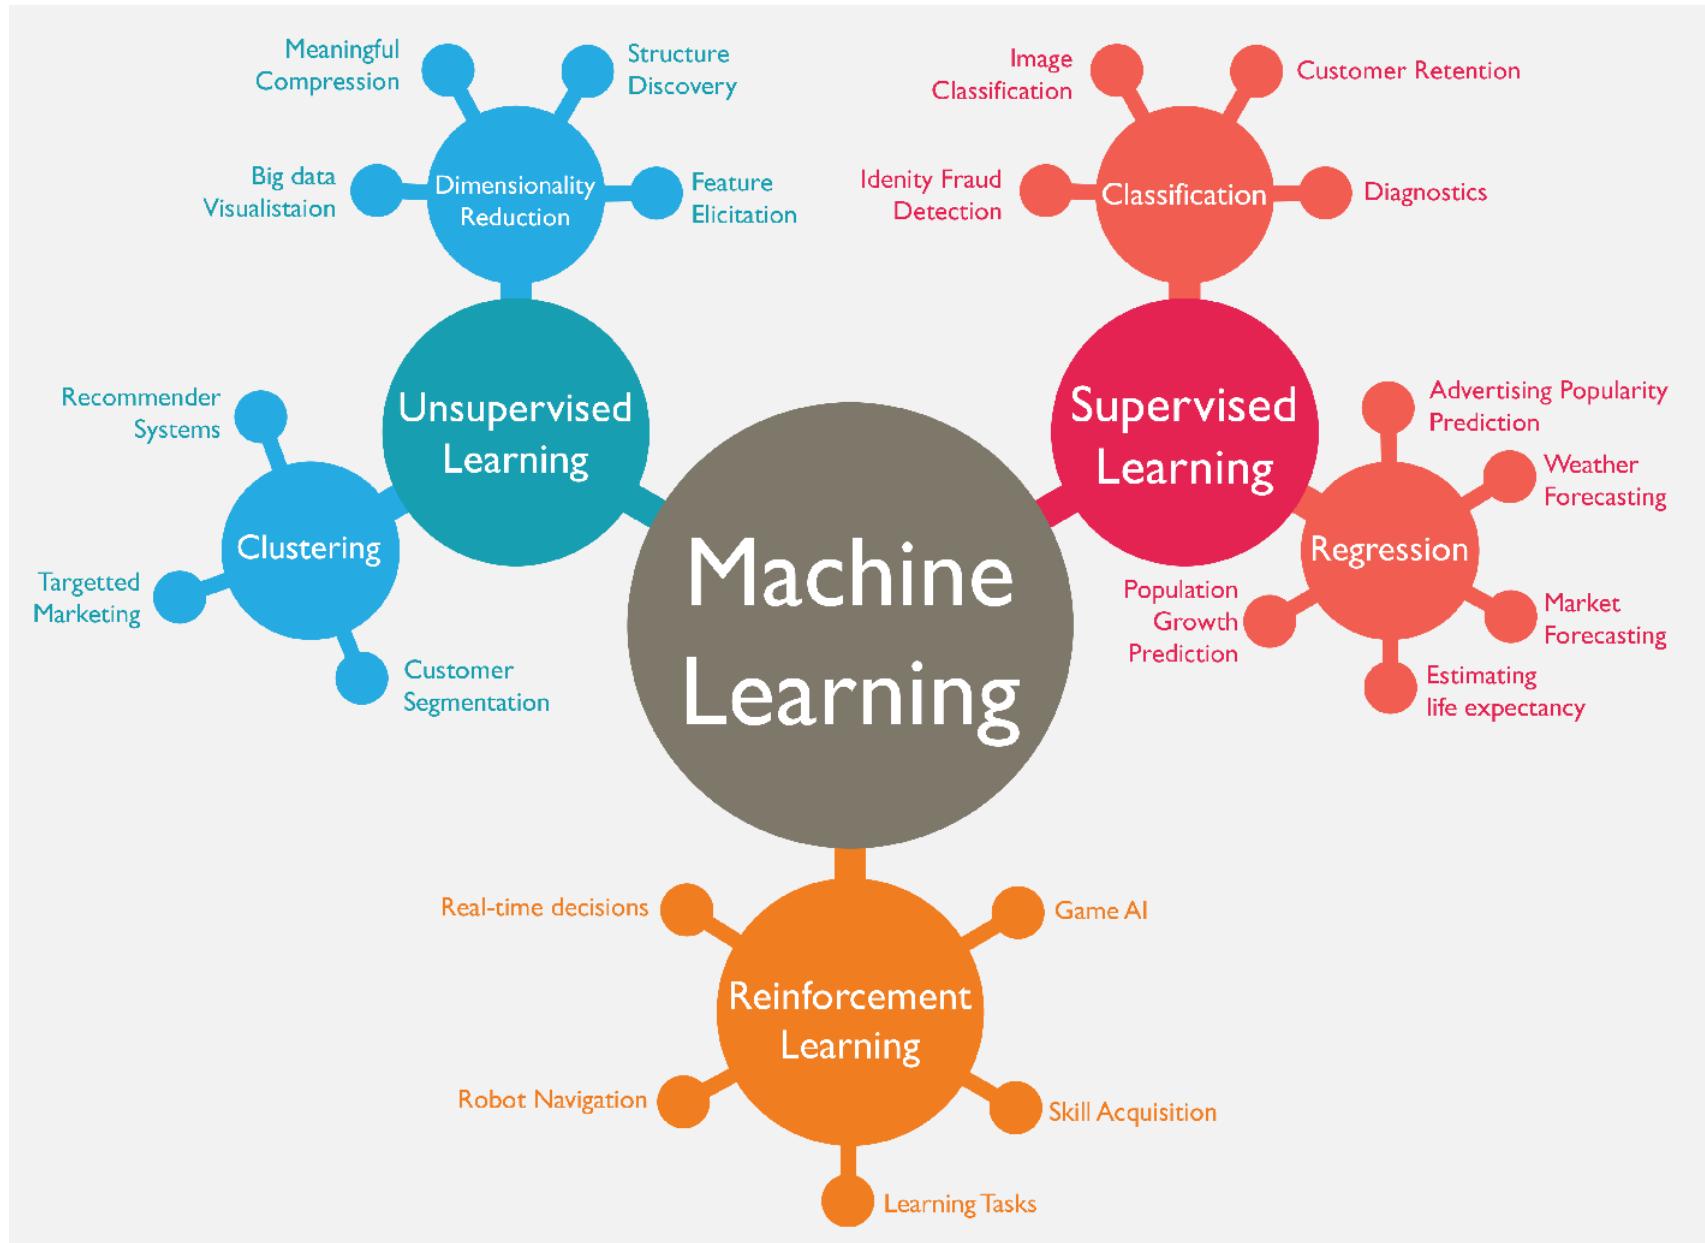
Semi-supervised
Learning

Reinforcement
Learning

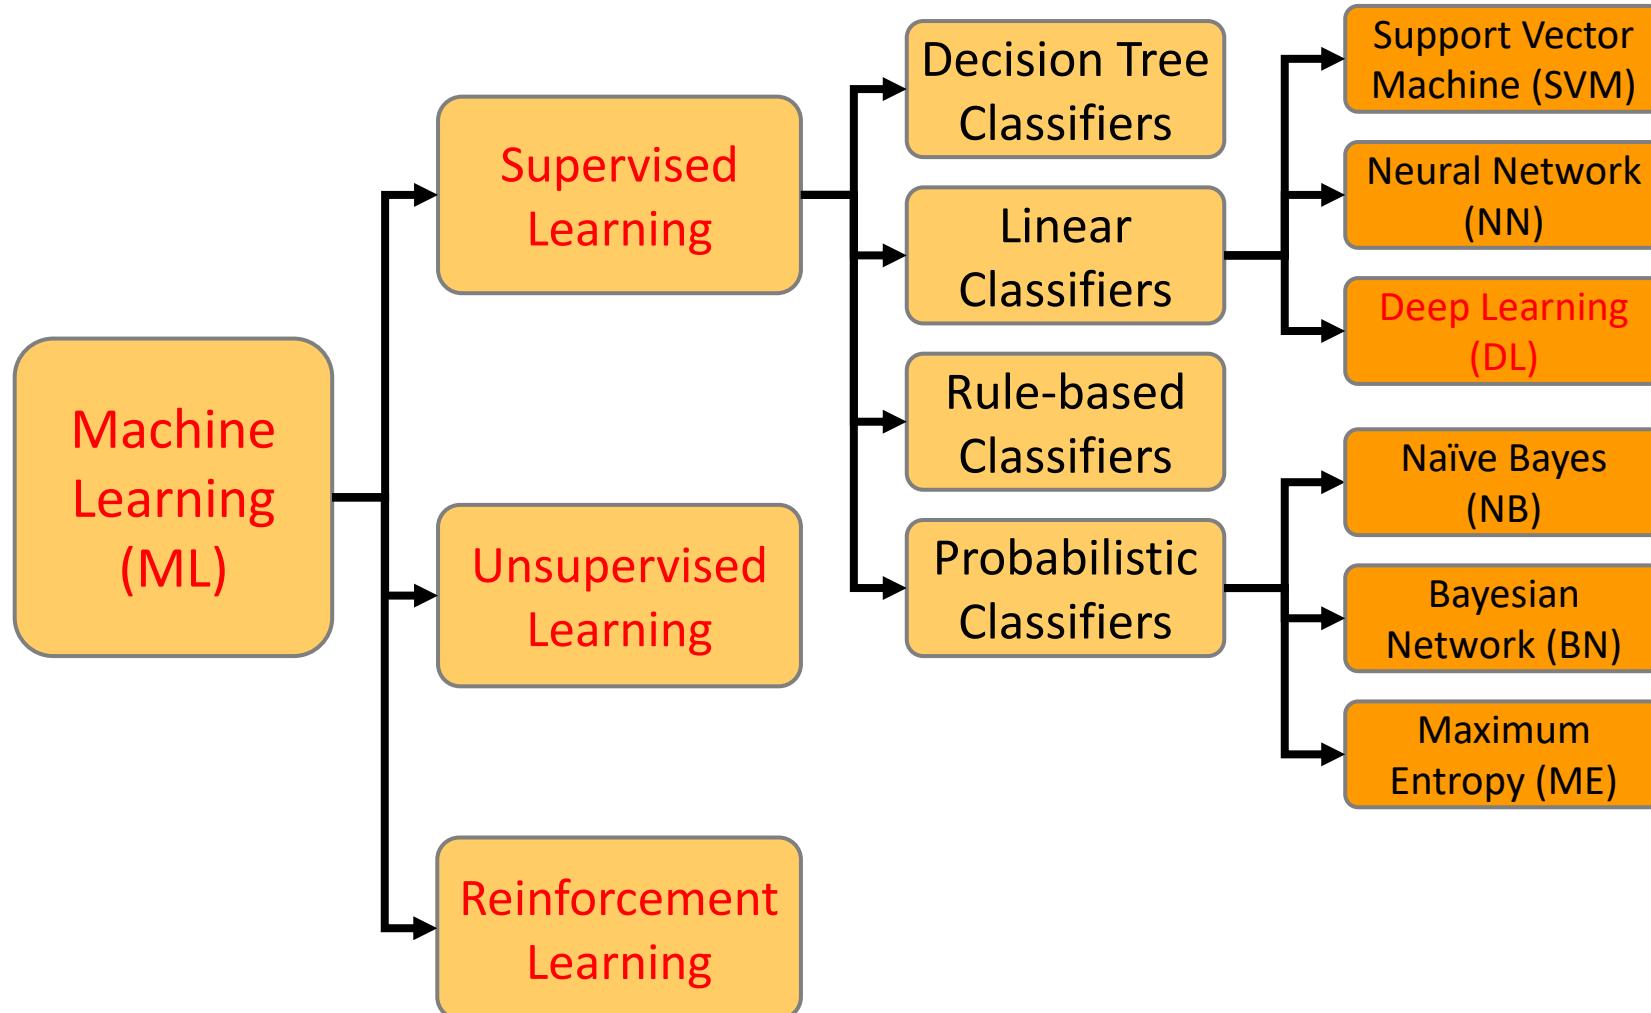
3 Machine Learning Algorithms



Machine Learning (ML)



Machine Learning (ML) / Deep Learning (DL)



Machine Learning Models

Deep Learning

Kernel

Association rules

Ensemble

Decision tree

Dimensionality reduction

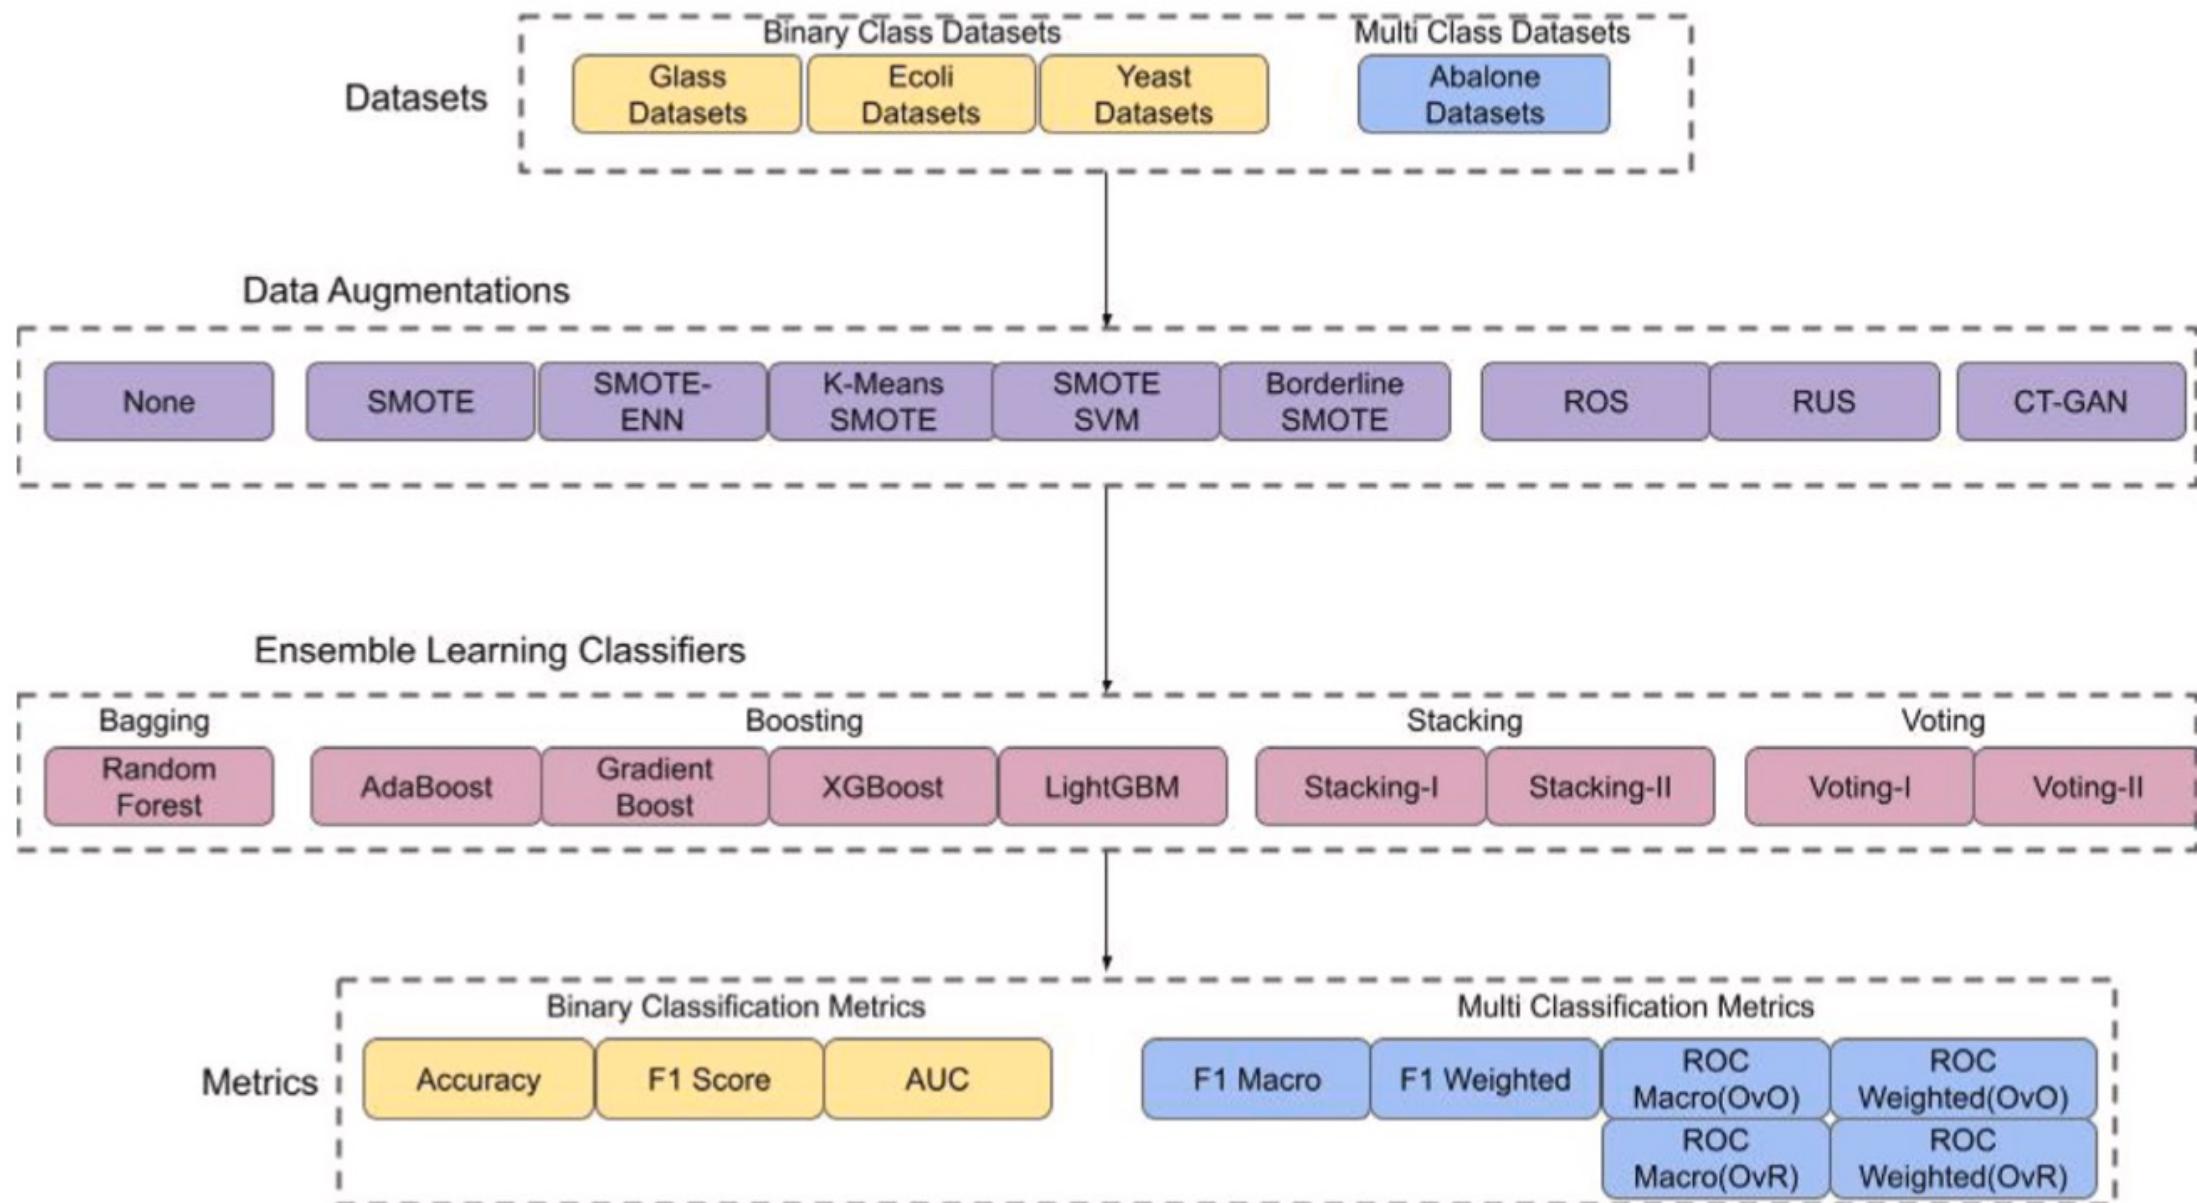
Clustering

Regression Analysis

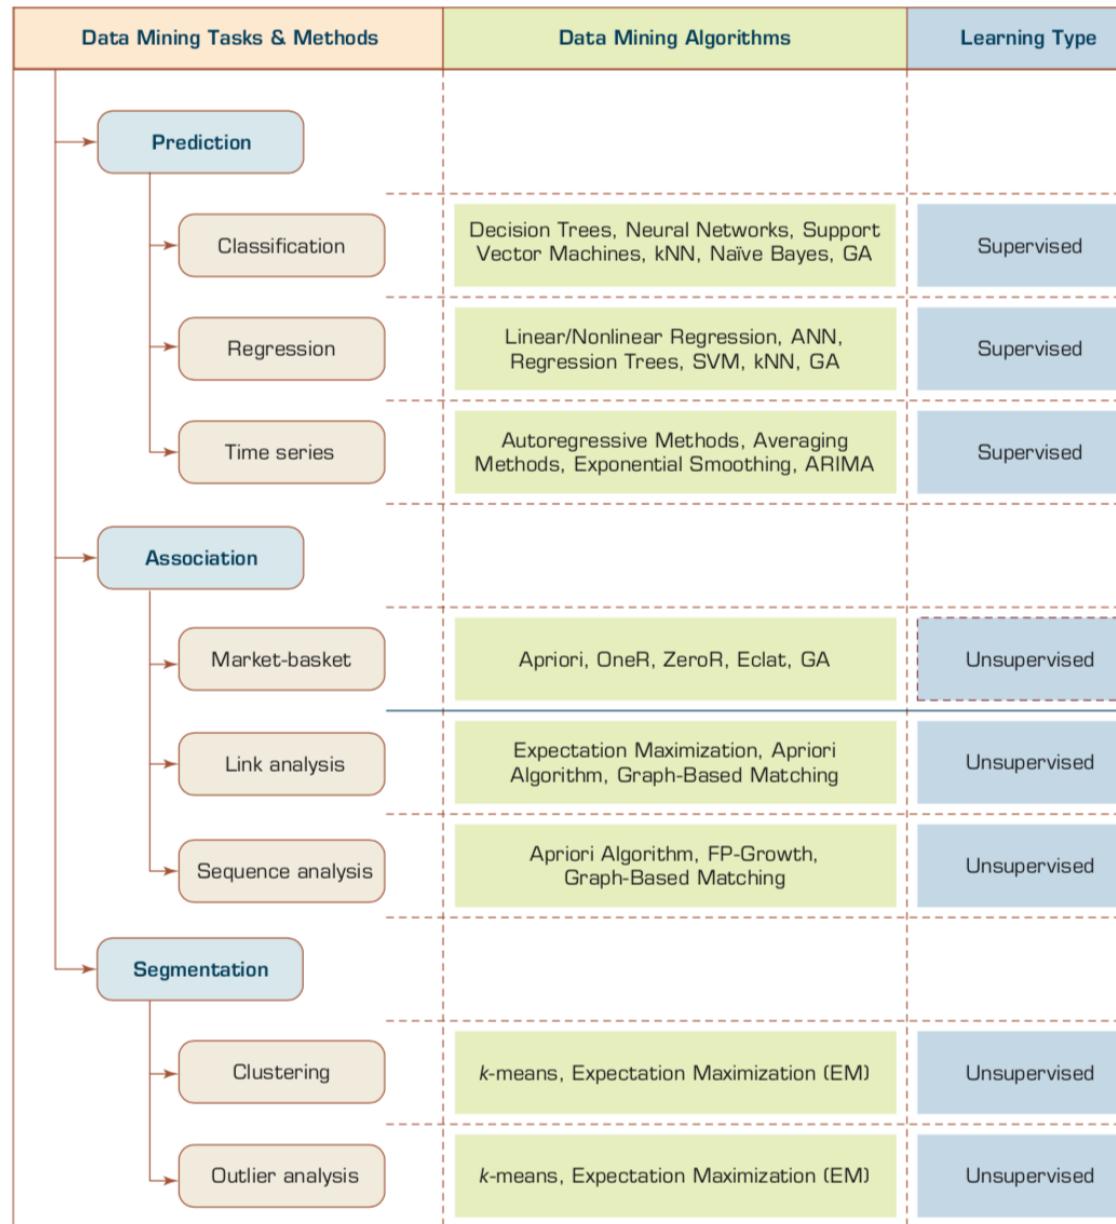
Bayesian

Instance based

Ensemble Learning and Data Augmentations (DA)



Machine Learning: Data Mining Tasks & Methods



Data Mining Methods

- **Supervised Learning**
 - **Classification**
 - Class Label Prediction
 - **Regression**
 - Numeric Value Prediction
- **Unsupervised Learning**
 - **Clustering**
 - **Association**

Scikit-Learn

Machine Learning in Python

Scikit-Learn

scikit-learn

Machine Learning in Python

Getting Started

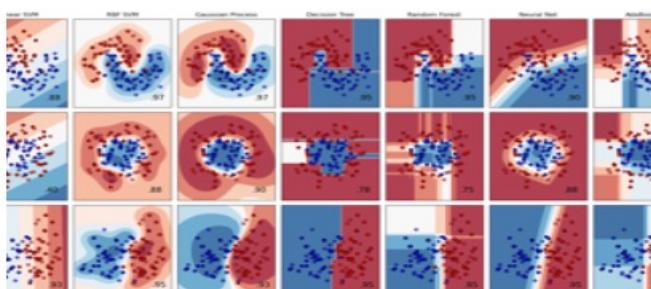
Release Highlights for 1.5

- Simple and efficient tools for predictive data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.
Algorithms: [Gradient boosting](#), [nearest neighbors](#), [random forest](#), [logistic regression](#), and [more...](#)



Examples

Dimensionality reduction

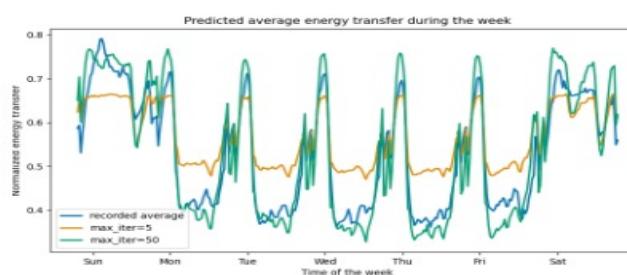
Reducing the number of random variables to consider.

Applications: Visualization, increased efficiency.
Algorithms: [PCA](#), [feature selection](#), [non-negative matrix factorization](#), and [more...](#)

Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, stock prices.
Algorithms: [Gradient boosting](#), [nearest neighbors](#), [random forest](#), [ridge](#), and [more...](#)



Examples

Model selection

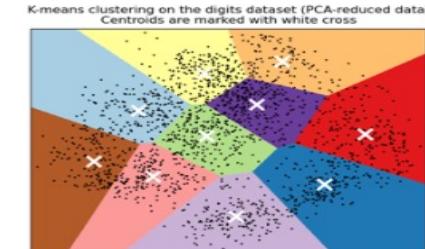
Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning.
Algorithms: [Grid search](#), [cross validation](#), [metrics](#), and [more...](#)

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, grouping experiment outcomes.
Algorithms: [k-Means](#), [HDBSCAN](#), [hierarchical clustering](#), and [more...](#)



Examples

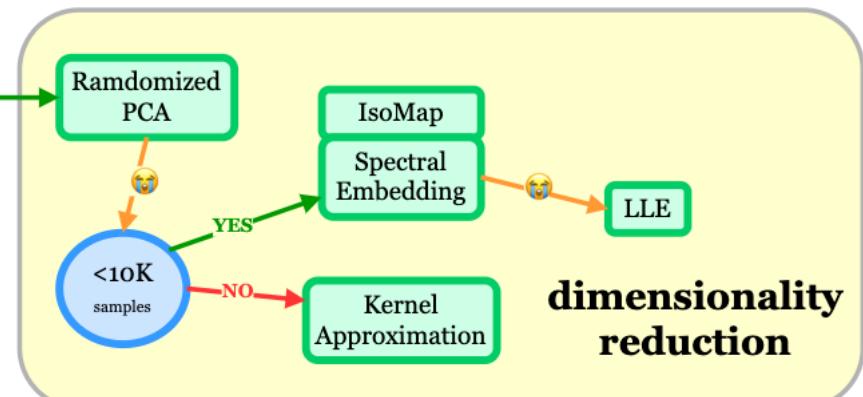
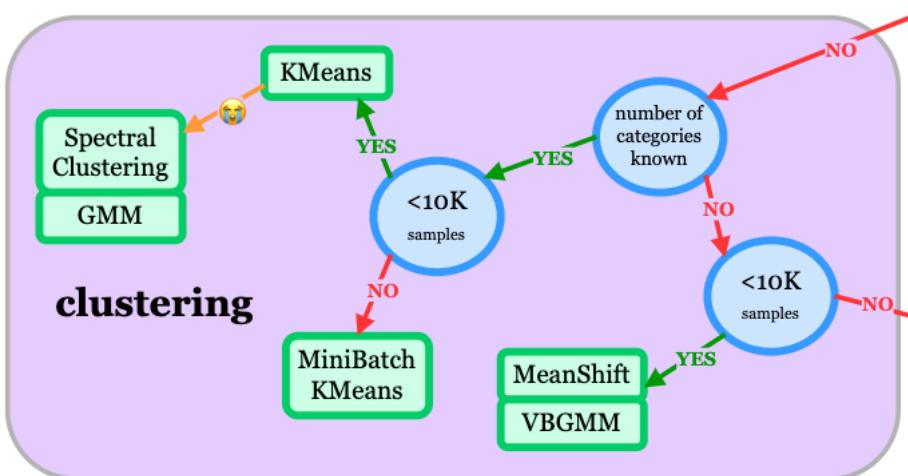
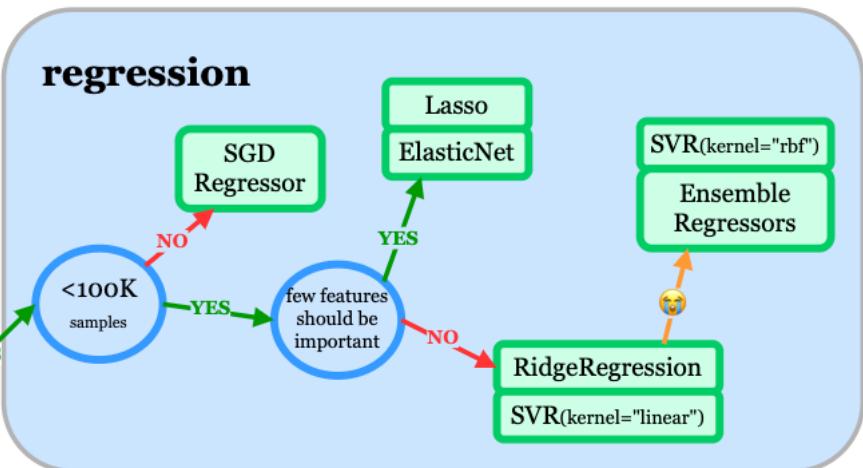
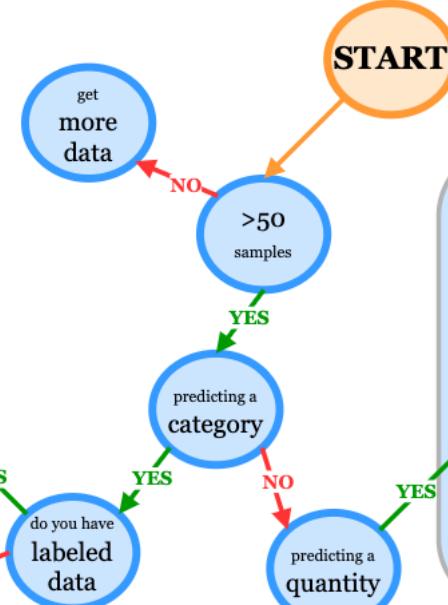
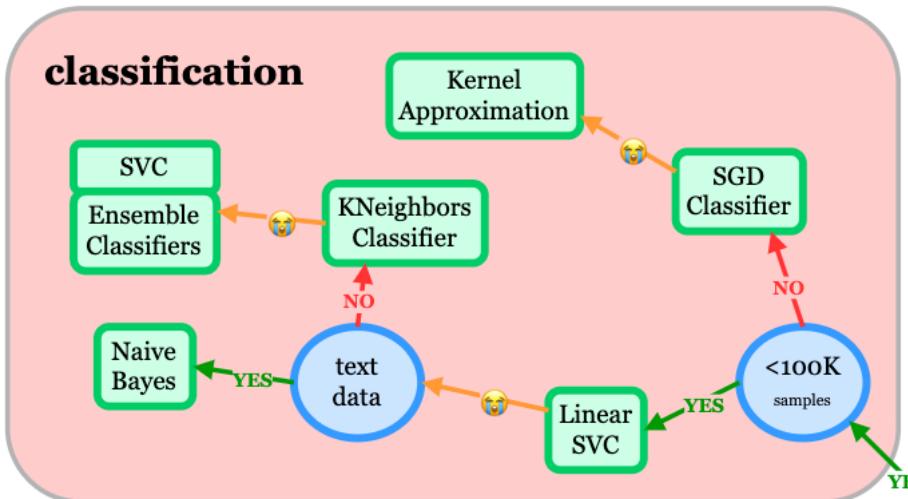
Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.
Algorithms: [Preprocessing](#), [feature extraction](#), and [more...](#)

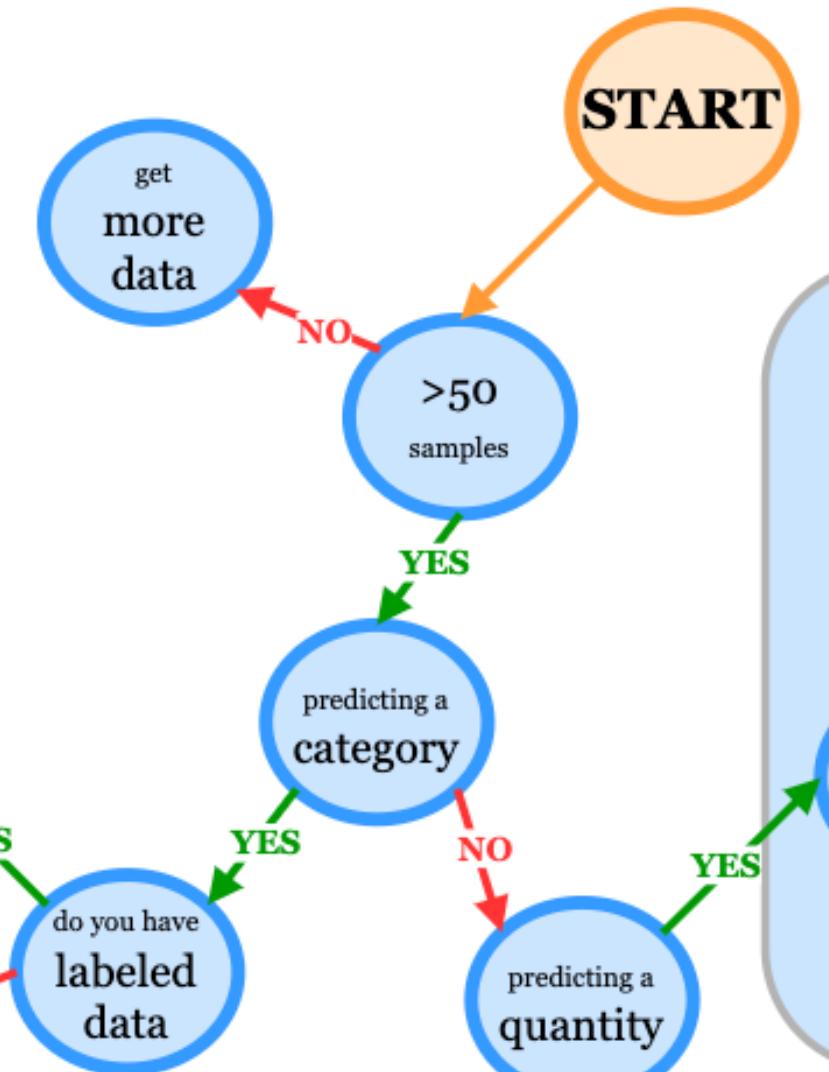
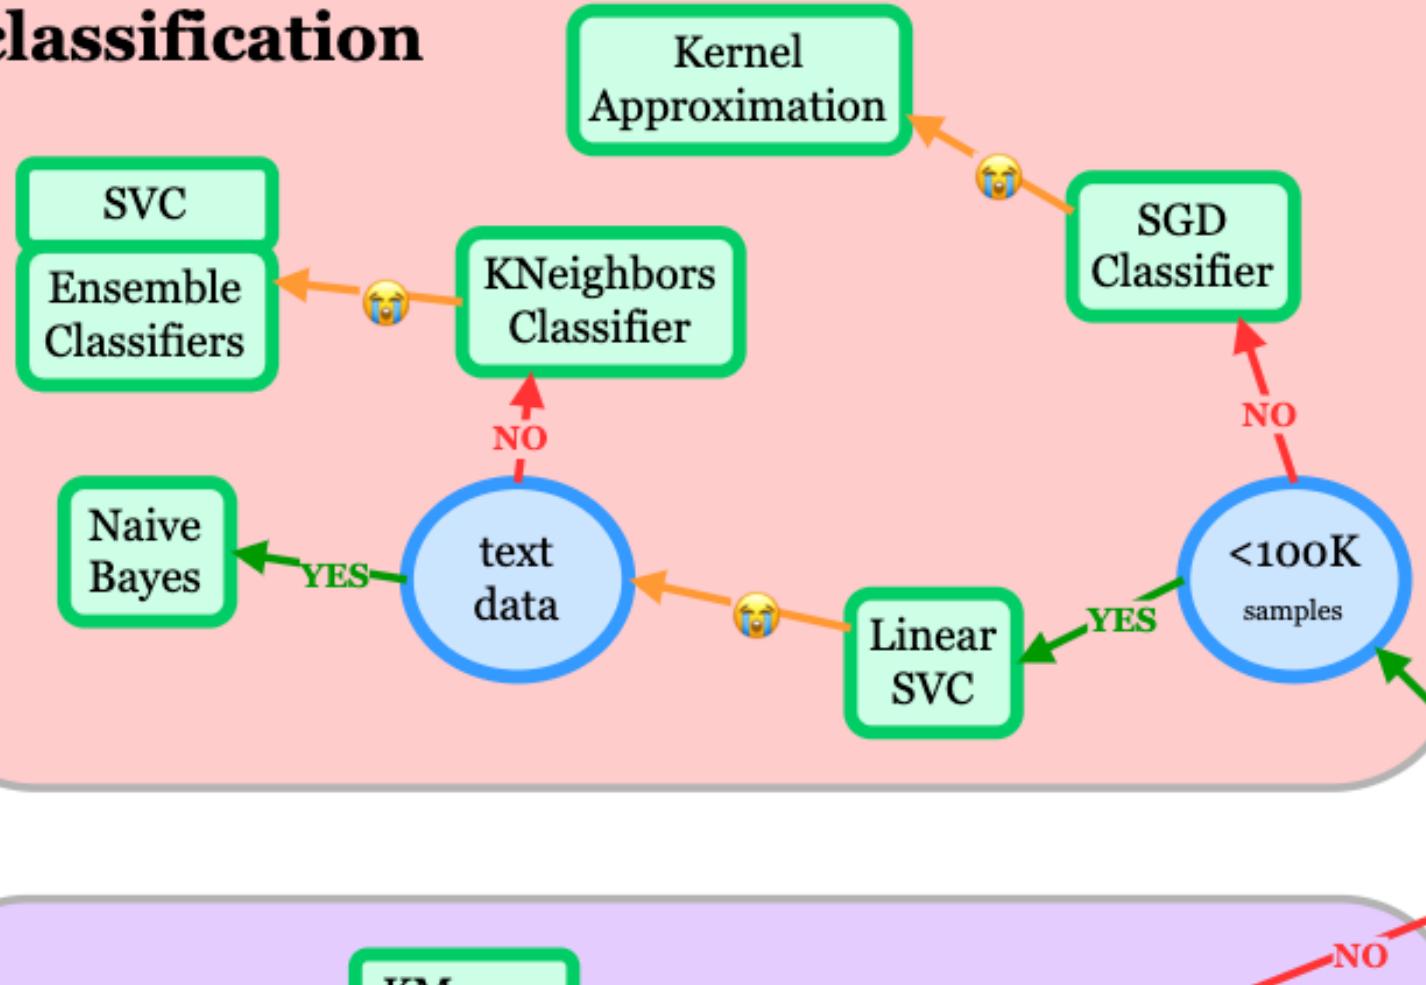
Scikit-Learn Machine Learning Map

scikit
learn
algorithm cheat sheet



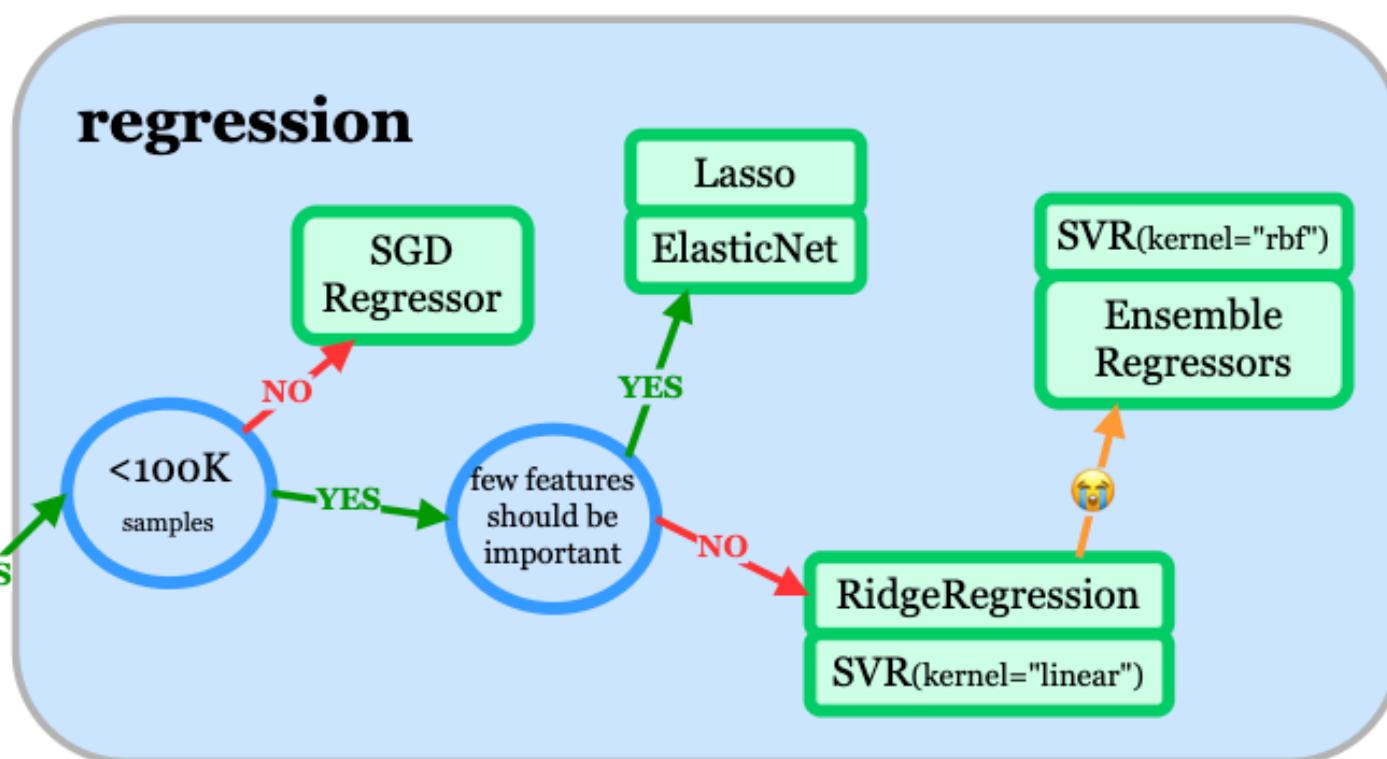
Scikit-Learn Machine Learning Map

classification

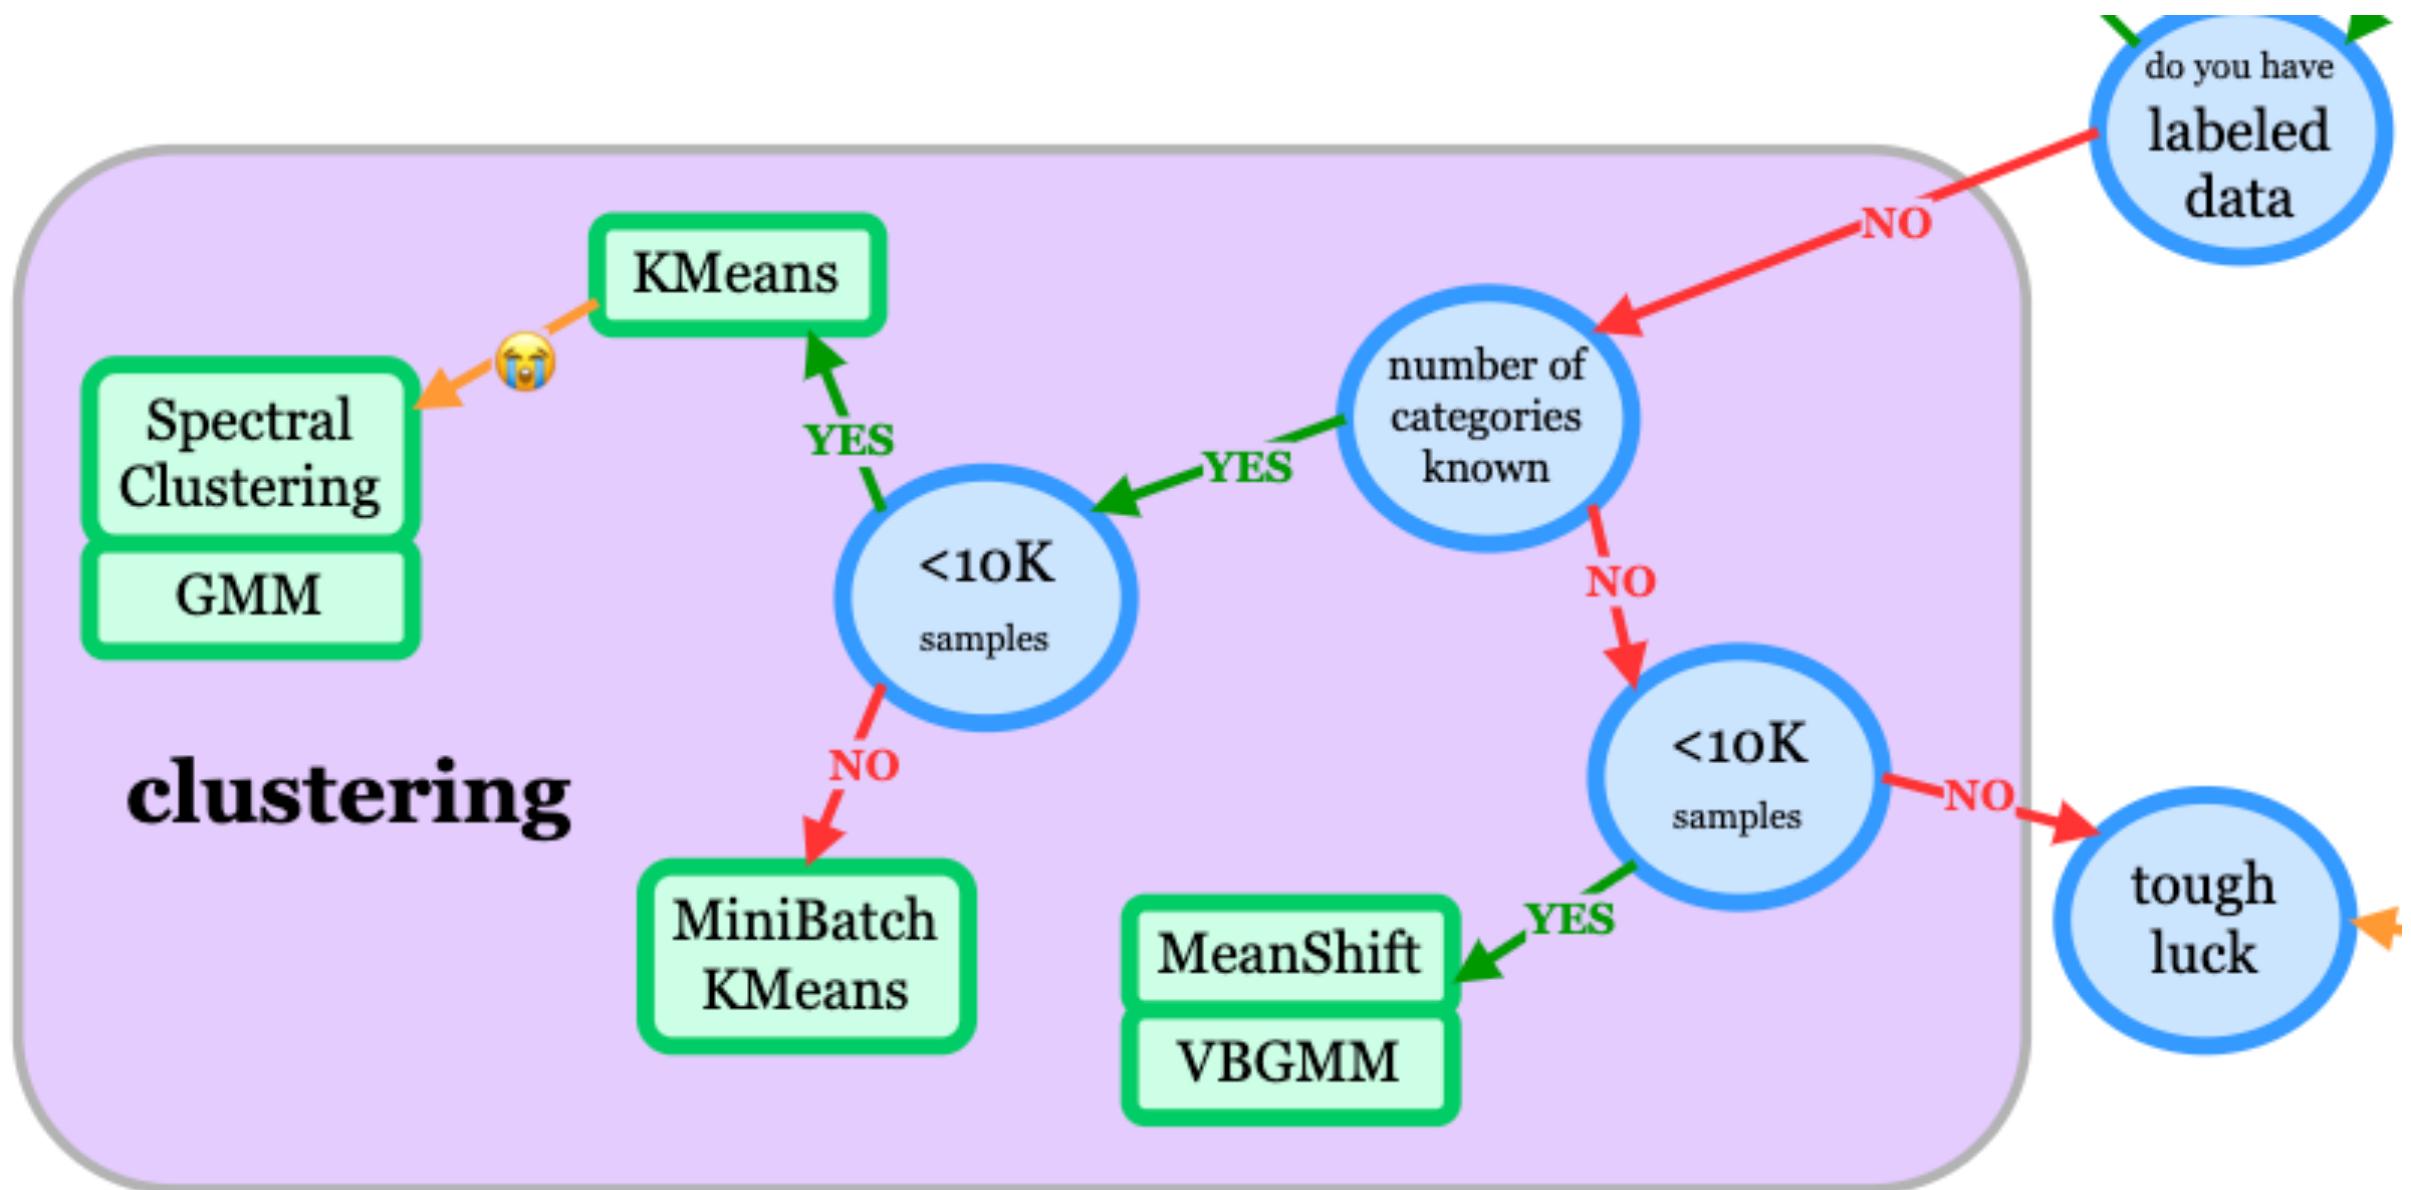


Scikit-Learn Machine Learning Map

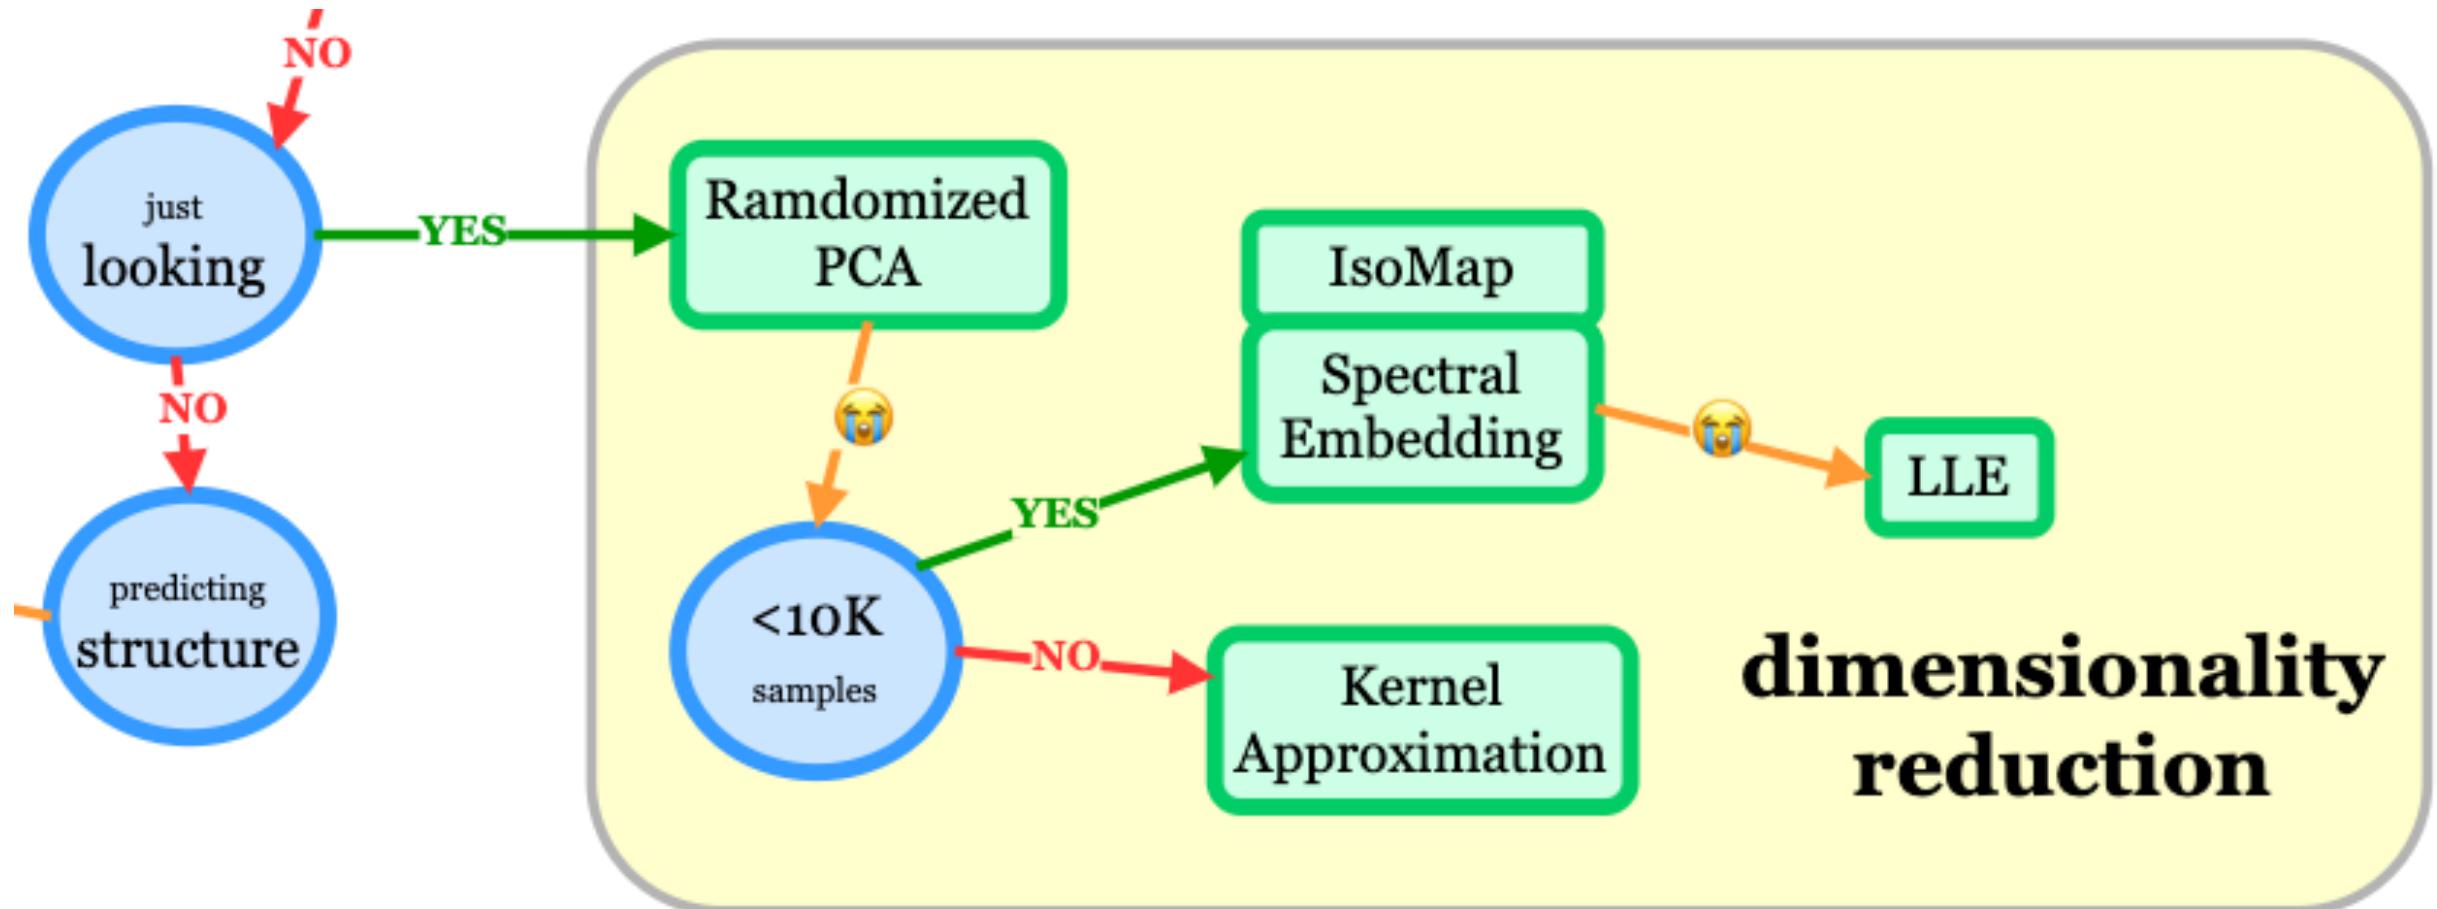
scikit
learn
algorithm cheat sheet



Scikit-Learn Machine Learning Map



Scikit-Learn Machine Learning Map



Iris flower data set

setosa



versicolor



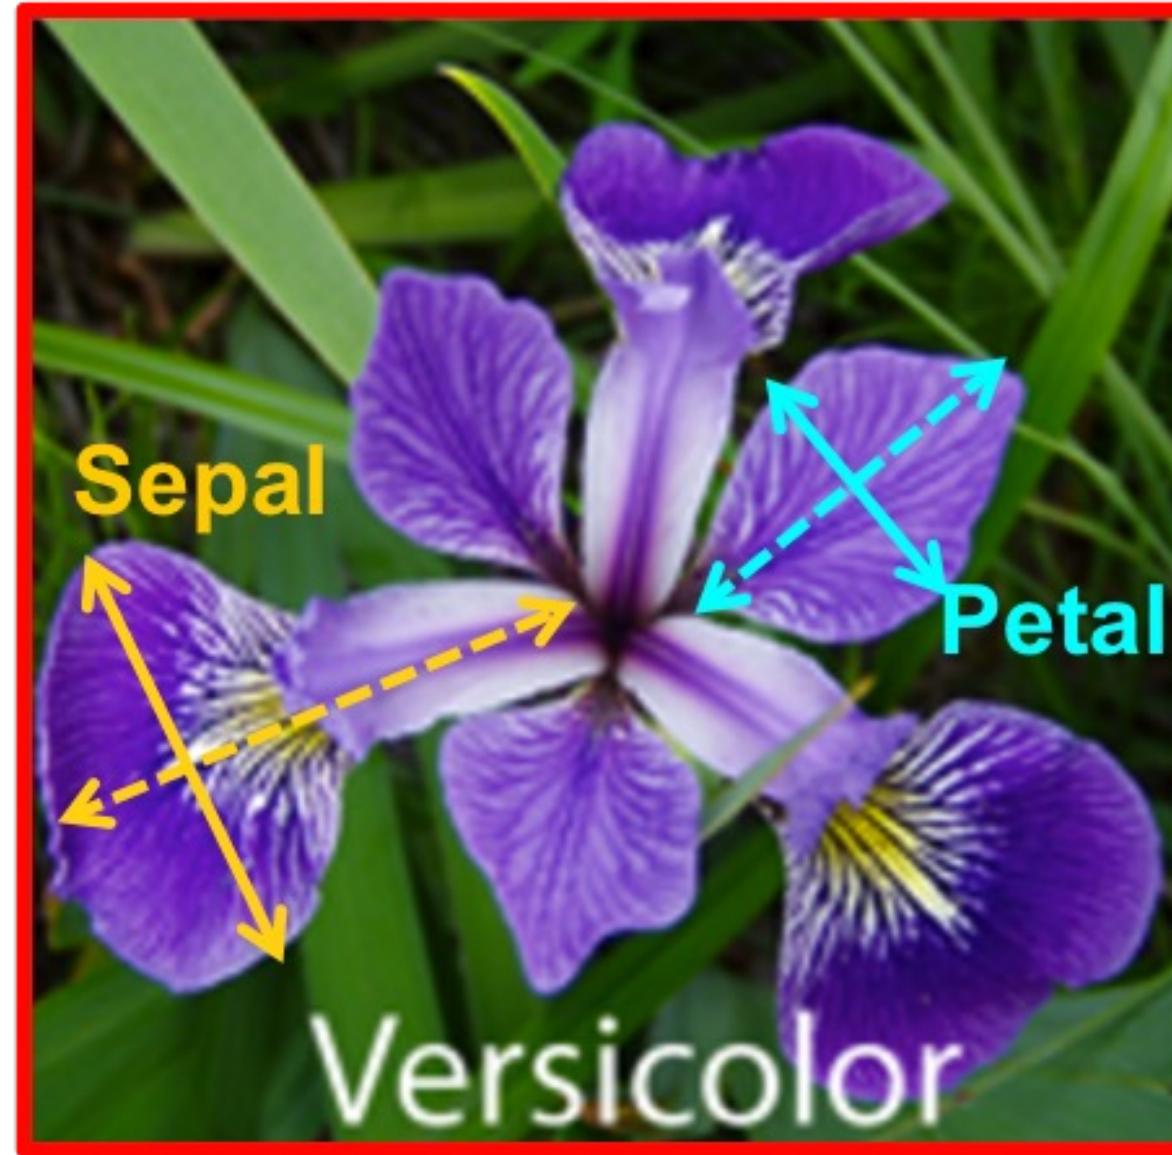
virginica



Source: https://en.wikipedia.org/wiki/Iris_flower_data_set

Source: <http://suruchifialoke.com/2016-10-13-machine-learning-tutorial-iris-classification/>

Iris Classification



iris.data

<https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data>

5.1,3.5,1.4,0.2,Iris-setosa
4.9,3.0,1.4,0.2,Iris-setosa
4.7,3.2,1.3,0.2,Iris-setosa
4.6,3.1,1.5,0.2,Iris-setosa
5.0,3.6,1.4,0.2,Iris-setosa
5.4,3.9,1.7,0.4,Iris-setosa
4.6,3.4,1.4,0.3,Iris-setosa
5.0,3.4,1.5,0.2,Iris-setosa
4.4,2.9,1.4,0.2,Iris-setosa
4.9,3.1,1.5,0.1,Iris-setosa
5.4,3.7,1.5,0.2,Iris-setosa
4.8,3.4,1.6,0.2,Iris-setosa
4.8,3.0,1.4,0.1,Iris-setosa
4.3,3.0,1.1,0.1,Iris-setosa
5.8,4.0,1.2,0.2,Iris-setosa
5.7,4.4,1.5,0.4,Iris-setosa
5.4,3.9,1.3,0.4,Iris-setosa
5.1,3.5,1.4,0.3,Iris-setosa
5.7,3.8,1.7,0.3,Iris-setosa
5.1,3.8,1.5,0.3,Iris-setosa
5.4,3.4,1.7,0.2,Iris-setosa
5.1,3.7,1.5,0.4,Iris-setosa
4.6,3.6,1.0,0.2,Iris-setosa
5.1,3.3,1.7,0.5,Iris-setosa
4.8,3.4,1.9,0.2,Iris-setosa
5.0,3.0,1.6,0.2,Iris-setosa
5.0,3.4,1.6,0.4,Iris-setosa

setosa



virginica



versicolor



Machine Learning

Supervised Learning (Classification)

Learning from Examples

5.1, 3.5, 1.4, 0.2, Iris-setosa
4.9, 3.0, 1.4, 0.2, Iris-setosa
4.7, 3.2, 1.3, 0.2, Iris-setosa
7.0, 3.2, 4.7, 1.4, Iris-versicolor
6.4, 3.2, 4.5, 1.5, Iris-versicolor
6.9, 3.1, 4.9, 1.5, Iris-versicolor
6.3, 3.3, 6.0, 2.5, Iris-virginica
5.8, 2.7, 5.1, 1.9, Iris-virginica
7.1, 3.0, 5.9, 2.1, Iris-virginica

Machine Learning

Supervised Learning (Classification)

Learning from Examples

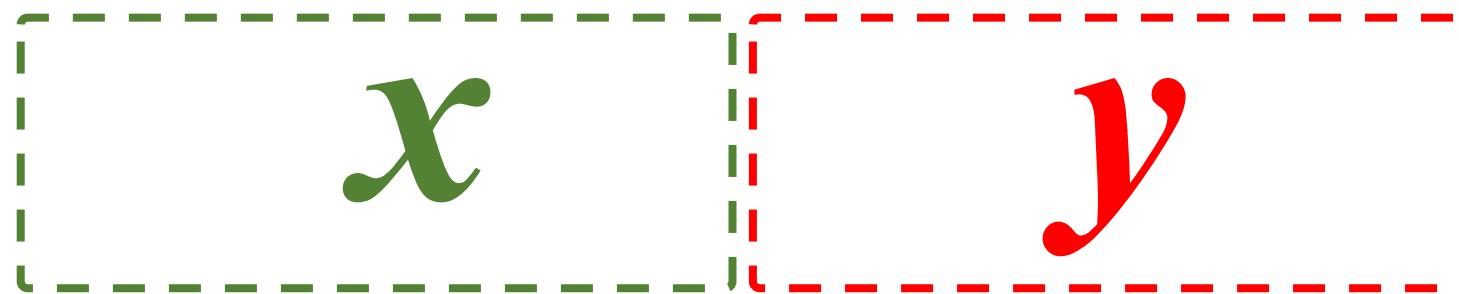
x	$y = f(x)$	y
5.1, 3.5, 1.4, 0.2	Iris-setosa	
4.9, 3.0, 1.4, 0.2	Iris-setosa	
4.7, 3.2, 1.3, 0.2	Iris-setosa	
7.0, 3.2, 4.7, 1.4	Iris-versicolor	
6.4, 3.2, 4.5, 1.5	Iris-versicolor	
6.9, 3.1, 4.9, 1.5	Iris-versicolor	
6.3, 3.3, 6.0, 2.5	Iris-virginica	
5.8, 2.7, 5.1, 1.9	Iris-virginica	
7.1, 3.0, 5.9, 2.1	Iris-virginica	

Machine Learning

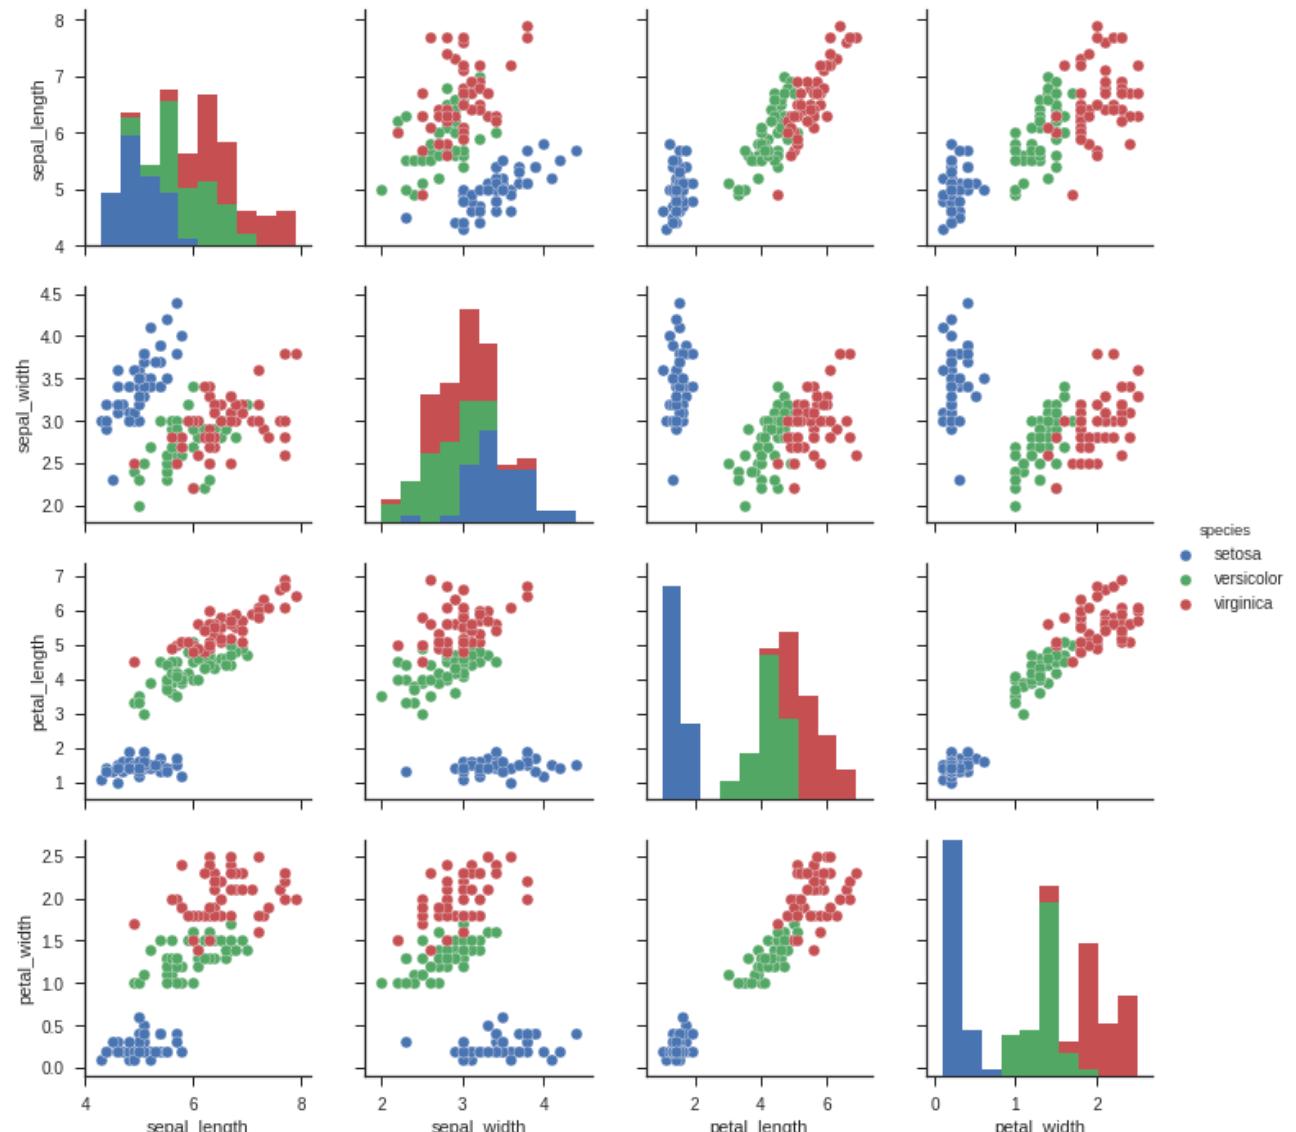
Supervised Learning (Classification)

Learning from Examples

$$y = f(x)$$



Iris Data Visualization



Python in Google Colab (Python101)

<https://colab.research.google.com/drive/1FEG6DnGvwfUbeo4zJ1zTunjMqf2RkCrT>

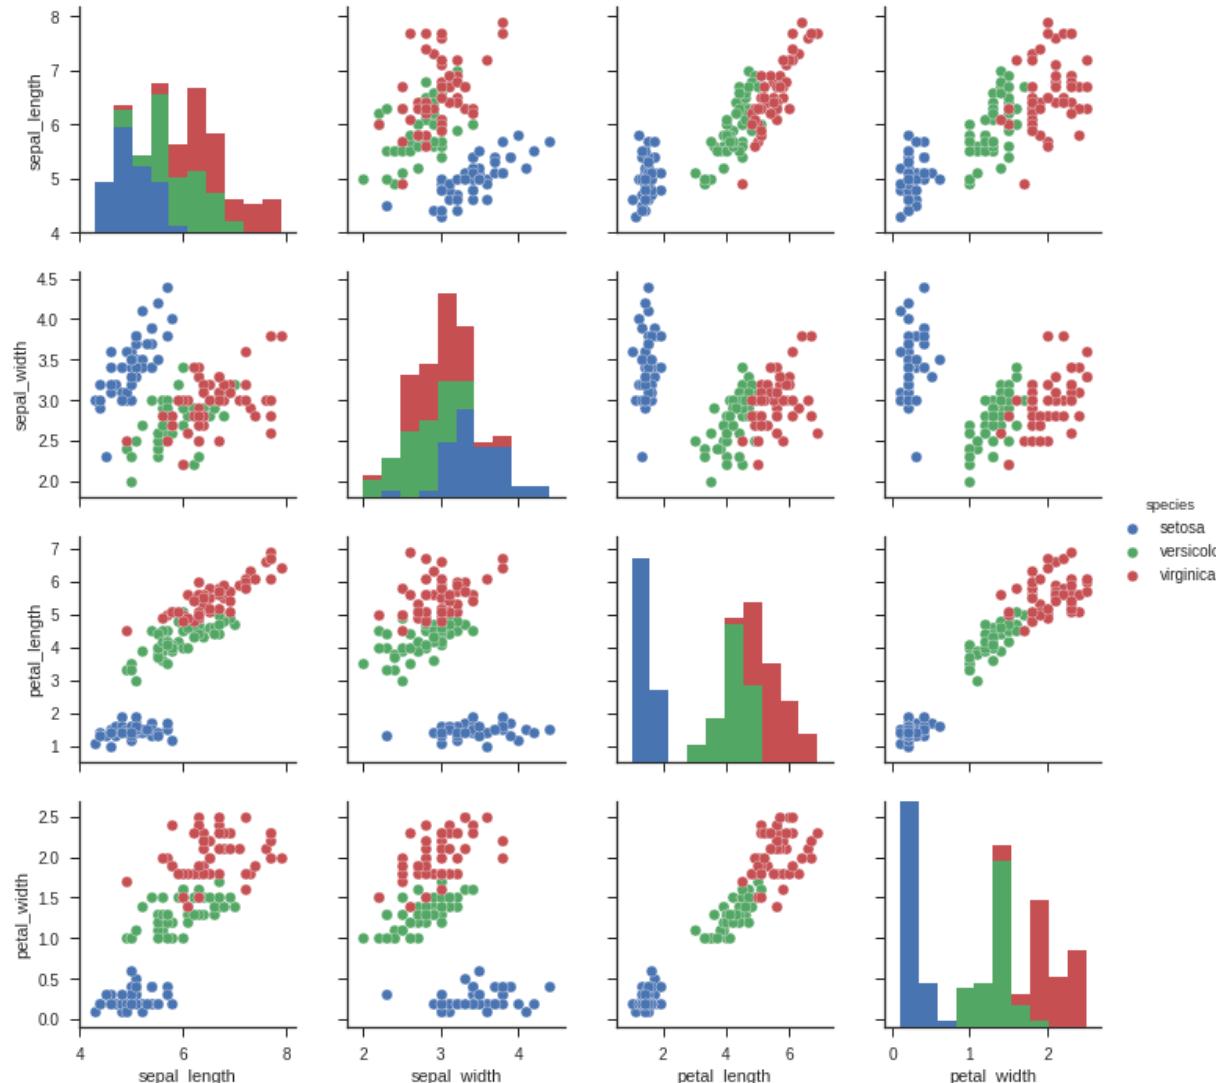
The screenshot shows a Google Colab interface with the following details:

- Title:** python101.ipynb
- File Menu:** File, Edit, View, Insert, Runtime, Tools, Help
- Toolbar:** Comment, Share, Settings, A (font size)
- Table of Contents:** Machine Learning with scikit-learn, Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python, Investment Portfolio Optimization, Text Analytics and Natural Language Processing (NLP), Python for Natural Language Processing, spaCy Chinese Model, Open Chinese Convert (OpenCC, 開放中文轉換), Jieba 結巴中文分詞, Natural Language Toolkit (NLTK), Stanza: A Python NLP Library for Many Human Languages, Text Processing and Understanding, NLTK (Natural Language Processing with Python – Analyzing Text with the
- Code Cell:** Contains Python code for importing libraries and loading a dataset.

```
1 # Import libraries
2 import numpy as np
3 import pandas as pd
4 %matplotlib inline
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from pandas.plotting import scatter_matrix
8
9 # Import sklearn
10 from sklearn import model_selection
11 from sklearn.metrics import classification_report
12 from sklearn.metrics import confusion_matrix
13 from sklearn.metrics import accuracy_score
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.tree import DecisionTreeClassifier
16 from sklearn.neighbors import KNeighborsClassifier
17 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
18 from sklearn.naive_bayes import GaussianNB
19 from sklearn.svm import SVC
20 from sklearn.neural_network import MLPClassifier
21 print("Imported")
22
23 # Load dataset
24 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
25 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
```

<https://tinyurl.com/aintpuppython101>

```
import seaborn as sns  
sns.set(style="ticks", color_codes=True)  
iris = sns.load_dataset("iris")  
g = sns.pairplot(iris, hue="species")
```



```
import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.plotting import scatter_matrix
```

```
# Import Libraries
import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.plotting import scatter_matrix
print('imported')
```

imported

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
df = pd.read_csv(url, names=names)
print(df.head(10))
```

```
# Load dataset
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
df = pd.read_csv(url, names=names)
print(df.head(10)).
```

	sepal-length	sepal-width	petal-length	petal-width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa

df.tail(10)

```
print(df.tail(10))
```

	sepal-length	sepal-width	petal-length	petal-width	class
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

df . describe ()

```
print(df.describe())
```

	sepal-length	sepal-width	petal-length	petal-width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

```
print(df.info())
print(df.shape)
```

```
print(df.info()).
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
sepal-length      150 non-null float64
sepal-width       150 non-null float64
petal-length      150 non-null float64
petal-width       150 non-null float64
class             150 non-null object
dtypes: float64(4), object(1)
memory usage: 5.9+ KB
None
```

```
print(df.shape)
```

```
(150, 5)
```

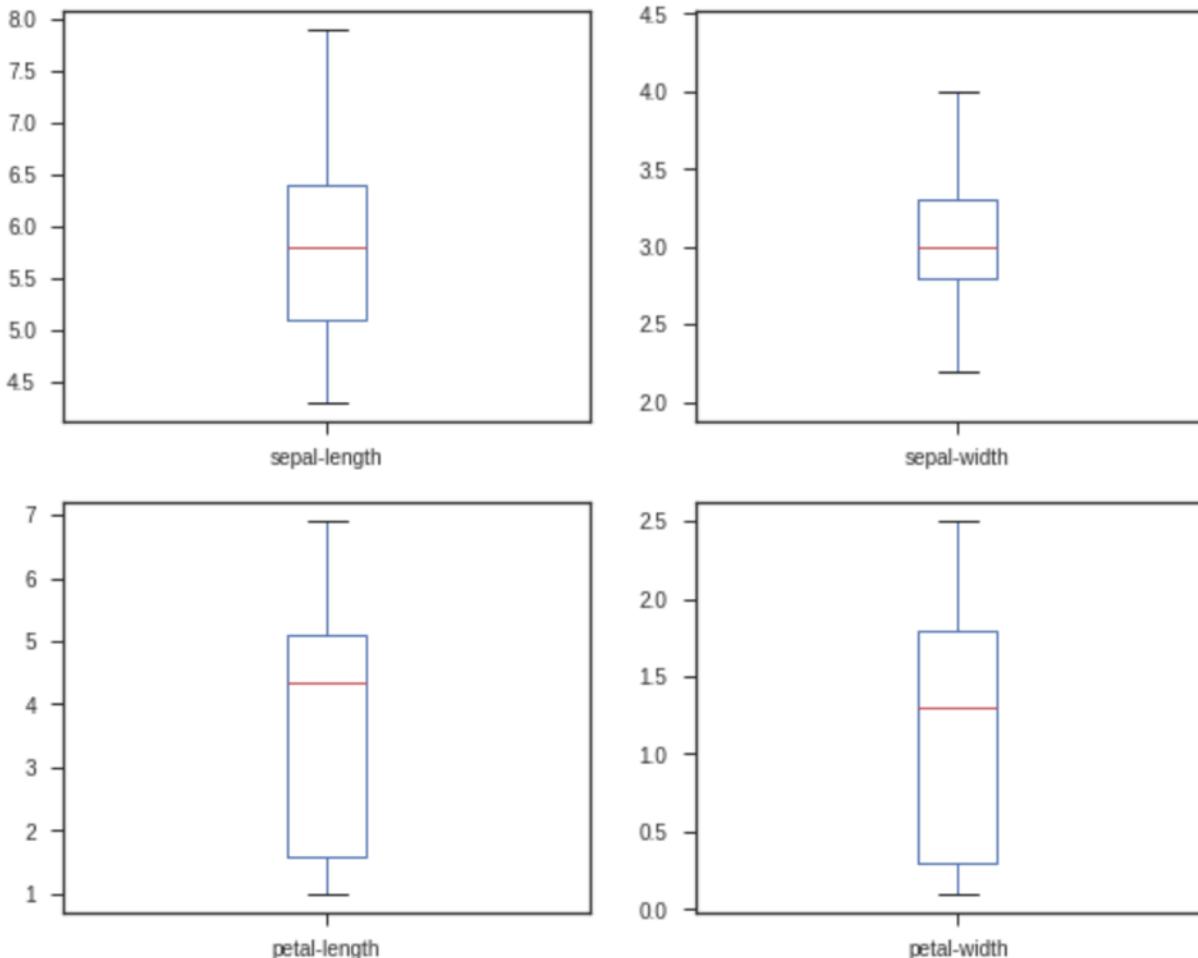
```
df.groupby('class').size()
```

```
print(df.groupby('class').size())
```

```
class
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
dtype: int64
```

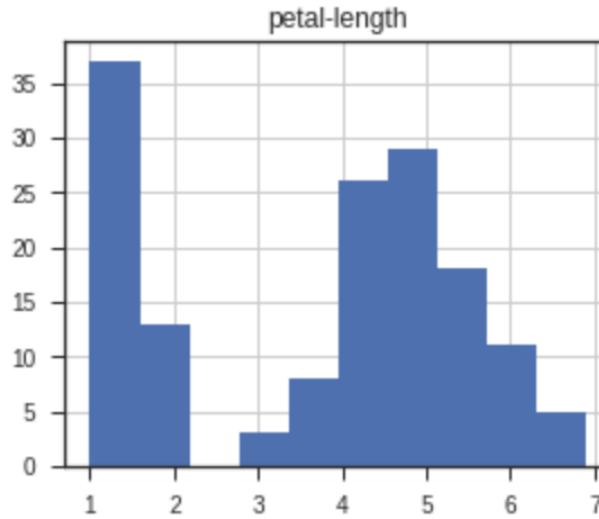
```
plt.rcParams["figure.figsize"] = (10,8)
df.plot(kind='box', subplots=True, layout=(2,2), sharex=False, sharey=False)
plt.show()
```

```
plt.rcParams["figure.figsize"] = (10,8)
df.plot(kind='box', subplots=True, layout=(2,2), sharex=False, sharey=False)
plt.show(..)
```

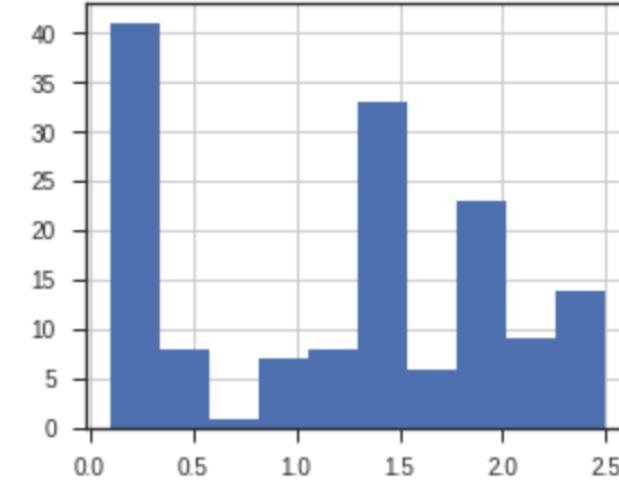


```
df.hist()  
plt.show()
```

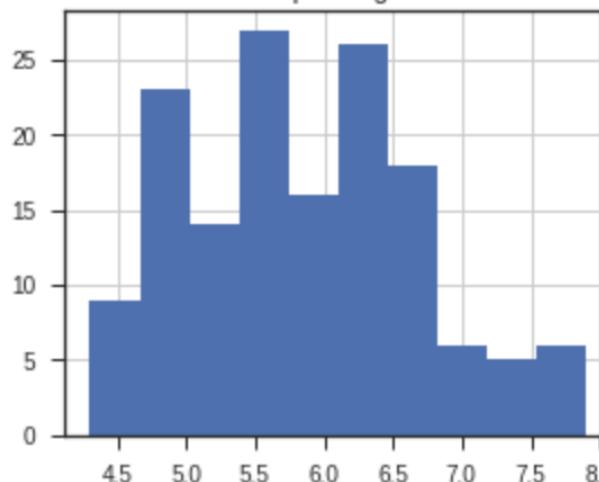
```
df.hist()  
plt.show()
```



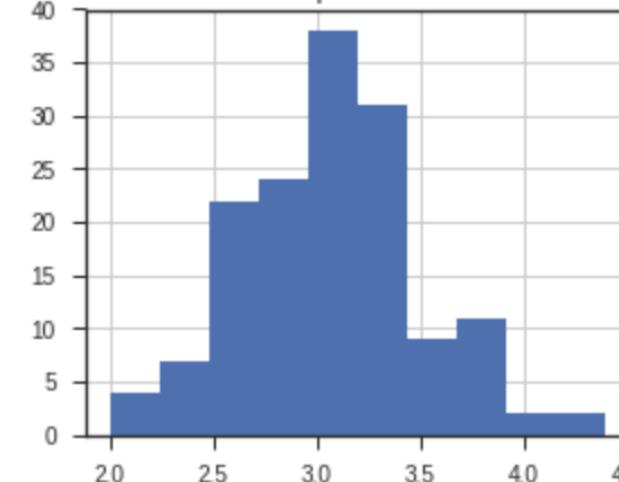
petal-width



sepal-length

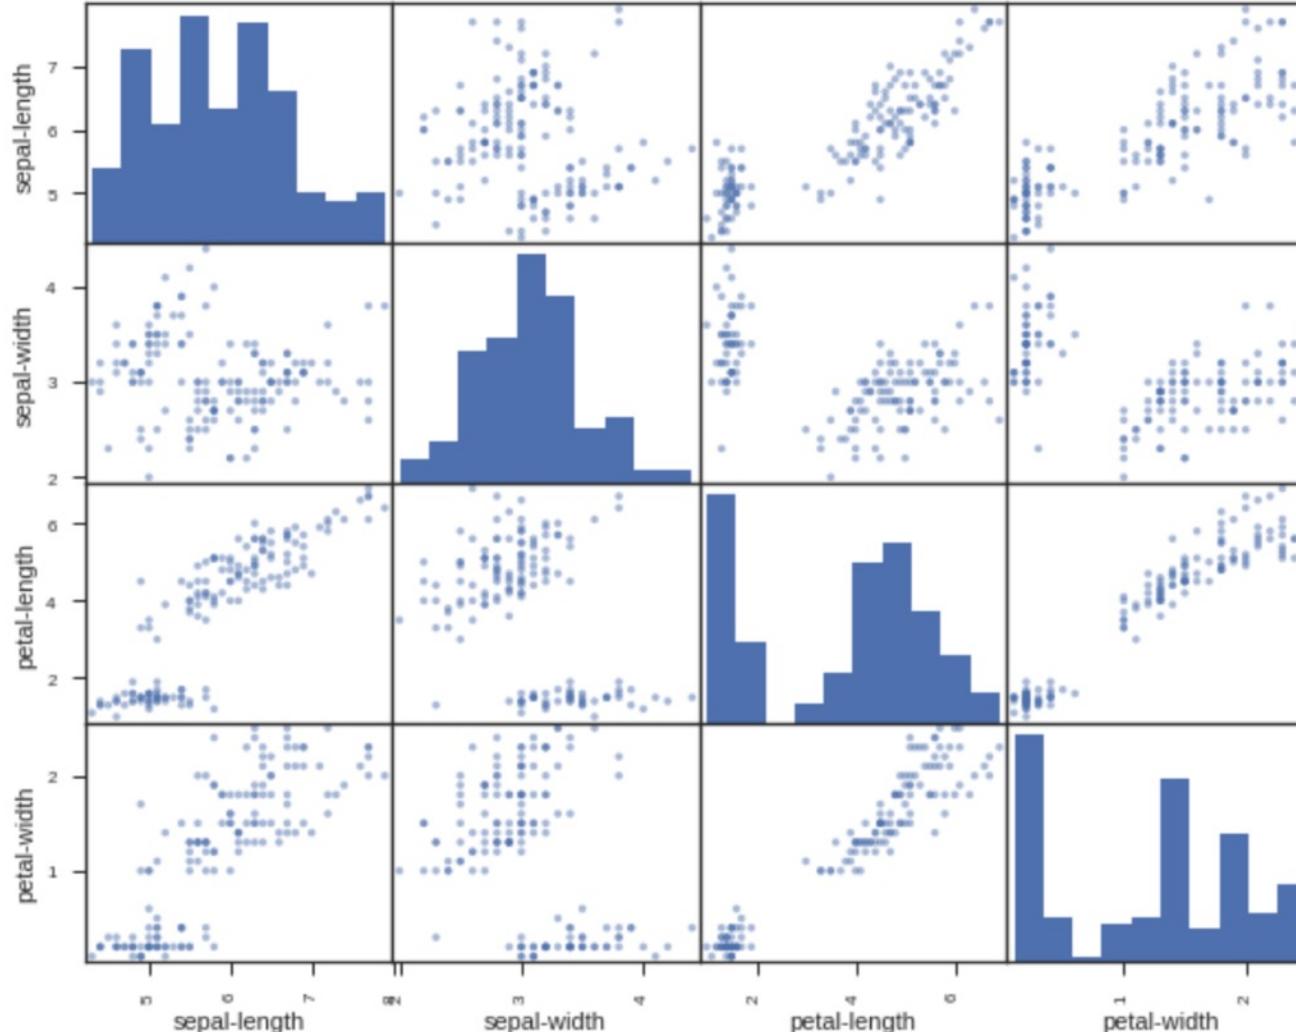


sepal-width



```
scatter_matrix(df)  
plt.show()
```

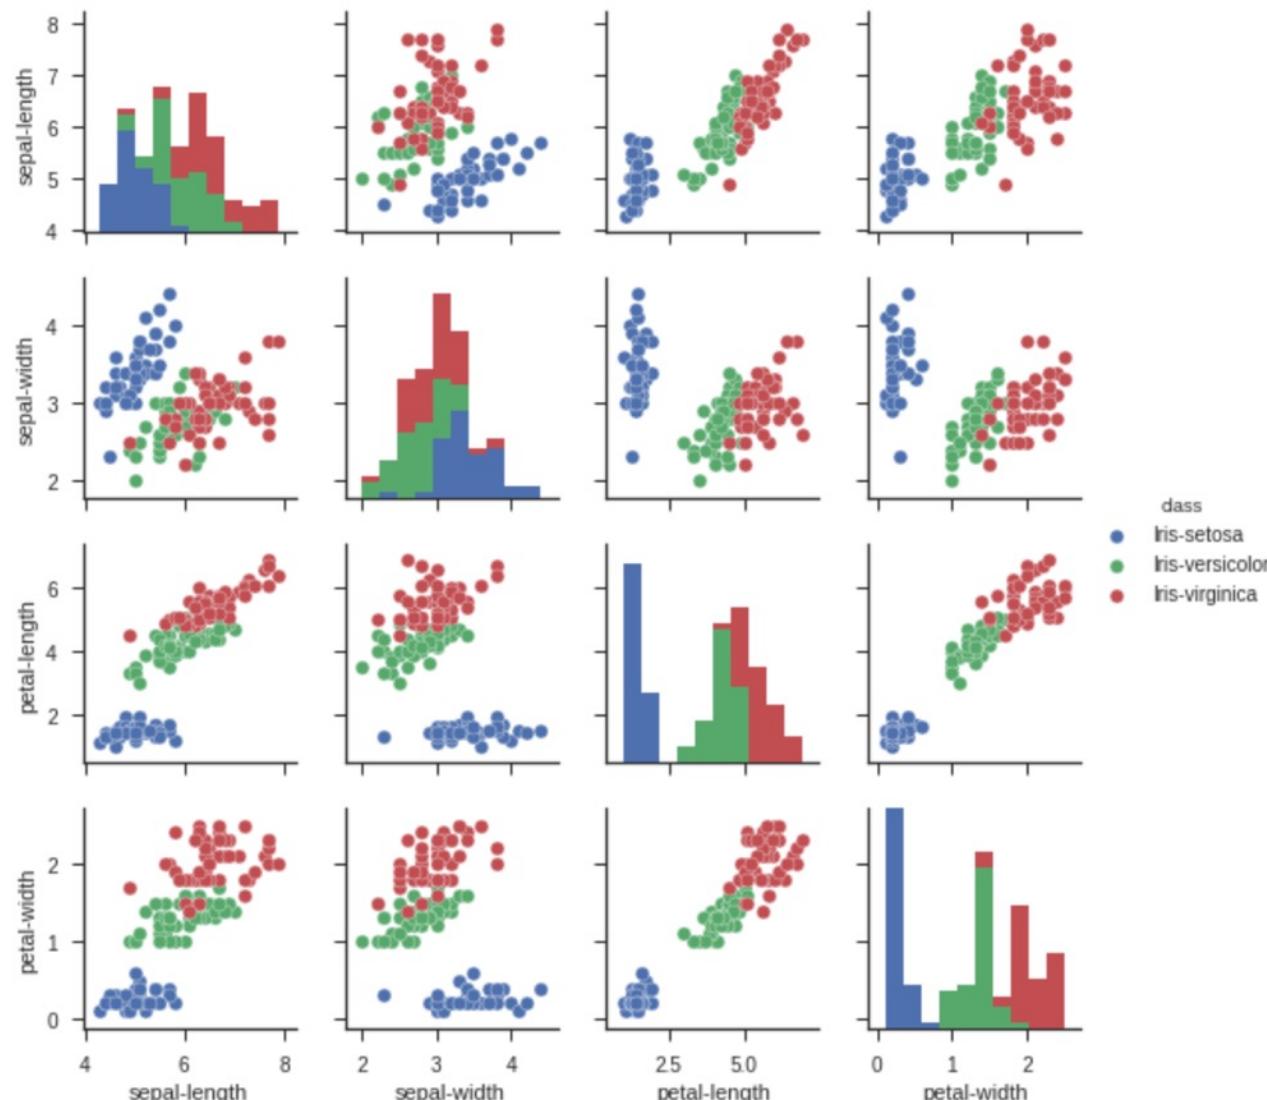
```
scatter_matrix(df)  
plt.show()
```



```
sns.pairplot(df, hue="class", size=2)
```

```
sns.pairplot(df, hue="class", size=2)
```

```
<seaborn.axisgrid.PairGrid at 0x7f1d21267390>
```

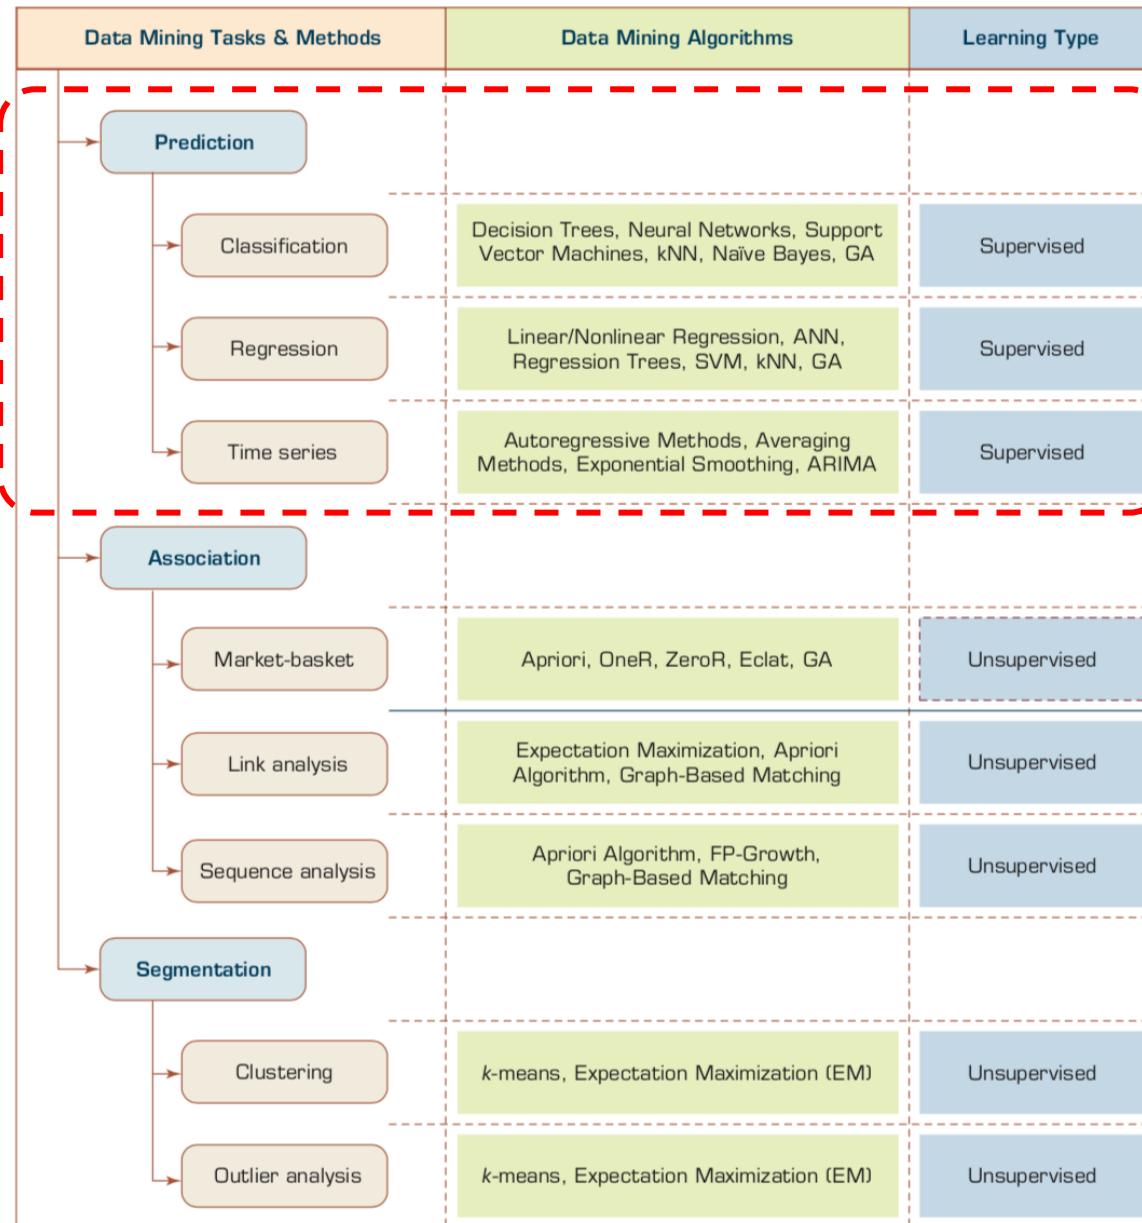


Machine Learning: Supervised Learning: Classification and Prediction

Machine Learning: Data Mining Tasks & Methods

Prediction Classification

Supervised Learning:
Classification
and
Prediction



Machine Learning: Supervised Learning Classification and Prediction

The screenshot shows a Jupyter Notebook interface with the following details:

- Title:** python101.ipynb
- File Menu:** File, Edit, View, Insert, Runtime, Tools, Help
- Toolbar:** Comment, Share, Settings, A (theme)
- Table of Contents:** Machine Learning with scikit-learn (selected), Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python, Investment Portfolio Optimization, Text Analytics and Natural Language Processing (NLP), Python for Natural Language Processing, spaCy Chinese Model, Open Chinese Convert (OpenCC, 開放中文轉換), Jieba 結巴中文分詞, Natural Language Toolkit (NLTK), Stanza: A Python NLP Library for Many Human Languages, Text Processing and Understanding, NLTK (Natural Language Processing with Python – Analyzing Text with the
- Code Cell:** Machine Learning with scikit-learn, Classification and Prediction
- Code Content:**

```
1 # Import libraries
2 import numpy as np
3 import pandas as pd
4 %matplotlib inline
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from pandas.plotting import scatter_matrix
8
9 # Import sklearn
10 from sklearn import model_selection
11 from sklearn.metrics import classification_report
12 from sklearn.metrics import confusion_matrix
13 from sklearn.metrics import accuracy_score
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.tree import DecisionTreeClassifier
16 from sklearn.neighbors import KNeighborsClassifier
17 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
18 from sklearn.naive_bayes import GaussianNB
19 from sklearn.svm import SVC
20 from sklearn.neural_network import MLPClassifier
21 print("Imported")
22
23 # Load dataset
24 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
25 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
```

```
# Import sklearn
from sklearn import model_selection
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
print("Imported")
```

```

1 # Load dataset
2 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
3 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
4 df = pd.read_csv(url, names=names)
5
6 print(df.head(10))
7 print(df.tail(10))
8 print(df.describe())
9 print(df.info())
10 print(df.shape)
11 print(df.groupby('class').size())
12
13 plt.rcParams["figure.figsize"] = (10,8)
14 df.plot(kind='box', subplots=True, layout=(2,2), sharex=False, sharey=False)
15 plt.show()
16
17 df.hist()
18 plt.show()
19
20 scatter_matrix(df)
21 plt.show()
22
23 sns.pairplot(df, hue="class", size=2).

```

	sepal-length	sepal-width	petal-length	petal-width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa
	sepal-length	sepal-width	petal-length	petal-width	class
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica

```

1 # Load dataset
2 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
3 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
4 df = pd.read_csv(url, names=names)
5
6 print(df.head(10))
7 print(df.tail(10))
8 print(df.describe())
9 print(df.info())
10 print(df.shape)
11 print(df.groupby('class').size())
12
13 plt.rcParams["figure.figsize"] = (10,8)
14 df.plot(kind='box', subplots=True, layout=(2,2), sharex=False, sharey=False)
15 plt.show()
16
17 df.hist()
18 plt.show()
19
20 scatter_matrix(df)
21 plt.show()
22
23 sns.pairplot(df, hue="class", size=2).

```

	sepal-length	sepal-width	petal-length	petal-width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa
	sepal-length	sepal-width	petal-length	petal-width	class
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica

df.corr()

```
1 | df.corr()
```

	sepal-length	sepal-width	petal-length	petal-width
sepal-length	1.000000	-0.109369	0.871754	0.817954
sepal-width	-0.109369	1.000000	-0.420516	-0.356544
petal-length	0.871754	-0.420516	1.000000	0.962757
petal-width	0.817954	-0.356544	0.962757	1.000000

```
# Split-out validation dataset
array = df.values
X = array[:,0:4]
Y = array[:,4]
validation_size = 0.20
seed = 7
X_train, X_validation, Y_train, Y_validation =
model_selection.train_test_split(X, Y,
test_size=validation_size, random_state=seed)
scoring = 'accuracy'
```

```
1 # Split-out validation dataset
2 array = df.values
3 X = array[:,0:4]
4 Y = array[:,4]
5 validation_size = 0.20
6 seed = 7
7 X_train, X_validation, Y_train, Y_validation = model_selection.train_test_split(X, Y, test_size=validation_size, random_state=seed)
8 scoring = 'accuracy'|
```

```
1 len(Y_validation)
```

```
# Models
models = []
models.append( ('LR' , LogisticRegression()))
models.append( ('LDA',
LinearDiscriminantAnalysis()))
models.append( ('KNN' , KNeighborsClassifier()))
models.append( ('DT',
DecisionTreeClassifier()))
models.append( ('NB' , GaussianNB()))
models.append( ('SVM', SVC()))
```

```
# evaluate each model in turn
results = []
names = []
for name, model in models:
    kfold = model_selection.KFold(n_splits=10,
random_state=seed)
    cv_results =
model_selection.cross_val_score(model,
X_train, Y_train, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %.4f (%.4f)" % (name,
cv_results.mean(), cv_results.std())
    print(msg)
```

```
1 # Models
2 models = []
3 models.append('LR', LogisticRegression())
4 models.append('LDA', LinearDiscriminantAnalysis())
5 models.append('KNN', KNeighborsClassifier())
6 models.append('DT', DecisionTreeClassifier())
7 models.append('NB', GaussianNB())
8 models.append('SVM', SVC())
9 # evaluate each model in turn
10 results = []
11 names = []
12 for name, model in models:
13     kfold = model_selection.KFold(n_splits=10, random_state=seed)
14     cv_results = model_selection.cross_val_score(model, X_train, Y_train, cv=kfold, scoring=scoring)
15     results.append(cv_results)
16     names.append(name)
17     msg = "%s: %.4f (%.4f)" % (name, cv_results.mean(), cv_results.std())
18     print(msg)
```

LR: 0.9667 (0.0408)
LDA: 0.9750 (0.0382)
KNN: 0.9833 (0.0333)
DT: 0.9750 (0.0382)
NB: 0.9750 (0.0534)
SVM: 0.9917 (0.0250)

```
# Make predictions on validation dataset
model = KNeighborsClassifier()
model.fit(X_train, Y_train)
predictions = model.predict(X_validation)
print("%.4f" % accuracy_score(Y_validation,
predictions))
print(confusion_matrix(Y_validation,
predictions))
print(classification_report(Y_validation,
predictions))
print(model)
```

```
1 # Make predictions on validation dataset
2 model = KNeighborsClassifier()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("%.4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)
```

0.9000

```
[[ 7  0  0]
 [ 0 11  1]
 [ 0  2  9]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Iris-setosa	1.00	1.00	1.00	7
-------------	------	------	------	---

Iris-versicolor	0.85	0.92	0.88	12
-----------------	------	------	------	----

Iris-virginica	0.90	0.82	0.86	11
----------------	------	------	------	----

avg / total	0.90	0.90	0.90	30
-------------	------	------	------	----

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                      metric_params=None, n_jobs=1, n_neighbors=5, p=2,
                      weights='uniform')
```

```
# Make predictions on validation dataset
model = SVC()
model.fit(X_train, Y_train)
predictions = model.predict(X_validation)
print("%.4f" % accuracy_score(Y_validation,
predictions))
print(confusion_matrix(Y_validation,
predictions))
print(classification_report(Y_validation,
predictions))
print(model)
```

```
model = SVC()
model.fit(X_train, Y_train)
predictions = model.predict(X_validation)
```

```
1 # Make predictions on validation dataset
2 model = SVC()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("% .4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)
```

0.9333

```
[[ 7  0  0]
 [ 0 10  2]
 [ 0  0 11]]
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	7
Iris-versicolor	1.00	0.83	0.91	12
Iris-virginica	0.85	1.00	0.92	11
avg / total	0.94	0.93	0.93	30

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```

1 # Make predictions on validation dataset
2 model = DecisionTreeClassifier()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("%.4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)

```

0.9000

[7	0	0
[0	11	1
[0	2	9

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	7
Iris-versicolor	0.85	0.92	0.88	12
Iris-virginica	0.90	0.82	0.86	11
avg / total	0.90	0.90	0.90	30

```

DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                      max_features=None, max_leaf_nodes=None,
                      min_impurity_decrease=0.0, min_impurity_split=None,
                      min_samples_leaf=1, min_samples_split=2,
                      min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                      splitter='best')

```

```

1 # Make predictions on validation dataset
2 model = GaussianNB()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("% .4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)

```

0.8333

`[[7 0 0]
 [0 9 3]
 [0 2 9]]`

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	7
Iris-versicolor	0.82	0.75	0.78	12
Iris-virginica	0.75	0.82	0.78	11
avg / total	0.84	0.83	0.83	30

`GaussianNB(priors=None)`

```
1 # Make predictions on validation dataset
2 model = LogisticRegression()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("% .4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)
```

0.8000

```
[[ 7  0  0]
 [ 0  7  5]
 [ 0  1 10]]
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	7
Iris-versicolor	0.88	0.58	0.70	12
Iris-virginica	0.67	0.91	0.77	11
avg / total	0.83	0.80	0.80	30

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                   intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                   penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                   verbose=0, warm_start=False)
```

```
1 # Make predictions on validation dataset
2 model = LinearDiscriminantAnalysis()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("%.4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)
```

0.9667

```
[[ 7  0  0]
 [ 0 11  1]
 [ 0  0 11]]
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	7
Iris-versicolor	1.00	0.92	0.96	12
Iris-virginica	0.92	1.00	0.96	11
avg / total	0.97	0.97	0.97	30

```
LinearDiscriminantAnalysis(n_components=None, priors=None, shrinkage=None,
                           solver='svd', store_covariance=False, tol=0.0001)
```

```
1 # Make predictions on validation dataset
2 model = MLPClassifier()
3 model.fit(X_train, Y_train)
4 predictions = model.predict(X_validation)
5 print("% .4f" % accuracy_score(Y_validation, predictions))
6 print(confusion_matrix(Y_validation, predictions))
7 print(classification_report(Y_validation, predictions))
8 print(model)
```

0.9000

```
[[ 7  0  0]
 [ 0  9  3]
 [ 0  0 11]]
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	7
Iris-versicolor	1.00	0.75	0.86	12
Iris-virginica	0.79	1.00	0.88	11
avg / total	0.92	0.90	0.90	30

```
MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
              beta_2=0.999, early_stopping=False, epsilon=1e-08,
              hidden_layer_sizes=(100,), learning_rate='constant',
              learning_rate_init=0.001, max_iter=200, momentum=0.9,
              nesterovs_momentum=True, power_t=0.5, random_state=None,
              shuffle=True, solver='adam', tol=0.0001, validation_fraction=0.1,
              verbose=False, warm_start=False)
```

Evaluation (Accuracy of Classification Model)

Assessing the Classification Model

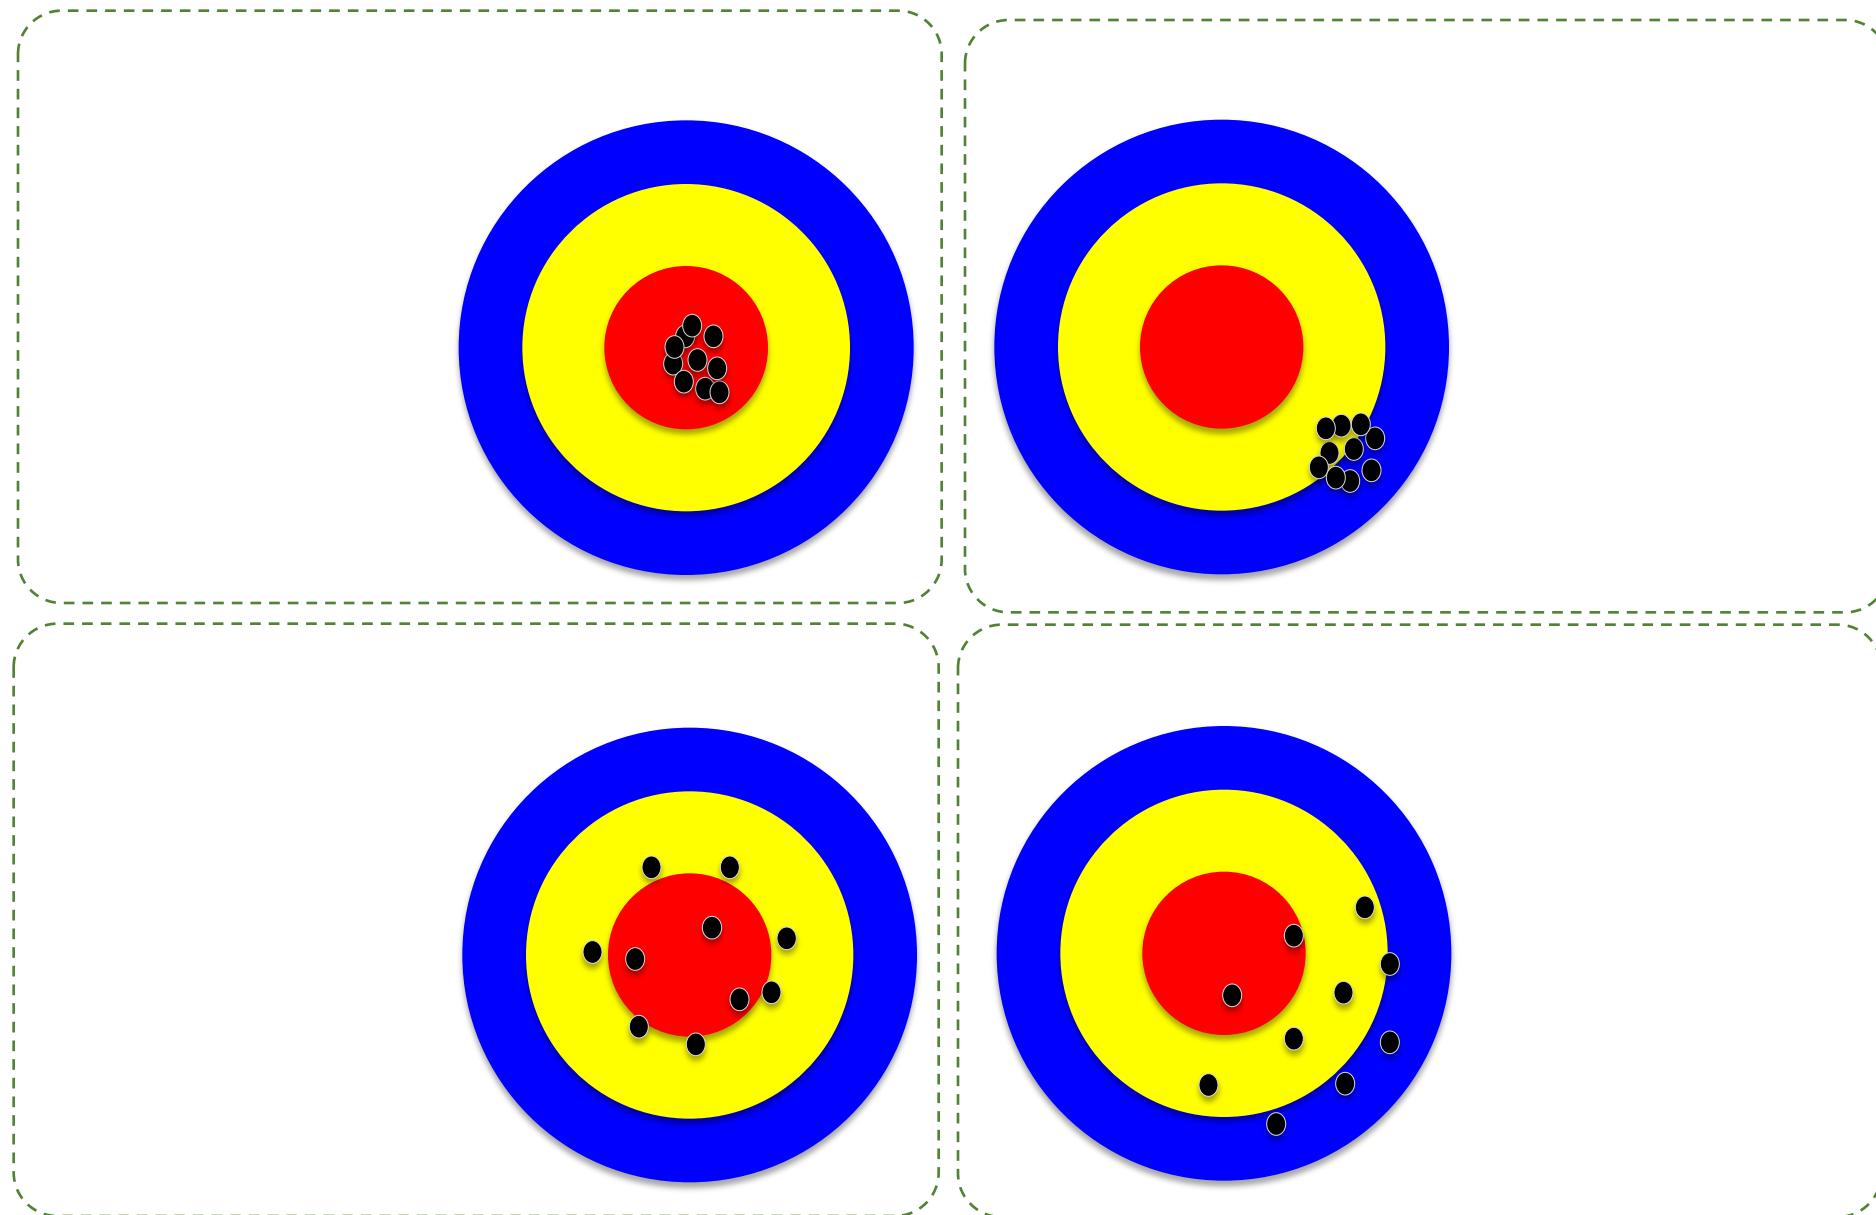
- Predictive accuracy
 - Hit rate
- Speed
 - Model building; predicting
- Robustness
- Scalability
- Interpretability
 - Transparency, explainability

Accuracy

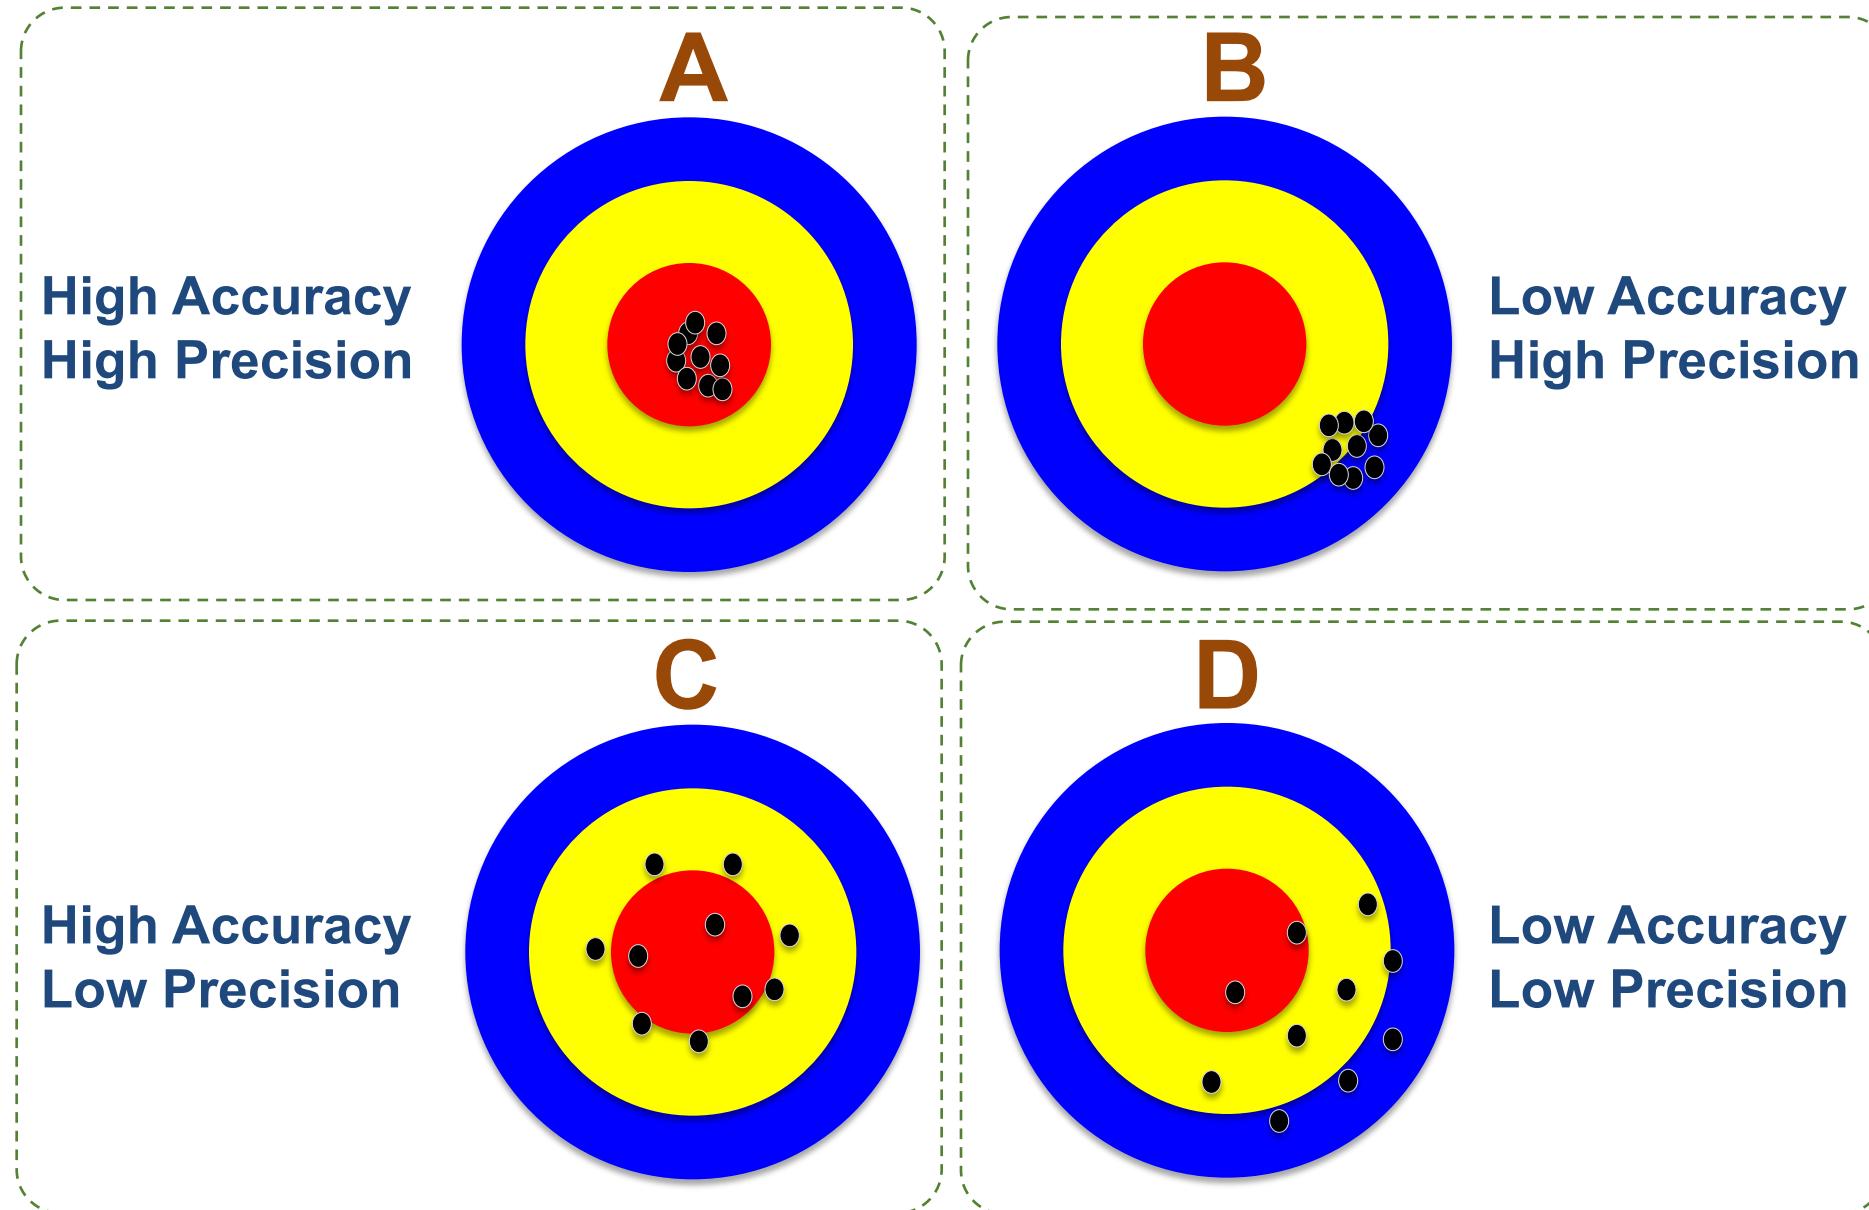
Validity

Precision

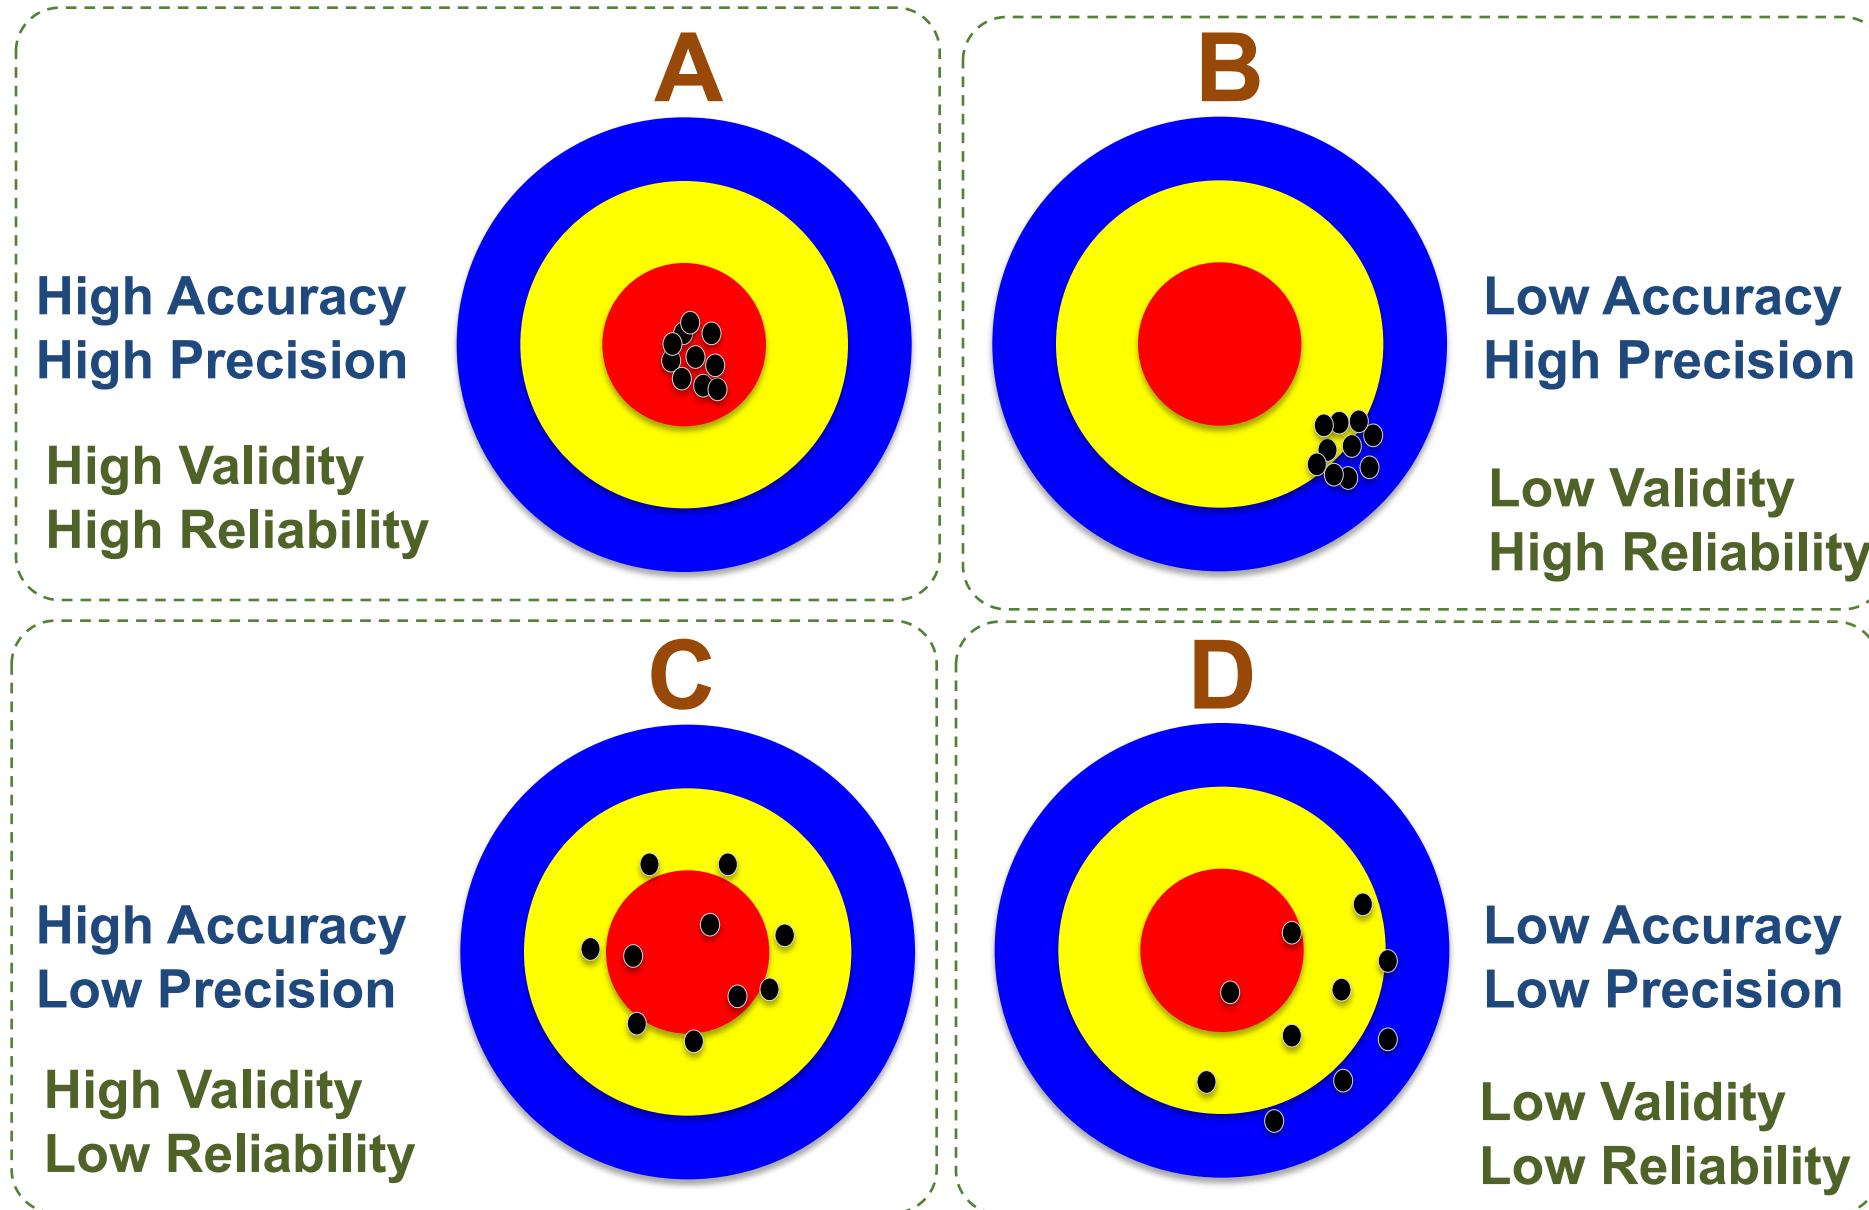
Reliability



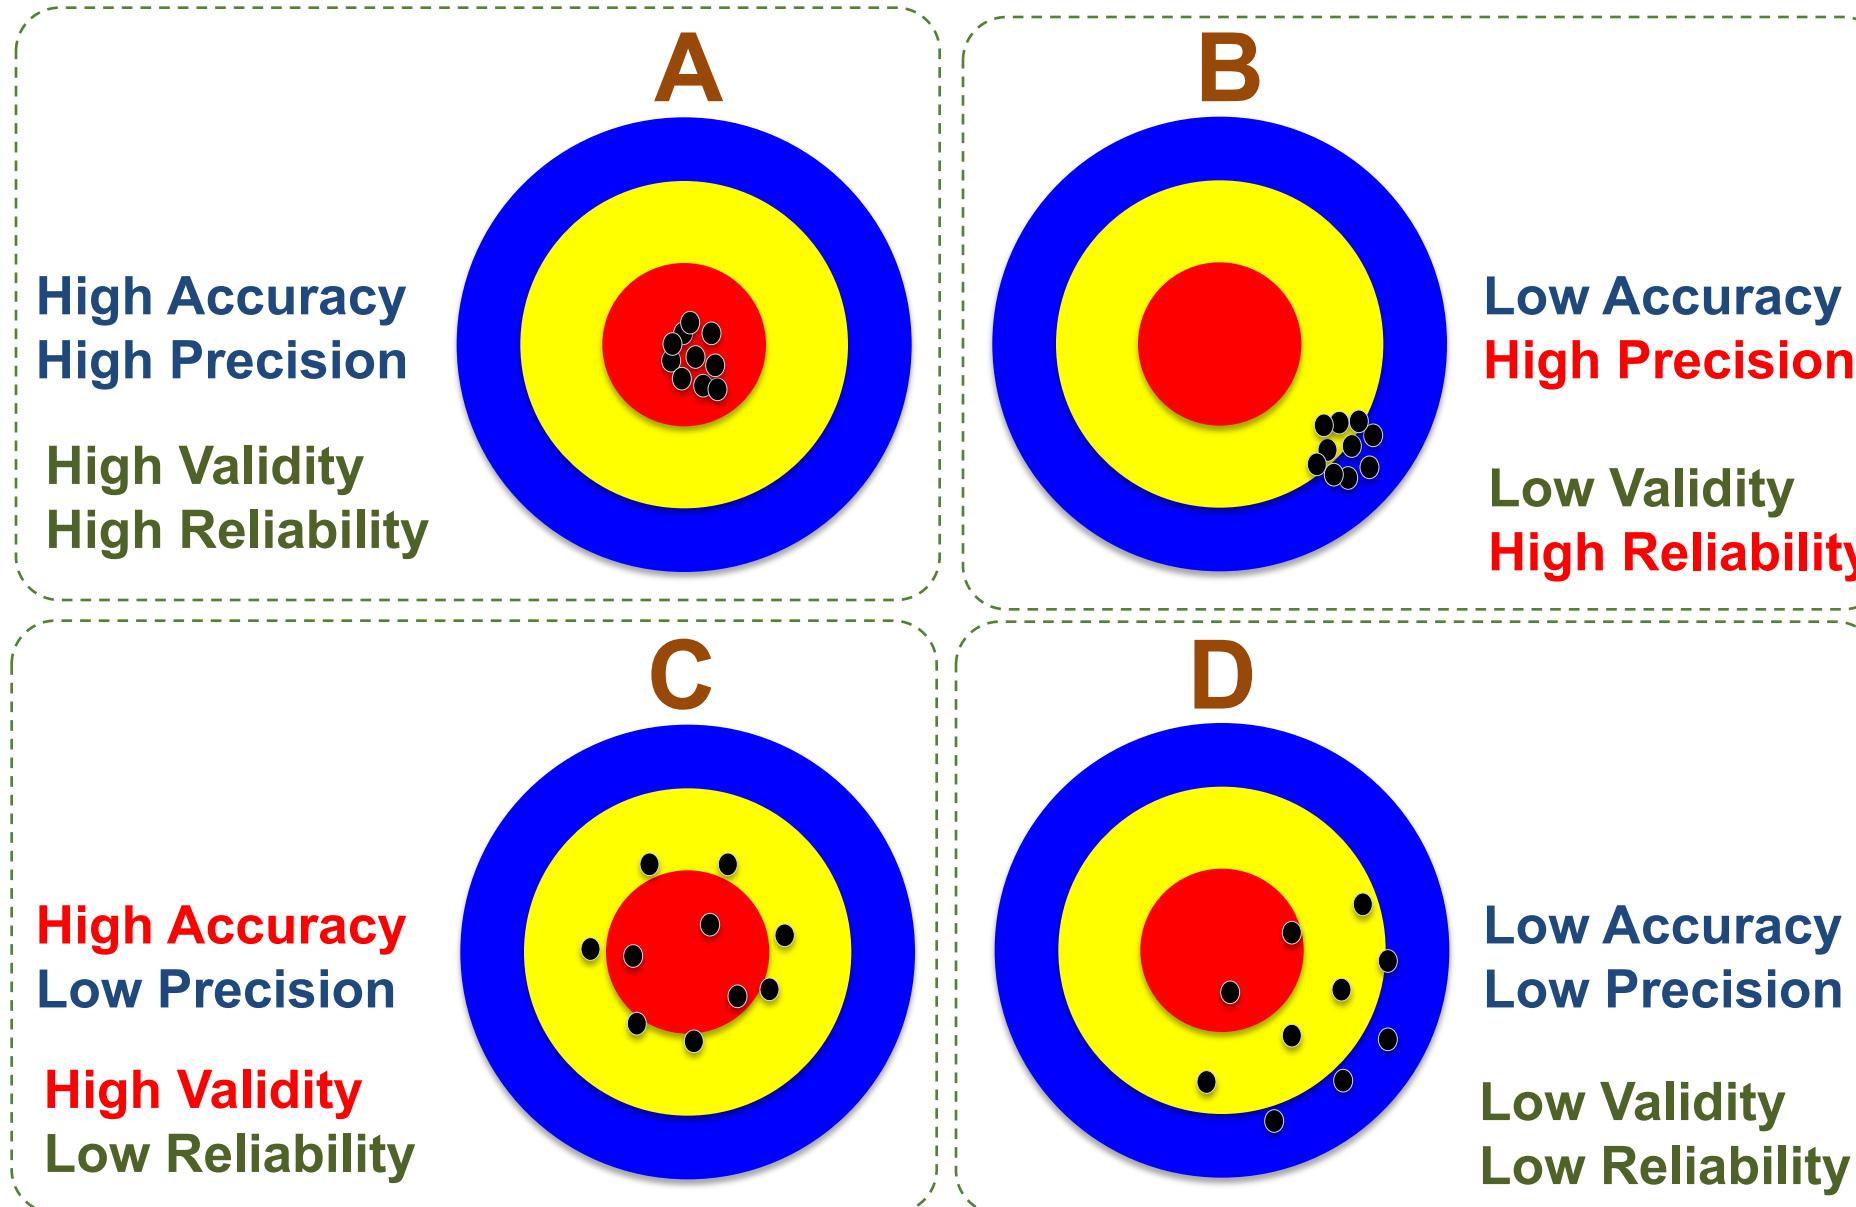
Accuracy vs. Precision



Accuracy vs. Precision



Accuracy vs. Precision



Confusion Matrix

for Tabulation of Two-Class Classification Results

		True/Observed Class	
		Positive	Negative
Predicted Class	Positive	True Positive Count (TP)	False Positive Count (FP)
	Negative	False Negative Count (FN)	True Negative Count (TN)

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{True Positive Rate} = \frac{TP}{TP + FN}$$

$$\text{True Negative Rate} = \frac{TN}{TN + FP}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

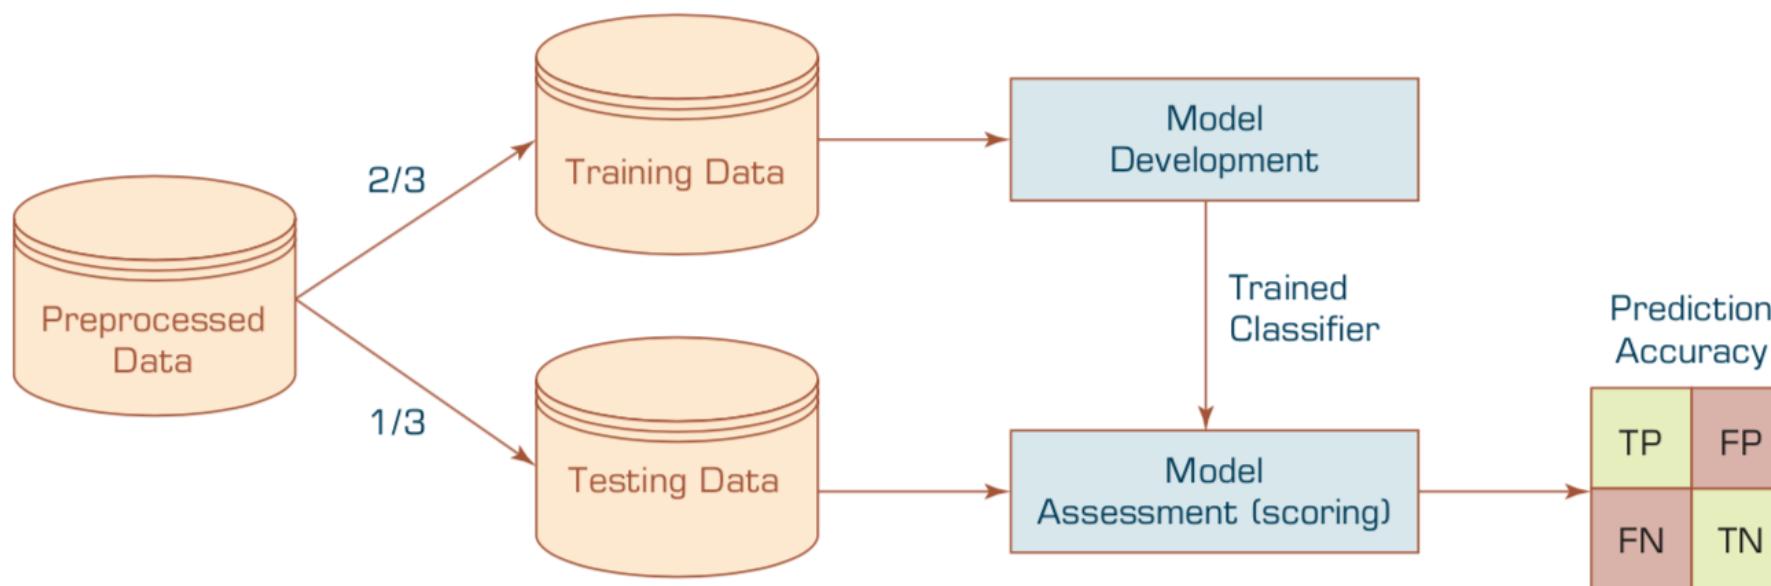
$$\text{Recall} = \frac{TP}{TP + FN}$$

Sensitivity =True Positive Rate

Specificity =True Negative Rate

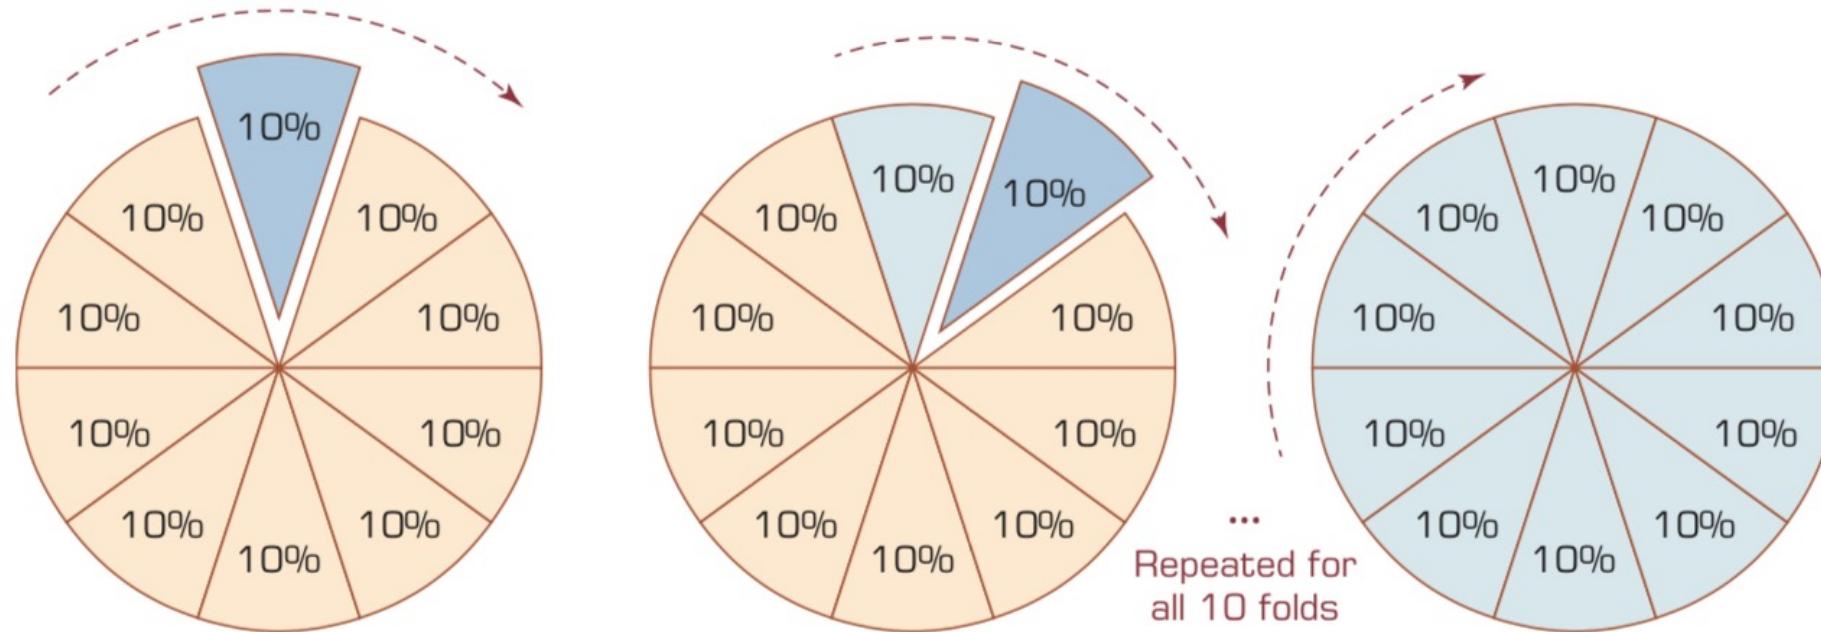
Estimation Methodologies for Classification

- **Simple split (or holdout or test sample estimation)**
 - Split the data into 2 mutually exclusive sets training (~70%) and testing (30%)



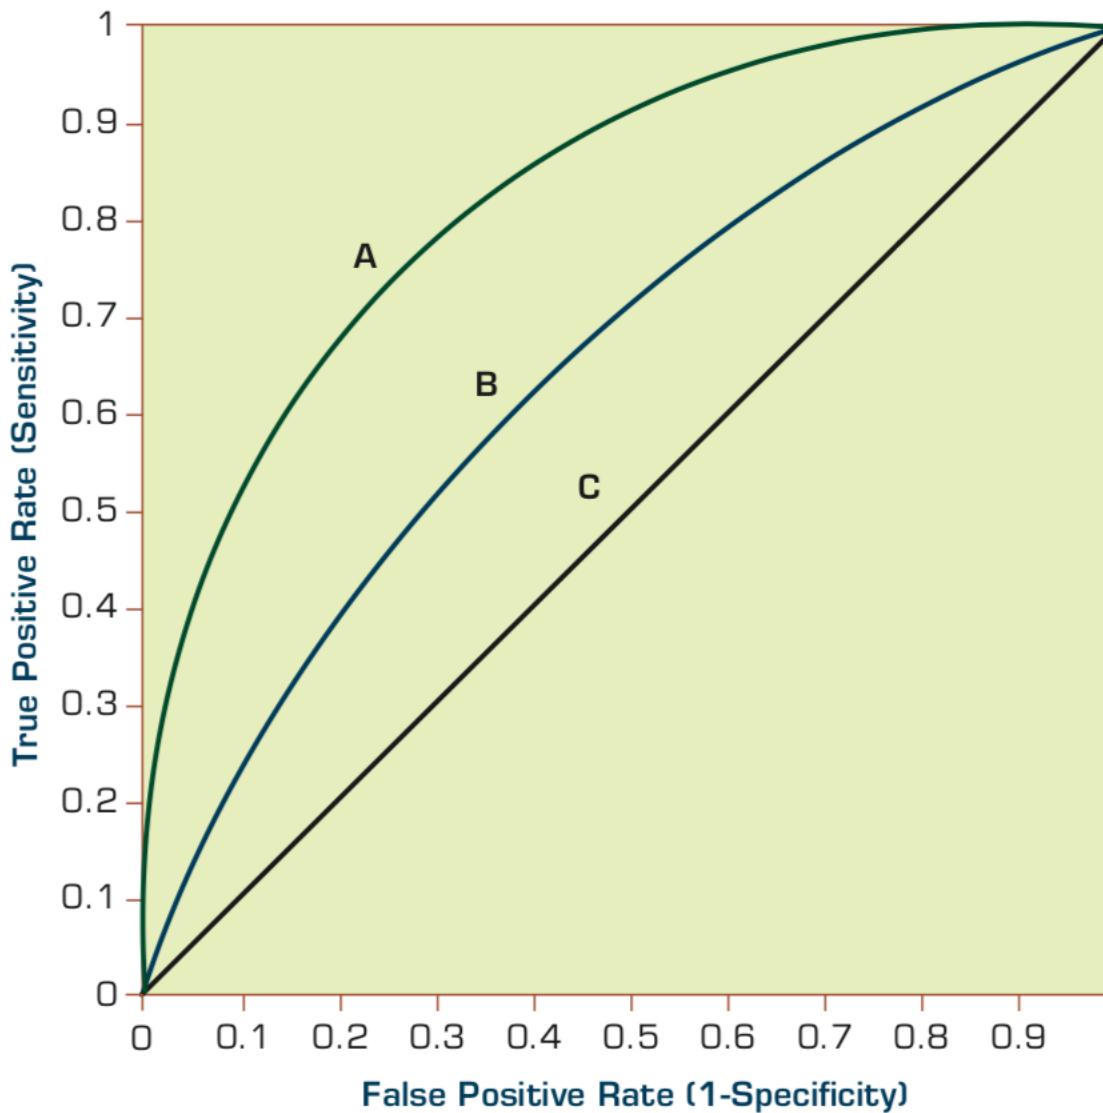
- For ANN, the data is split into three sub-sets (training [~60%], validation [~20%], testing [~20%])

k-Fold Cross-Validation



Estimation Methodologies for Classification

Area under the ROC curve



		True Class (actual value)		
		Positive	Negative	total
Predictive Class (prediction outcome)	Positive	True Positive (TP)	False Positive (FP)	P'
	Negative	False Negative (FN)	True Negative (TN)	N'
total		P	N	

$$\text{True Positive Rate (Sensitivity)} = \frac{TP}{TP + FN}$$

$$\text{True Negative Rate (Specificity)} = \frac{TN}{TN + FP}$$

$$\text{False Positive Rate} = \frac{FP}{FP + TN}$$

$$\text{False Positive Rate (1-Specificity)} = \frac{FP}{FP + TN}$$

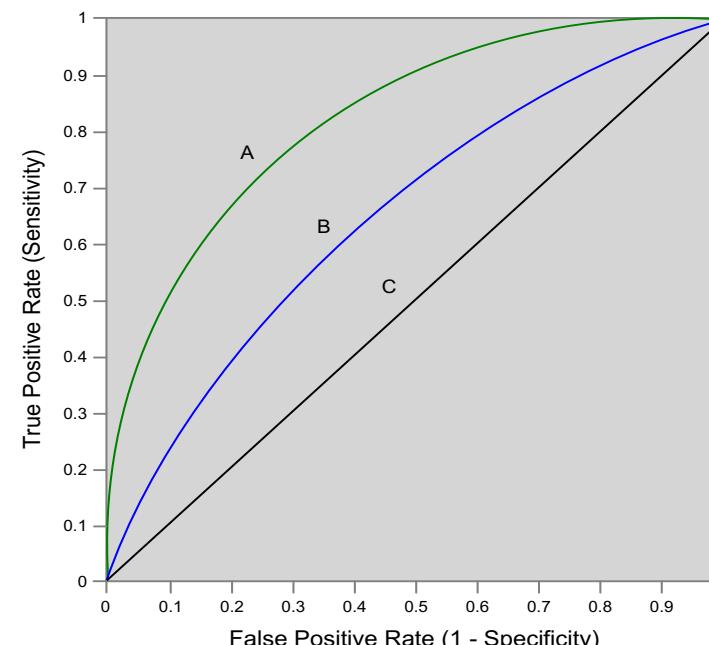
$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

$$\text{True Positive Rate} = \frac{TP}{TP + FN}$$

$$\text{True Negative Rate} = \frac{TN}{TN + FP}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$



Source: http://en.wikipedia.org/wiki/Receiver_operating_characteristic

		True Class (actual value)		total
		Positive	Negative	
Predictive Class (prediction outcome)	Positive	True Positive (TP)	False Positive (FP)	P'
	Negative	False Negative (FN)	True Negative (TN)	N'
total		P	N	

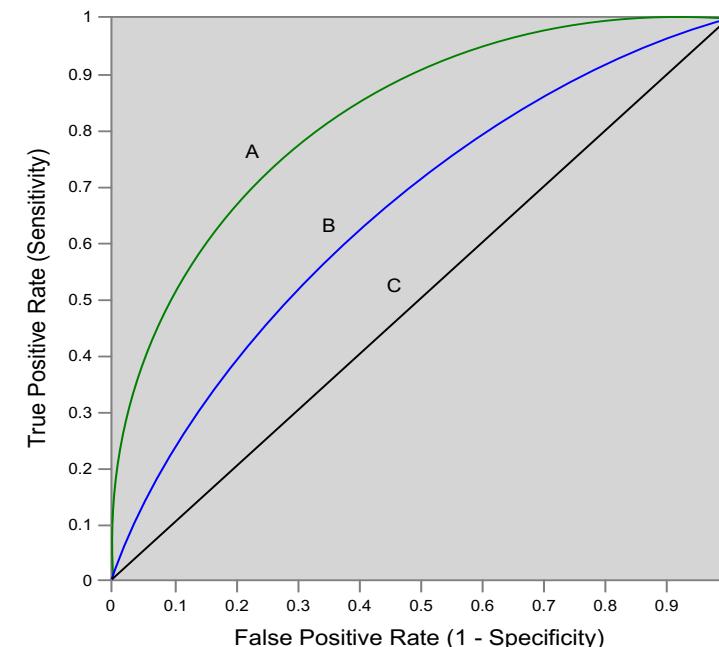
$$\text{True Positive Rate (Sensitivity)} = \frac{TP}{TP + FN}$$

Sensitivity

- = True Positive Rate
- = Recall
- = Hit rate
- = $TP / (TP + FN)$

$$\text{True Positive Rate} = \frac{TP}{TP + FN}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$



Source: http://en.wikipedia.org/wiki/Receiver_operating_characteristic

		True Class (actual value)		total
		Positive	Negative	
Predictive Class (prediction outcome)	Positive	True Positive (TP)	False Positive (FP)	P'
	Negative	False Negative (FN)	True Negative (TN)	N'

Specificity

= True Negative Rate

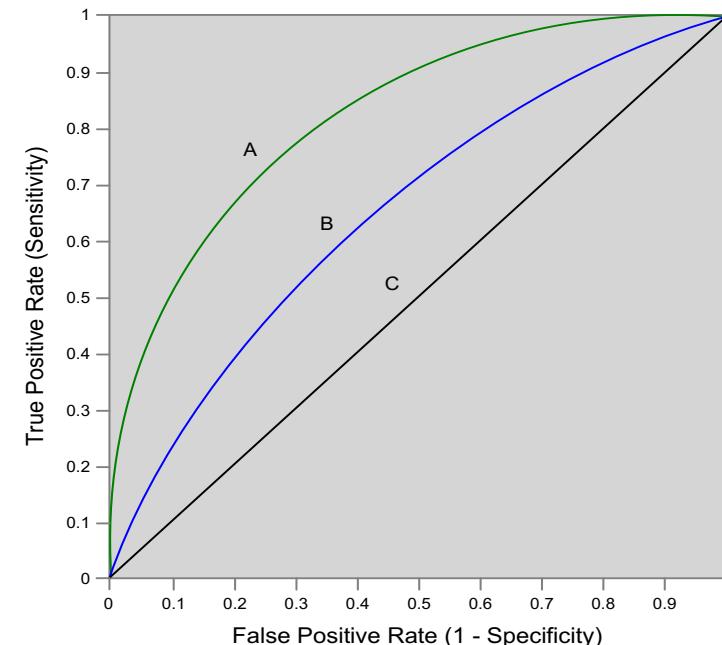
$$= TN / N$$

$$= TN / (TN + FP)$$

$$True\ Negative\ Rate\ (Specificity) = \frac{TN}{TN + FP}$$

$$False\ Positive\ Rate\ (1 - Specificity) = \frac{FP}{FP + TN}$$

$$True\ Negative\ Rate = \frac{TN}{TN + FP}$$



		True Class (actual value)		total
		Positive	Negative	
Predictive Class (prediction outcome)	Positive	True Positive (TP)	False Positive (FP)	P'
	Negative	False Negative (FN)	True Negative (TN)	N'
total		P	N	

Precision
= Positive Predictive Value (PPV)

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall
= True Positive Rate (TPR)
= Sensitivity
= Hit Rate

$$\text{Recall} = \frac{TP}{TP + FN}$$

F1 score (F-score)(F-measure)
is the harmonic mean of
precision and recall
= $2TP / (P + P')$
= $2TP / (2TP + FP + FN)$

$$F = 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

	A		
	63 (TP)	28 (FP)	91
	37 (FN)	72 (TN)	109
	100	100	200

$$\text{TPR} = 0.63$$

$$\text{FPR} = 0.28$$

$$\text{PPV} = 0.69$$

$$= 63 / (63 + 28)$$

$$= 63 / 91$$

$$\text{F1} = 0.66$$

$$= 2 * (0.63 * 0.69) / (0.63 + 0.69)$$

$$= (2 * 63) / (100 + 91)$$

$$= (0.63 + 0.69) / 2 = 1.32 / 2 = 0.66$$

$$\text{ACC} = 0.68$$

$$= (63 + 72) / 200$$

$$= 135 / 200 = 67.5$$

Recall

= True Positive Rate (TPR)

= Sensitivity

= Hit Rate

= $\text{TP} / (\text{TP} + \text{FN})$

Specificity

= True Negative Rate

= TN / N

= $\text{TN} / (\text{TN} + \text{FP})$

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{True Negative Rate (Specificity)} = \frac{\text{TN}}{\text{TN} + \text{FP}}$$

$$\text{False Positive Rate (1-Specificity)} = \frac{\text{FP}}{\text{FP} + \text{TN}}$$

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Precision

= Positive Predictive Value (PPV)

$$F = 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

F1 score (F-score)

(F-measure)

is the harmonic mean of precision and recall

$$= 2\text{TP} / (\text{P} + \text{P}')$$

$$= 2\text{TP} / (2\text{TP} + \text{FP} + \text{FN})$$

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

	A	
63 (TP)	28 (FP)	91
37 (FN)	72 (TN)	109
100	100	200
TPR = 0.63		

$$\text{FPR} = 0.28$$

$$\begin{aligned}\text{PPV} &= 0.69 \\ &= 63/(63+28) \\ &= 63/91\end{aligned}$$

$$\begin{aligned}\text{F1} &= 0.66 \\ &= 2*(0.63*0.69)/(0.63+0.69) \\ &= (2 * 63) / (100 + 91) \\ &= (0.63 + 0.69) / 2 = 1.32 / 2 = 0.66\end{aligned}$$

$$\begin{aligned}\text{ACC} &= 0.68 \\ &= (63 + 72) / 200 \\ &= 135/200 = 67.5\end{aligned}$$

	B	
77 (TP)	77 (FP)	154
23 (FN)	23 (TN)	46
100	100	200
TPR = 0.77		

$$\text{FPR} = 0.77$$

$$\text{PPV} = 0.50$$

$$\text{F1} = 0.61$$

$$\text{ACC} = 0.50$$

Recall

- = True Positive Rate (TPR)
- = Sensitivity
- = Hit Rate

$$\text{Recall} = \frac{TP}{TP + FN}$$

Precision

- = Positive Predictive Value (PPV)

$$\text{Precision} = \frac{TP}{TP + FP}$$

C		
TP	FP	112
FN	TN	88
24	88	
(TP)	(FP)	
76	12	
(FN)	(TN)	
100	100	200

$$\text{TPR} = 0.24$$

$$\text{FPR} = 0.88$$

$$\text{PPV} = 0.21$$

$$\text{F1} = 0.22$$

$$\text{ACC} = 0.18$$

C'		
TP	FP	88
FN	TN	112
76	12	
(TP)	(FP)	
24	88	
(FN)	(TN)	
100	100	200

$$\text{TPR} = 0.76$$

$$\text{FPR} = 0.12$$

$$\text{PPV} = 0.86$$

$$\text{F1} = 0.81$$

$$\text{ACC} = 0.82$$

Recall

= True Positive Rate (TPR)

= Sensitivity

= Hit Rate

$$\text{Recall} = \frac{TP}{TP + FN}$$

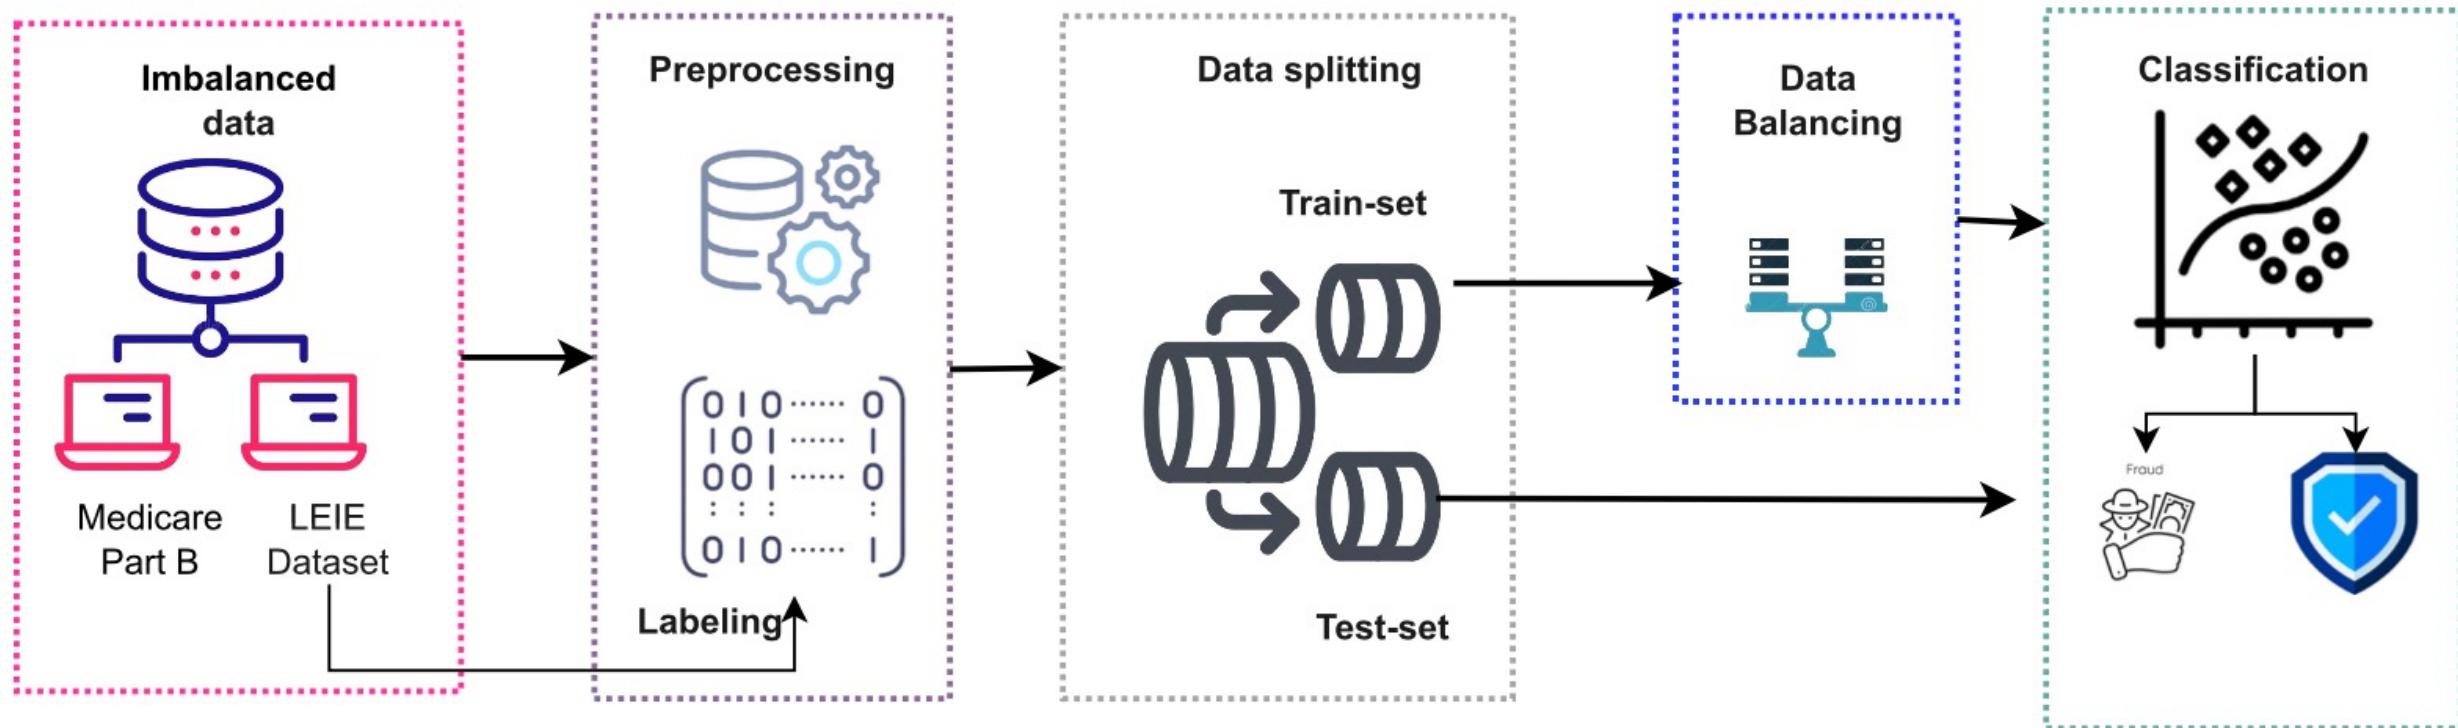
Precision

= Positive Predictive Value (PPV)

$$\text{Precision} = \frac{TP}{TP + FP}$$

Enhancing Medicare Fraud Detection Through Machine Learning: Addressing Class Imbalance With SMOTE-ENN

Architecture for healthcare fraud detection based on SMOTE-ENN
Synthetic Minority Over-sampling technique with Edited Nearest Neighbors (SMOTE-ENN)



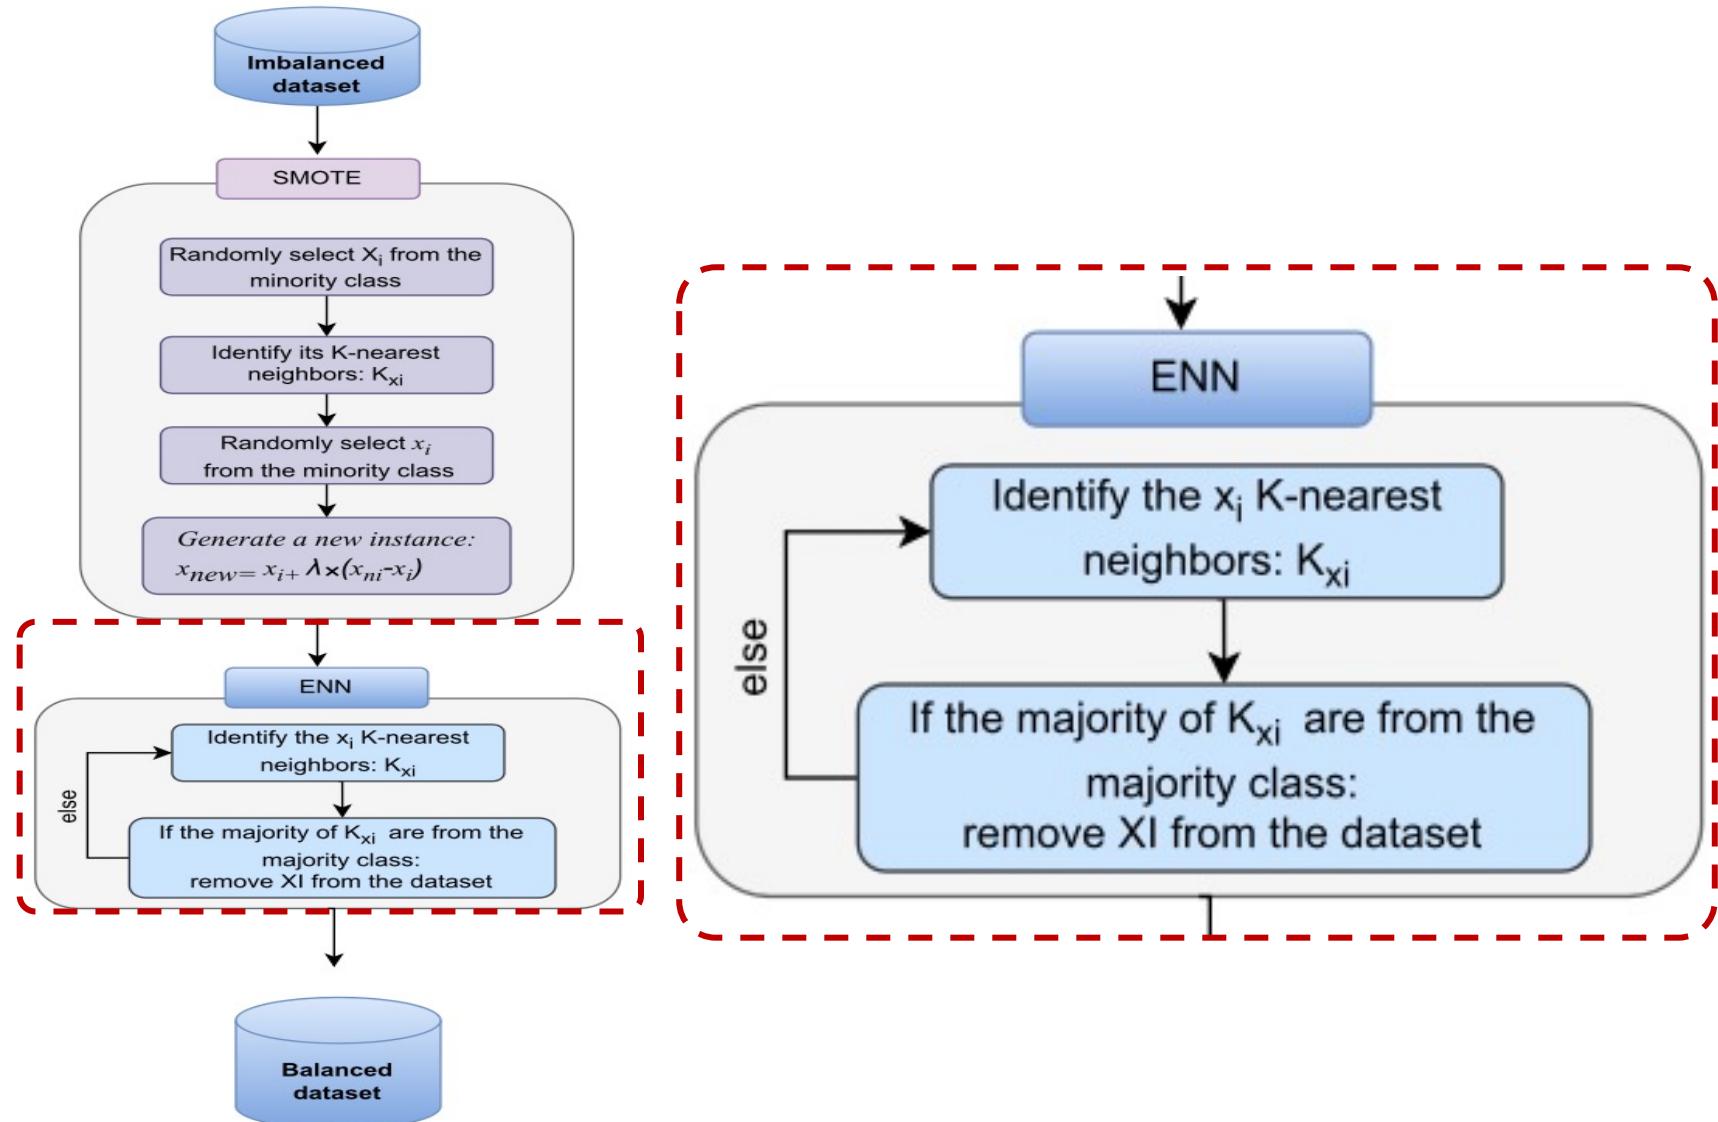
Medicare Fraud Detection Through Machine Learning: Addressing Class Imbalance

Ref	Dataset	ML Methods	Data Balancing Method	Evaluation
[21]	Medicare	Logistic Regression(LR), Random Forest (RF), Gradient Boosting Trees (GBT)	ROS, RUS, SMOTE, SMOTE variants, ADASYN	Area Under the Curve (AUC)= 0.82
[20]	Medicare Part B	Naive Bayes(NB), LR, Decision Trees (DT), K-Nearest Neighbors (KNN), Support Vector Machine (SVM), RF Word2Vec (Skip-gram, Continuous Bag Of Words (CBOW))	RUS	AUC
[22]	Medicare		Undersampling	AUC=0.870, Geometric Mean (G-mean)=0.783
[23]	Medicare	Logistic Regression (LR), RF, GBT, Multi-Layer Perceptron (MLP)	ROS, RUS	AUC=0.830
[14]	Healthcare Transactions	NB, LR, KNN, RF, Convolutional Neural Network (CNN)	Hybrid Resampling	Accuracy=97.58
[16]	Part D Medicare	eXtreme Gradient Boosting (XGBoost), RF	-	AUC= 0.97
[17]	Prescription Claims	LR, RF, Principal Component Analysis(PCA)	-	Receiver Operating Characteristic (ROC)= 0.76, F1-score= 0.88
[11]	Healthcare insurance	LR, DT, RF, XGBoost	CWS, ADASYN	AUC=0.95
[2]	Medicare	Category Boosting (CatBoost), XGBoost, RF, Extremely Randomized Trees(ET), Light Gradient Boosting Machine (LightGBM), DT, LR, Ensemble Feature Selection	-	AUC= 0.95, Area Under the Precision-Recall Curve (AUPRC)=0.78
[25]	Medicare	CatBoost, XGBoost, LightGBM, RF, ET	RUS	AUC=0.97, AUPRC=0.92
[26]	Medicare U.S. Medicare	CatBoost, XGBoost, RF, ET	RUS	AUC=0.99
[19]		XGBoost, RF	-	G- mean = 0.90, AUC = 0.962
[18]	Texas Medicaid	Bayesian Belief Network(BNN)	-	F-score=0.94
[6]	Healthcare Insurance	SVM, DT, RF, MLP	-	F-score=0.95
[24]	Healthcare Claims	Deep Autoencoders	-	precision=0.87, recall=1.00, F-score=0.93

Enhancing Medicare Fraud Detection Through Machine Learning: Addressing Class Imbalance With SMOTE-ENN

SMOTE-ENN process

Synthetic Minority Over-sampling technique with Edited Nearest Neighbors (SMOTE-ENN)



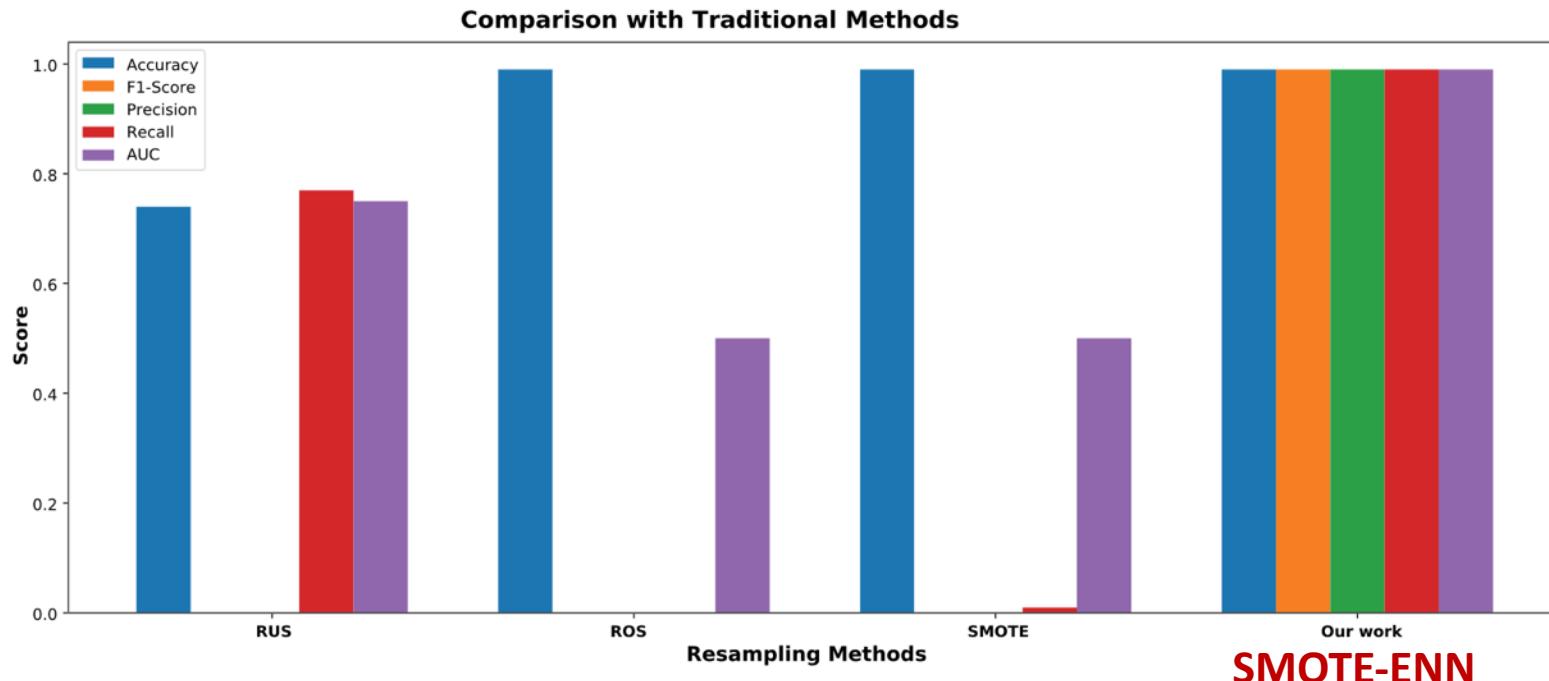
Enhancing Medicare Fraud Detection Through Machine Learning: Addressing Class Imbalance With SMOTE-ENN

Synthetic Minority Over-sampling technique with Edited Nearest Neighbors (SMOTE-ENN)

Classification results using SMOTE-ENN and cross-validation

Classifier	Accuracy	F1-Score	Precision	Recall	AUC
LR	0.65	0.65	0.69	0.67	0.73
DT	1.00	1.00	0.99	1.00	0.95
RF	0.95	0.95	0.95	0.95	0.99
XGBoost	0.96	0.96	0.96	0.96	0.99
Adaboost	0.65	0.64	0.70	0.67	0.68
LGBM	0.91	0.91	0.90	0.91	0.97

Enhancing Medicare Fraud Detection Through Machine Learning: Addressing Class Imbalance With SMOTE-ENN



Classifier	Accuracy	F1-Score	Precision	Recall	AUC
RUS	0.74	0.00	0.00	0.77	0.75
ROS	0.99	0.00	0.00	0.00	0.50
SMOTE	0.99	0.00	0.00	0.01	0.50
SMOTE-ENN	0.99	0.99	0.99	0.99	0.99

ML Evaluation of Imbalanced Dataset: Ensemble Learning and Data Augmentations (DA)

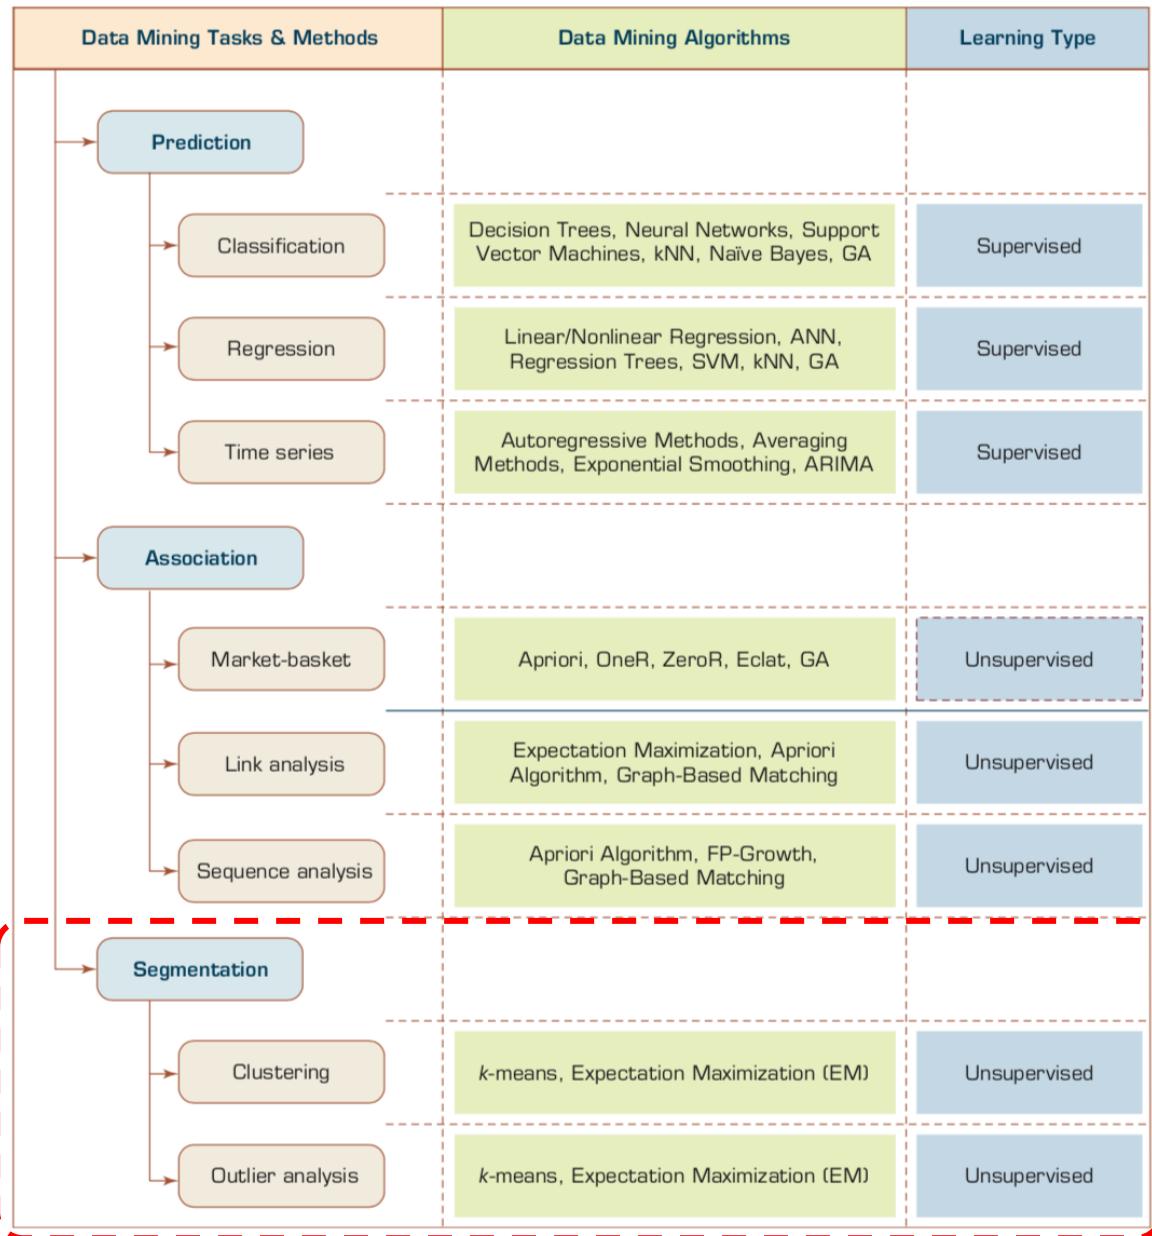
Data augmentation	Ensemble model	Accuracy (Best, Std)	F1 (Best, Std)	AUC (Best, Std)
Yeast-v6				
Borderline-SMOTE	LightGBM	98.490(98.653, 0.150)	99.230(99.314, 0.077)	76.668(76.751, 0.077)
No augmentation	Stacking-I	98.185(98.653, 0.320)	99.076(99.315, 0.165)	69.757(76.579, 3.658)
SVM-SMOTE	AdaBoost	98.072(98.653, 0.466)	99.015(99.314, 0.240)	75.577(80.253, 2.014)
ROS	Stacking-I	97.997(98.822, 0.415)	98.977(99.401, 0.214)	74.884(76.837, 2.607)
Borderline-SMOTE	Voting-Soft	97.949(98.653, 0.501)	98.951(99.314, 0.259)	75.854(80.253, 1.923)
ROS	XGBoost	97.866(98.485, 0.303)	98.908(99.227, 0.157)	76.348(76.665, 0.155)
SMOTE	Stacking-II	97.539(98.485, 0.708)	98.737(99.227, 0.370)	77.385(83.841, 2.273)
SMOTE	Voting-Hard	97.386(98.485, 0.707)	98.657(99.227, 0.370)	77.607(83.841, 2.492)
SMOTE-ENN	Random Forests	92.917(97.306, 1.932)	96.273(98.618, 1.048)	73.380(78.962, 1.694)
RUS	Stacking-II	87.345(94.444, 3.405)	93.087(97.093, 2.014)	81.085(88.581, 3.112)

Machine Learning: Unsupervised Learning: Cluster Analysis, Market Segmentation

Machine Learning: Data Mining Tasks & Methods

Unsupervised Learning: Cluster Analysis, Market Segmentation

Segmentation



Example of Cluster Analysis

Point	P	P(x,y)
p01	a	(3, 4)
p02	b	(3, 6)
p03	c	(3, 8)
p04	d	(4, 5)
p05	e	(4, 7)
p06	f	(5, 1)
p07	g	(5, 5)
p08	h	(7, 3)
p09	i	(7, 5)
p10	j	(8, 5)

K-Means Clustering

Point	P	P(x,y)	m1 distance	m2 distance	Cluster
p01	a	(3, 4)	1.95	3.78	Cluster1
p02	b	(3, 6)	0.69	4.51	Cluster1
p03	c	(3, 8)	2.27	5.86	Cluster1
p04	d	(4, 5)	0.89	3.13	Cluster1
p05	e	(4, 7)	1.22	4.45	Cluster1
p06	f	(5, 1)	5.01	3.05	Cluster2
p07	g	(5, 5)	1.57	2.30	Cluster1
p08	h	(7, 3)	4.37	0.56	Cluster2
p09	i	(7, 5)	3.43	1.52	Cluster2
p10	j	(8, 5)	4.41	1.95	Cluster2

m1 (3.67, 5.83)

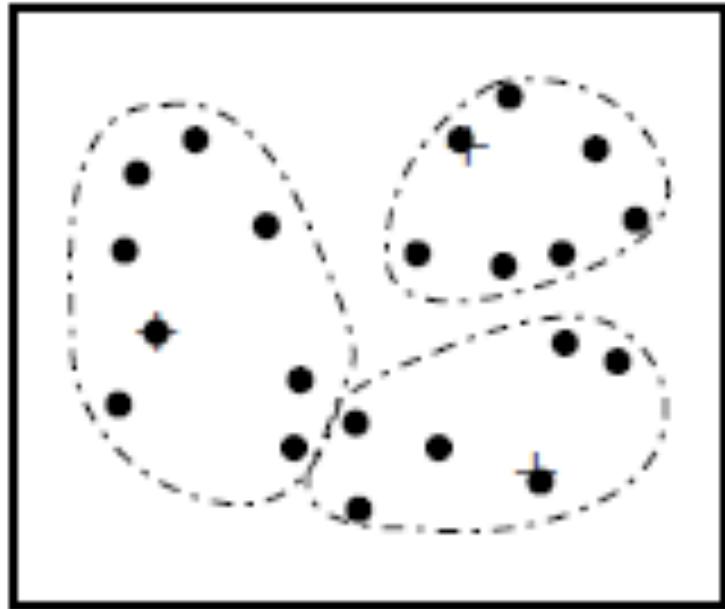
m2 (6.75, 3.50)

Cluster Analysis

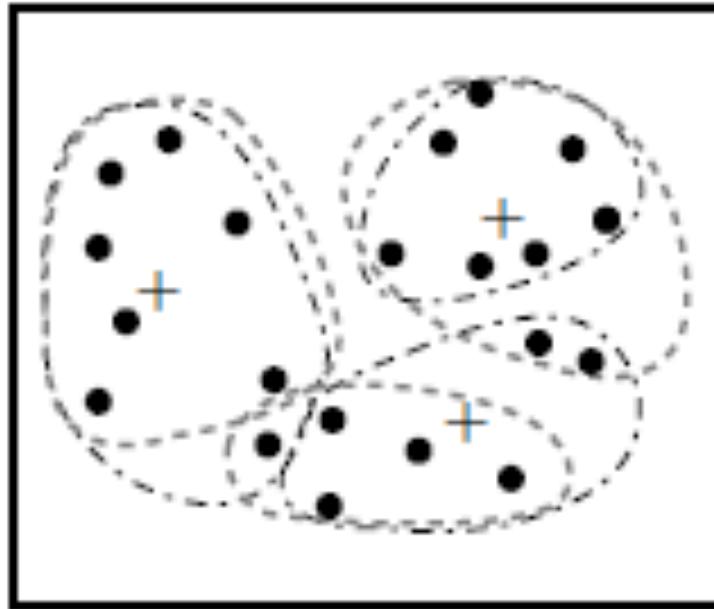
Cluster Analysis

- Used for automatic identification of **natural groupings** of things
- Part of the machine-learning family
- Employ **unsupervised learning**
- Learns the clusters of things from past data, then assigns new instances
- There is not an output variable
- Also known as **segmentation**

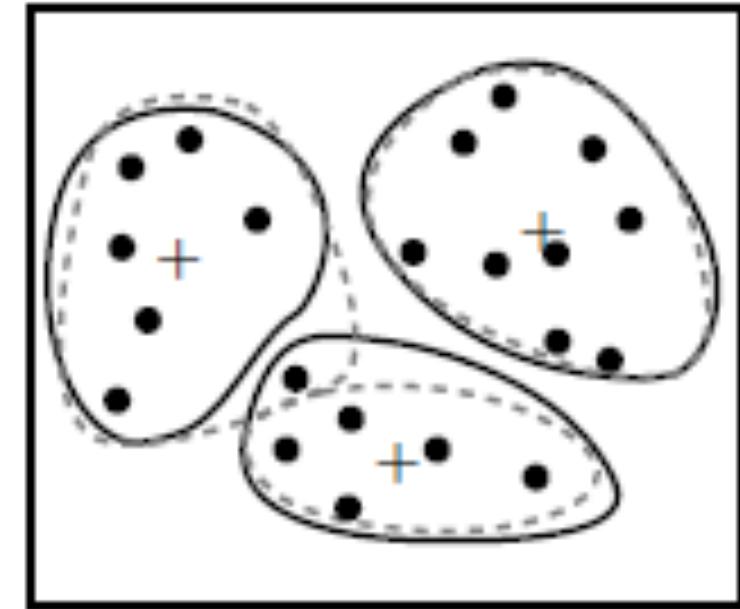
Cluster Analysis



(a)



(b)



(c)

Clustering of a set of objects based on the *k-means method*.
(*The mean of each cluster is marked by a “+”.*)

Cluster Analysis

- Clustering results may be used to
 - Identify natural **groupings of customers**
 - Identify rules for assigning new cases to classes for targeting/diagnostic purposes
 - Provide characterization, definition, labeling of populations
 - Decrease the size and complexity of problems for other data mining methods
 - Identify **outliers** in a specific domain (e.g., rare-event detection)

***k*-Means Clustering Algorithm**

- k : pre-determined number of clusters
- Algorithm (**Step 0:** determine value of k)

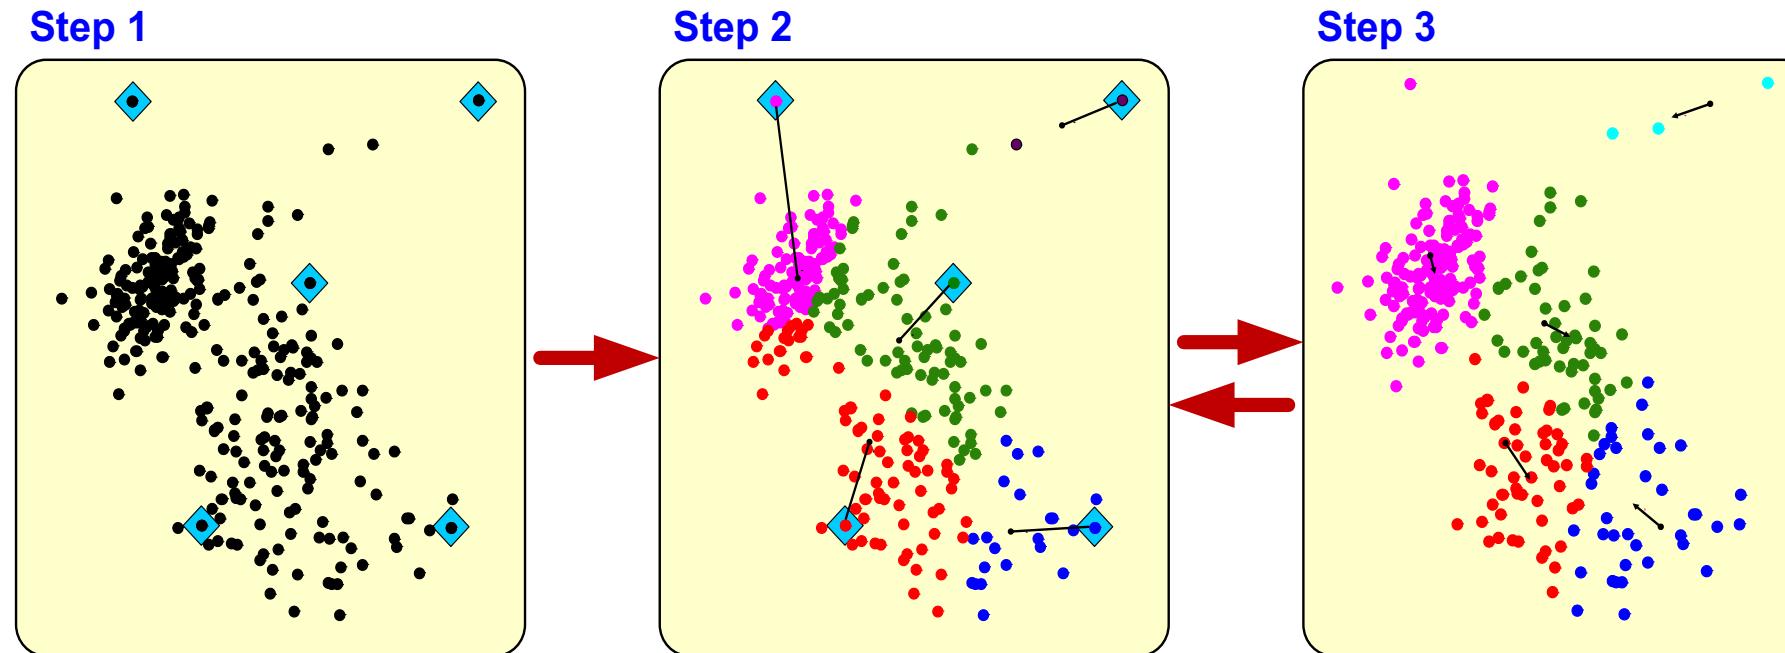
Step 1: Randomly generate k random points as initial cluster centers

Step 2: Assign each point to the nearest cluster center

Step 3: Re-compute the new cluster centers

Repetition step: Repeat steps 2 and 3 until some convergence criterion is met (usually that the assignment of points to clusters becomes stable)

Cluster Analysis for Data Mining - k -Means Clustering Algorithm



Similarity

Distance

Similarity and Dissimilarity Between Objects

- Distances are normally used to measure the similarity or dissimilarity between two data objects
- Some popular ones include: *Minkowski distance*:

$$d(i,j) = \sqrt[q]{(|x_{i1} - x_{j1}|^q + |x_{i2} - x_{j2}|^q + \dots + |x_{ip} - x_{jp}|^q)}$$

where $i = (x_{i1}, x_{i2}, \dots, x_{ip})$ and $j = (x_{j1}, x_{j2}, \dots, x_{jp})$ are two p -dimensional data objects, and q is a positive integer

- If $q = 1$, d is *Manhattan distance*

$$d(i,j) = |x_{i1} - x_{j1}| + |x_{i2} - x_{j2}| + \dots + |x_{ip} - x_{jp}|$$

Similarity and Dissimilarity Between Objects (Cont.)

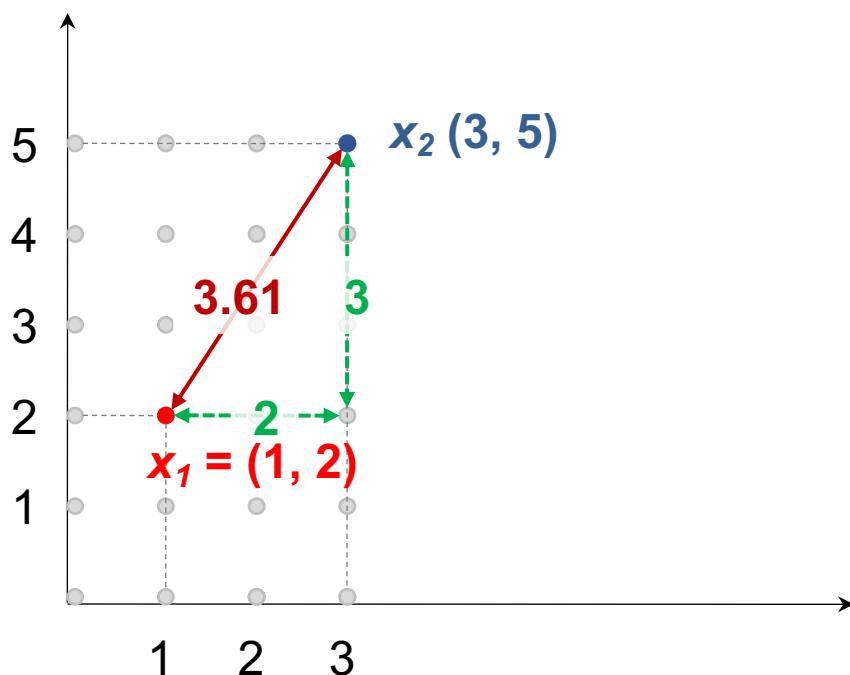
- If $q = 2$, d is **Euclidean distance**:

$$d(i,j) = \sqrt{(|x_{i_1} - x_{j_1}|^2 + |x_{i_2} - x_{j_2}|^2 + \dots + |x_{i_p} - x_{j_p}|^2)}$$

- **Properties**
 - $d(i,j) \geq 0$
 - $d(i,i) = 0$
 - $d(i,j) = d(j,i)$
 - $d(i,j) \leq d(i,k) + d(k,j)$
- **Also, one can use weighted distance, parametric Pearson product moment correlation, or other disimilarity measures**

Euclidean distance vs Manhattan distance

- Distance of two point $x_1 = (1, 2)$ and $x_2 (3, 5)$

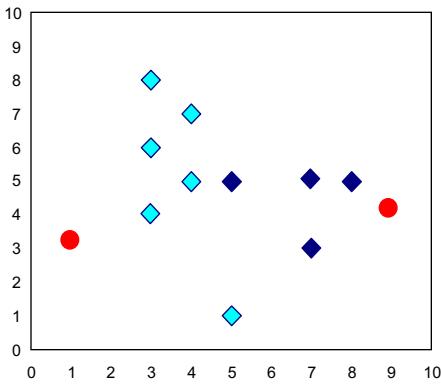


Euclidean distance:
 $= ((3-1)^2 + (5-2)^2)^{1/2}$
 $= (2^2 + 3^2)^{1/2}$
 $= (4 + 9)^{1/2}$
 $= (13)^{1/2}$
 $= 3.61$

Manhattan distance:
 $= (3-1) + (5-2)$
 $= 2 + 3$
 $= 5$

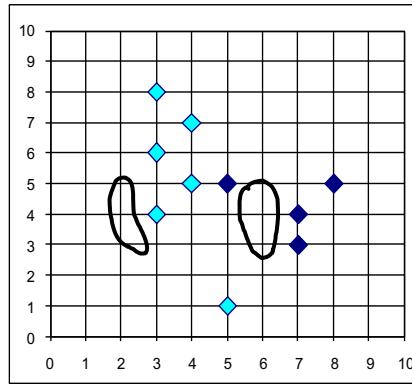
The *K*-Means Clustering Method

- Example

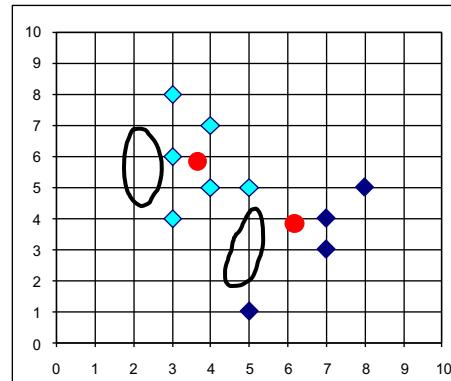


K=2
Arbitrarily choose K
object as initial
cluster center

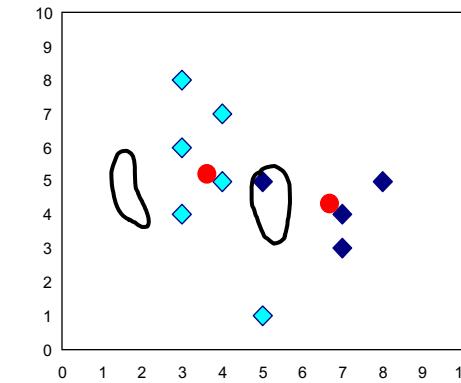
Assign
each
objects
to most
similar
center



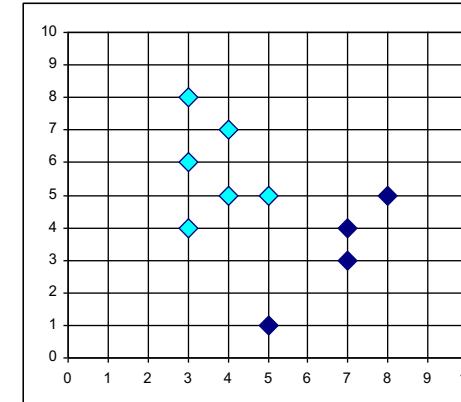
↑ reassign



Update
the
cluster
means



↓ reassign



Update
the
cluster
means

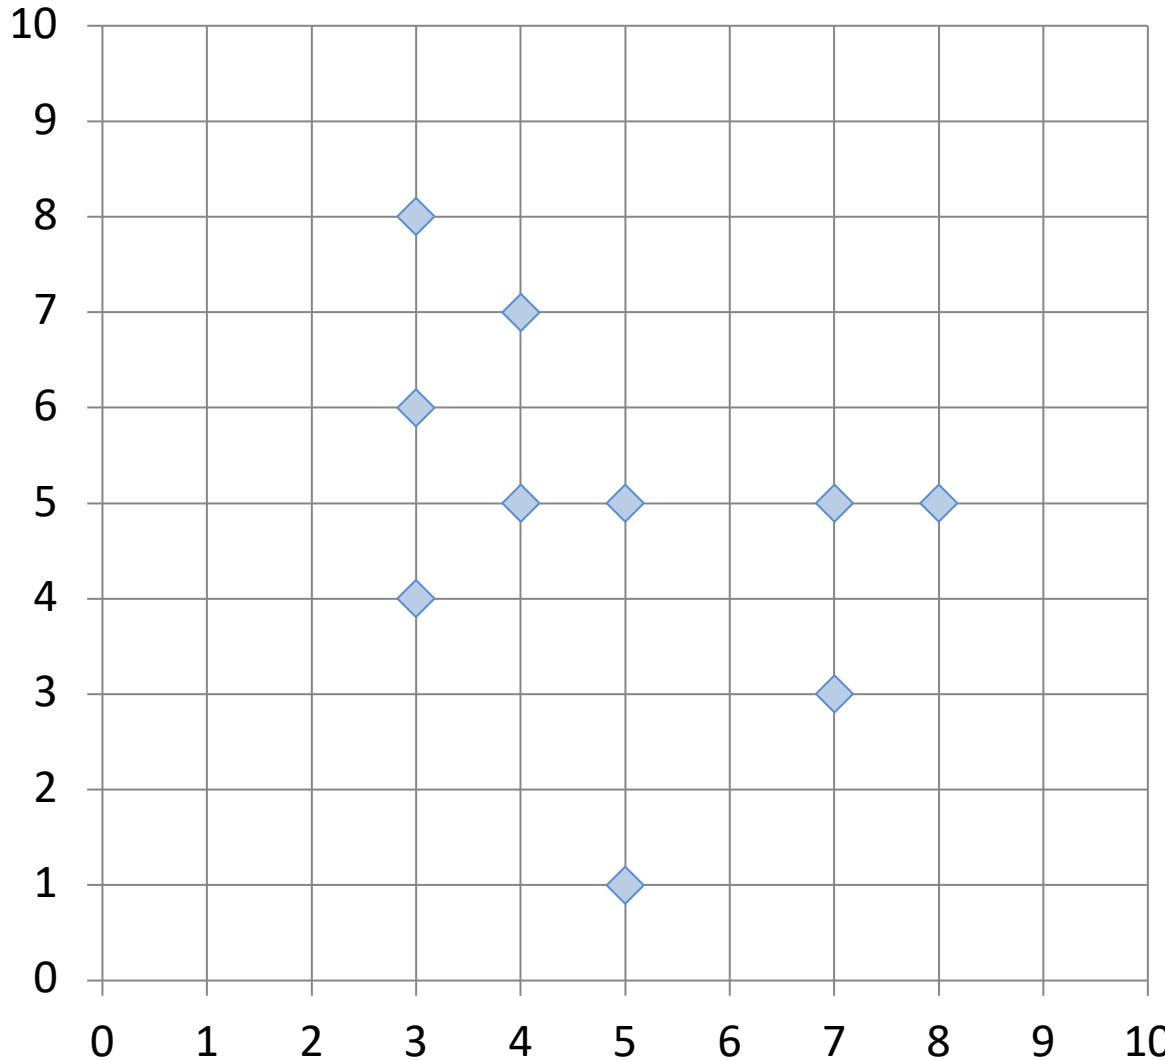
K-Means Clustering

Example of Cluster Analysis

Point	P	P(x,y)
p01	a	(3, 4)
p02	b	(3, 6)
p03	c	(3, 8)
p04	d	(4, 5)
p05	e	(4, 7)
p06	f	(5, 1)
p07	g	(5, 5)
p08	h	(7, 3)
p09	i	(7, 5)
p10	j	(8, 5)

K-Means Clustering

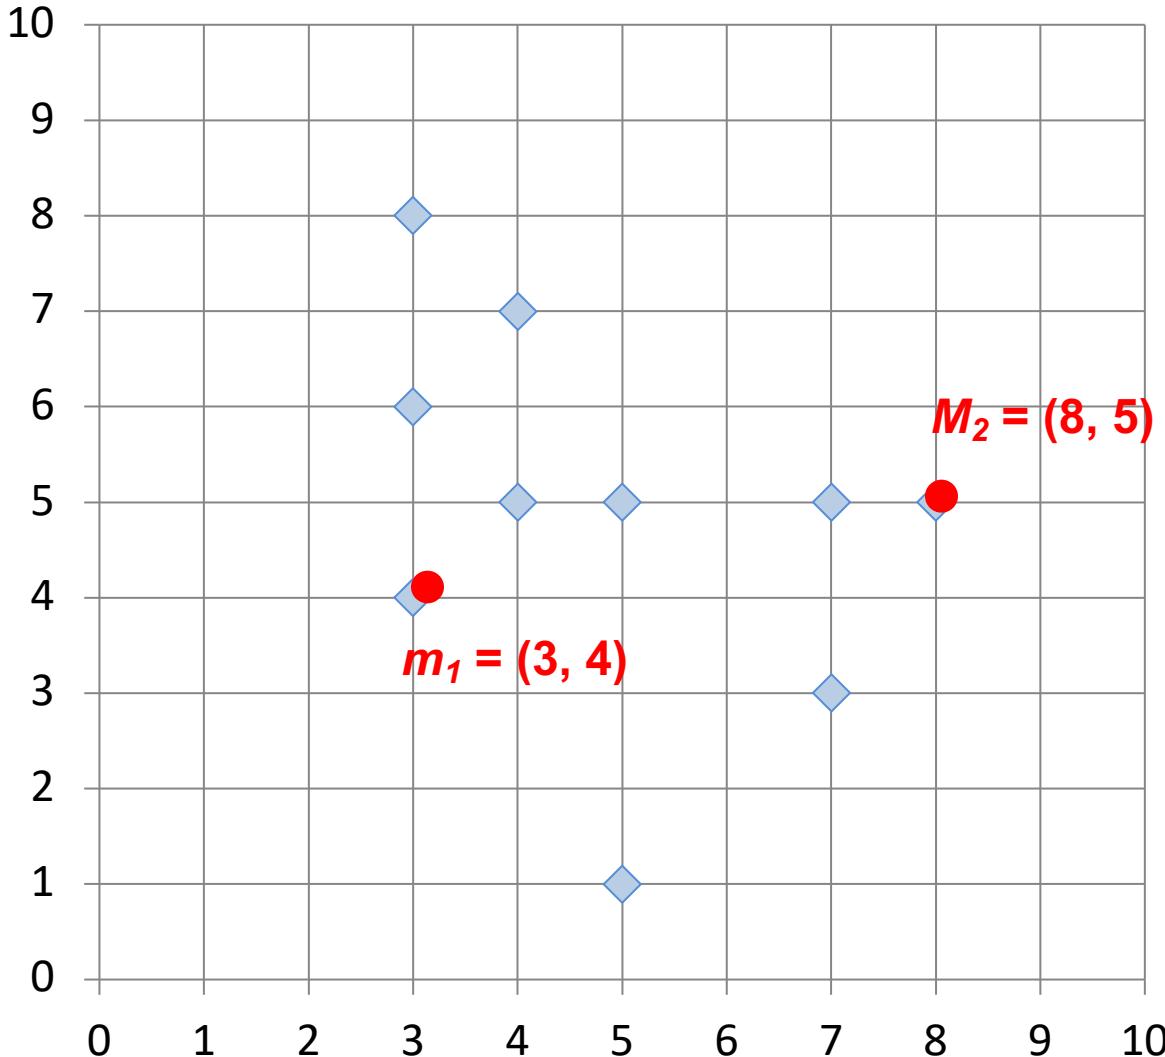
Step by Step



Point	P	P(x,y)
p01	a	(3, 4)
p02	b	(3, 6)
p03	c	(3, 8)
p04	d	(4, 5)
p05	e	(4, 7)
p06	f	(5, 1)
p07	g	(5, 5)
p08	h	(7, 3)
p09	i	(7, 5)
p10	j	(8, 5)

K-Means Clustering

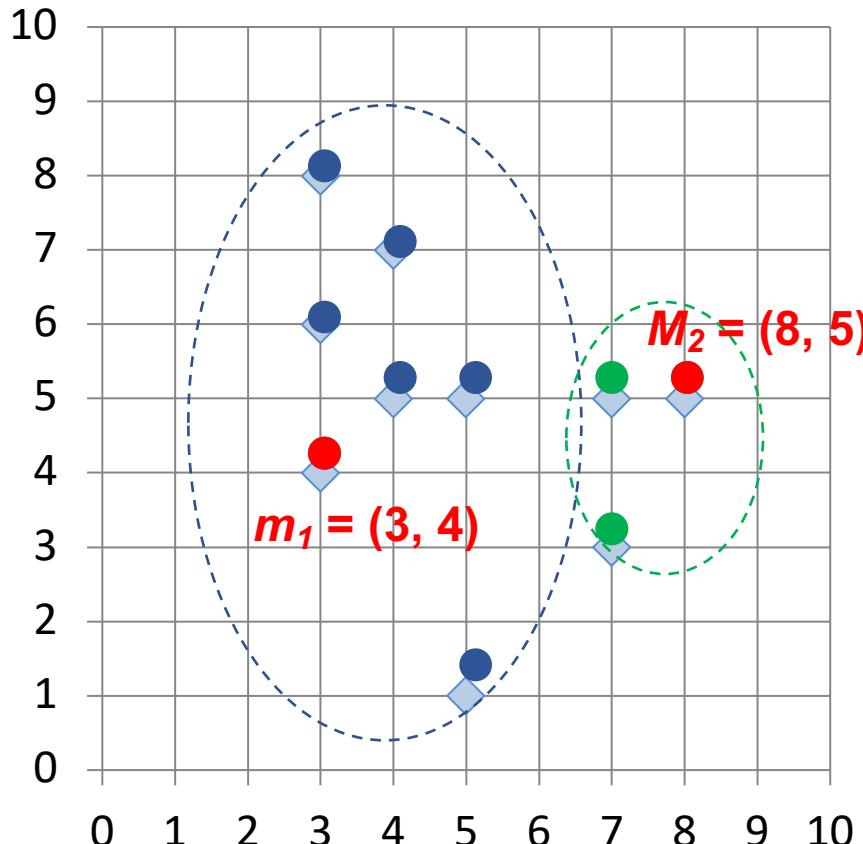
Step 1: K=2, Arbitrarily choose K object as initial cluster center



Initial m_1 (3, 4)
Initial m_2 (8, 5)

Step 2: Compute seed points as the centroids of the clusters of the current partition

Step 3: Assign each objects to most similar center



K-Means Clustering

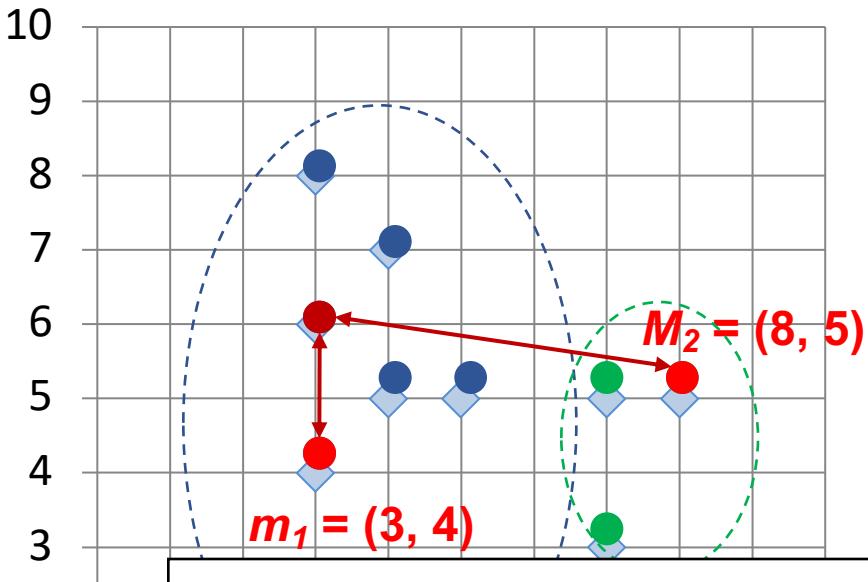
Point	P	P(x,y)	m1 distance	m2 distance	Cluster
p01	a	(3, 4)	0.00	5.10	Cluster1
p02	b	(3, 6)	2.00	5.10	Cluster1
p03	c	(3, 8)	4.00	5.83	Cluster1
p04	d	(4, 5)	1.41	4.00	Cluster1
p05	e	(4, 7)	3.16	4.47	Cluster1
p06	f	(5, 1)	3.61	5.00	Cluster1
p07	g	(5, 5)	2.24	3.00	Cluster1
p08	h	(7, 3)	4.12	2.24	Cluster2
p09	i	(7, 5)	4.12	1.00	Cluster2
p10	j	(8, 5)	5.10	0.00	Cluster2

Initial m_1 (3, 4)

Initial m_2 (8, 5)

Step 2: Compute seed points as the centroids of the clusters of the current partition

Step 3: Assign each objects to most similar center



K-1

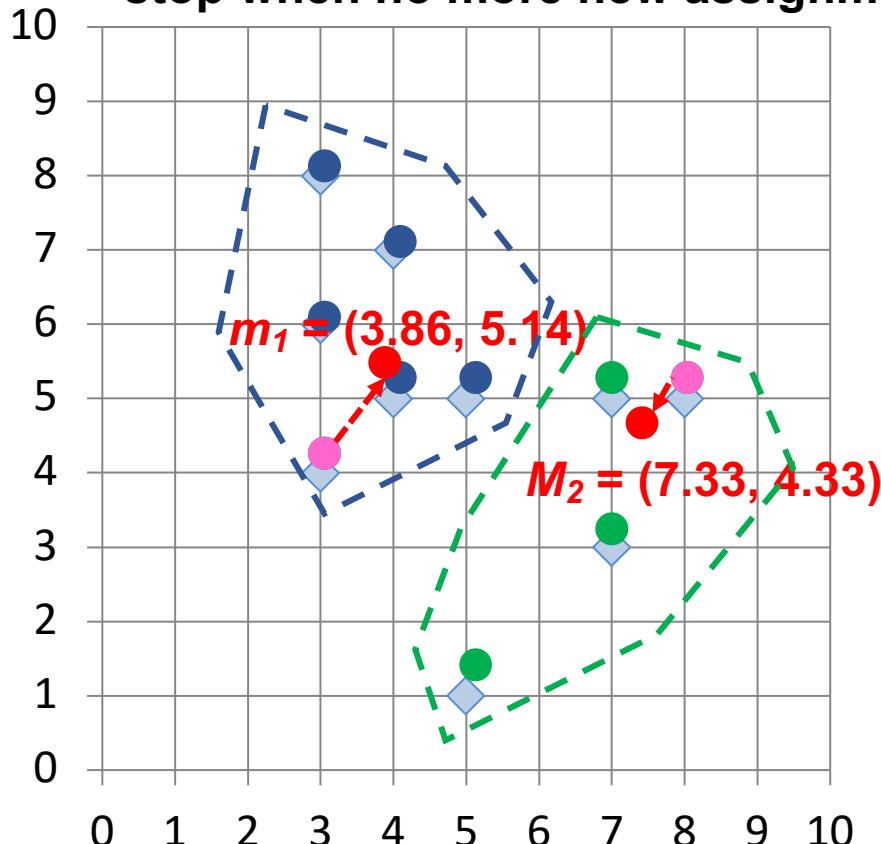
Euclidean distance
 $b(3,6) \leftrightarrow m1(3,4)$
 $= ((3-3)^2 + (4-6)^2)^{1/2}$
 $= (0^2 + (-2)^2)^{1/2}$
 $= (0 + 4)^{1/2}$
 $= (4)^{1/2}$
 $= 2.00$

Point	P	P(x,y)	m1 distance	m2 distance	Cluster
p01	a	(3, 4)	0.00	5.10	Cluster1
p02	b	(3, 6)	2.00	5.10	Cluster1
p03	c	(3, 8)	4.00	5.83	Cluster1
p04	d	(4, 5)	1.41	4.00	Cluster1
p05			Euclidean distance		
p06			$b(3,6) \leftrightarrow m2(8,5)$		ster1
p07			$= ((8-3)^2 + (5-6)^2)^{1/2}$		ster1
p08			$= (5^2 + (-1)^2)^{1/2}$		ster2
p09			$= (25 + 1)^{1/2}$		ster2
p10			$= (26)^{1/2}$		ster2

Initial $m_1 (3, 4)$

Initial $m_2 (8, 5)$

**Step 4: Update the cluster means,
Repeat Step 2, 3,
stop when no more new assignment**



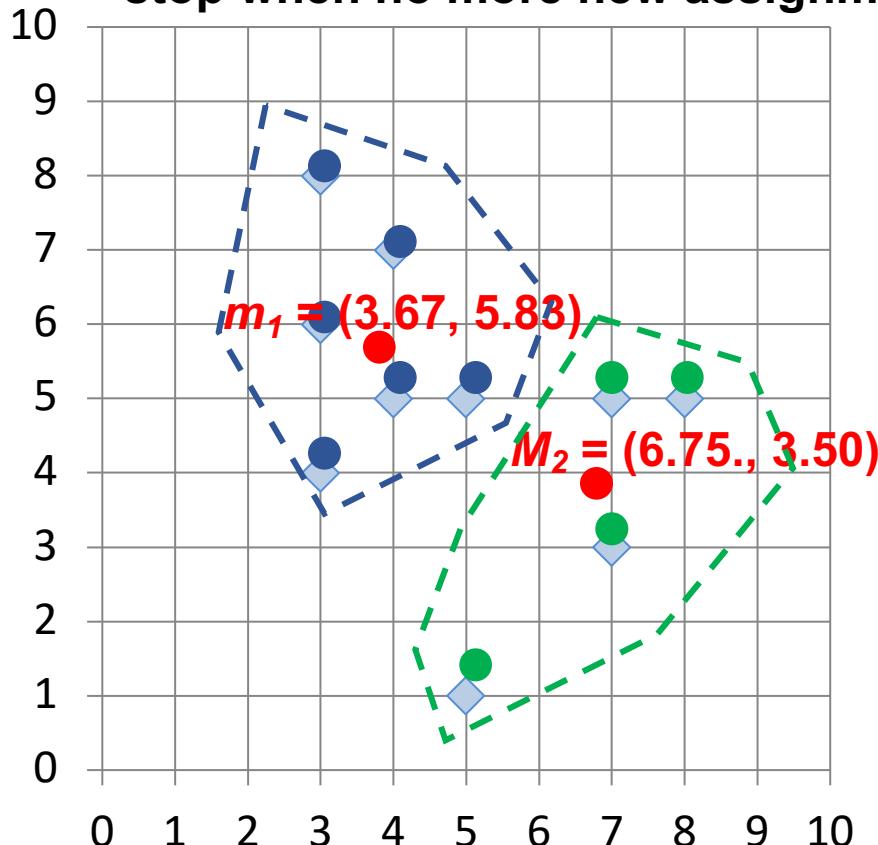
K-Means Clustering

Point	P	P(x,y)	m_1 distance	m_2 distance	Cluster
p01	a	(3, 4)	1.43	4.34	Cluster1
p02	b	(3, 6)	1.22	4.64	Cluster1
p03	c	(3, 8)	2.99	5.68	Cluster1
p04	d	(4, 5)	0.20	3.40	Cluster1
p05	e	(4, 7)	1.87	4.27	Cluster1
p06	f	(5, 1)	4.29	4.06	Cluster2
p07	g	(5, 5)	1.15	2.42	Cluster1
p08	h	(7, 3)	3.80	1.37	Cluster2
p09	i	(7, 5)	3.14	0.75	Cluster2
p10	j	(8, 5)	4.14	0.95	Cluster2

$$m_1 \quad (3.86, 5.14)$$

$$m_2 \quad (7.33, 4.33)$$

**Step 4: Update the cluster means,
Repeat Step 2, 3,
stop when no more new assignment**

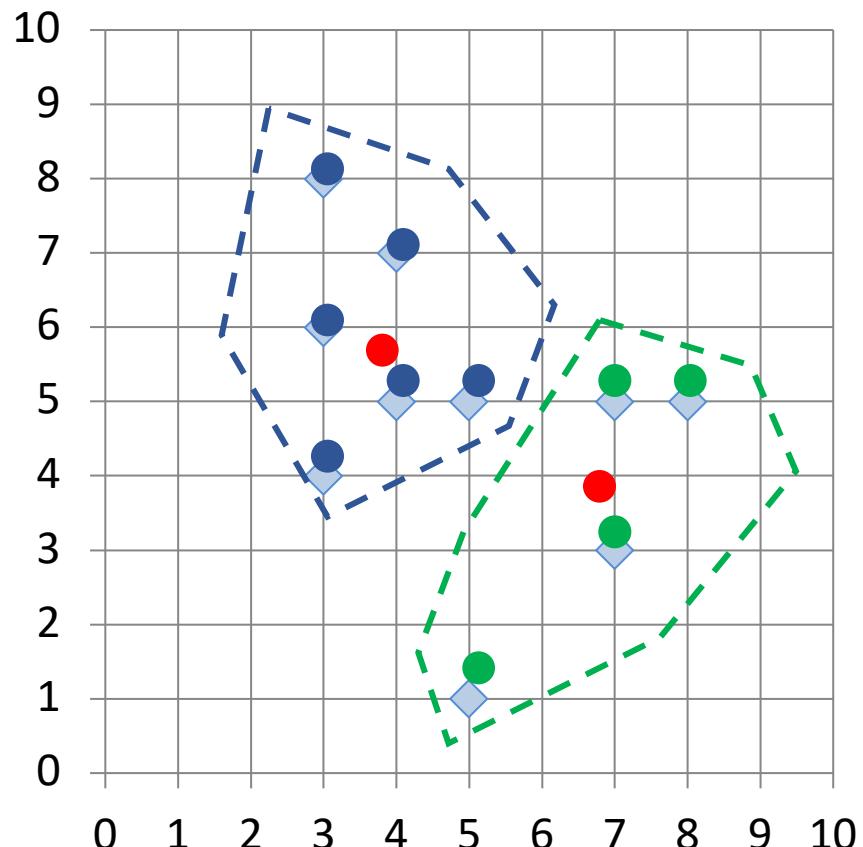


K-Means Clustering

Point	P	P(x,y)	m_1 distance	m_2 distance	Cluster
p01	a	(3, 4)	1.95	3.78	Cluster1
p02	b	(3, 6)	0.69	4.51	Cluster1
p03	c	(3, 8)	2.27	5.86	Cluster1
p04	d	(4, 5)	0.89	3.13	Cluster1
p05	e	(4, 7)	1.22	4.45	Cluster1
p06	f	(5, 1)	5.01	3.05	Cluster2
p07	g	(5, 5)	1.57	2.30	Cluster1
p08	h	(7, 3)	4.37	0.56	Cluster2
p09	i	(7, 5)	3.43	1.52	Cluster2
p10	j	(8, 5)	4.41	1.95	Cluster2

$$\begin{aligned}m_1 &= (3.67, 5.83) \\m_2 &= (6.75, 3.50)\end{aligned}$$

stop when no more new assignment

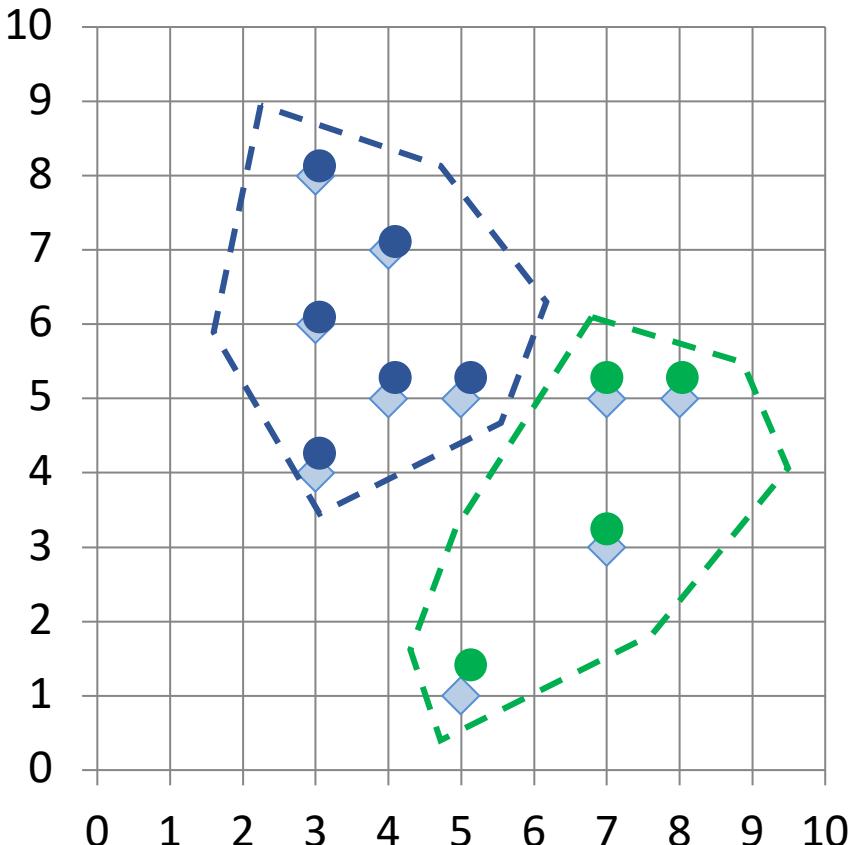


K-Means Clustering

$$\begin{aligned} m_1 & (3.67, 5.83) \\ m_2 & (6.75, 3.50) \end{aligned}$$

K-Means Clustering ($K=2$, two clusters)

stop when no more new assignment



Point	P	P(x,y)	m ₁ distance	m ₂ distance	Cluster
p01	a	(3, 4)	1.95	3.78	Cluster1
p02	b	(3, 6)	0.69	4.51	Cluster1
p03	c	(3, 8)	2.27	5.86	Cluster1
p04	d	(4, 5)	0.89	3.13	Cluster1
p05	e	(4, 7)	1.22	4.45	Cluster1
p06	f	(5, 1)	5.01	3.05	Cluster2
p07	g	(5, 5)	1.57	2.30	Cluster1
p08	h	(7, 3)	4.37	0.56	Cluster2
p09	i	(7, 5)	3.43	1.52	Cluster2
p10	j	(8, 5)	4.41	1.95	Cluster2

$$m_1 \ (3.67, 5.83)$$

$$m_2 \ (6.75, 3.50)$$

K-Means Clustering

K-Means Clustering

Point	P	P(x,y)	m1 distance	m2 distance	Cluster
p01	a	(3, 4)	1.95	3.78	Cluster1
p02	b	(3, 6)	0.69	4.51	Cluster1
p03	c	(3, 8)	2.27	5.86	Cluster1
p04	d	(4, 5)	0.89	3.13	Cluster1
p05	e	(4, 7)	1.22	4.45	Cluster1
p06	f	(5, 1)	5.01	3.05	Cluster2
p07	g	(5, 5)	1.57	2.30	Cluster1
p08	h	(7, 3)	4.37	0.56	Cluster2
p09	i	(7, 5)	3.43	1.52	Cluster2
p10	j	(8, 5)	4.41	1.95	Cluster2

m1 (3.67, 5.83)

m2 (6.75, 3.50)

Machine Learning: Unsupervised Learning: Cluster Analysis K-Means Clustering

Python in Google Colab (Python101)

<https://colab.research.google.com/drive/1FEG6DnGvwfUbeo4zJ1zTunjMqf2RkCrT>

python101.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share A

RAM Disk Editing

Table of contents

- Python101
- Python File Input / Output
- OS, IO, files, and Google Drive
- Python Try Except
- Python Class
- Python Programming
- Pythong String and Text
- Python Numpy
- Python Pandas
- Python Data Visualization
- Unsupervised Learning:
Association Analysis, Market Basket Analysis
 - Association Rules Generation from Frequent Itemsets
 - Market Basket Analysis
- Unsupervised Learning: Cluster Analysis, Market Segmentation
 - Cluster Analysis: K-Means Clustering**
 - Market Segmentation
 - Mall Customer Segmentation

+ Code + Text

```
1 import pandas as pd
2 from sklearn.cluster import KMeans
3 import plotly.express as px
4 data = {'X': [3, 3, 3, 4, 4, 5, 5, 7, 7, 8],
5          'Y': [4, 6, 8, 5, 7, 1, 5, 3, 5, 5]
6         }
7 df = pd.DataFrame(data, columns=['X', 'Y'])
8 print(df)
9 kmeans = KMeans(n_clusters=2)
10 cluster = kmeans.fit_predict(df[['X', 'Y']])
11 df['Cluster'] = cluster
12 print(df)
13 px.scatter(data_frame=df, x=df['X'], y=df['Y'], color=df['cluster'], range_x = (0,10), range_y = (0,10))
```

X Y

	X	Y
0	3	4
1	3	6
2	3	8
3	4	5
4	4	7
5	5	1
6	5	5
7	7	3
8	7	5
9	8	5

K-Means Clustering

color

0 0.2 0.4 0.6 0.8 0.8 1

	X	Y	Cluster
0	3	4	0
1	3	6	0
2	3	8	0
3	4	5	0
4	4	7	0
5	5	1	1
6	5	5	0

<https://tinyurl.com/aintpuppython101>

```
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=2)
cluster = kmeans.fit_predict(df[['X', 'Y']])
```

```
1 import pandas as pd
2 from sklearn.cluster import KMeans
3 import plotly.express as px
4 data = {'X': [3, 3, 3, 4, 4, 5, 5, 7, 7, 8],
5          'Y': [4, 6, 8, 5, 7, 1, 5, 3, 5, 5]
6         }
7 df = pd.DataFrame(data, columns=['X', 'Y'])
8 print(df)
9 kmeans = KMeans(n_clusters=2)
10 cluster = kmeans.fit_predict(df[['X', 'Y']])
11 df['Cluster'] = cluster
12 print(df)
13 px.scatter(data_frame=df, x=df['X'], y=df['Y'], color=df['cluster'], range_x = (0,10), range_y = (0,10), title='K-Means Clustering')
```

```
from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=2)
cluster = kmeans.fit_predict(df[['X', 'Y']])
```

```
import pandas as pd
from sklearn.cluster import KMeans
import plotly.express as px
data = {'X': [3, 3, 3, 4, 4, 5, 5, 7, 7, 8],
        'Y': [4, 6, 8, 5, 7, 1, 5, 3, 5, 5]}
df = pd.DataFrame(data, columns=['X', 'Y'])
print(df)
kmeans = KMeans(n_clusters=2)
cluster = kmeans.fit_predict(df[['X', 'Y']])
df['Cluster'] = cluster
print(df)
px.scatter(data_frame=df, x=df['X'], y=df['Y'],
           color=df['cluster'], range_x = (0,10), range_y = (0,10),
           title='K-Means Clustering')
```

K-Means Clustering

```
1 #importing the libraries
2 import numpy as np
3 import matplotlib.pyplot as plt
4 %matplotlib inline
5 import pandas as pd
6
7 #importing the Iris dataset with pandas
8 # Load dataset
9 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
10 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
11 df = pd.read_csv(url, names=names)
12
13 array = df.values
14 X = array[:,0:4]
15 Y = array[:,4]
16
17 #Finding the optimum number of clusters for k-means classification
18 from sklearn.cluster import KMeans
19 wcss = []
20
21 for i in range(1, 8):
22     kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)
23     kmeans.fit(X)
24     wcss.append(kmeans.inertia_)
25
26 #Plotting the results onto a line graph, allowing us to observe 'The elbow'
27 plt.rcParams["figure.figsize"] = (10,8)
28 plt.plot(range(1, 8), wcss)
29 plt.title('The elbow method')
30 plt.xlabel('Number of clusters')
31 plt.ylabel('WCSS') #within cluster sum of squares
32 plt.show()
```

```
#importing the libraries
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import pandas as pd

#importing the Iris dataset with pandas
# Load dataset
url = "https://archive.ics.uci.edu/ml/machine-
learning-databases/iris/iris.data"
names = ['sepal-length', 'sepal-width',
'petal-length', 'petal-width', 'class']
df = pd.read_csv(url, names=names)

array = df.values
X = array[:,0:4]
Y = array[:,4]
```

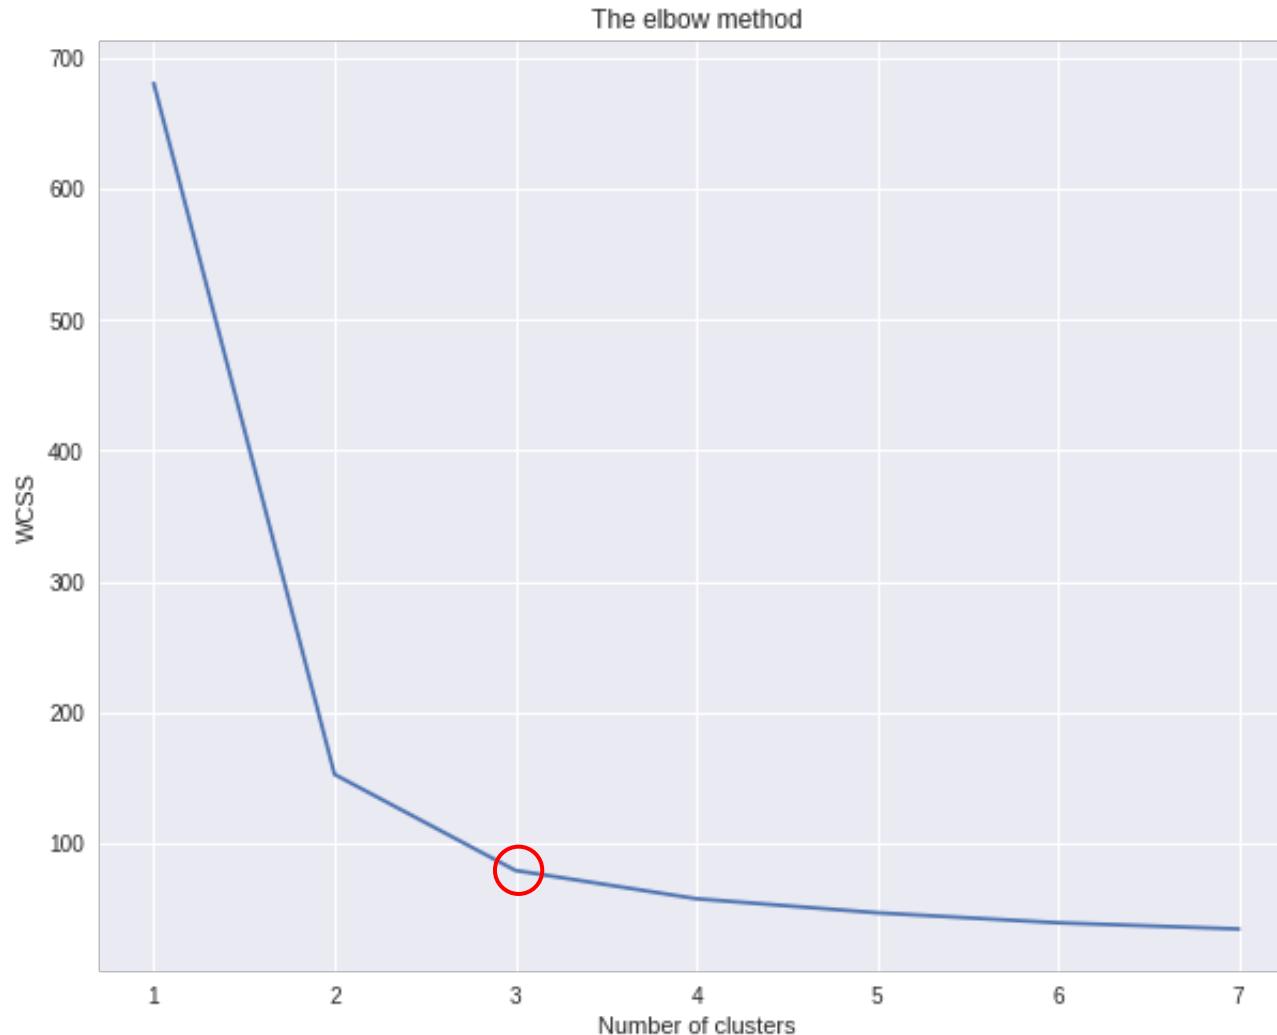
```
#Finding the optimum number of clusters for k-means classification
from sklearn.cluster import KMeans
wcss = []

for i in range(1, 8):
    kmeans = KMeans(n_clusters = i, init = 'k-means++',
max_iter = 300, n_init = 10, random_state = 0)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

#Plotting the results onto a line graph, allowing us to observe 'The elbow'
plt.rcParams["figure.figsize"] = (10,8)
plt.plot(range(1, 8), wcss)
plt.title('The elbow method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS') #within cluster sum of squares
plt.show()
```

K-Means Clustering

The elbow method ($k=3$)



```
kmeans = KMeans(n_clusters = 3,  
init = 'k-means++', max_iter = 300,  
n_init = 10, random_state = 0)  
y_kmeans = kmeans.fit_predict(X)
```

```
1 #Applying kmeans to the dataset / Creating the kmeans classifier  
2 kmeans = KMeans(n_clusters = 3, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)  
3 y_kmeans = kmeans.fit_predict(X).
```

```

#Visualising the clusters
plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 100,
c = 'red', label = 'Iris-setosa')
plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 100,
c = 'blue', label = 'Iris-versicolour')
plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 100,
c = 'green', label = 'Iris-virginica')

#Plotting the centroids of the clusters
plt.scatter(kmeans.cluster_centers_[:, 0],
kmeans.cluster_centers_[:,1], s = 100, c = 'yellow', label =
'Centroids')

plt.legend()

```

```

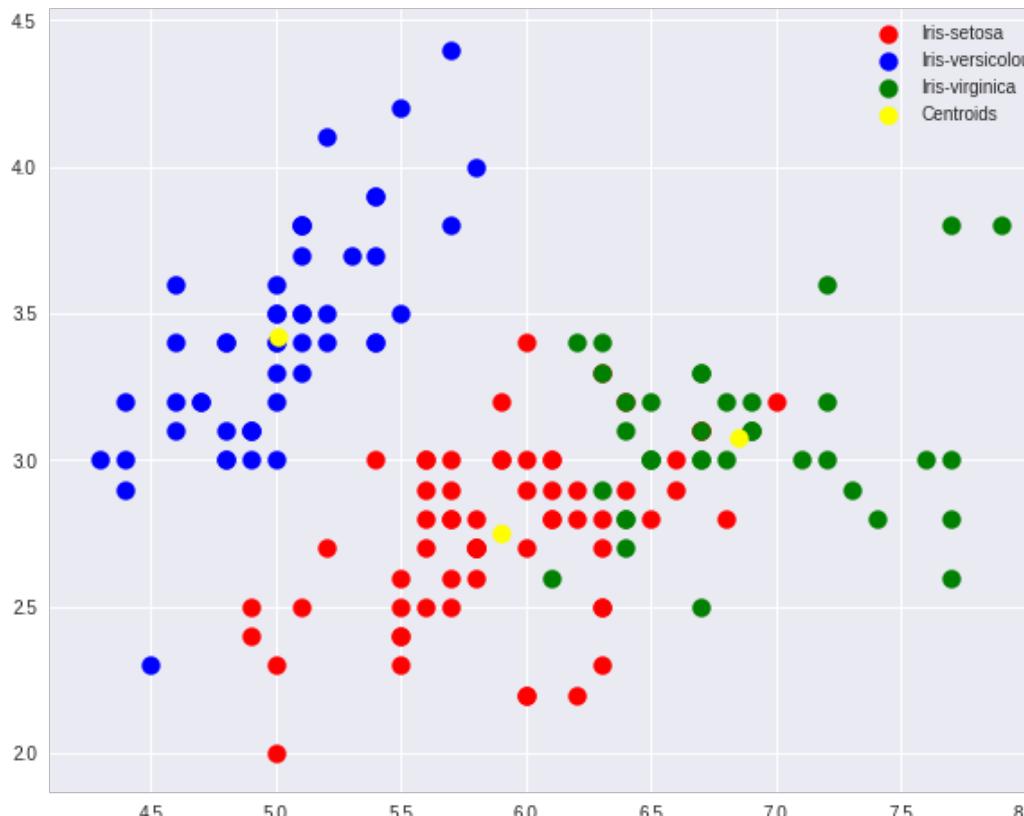
1 #Visualising the clusters
2 plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 100, c = 'red', label = 'Iris-setosa')
3 plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 100, c = 'blue', label = 'Iris-versicolour')
4 plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 100, c = 'green', label = 'Iris-virginica')
5
6 #Plotting the centroids of the clusters
7 plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:,1], s = 100, c = 'yellow', label = 'Centroids')
8
9 plt.legend()

```

K-Means Clustering

```
1 #Applying kmeans to the dataset / Creating the kmeans classifier
2 kmeans = KMeans(n_clusters = 3, init = 'k-means++', max_iter = 300, n_init = 10, random_state = 0)
3 y_kmeans = kmeans.fit_predict(X)

1 #Visualising the clusters
2 plt.scatter(X[y_kmeans == 0, 0], X[y_kmeans == 0, 1], s = 100, c = 'red', label = 'Iris-setosa')
3 plt.scatter(X[y_kmeans == 1, 0], X[y_kmeans == 1, 1], s = 100, c = 'blue', label = 'Iris-versicolour')
4 plt.scatter(X[y_kmeans == 2, 0], X[y_kmeans == 2, 1], s = 100, c = 'green', label = 'Iris-virginica')
5
6 #Plotting the centroids of the clusters
7 plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:,1], s = 100, c = 'yellow', label = 'Centroids')
8
9 plt.legend()
```



Market Segmentation

The screenshot shows a Jupyter Notebook interface with the following details:

- Title:** python101.ipynb
- Toolbar:** File, Edit, View, Insert, Runtime, Tools, Help, All changes saved
- Table of contents:** Python Pandas, Python Data Visualization, Unsupervised Learning: Association Analysis, Market Basket Analysis, Association Rules Generation from Frequent Itemsets, Market Basket Analysis, Unsupervised Learning: Cluster Analysis, Market Segmentation, Cluster Analysis: K-Means Clustering, **Market Segmentation** (selected), Mall Customer Segmentation, Machine Learning with scikit-learn, Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python, Investment Portfolio Optimization, Text Analytics and Natural Language Processing (NLP), Python for Natural Language.
- Code Cell:** # Source: <https://www.kaggle.com/amanjarvis1704/k-means-clustering>

```
1 # Source: https://www.kaggle.com/amanjarvis1704/k-means-clustering
2 import pandas as pd
3 import numpy as np
4 import plotly.express as px
5 import matplotlib.pyplot as plt
6 from sklearn.cluster import KMeans
7 %matplotlib inline
8 url='https://raw.githubusercontent.com/imamanmehrotra/Datasets/main/income_kmeans.csv'
9 df=pd.read_csv(url)
10 print(df.shape)
11 print(df.describe())
12 print(df)
13 px.scatter(data_frame=df, x='Age', y='Income($)', hover_data=['Name'])
```
- Output Cell:** (22, 3)

	Age	Income(\$)
count	22.000000	22.000000
mean	34.818182	90431.818182
std	5.901060	43505.964412
min	26.000000	45000.000000
25%	29.000000	58500.000000
50%	36.500000	67500.000000
75%	39.750000	135250.000000
max	43.000000	162000.000000

	Name	Age	Income(\$)
0	Rob	27	70000
1	Michael	29	90000
2	Mohan	29	61000
3	Ismail	28	60000
4	Kory	42	150000
5	Gautam	39	155000

Mall Customer Segmentation

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** python101.ipynb
- File Menu:** File Edit View Insert Runtime Tools Help All changes saved
- Toolbar:** Comment Share Settings (A)
- Table of Contents:** Python Pandas, Python Data Visualization, Unsupervised Learning: Association Analysis, Market Basket Analysis, Association Rules Generation from Frequent Itemsets, Market Basket Analysis, Unsupervised Learning: Cluster Analysis, Market Segmentation, Cluster Analysis: K-Means Clustering, Market Segmentation, **Mall Customer Segmentation** (selected), Machine Learning with scikit-learn, Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python, Investment Portfolio Optimization, Text Analytics and Natural Language Processing (NLP), Python for Natural Language.
- Code Cell:** + Code + Text
- RAM/Disk:** RAM Disk
- Cell Content:**
 - Mall Customer Segmentation**
 - Source: <https://www.kaggle.com/vjchoudhary7/customer-segmentation-tutorial-in-python>
 - Code snippet:

```
1 #Source: Agustin Pugliese (2020), Clustering Model Comparison with Plotly, https://www.kaggle.com/agustinpugliese/c
2 import os
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.cluster import KMeans, AgglomerativeClustering, AffinityPropagation, DBSCAN
8 import scipy.cluster.hierarchy as sch
9 import plotly.figure_factory as ff
10 import plotly.express as px
11 import plotly.graph_objects as go
12 sns.set()
13 %matplotlib inline
14 url='https://web.ntpu.edu.tw/~myday/data/example/Mall_Customers.csv'
15 df=pd.read_csv(url)
16 print(df.shape)
17 print(df.describe())
18 print(df)
19 px.scatter(data_frame=df, x='Age',y='Annual Income (k$)')
```
 - Output:

```
(200, 5)
   CustomerID      Age  Annual Income (k$)  Spending Score (1-100)
count  200.000000  200.000000    200.000000  200.000000
mean   100.500000  38.850000    60.560000   50.200000
std    57.879185  13.969007   26.264721   25.823522
min     1.000000  18.000000    15.000000   1.000000
25%    50.750000  28.750000    41.500000   34.750000
50%   100.500000  36.000000    61.500000   50.000000
75%   150.250000  49.000000    78.000000   73.000000
```

Mall Customer Segmentation

python101.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share A

RAM Disk Editing

Table of contents

- Python Pandas
- Python Data Visualization
- Unsupervised Learning: Association Analysis, Market Basket Analysis
- Association Rules Generation from Frequent Itemsets
- Market Basket Analysis
- Unsupervised Learning: Cluster Analysis, Market Segmentation
 - Cluster Analysis: K-Means Clustering
 - Market Segmentation
- Mall Customer Segmentation**
- Machine Learning with scikit-learn
 - Classification and Prediction
 - K-Means Clustering
- Deep Learning for Financial Time Series Forecasting
- Portfolio Optimization and Algorithmic Trading
 - Investment Portfolio Optimisation with Python
 - Efficient Frontier Portfolio Optimisation in Python
 - Investment Portfolio Optimization
- Text Analytics and Natural Language Processing (NLP)

+ Code + Text

```
1 fig = px.scatter(df2, x="Age", y = "Annual Income (k$)", size='Spending Score (1-100)', color="Gender")
2 fig.show()
```

Annual Income (k\$)

Age

Gender=Male
Gender=Female

Mall Customer Segmentation

python101.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Comment Share A

RAM Disk Editing

Table of contents

- Python Pandas
- Python Data Visualization
- Unsupervised Learning: Association Analysis, Market Basket Analysis
- Association Rules Generation from Frequent Itemsets
- Market Basket Analysis
- Unsupervised Learning: Cluster Analysis, Market Segmentation
 - Cluster Analysis: K-Means Clustering
 - Market Segmentation
- Mall Customer Segmentation**
- Machine Learning with scikit-learn
 - Classification and Prediction
 - K-Means Clustering
- Deep Learning for Financial Time Series Forecasting
- Portfolio Optimization and Algorithmic Trading
 - Investment Portfolio Optimisation with Python
 - Efficient Frontier Portfolio Optimisation in Python
 - Investment Portfolio Optimization
- Text Analytics and Natural Language Processing (NLP)

+ Code + Text

```
1 kmeans = KMeans(n_clusters = 5, init="k-means++", max_iter = 500, n_init = 10, random_state = 123)
2 identified_clusters = kmeans.fit_predict(X)
3
4 data_with_clusters = df2.copy()
5 data_with_clusters[ 'Cluster' ] = identified_clusters
6 fig = px.scatter_3d(data_with_clusters, x = 'Age', y='Annual Income (k$)', z='Spending Score (1-100)',
7                      color='Cluster', opacity = 0.8, size='Age', size_max=30)
8 fig.show()
```

Cluster

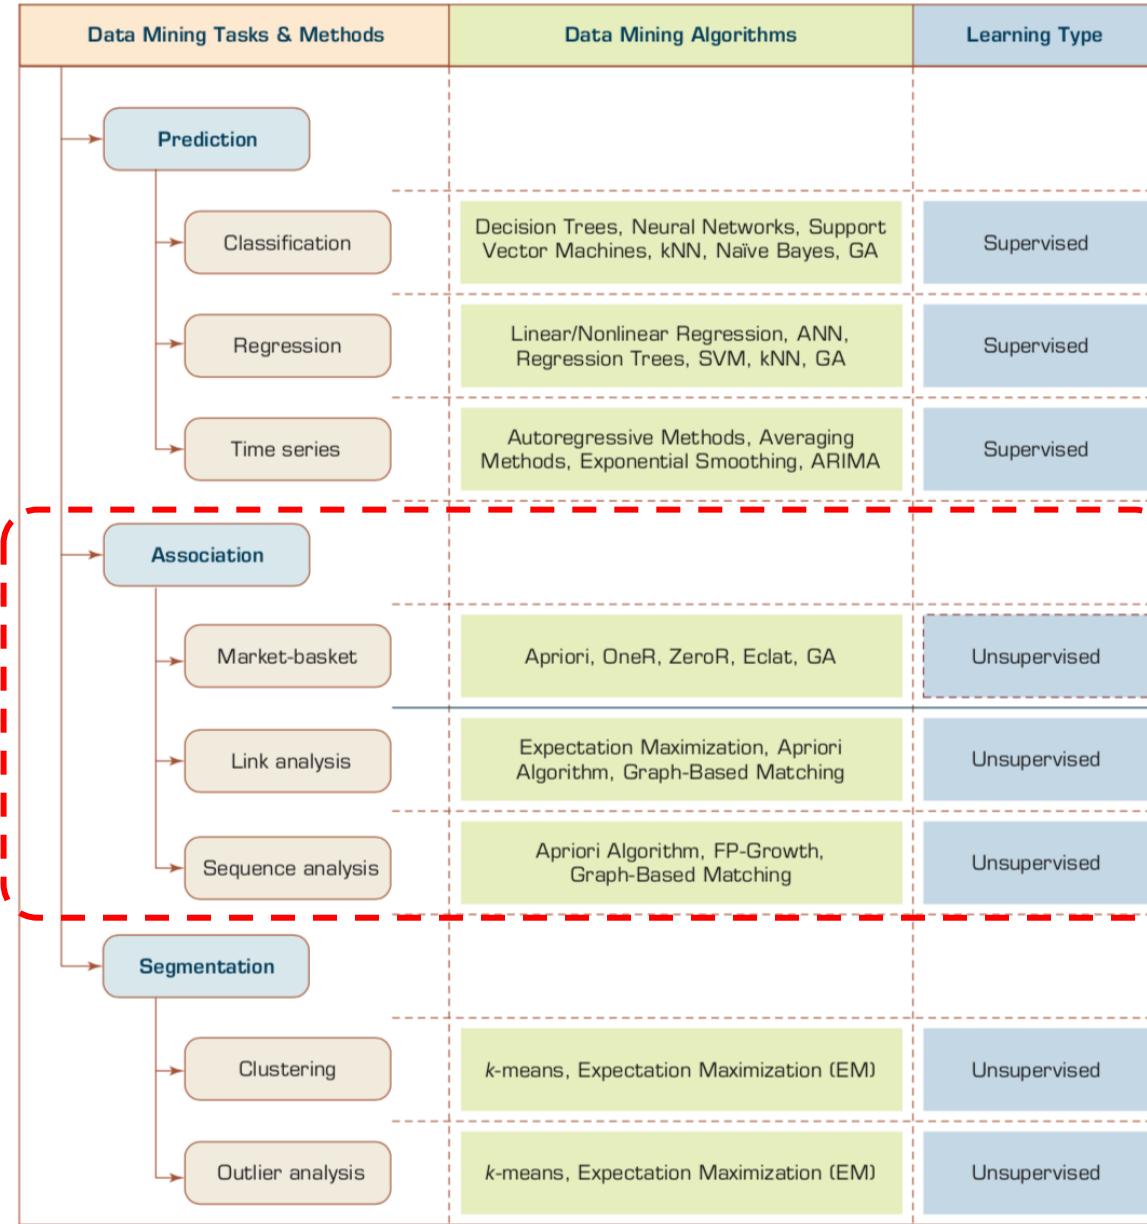
4
3.5
3
2.5
2
1.5
1
0.5
0

Machine Learning: Unsupervised Learning: Association Analysis, Market Basket Analysis

Machine Learning: Data Mining Tasks & Methods

Unsupervised Learning: Association Analysis Market-basket

Association

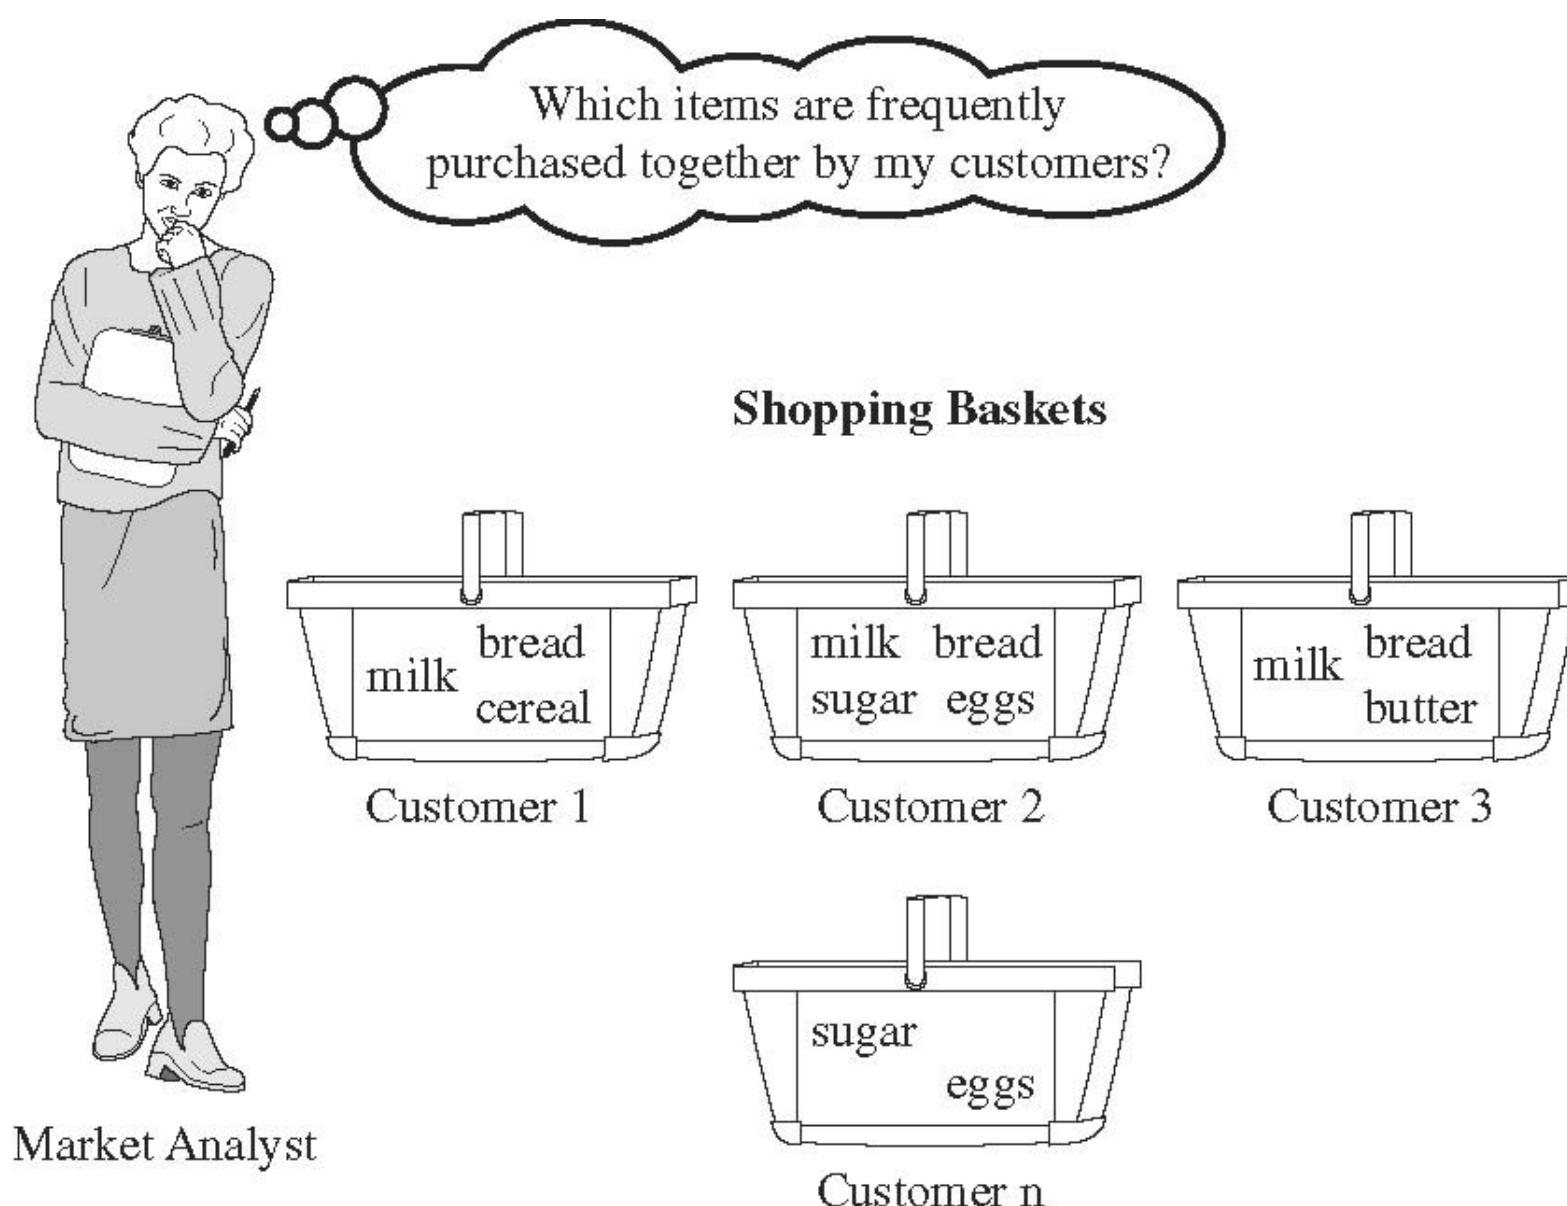


Transaction Database

Transaction ID	Items bought
T01	A, B, D
T02	A, C, D
T03	B, C, D, E
T04	A, B, D
T05	A, B, C, E
T06	A, C
T07	B, C, D
T08	B, D
T09	A, C, E
T10	B, D

Association Analysis

Market Basket Analysis



Python mlxtend Association Rules

```
# !pip install mlxtend
import pandas as pd
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori
from mlxtend.frequent_patterns import association_rules

dataset = [['A', 'B', 'D'],
           ['A', 'C', 'D'],
           ['B', 'C', 'D', 'E'],
           ['A', 'B', 'D'],
           ['A', 'B', 'C', 'E'],
           ['A', 'C'],
           ['B', 'C', 'D'],
           ['B', 'D'],
           ['A', 'C', 'E'],
           ['B', 'D']]

te = TransactionEncoder()
te_ary = te.fit(dataset).transform(dataset)
df = pd.DataFrame(te_ary, columns=te.columns_)
frequent_itemsets = apriori(df, min_support=0.2, use_colnames=True)
association_rules(frequent_itemsets, metric="confidence", min_threshold=0.8)
```

Python mlxtend Association Rules

```
1 # !pip install mlxtend
2 import pandas as pd
3 from mlxtend.preprocessing import TransactionEncoder
4 from mlxtend.frequent_patterns import apriori
5 from mlxtend.frequent_patterns import association_rules
6 dataset = [['A', 'B', 'D'],
7             ['A', 'C', 'D'],
8             ['B', 'C', 'D', 'E'],
9             ['A', 'B', 'D'],
10            ['A', 'B', 'C', 'E'],
11            ['A', 'C'],
12            ['B', 'C', 'D'],
13            ['B', 'D'],
14            ['A', 'C', 'E'],
15            ['B', 'D']]
16 te = TransactionEncoder()
17 te_ary = te.fit(dataset).transform(dataset)
18 df = pd.DataFrame(te_ary, columns=te.columns_)
19 frequent_itemsets = apriori(df, min_support=0.2, use_colnames=True)
20 rules = association_rules(frequent_itemsets, metric="confidence", min_threshold=0.8)
21 rules
```

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction
0	(B)	(D)	0.7	0.7	0.6	0.857143	1.224490	0.11	2.1
1	(D)	(B)	0.7	0.7	0.6	0.857143	1.224490	0.11	2.1
2	(E)	(C)	0.3	0.6	0.3	1.000000	1.666667	0.12	inf
3	(E, A)	(C)	0.2	0.6	0.2	1.000000	1.666667	0.08	inf
4	(E, B)	(C)	0.2	0.6	0.2	1.000000	1.666667	0.08	inf

Python in Google Colab (Python101)

<https://colab.research.google.com/drive/1FEG6DnGvwfUbeo4zJ1zTunjMqf2RkCrT>

The screenshot shows a Google Colab interface with the following details:

- Title:** python101.ipynb
- Toolbar:** File, Edit, View, Insert, Runtime, Tools, Help, Saving...
- Table of Contents:** Python Pandas, Python Data Visualization, Unsupervised Learning: Association Analysis, Market Basket Analysis, Association Rules Generation from Frequent Itemsets (highlighted), Market Basket Analysis, Machine Learning with scikit-learn, Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python.
- Section:** Association Rules Generation from Frequent Itemsets
- Code:**

```
1 # !pip install mlxtend
2 import pandas as pd
3 from mlxtend.preprocessing import TransactionEncoder
4 from mlxtend.frequent_patterns import apriori
5 from mlxtend.frequent_patterns import association_rules
6 dataset = [['Milk', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
7             ['Dill', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
8             ['Milk', 'Apple', 'Kidney Beans', 'Eggs'],
9             ['Milk', 'Unicorn', 'Corn', 'Kidney Beans', 'Yogurt'],
10            ['Corn', 'Onion', 'Unicorn', 'Kidney Beans', 'Ice cream', 'Eggs']]
11 te = TransactionEncoder()
12 te_ary = te.fit(dataset).transform(dataset)
13 df = pd.DataFrame(te_ary, columns=te.columns_)
14 frequent_itemsets = apriori(df, min_support=0.6, use_colnames=True)
15 frequent_itemsets
```
- Output:** A table showing frequent itemsets with support values:

	support	itemsets
0	0.8	(Eggs)
1	1.0	(Kidney Beans)
2	0.6	(Milk)
3	0.6	(Onion)
4	0.6	(Yogurt)

<https://tinyurl.com/aintpuppython101>

```
# ! pip install mlxtend
from mlxtend.frequent_patterns import apriori
from mlxtend.frequent_patterns import association_rules

frequent_itemsets = apriori(df, min_support=0.6,
use_colnames=True)
```

```
1 # ! pip install mlxtend
2 import pandas as pd
3 from mlxtend.preprocessing import TransactionEncoder
4 from mlxtend.frequent_patterns import apriori
5 from mlxtend.frequent_patterns import association_rules
6
7 dataset = [['Milk', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
8             ['Dill', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
9             ['Milk', 'Apple', 'Kidney Beans', 'Eggs'],
10            ['Milk', 'Unicorn', 'Corn', 'Kidney Beans', 'Yogurt'],
11            ['Corn', 'Onion', 'Onion', 'Kidney Beans', 'Ice cream', 'Eggs']]
12
13 te = TransactionEncoder()
14 te_ary = te.fit(dataset).transform(dataset)
15 df = pd.DataFrame(te_ary, columns=te.columns_)
16 frequent_itemsets = apriori(df, min_support=0.6, use_colnames=True)
17
18 frequent_itemsets
```

```
# ! pip install mlxtend
import pandas as pd
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori
from mlxtend.frequent_patterns import association_rules

dataset = [['Milk', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
           ['Dill', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],
           ['Milk', 'Apple', 'Kidney Beans', 'Eggs'],
           ['Milk', 'Unicorn', 'Corn', 'Kidney Beans', 'Yogurt'],
           ['Corn', 'Onion', 'Onion', 'Kidney Beans', 'Ice cream', 'Eggs']]

te = TransactionEncoder()
te_ary = te.fit(dataset).transform(dataset)
df = pd.DataFrame(te_ary, columns=te.columns_)
frequent_itemsets = apriori(df, min_support=0.6,
                             use_colnames=True)

frequent_itemsets
```

```
frequent_itemsets = apriori(df,  
min_support=0.6,  
use_colnames=True)
```

	support	itemsets
0	0.8	(Eggs)
1	1.0	(Kidney Beans)
2	0.6	(Milk)
3	0.6	(Onion)
4	0.6	(Yogurt)
5	0.8	(Eggs, Kidney Beans)
6	0.6	(Onion, Eggs)
7	0.6	(Milk, Kidney Beans)
8	0.6	(Onion, Kidney Beans)
9	0.6	(Yogurt, Kidney Beans)
10	0.6	(Onion, Eggs, Kidney Beans)

```
association_rules(frequent_itemsets,  
metric="confidence", min_threshold=0.7)
```

1 | association_rules(frequent_itemsets, metric="confidence", min_threshold=0.7.)

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction
0	(Eggs)	(Kidney Beans)	0.8	1.0	0.8	1.00	1.00	0.00	inf
1	(Kidney Beans)	(Eggs)	1.0	0.8	0.8	0.80	1.00	0.00	1.000000
2	(Onion)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf
3	(Eggs)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000
4	(Milk)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
5	(Onion)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
6	(Yogurt)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
7	(Onion, Eggs)	(Kidney Beans)	0.6	1.0	0.6	1.00	1.00	0.00	inf
8	(Onion, Kidney Beans)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf
9	(Eggs, Kidney Beans)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000
10	(Onion)	(Eggs, Kidney Beans)	0.6	0.8	0.6	1.00	1.25	0.12	inf
11	(Eggs)	(Onion, Kidney Beans)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000

```
rules =  
association_rules(frequent_itemsets,  
metric="lift", min_threshold=1.2)  
rules
```

```
1 rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1.2)  
2 rules
```

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction
0	(Onion)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf
1	(Eggs)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000
2	(Onion, Kidney Beans)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf
3	(Eggs, Kidney Beans)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000
4	(Onion) (Eggs, Kidney Beans)		0.6	0.8	0.6	1.00	1.25	0.12	inf
5	(Eggs) (Onion, Kidney Beans)		0.8	0.6	0.6	0.75	1.25	0.12	1.600000

```
rules["antecedent_len"] =  
rules["antecedents"].apply(lambda x: len(x))  
rules
```

The screenshot shows a Jupyter Notebook cell with the following content:

```
1 | rules["antecedent_len"] = rules["antecedents"].apply(lambda x: len(x))  
2 | rules
```

Below the code, the resulting DataFrame is displayed:

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction	antecedent_len
0	(Onion)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf	1
1	(Eggs)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000	1
2	(Onion, Kidney Beans)	(Eggs)	0.6	0.8	0.6	1.00	1.25	0.12	inf	2
3	(Eggs, Kidney Beans)	(Onion)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000	2
4	(Onion)	(Eggs, Kidney Beans)	0.6	0.8	0.6	1.00	1.25	0.12	inf	1
5	(Eggs)	(Onion, Kidney Beans)	0.8	0.6	0.6	0.75	1.25	0.12	1.600000	1

```
rules[ (rules['antecedent_len'] >= 2) &
      (rules['confidence'] > 0.75) &
      (rules['lift'] > 1.2) ]
```

The screenshot shows a Jupyter Notebook cell with the following code:

```
1 rules[ (rules['antecedent_len'] >= 2) &
2     (rules['confidence'] > 0.75) &
3     (rules['lift'] > 1.2) ]
```

Below the code, a Pandas DataFrame is displayed:

	antecedents	consequents	antecedent support	consequent support	support	confidence	lift	leverage	conviction	antecedent_len
2	(Onion, Kidney Beans)	(Eggs)	0.6	0.8	0.6	1.0	1.25	0.12	inf	2

```
rules[rules['antecedents'] ==  
{'Eggs', 'Kidney Beans'}]
```

```
1 | rules[rules['antecedents'] == {'Eggs', 'Kidney Beans'}] :  
2 |  
3 |   antecedents  consequents      antecedent support  consequent support  support  confidence  lift  leverage  conviction  antecedent_len  
3 |   (Eggs, Kidney Beans)  (Onion)          0.8            0.6           0.6       0.75     1.25    0.12        1.6             2
```

The Theory of Learning

The Theory of Learning

- Computational Learning Theory
 - Probably Approximately Correct (PAC) Learning
 - Vapnik-Chervonenkis (VC) Dimension
 - Bias-Variance Trade-off
- Overfitting and Underfitting
 - Avoid overfitting:
Regularization, Cross-Validation

What is Learning in Machine Learning?

- Learning involves finding patterns from data to make predictions.
- Performance is measured through generalization to unseen data.
- Three paradigms of learning:
 - Supervised Learning (e.g., classification, regression)
 - Unsupervised Learning (e.g., clustering, dimensionality reduction)
 - Reinforcement Learning (e.g., learning through rewards)

The Theory of Learning

- How can we be sure that our learned hypothesis will predict well for previously unseen inputs?
 - How do we know that the hypothesis h is close to the target function f if we don't know what is?
- How many examples do we need to get a good h ?
- What hypothesis space should we use?
- If the hypothesis space is very complex, can we even find the best h or do we have to settle for a local maximum?
- How complex should h be?
- How do we avoid overfitting?

Computational Learning Theory

- Intersection of **AI, statistics, and theoretical computer science.**
- Any **hypothesis** that is seriously wrong will almost certainly be “found out” with high probability after a small number of examples.

Probably Approximately Correct (PAC) Learning

- Any **hypothesis** that is consistent with a sufficiently large set of training examples is unlikely to be seriously wrong.
- **PAC learning algorithm:**
 - Any learning algorithm that returns hypotheses that are probably approximately correct.

Probably Approximately Correct (PAC) Learning

- **PAC Learning** provides a way to understand generalization
- Key Components:
 - Error threshold (ϵ)
 - Confidence level ($1 - \delta$)
- Example: Linear models are PAC-learnable with enough data

Linear function

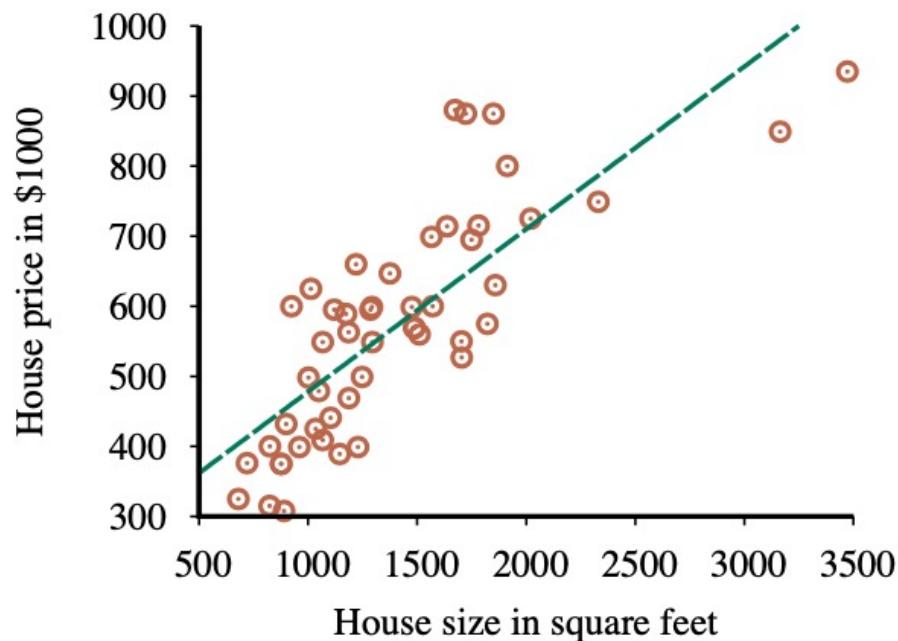
$$y = f(x)$$

$$y = w_1x + w_0$$

$$h_w(x) = w_1x + w_0$$

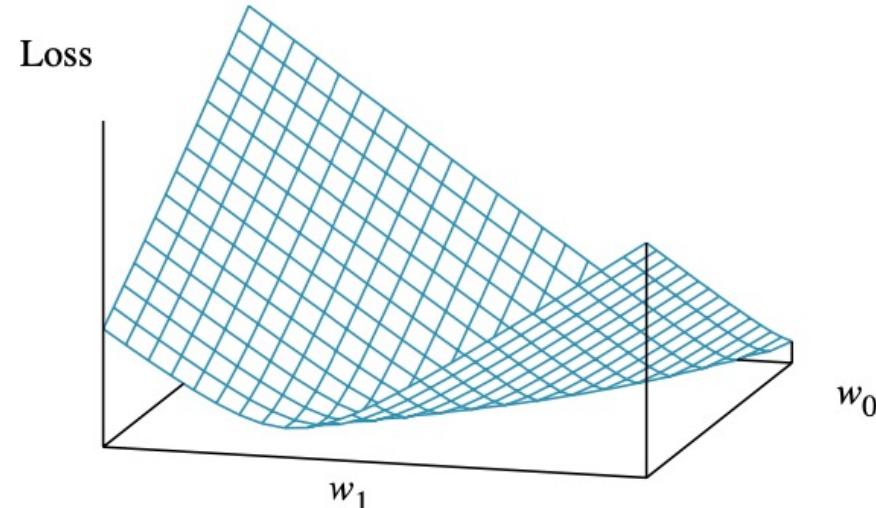
Linear Regression Weight Space

$$h_w(x) = w_1x + w_0$$



$$y = 0.232x + 246$$

$$w^* = \operatorname{argmin}_w \text{Loss}(h_w)$$



Loss function for Weights (w_1, w_0)

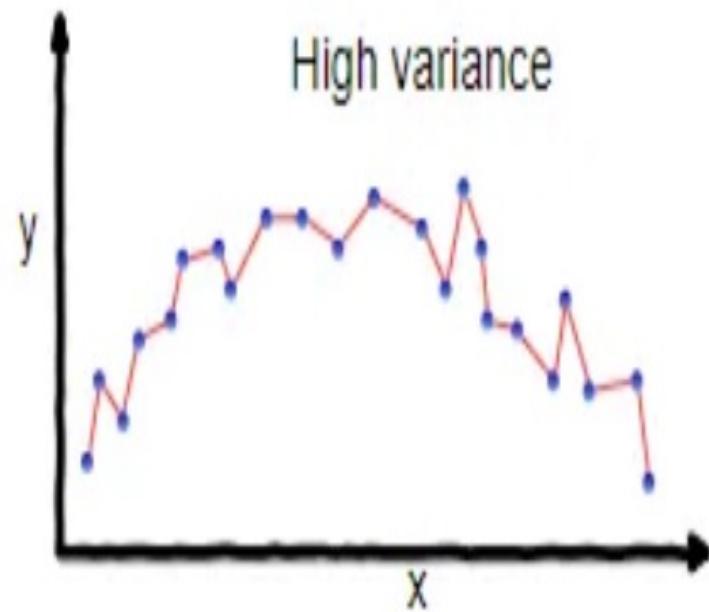
Vapnik-Chervonenkis (VC) Dimension

- VC Dimension measures a model's capacity to fit various patterns
- Higher VC Dimension
Higher model complexity
- Balancing model complexity helps improve generalization

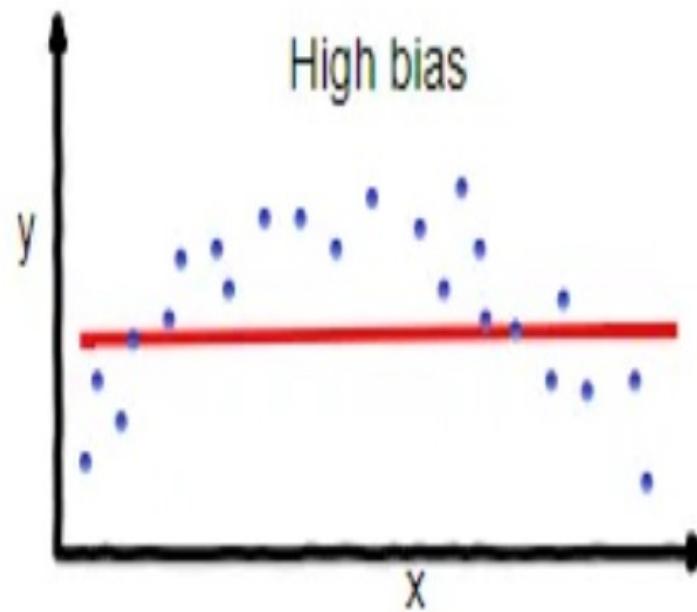
Bias-Variance Trade-off

- **Bias**: Error due to simple models (e.g., linear models)
- **Variance**: Error from models too sensitive to noise (e.g., deep trees)
- Bias-Variance Trade-off:
 - Low bias-high variance: Risk of **overfitting**
 - High bias-low variance: Risk of **underfitting**
- Selecting the right model **balances bias and variance**

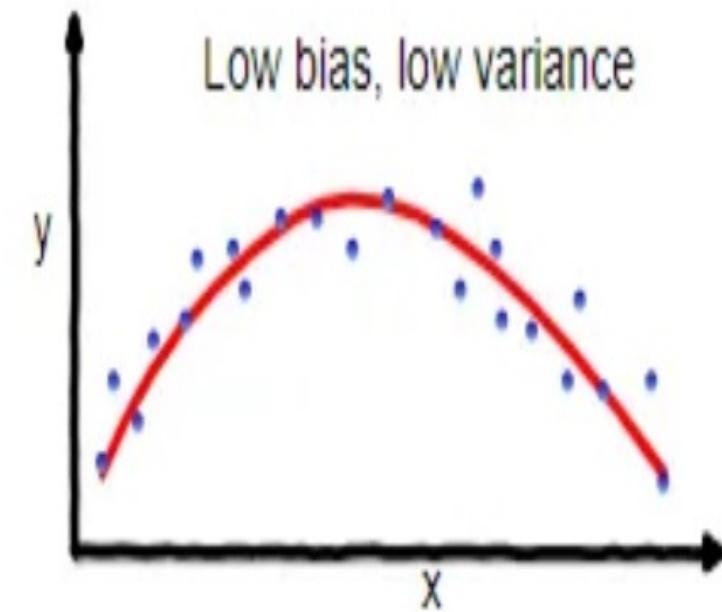
Overfitting vs. Underfitting



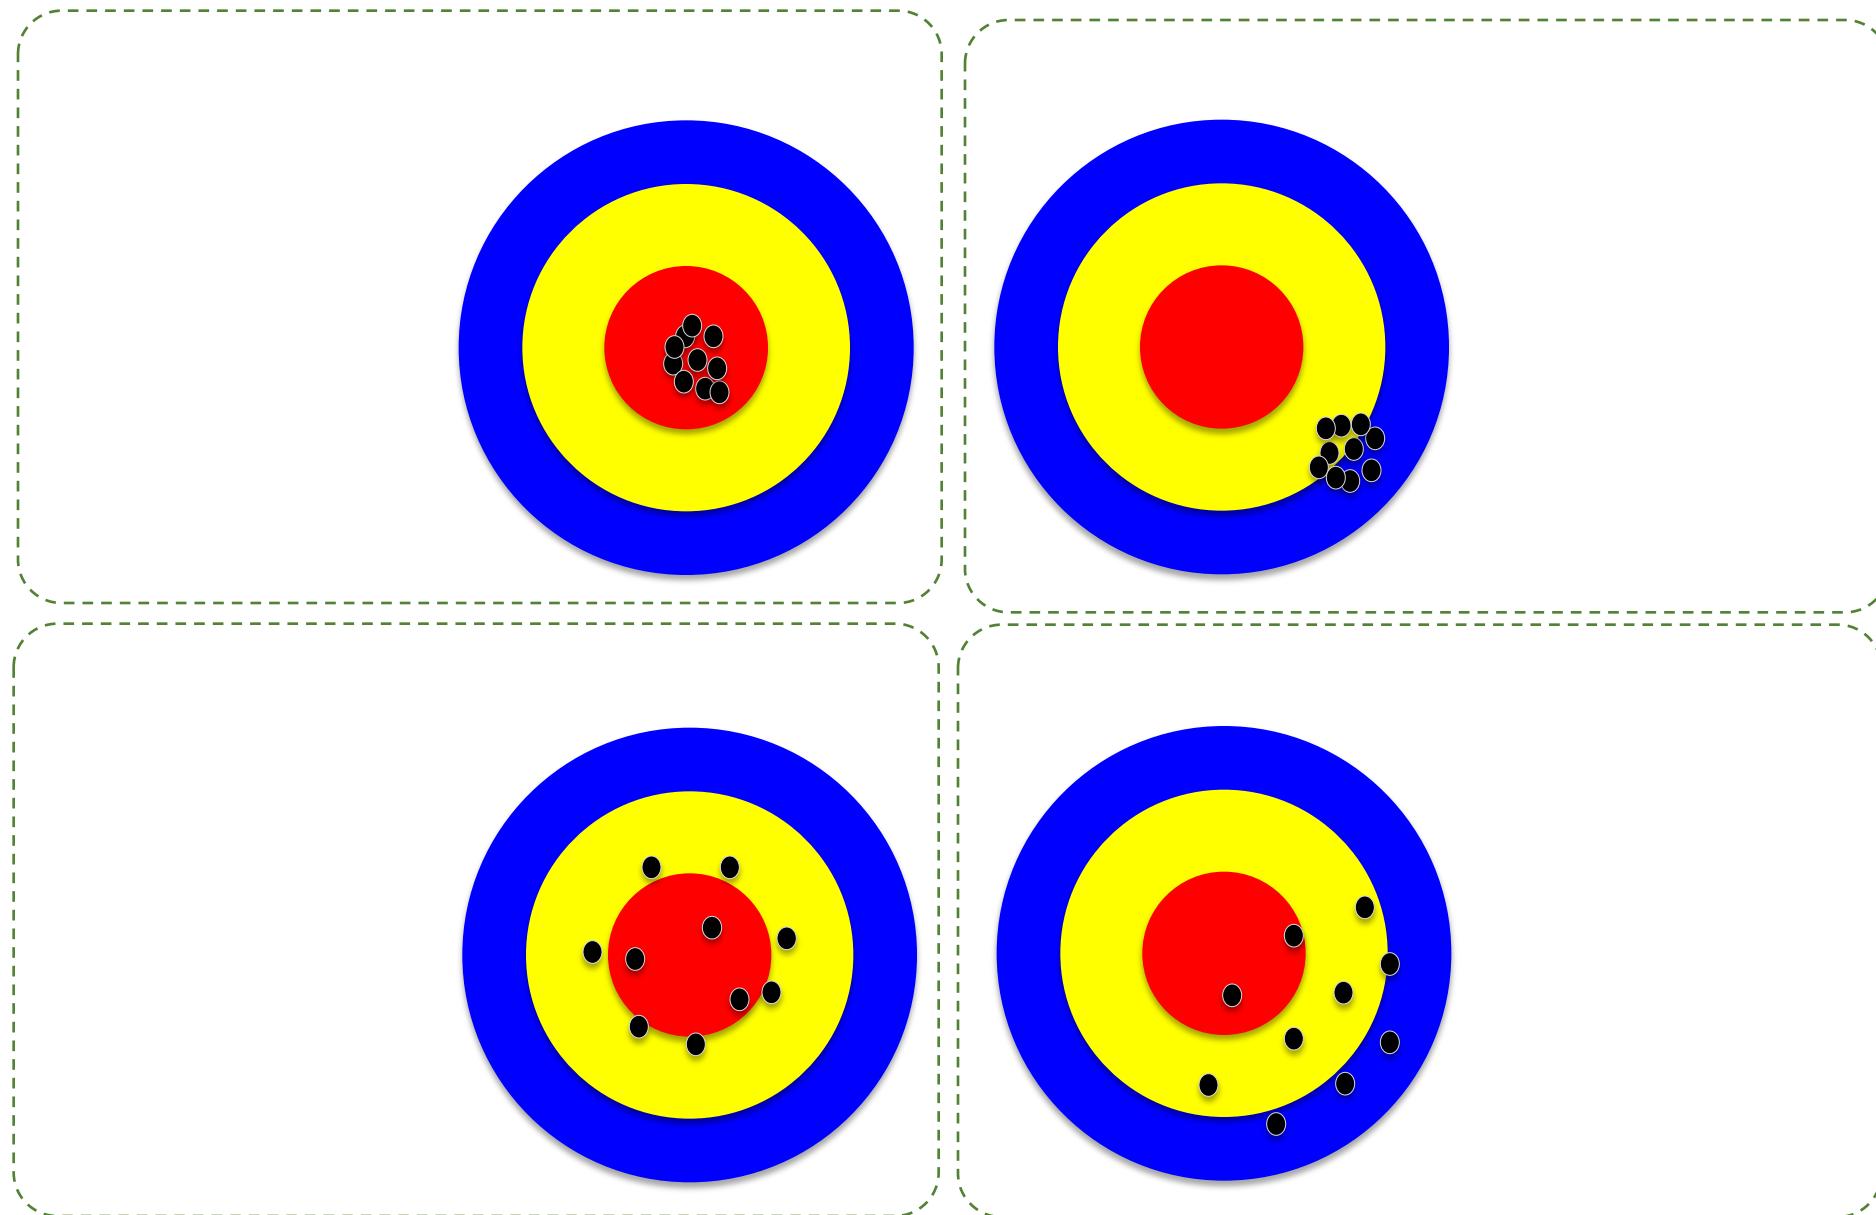
Overfitting



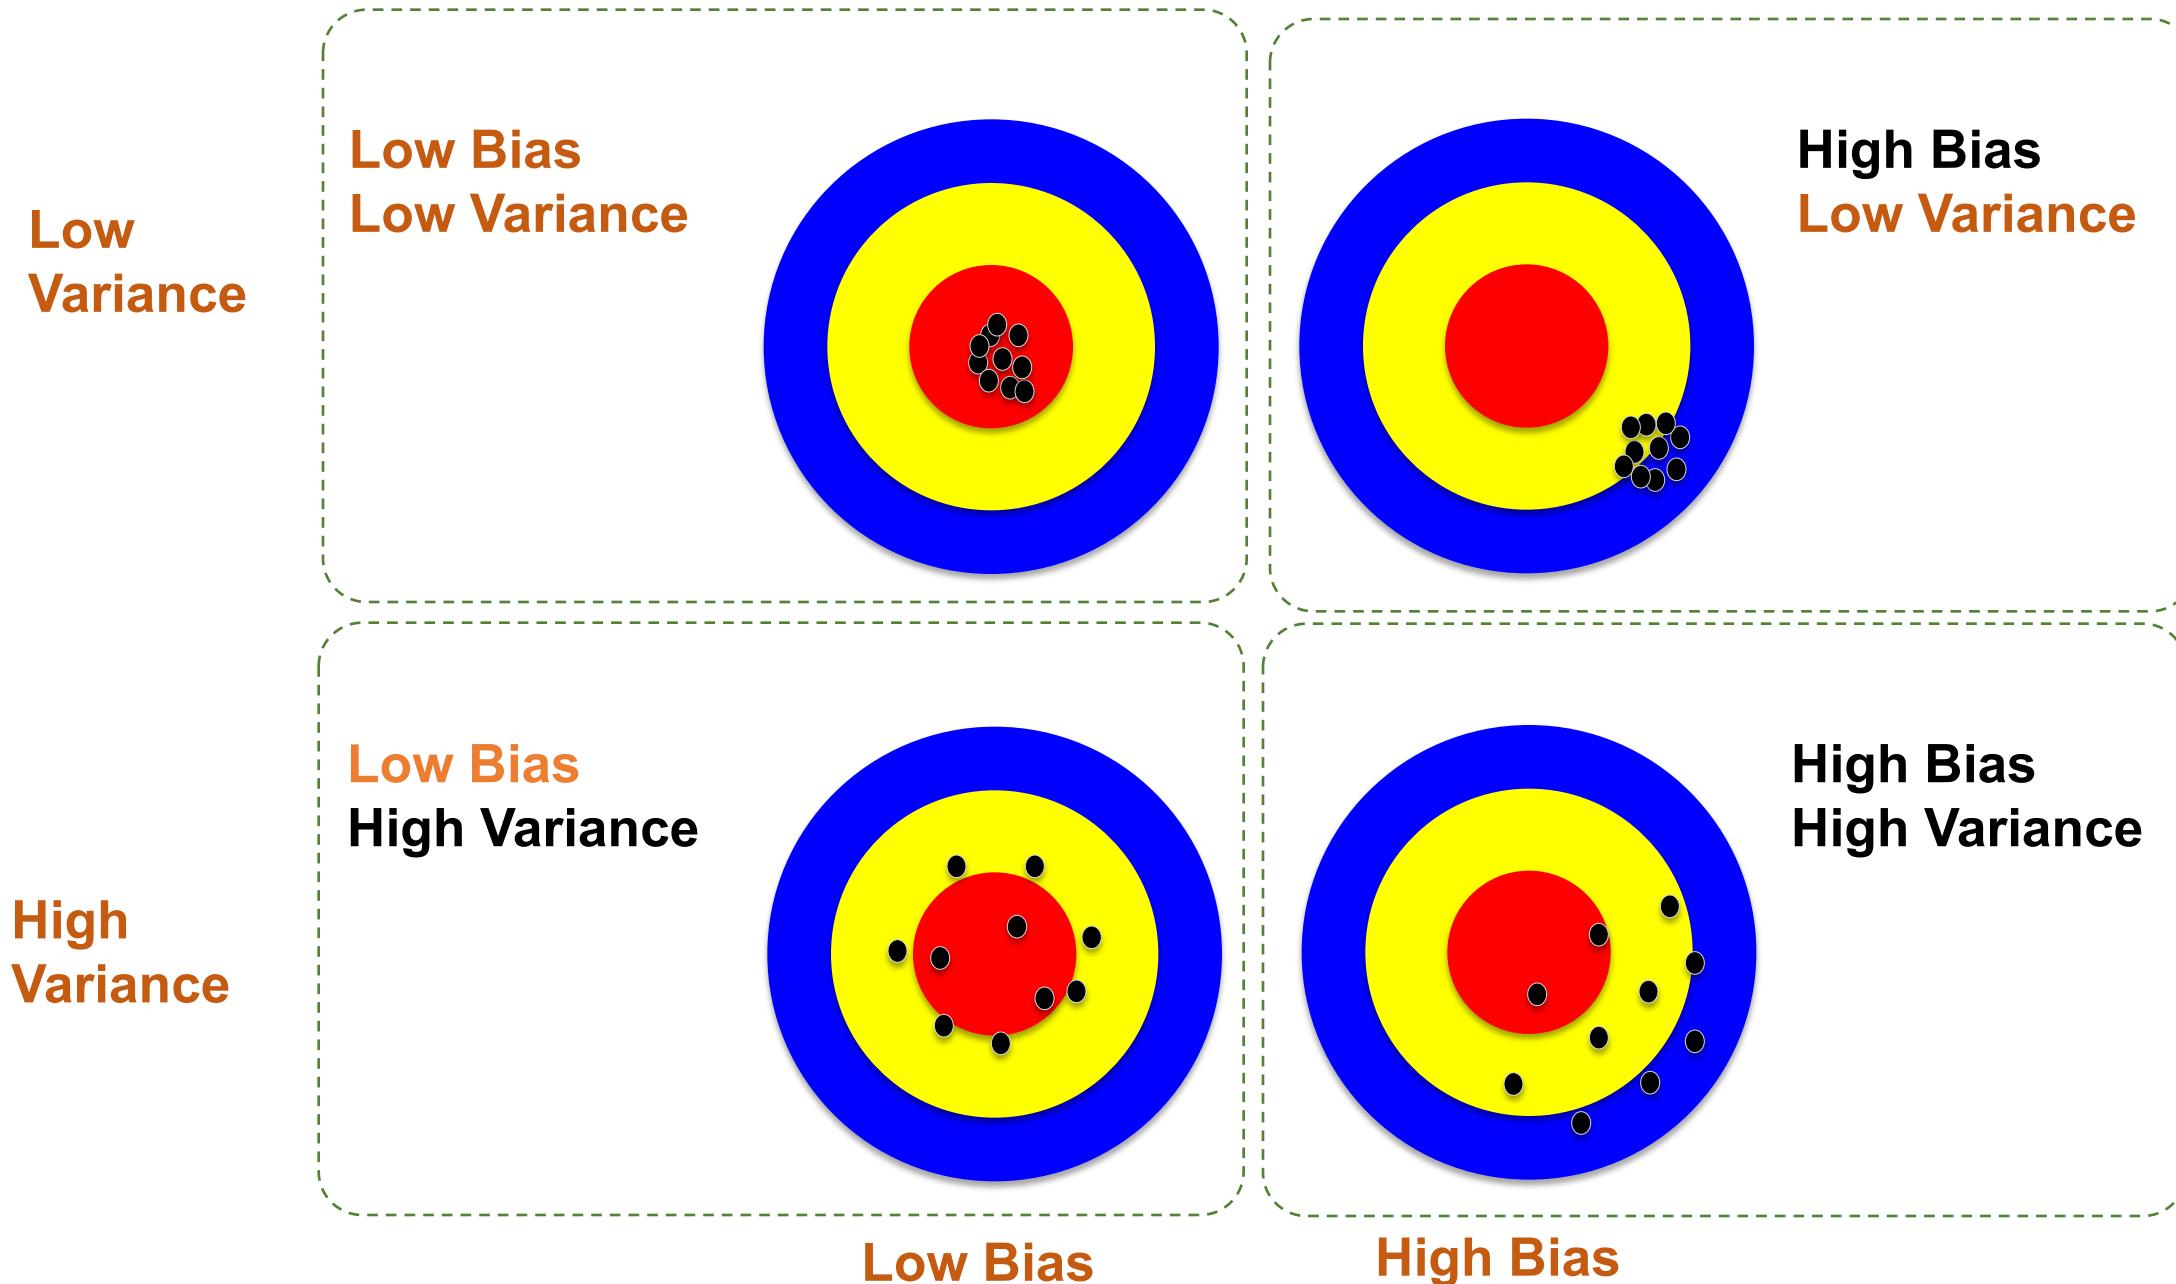
Underfitting



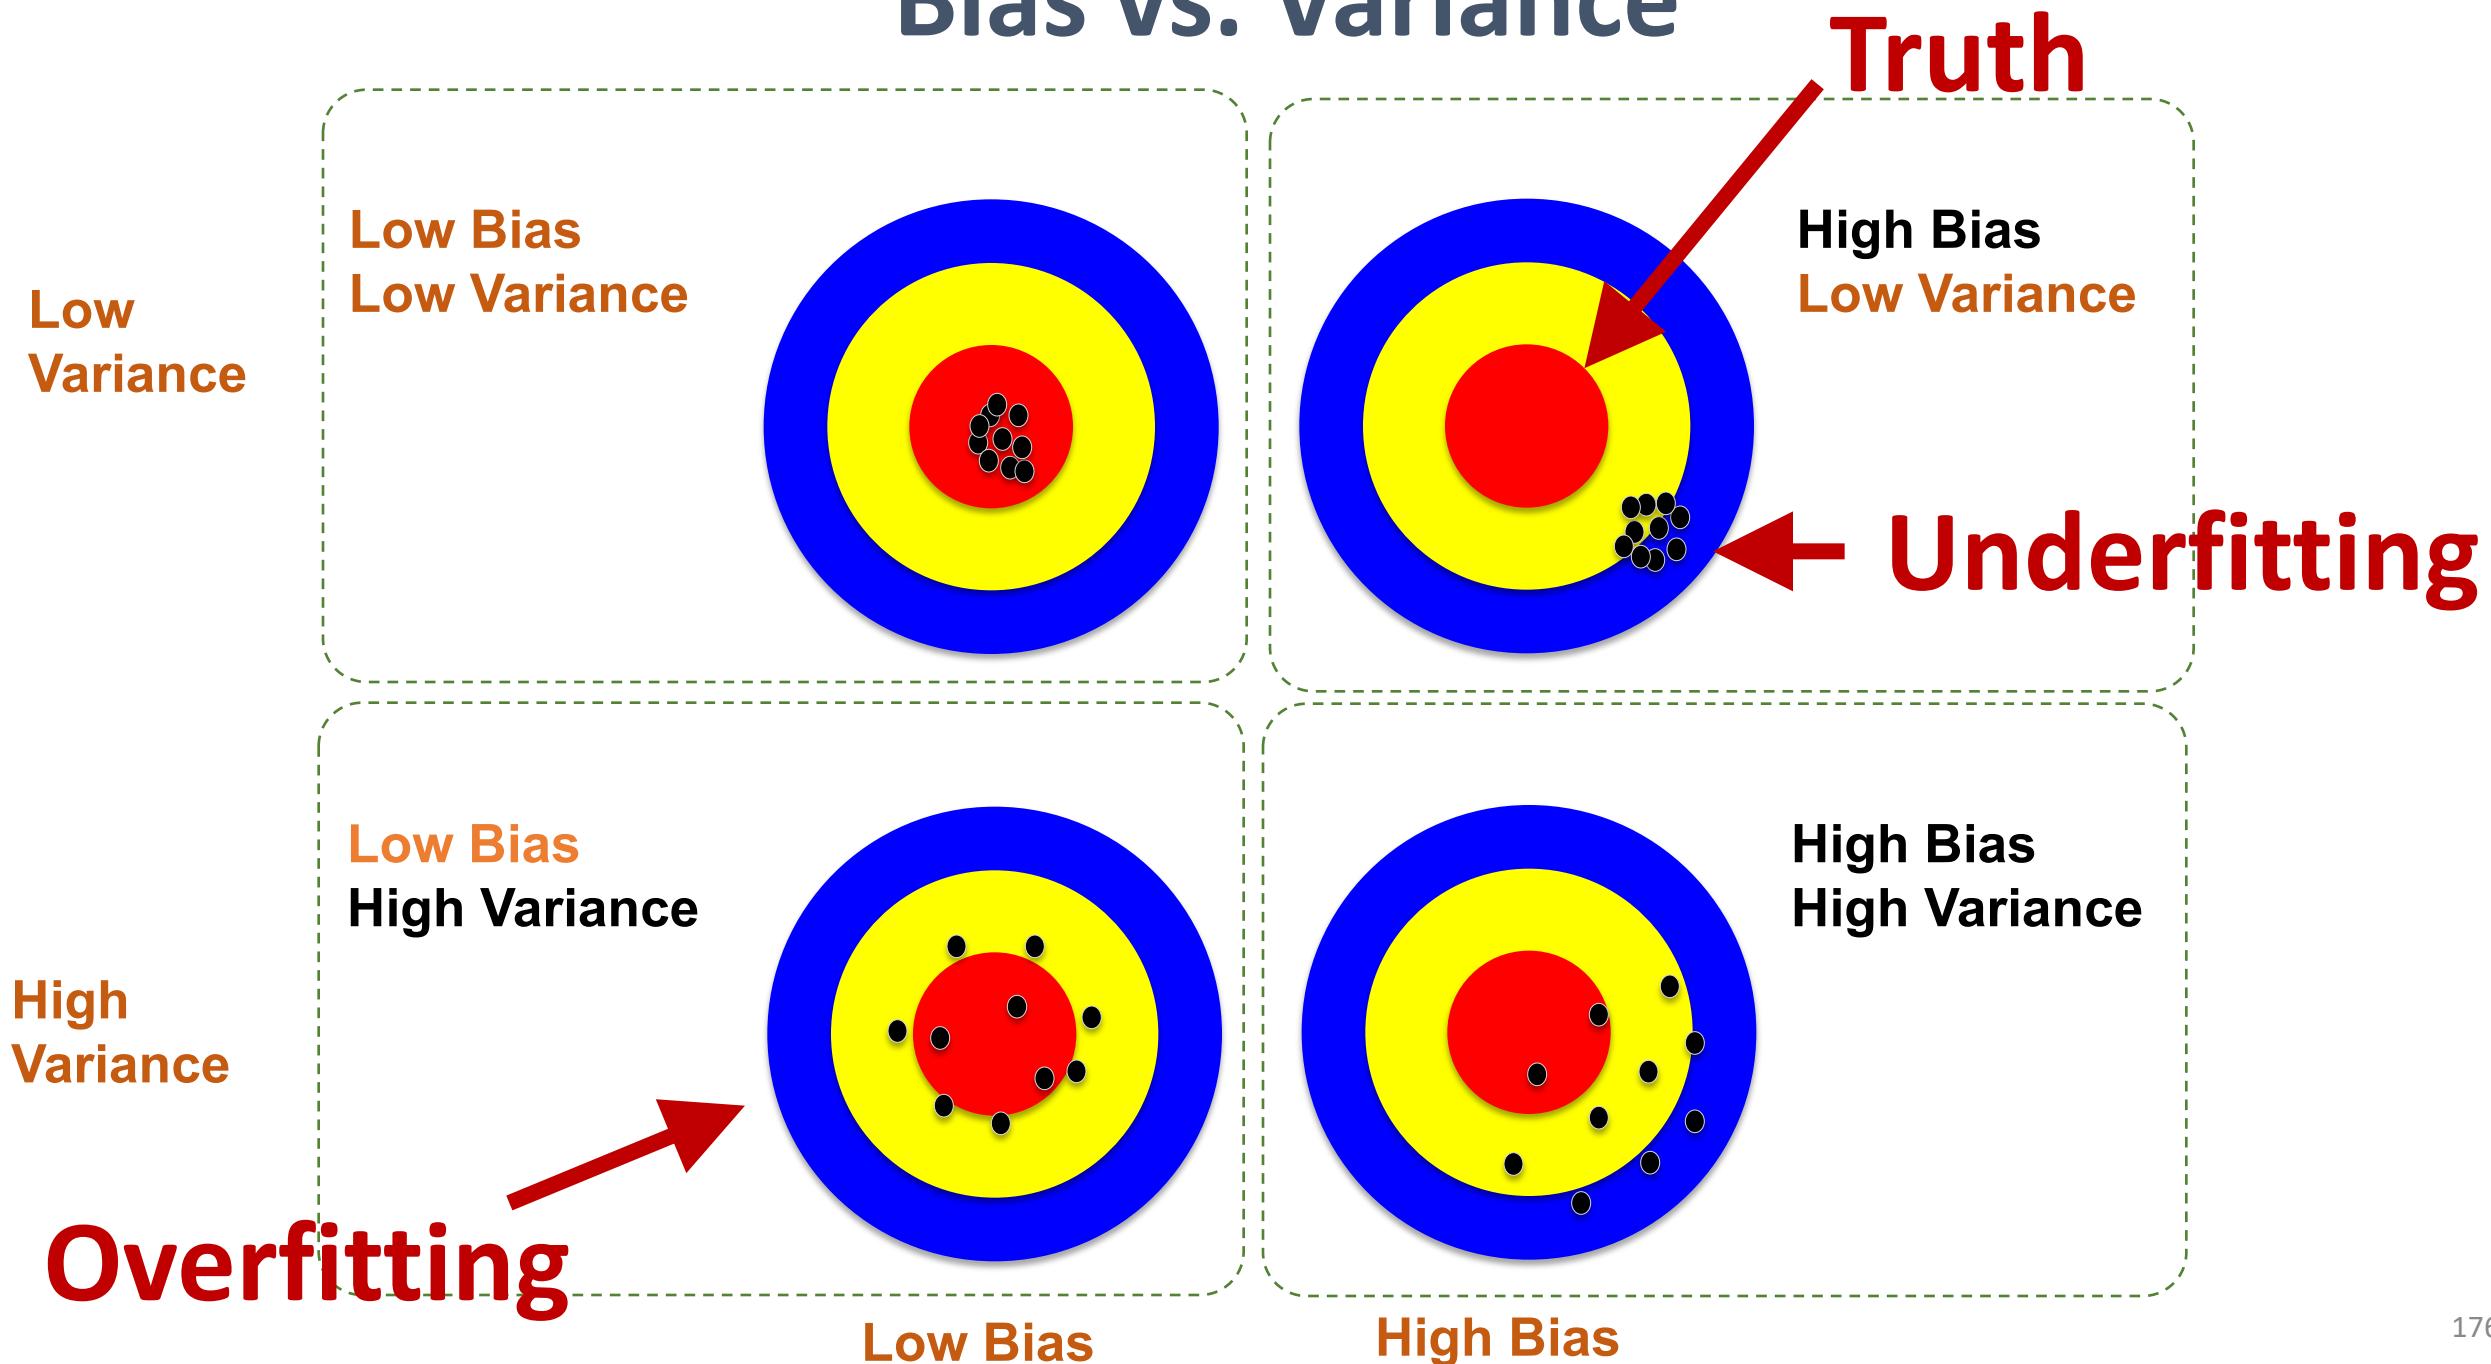
Good Balance



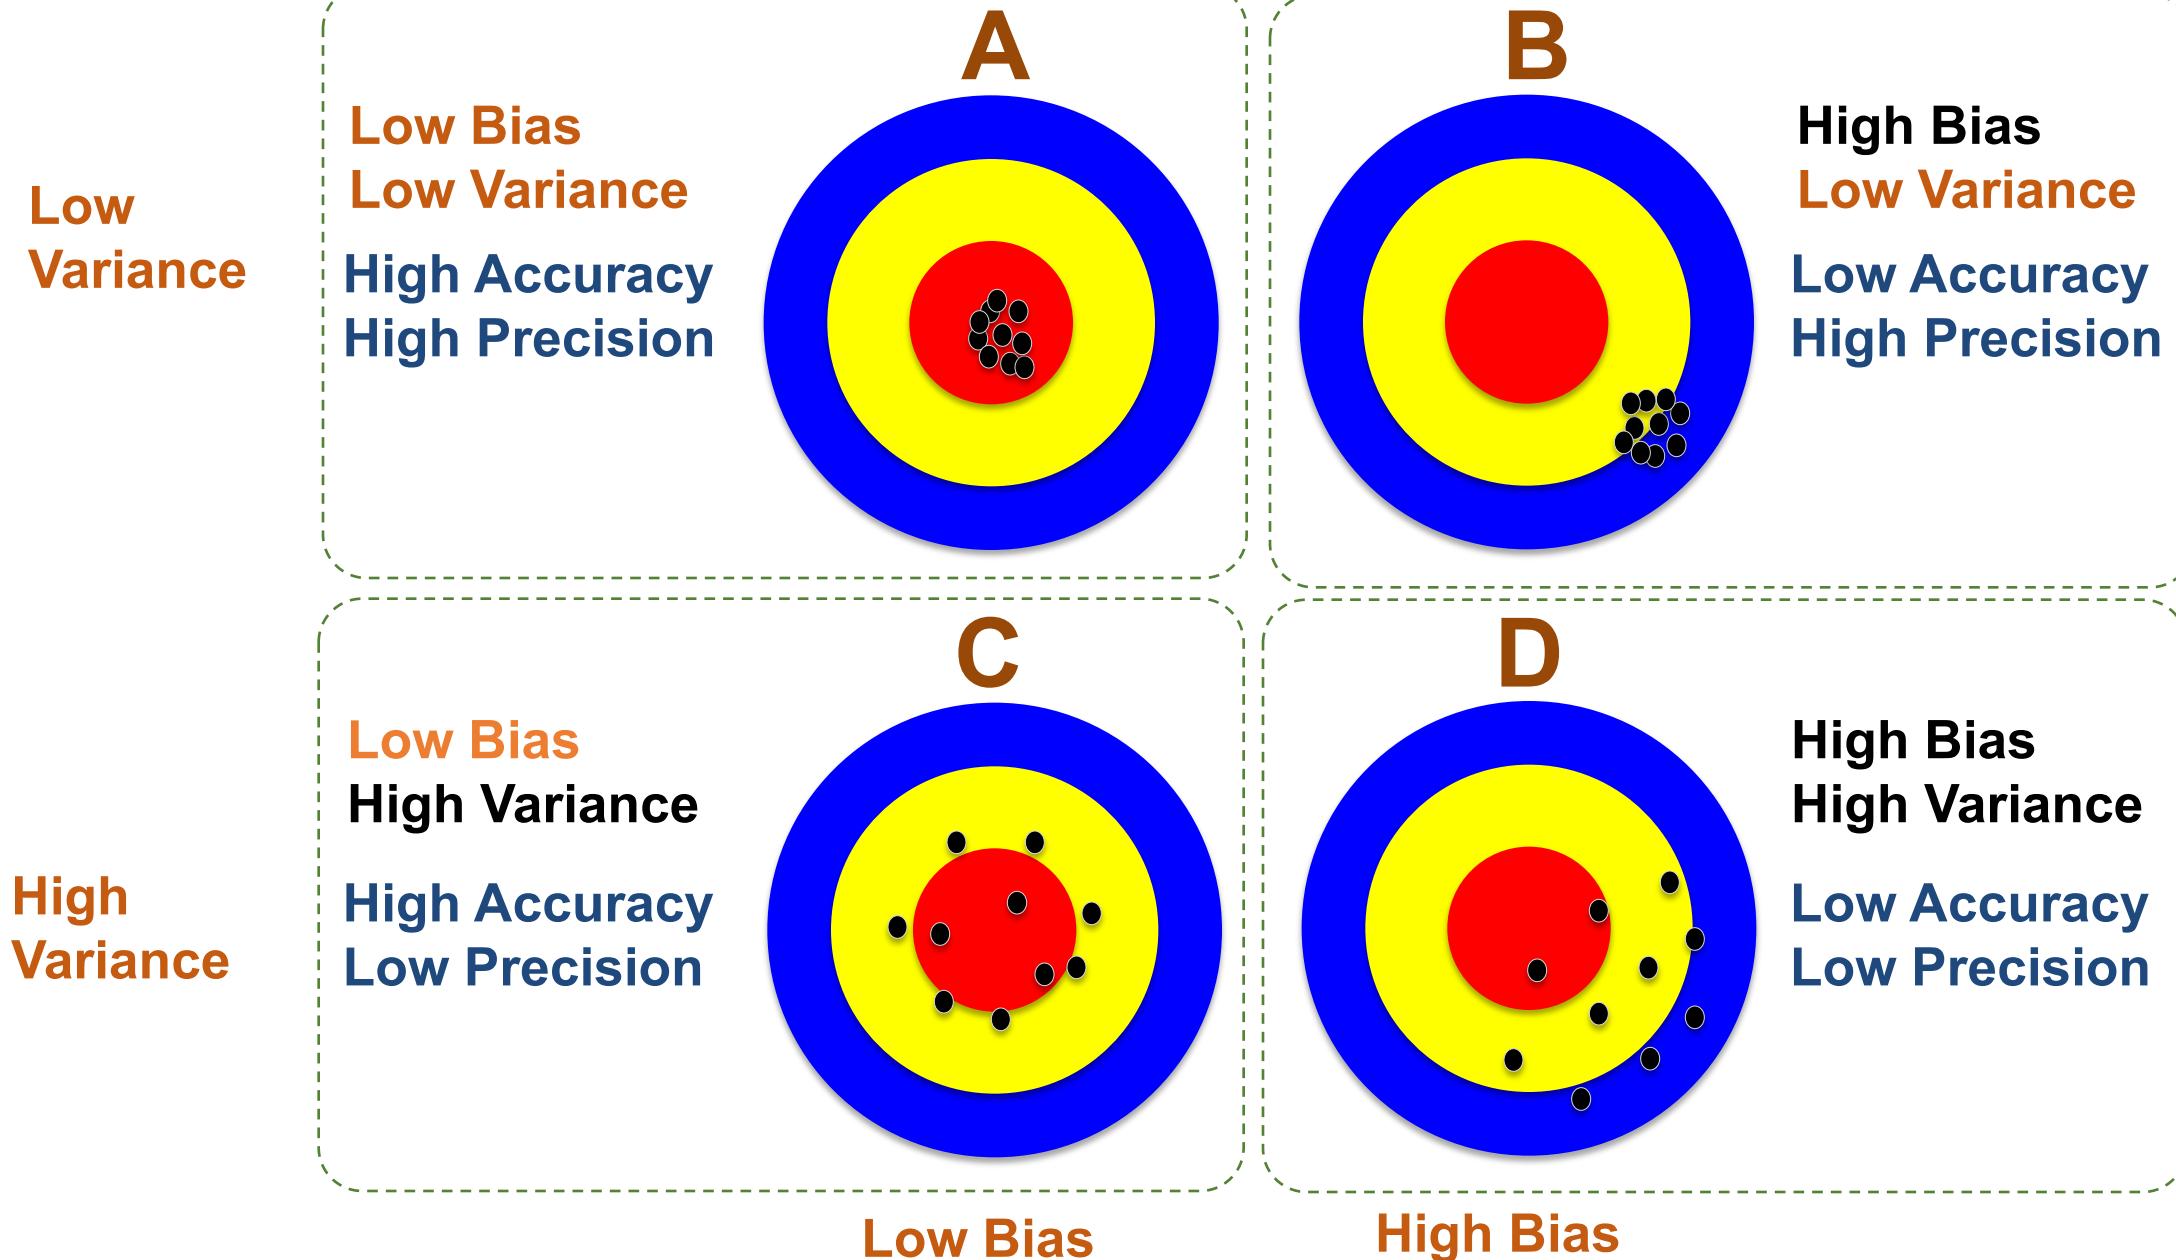
Bias vs. Variance



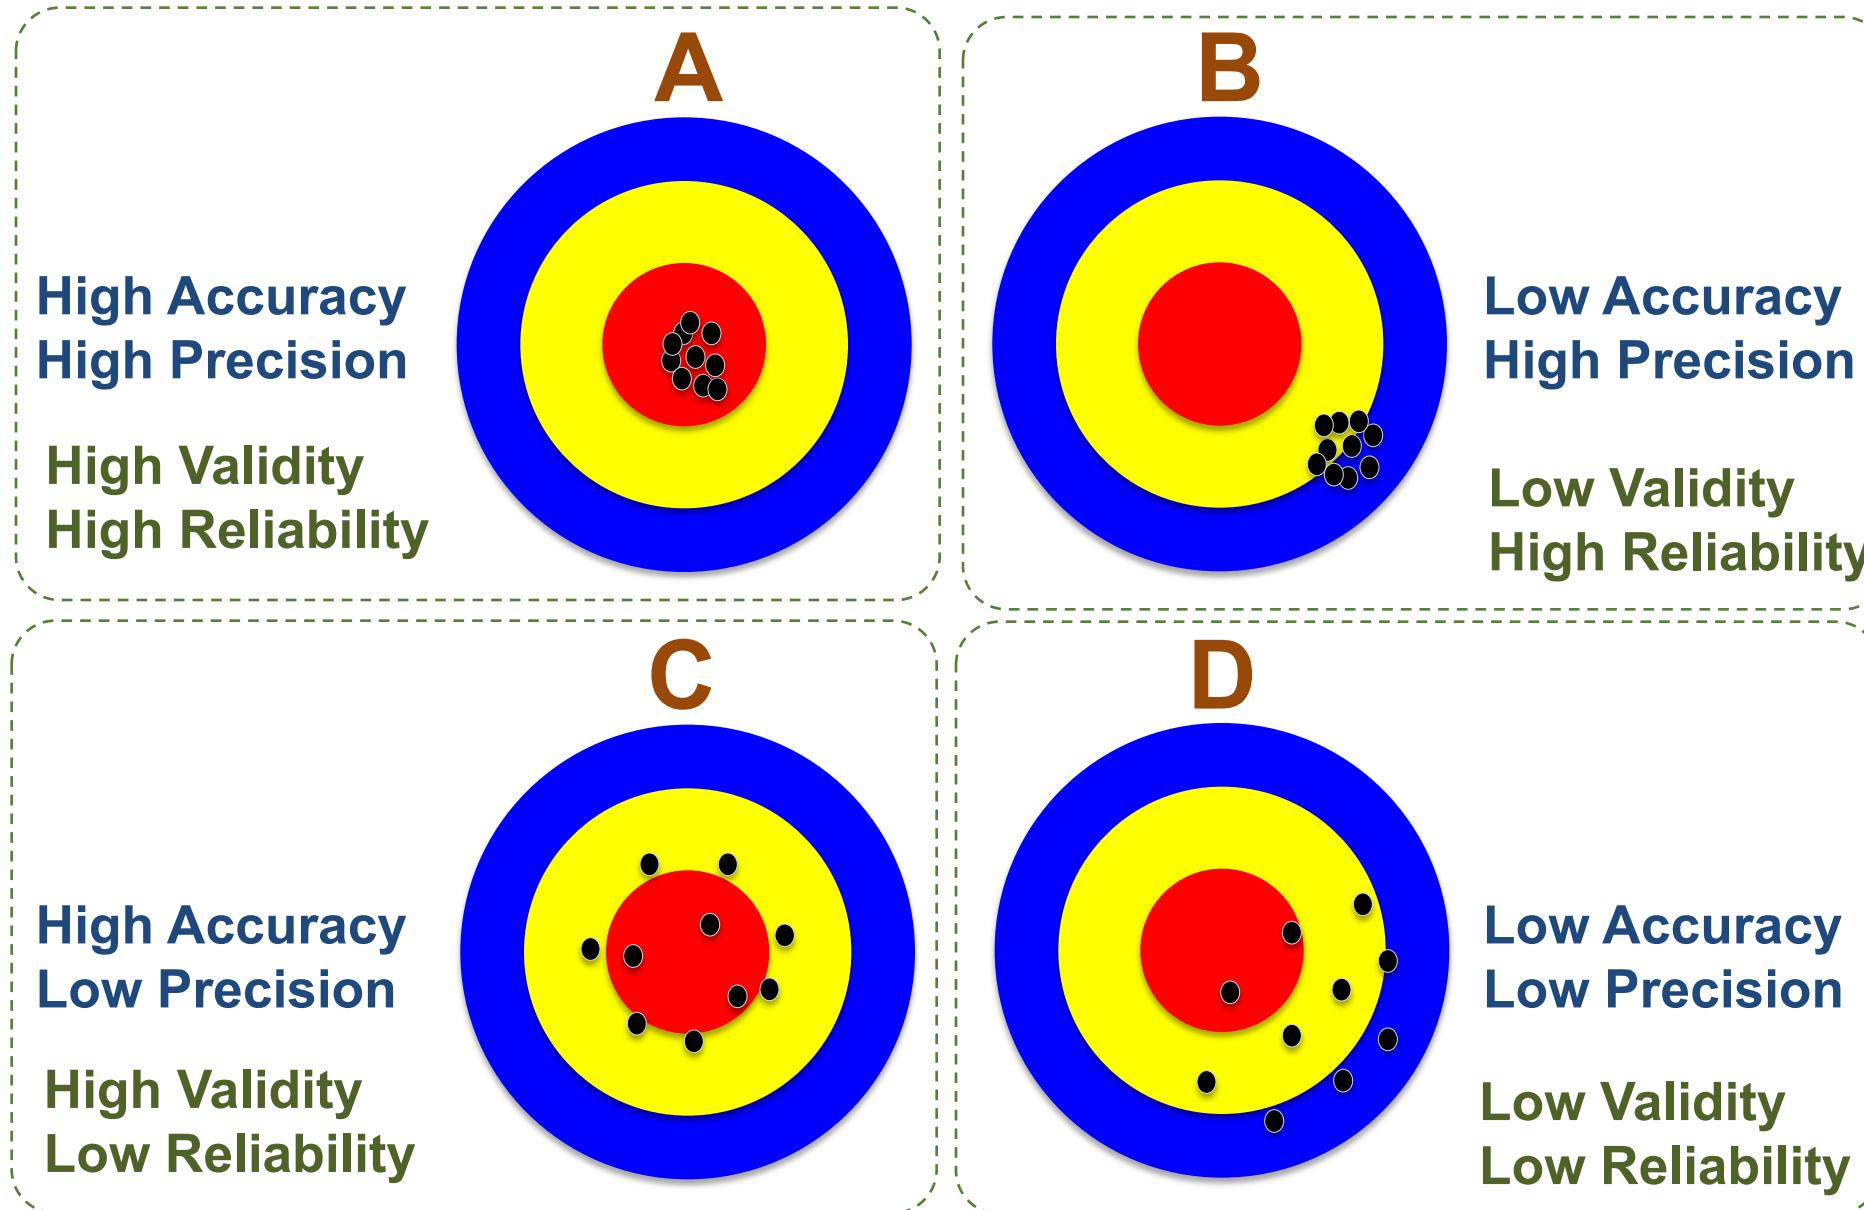
Bias vs. Variance



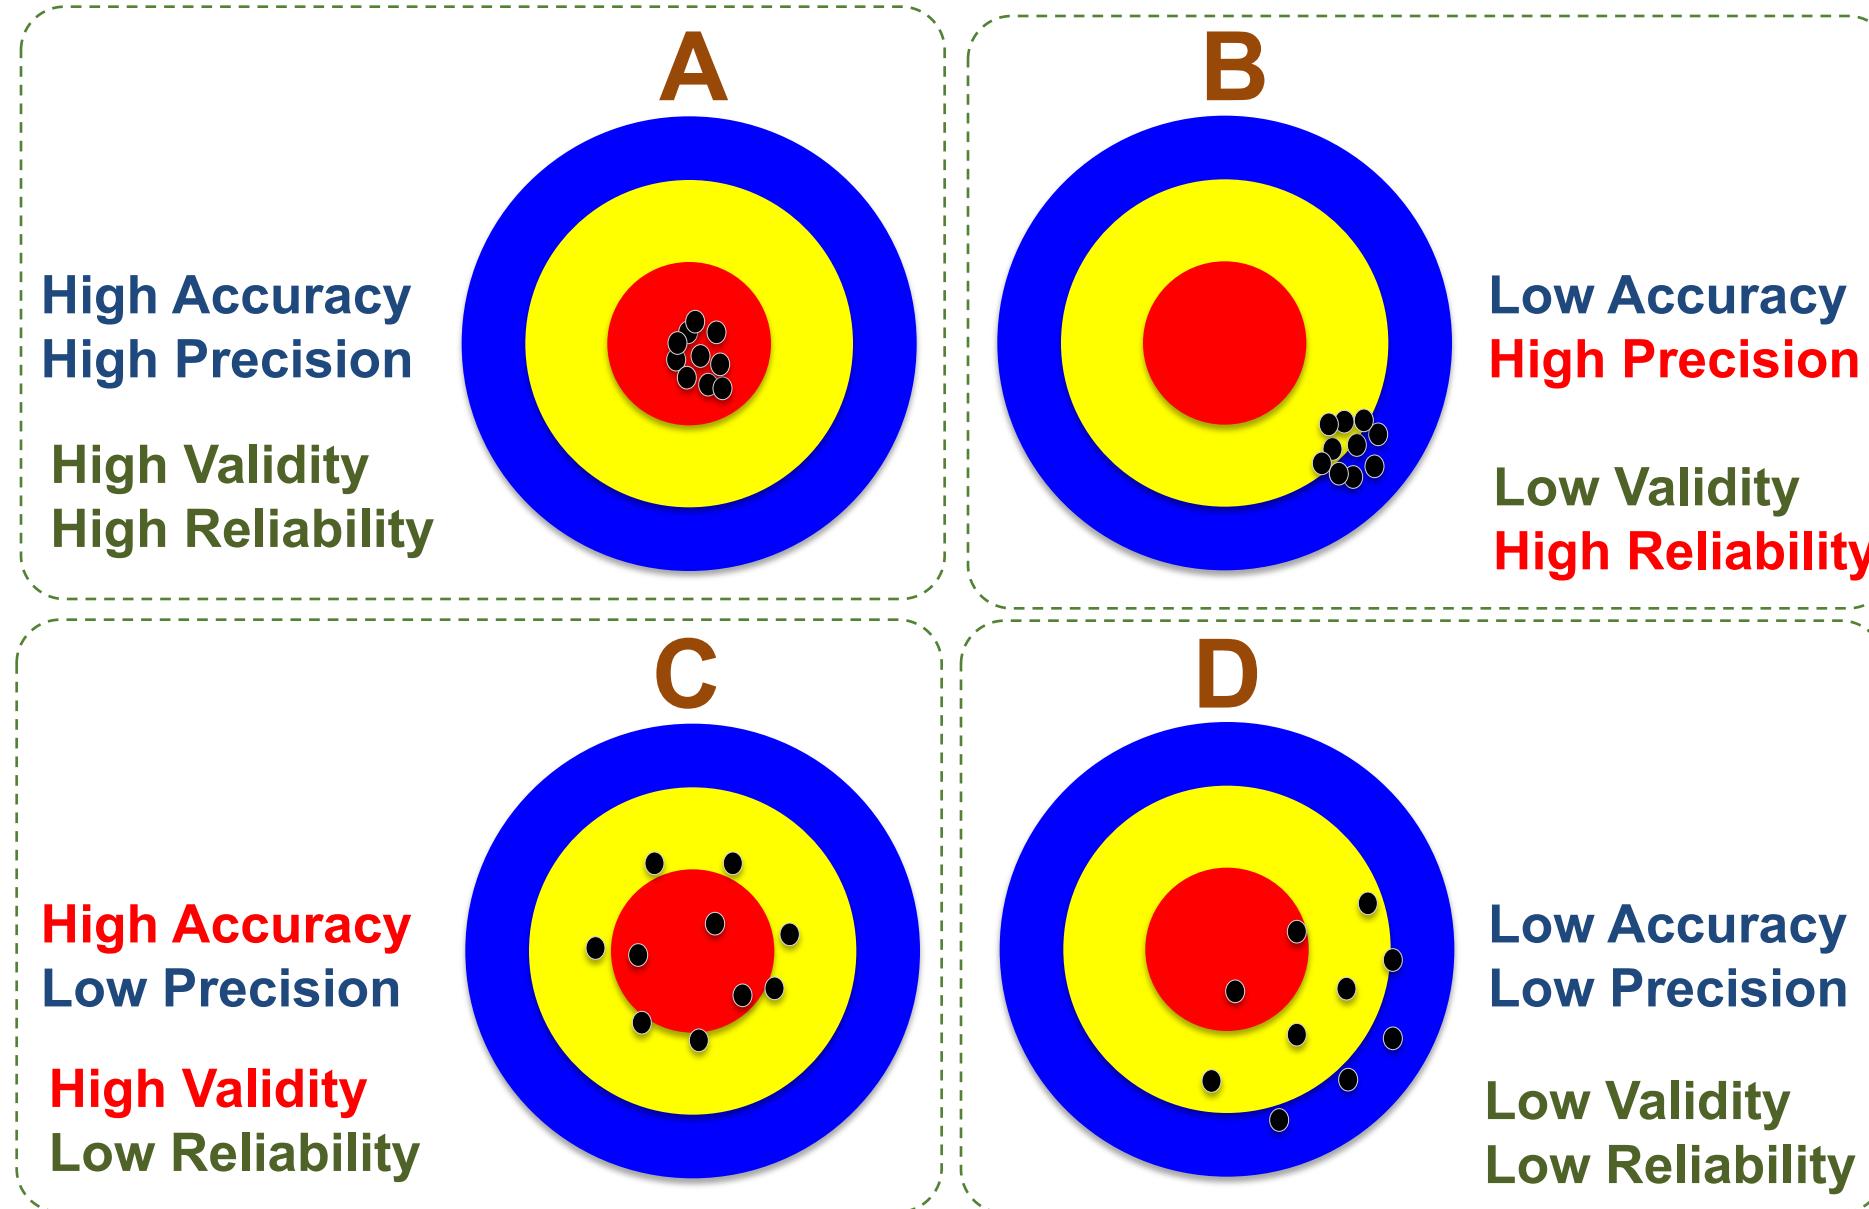
Bias vs. Variance



Accuracy vs. Precision



Accuracy (Validity) vs. Precision (Reliability)



Learning Theory

- **PAC Learning:**
 - Focuses on **generalization**
- **VC Dimension:**
 - Measures model capacity and impacts **generalization**
- **Bias-Variance Trade-off:**
 - Helps in **model selection and tuning**

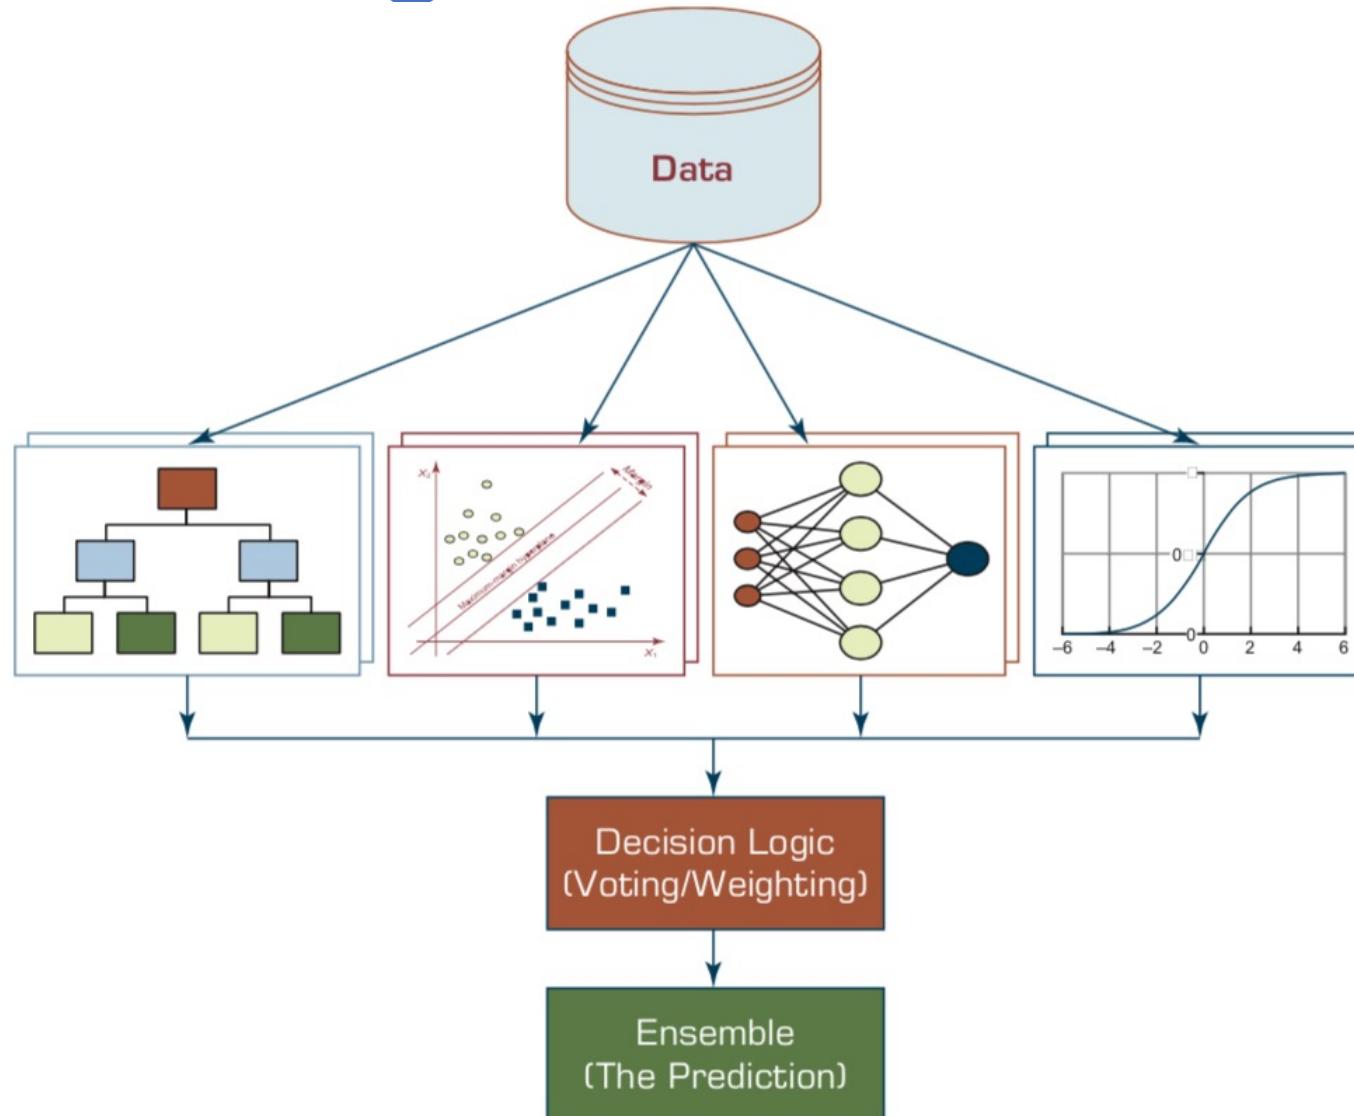
Ensemble Learning

Ensemble Learning

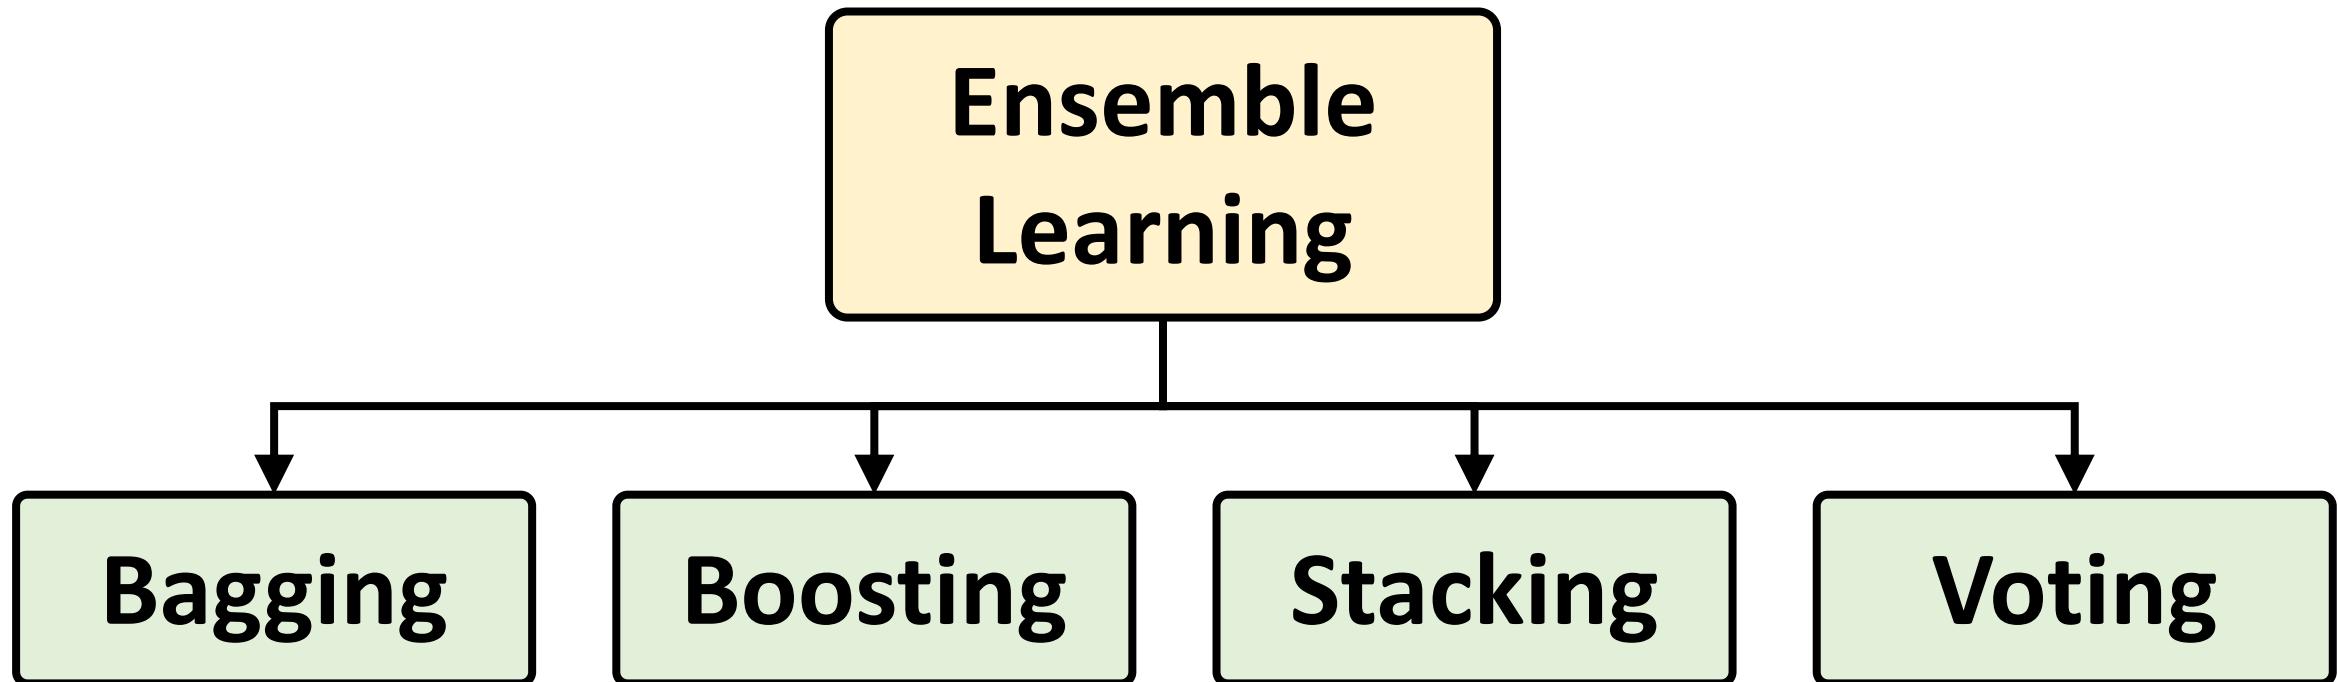
- Select a collection, or **ensemble**, of hypotheses, h_1, h_2, \dots, h_n , and **combine** their predictions by **averaging**, **voting**, or by another level of machine learning.

Ensemble Models

Heterogeneous Ensemble



Ensemble Learning: Bagging, Boosting, Stacking, Voting



Ensemble Learning

- Base model
 - individual hypotheses
 - h_1, h_2, \dots, h_n
- Ensemble model
 - hypotheses combination

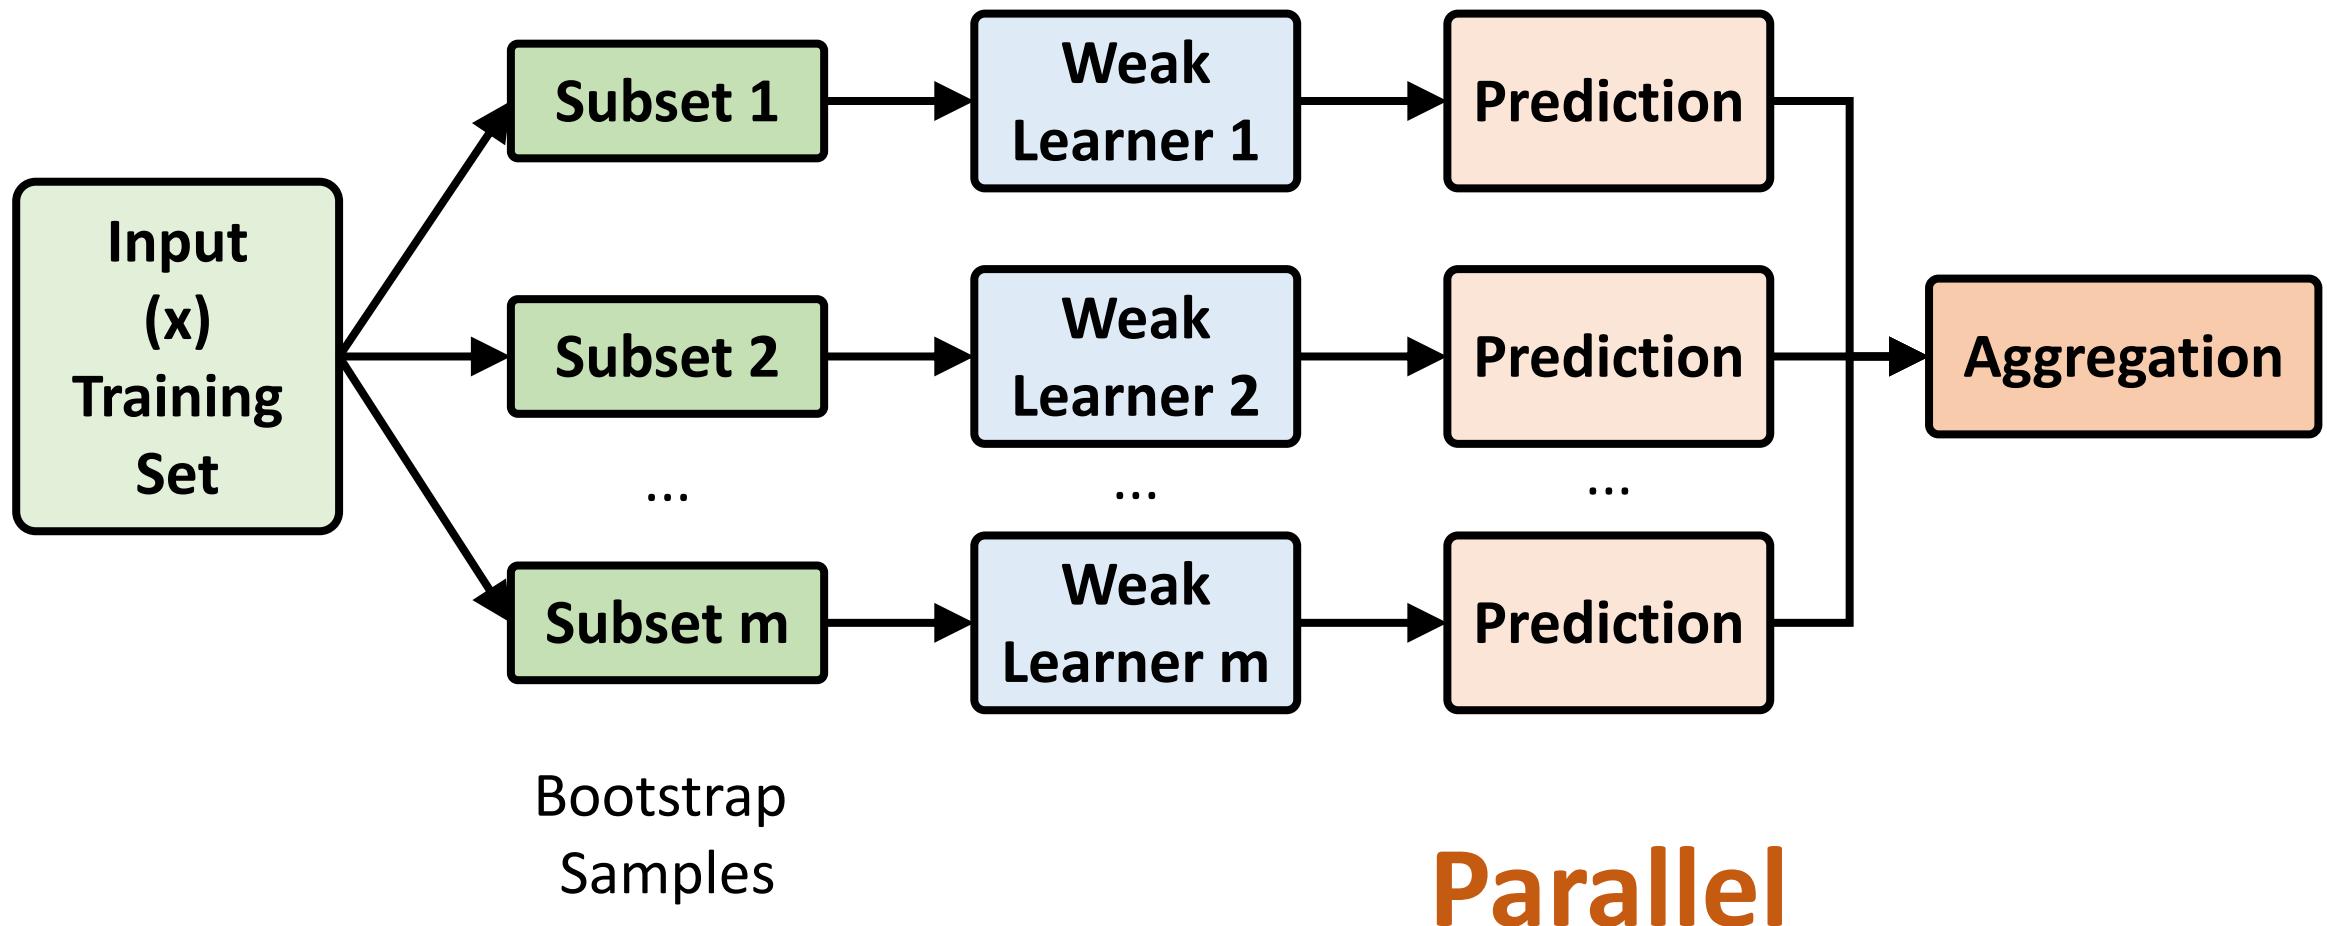
Why Ensemble Learning

- Reduce bias
- Reduce variance

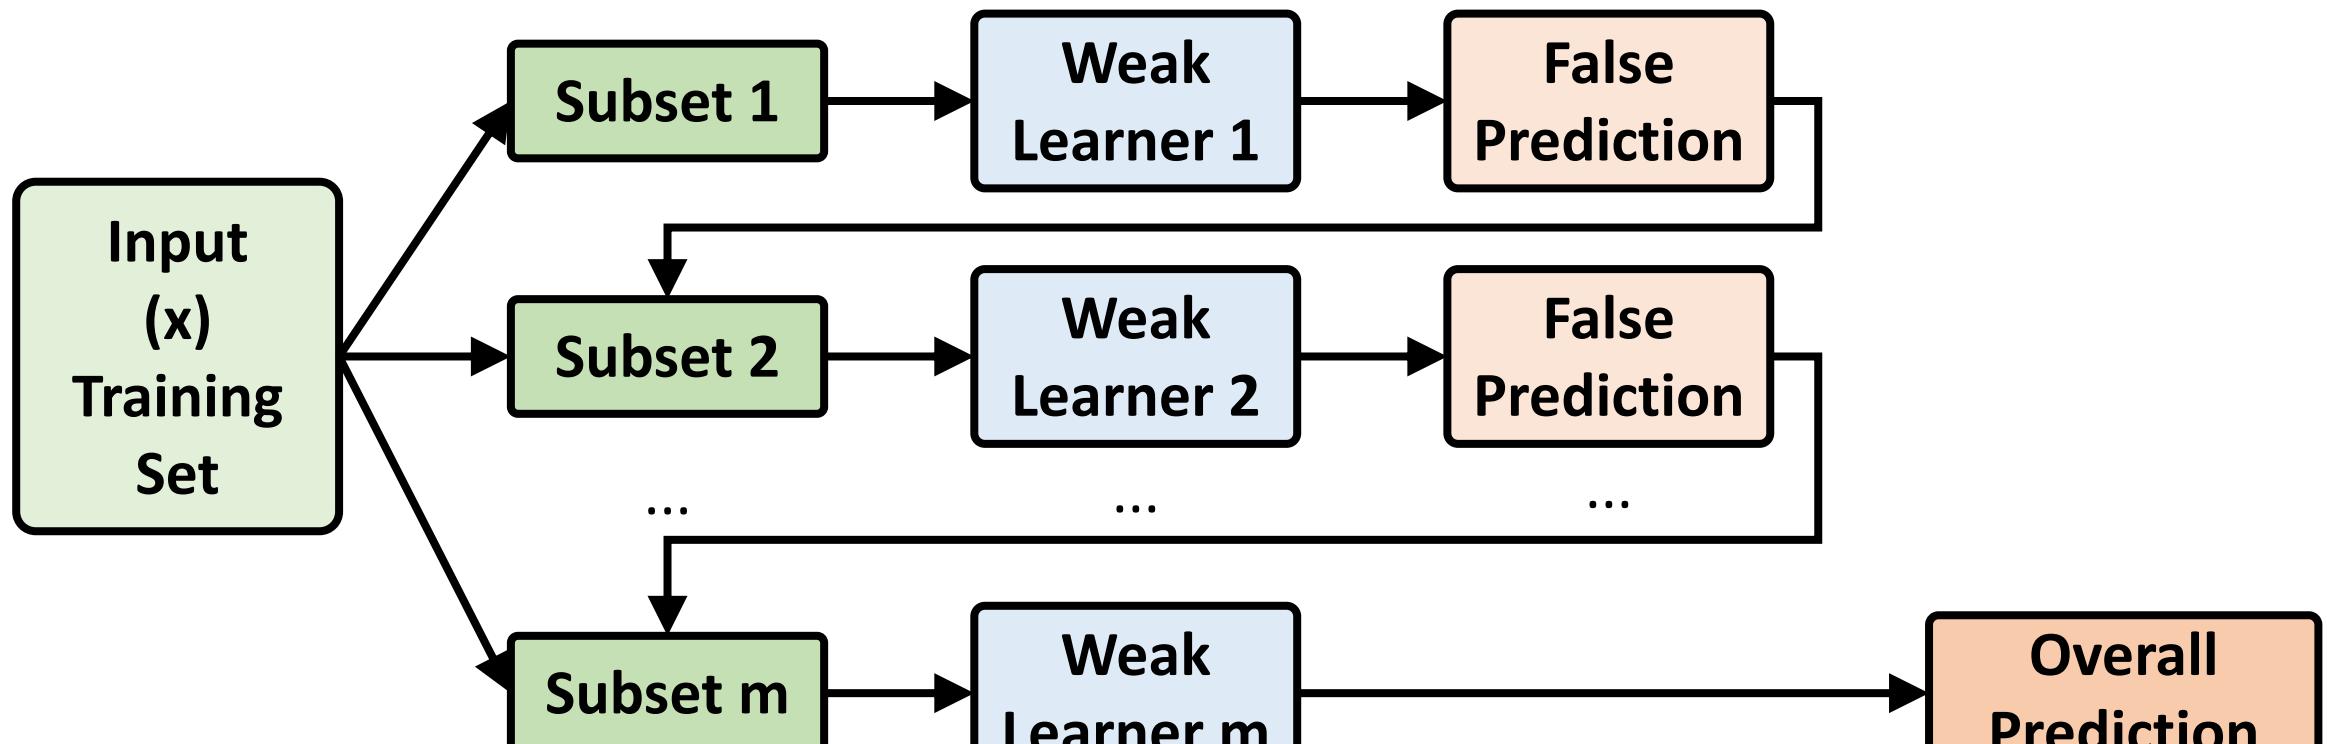
Ensemble Learning

- **Bagging**
 - Random Forests (RF)
- **Boosting**
 - Gradient Boosting, XGBoost, LightGBM, CatBoost
- **Stacking**
- **Online learning**

Ensemble Learning: Bagging (Bootstrap Aggregation)

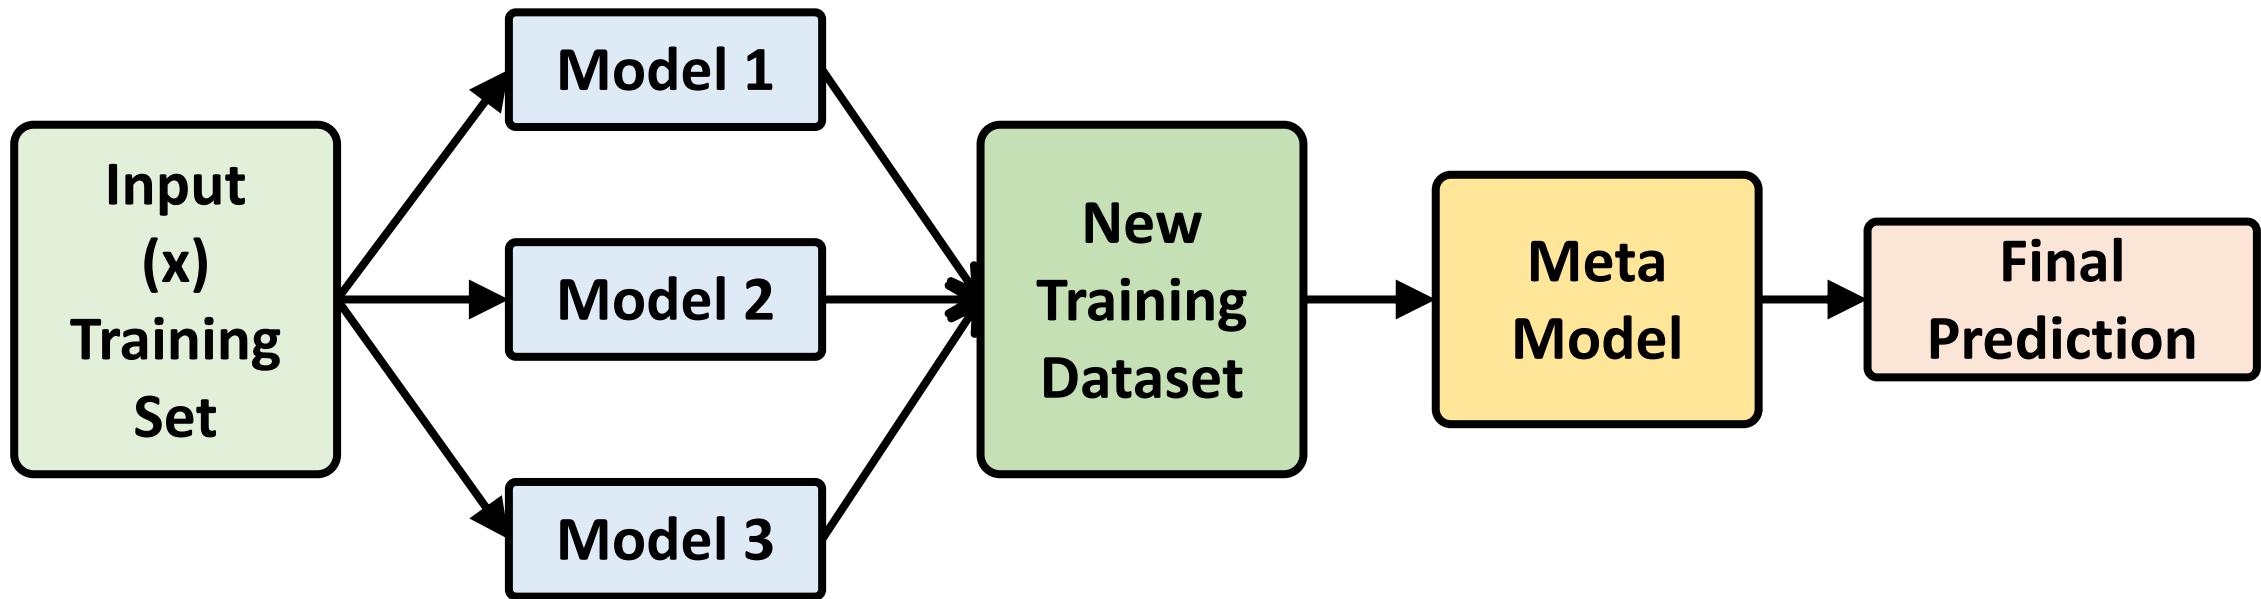


Ensemble Learning: Boosting

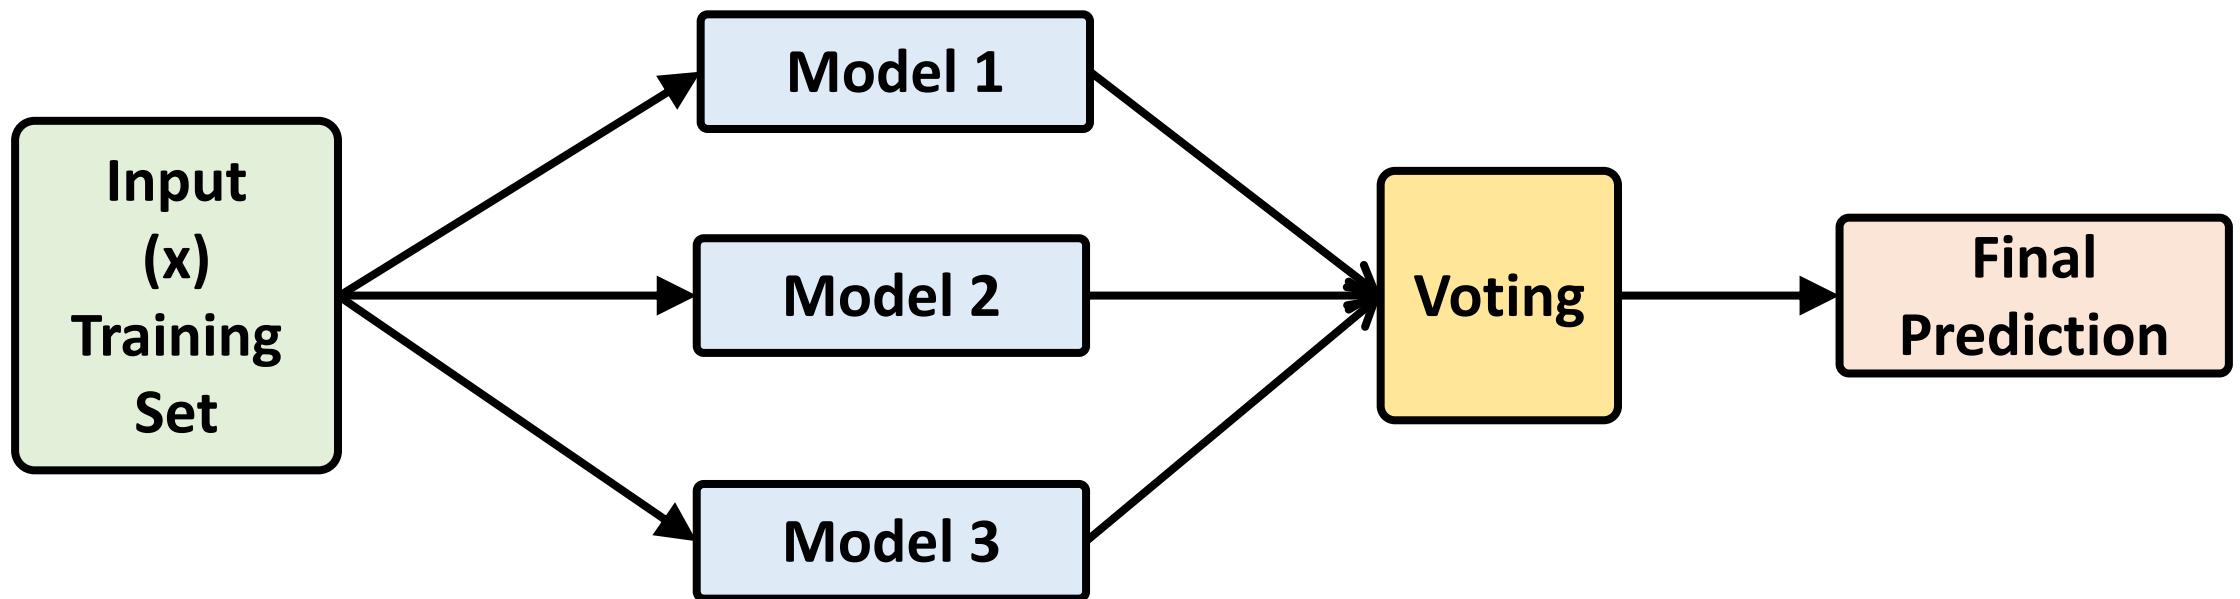


Sequential

Ensemble Learning: Stacking



Ensemble Learning: Voting



Ensemble Learning: Bagging

- **Bagging**
 - Generate distinct training sets by sampling with replacement from the original training set.
- Classification:
 - Plurality Vote (Majority Vote)
- Regression:
 - Average

Ensemble Learning: Random forests

- **Random forest model** is a form of **decision tree bagging** in which we take extra steps to make the ensemble of trees more diverse, to reduce variance.
- The key idea is to randomly vary the **attribute choices** (rather than the training examples)

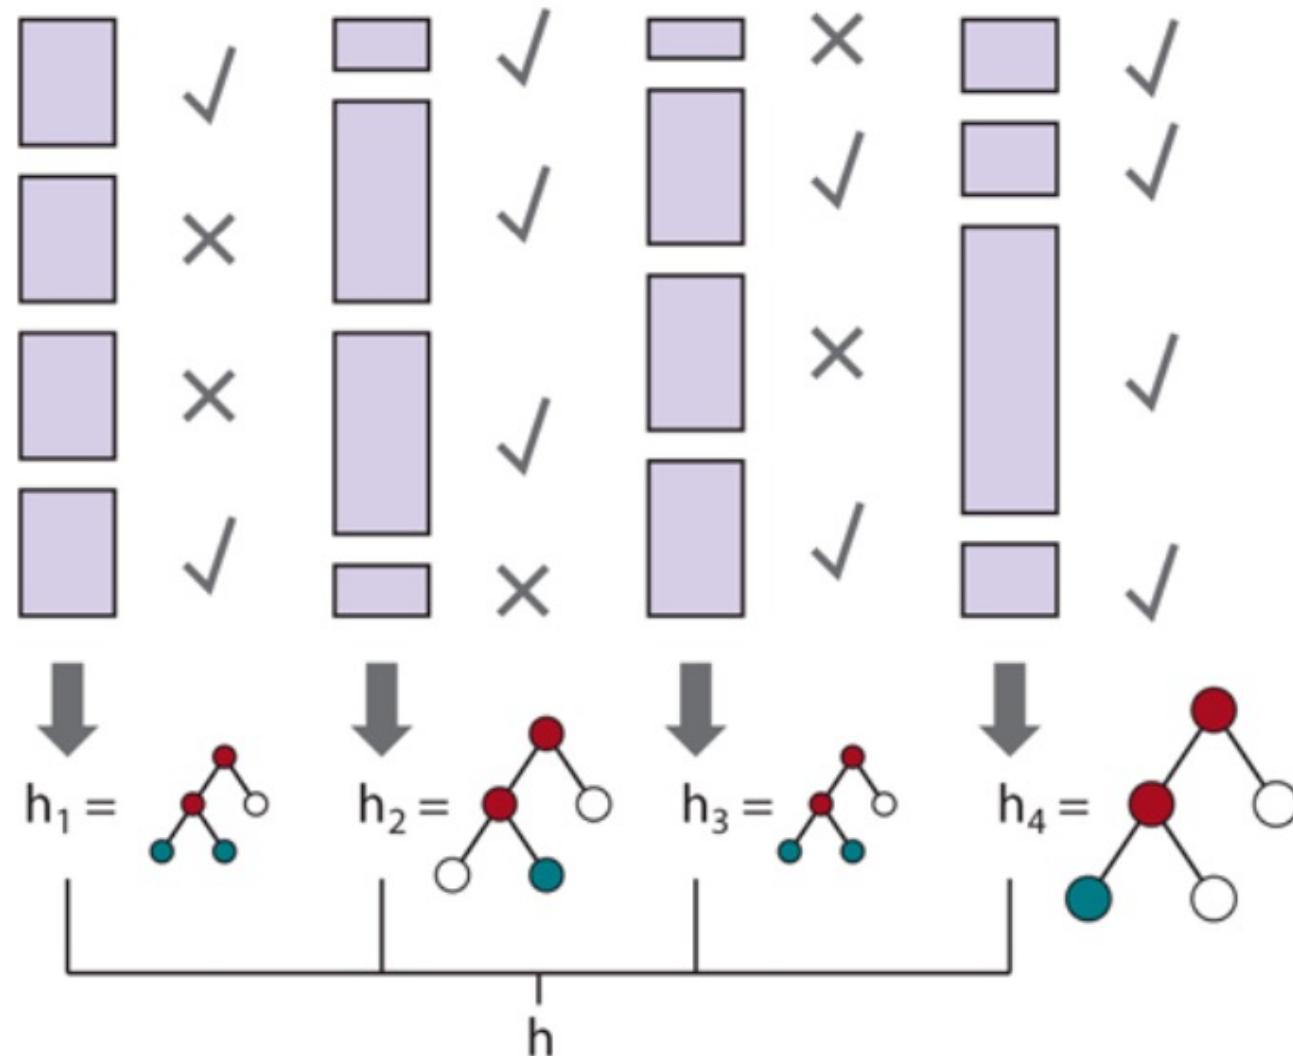
Ensemble Learning: Random forests

- **Extremely randomized trees (ExtraTrees)**
 - Use randomness in selecting the split point value
 - for each selected attribute, we randomly sample several candidate values from a uniform distribution over the attribute's range

Ensemble Learning: Boosting

- **Boosting**
 - The most popular ensemble method
- **Weighted training set**

Ensemble Learning: Boosting



Ensemble Learning: Gradient boosting

- **Gradient boosting**
 - Gradient boosting is a form of boosting using gradient descent
- **Gradient boosting machines (GBM)**
- **Gradient boosted regression trees (GBRT)**
- Popular method for regression and classification of factored tabular data

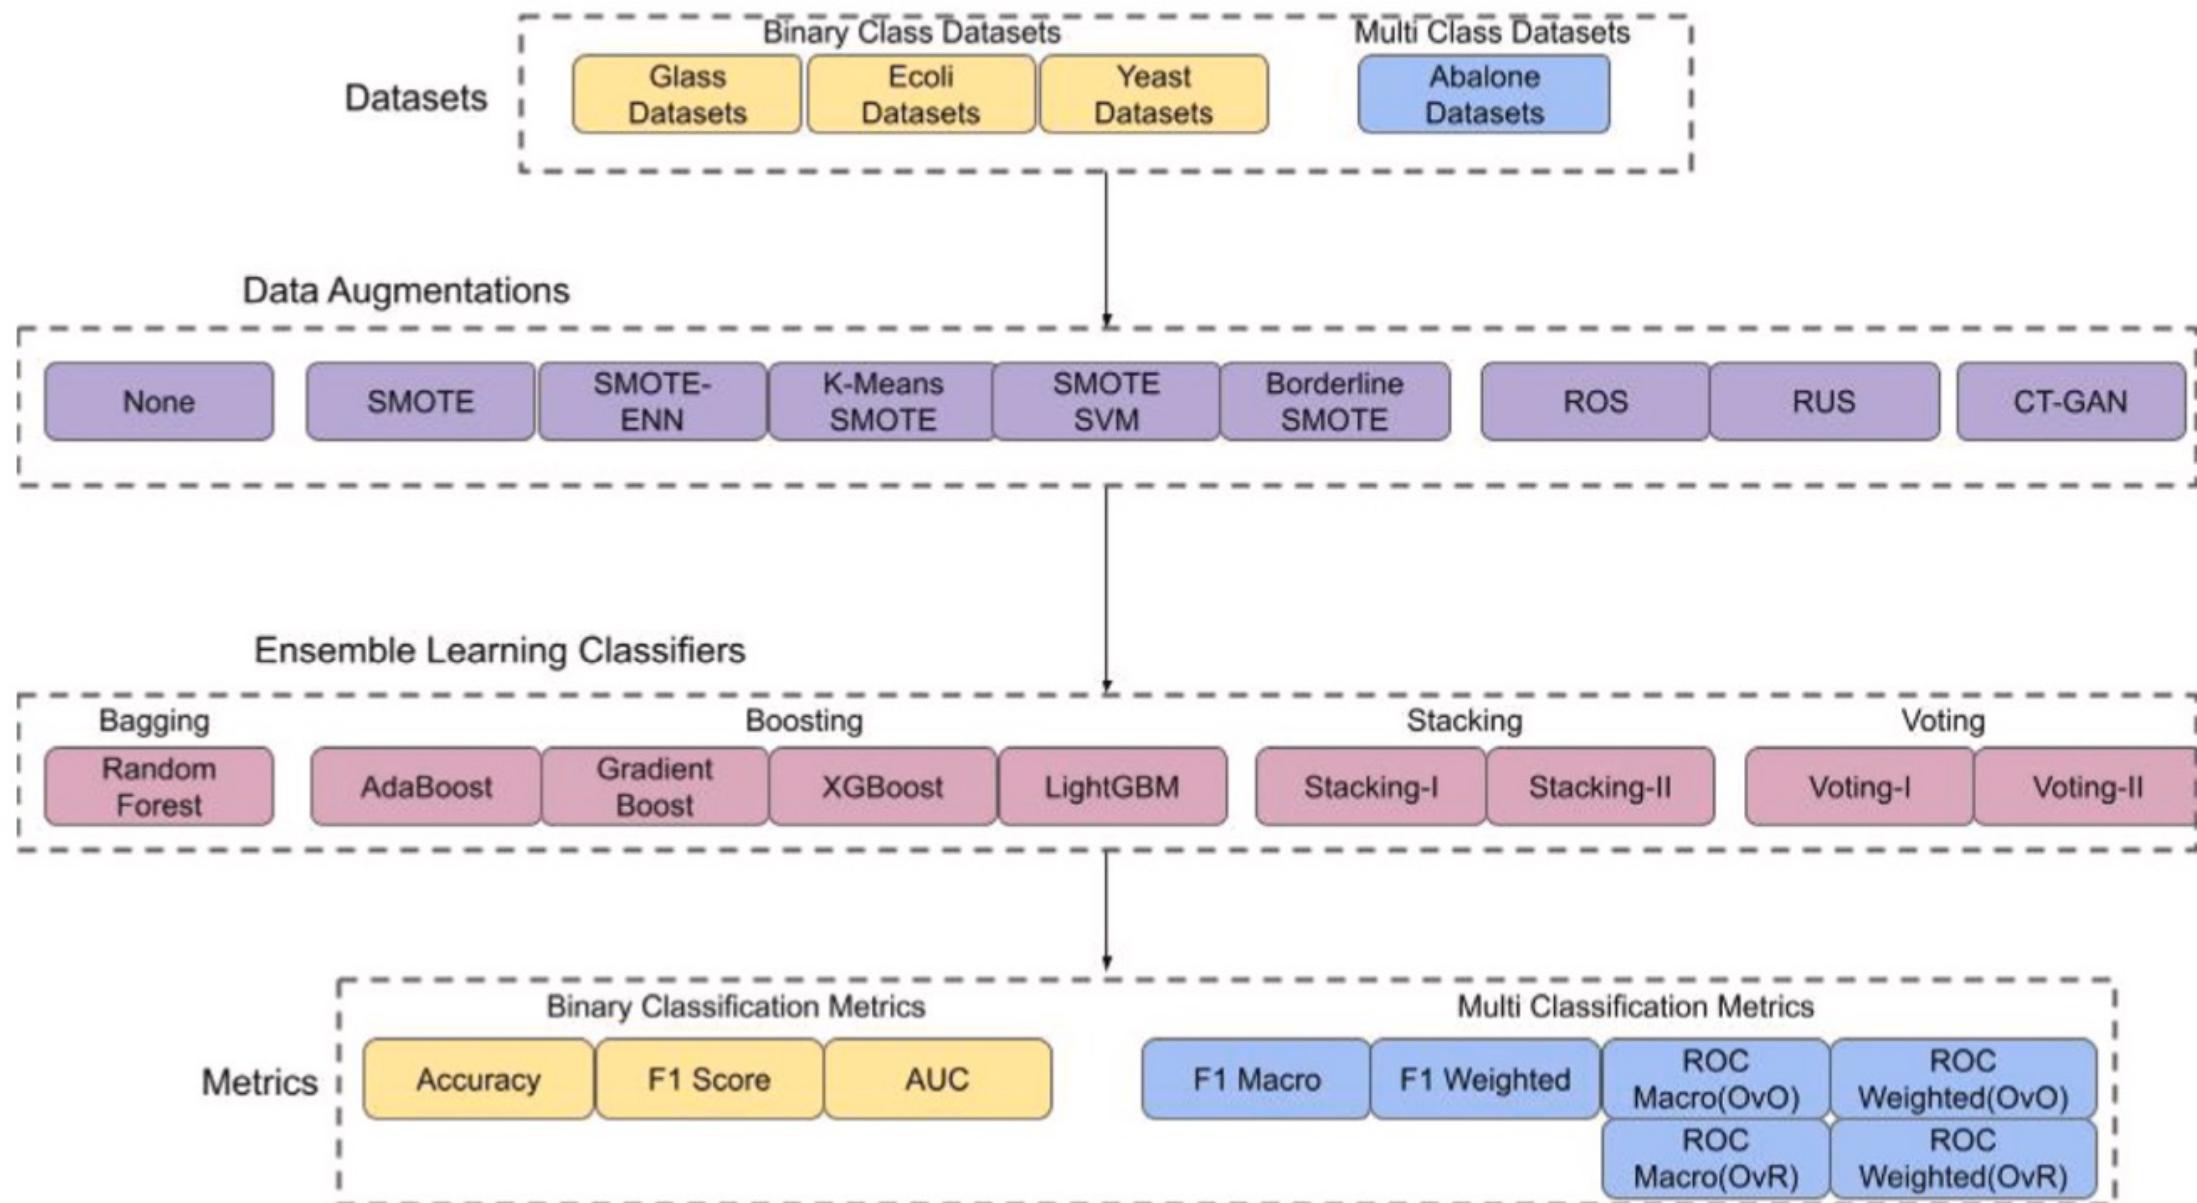
Ensemble Learning: Stacking

- **Stacking**
 - Stacked generalization combines **multiple base models from different model classes** trained on the **same data**.
- **Bagging**
 - Combines **multiple base models of the same model class** trained on **different data**.

Ensemble Learning: Online learning

- Online learning
 - Data are not i.i.d.
(independent and identically distributed)
 - An agent receives an input x_i from nature, predicts the corresponding y_i and then is told the correct answer.

Ensemble Learning and Data Augmentations (DA)



ML Evaluation of Imbalanced Dataset: Ensemble Learning and Data Augmentations (DA)

Data augmentation	Ensemble model	Accuracy (Best, Std)	F1 (Best, Std)	AUC (Best, Std)
Yeast-v6				
Borderline-SMOTE	LightGBM	98.490(98.653, 0.150)	99.230(99.314, 0.077)	76.668(76.751, 0.077)
No augmentation	Stacking-I	98.185(98.653, 0.320)	99.076(99.315, 0.165)	69.757(76.579, 3.658)
SVM-SMOTE	AdaBoost	98.072(98.653, 0.466)	99.015(99.314, 0.240)	75.577(80.253, 2.014)
ROS	Stacking-I	97.997(98.822, 0.415)	98.977(99.401, 0.214)	74.884(76.837, 2.607)
Borderline-SMOTE	Voting-Soft	97.949(98.653, 0.501)	98.951(99.314, 0.259)	75.854(80.253, 1.923)
ROS	XGBoost	97.866(98.485, 0.303)	98.908(99.227, 0.157)	76.348(76.665, 0.155)
SMOTE	Stacking-II	97.539(98.485, 0.708)	98.737(99.227, 0.370)	77.385(83.841, 2.273)
SMOTE	Voting-Hard	97.386(98.485, 0.707)	98.657(99.227, 0.370)	77.607(83.841, 2.492)
SMOTE-ENN	Random Forests	92.917(97.306, 1.932)	96.273(98.618, 1.048)	73.380(78.962, 1.694)
RUS	Stacking-II	87.345(94.444, 3.405)	93.087(97.093, 2.014)	81.085(88.581, 3.112)

Kaggle Datasets for Machine Learning

≡

Search



Datasets

+ New Dataset Your Work

🔍 ESG

Filters

All datasets X Computer Science Education Classification Computer Vision NLP Data Visualization Pre-Trained Model

89 Datasets

Hotness ▾

Hotness ▾

50

Bronze ...

67

Gold ...

246

Silver ...

**S&P 500 ESG Risk Ratings**
Prithish Dugar · Updated 4 months ago
Usability 10.0 · 1 File (CSV) · 252 kB

**Public Company ESG Ratings Dataset**
Alistair King · Updated 8 months ago
Usability 10.0 · 1 File (CSV) · 43 kB

**US Funds dataset from Yahoo Finance**
Stefano Leone · Updated 3 years ago
Usability 10.0 · 7 Files (CSV) · 371 MB

Kaggle Datasets for Machine Learning

≡

🔍 Search

👤

Datasets

+ New Dataset

Your Work

📍

🔍 credit card

Filters

🏆

All datasets X Computer Science Education Classification Computer Vision NLP Data Visualization Pre-Trained Model

📄 1,125 Datasets

Most Votes ▾

Grid

Line

Machine Learning Group - ULB · Updated 7 years ago

Usability 8.5 · 1 File (CSV) · 69 MB

11606

Gold ...

Machine Learning Group - ULB · Updated 7 years ago

Usability 8.5 · 1 File (CSV) · 69 MB

2580

Gold ...

Aung Pyae · Updated 5 years ago

Usability 8.8 · 1 File (CSV) · 37 kB

Credit Card Fraud Detection

Sakshi Goyal · Updated 4 years ago

Usability 10.0 · 1 File (CSV) · 388 kB

2256

Gold ...

Credit Card customers

Source: <https://www.kaggle.com/datasets?search=credit+card&sort=votes>

Kaggle Datasets: Credit Card Fraud Detection

≡

Search

MACHINE LEARNING GROUP - ULB AND 1 COLLABORATOR · UPDATED 7 YEARS AGO

▲ 11606 New Notebook Download ⋮

Credit Card Fraud Detection

Anonymized credit card transactions labeled as fraudulent or genuine

Data Card Code (4966) Discussion (107) Suggestions (0)

About Dataset

Context

It is important that credit card companies are able to recognize fraudulent credit card transactions so that customers are not charged for items that they did not purchase.

Content

The dataset contains transactions made by credit cards in September 2013 by European cardholders.



Usability ⓘ
8.53

License
Database: Open Database, Cont...

Expected update frequency
Not specified

Tags

Kaggle Code: Credit Card Fraud Detection



Credit Card Fraud Detection



11606

New Notebook

Download

[Data Card](#) [Code \(4966\)](#) [Discussion \(107\)](#) [Suggestions \(0\)](#)[All](#) [Your Work](#) [Shared With You](#) [Bookmarks](#)

Most Votes ▾



Credit Fraud || Dealing with Imbalanced Datasets

Updated 5y ago

[657 comments](#) · Credit Card Fraud Detection

5453

Gold ⋮



Outlier!!! The Silent Killer

Updated 3y ago

[138 comments](#) · Titanic - Machine Learning from Disaster +16

1070

Gold ⋮



In depth skewed data classif. (93% recall acc now)

Updated 8y ago

[125 comments](#) · Credit Card Fraud Detection

796

Gold ⋮



Credit Card Fraud Detection Predictive Models

Updated 4y ago

[41 comments](#) · Credit Card Fraud Detection

489

Gold ⋮

Kaggle Code: Credit Card Fraud Detection

The screenshot shows a Kaggle kernel page for a notebook titled "Credit Fraud || Dealing with Imbalanced Datasets". The page includes a sidebar with various icons, a search bar, user information, and navigation controls. The main content area displays the notebook's title, Python version, run time, and categories.

Credit Fraud || Dealing with Imbalanced Datasets

Python · Credit Card Fraud Detection

Notebook Input Output Logs Comments (657)

Run
861.1s Version 70 of 70

Finance Banking Data Visualization Classification Dimensionality Reduction

Credit Fraud Detector

Note: There are still aspects of this kernel that will be subjected to changes. I've noticed a recent increase of interest towards this kernel so I will focus more on the steps I took and why I took them to make it clear why I took those steps.

Before we Begin:

Meta Learning: Learning to Learn

Deep Learning

Transfer Learning

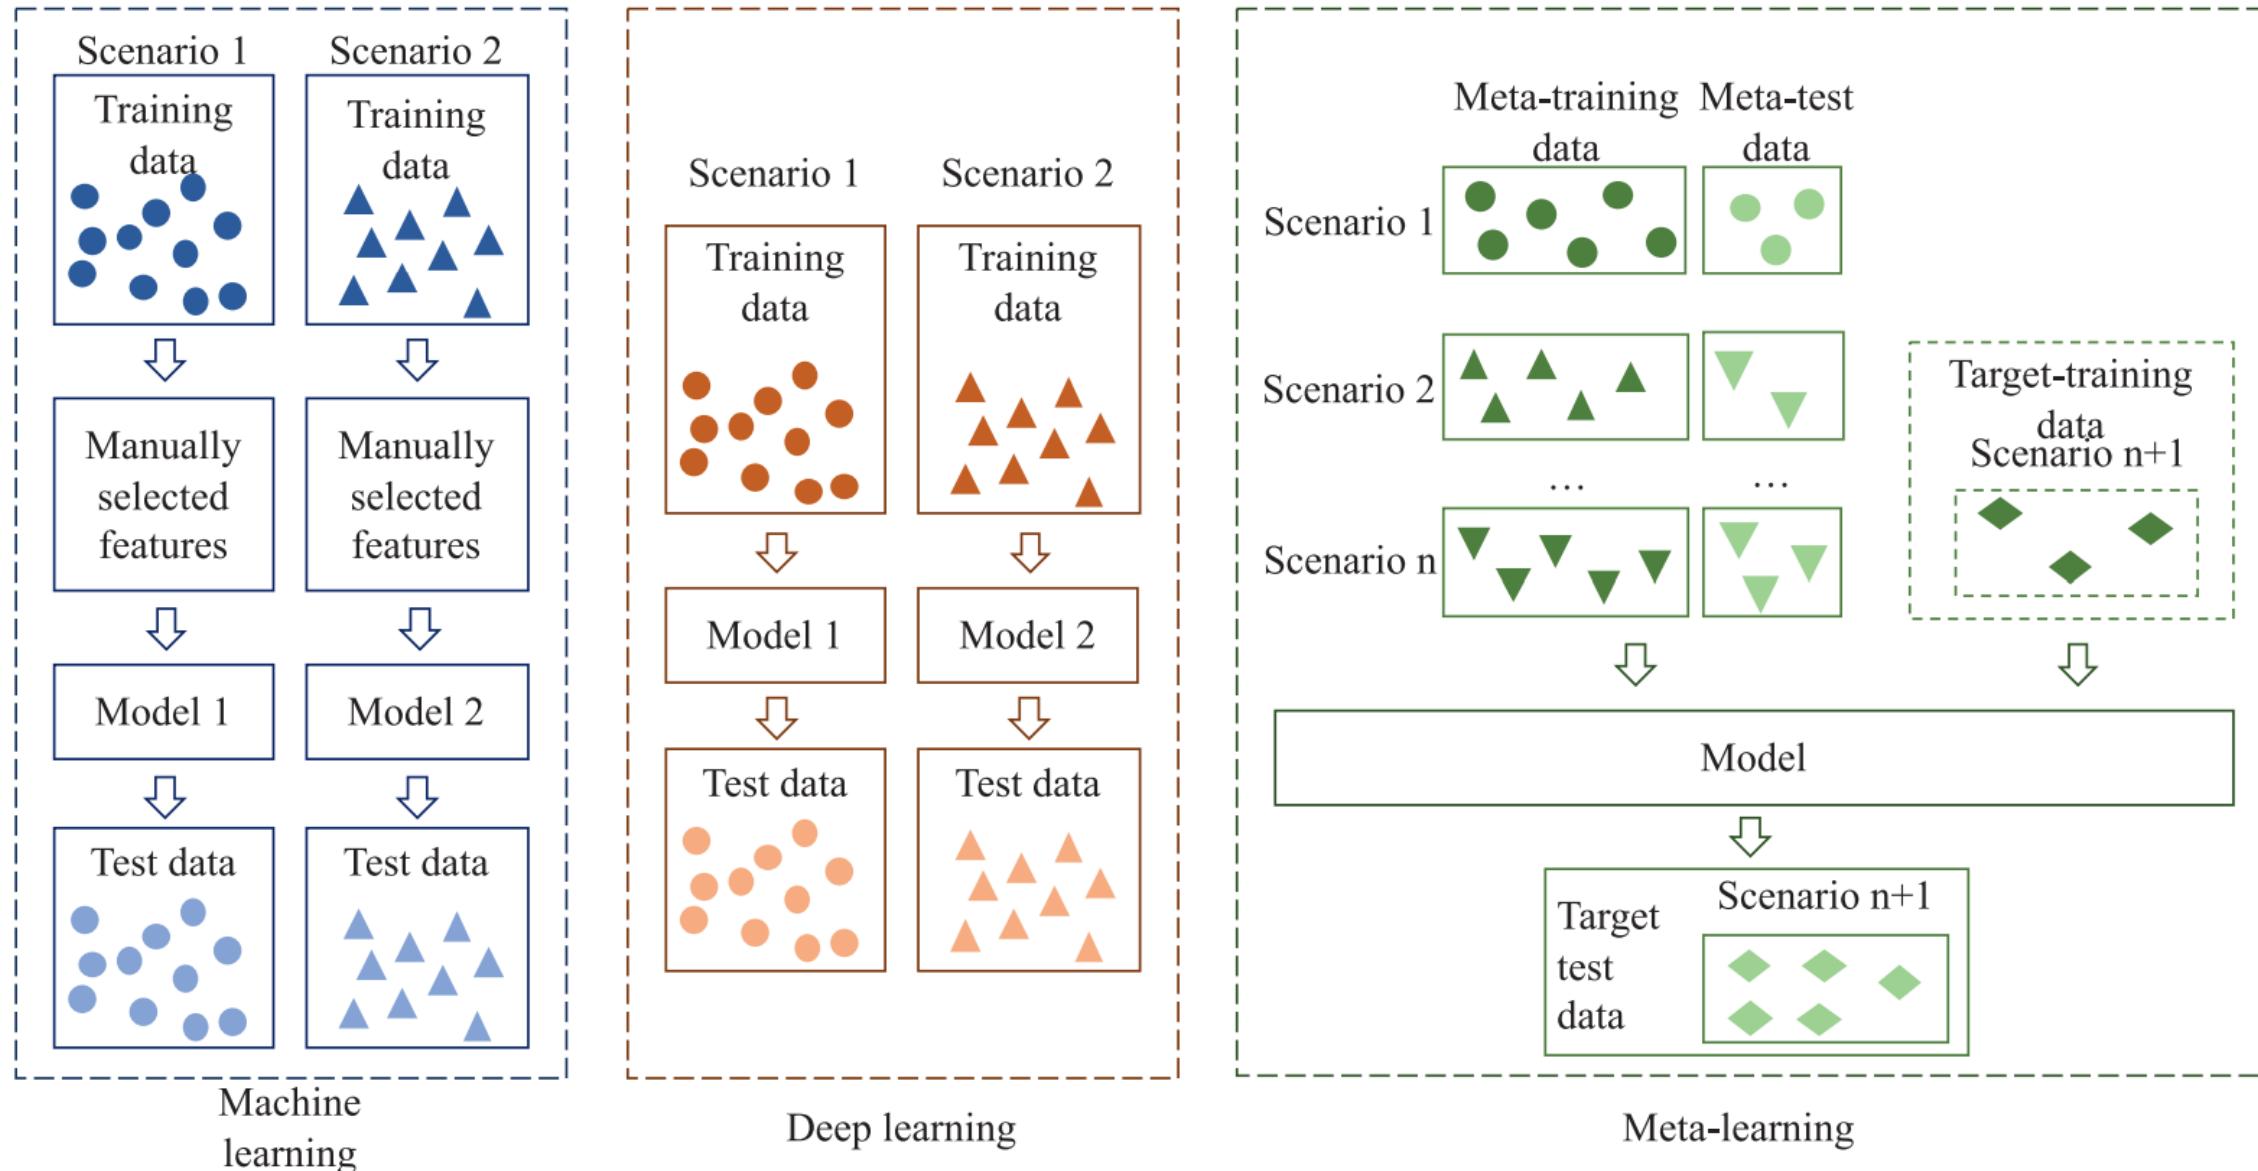
Few-Shot Learning

Meta Learning

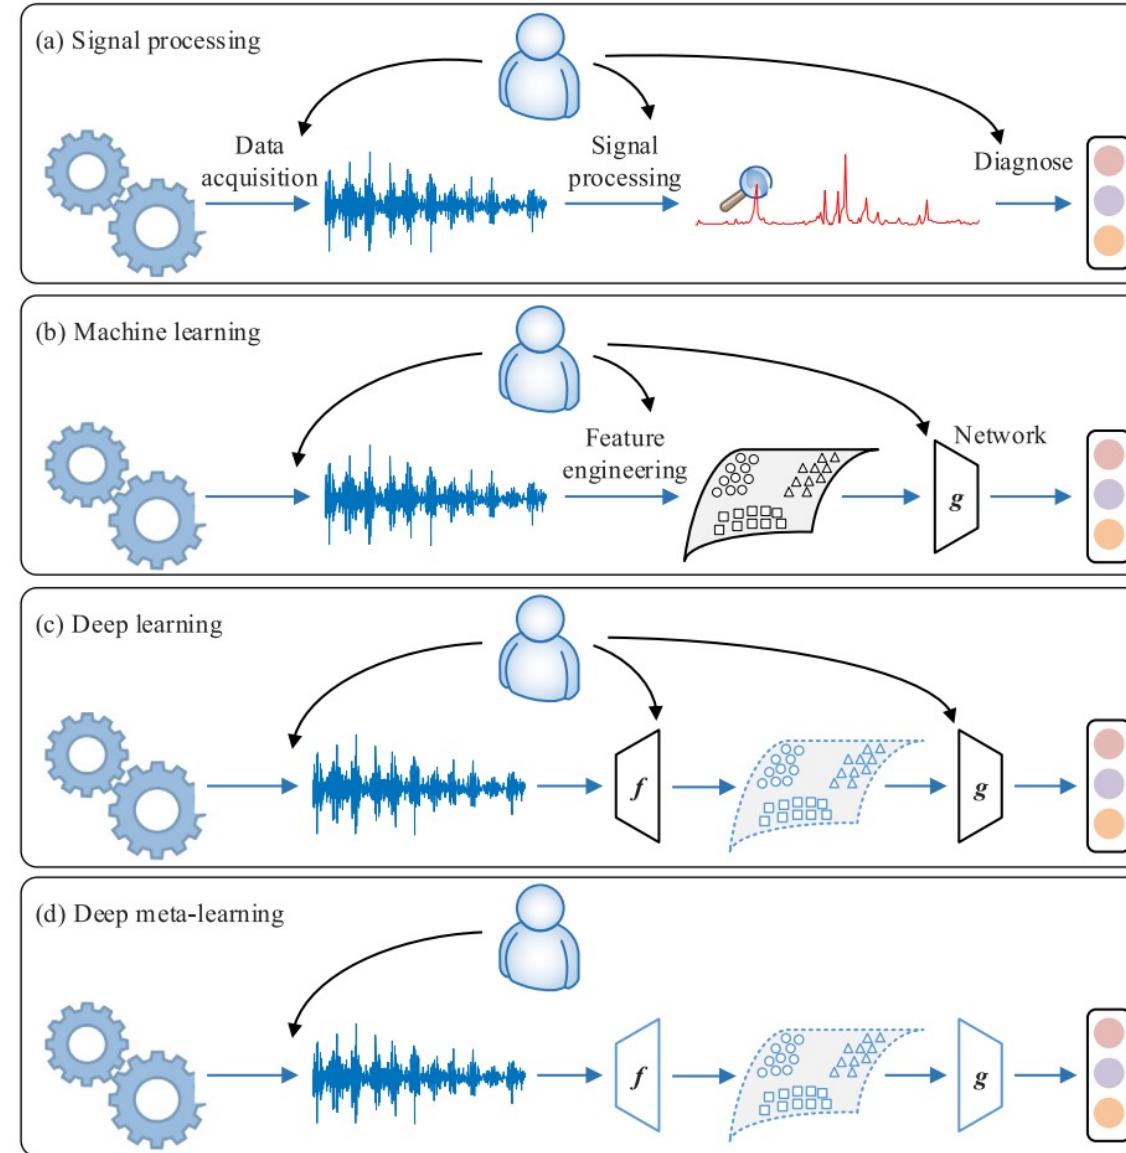
Deep Learning, Transfer Learning, Few-Shot Learning, Meta Learning

- Deep Learning
 - Transfer Learning
 - Pre-training, Fine-Tuning (FT)
- Meta Learning: Learning to Learn
- Few-Shot Learning (FSL)
- One-Shot Learning (1SL)
- Zero-Shot Learning (OSL)(ZSL)

Machine Learning, Deep Learning, Meta Learning

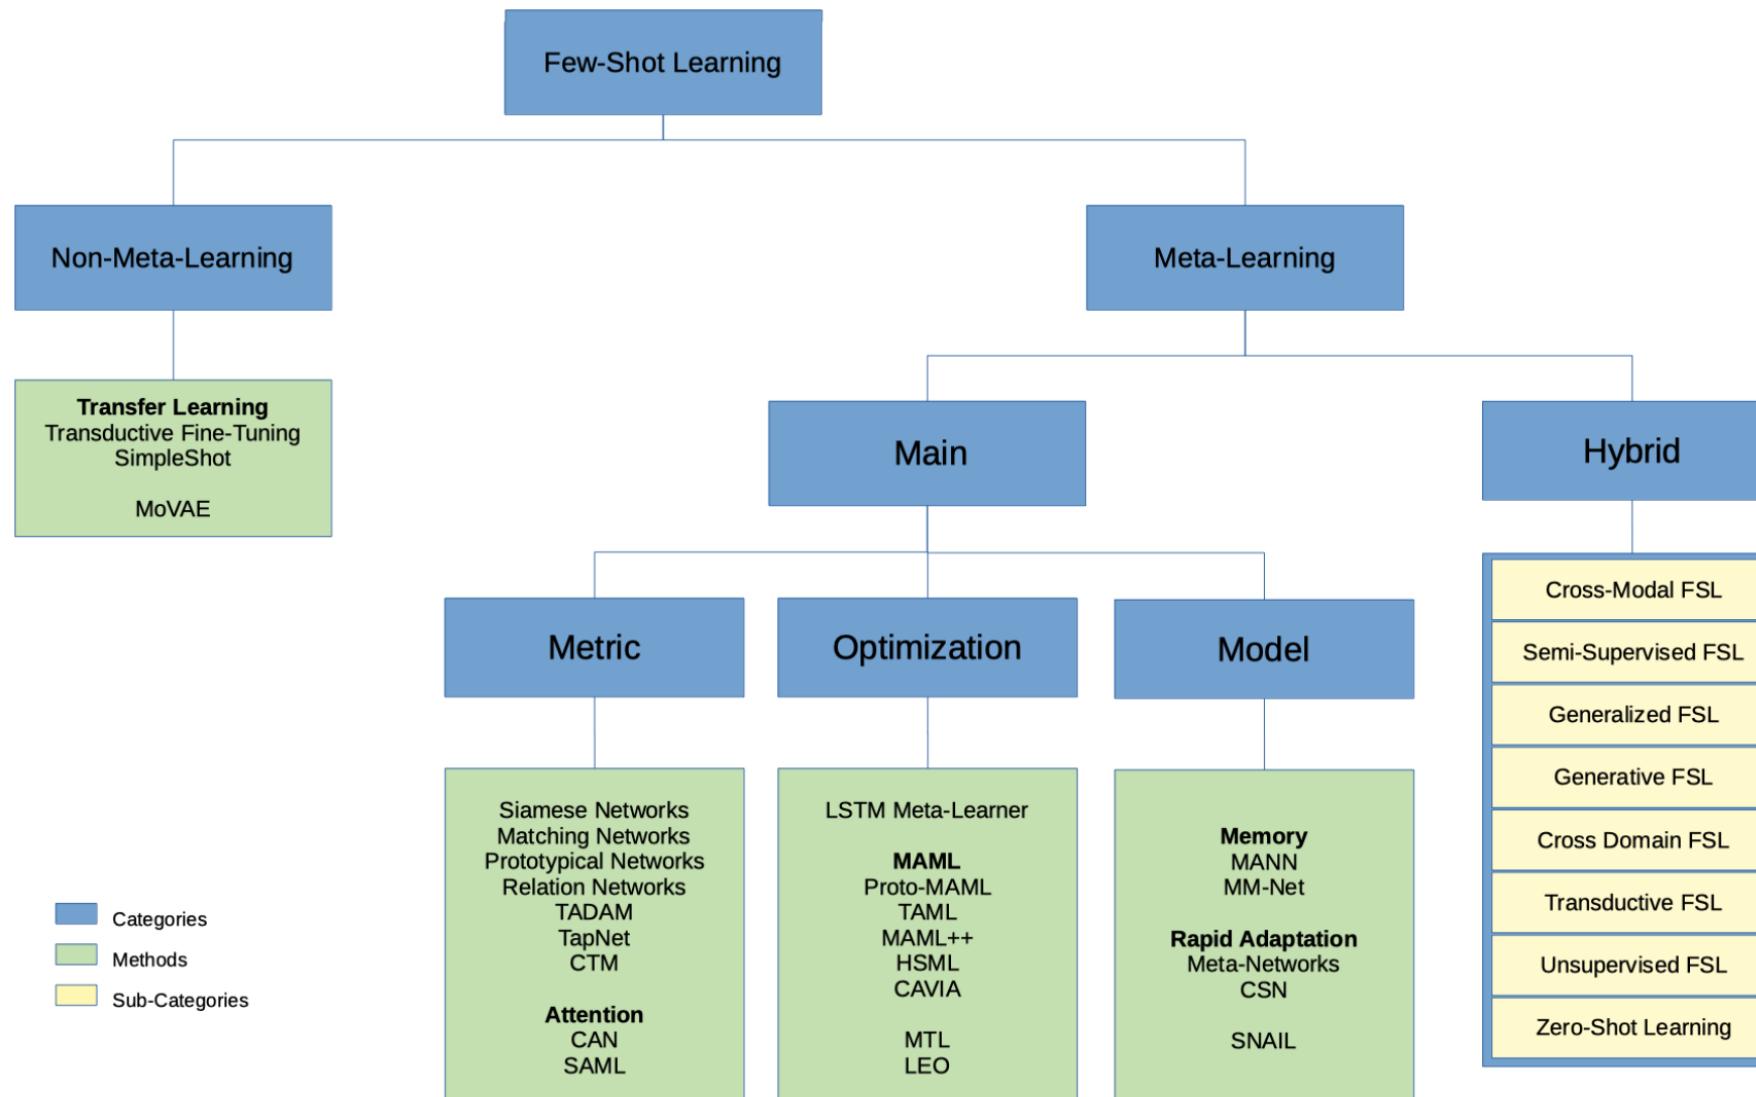


Machine Learning, Deep Learning, Meta Learning



Few-Shot Learning (FSL) and Meta Learning

Machine learning from few training examples



Meta Learning, Transfer Learning, Ensemble Learning, Continual Learning, Multi-Task Learning

Features	Method					
	Transfer learning	Ensemble learning	Continual learning	Multi-task learning	Hierarchical Bayesian models	
Meta-learning						
Learning from prior experience	✓	✓	✗	✓	✗	✓
Relationship between source tasks	No limitation	Related	Same	Task streams	Related	Related
Relationship between source tasks and target tasks	No limitation	Related	Same	Related	Related	Related
Considering the requirements of the target task	✓	✗	✗	✗	✗	✗

Meta-Learning and Few-shot Learning

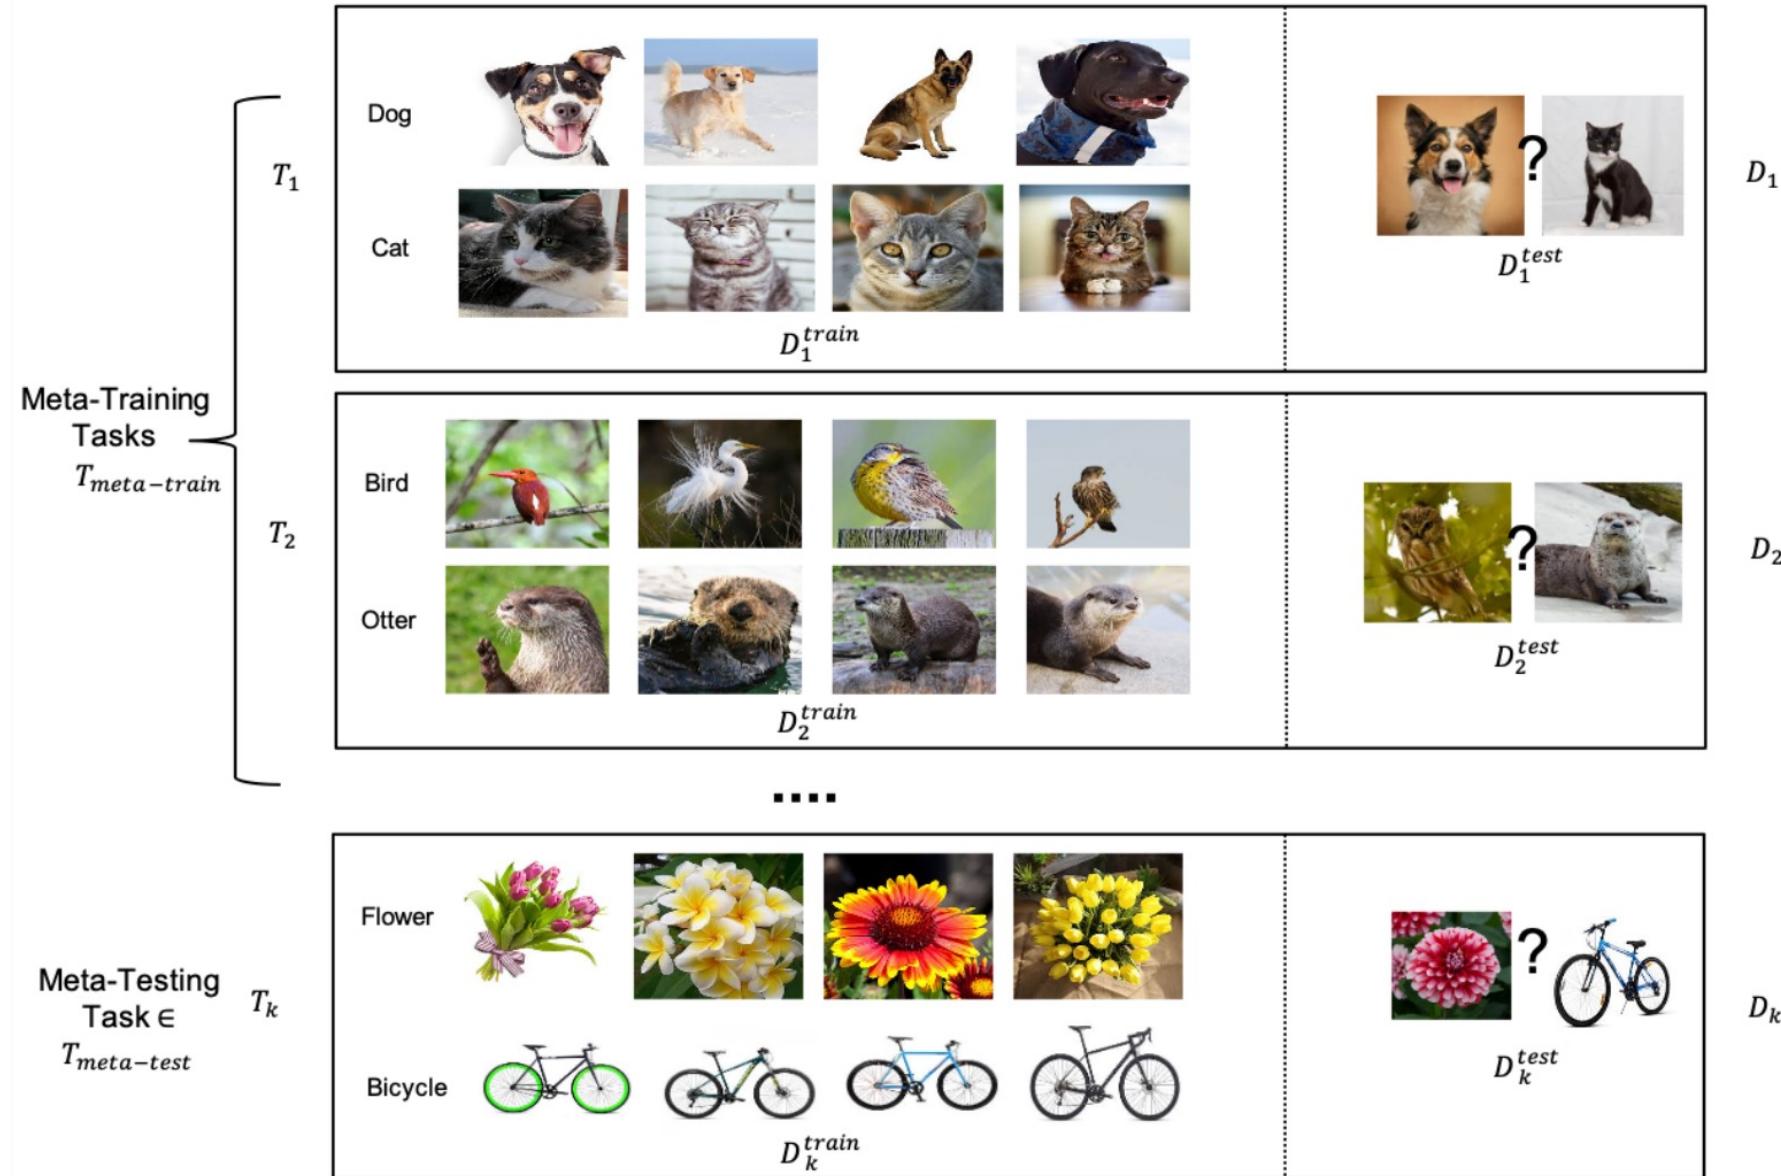
Notations and Terms

Optimization-based Meta-learning		Metric-based Meta-learning	
Notation A	Term A	Notation B	Term B
\mathcal{D}_i^{train}	Training set for task \mathcal{T}_i	S_i	Support Set for task \mathcal{T}_i
\mathcal{D}_i^{test}	Test set for task \mathcal{T}_i	Q_i	Query Set for task \mathcal{T}_i
$\mathcal{D}_{meta-train}$	Meta-training set	\mathcal{D}_{train}	Training Set
$\mathcal{D}_{meta-test}$	Meta-testing set	\mathcal{D}_{test}	Test Set

Meta-Learning Symbols

Symbol	Meaning
\mathcal{T}_i	Task i
\mathcal{L}	Loss function
(x_k, y_k)	Input-Output pair
f_θ	Model (function) with parameters θ
g_{θ_1}	Embedding function
d_{θ_2} or d	Distance function
g_ϕ	Meta-Learning model with parameters ϕ
$P_\theta(y x)$	Output probability of y for input x using model parameters θ
$k_\theta(x_1, x_2)$	Kernel function measuring similarity between two vectors x_1 and x_2
σ	Softmax function
α, β	Learning rates
w	Weights
\mathbf{v}_c	Prototype of class c
C	Set of classes present in S
S^c	Subset of S containing all elements (x_k, y_k) such that $y_k = c$
\oplus	Concatenation operator
B	Number of batches (X_b, Y_b) sampled in inner-loop for a randomly sampled task \mathcal{T}_i
I	Number of tasks \mathcal{T}_i sampled in inner-loop
J	Number of outer-loop iterations

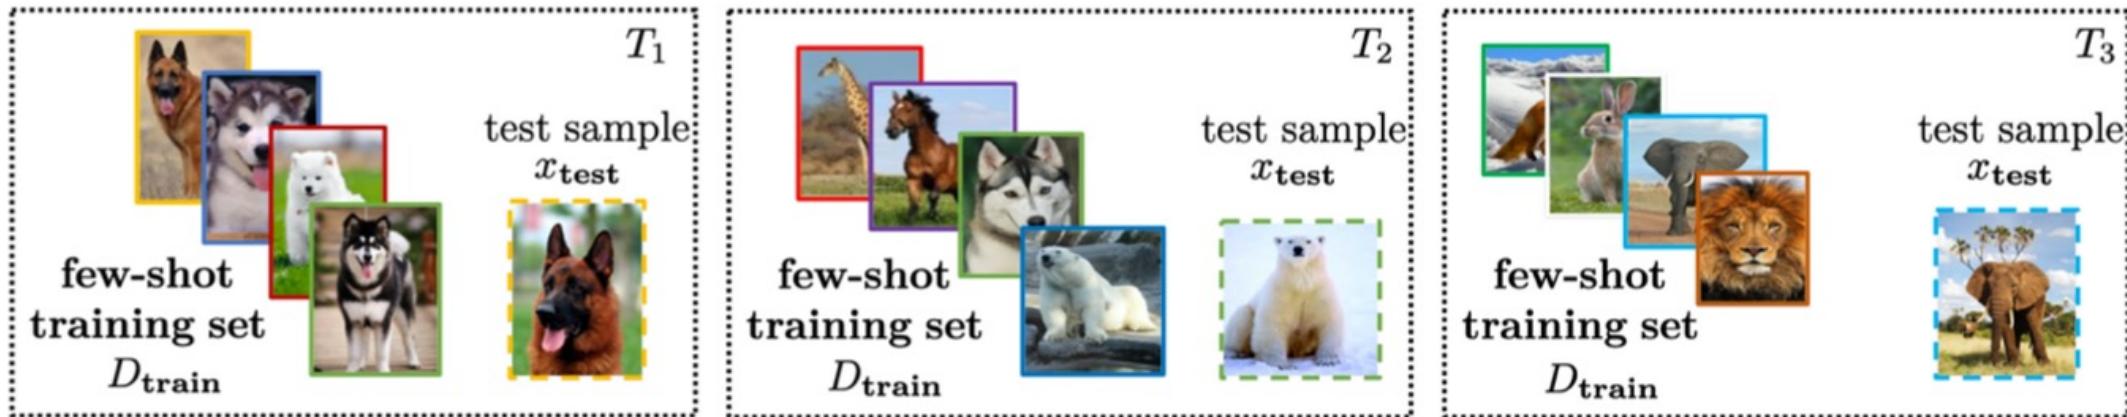
Meta-Learning Example Setup



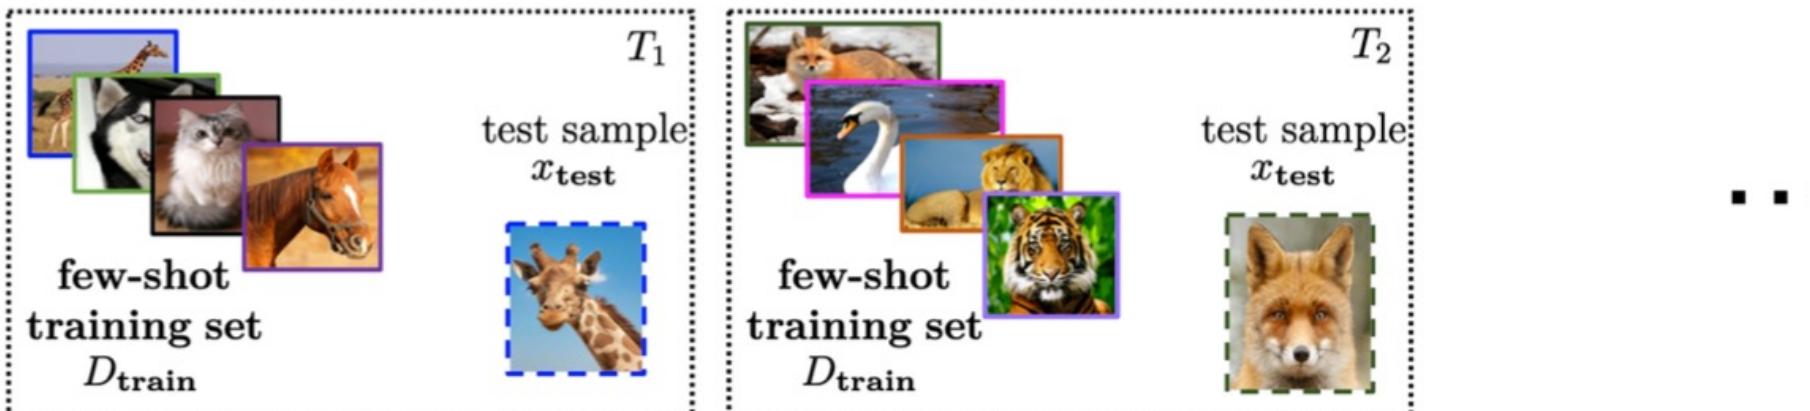
Few-Shot Learning (FSL)

Solving the FSL problem by meta-learning

meta-training tasks T_s 's



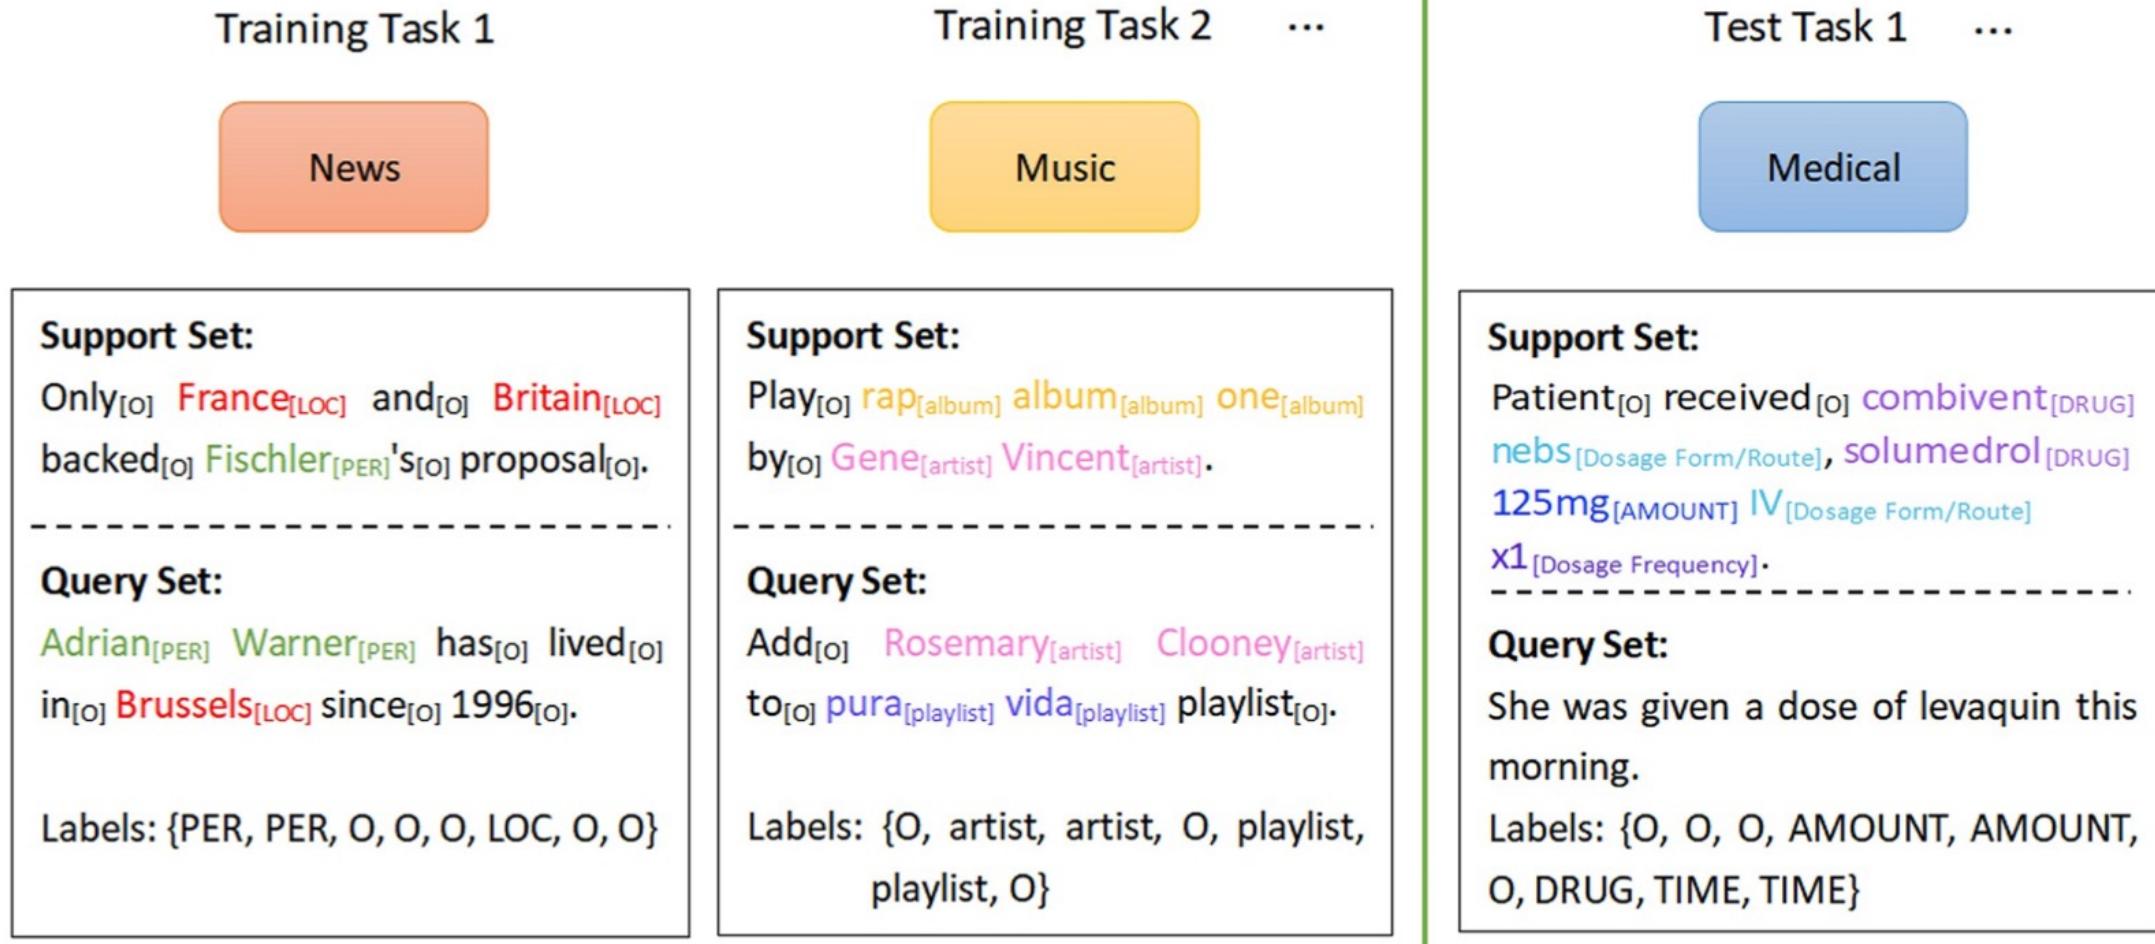
meta-testing tasks T_t 's



Few-Shot Learning (FSL)

Meta-learning

Each task mimics the few-shot scenario, and can be completely non-overlapping.
Support sets are used to train; query sets are used to evaluate the model



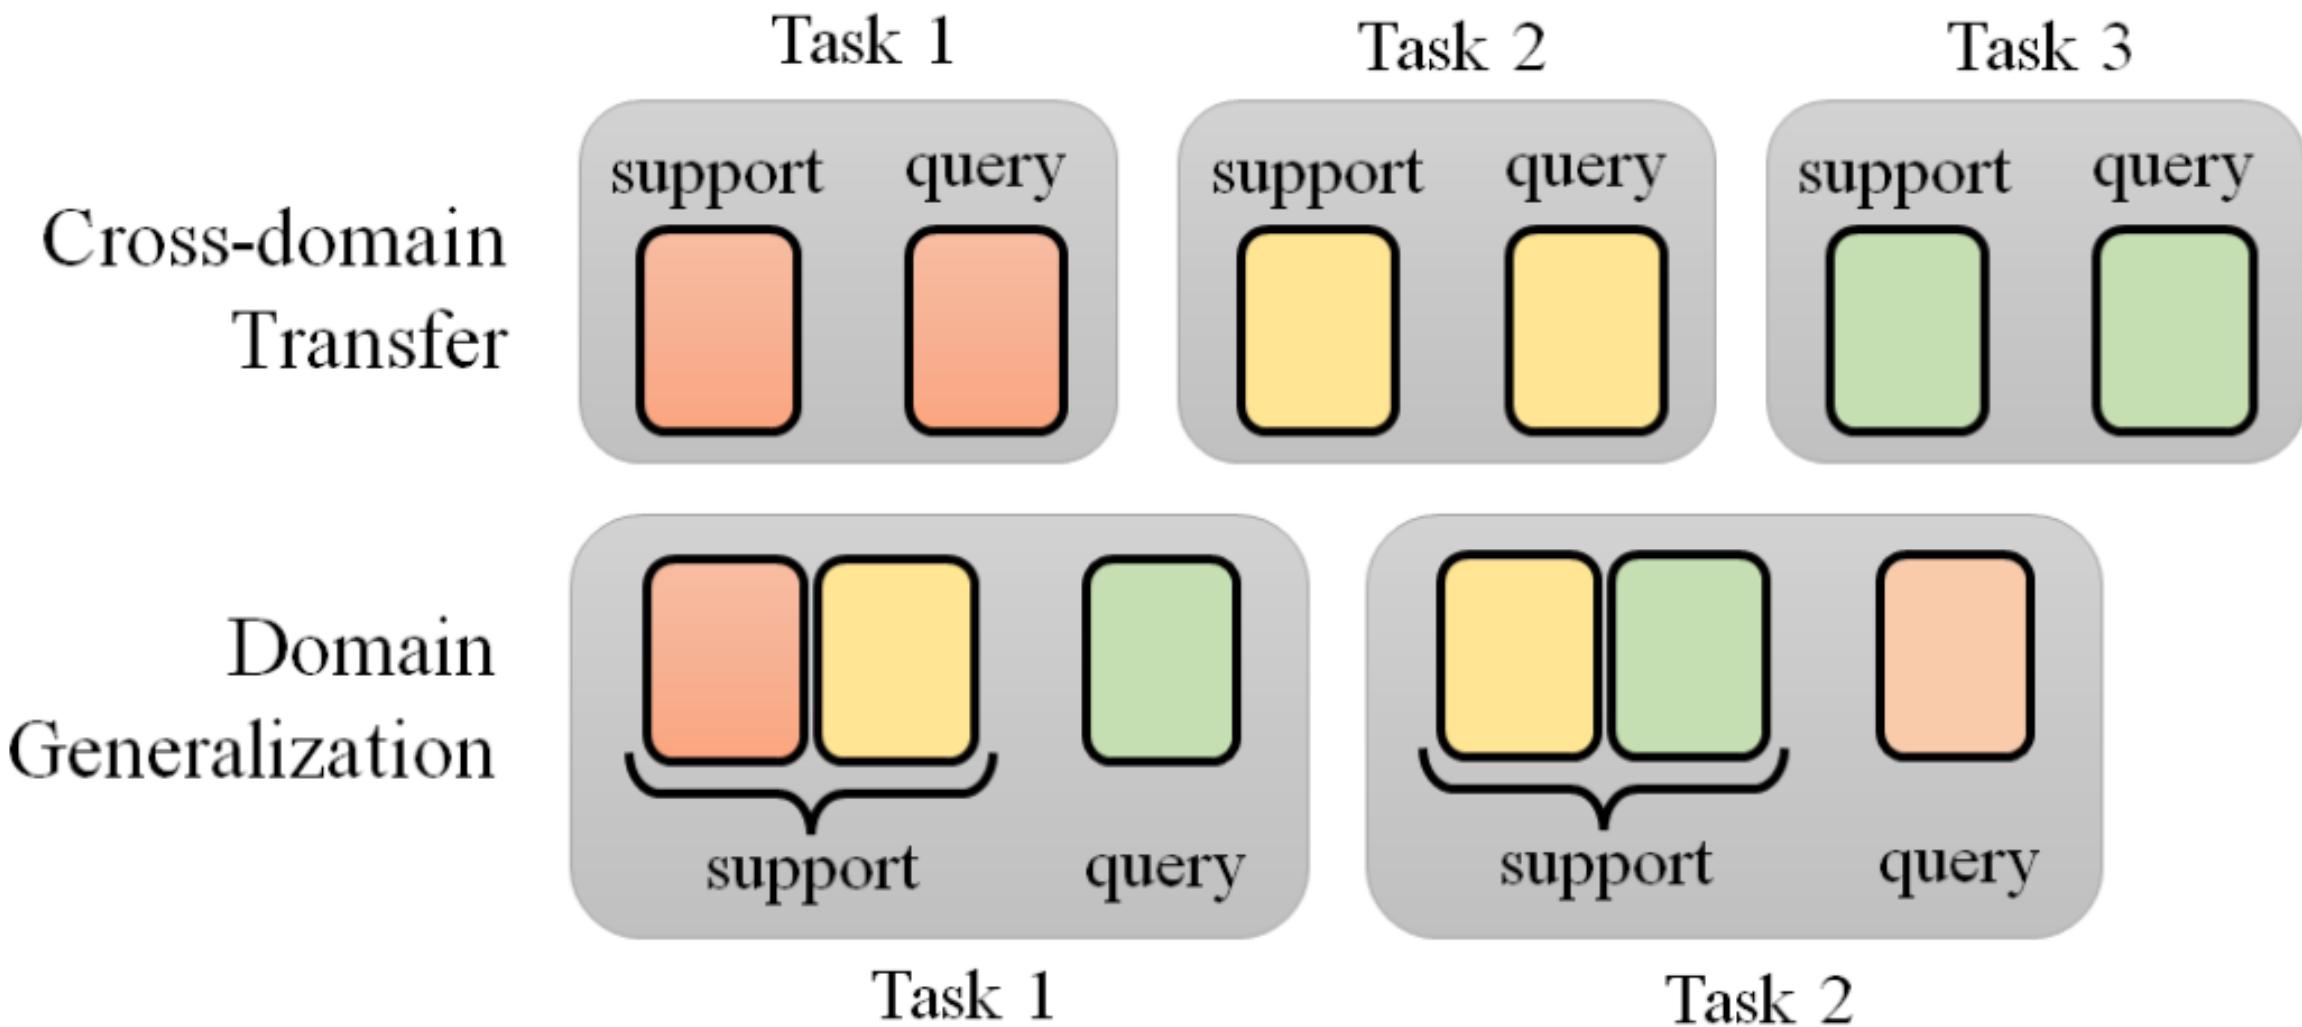
Meta-Task Learning (MTL)

Transfer Learning Strategy for Meta-Learning

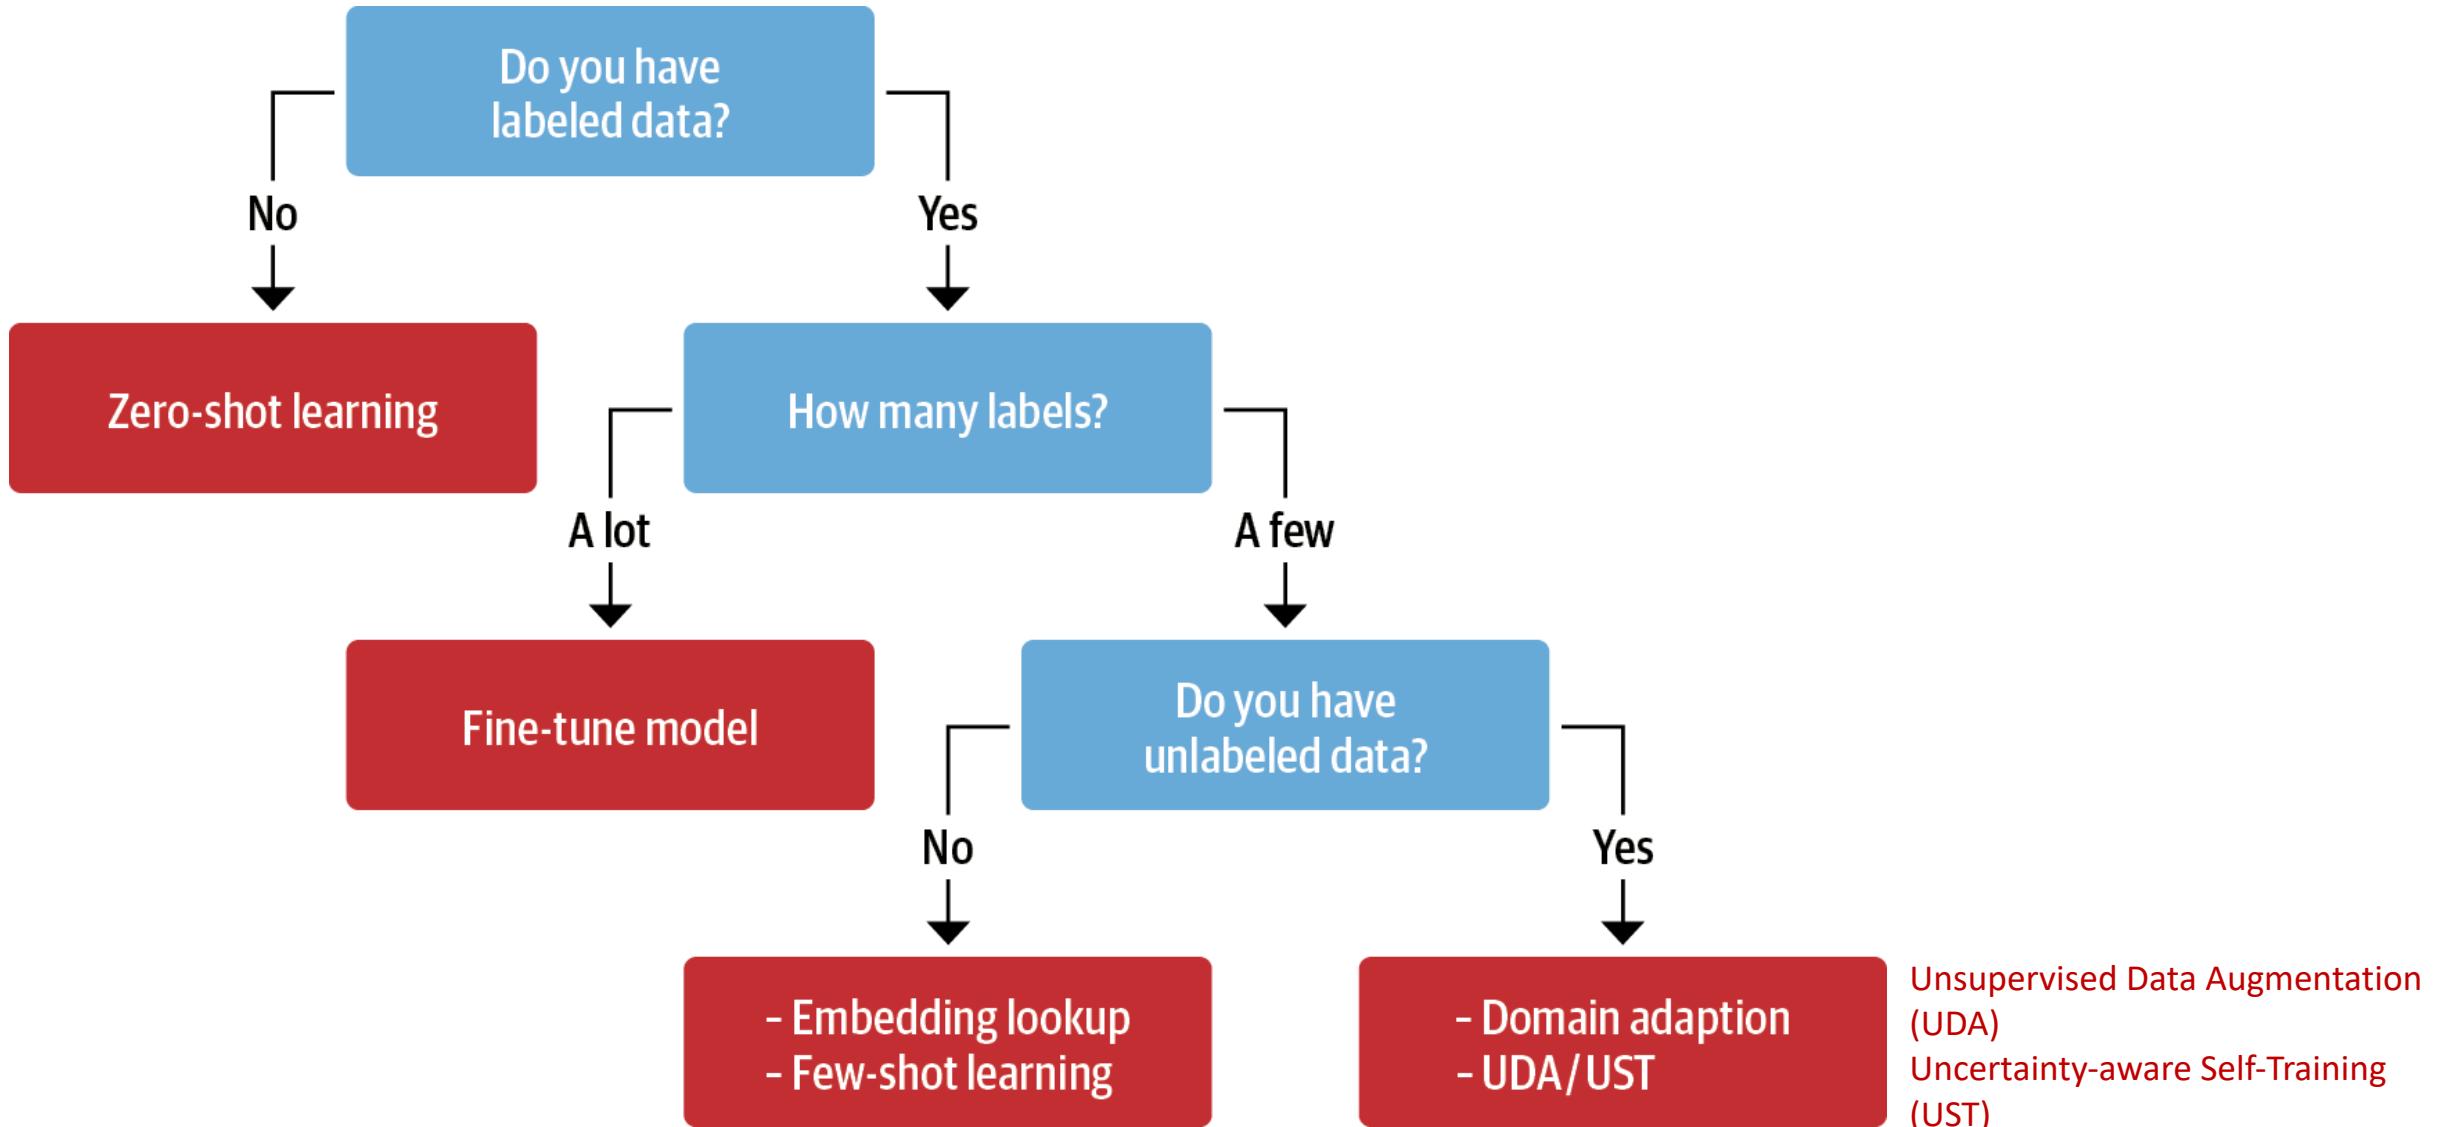
	large-scale training	meta-training	meta-test
Transfer Learning		task ₁ model ₁	task ₂ model ₁ + FT
Meta-Learning		task ₁ model ₁ ... task _N model _N	task _{N+1} model _{N+1}
Meta-Transfer Learning	task model	task ₁ model + SS ₁ + FT ₁ ⋮ task _N model + SS _N + FT _N	task _{N+1} model + SS _N + FT _{N+1}

Meta Learning

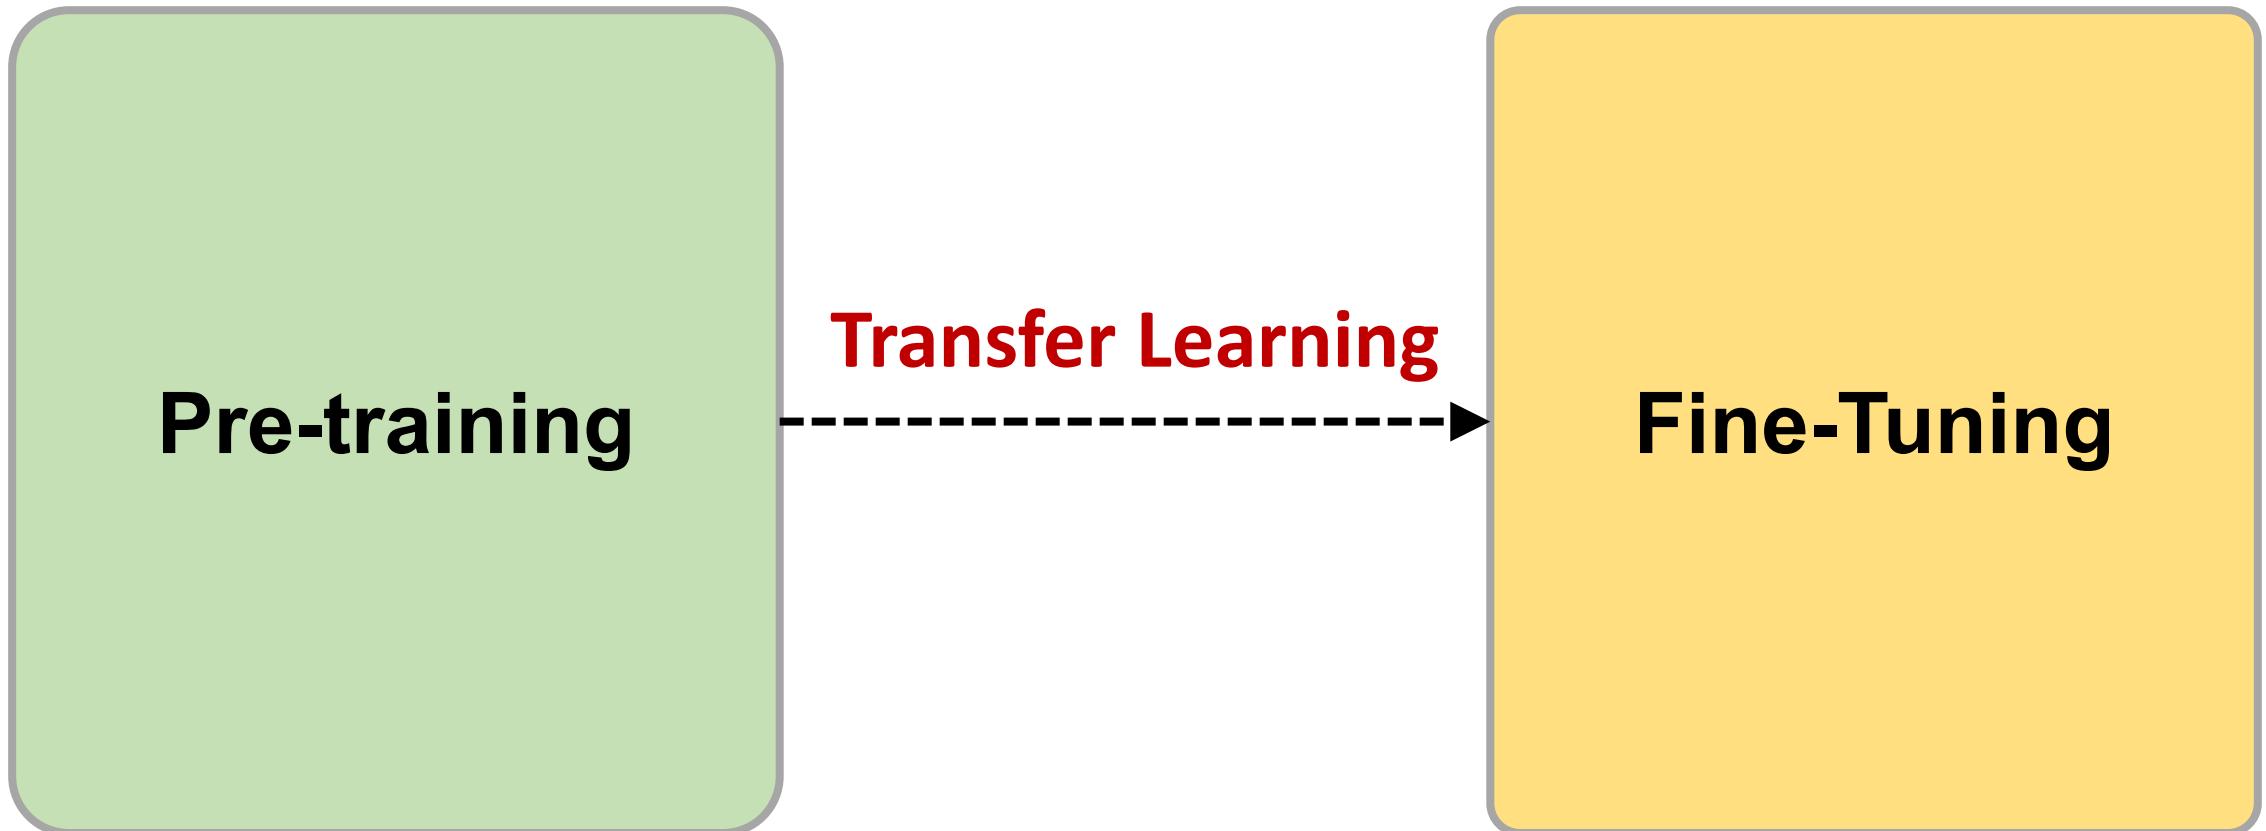
The task construction of cross-domain transfer and domain generalization



Transfer Learning, Fine-tuning, Few-shot learning



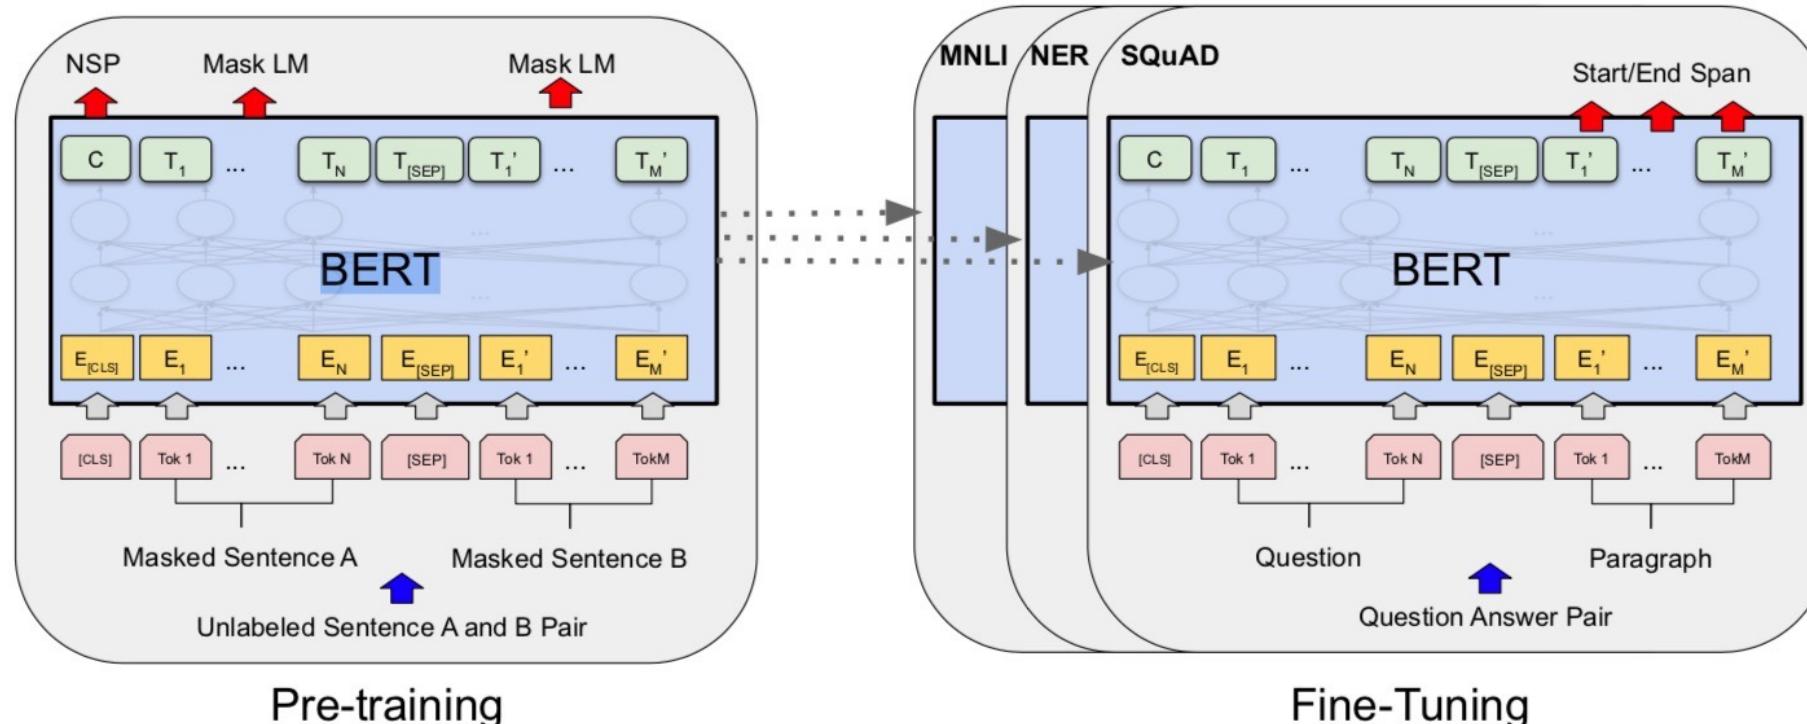
Transfer Learning



BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

BERT (Bidirectional Encoder Representations from Transformers)

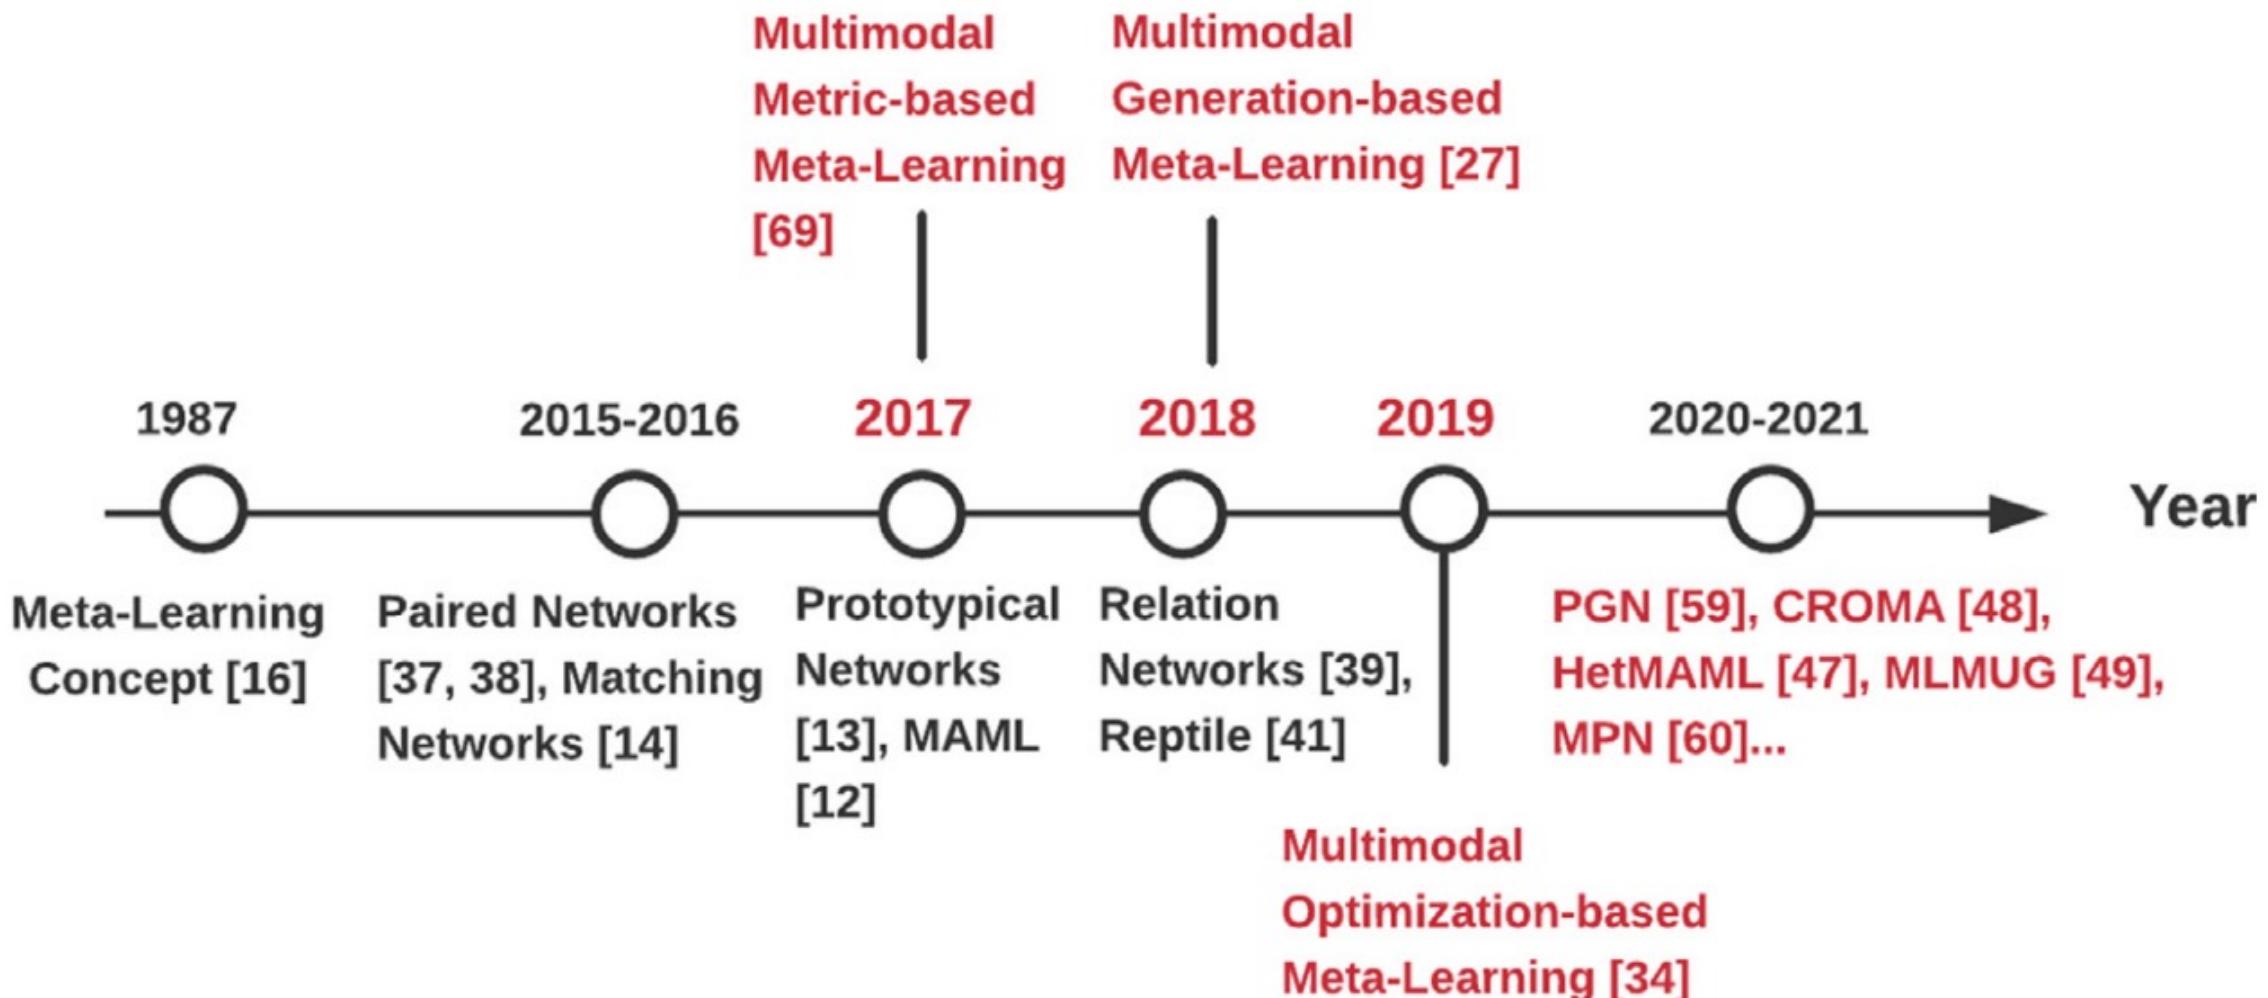
Overall pre-training and fine-tuning procedures for BERT



Source: Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova (2018).

"Bert: Pre-training of deep bidirectional transformers for language understanding." arXiv preprint arXiv:1810.04805.

Meta Learning



Meta Learning

Year	Achievement	Ref.
1987	(1) A new framework of “learning how to learn” with self-referential learning was proposed. The neural networks in self-referential learning can regard their weights as inputs and update them continuously. (2) Based on the conventional neural network, two types of wights were used to connect the neurons. Each type of weight presents a different learning speed.	[34,35]
1990	A synaptic learning rule, which is biologically plausible, was proposed to automatically study the learning rules.	[36]
1993	A chain of meta-networks was introduced to improve the learning capacity of a recurrent neural network for a dynamic environment.	[37]
1995	A framework was proposed to optimize the learning rule within a parametric learning rule space.	[38]
1996	An improved self-referential model was proposed. Time ratios were used to measure the effects of learning processes on the later learning processes.	[39]
1998	The term “Learning to learn” was proposed to equally represent the concept of meta-learning.	[40]
2001	Gradient descent methods were firstly used in meta-learning instead of evolutionary methods, which were widely used in previous research.	[41,42]
2003	A biologically plausible meta-reinforcement learning algorithm was proposed to tune the parameters of the meta-learning model dynamically and adaptively.	[43]
2004	A new perspective of meta-learning was proposed: exploring the interaction between the learning mechanism and the specific contexts to which the mechanism applies.	[9]
2008	The zero-data learning problem was addressed.	[44]
2010–2012	The breakthrough of deep neural networks marks the beginning of the era of meta-learning.	[45–47]
2013	The relationship between transfer learning and meta-learning was described.	[48]
2016	A meta-learning algorithm named gradient descent by gradient descent was proposed.	[49]
2017	(1) MAML was proposed. (2) A doctoral thesis systematically introduced the concept of meta-learning and corresponding methods.	[50,51]
2018	Reptile, an improved version of MAML, was proposed.	[52]
2019	The Capsule network provides a new method to improve the learning capacity of meta-learning, especially in computer vision.	[53]
2020	Combining auto-encoder and capsule network to focus on the zero-shot learning problem.	[54]

Meta-learning Approaches

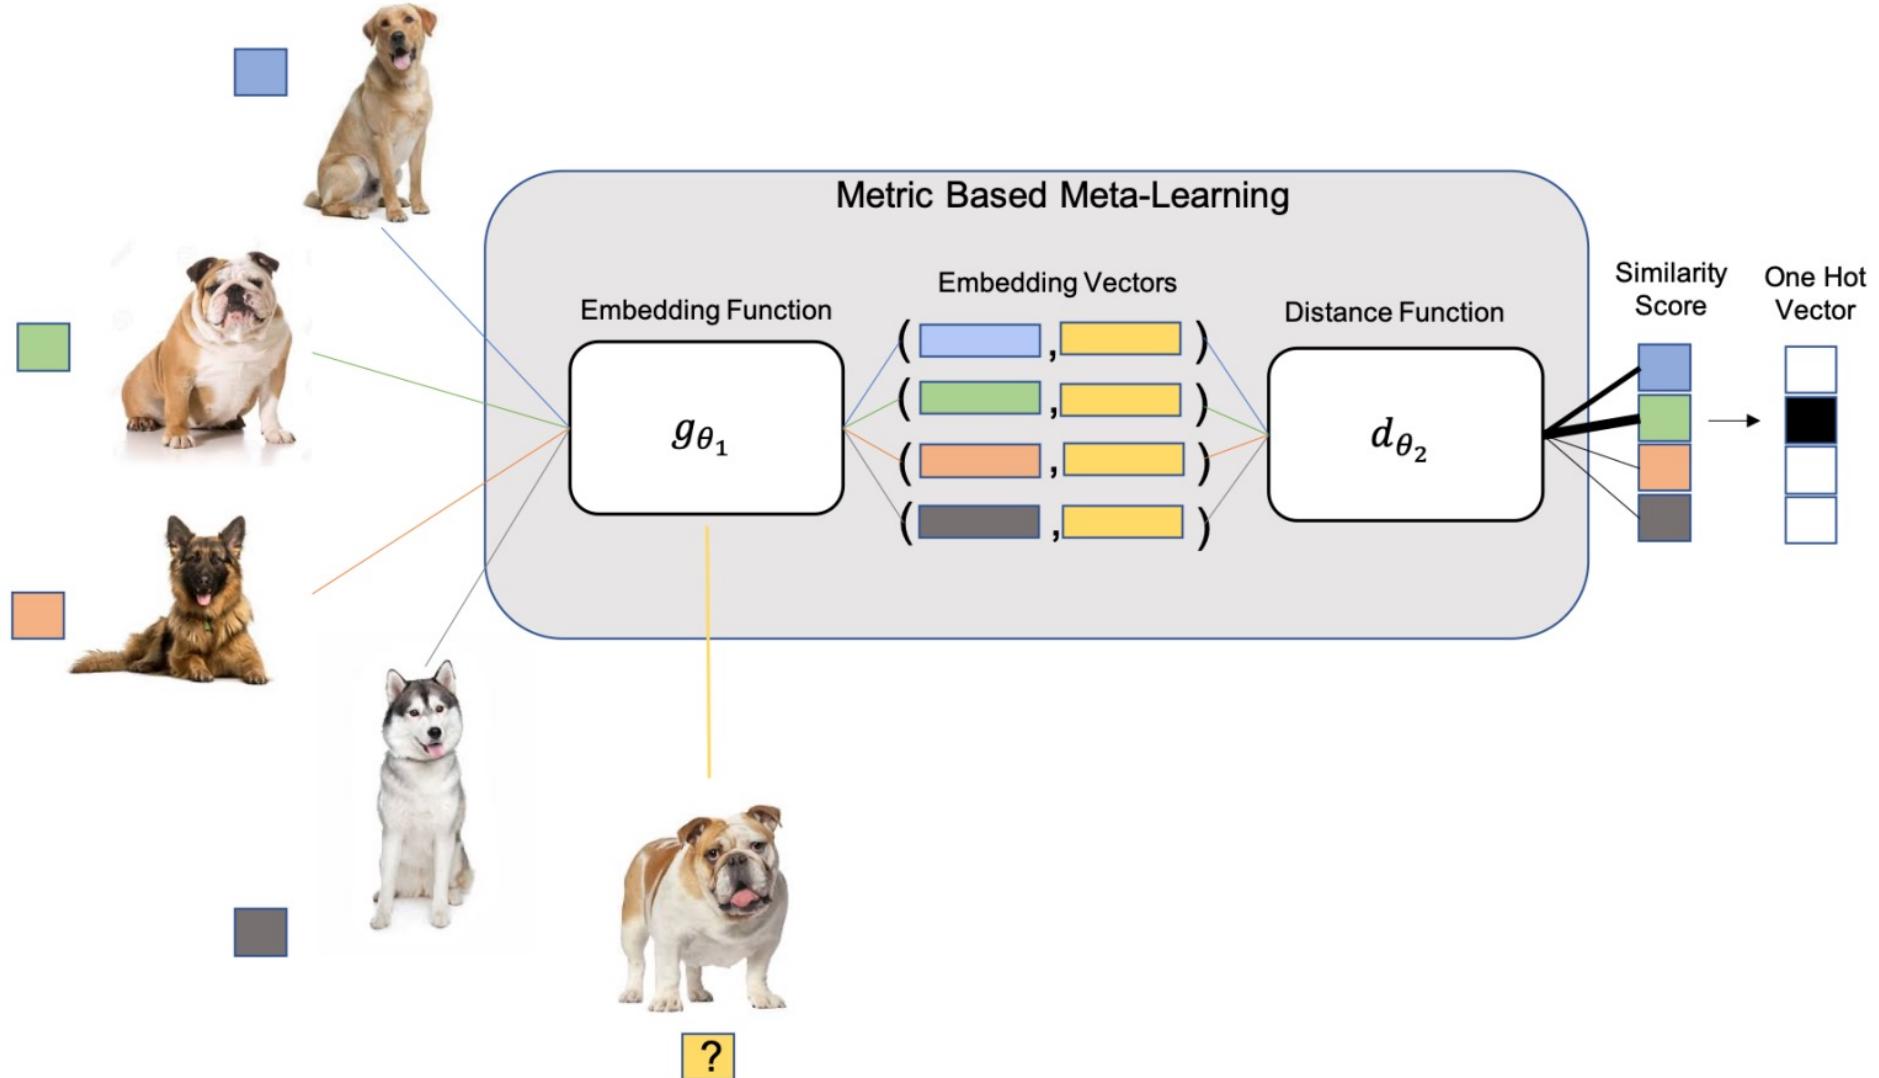
	Metric-based	Optimization-based	Model-based
Key idea	Metric Learning [19]	Gradient Descent	Memory; RNN
How $P_\theta(y x)$ is modeled?	$\sum_{(x_k, y_k) \in S} k_\theta(x, x_k) y_k,$	$P_{\theta'}(y x),$ where $\theta' = g_\phi(\theta, S)$	$f_\theta(x, S).$
Advantages	Faster Inference. Easy to deploy.	Offers flexibility to optimize in dynamic environments. S can be discarded post-optimization.	Faster inference with memory models. Eliminates the need for defining a metric or optimizing at test.
Disadvantages	Less adaptive to optimization in dynamic environments. Computational complexity grows linearly with size of S at test.	Optimization at inference is undesirable for real-world deployment. Prone to overfitting.	Less efficient to hold data in memory as S grows. Hard to design.

Meta Learning: Learning to Learn

Class	Methods	Reference	Summary
Metric-Based	Siamese Neural Networks	[32–36]	We show four metric-based meta-learning algorithms, focusing on feature extractors, similarity metrics, and automatic algorithm selection
	Matching Networks	[37–41]	
	Prototype Networks	[42–46]	However, the metric-based approaches are sensitive to the dataset and
	Relation Networks	[47–53]	increase the computational expenditure when the number of tasks is large.
	Memory-Augmented Neural Networks	[54–56] [57,58]	We display three model-based approaches. MANN combines neural networks with external memory modules, but the model is complex. Meta-Net is computationally intensive and has high memory requirements. SNAIL is relatively simplified, but has to
Model-Based	Meta Networks	[59–65]	be optimized in terms of automatic parameter tuning and reducing computation.
	Simple Neural Attentive Meta-Learner	[66–71]	
	MAML	[72–80]	We present three methods of optimization-based meta-learning. MAML is relatively simple to implement, but the capacity of the model is limited.
Optimization-Based	META- LSTM	[81–86]	Meta-LSTM has a large capacity, but a complicated training process. Meta-SGD has improved capacity but still has difficulties in generalization ability.
	META- SGD	[87–93]	

Metric-based Meta-learning

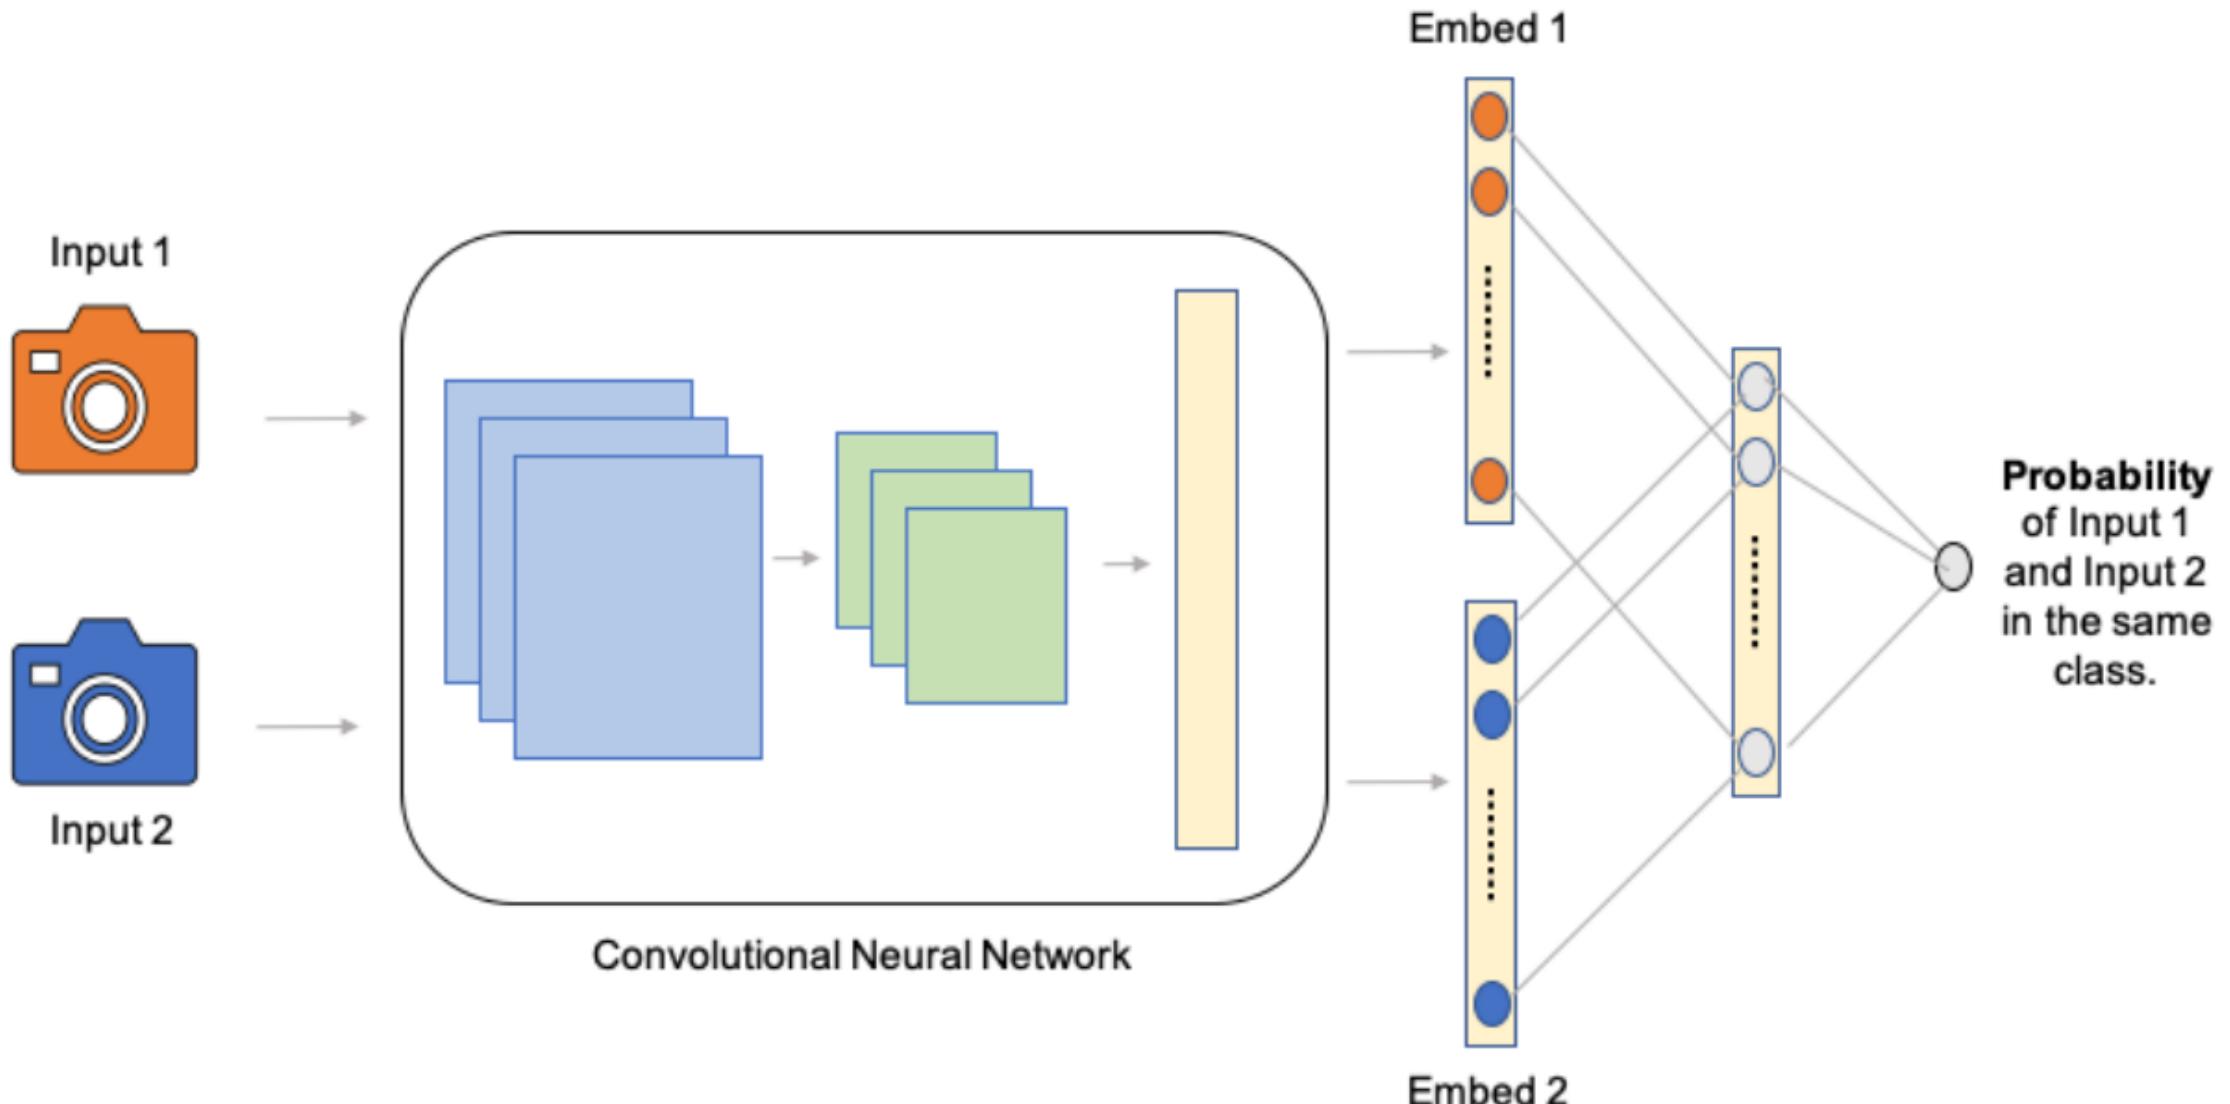
M-Way K-Shot Task (4-way-1-shot classification task)



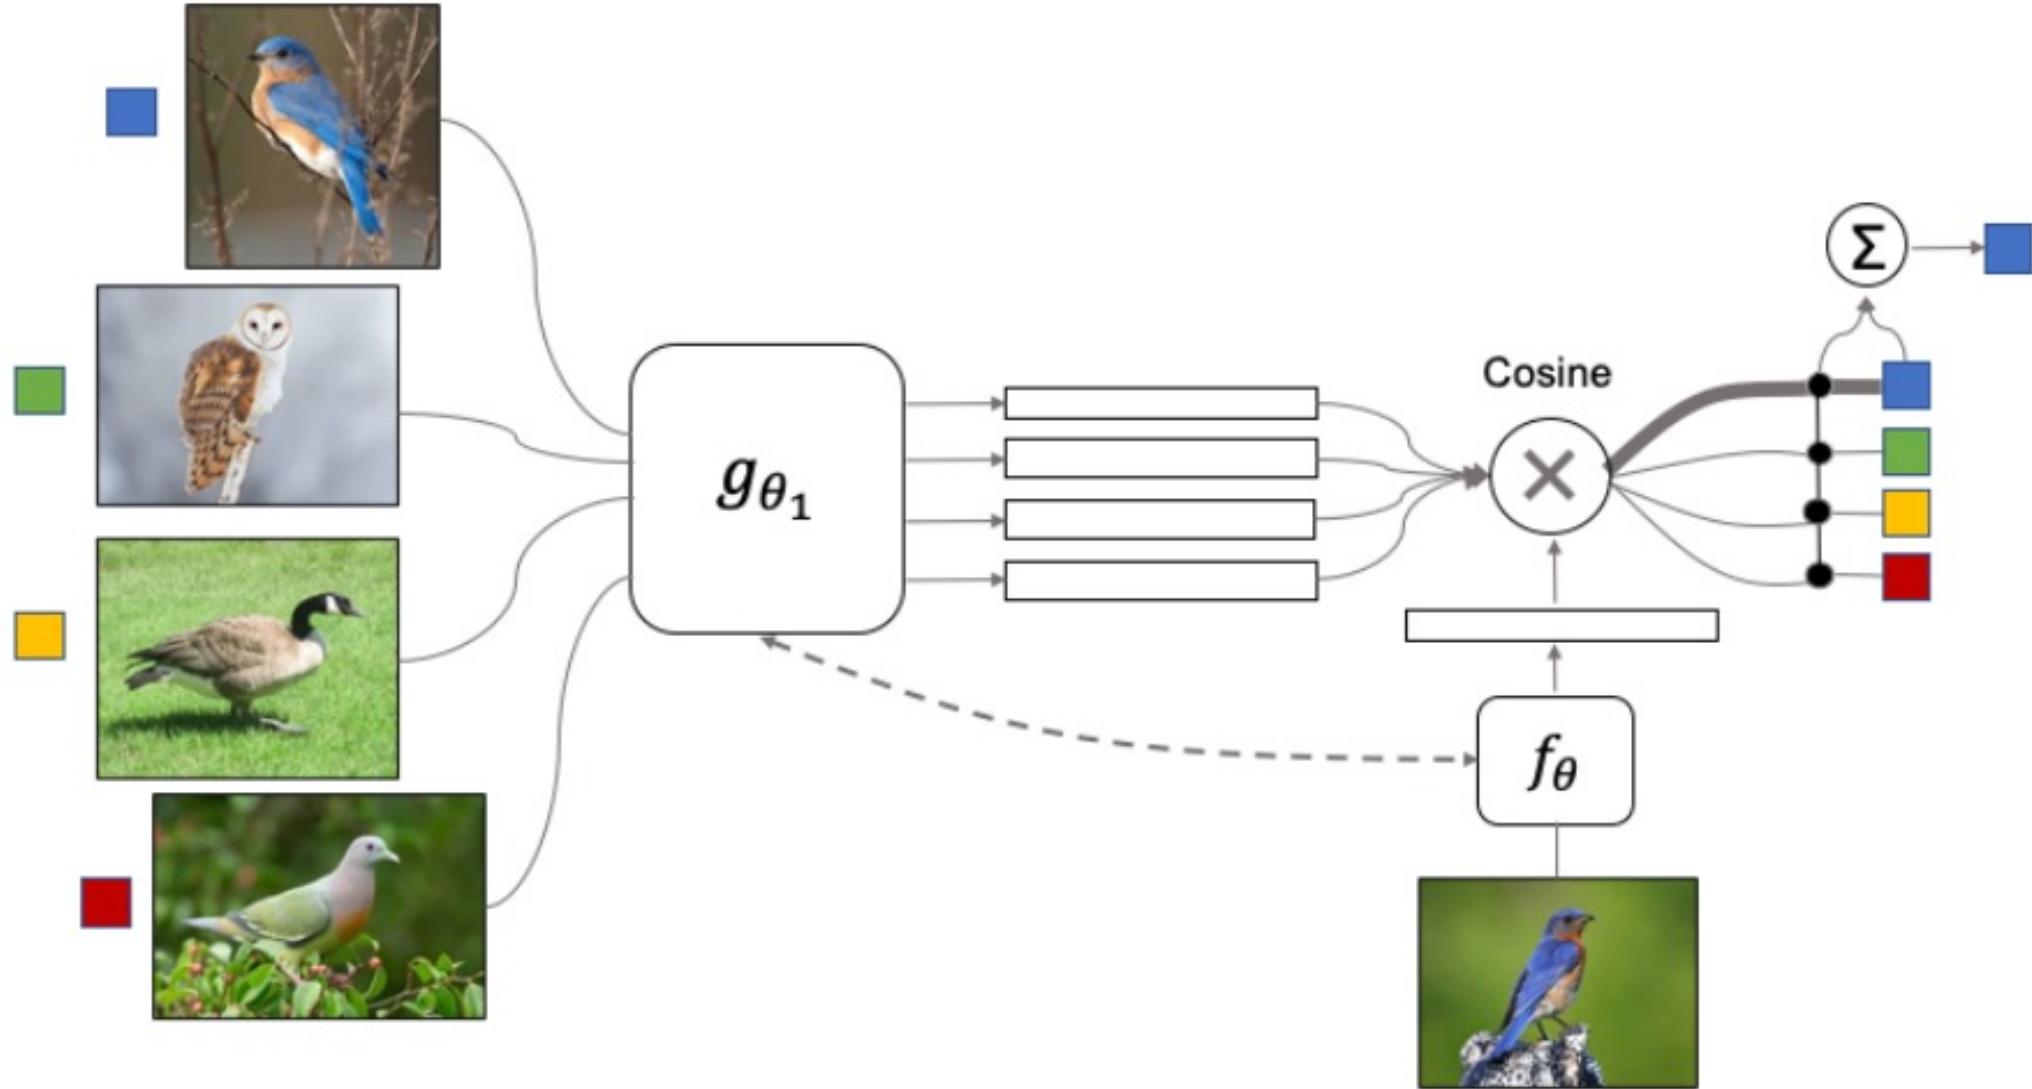
Metric-based Meta-Learning Methods

Method	T.I	g_{θ_1}	d_{θ_2}	Prediction	Loss
Siamese Networks [20]	Yes	CNN	L1	$v = w \cdot d(g_{\theta_1}(x_1), g_{\theta_2}(x_2))$ $p = \text{sigmoid}(\sum_j v_j)$	$-(y \log(p) + (1 - y)(\log(1 - p)))$
Matching Networks [13]	Yes	CNN + LSTM w/ attention	Cosine Similarity	$\hat{y} = \sum_{k=1}^t \sigma(d(f_{\theta}(\hat{x}), g_{\theta_1}(x_k))y_k)$ $P(y = c \hat{x}) = \hat{y}_c$	$-\log P$
Prototypical Networks [21]	Yes	CNN	Euclidean	$P(y = c x) = \sigma(-d(g_{\theta_1}(x), \mathbf{v}_c))$	$-\log P$
Relation Networks [22]	Yes	CNN	Learned by CNN	$r_c = d_{\theta_2}(g_{\theta_1}(x) \oplus \mathbf{v}_c))$	$\sum_{c \in C} (r_c - 1(y == c))^2$
TADAM [16]	No	ResNet-12	Cosine / Euclidean	$P_{\lambda}(y = c x) = \sigma(-\lambda d(g_{\theta_1}(x, \Gamma), \mathbf{v}_c))$	$-\log P$
TapNet [23]	No	Resnet-12	Euclidean	$P(y = c x) = \sigma(-d(\mathbf{M}(g_{\theta_1}(x)), \mathbf{M}(\Phi_c)))$	$-\log P$
CTM [24]	No	Any	Any	-	-

Convolutional Siamese Network

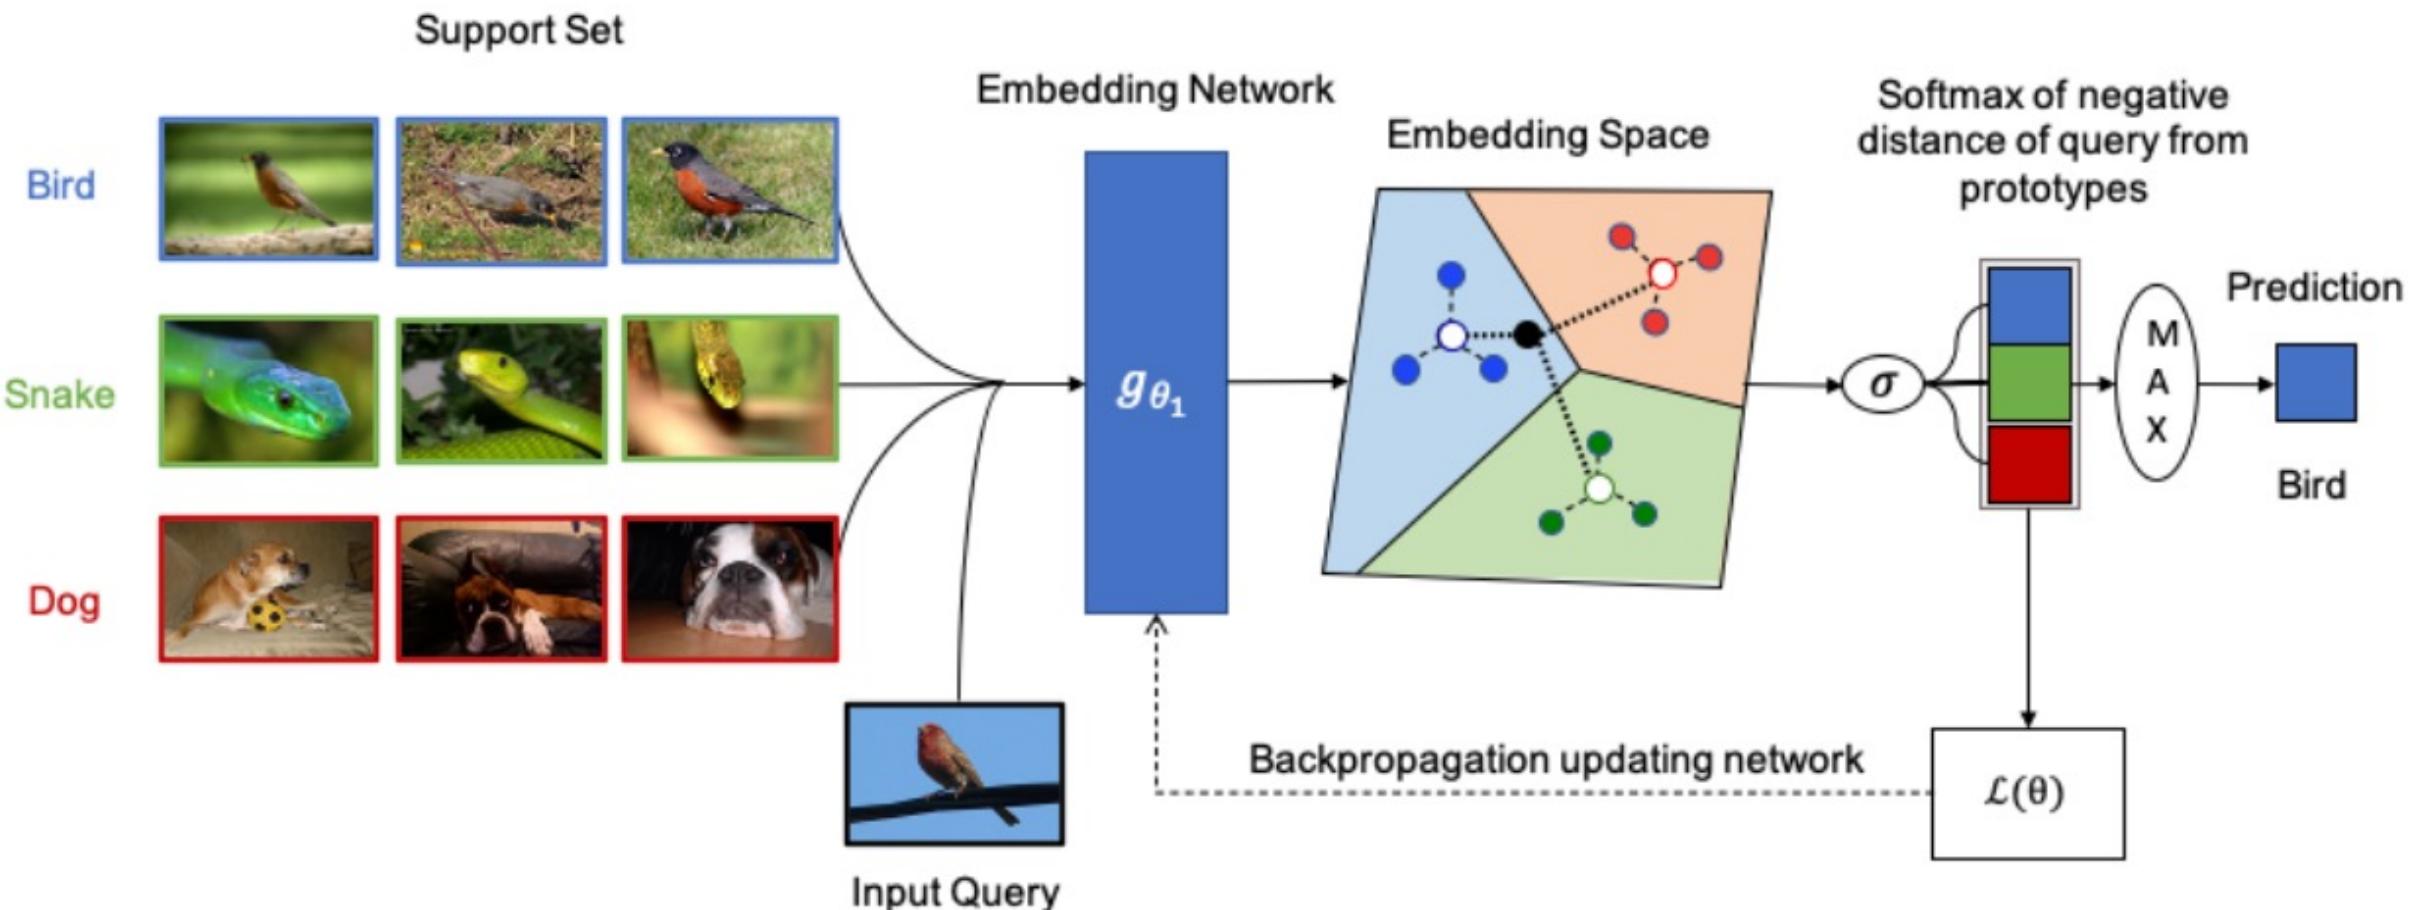


Meta Learning: Matching Networks

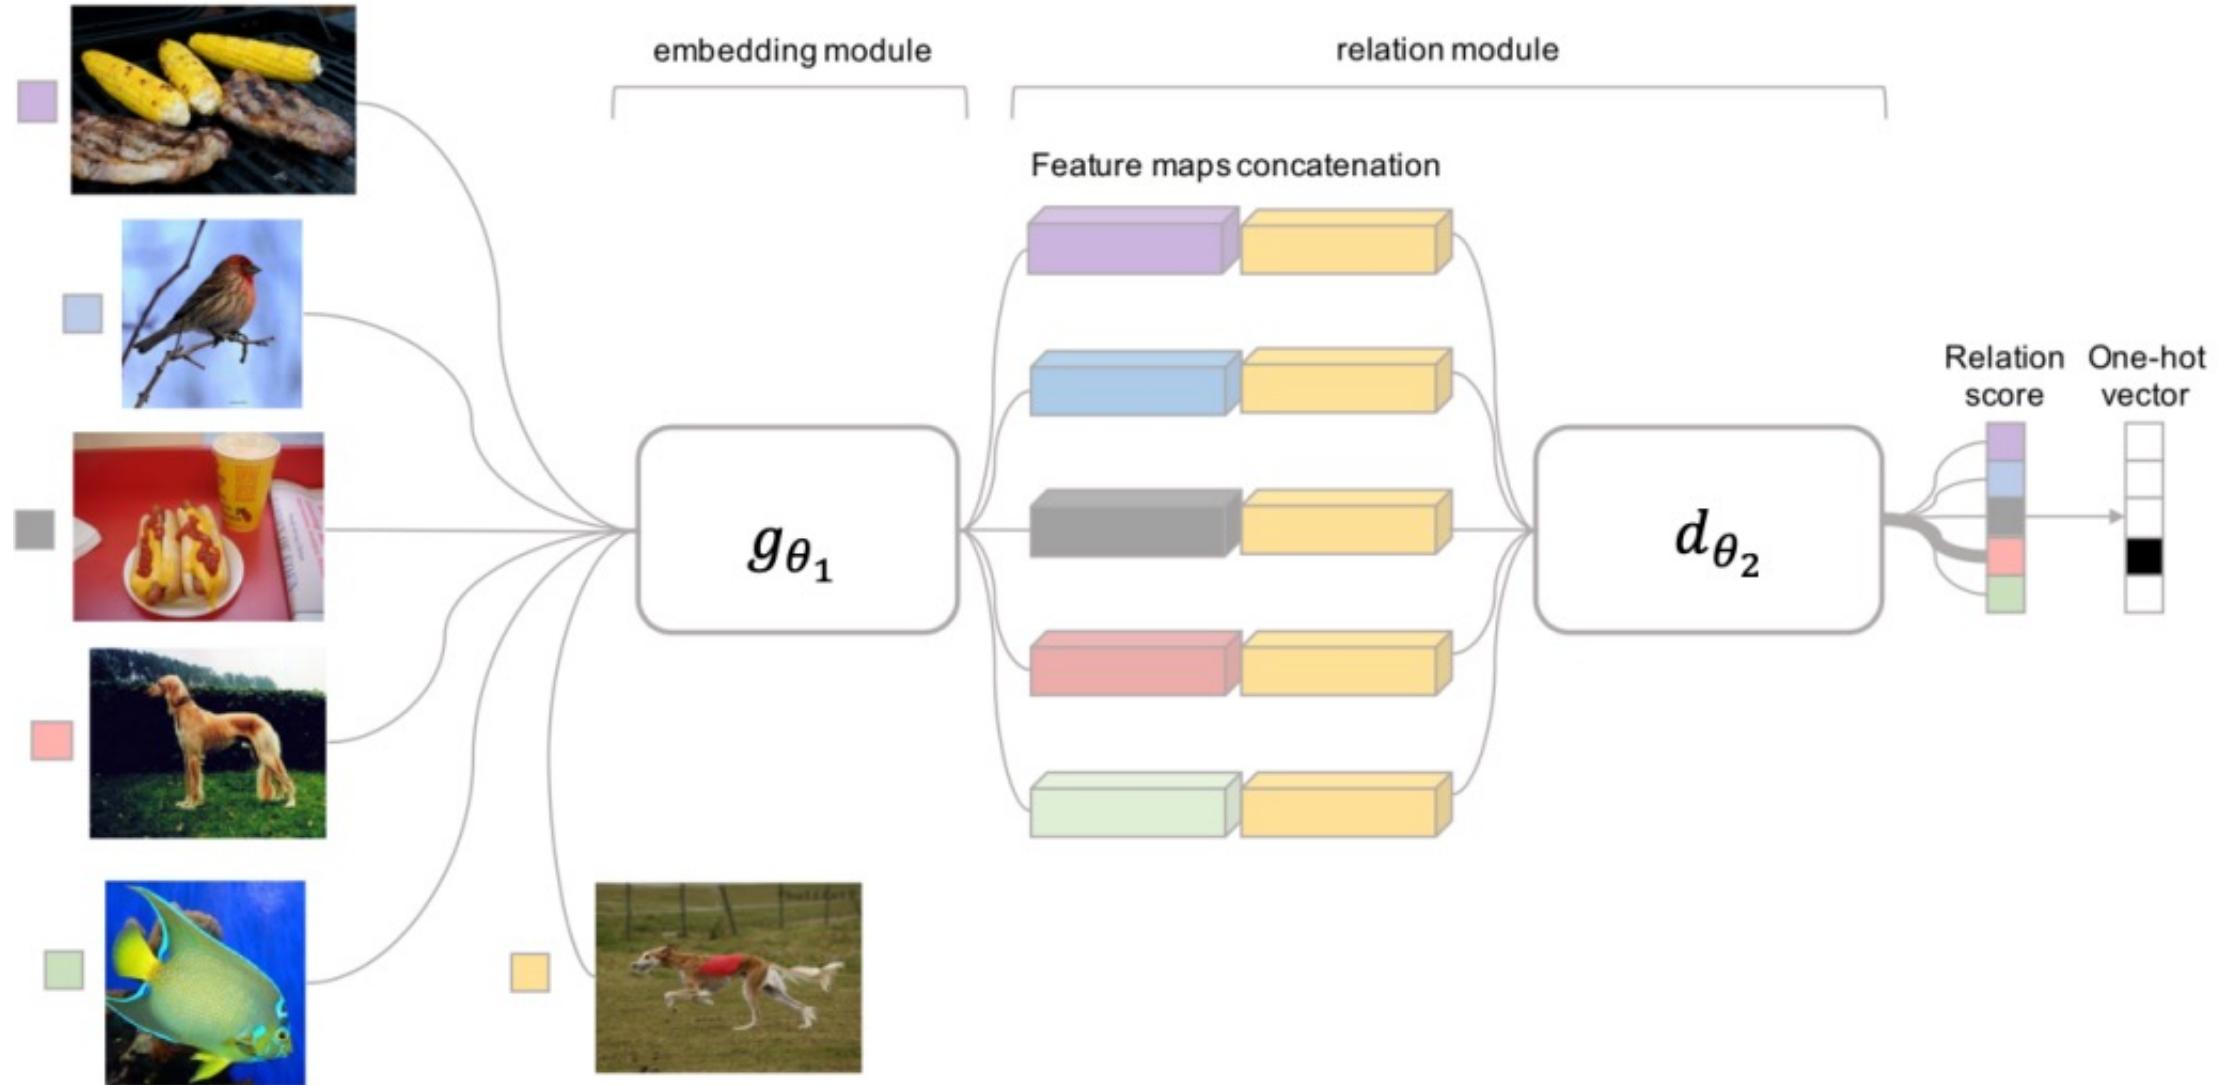


Few-shot Prototypes

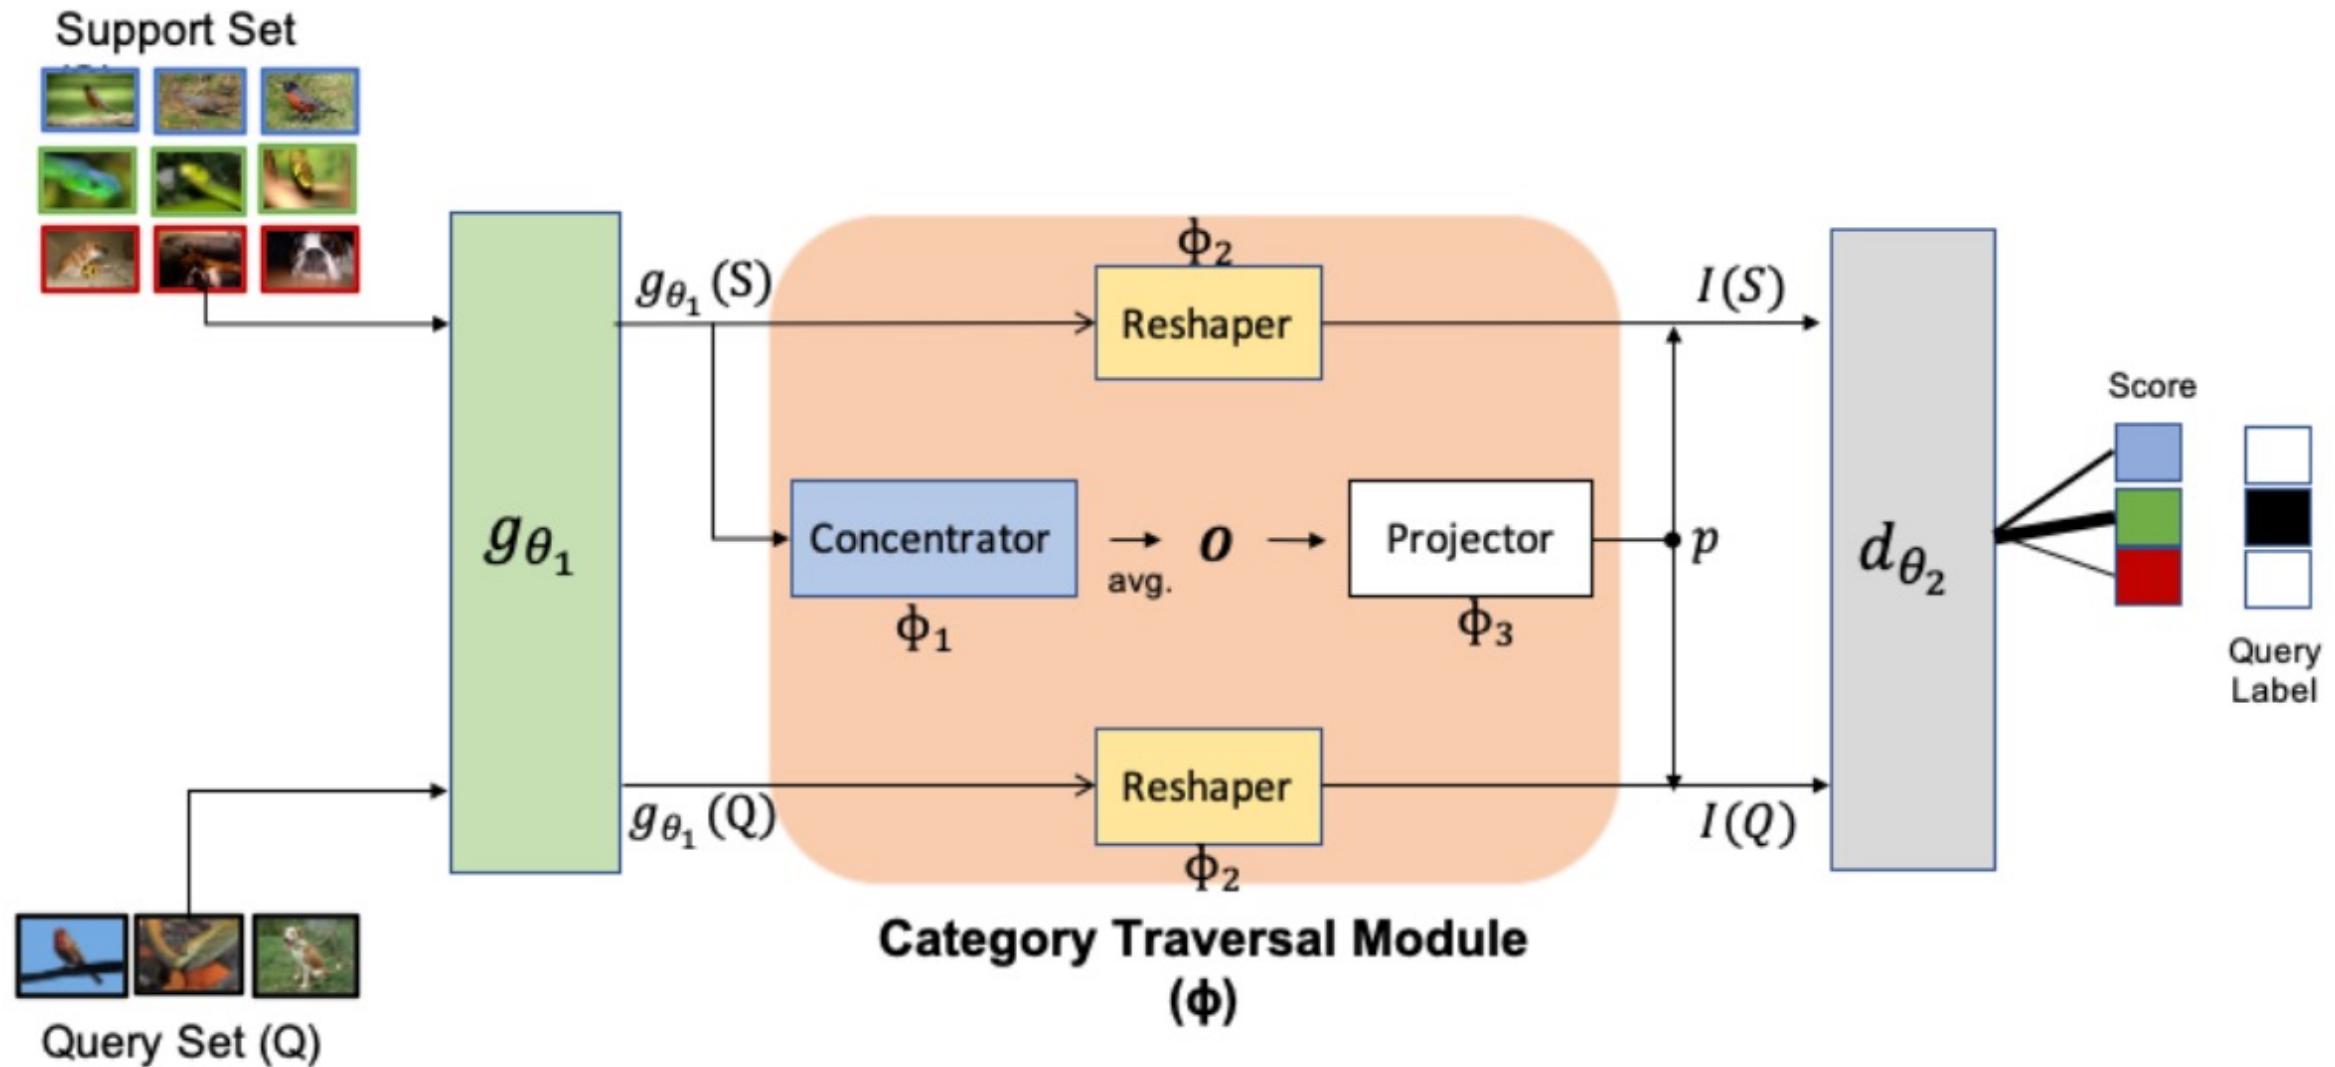
v_c are computed as the mean of embedded support examples for each class



Meta Learning: Relation Network



Meta Learning: Category Traversal Module (CTM)

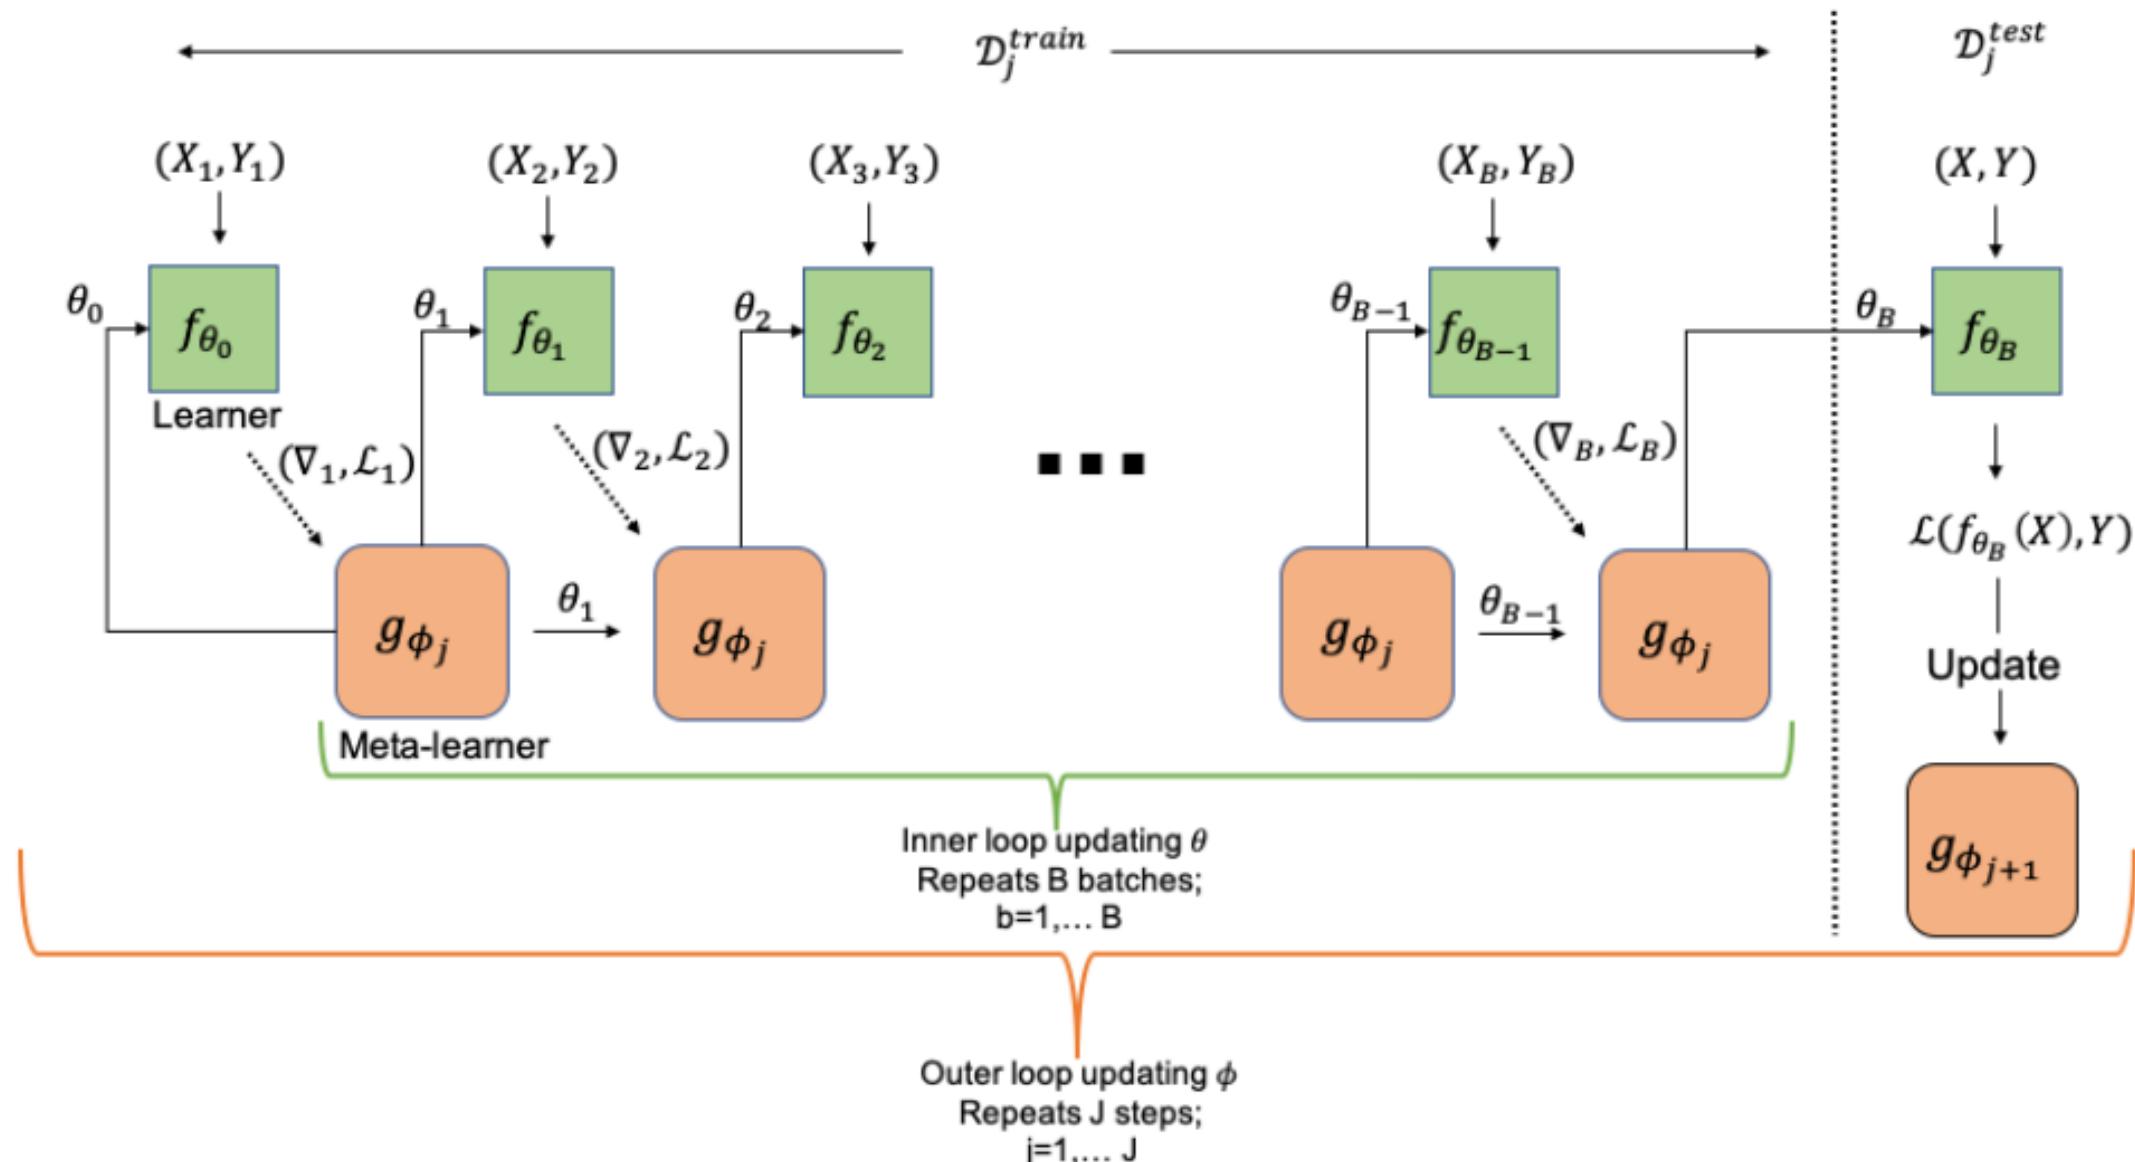


Optimization-based Meta-Learning Methods

Method	Learner	Meta-Learner
LSTM Meta-Learner [35]	Repeat $\forall b \in [1..B]$ $\mathcal{L}_b \leftarrow \mathcal{L}(f(X_b; \theta_{b-1}), Y_b)$ $\theta_b \leftarrow g((\nabla_{\theta_{b-1}} \mathcal{L}_b, \mathcal{L}_b); \phi_{j-1})$	Repeat $\forall j \in [1..J]$ $\mathcal{L}_j^{test} \leftarrow \mathcal{L}(f(X; \theta_B), Y)$ $\phi_j \leftarrow \phi_{j-1} - \alpha \nabla_{\phi_{j-1}} \mathcal{L}_j^{test}$
MAML [14]	Repeat $\forall i \in [1..I]$ $\mathcal{L}_i^{train} \leftarrow \mathcal{L}(f(\mathcal{D}_i^{train}; \theta_{j-1}))$ $\theta_i^* \leftarrow \theta_{j-1} - \alpha \nabla_{\theta_{j-1}} \mathcal{L}_i^{train}$ $\mathcal{L}_i^{test} \leftarrow \mathcal{L}(f(\mathcal{D}_i^{test}; \theta_i^*))$	Repeat $\forall j \in [1..J]$ $\theta_j \leftarrow \theta_{j-1} - \beta \nabla_{\theta_{j-1}} \sum_{i=1}^I \mathcal{L}_i^{test}$
MTL [37]	$\mathcal{L}_i^{train} \leftarrow \mathcal{L}(f(\mathcal{D}_i^{train}; [\theta_{j-1}, \phi_{j-1}, \Theta]))$ $\theta_i^* \leftarrow \theta_{j-1} - \alpha \nabla_{\theta_{j-1}} \mathcal{L}_i^{train}$ $\mathcal{L}_i^{test} \leftarrow \mathcal{L}(f(\mathcal{D}_i^{test}; \theta_i^*))$	$\theta_j \leftarrow \theta_{j-1} - \beta \nabla_{\theta_{j-1}} \sum_{i=1}^I \mathcal{L}_i^{test}$ $\phi_j \leftarrow \phi_{j-1} - \beta \nabla_{\phi_{j-1}} \sum_{i=1}^I \mathcal{L}_i^{test}$
LEO [38]	$\phi_{j-1} = \{\phi_e, \phi_r, \phi_d, \alpha\}$ $\mathbf{z}_i \leftarrow g(\mathcal{D}_i^{train}; [\phi_e, \phi_r, \Theta])$ $\theta_i \leftarrow g(\mathbf{z}_i; \phi_d)$ $\mathcal{L}_i^{train} \leftarrow \mathcal{L}(f(\mathcal{D}_i^{train}; \theta_i))$ $\mathbf{z}_i^* \leftarrow \mathbf{z}_i - \alpha \nabla_{\mathbf{z}_i} \mathcal{L}_i^{train}$ $\theta_i^* \leftarrow g(\mathbf{z}_i^*; \phi_d)$ $\mathcal{L}_i^{test} \leftarrow \mathcal{L}(f(\mathcal{D}_i^{test}; \theta_i^*))$	$\phi_j \leftarrow \phi_{j-1} - \beta \nabla_{\phi_{j-1}} \sum_{i=1}^I \mathcal{L}_i^{test}$

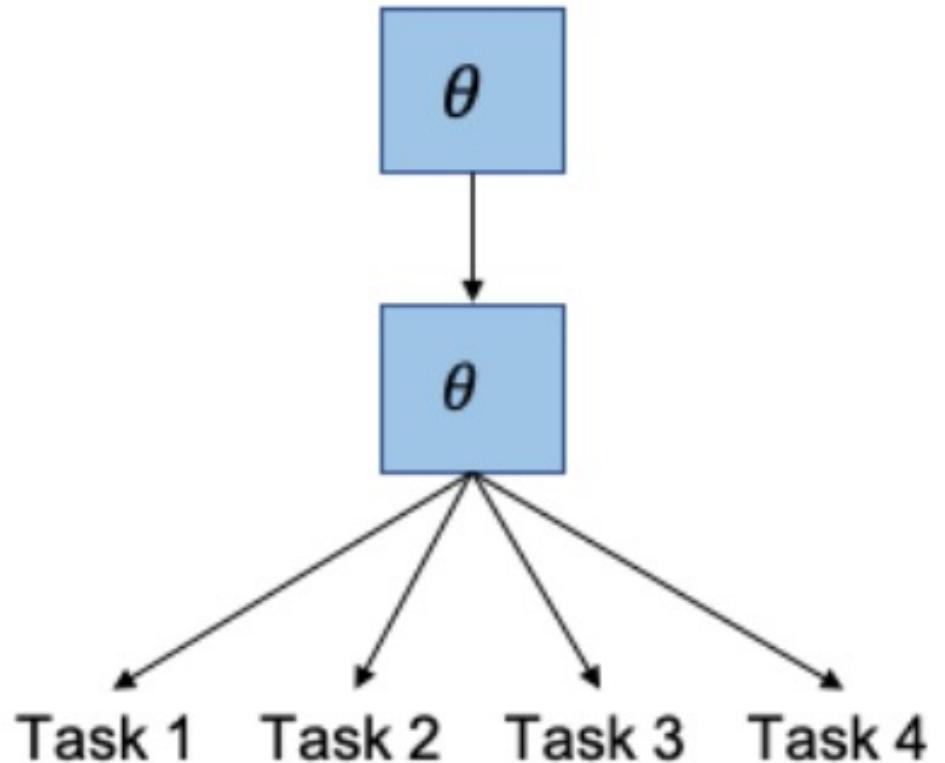
Source: Parnami, Archit, and Minwoo Lee. "Learning from Few Examples: A Summary of Approaches to Few-Shot Learning." arXiv preprint arXiv:2203.04291 (2022).

Computational Graph for the Forward Pass of the Meta-learner

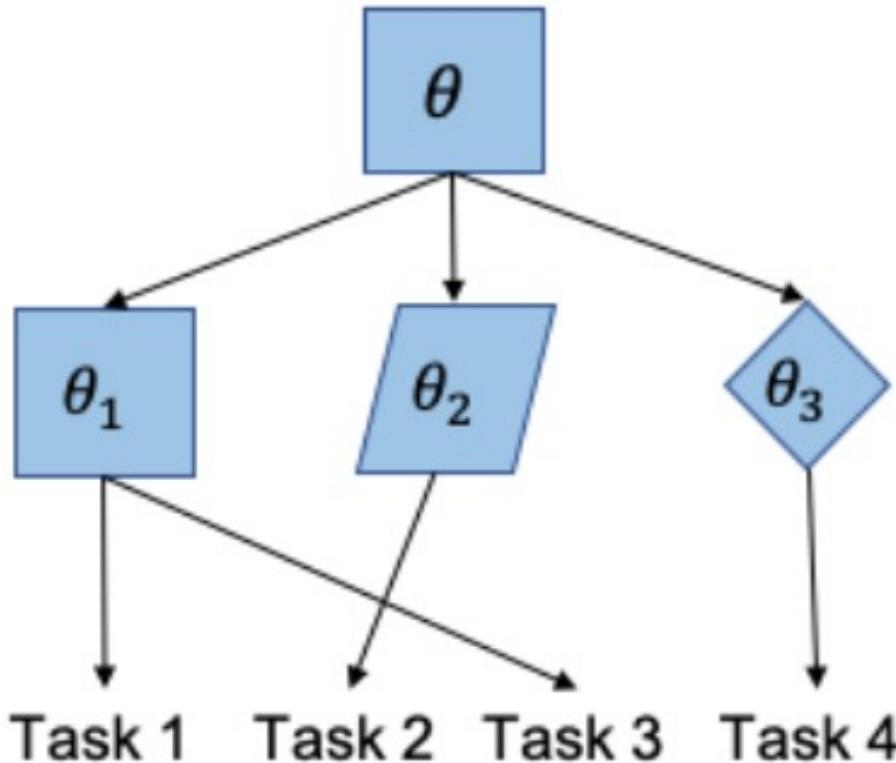


Model-Agnostic Meta-Learning (MAML)

Hierarchically Structured Meta-Learning (HSML)

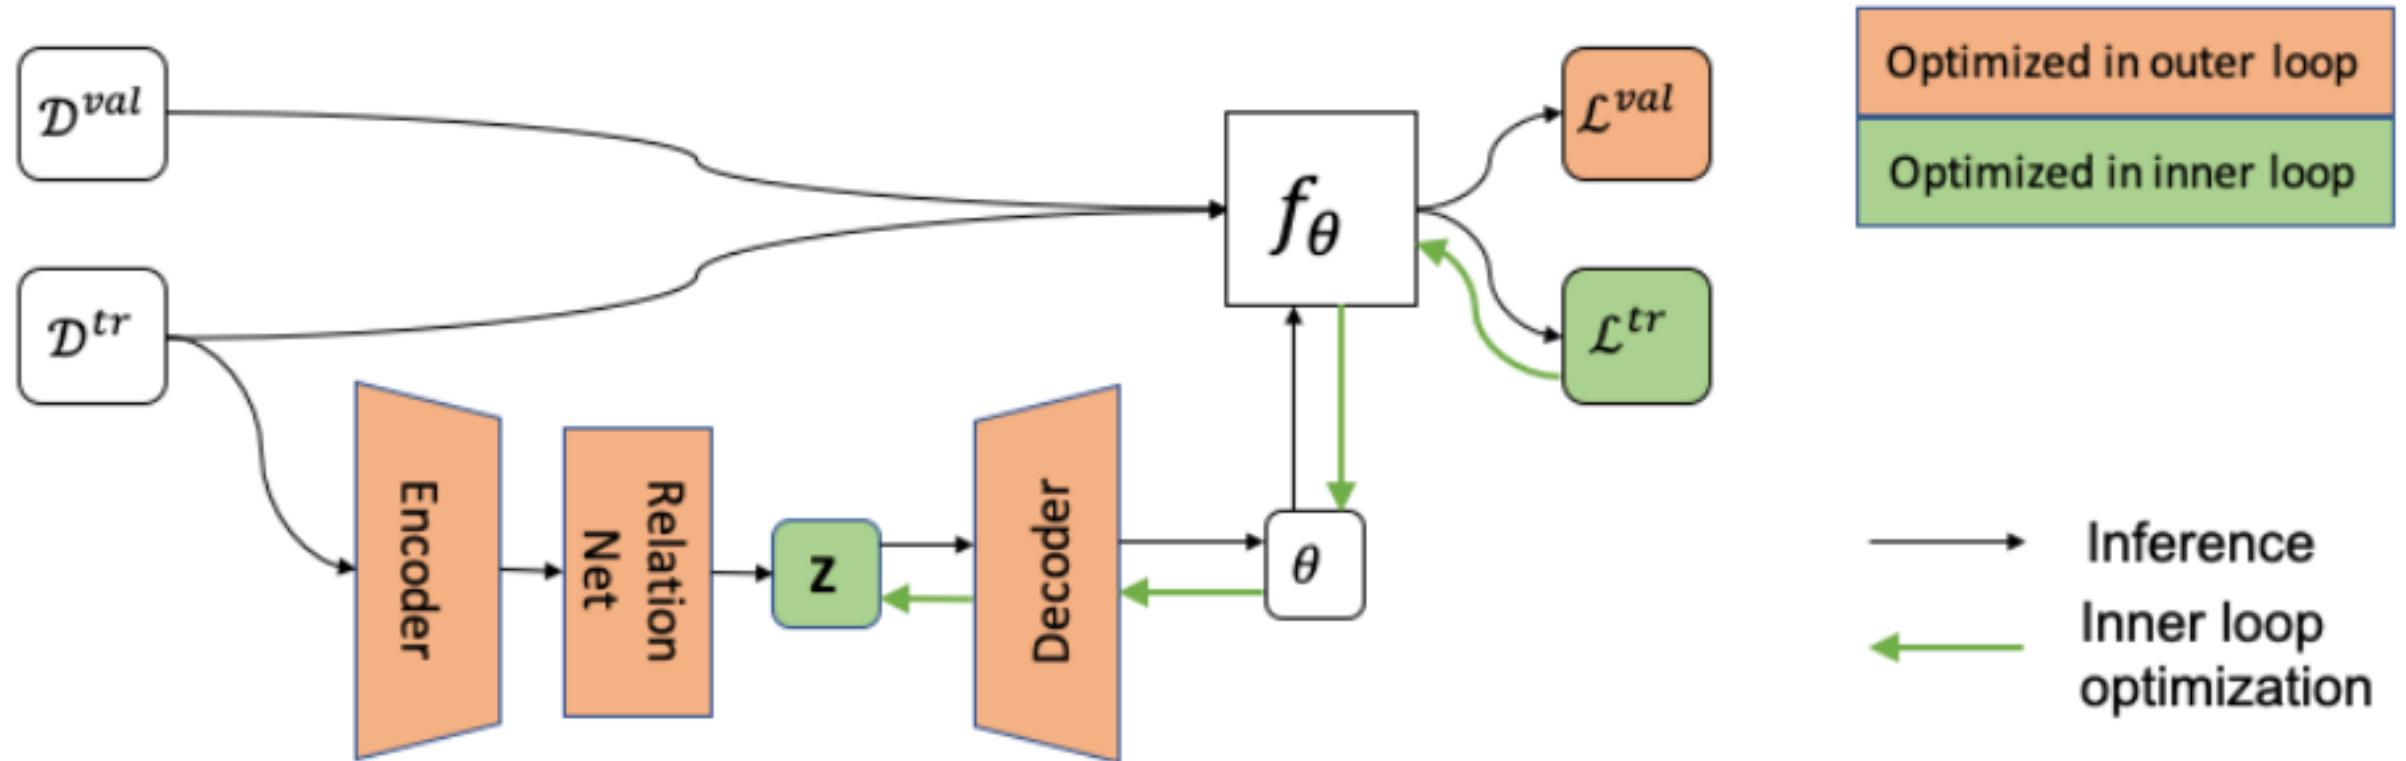


Globally Shared Initialization
(MAML)

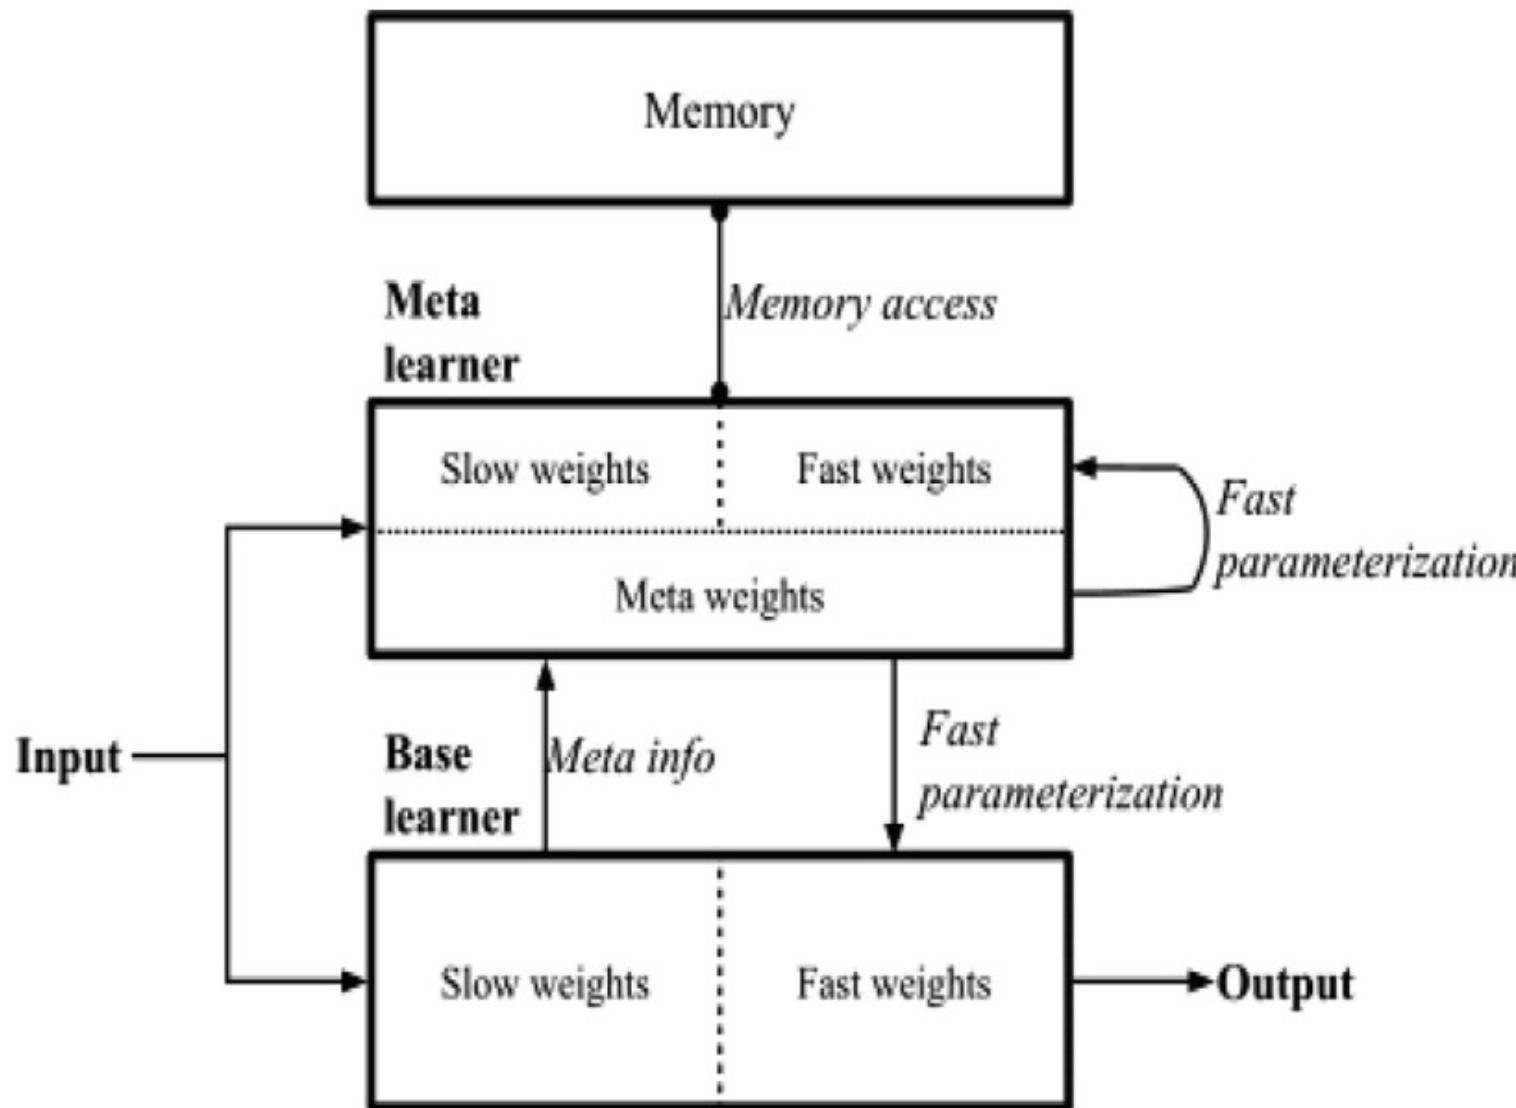


Hierarchically Clustered Initialization
(HSML)

Latent Embedding Optimization (LEO)

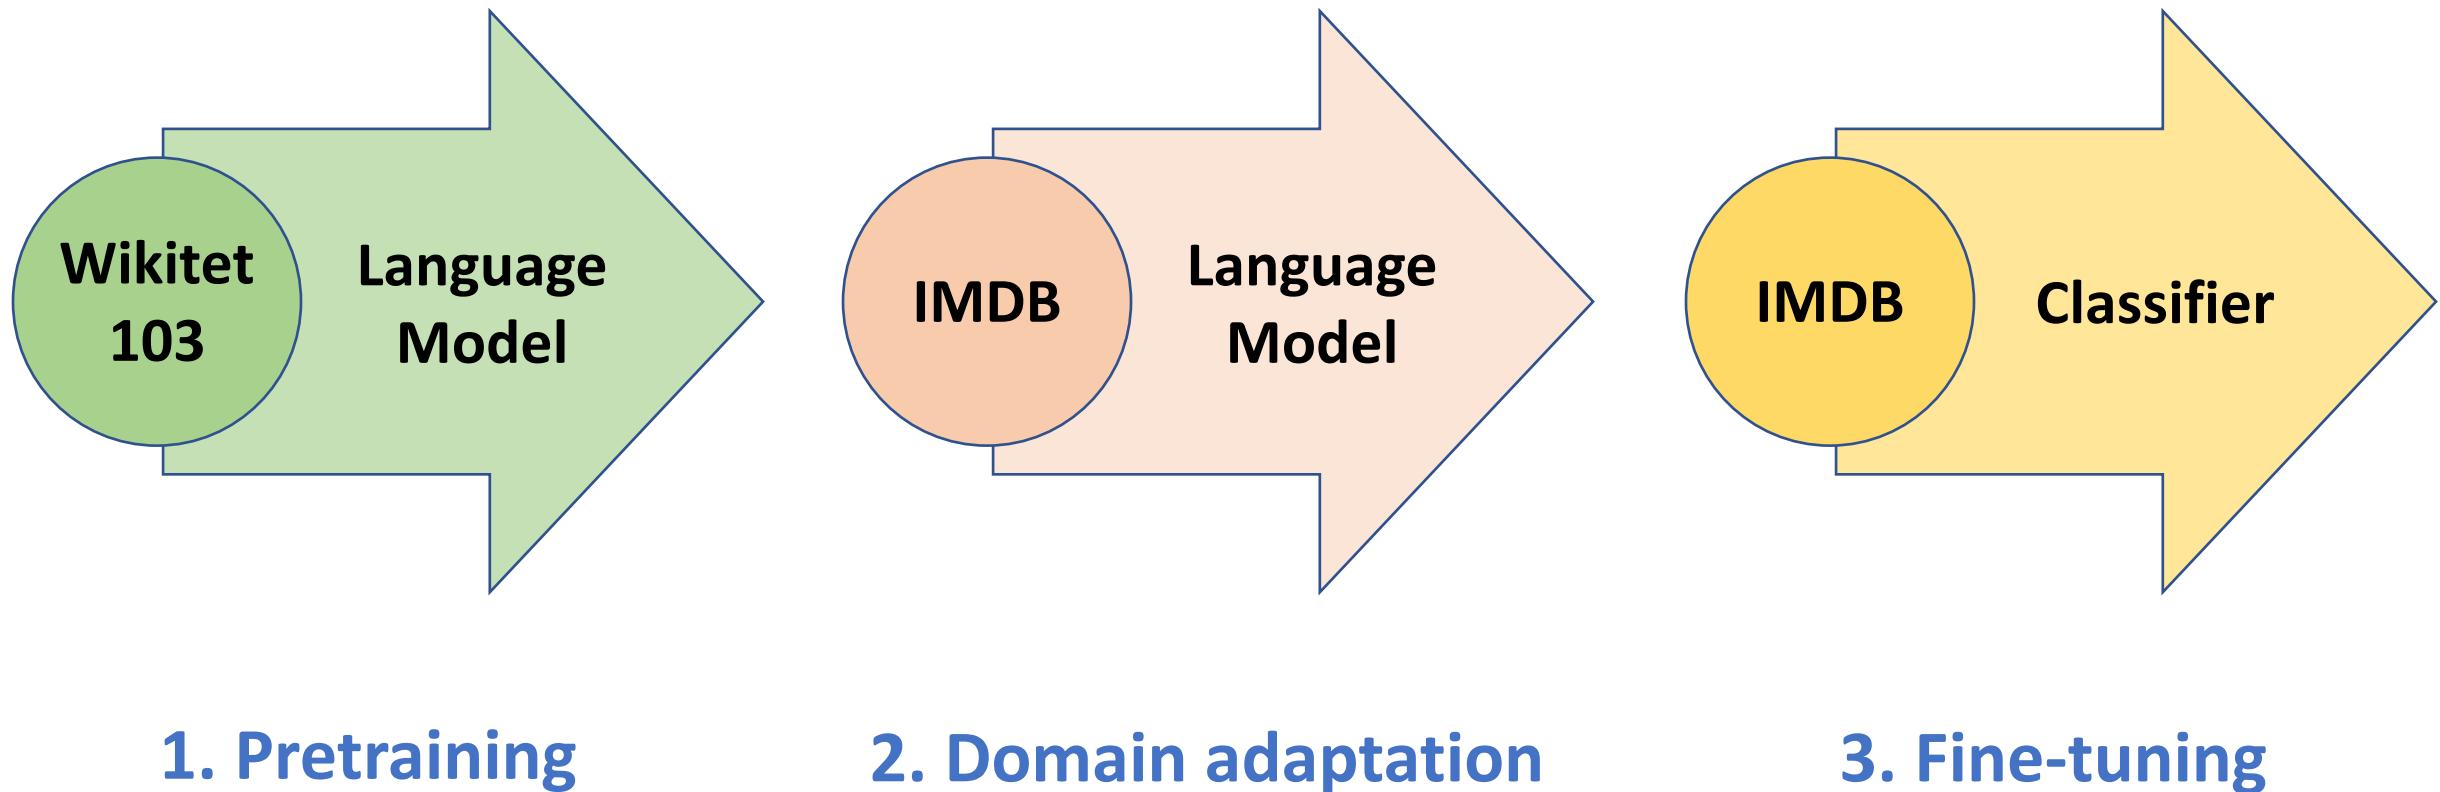


Overall Architecture of Meta Networks



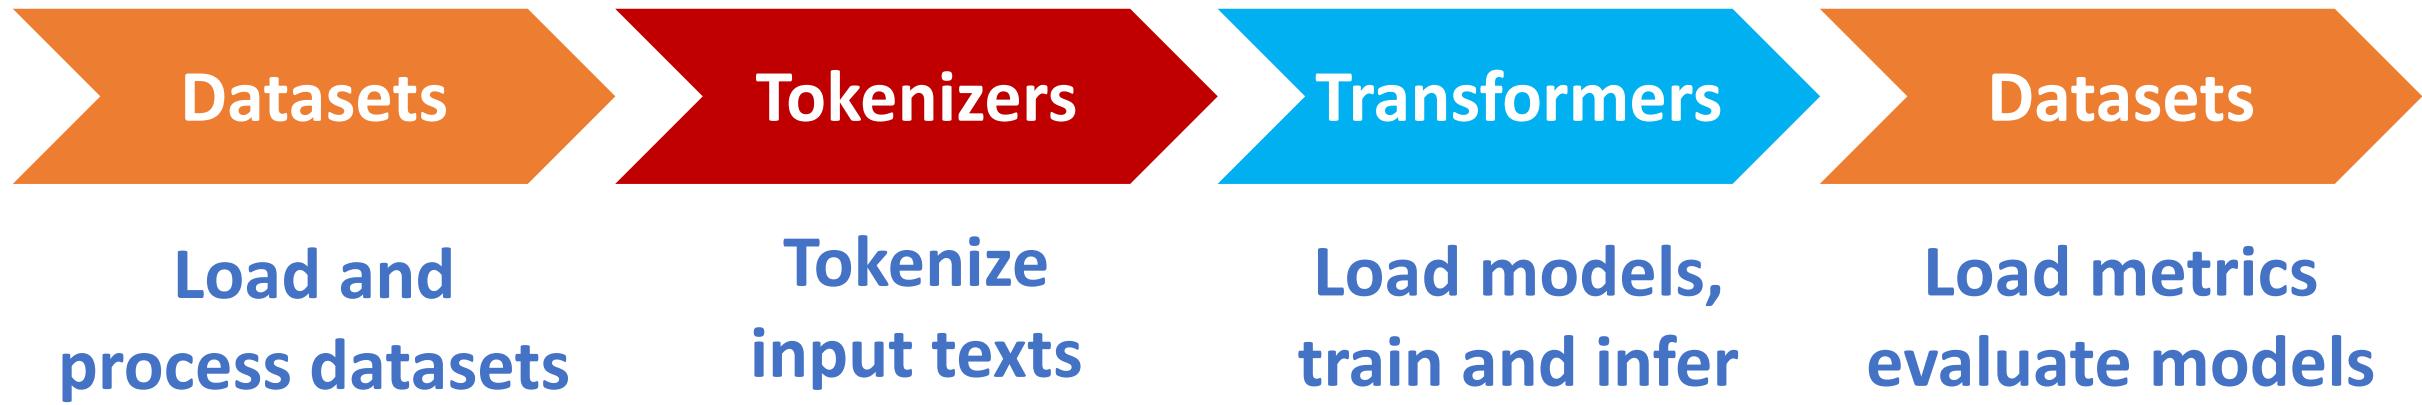
ULMFiT: 3 Steps

Transfer Learning in NLP



A typical pipeline for training transformer models

with the Datasets, Tokenizers, and Transformers libraries



Few-Shot Learning (FSL)

Typical Scenarios

- Acting as a test bed for learning like human
- Learning for rare cases
- Reducing data gathering effort and computational cost

Few-Shot Learning (FSL)

- Few-Shot Learning (FSL) is a sub-area in machine learning.
- Machine Learning Definition
 - A computer program is said to learn from experience **E** with respect to some classes of task **T** and performance measure **P** if its performance can improve with **E** on **T** measured by **P**.
 - Example: Image classification task (**T**), a machine learning program can improve its classification accuracy (**P**) through **E** obtained by training on a large number of labeled images (e.g., the ImageNet data set).

Machine Learning

task T	experience E	performance P
image classification [73]	large-scale labeled images for each class	classification accuracy
the ancient game of Go [120]	a database containing around 30 million recorded moves of human experts and self-play records	winning rate

Few-Shot Learning (FSL)

- Few-shot Learning (FSL) is a type of machine learning problems (specified by E, T, and P), where E contains only a limited number of examples with supervised information for the target T.
 - Existing FSL problems are mainly supervised learning problems.
 - Few-shot classification learns classifiers given only a few labeled examples of each class.
 - image classification
 - sentiment classification from short text
 - object recognition

Few-Shot Learning (FSL)

- Few-shot classification learns a classifier h , which predicts label y_i for each input x_i .
- Usually, one considers the **N -way- K -shot** classification, in which D_{train} contains $I = KN$ examples from N classes each with K examples

Few-Shot Learning (FSL)

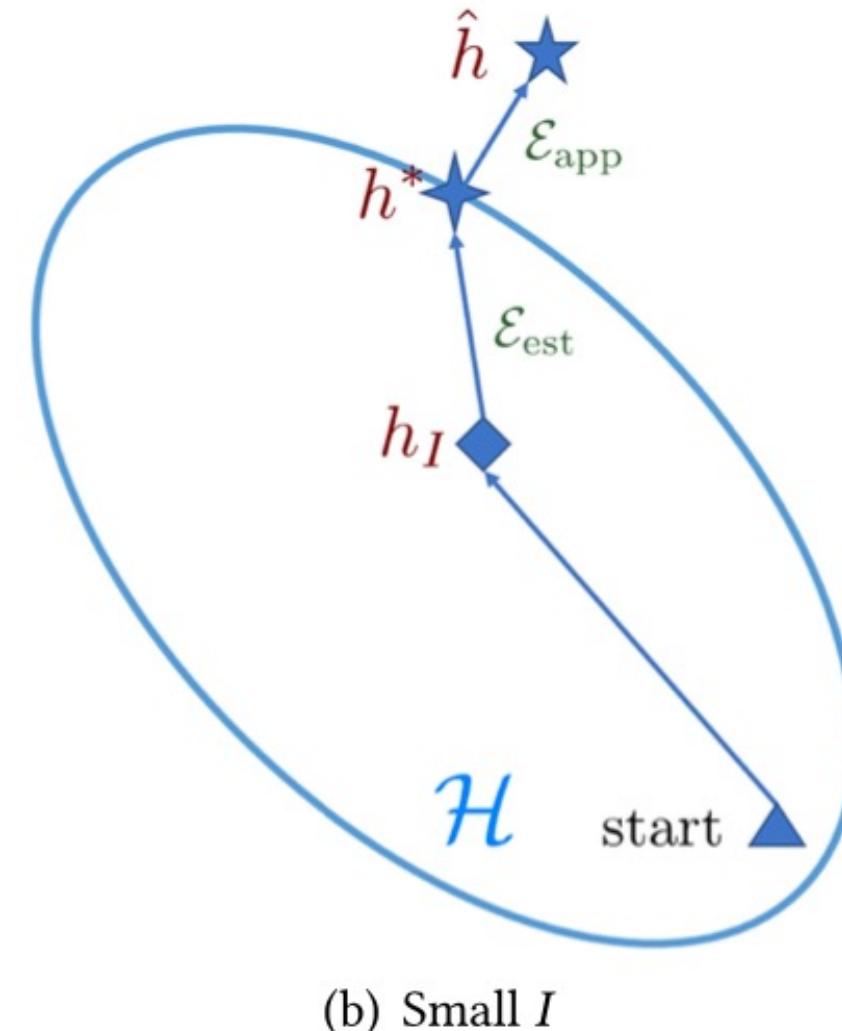
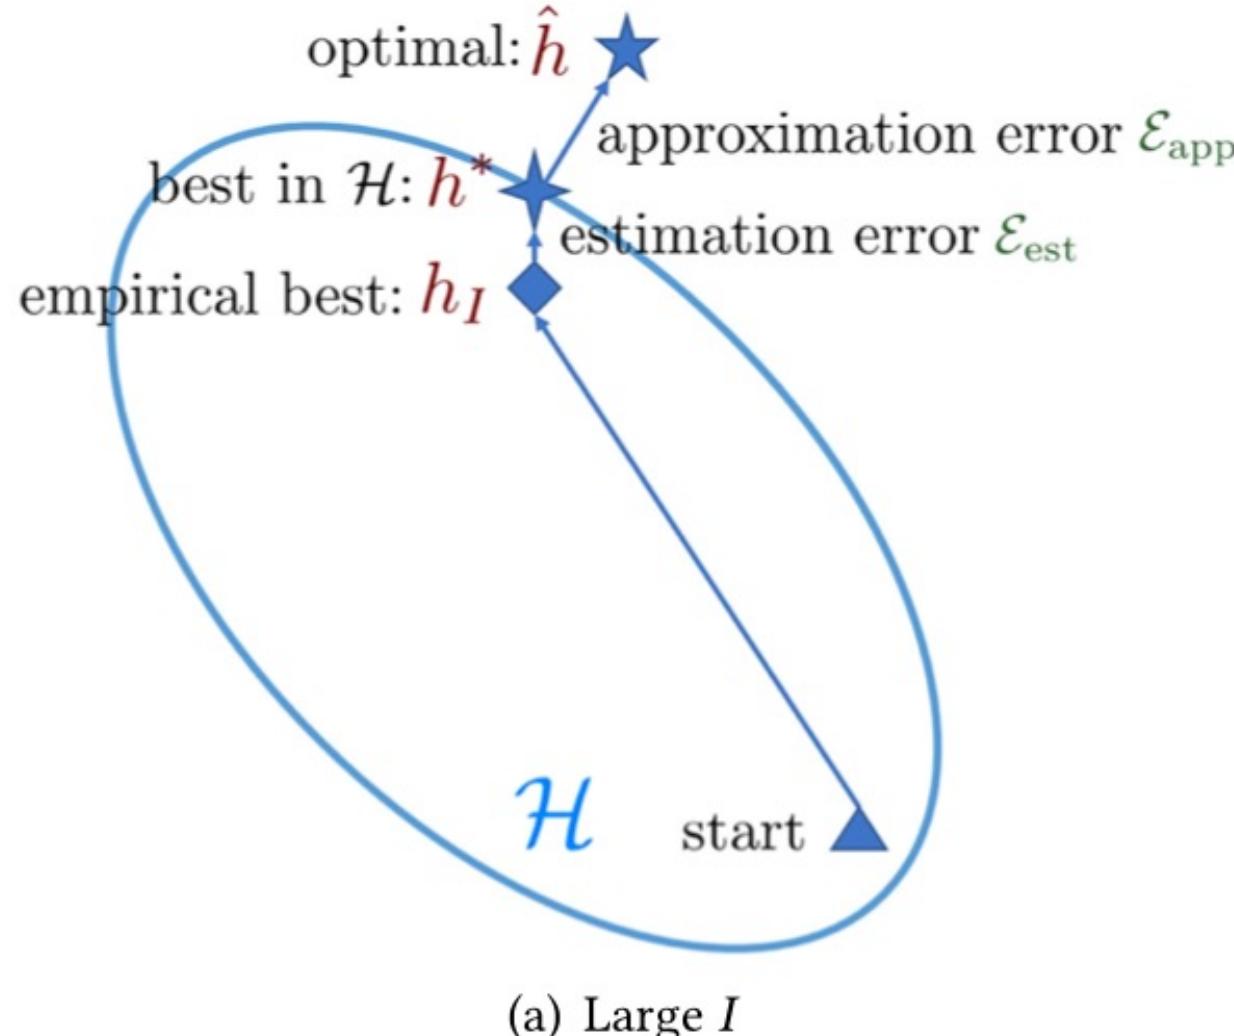
- Few-Shot Learning (FSL)
 - $K = 10 \sim 100$ examples
- One-Shot Learning (1SL)
 - $K = 1$ example
- Zero-Shot Learning (OSL)(ZSL)
 - $K = 0$

Few-Shot Learning (FSL)

task T	experience E		performance P
	supervised information	prior knowledge	
character generation [76]	a few examples of new character	pre-learned knowledge of parts and relations	pass rate of visual Turing test
drug toxicity discovery [4]	new molecule's limited assay	similar molecules' assays	classification accuracy
image classification [70]	a few labeled images for each class of the target T	raw images of other classes, or pre-trained models	classification accuracy

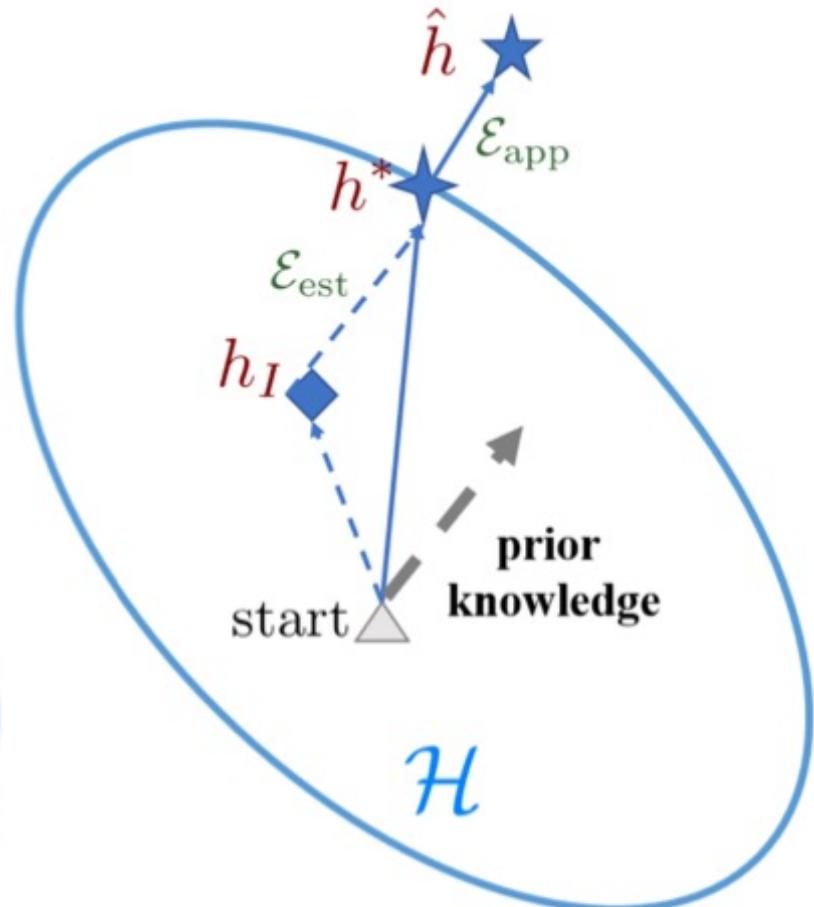
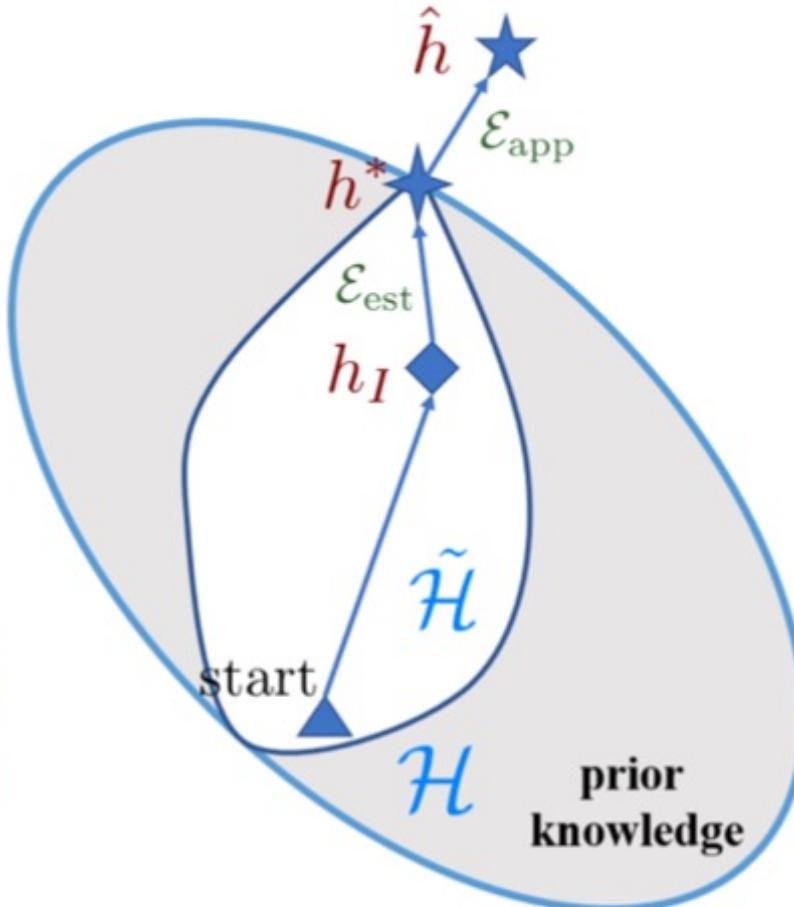
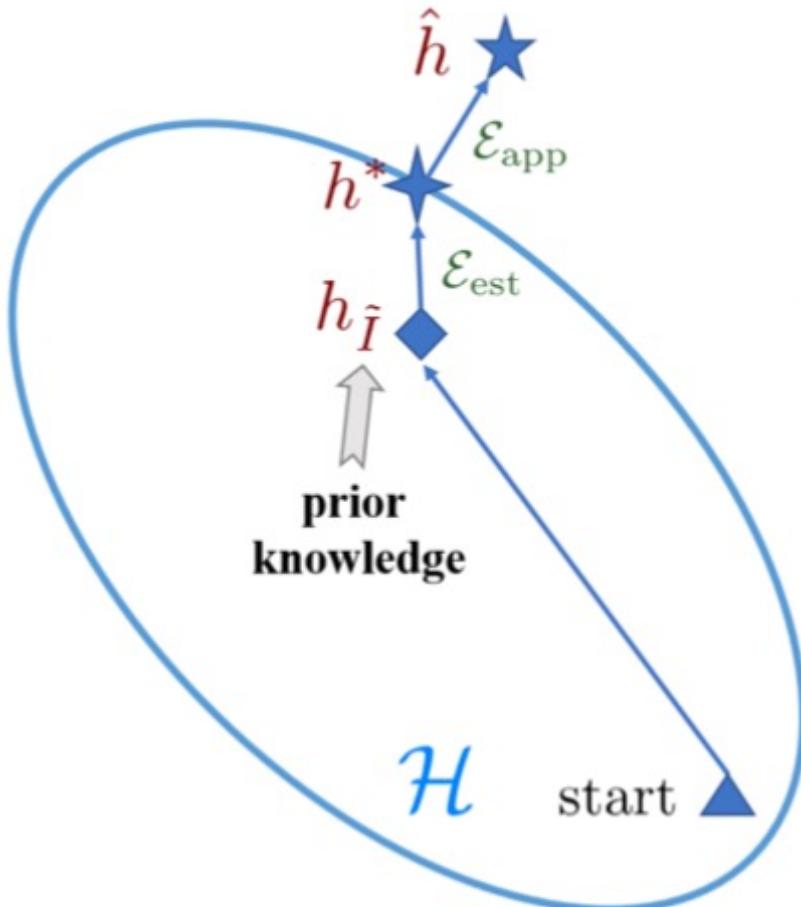
Few-Shot Learning (FSL)

Comparison of learning with sufficient and few training samples



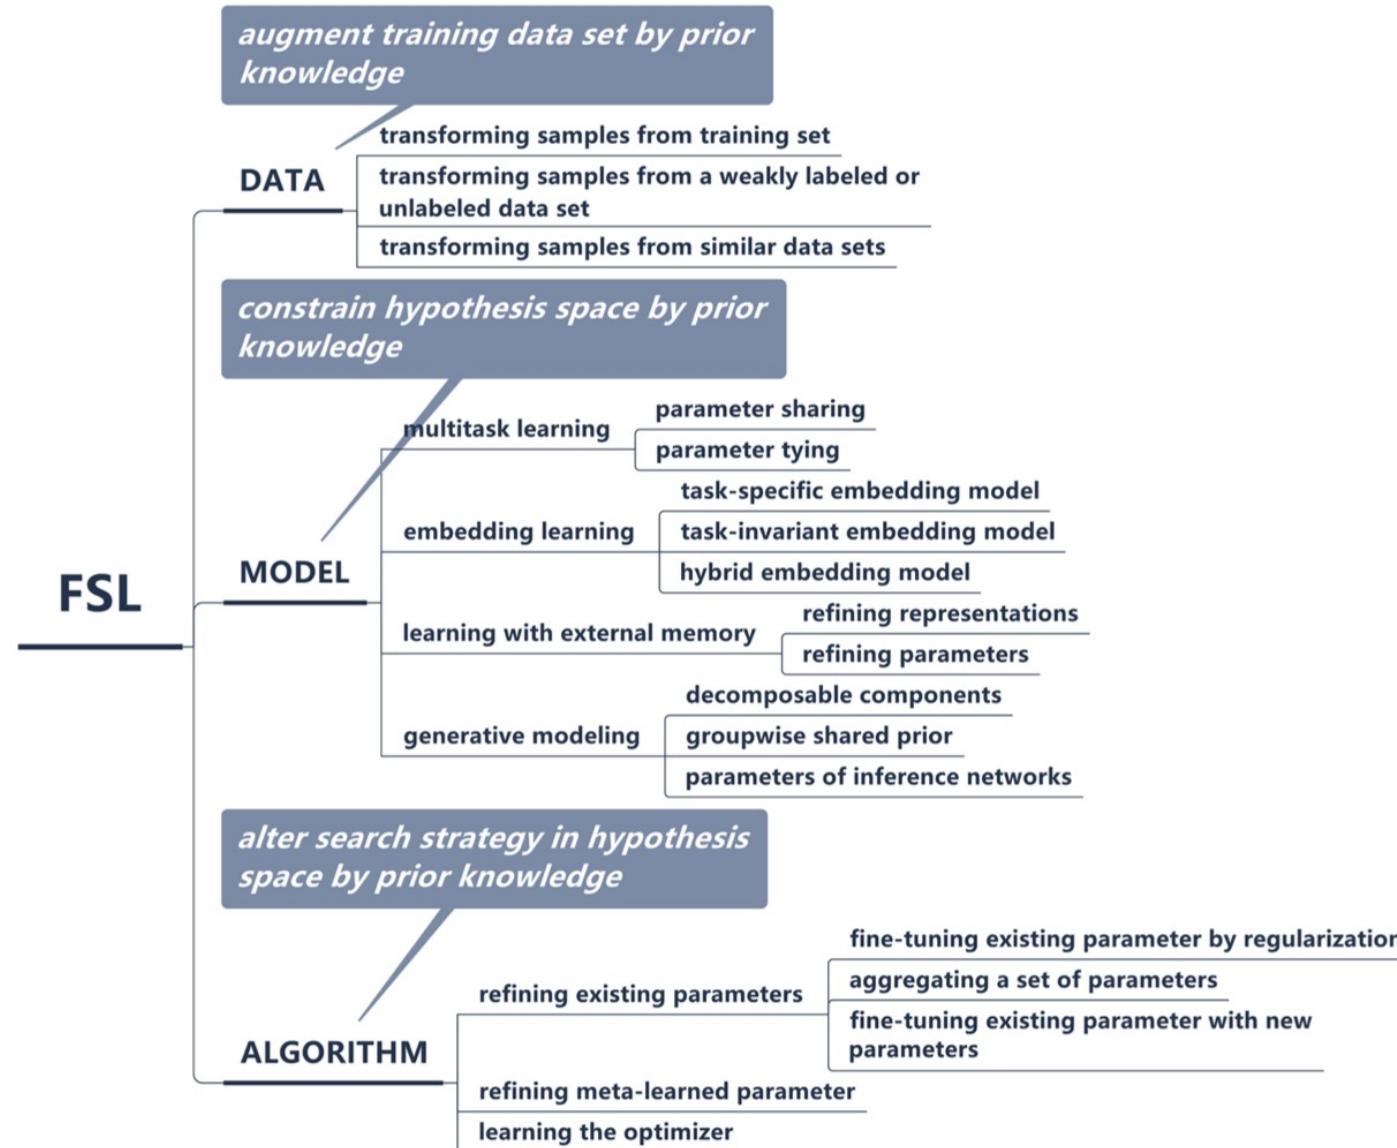
Few-Shot Learning (FSL)

Different perspectives on how FSL methods solve the few-shot problem



Few-Shot Learning (FSL)

A taxonomy of FSL methods



Few-Shot Learning (FSL)

augment training data set by prior knowledge

DATA

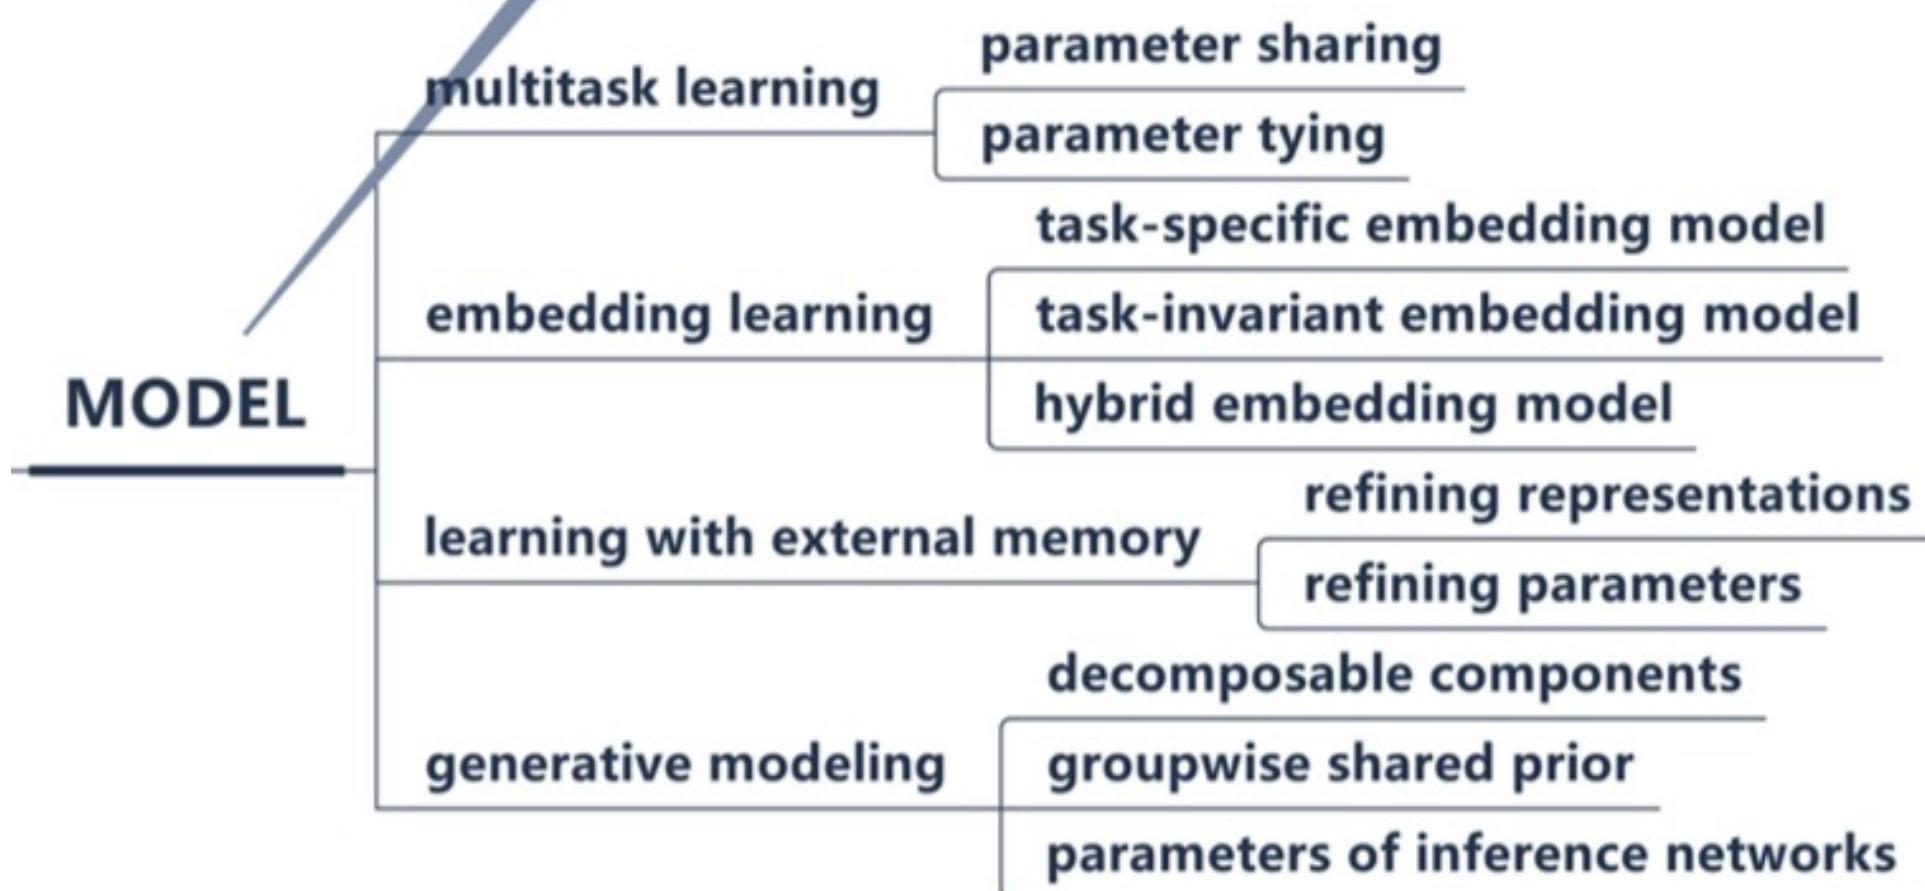
transforming samples from training set

transforming samples from a weakly labeled or unlabeled data set

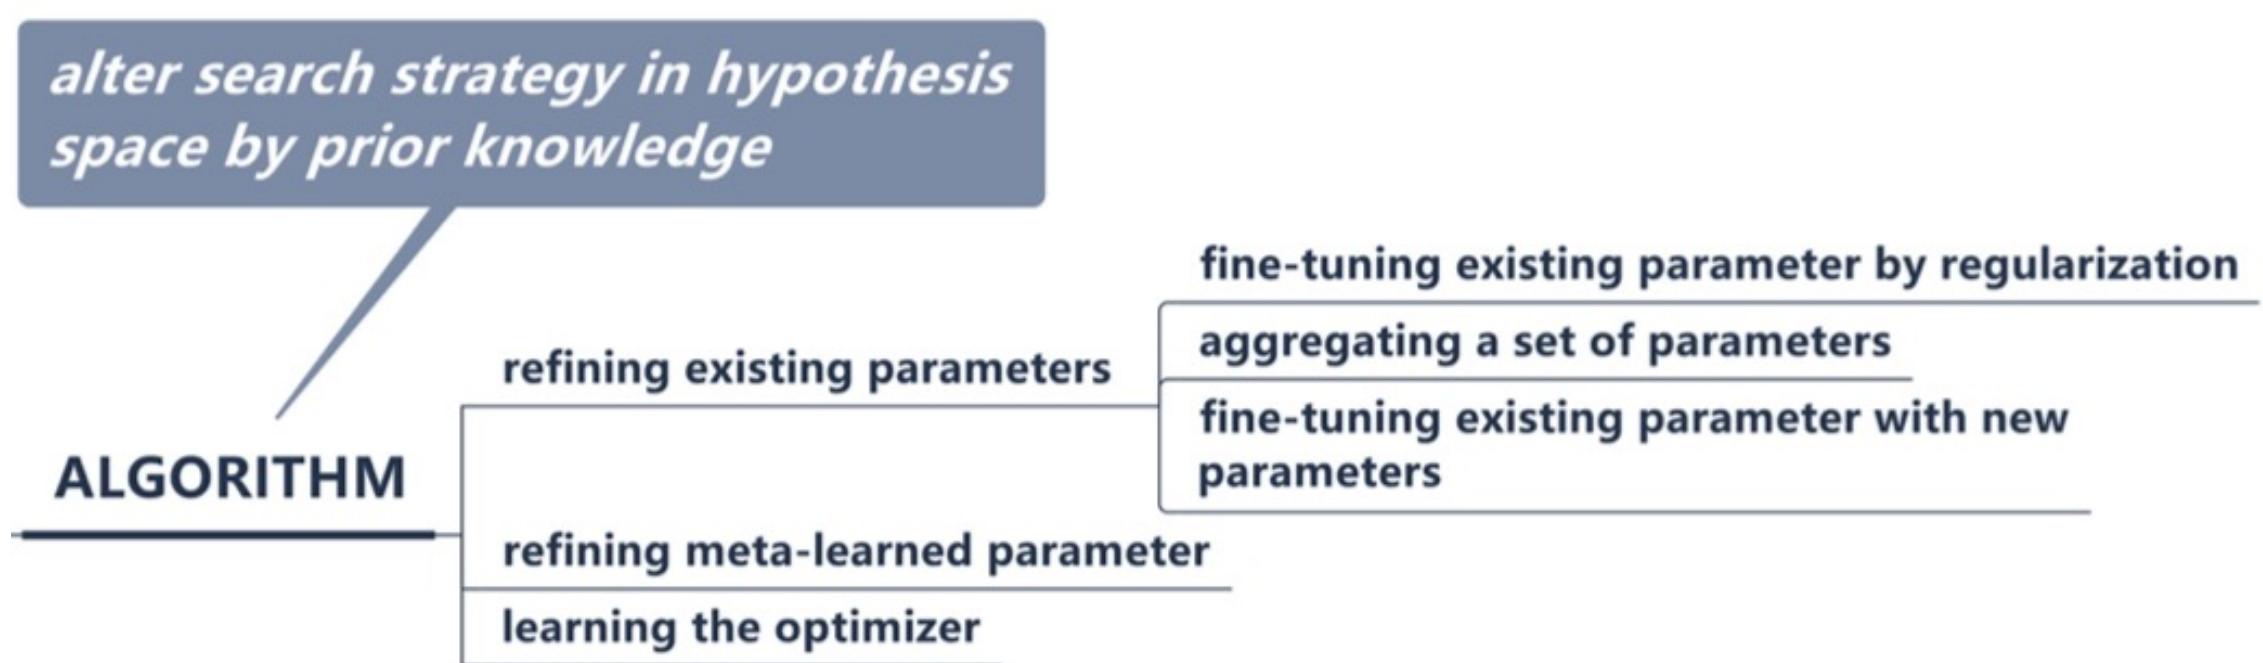
transforming samples from similar data sets

Few-Shot Learning (FSL)

constrain hypothesis space by prior knowledge

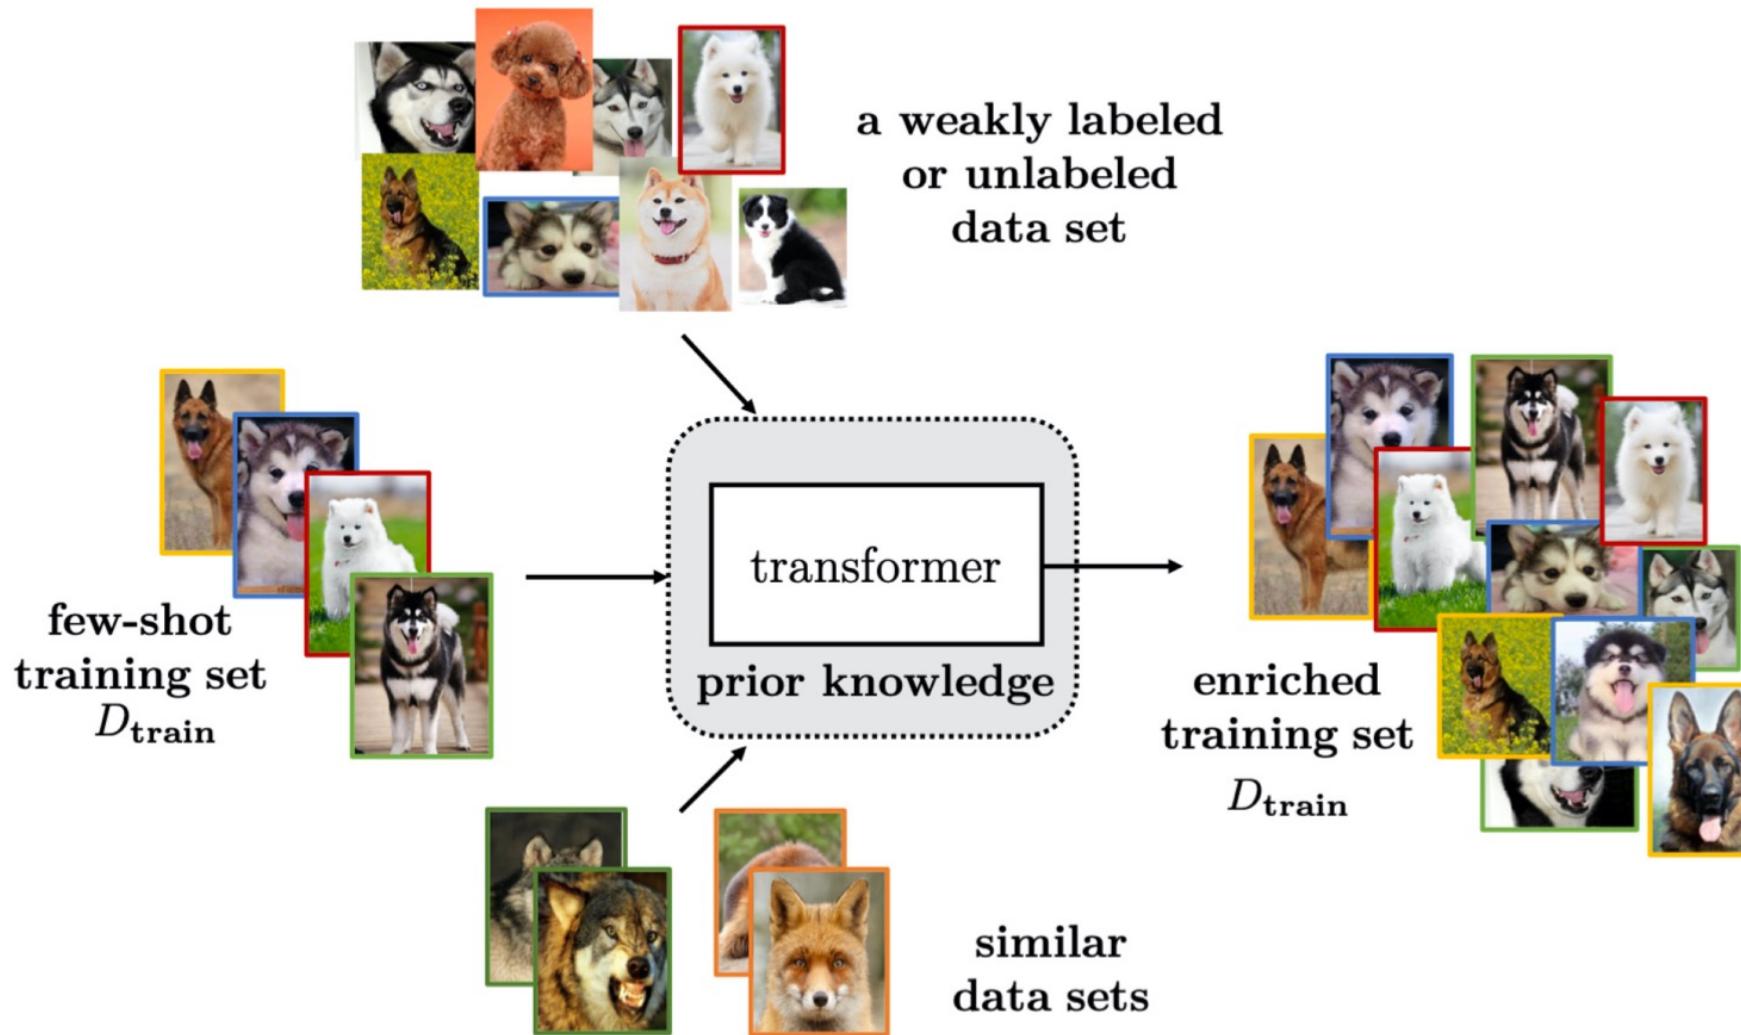


Few-Shot Learning (FSL)



Few-Shot Learning (FSL)

Solving the FSL problem by data augmentation



Few-Shot Learning (FSL)

Characteristics for FSL Methods Focusing on the Data Perspective

category	input (x, y)	transformer t	output (\tilde{x}, \tilde{y})
transforming samples from D_{train}	original (x_i, y_i)	learned transformation function on x_i	$(t(x_i), y_i)$
transforming samples from a weakly labeled or unlabeled data set	weakly labeled or unlabeled $(\bar{x}, -)$	a predictor trained from D_{train}	$(\bar{x}, t(\bar{x}))$
transforming samples from similar data sets	samples $\{(\hat{x}_j, \hat{y}_j)\}$ from similar data sets	an aggregator to combine $\{(\hat{x}_j, \hat{y}_j)\}$	$(t(\{\hat{x}_j\}), t(\{\hat{y}_j\}))$

The transformer $t(\cdot)$ takes input (x, y) and returns synthesized sample (\tilde{x}, \tilde{y}) to augment the few-shot D_{train} .

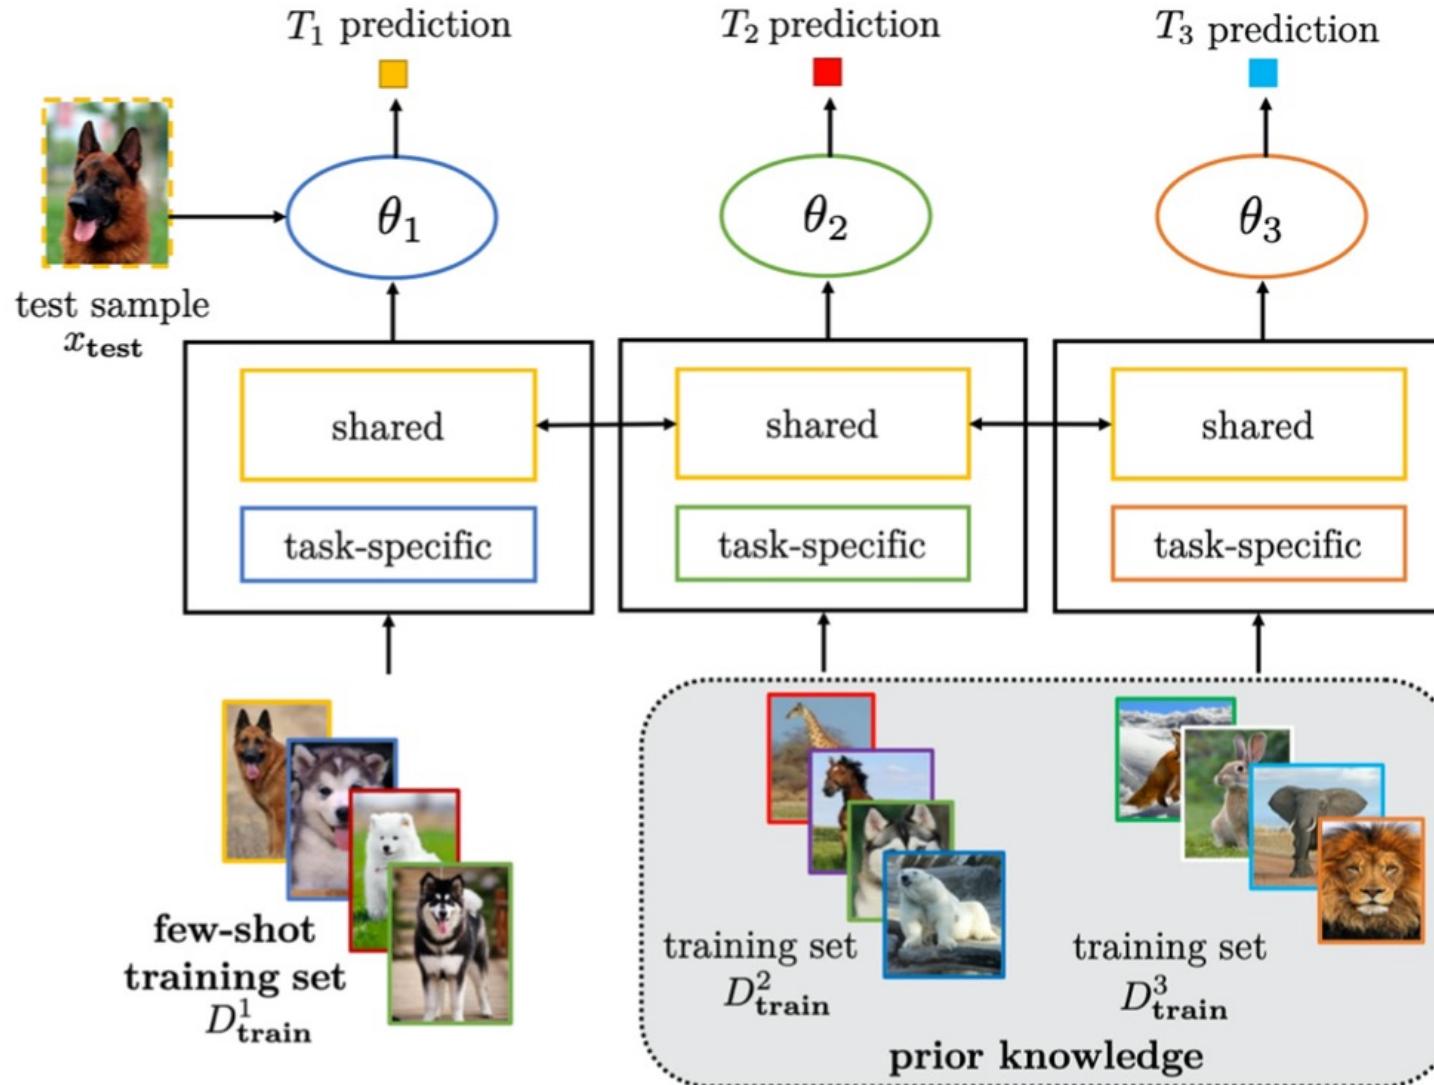
Few-Shot Learning (FSL)

Characteristics for FSL Methods Focusing on the Model Perspective

strategy	prior knowledge	how to constrain \mathcal{H}
multitask learning	other T 's with their data sets D 's	share/tie parameter
embedding learning	embedding learned from/together with other T 's	project samples to a smaller embedding space in which similar and dissimilar samples can be easily discriminated
learning with external memory	embedding learned from other T 's to interact with memory	refine samples using key-value pairs stored in memory
generative modeling	prior model learned from other T 's	restrict the form of distribution

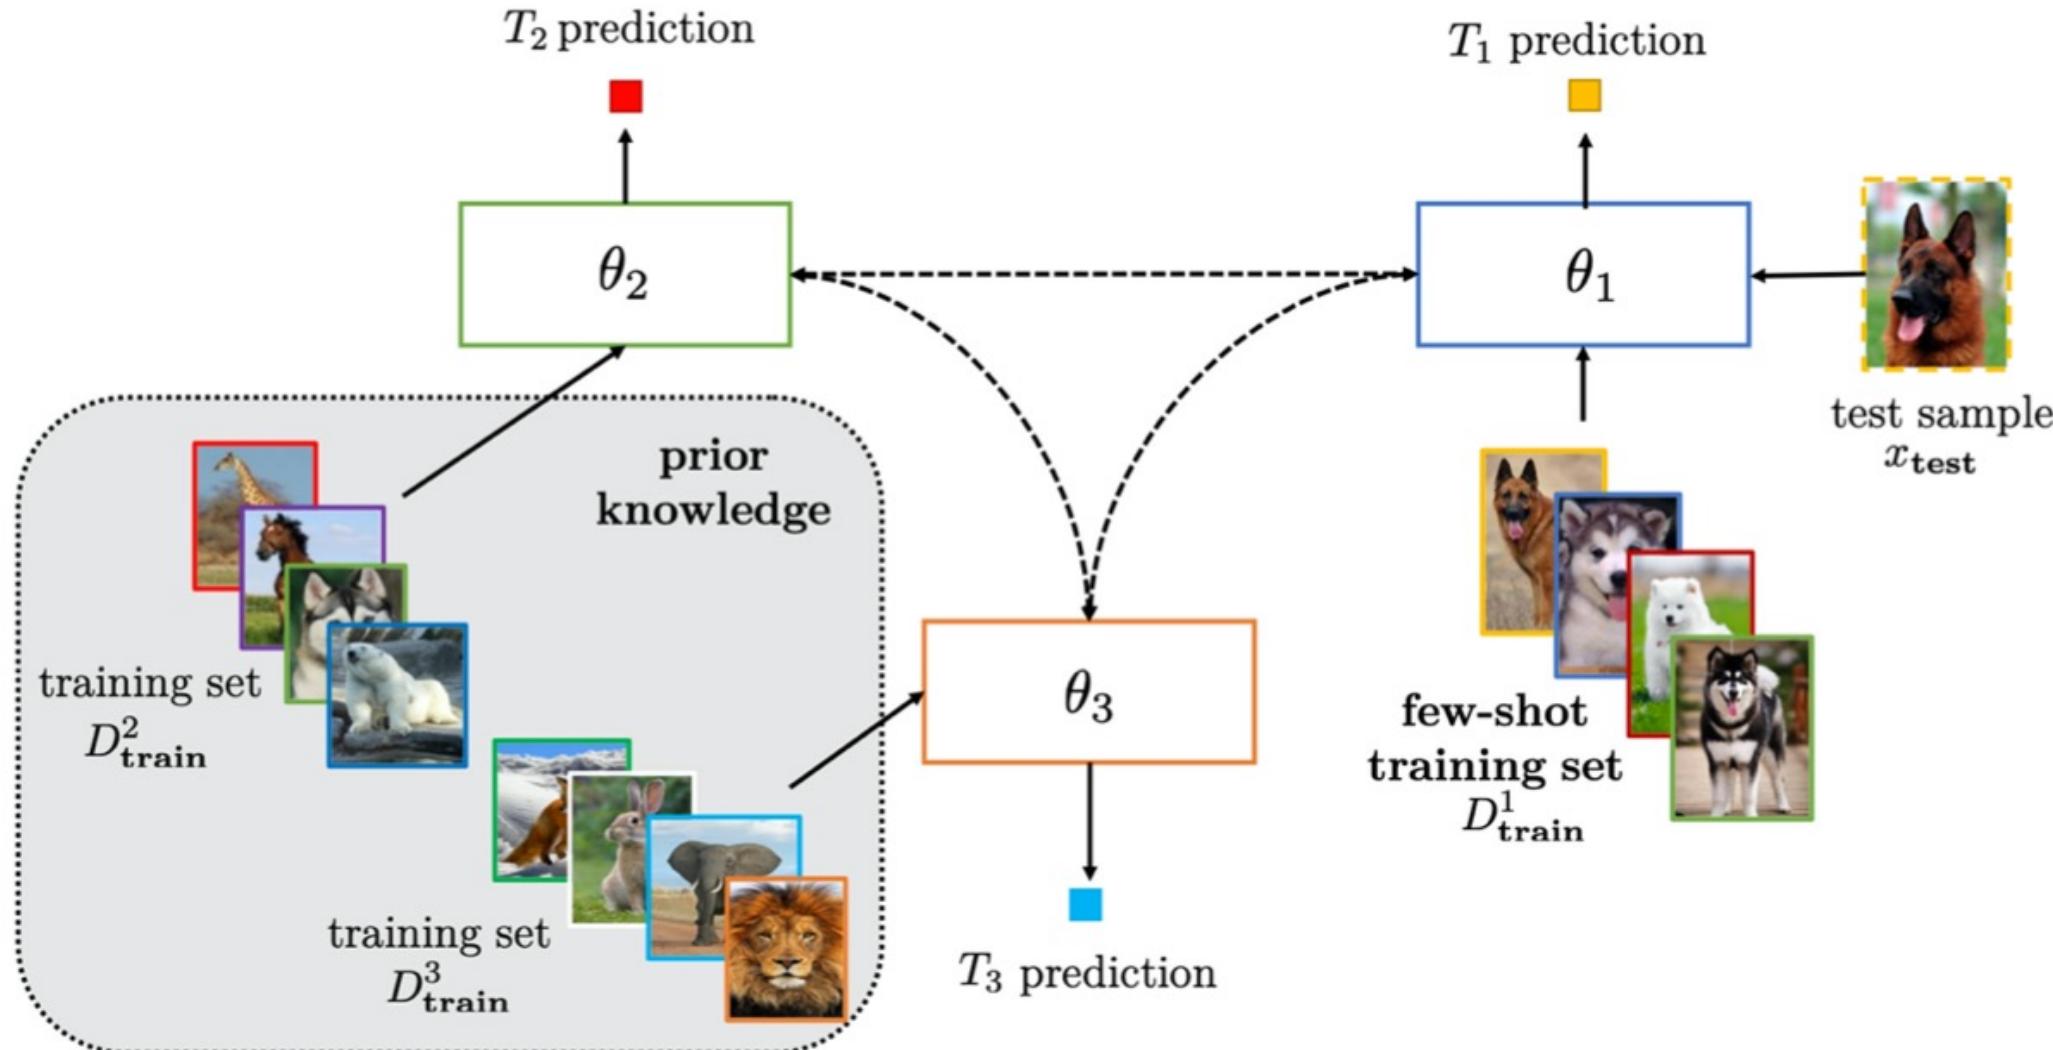
Few-Shot Learning (FSL)

Solving the FSL problem by multitask learning with parameter sharing



Few-Shot Learning (FSL)

Solving the FSL problem by multitask learning with parameter tying



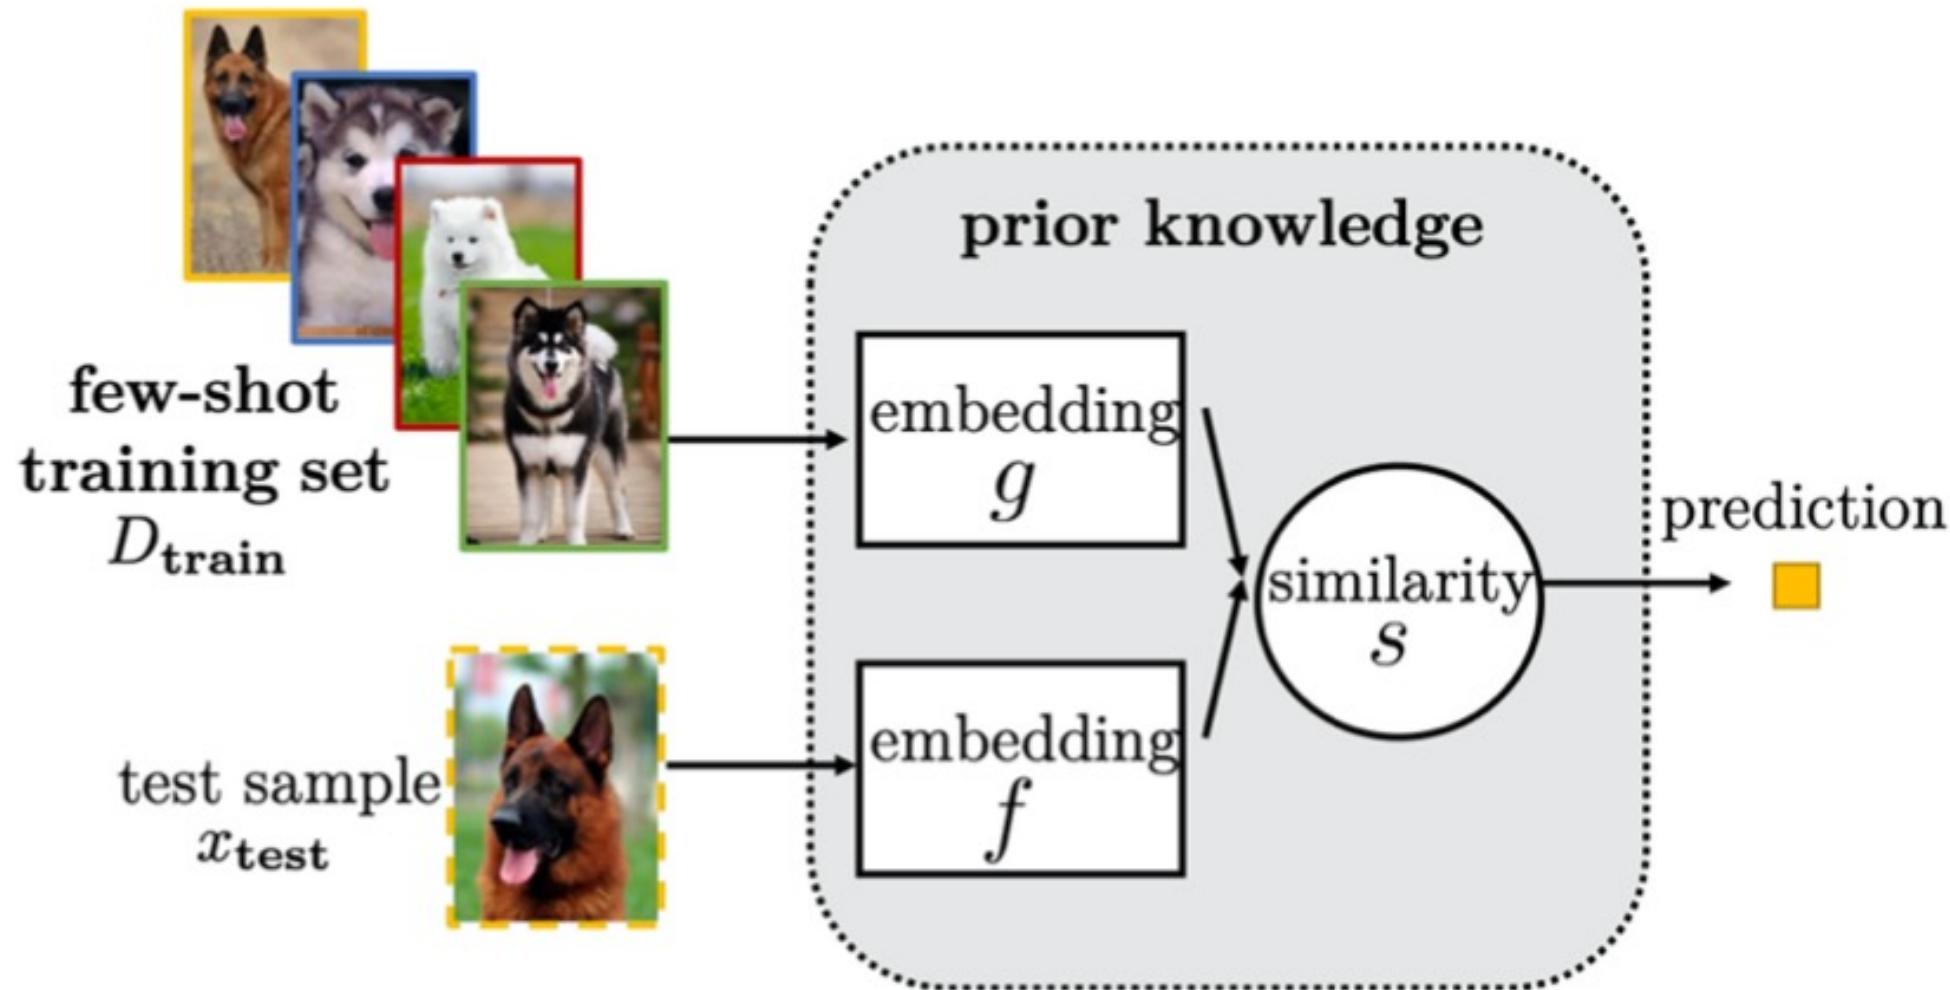
Few-Shot Learning (FSL)

Characteristics of Embedding Learning Methods

category	method	embedding function f for x_{test}	embedding function g for D_{train}	similarity measure s
task-specific	mAP-DLM/SSVM [130]	CNN	the same as f	cosine similarity
	class relevance pseudo-metric [36]	kernel	the same as f	squared ℓ_2 distance
	convolutional siamese net [70]	CNN	the same as f	weighted ℓ_1 distance
	Micro-Set [127]	logistic projection	the same as f	ℓ_2 distance
	Matching Nets [138]	CNN, LSTM	CNN, biLSTM	cosine similarity
	resLSTM [4]	GNN, LSTM	GNN, LSTM	cosine similarity
	Active MN [8]	CNN	biLSTM	cosine similarity
	SSMN [24]	CNN	another CNN	learned distance
	ProtoNet [121]	CNN	the same as f	squared ℓ_2 distance
	semi-supervised ProtoNet [108]	CNN	the same as f	squared ℓ_2 distance
task-invariant	PMN [141]	CNN, LSTM	CNN, biLSTM	cosine similarity
	ARC [119]	LSTM, biLSTM	the same as f	-
	Relation Net [126]	CNN	the same as f	-
	GNN [115]	CNN, GNN	the same as f	learned distance
	TPN [84]	CNN	the same as f	Gaussian similarity
	SNAIL [91]	CNN	the same as f	-
	Learnet [14]	adaptive CNN	CNN	weighted ℓ_1 distance
	DCCN [162]	adaptive CNN	CNN	-
	R2-D2 [13]	adaptive CNN	CNN	-
	TADAM [100]	adaptive CNN	the same as f	squared ℓ_2 distance
hybrid				

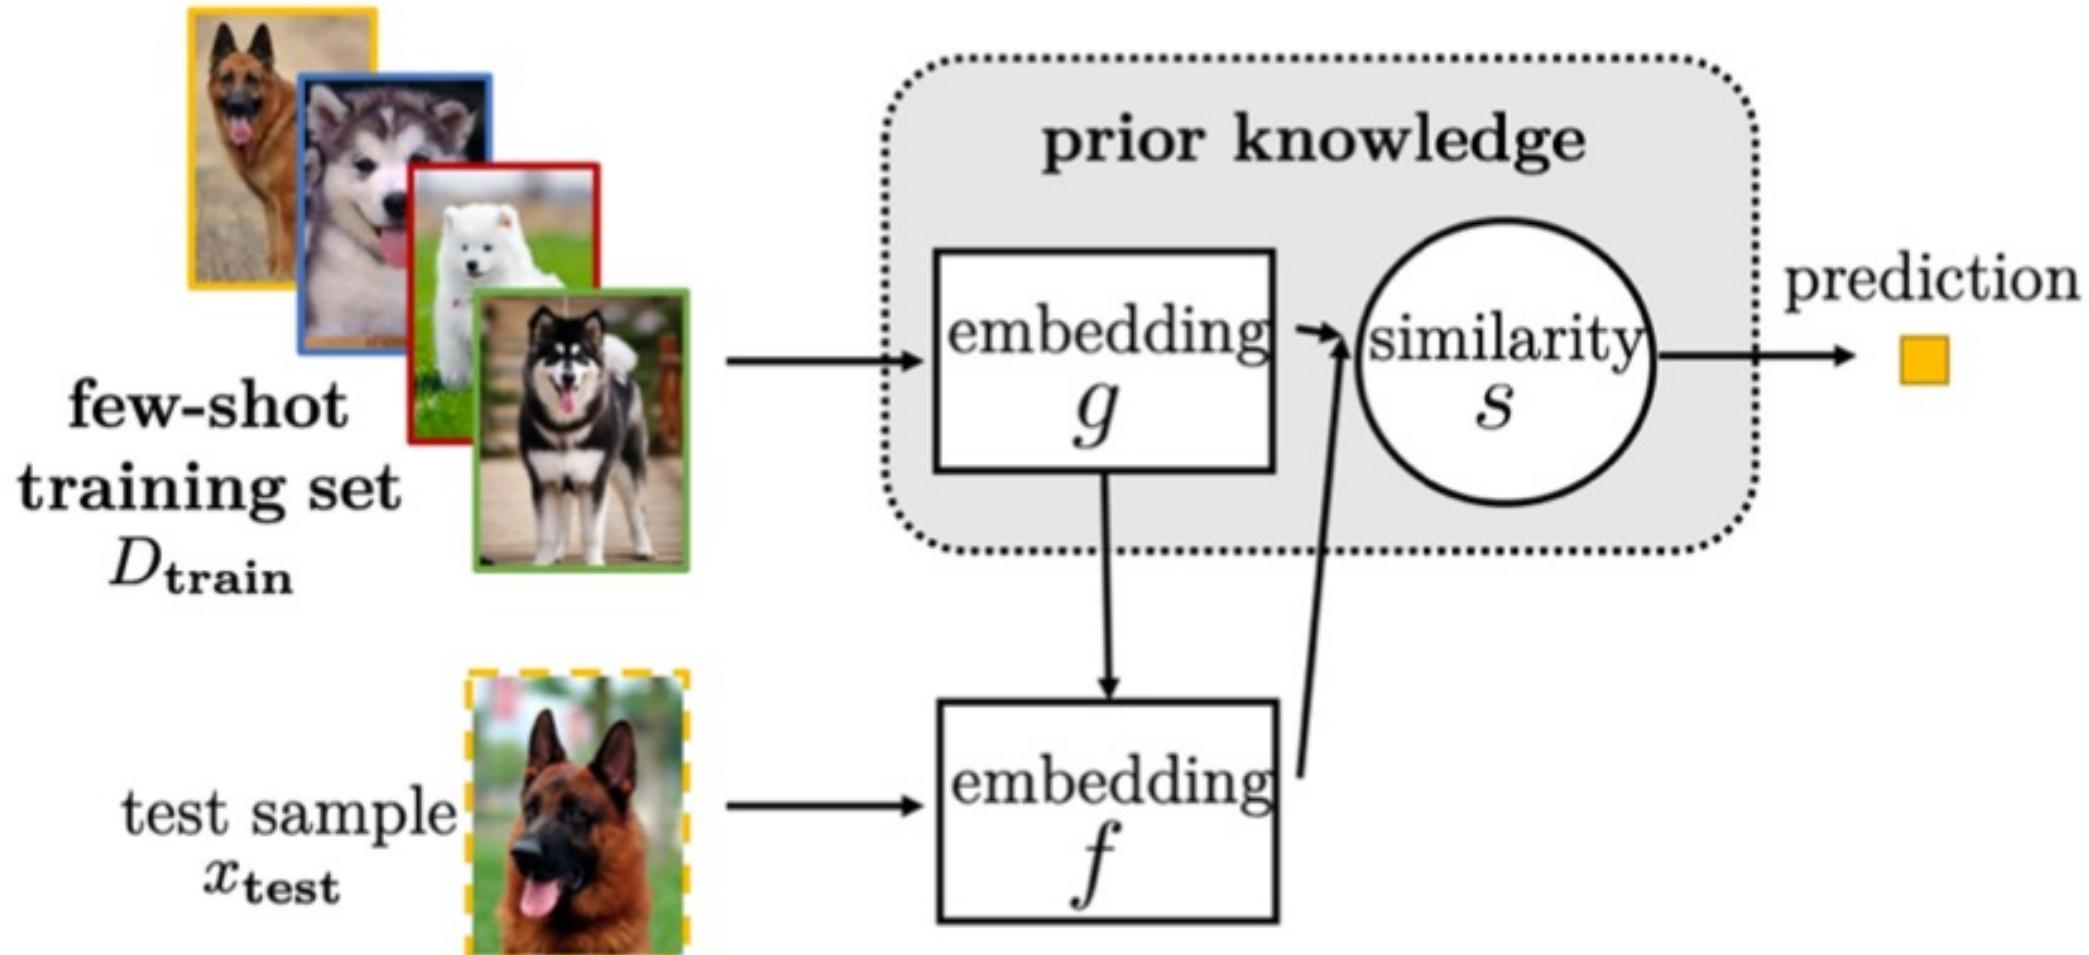
Few-Shot Learning (FSL)

Solving the FSL problem by task-invariant embedding model



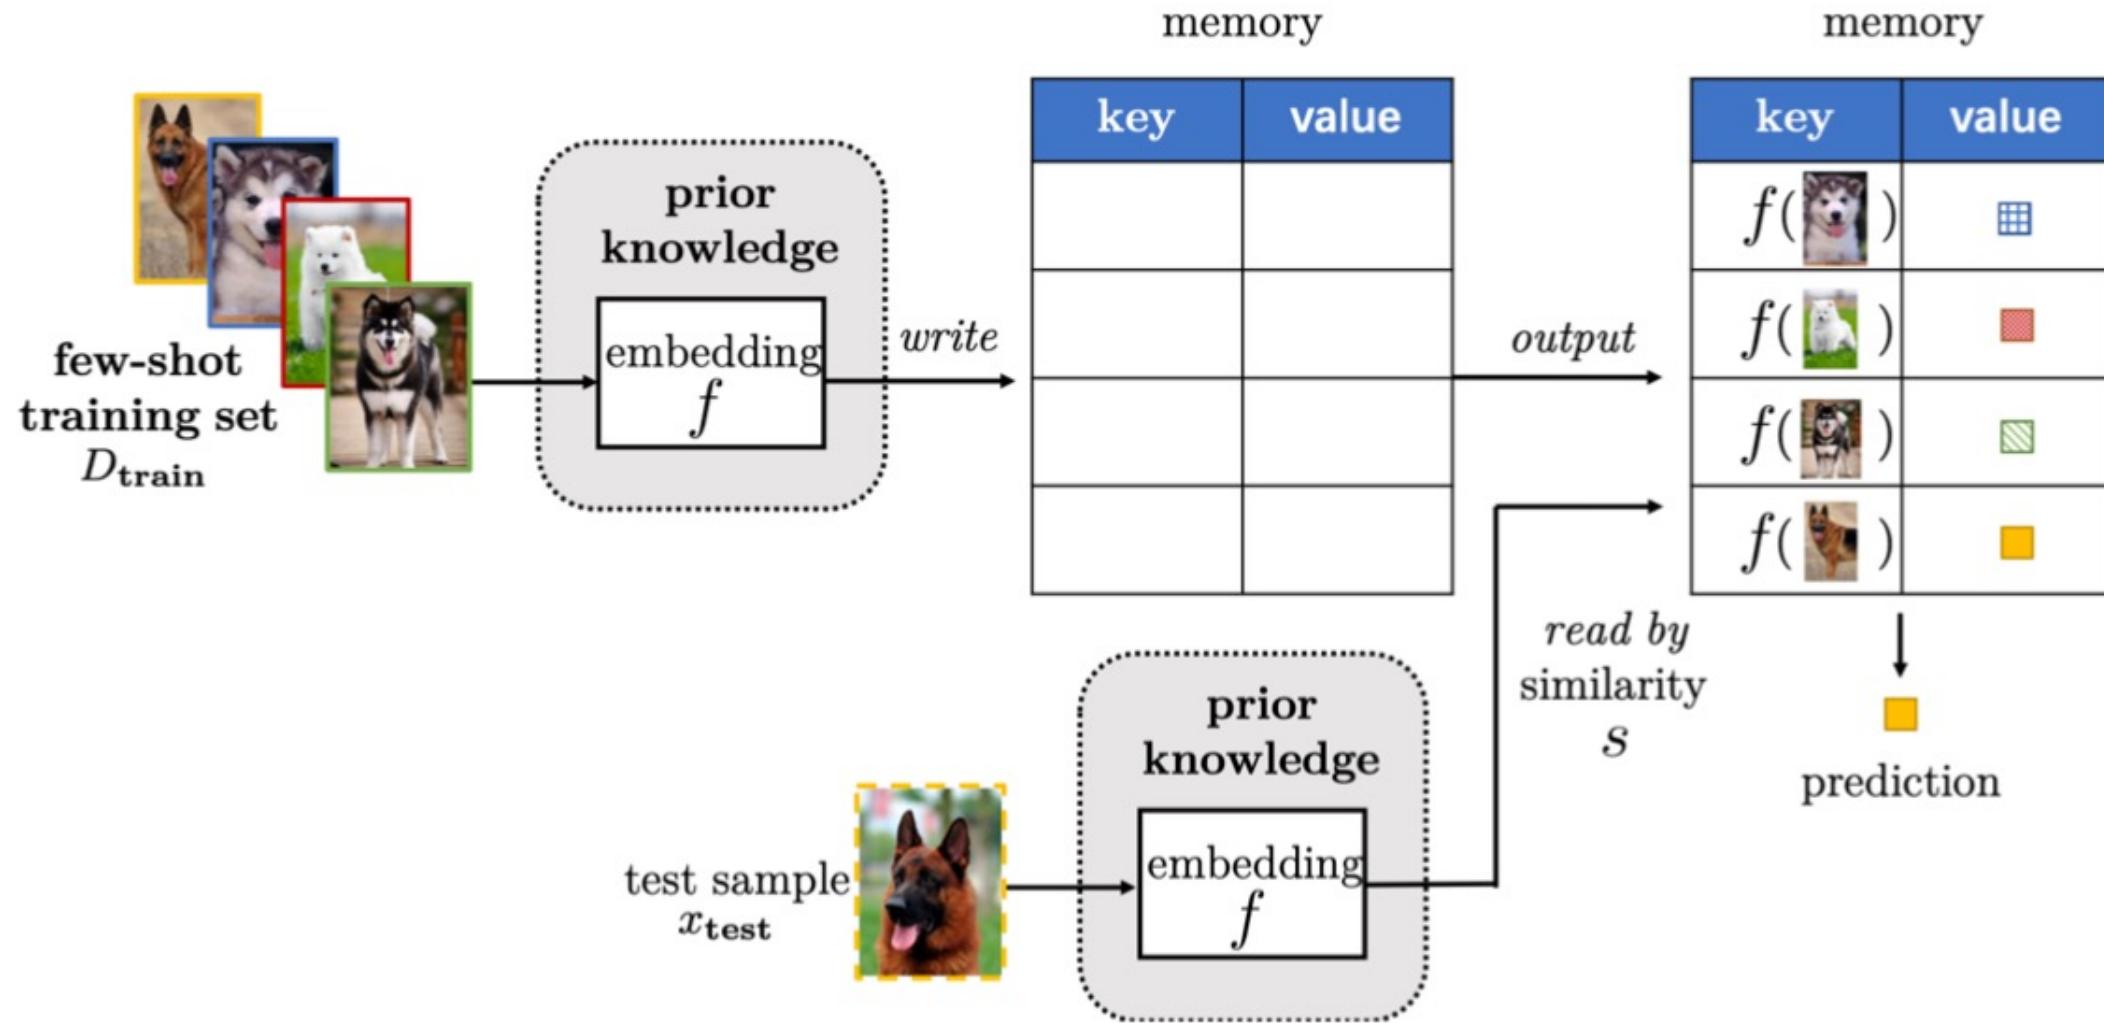
Few-Shot Learning (FSL)

Solving the FSL problem by hybrid embedding model



Few-Shot Learning (FSL)

Solving the FSL problem by learning with external memory



Few-Shot Learning (FSL)

Characteristics of FSL Methods

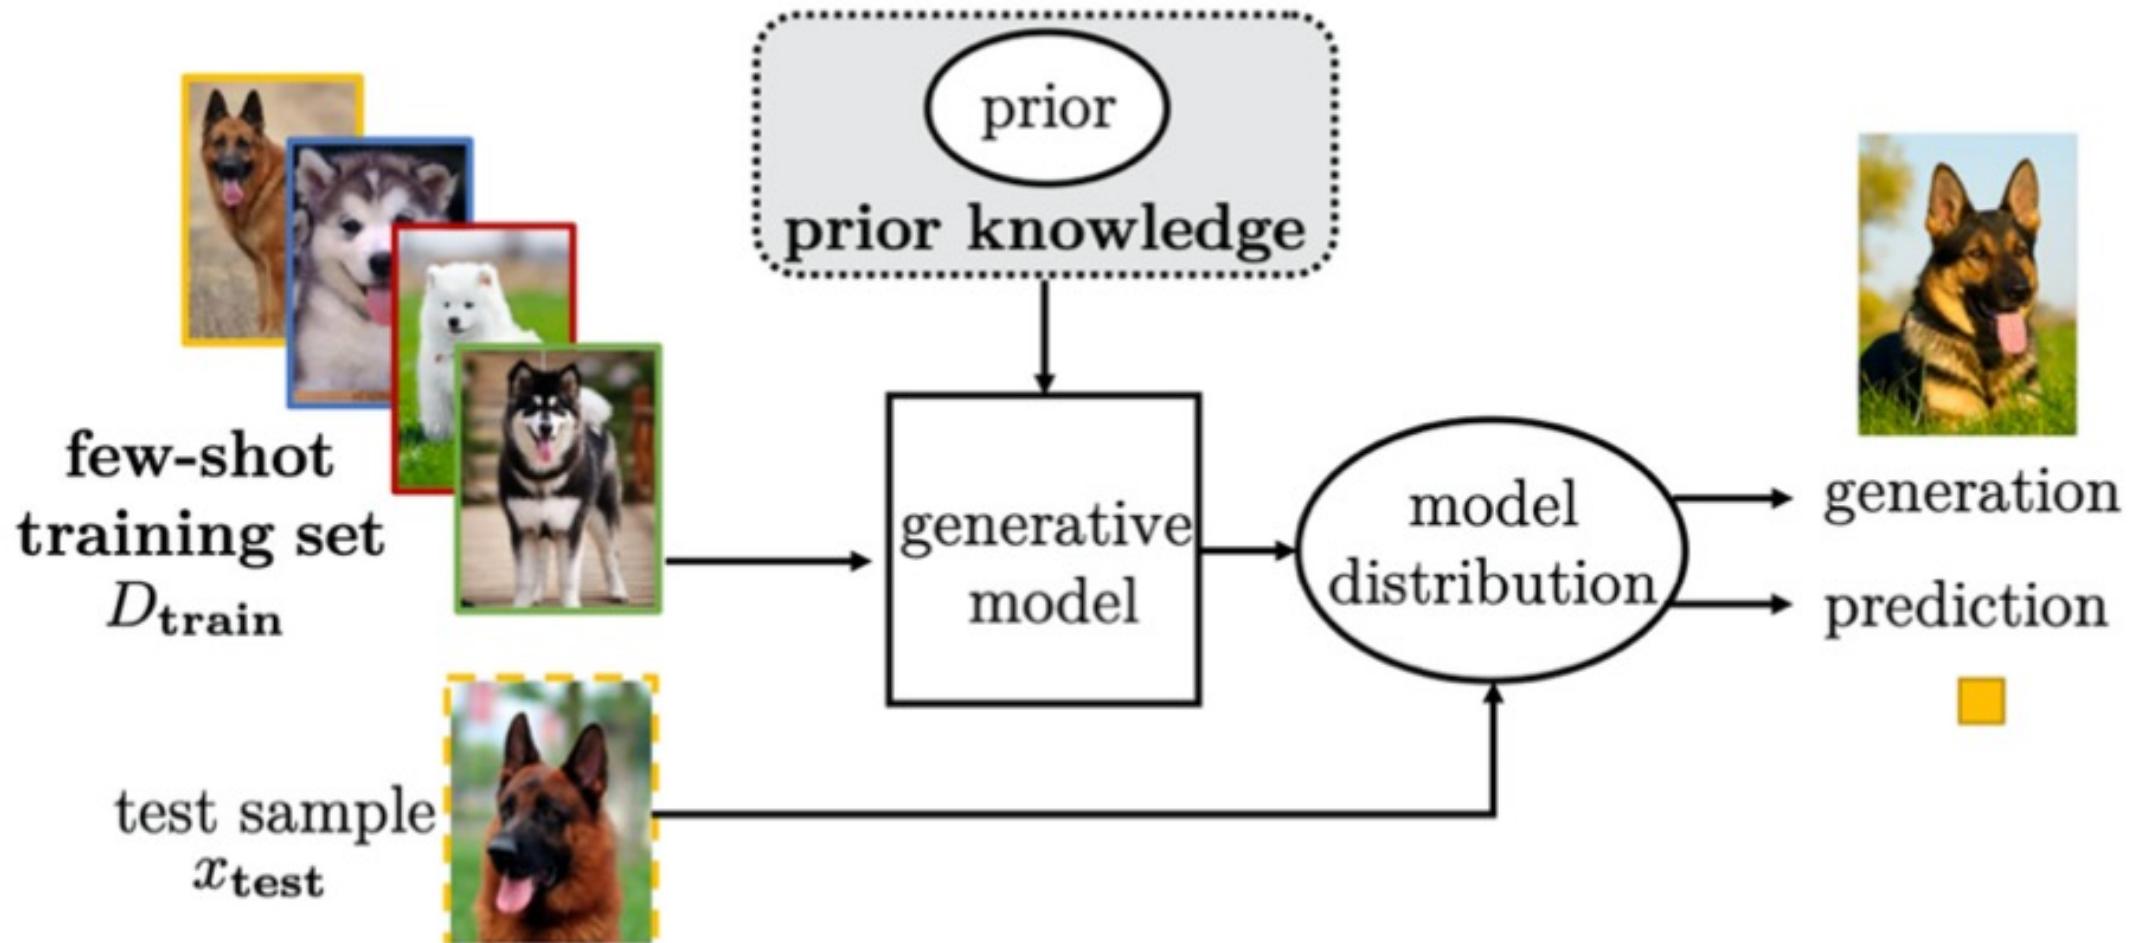
Based on Learning with External Memory

category	method	memory M		similarity s
		key M_{key}	value M_{value}	
refining representations	MANN [114]	$f(x_i, y_{i-1})$	$f(x_i, y_{i-1})$	cosine similarity
	APL [104]	$f(x_i)$	y_i	squared ℓ_2 distance
	abstraction memory [149]	$f(x_i)$	word embedding of y_i	dot product
	CMN [164]	$f(x_i)$	y_i , age	dot product
	life-long memory [65]	$f(x_i)$	y_i , age	cosine similarity
	Mem2Vec [125]	$f(x_i)$	word embedding of y_i , age	dot product
refining parameters	MetaNet [96]	$f(x_i)$	fast weight	cosine similarity
	CSNs [97]	$f(x_i)$	fast weight	cosine similarity
	MN-Net [22]	$f(x_i)$	y_i	dot product

Here, f is an embedding function usually pre-trained by CNN or LSTM.

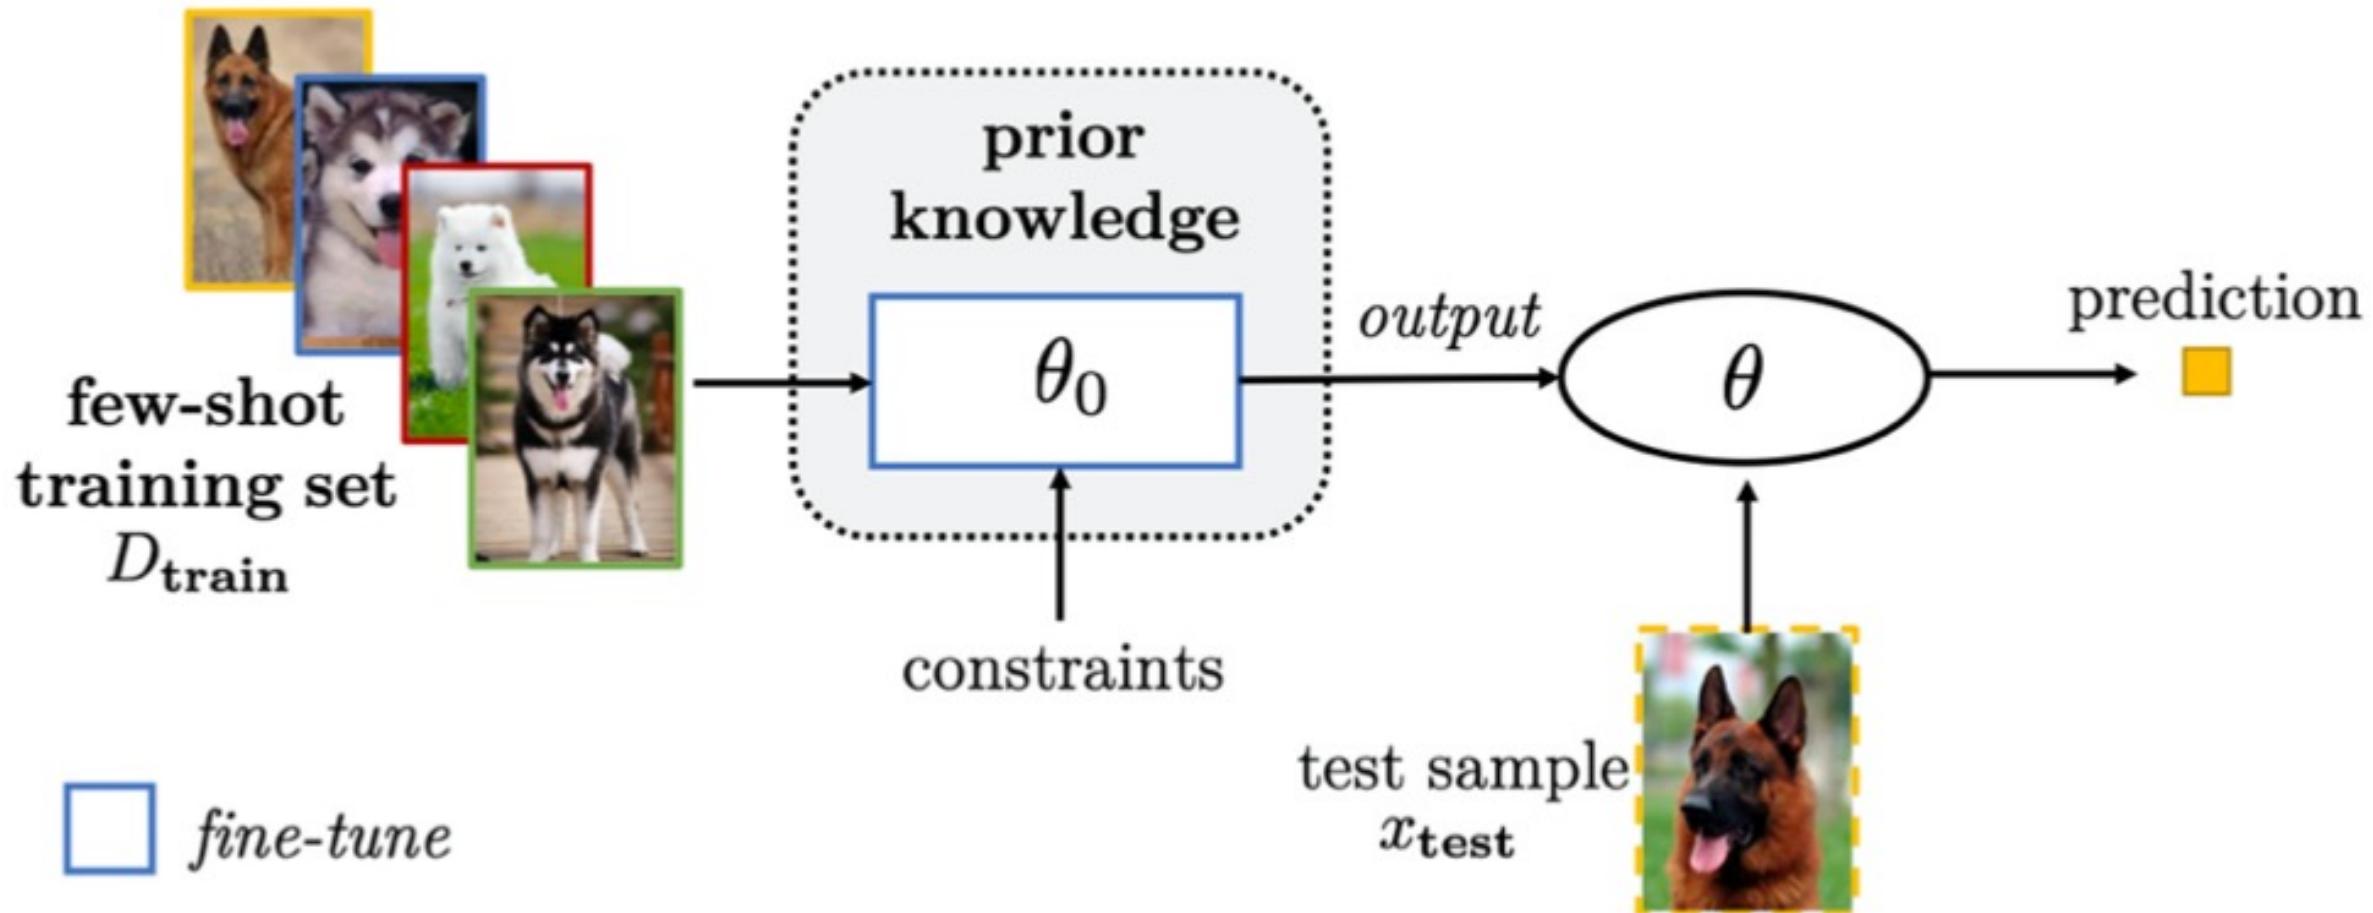
Few-Shot Learning (FSL)

Solving the FSL problem by generative modeling



Few-Shot Learning (FSL)

Solving the FSL problem by fine-tuning existing parameter θ_0 by regularization



Few-Shot Learning (FSL)

Characteristics for FSL Methods

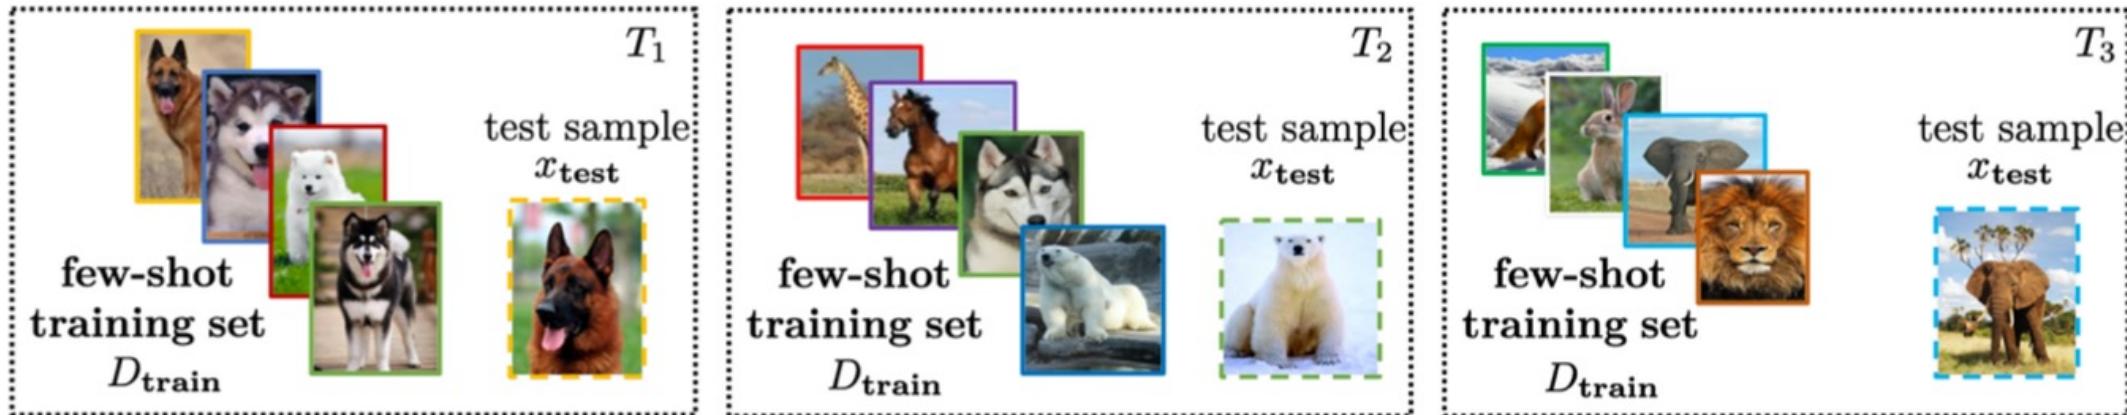
Focusing on the Algorithm Perspective

strategy	prior knowledge	how to search θ of the h^* in \mathcal{H}
refining existing parameters	learned θ_0	refine θ_0 by D_{train}
refining meta-learned parameters	meta-learner	refine θ_0 by D_{train}
learning the optimizer	meta-learner	use search steps provided by the meta-learner

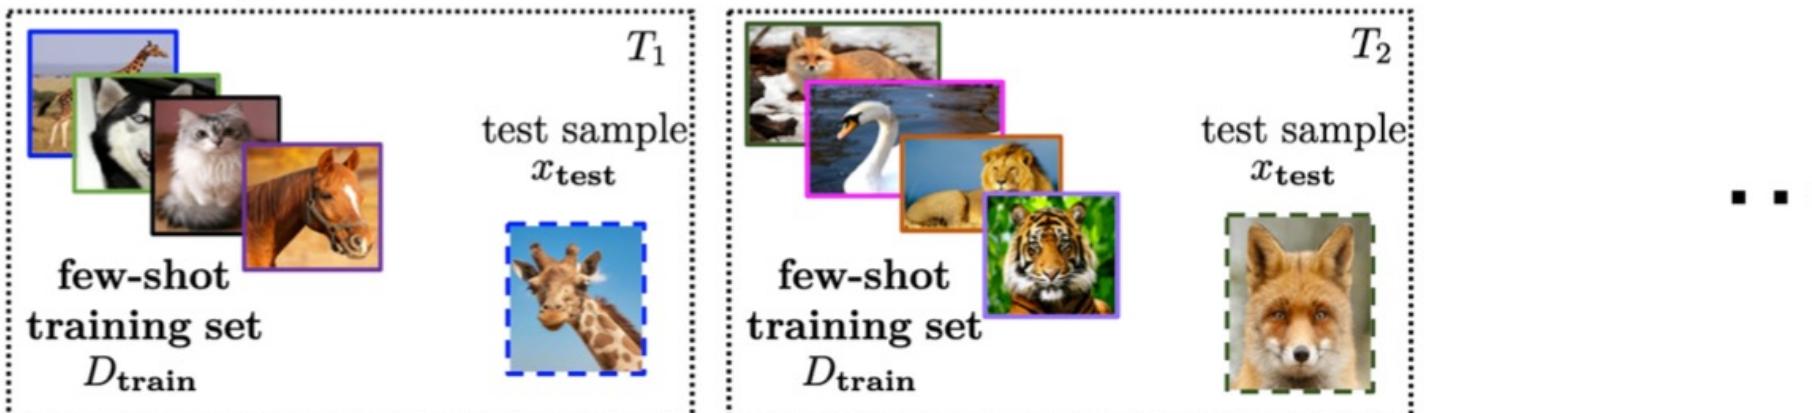
Few-Shot Learning (FSL)

Solving the FSL problem by meta-learning

meta-training tasks T_s 's



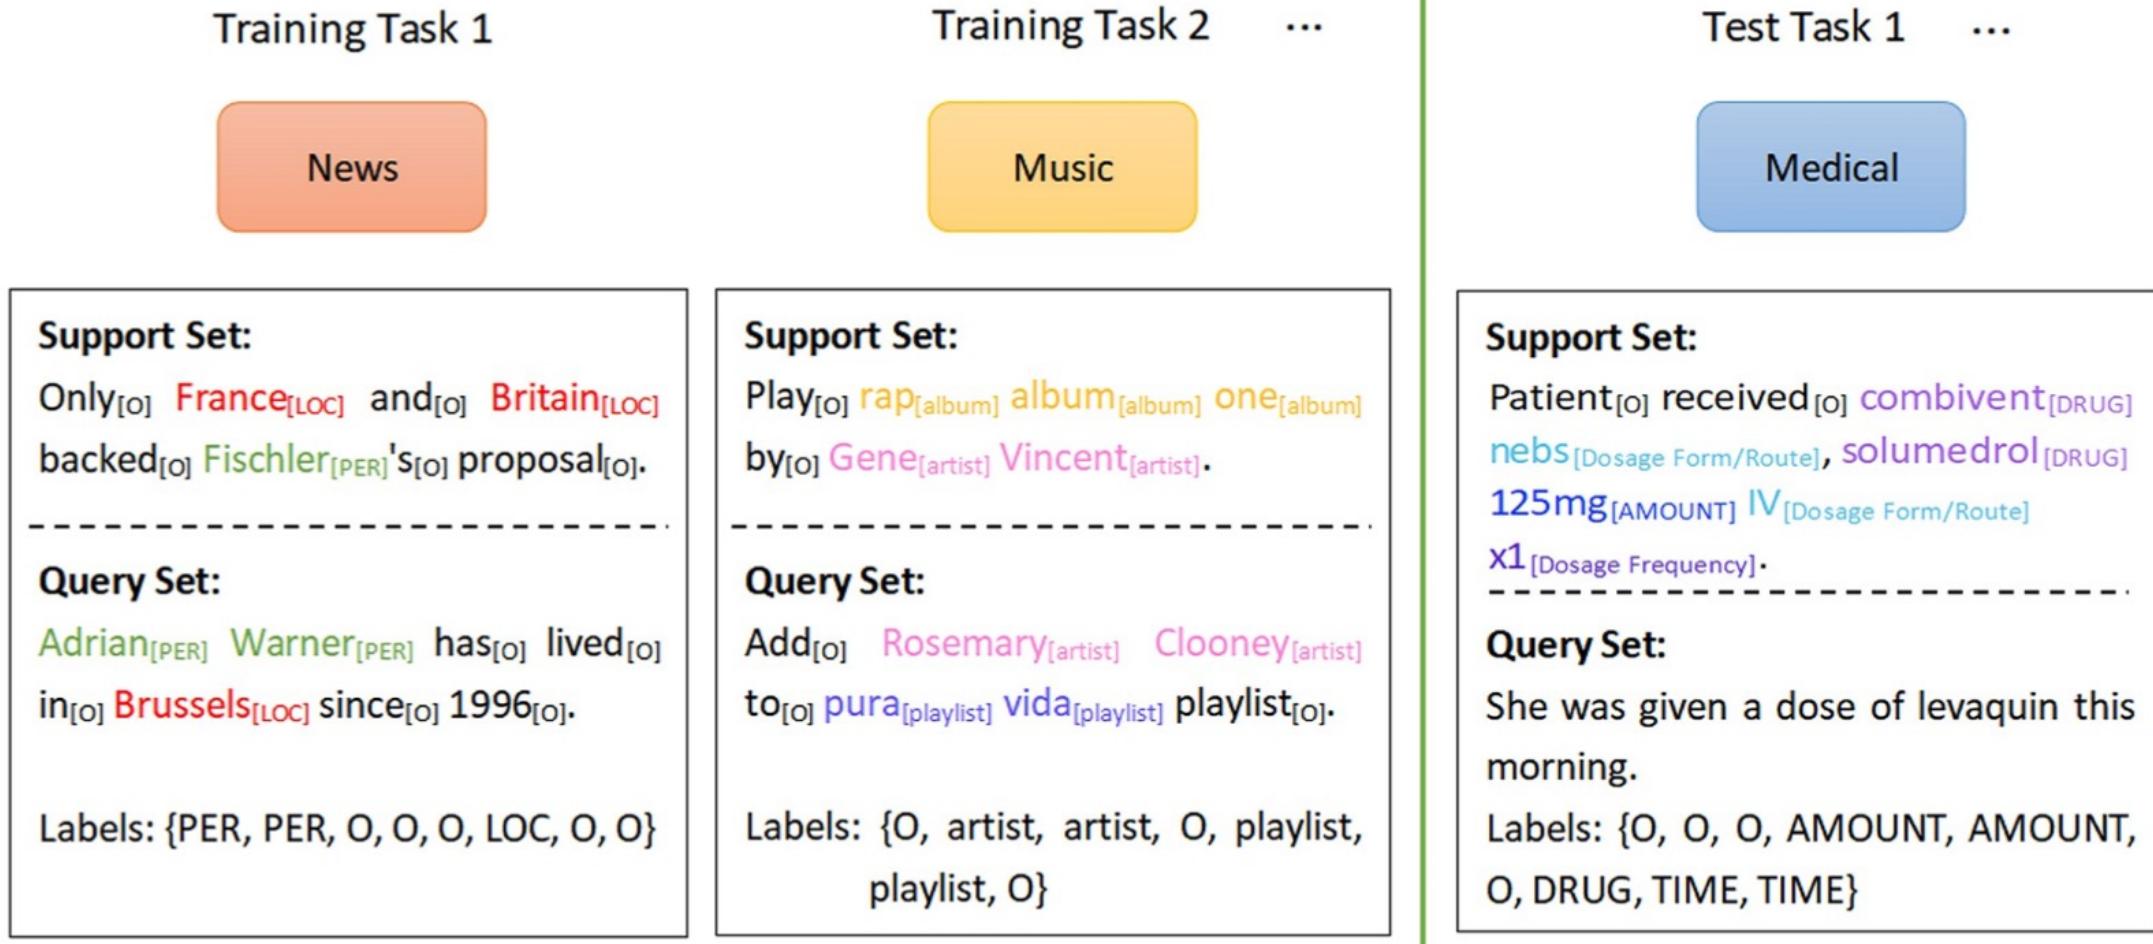
meta-testing tasks T_t 's



Few-Shot Learning (FSL)

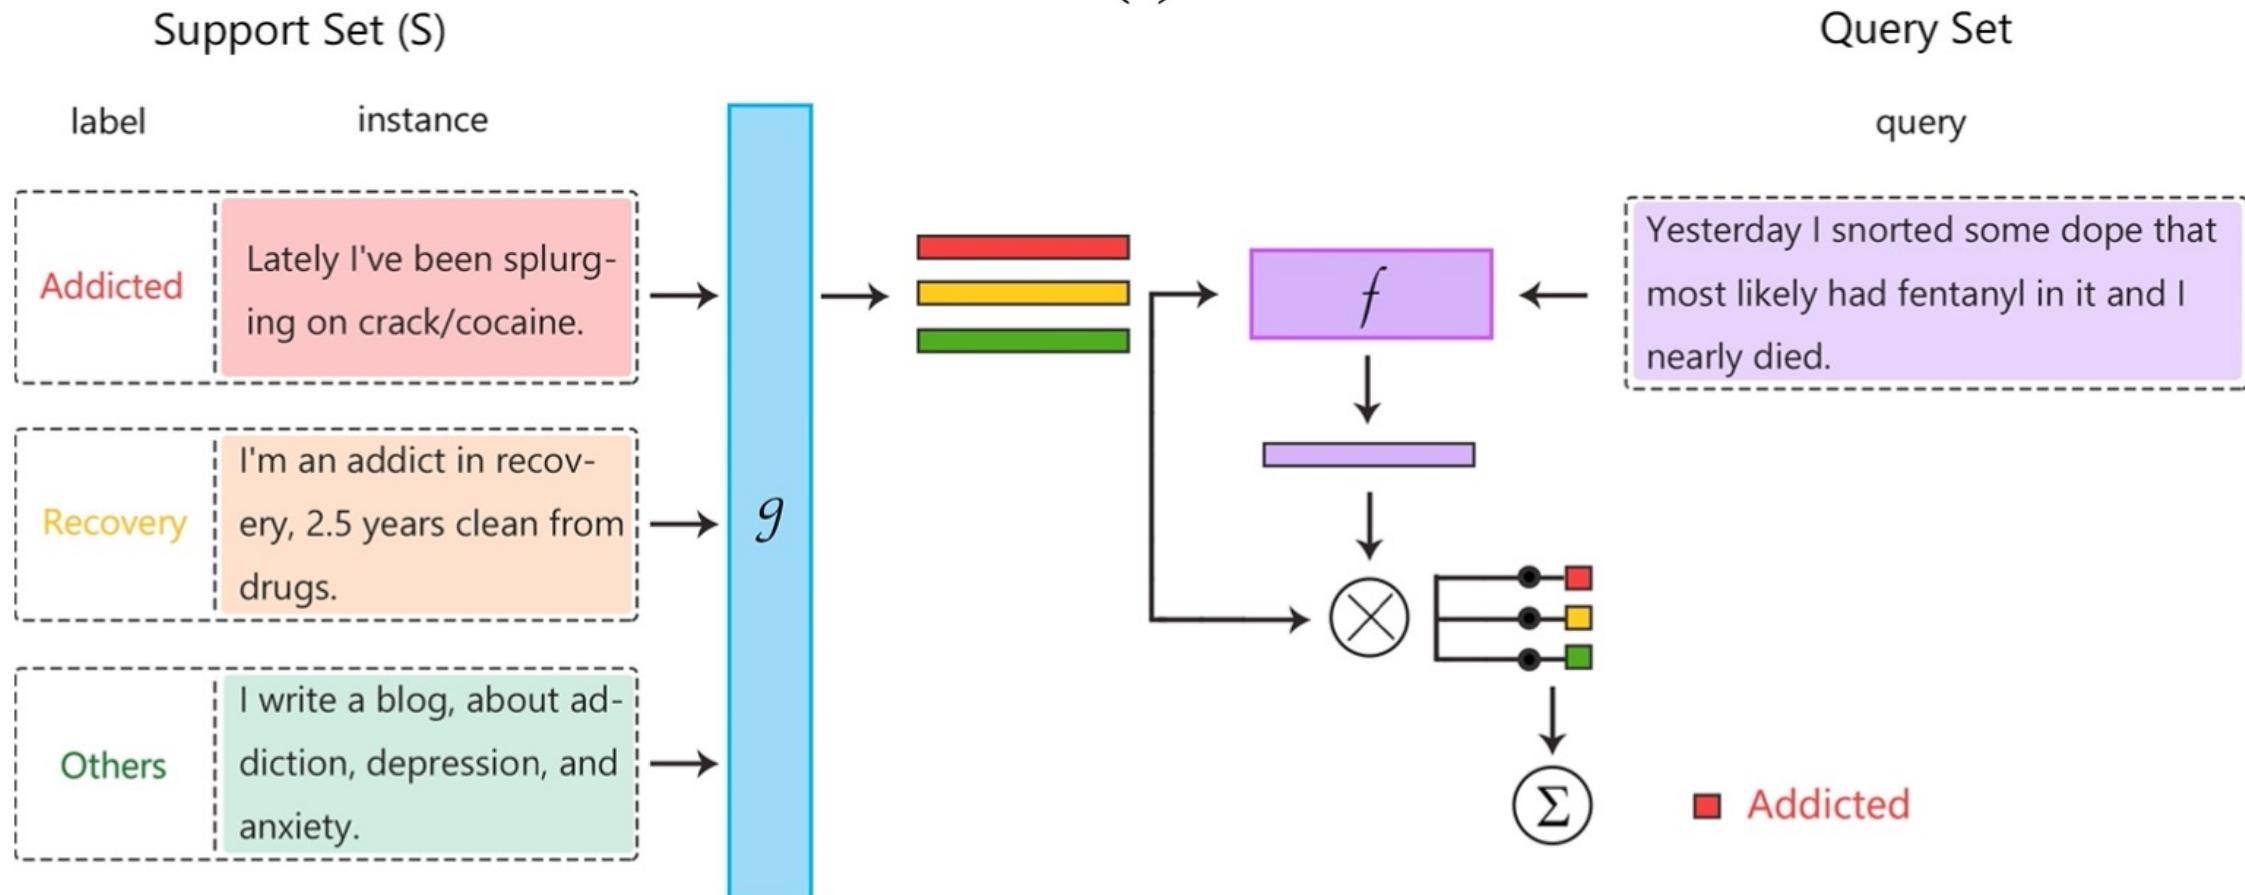
Meta-learning

Each task mimics the few-shot scenario, and can be completely non-overlapping.
Support sets are used to train; query sets are used to evaluate the model



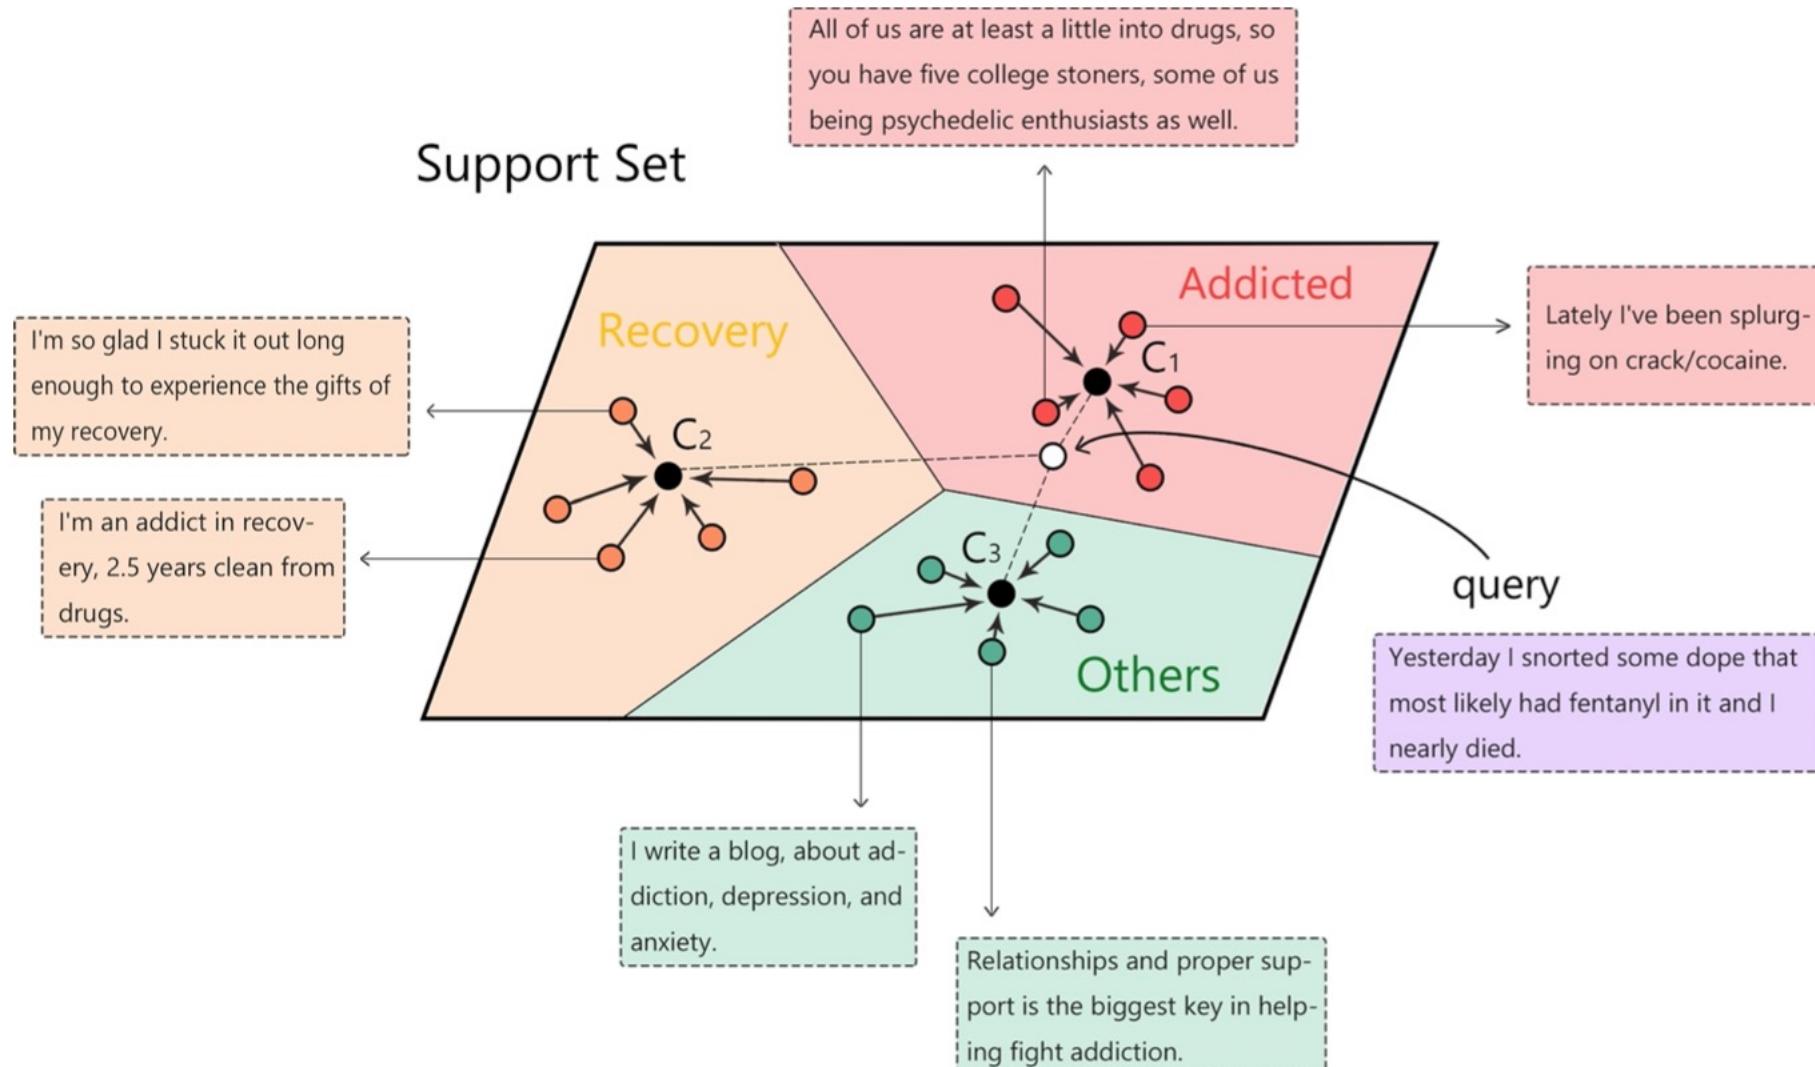
Few-Shot Learning (FSL)

Matching networks



Few-Shot Learning (FSL)

Prototypical network



Few-Shot Learning (FSL) for medical text

Study	Year	Data source	Research aim	Size of training set	Number of entities / classes	Entity type of training domain	Entity type of test domain
Alicia Lara-Clares and Ana Garcia-Serrano ⁴⁴	2019	MEDDOCAN shared task dataset ⁴⁵	NER	500 clinical cases, with no reconstruction	29	Clinical	Clinical
Ferré et al. ⁴⁶	2019	BB-norm dataset from the Bacteria Biotope 2019 Task ⁴⁷	Entity Normalization	Original dataset with no reconstruction and zero-shot	Not mentioned *	Biological	Biological
Hou et al. ⁴⁸	2020	Snips dataset ⁴⁹	Slot Tagging (NER)	1-shot and 5-shot	7	Six of Weather, Music, PlayList, Book (including biomedical), Search Screen (including biomedical), Restaurant and Creative Work.	The remaining one
Sharaf et al. ⁵⁰	2020	ten different datasets collected from the Open Parallel Corpus (OPUS) ⁵¹	Neural Machine Translation (NMT)	Sizes ranging from 4k to 64k training words (200 to 3200 sentences), but reconstructed	N/A †	Bible, European Central Bank, KDE, Quran, WMT news test sets, Books, European Medicines Agency (EMEA), Global Voices, Medical (ufal-Med), TED talks	Bible, European Central Bank, KDE, Quran, WMT news test sets, Books, European Medicines Agency (EMEA), Global Voices, Medical (ufal-Med), TED talks
Lu et al. ⁵²	2020	MIMIC II ²² and MIMIC III ²³ , and EU legislation dataset ⁵³	Multi-label Text Classification	5-shot for MIMIC II and III, 50-shot for EU legislation	MIMIC II: 9 MIMIC III: 15 EU legislation: 5	Medical	Medical

Few-Shot Learning (FSL) for medical text

Study	Year	Data source	Research aim	Size of training set	Number of entities / classes	Entity type of training domain	Entity type of test domain
Lu et al. ⁸⁰	2021	Constructed and shared a novel dataset ^{††} based on Weibo for the research of few-shot rumor detection, and use PHEME dataset ⁸¹	Rumor Detection (NER)	For the Weibo dataset: 2-way 3-event 5-shot 9-query; for PHEME dataset: 2-way 2-event 5-shot 9-query	Weibo: 14 PHEME: 5	Source posts and comments from Sina Weibo related to COVID-19	Source posts and comments from Sina Weibo related to COVID-19
Ma et al. ⁸²	2021	CCLE, CERES-corrected CRISPR gene disruption scores, GDSC1000 dataset, PDTc dataset and PDX dataset ^{‡‡}	Drug-response Predictions	1-shot, 2-shot, 5-shot and 10-shot	N/A [†]	Biomedical	Biomedical
Kormilitzin et al. ⁸³	2021	MIMIC-III ²³ and UK-CRIS datasets ^{30,31}	NER	25%, 50%, 75% and 100% of the training set, with no reconstruction	7	Electronic health record	Electronic health record
Guo et al. ⁸⁴	2021	Abstracts of biomedical literatures (from relation extraction task of BioNLP Shared Task 2011 and 2019 ⁴⁷) and structured biological datasets	NER	100%, 75%, 50%, 25%, 0% of training set, with no reconstruction	Not mentioned *	Biomedical entities	Biomedical entities
Lee et al. ⁸⁵	2021	COVID19-Scientific ⁸⁶ , COVID19-Social ⁸⁷ (fact-checked by journalists from a website called Politi-fact.com), FEVER ⁸⁸ (Fact Extraction and Verification, generated by altering sentences extracted from Wikipedia to promote research on fact-checking systems)	Fact-Checking (close to Text Classification)	2-shot, 10-shot and 50-shot	Not mentioned *	Facts about COVID-19	Facts about COVID-19

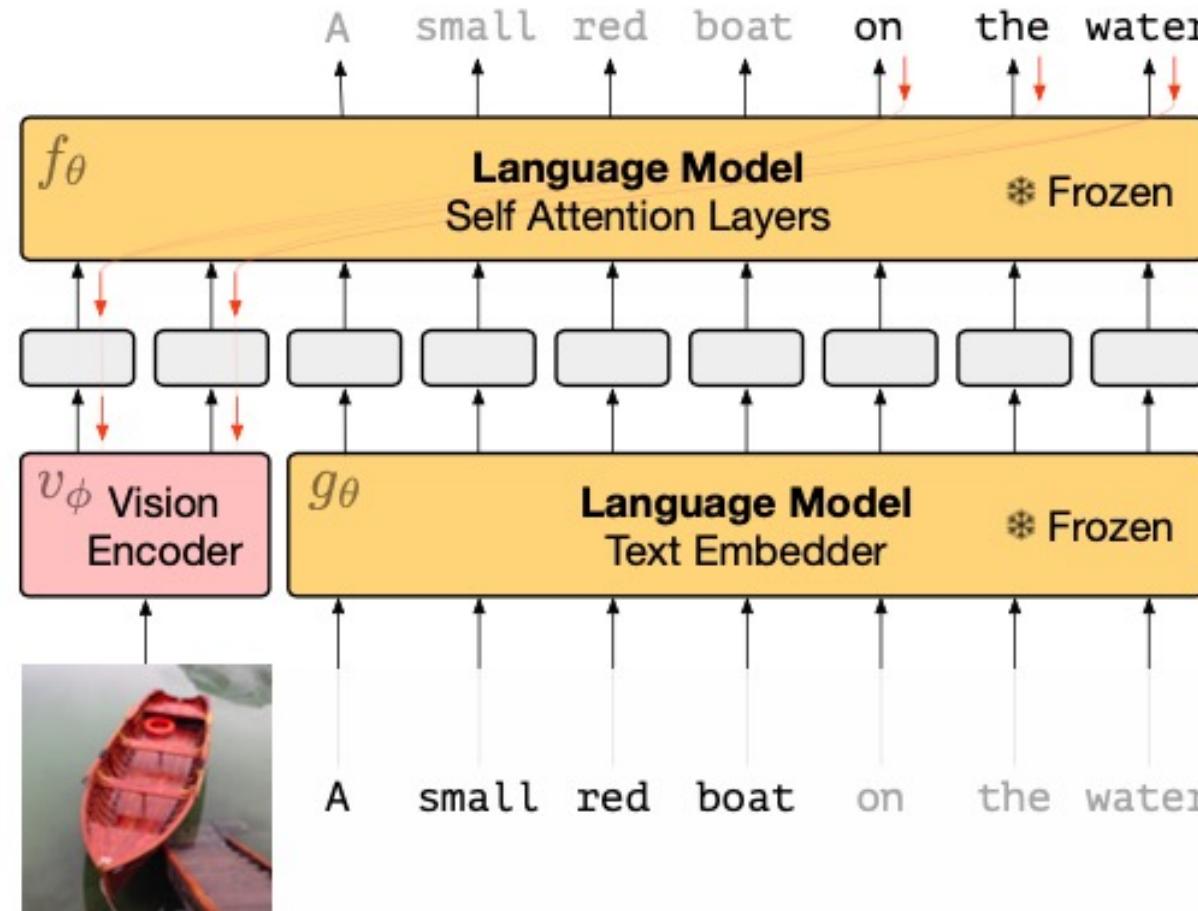
Multimodal Few-Shot Learning with Frozen Language Models

	This person is like 😊.		This person is like 😒.		This person is like	Model Completion 选拨. <EOS>
	This was invented by Zacharias Janssen.		This was invented by Thomas Edison.		This was invented by	Model Completion the Wright brothers. <EOS>
	With one of these I can drive around a track, overtaking other cars and taking corners at speed		With one of these I can take off from a city and fly across the sky to somewhere on the other side of the world		With one of these I can	Model Completion break into a secure building, unlock the door and walk right in <EOS>

Curated samples with about five seeds required to get past well-known language model failure modes of either repeating text for the prompt or emitting text that does not pertain to the image.

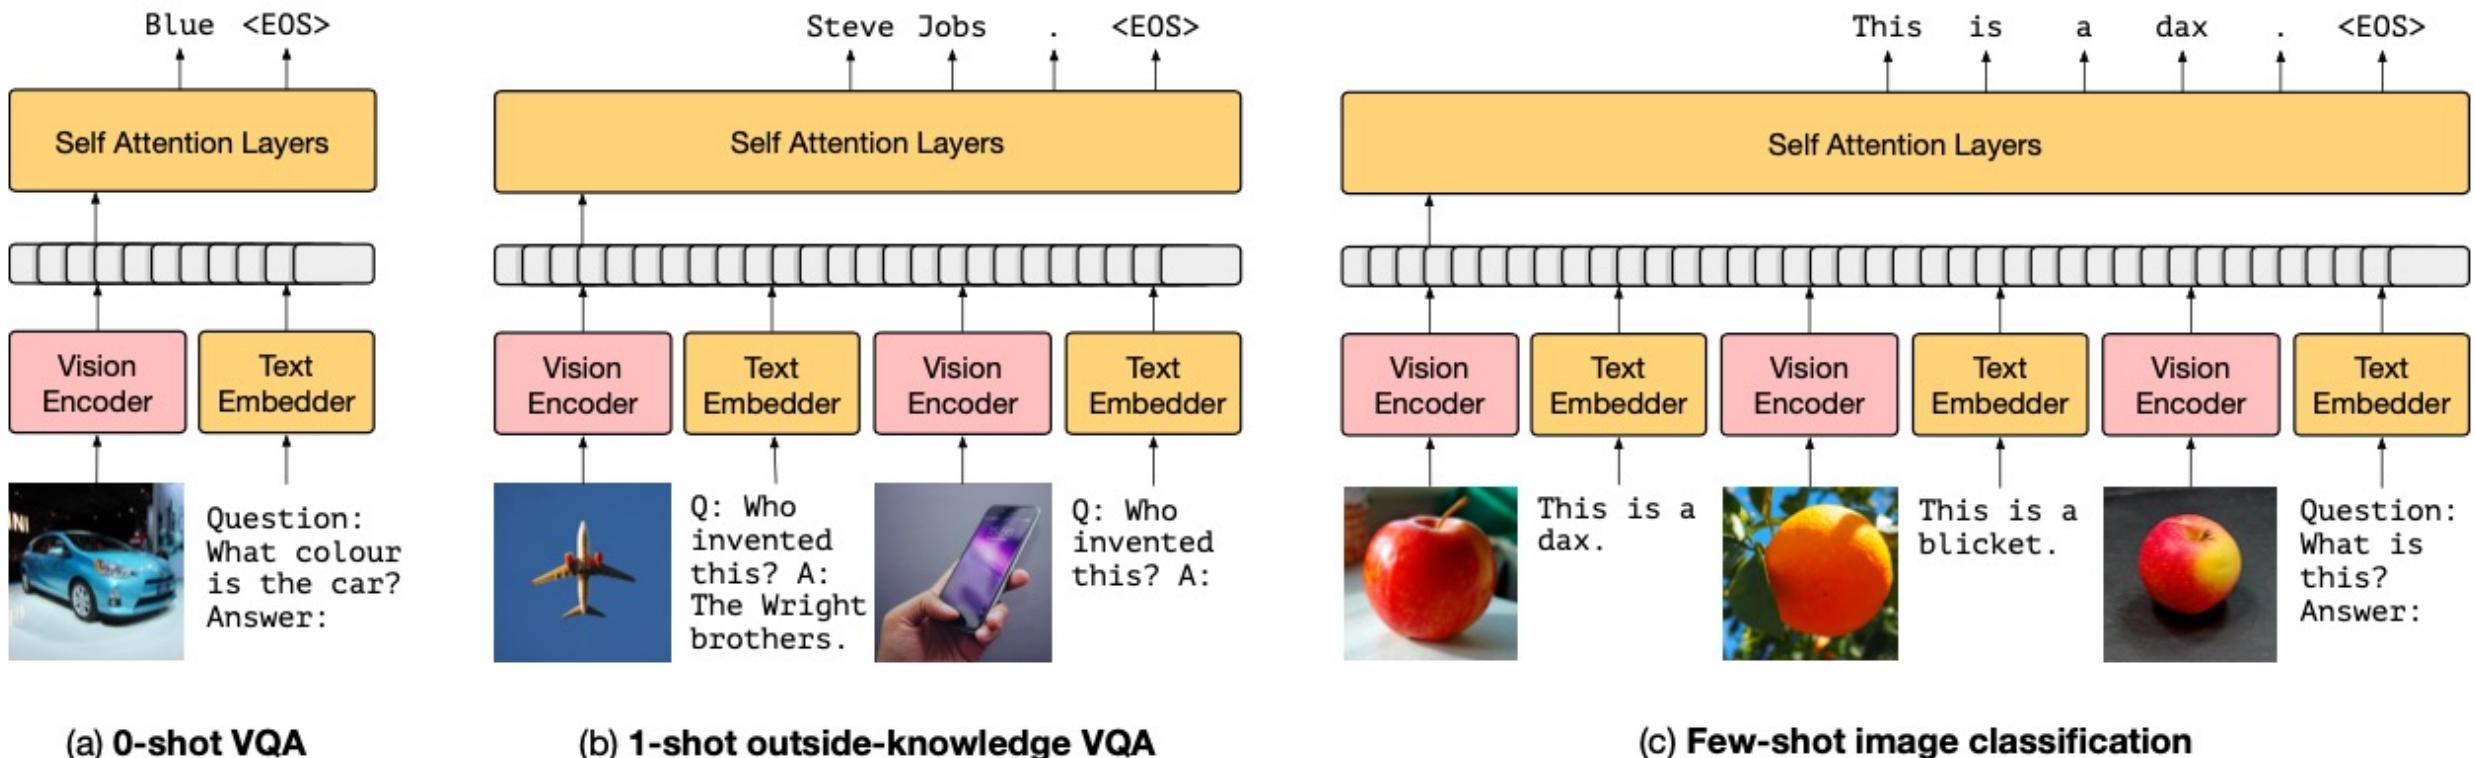
These samples demonstrate the ability to generate open-ended outputs that adapt to both images and text, and to make use of facts that it has learned during language-only pre-training.

Multimodal Few-Shot Learning with Frozen Language Models



Gradients through a frozen language model's self attention layers are used to train the vision encoder.

Multimodal Few-Shot Learning with Frozen Language Models



(a) 0-shot VQA

(b) 1-shot outside-knowledge VQA

(c) Few-shot image classification

Inference-Time interface for *Frozen*. The figure demonstrates how we can support (a) visual question answering, (b) outside-knowledge question answering and (c) few-shot image classification via in-context learning.

Multimodal Few-Shot Learning with Frozen Language Models

(a) minilmageNet



(b) Fast VQA



Examples of (a) the Open-Ended minilmageNet evaluation (b) the Fast VQA evaluation.

Source: Maria Tsimpoukelli, Jacob L. Menick, Serkan Cabi, S. M. Eslami, Oriol Vinyals, and Felix Hill (2021). "Multimodal few-shot learning with frozen language models."

Advances in Neural Information Processing Systems 34 (2021): 200-212.

Machine Learning: Ensemble Learning

Random Forest

The screenshot shows a Jupyter Notebook interface with the following details:

- Title Bar:** python101.ipynb
- Toolbar:** File, Edit, View, Insert, Runtime, Tools, Help, All changes saved
- Header Buttons:** Comment, Share, Settings, A (font size)
- Table of Contents:** Shows a hierarchical list of topics, with "Random Forest" selected.
- Code Cell:** Displays Python code for a Random Forest classifier using scikit-learn. The code includes importing libraries, loading the iris dataset, creating a DataFrame, and adding a new column for species names.

```
# Randon Forest: https://chrisalbon.com/machine_learning/trees_and_forests/random_forest_classifier/
# Load the library with the iris dataset
from sklearn.datasets import load_iris
# Load scikit's random forest classifier library
from sklearn.ensemble import RandomForestClassifier
# Load pandas
import pandas as pd
# Load numpy
import numpy as np
# Set random seed
np.random.seed(0)
# Create an object called iris with the iris data
iris = load_iris()
# Create a dataframe with the four feature variables
if = pd.DataFrame(iris.data, columns=iris.feature_names)
# View the top 5 rows
if.head()
# Add a new column with the species names, this is what we are going to try to predict
if['species'] = pd.Categorical.from_codes(iris.target, iris.target_names)
```

Machine Learning: Supervised Learning Classification and Prediction

The screenshot shows a Jupyter Notebook interface with the following details:

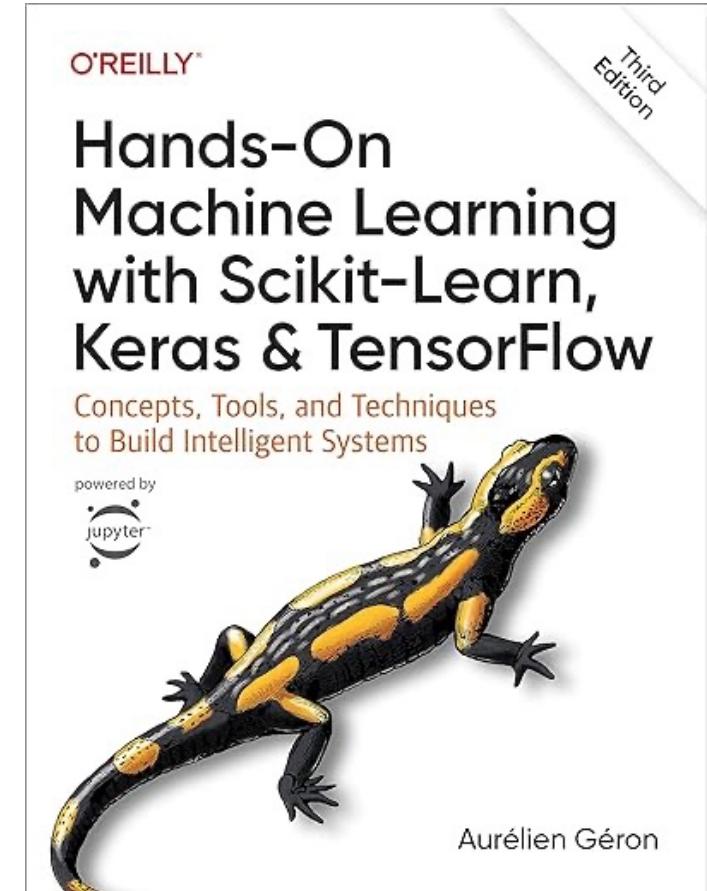
- Title:** python101.ipynb
- File Menu:** File, Edit, View, Insert, Runtime, Tools, Help
- Toolbar:** Comment, Share, Settings, A (theme)
- Table of Contents:** Machine Learning with scikit-learn, Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python, Investment Portfolio Optimization, Text Analytics and Natural Language Processing (NLP), Python for Natural Language Processing, spaCy Chinese Model, Open Chinese Convert (OpenCC, 開放中文轉換), Jieba 結巴中文分詞, Natural Language Toolkit (NLTK), Stanza: A Python NLP Library for Many Human Languages, Text Processing and Understanding, NLTK (Natural Language Processing with Python – Analyzing Text with the
- Code Cell:** Displays Python code for importing libraries and loading the Iris dataset.

```
1 # Import libraries
2 import numpy as np
3 import pandas as pd
4 %matplotlib inline
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from pandas.plotting import scatter_matrix
8
9 # Import sklearn
10 from sklearn import model_selection
11 from sklearn.metrics import classification_report
12 from sklearn.metrics import confusion_matrix
13 from sklearn.metrics import accuracy_score
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.tree import DecisionTreeClassifier
16 from sklearn.neighbors import KNeighborsClassifier
17 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
18 from sklearn.naive_bayes import GaussianNB
19 from sklearn.svm import SVC
20 from sklearn.neural_network import MLPClassifier
21 print("Imported")
22
23 # Load dataset
24 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
25 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
```

Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow

Notebooks

1. [The Machine Learning landscape](#)
2. [End-to-end Machine Learning project](#)
3. [Classification](#)
4. [Training Models](#)
5. [Support Vector Machines](#)
6. [Decision Trees](#)
7. [Ensemble Learning and Random Forests](#)
8. [Dimensionality Reduction](#)
9. [Unsupervised Learning Techniques](#)
10. [Artificial Neural Nets with Keras](#)
11. [Training Deep Neural Networks](#)
12. [Custom Models and Training with TensorFlow](#)
13. [Loading and Preprocessing Data](#)
14. [Deep Computer Vision Using Convolutional Neural Networks](#)
15. [Processing Sequences Using RNNs and CNNs](#)
16. [Natural Language Processing with RNNs and Attention](#)
17. [Autoencoders, GANs, and Diffusion Models](#)
18. [Reinforcement Learning](#)
19. [Training and Deploying TensorFlow Models at Scale](#)



Python in Google Colab (Python101)

<https://colab.research.google.com/drive/1FEG6DnGvwfUbeo4zJ1zTunjMqf2RkCrT>

The screenshot shows a Google Colab interface with the following details:

- Title:** python101.ipynb
- File Menu:** File, Edit, View, Insert, Runtime, Tools, Help
- Toolbar:** Comment, Share, Settings, A (font size)
- Table of Contents:** Machine Learning with scikit-learn, Classification and Prediction, K-Means Clustering, Deep Learning for Financial Time Series Forecasting, Portfolio Optimization and Algorithmic Trading, Investment Portfolio Optimisation with Python, Efficient Frontier Portfolio Optimisation in Python, Investment Portfolio Optimization, Text Analytics and Natural Language Processing (NLP), Python for Natural Language Processing, spaCy Chinese Model, Open Chinese Convert (OpenCC, 開放中文轉換), Jieba 結巴中文分詞, Natural Language Toolkit (NLTK), Stanza: A Python NLP Library for Many Human Languages, Text Processing and Understanding, NLTK (Natural Language Processing with Python – Analyzing Text with the
- Code Cell:** Contains Python code for importing libraries and loading a dataset.

```
1 # Import libraries
2 import numpy as np
3 import pandas as pd
4 %matplotlib inline
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from pandas.plotting import scatter_matrix
8
9 # Import sklearn
10 from sklearn import model_selection
11 from sklearn.metrics import classification_report
12 from sklearn.metrics import confusion_matrix
13 from sklearn.metrics import accuracy_score
14 from sklearn.linear_model import LogisticRegression
15 from sklearn.tree import DecisionTreeClassifier
16 from sklearn.neighbors import KNeighborsClassifier
17 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
18 from sklearn.naive_bayes import GaussianNB
19 from sklearn.svm import SVC
20 from sklearn.neural_network import MLPClassifier
21 print("Imported")
22
23 # Load dataset
24 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
25 names = ['sepal-length', 'sepal-width', 'petal-length', 'petal-width', 'class']
```

<https://tinyurl.com/aintpuppython101>

Summary

- The Theory of Learning
 - Computational Learning Theory
 - Probably Approximately Correct (PAC) Learning
- Ensemble Learning
 - Bagging: Random Forests (RF)
 - Boosting: Gradient Boosting, XGBoost, LightGBM, CatBoost
 - Stacking
 - Online learning
- Meta Learning: Learning to Learn

References

- Stuart Russell and Peter Norvig (2020), Artificial Intelligence: A Modern Approach, 4th Edition, Pearson.
- Numa Dhamani and Maggie Engler (2024), Introduction to Generative AI, Manning
- Denis Rothman (2024), Transformers for Natural Language Processing and Computer Vision - Third Edition: Explore Generative AI and Large Language Models with Hugging Face, ChatGPT, GPT-4V, and DALL-E 3, 3rd ed. Edition, Packt Publishing
- Thomas R. Caldwell (2025), The Agentic AI Bible: The Complete and Up-to-Date Guide to Design, Build, and Scale Goal-Driven, LLM-Powered Agents that Think, Execute and Evolve, Independently published
- Ben Auffarth (2023), Generative AI with LangChain: Build large language model (LLM) apps with Python, ChatGPT and other LLMs, Packt Publishing.
- Aurélien Géron (2022), Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems, 3rd Edition, O'Reilly Media.
- Steven D'Ascoli (2022), Artificial Intelligence and Deep Learning with Python: Every Line of Code Explained For Readers New to AI and New to Python, Independently published.
- Nithin Buduma, Nikhil Buduma, Joe Papa (2022), Fundamentals of Deep Learning: Designing Next-Generation Machine Intelligence Algorithms, 2nd Edition, O'Reilly Media.
- Aske Plaat, Annie Wong, Suzan Verberne, Joost Broekens, Niki van Stein, and Thomas Back. (2024) "Reasoning with Large Language Models, a Survey." arXiv preprint arXiv:2407.11511.
- Madaan, Aman, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon et al. (2024) "Self-refine: Iterative refinement with self-feedback." Advances in Neural Information Processing Systems 36.